

How to sleep at night while having a cloud service.

Common Architecture Do's and Don'ts

DANIEL
SADA

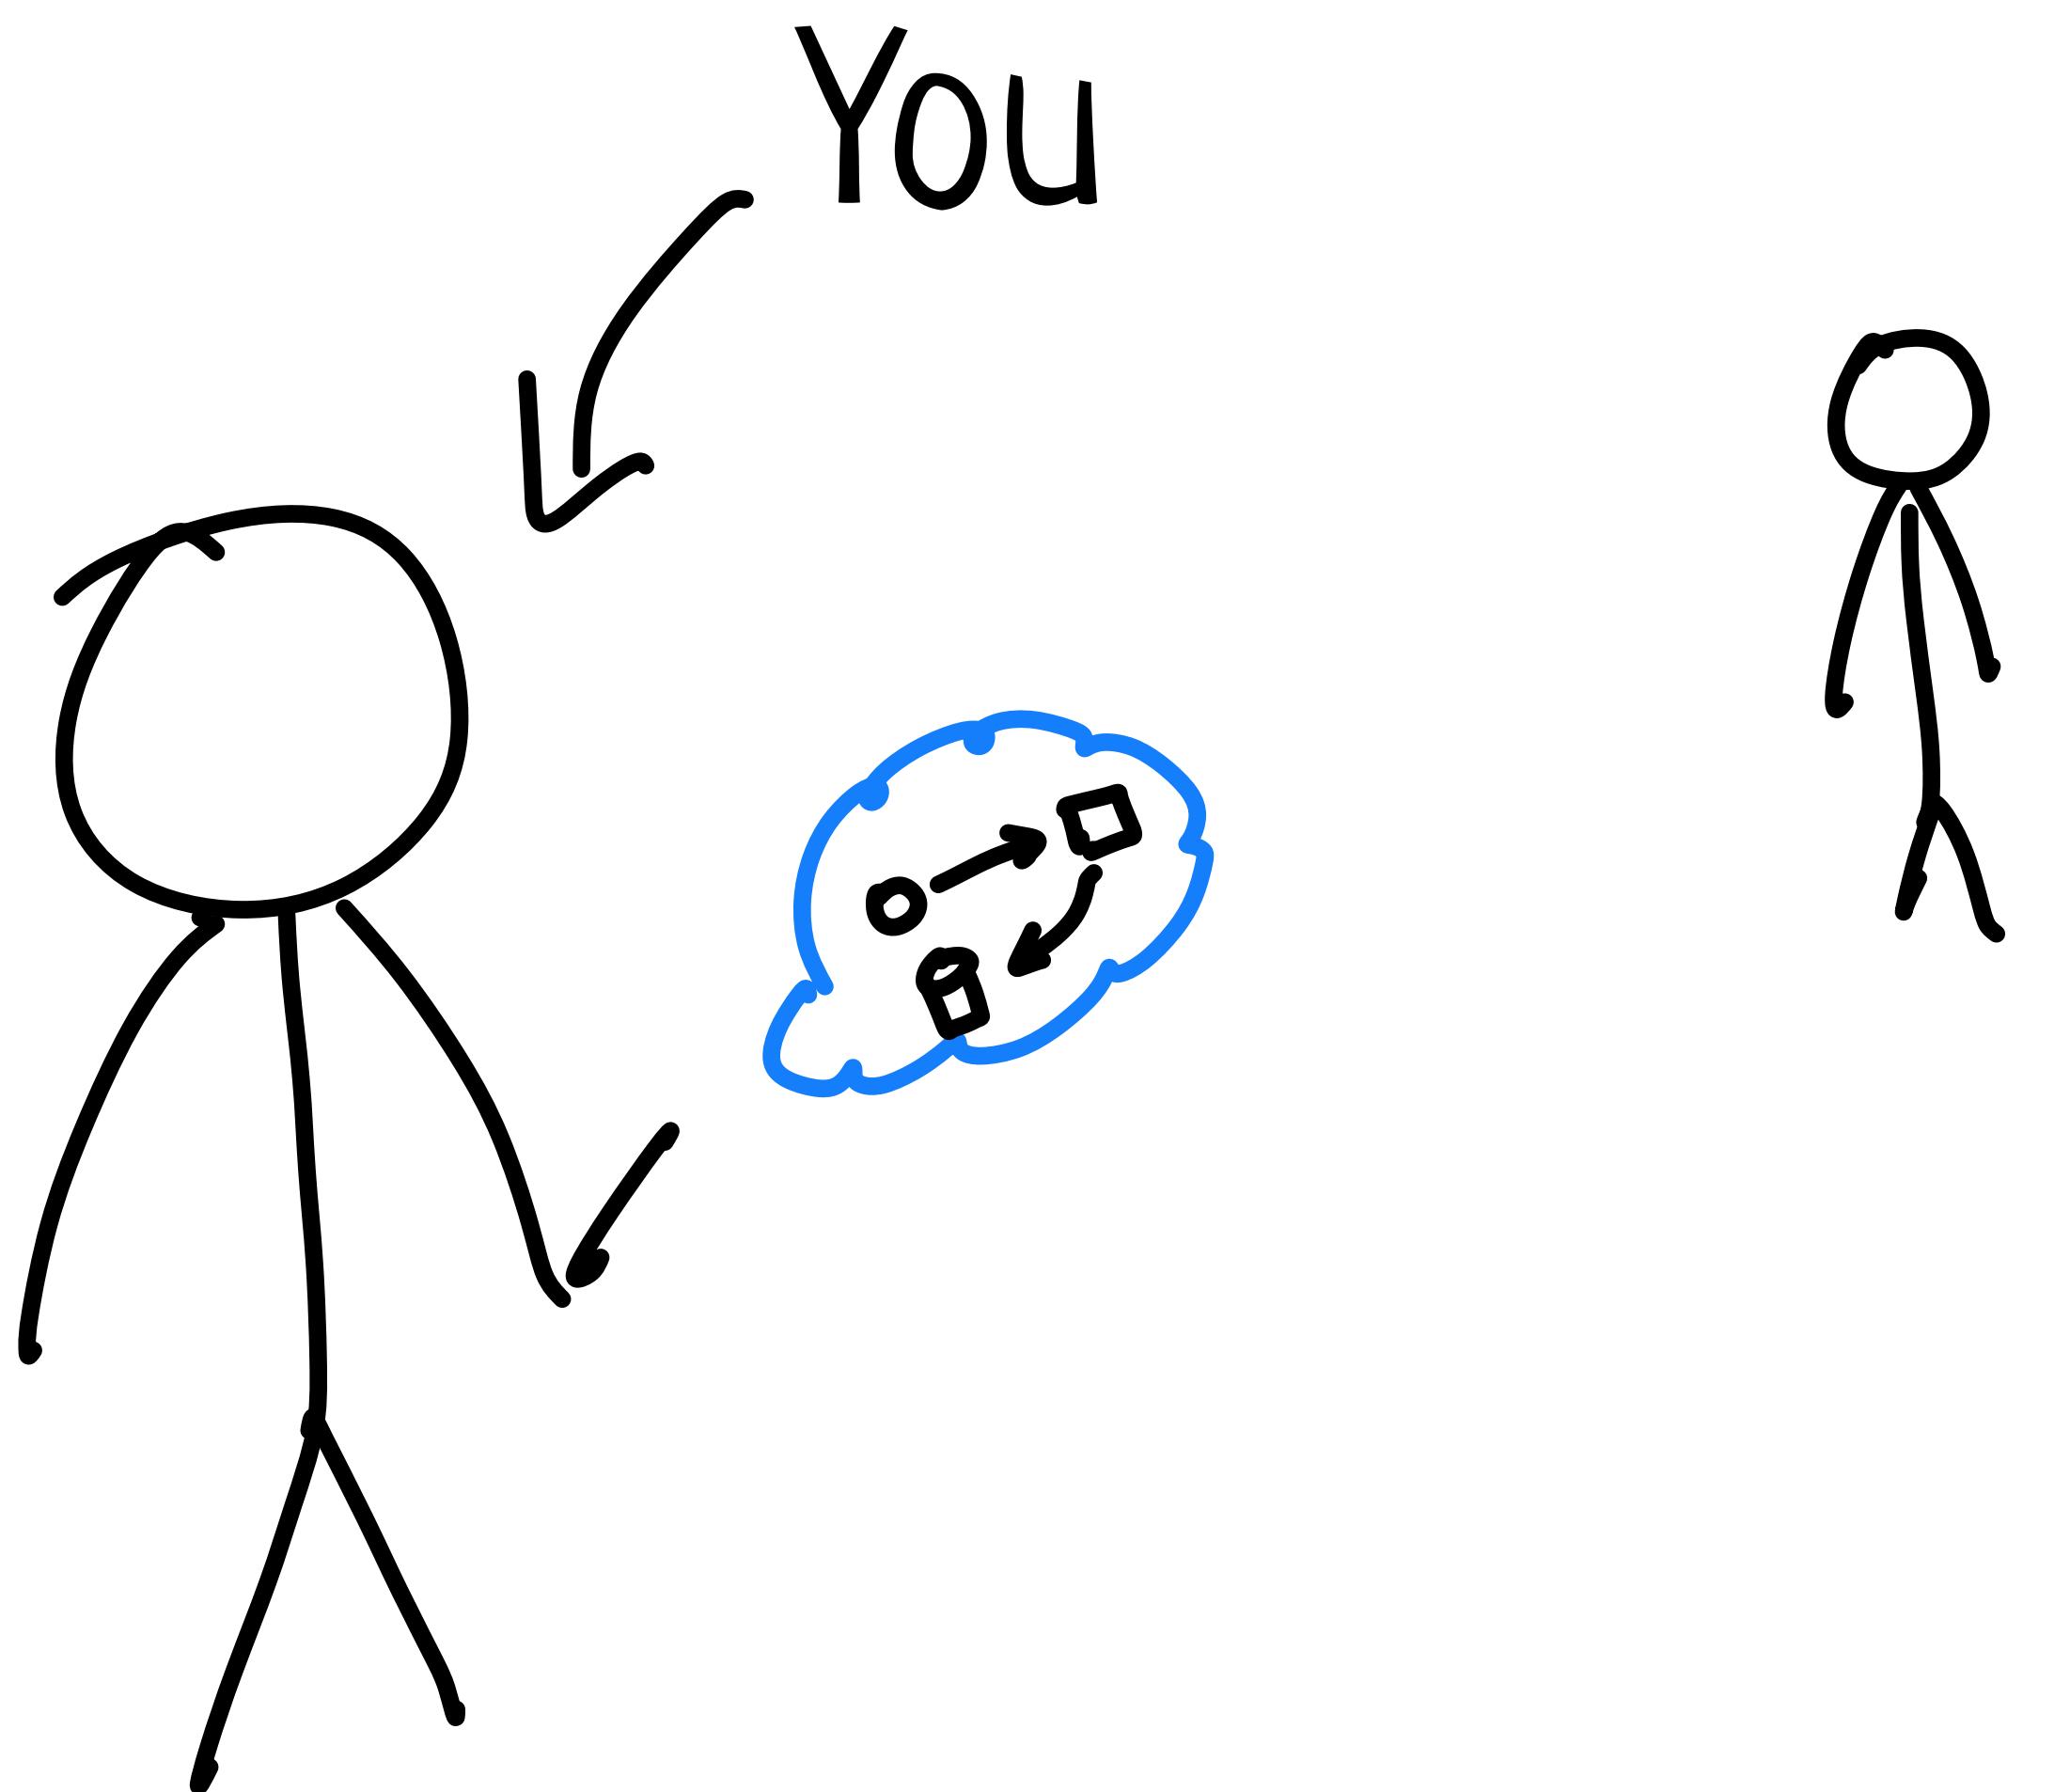
The logo consists of the name "DANIEL" in a black, hand-drawn style font above the word "SADA". The "SADA" is written in a larger, stylized font where each letter has a different color: "S" is blue, "A" is purple, "D" is orange, and "A" is red.

Developer Productivity & Culture



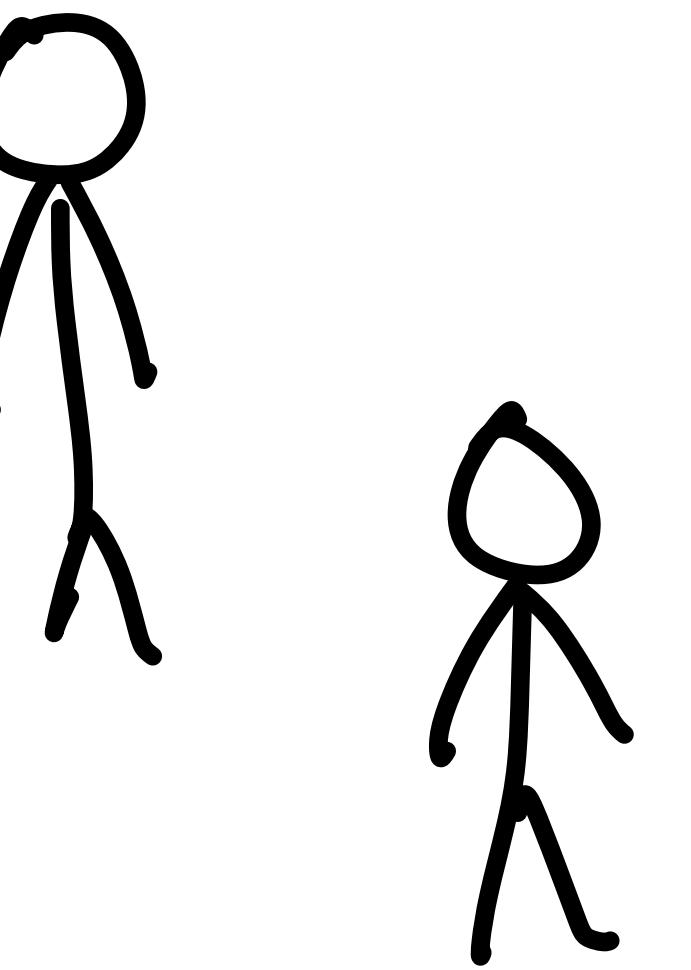
Want to deploy a service.

Someone shows up.

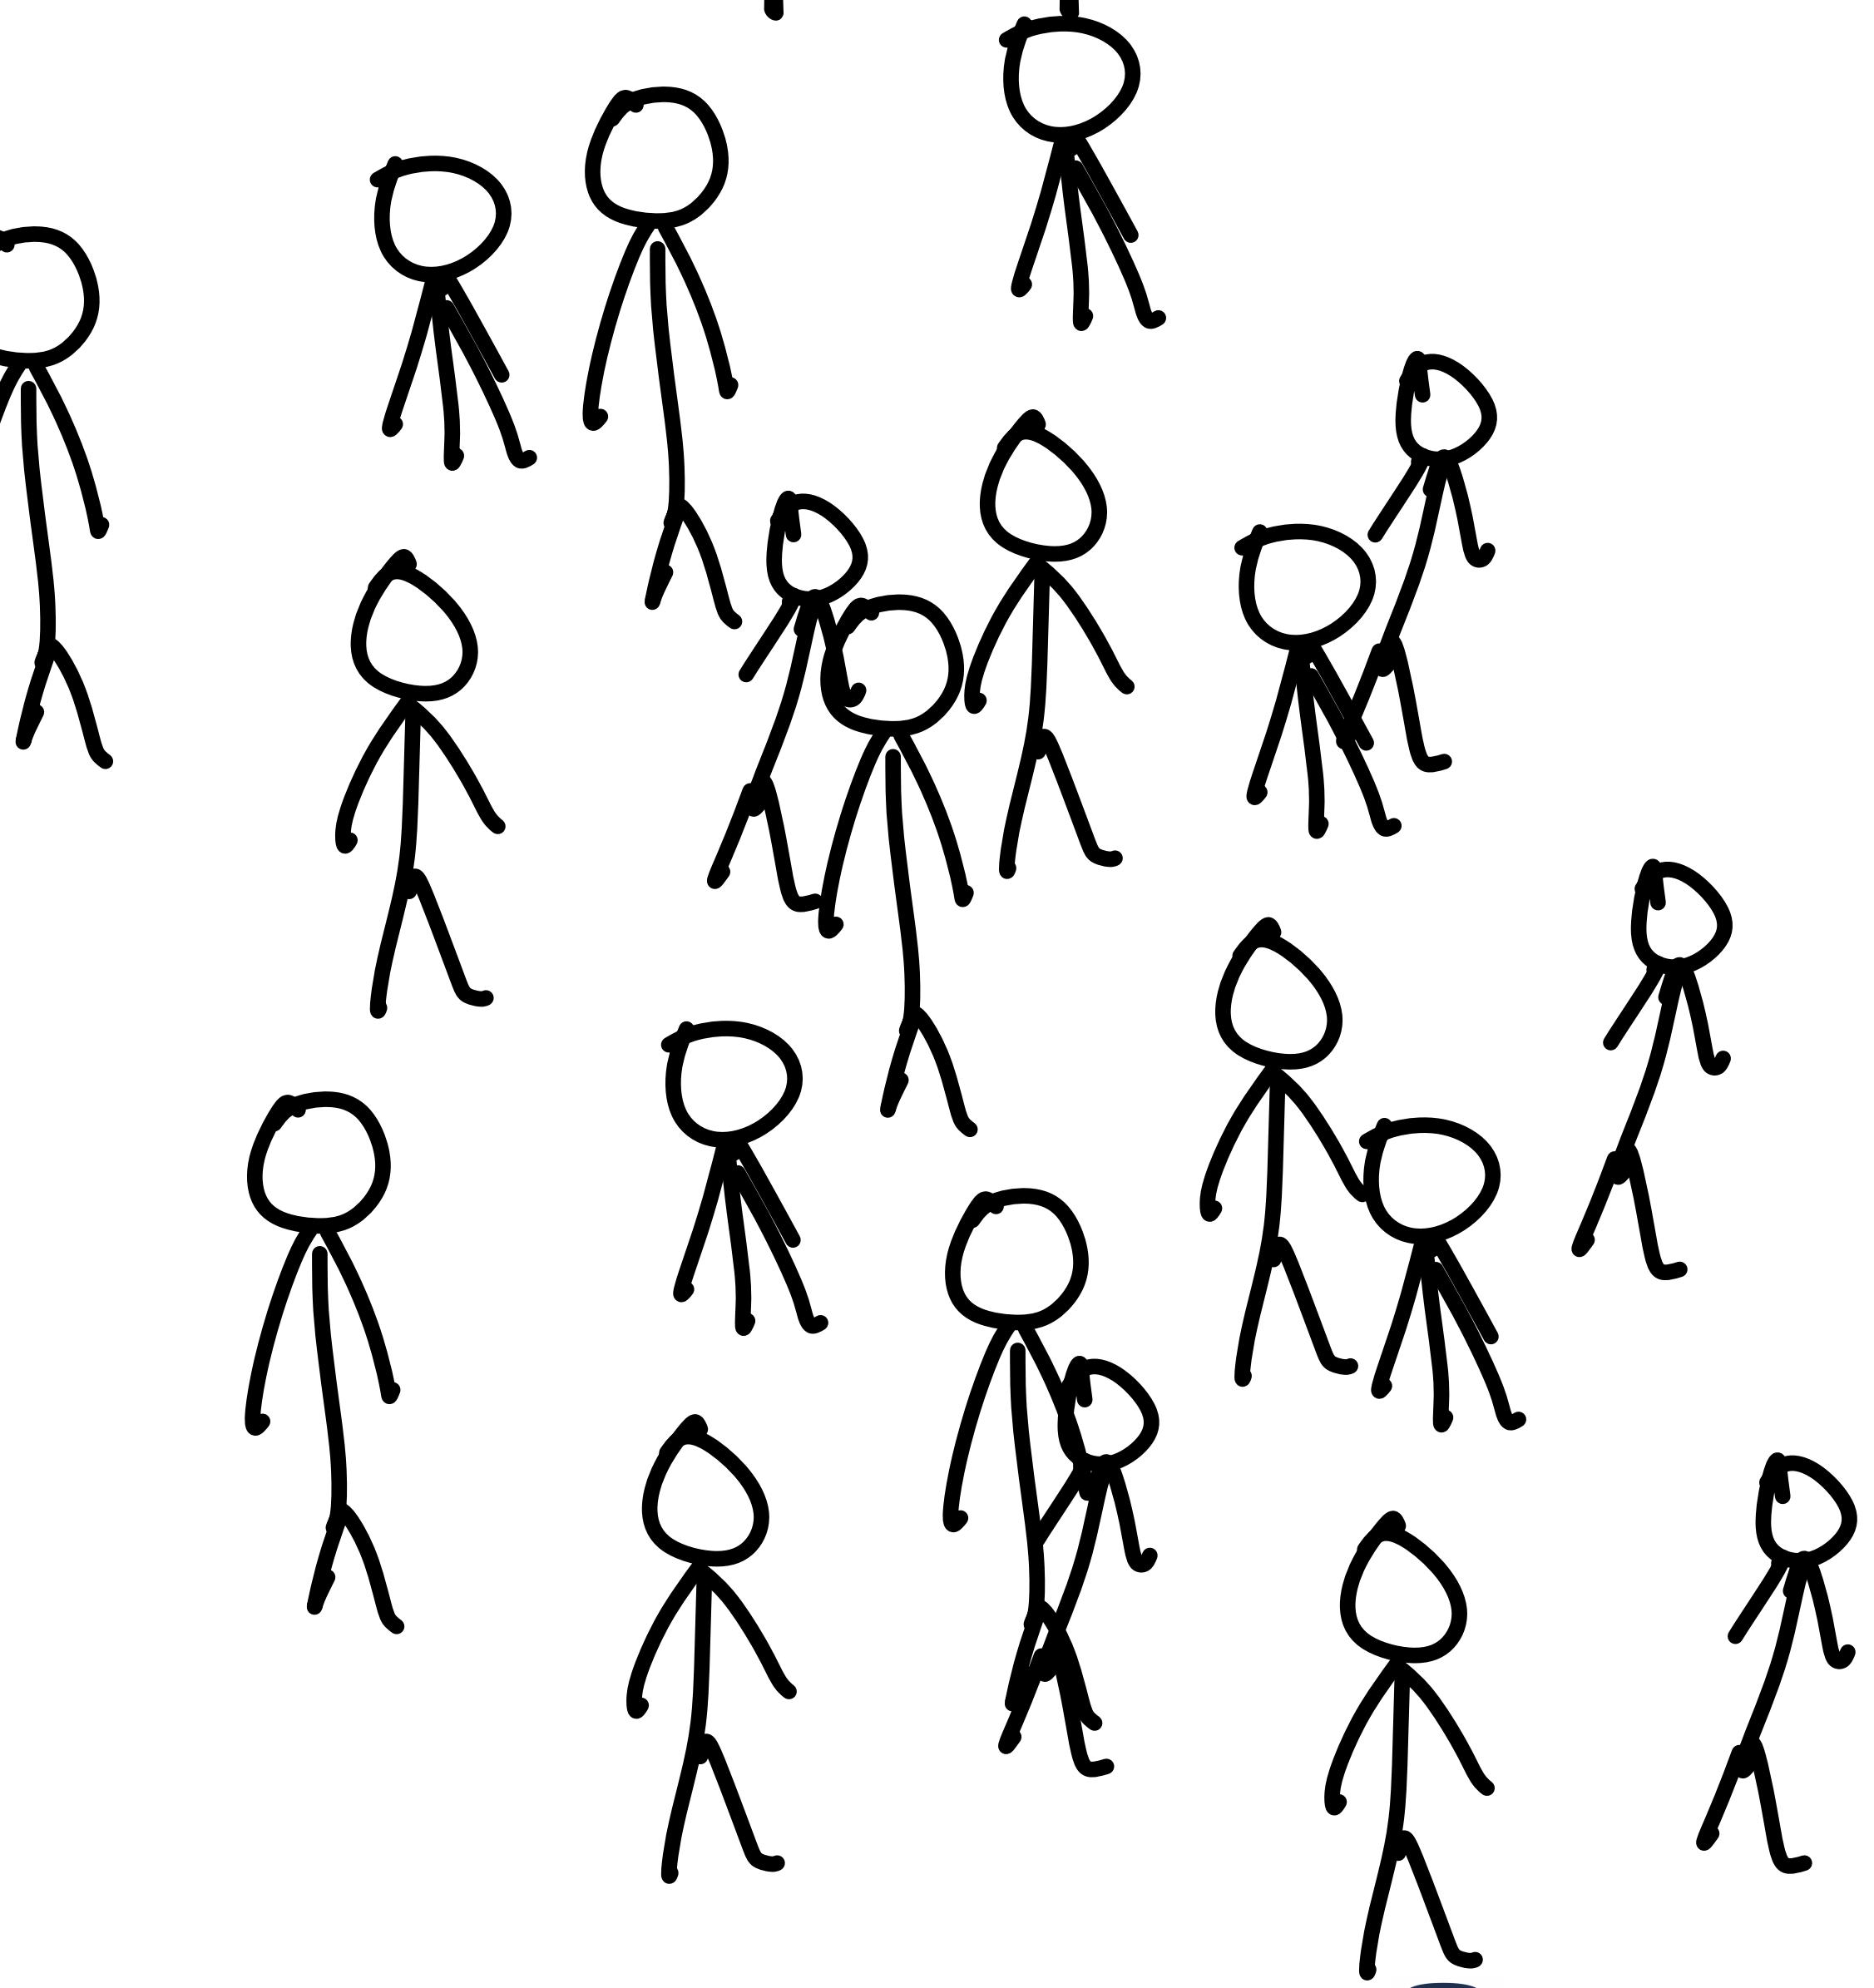




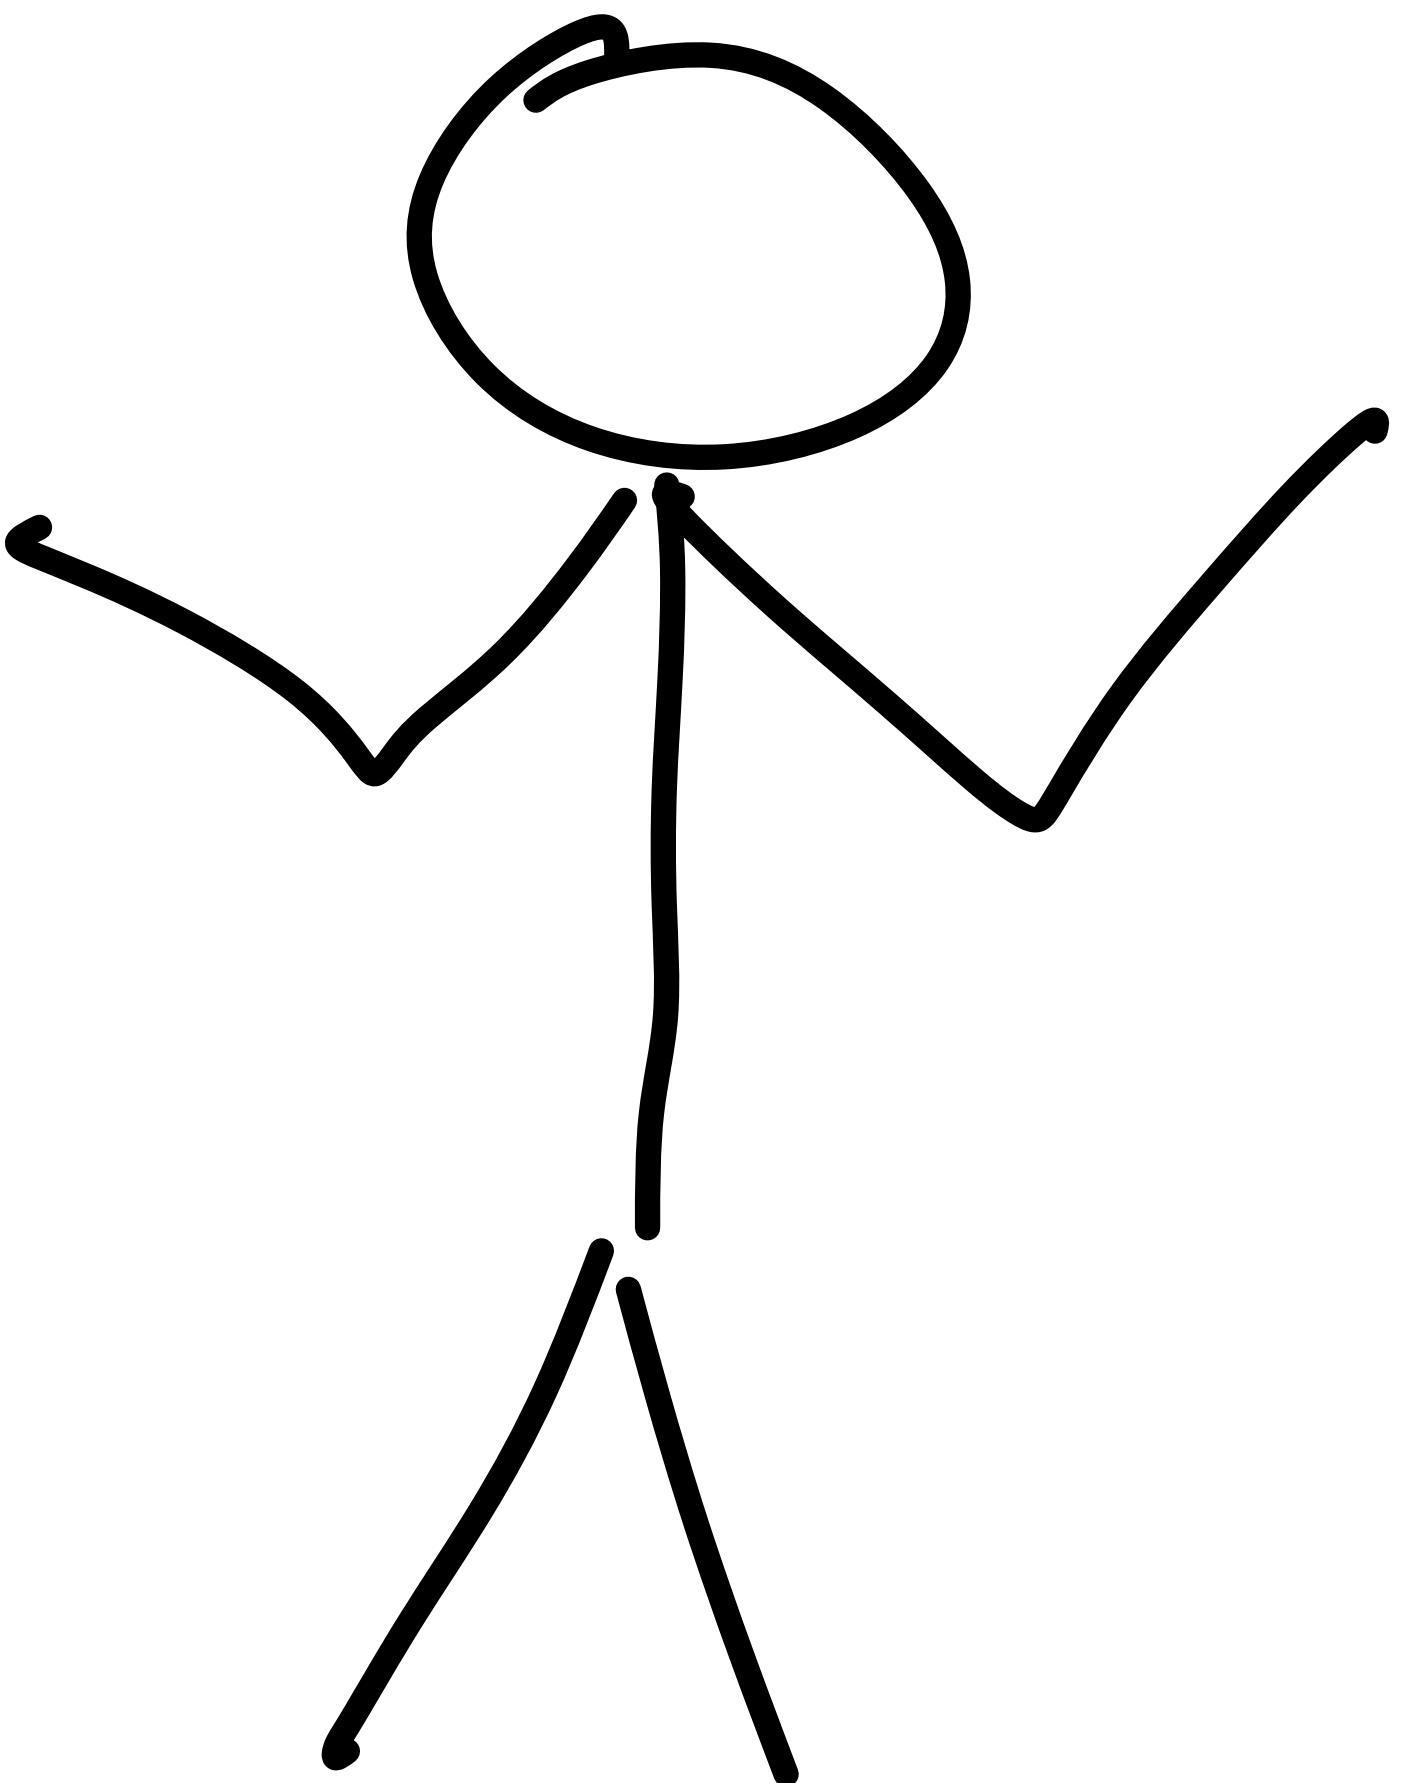
Another person shows up.



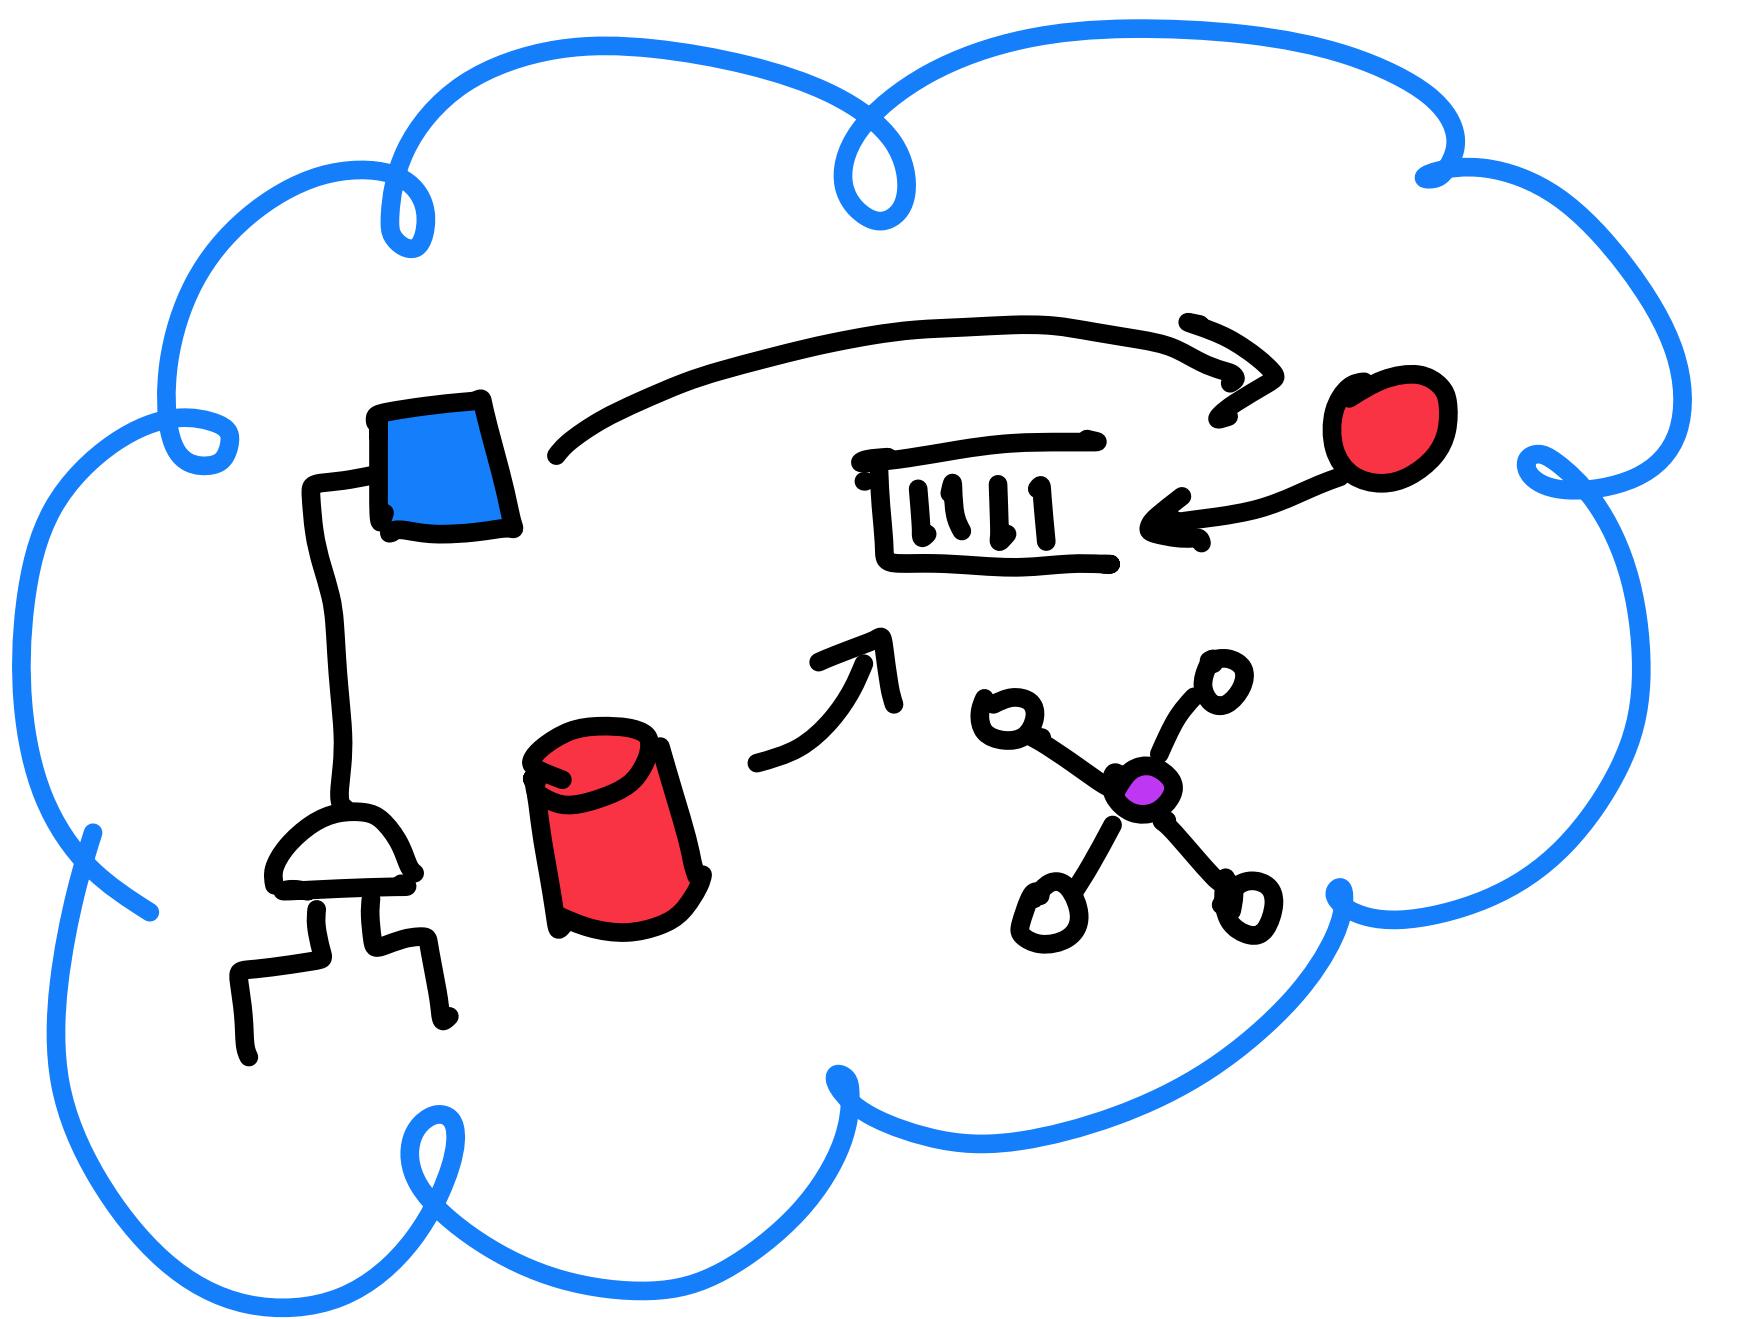
1,000,000 people show up



What do you do now?

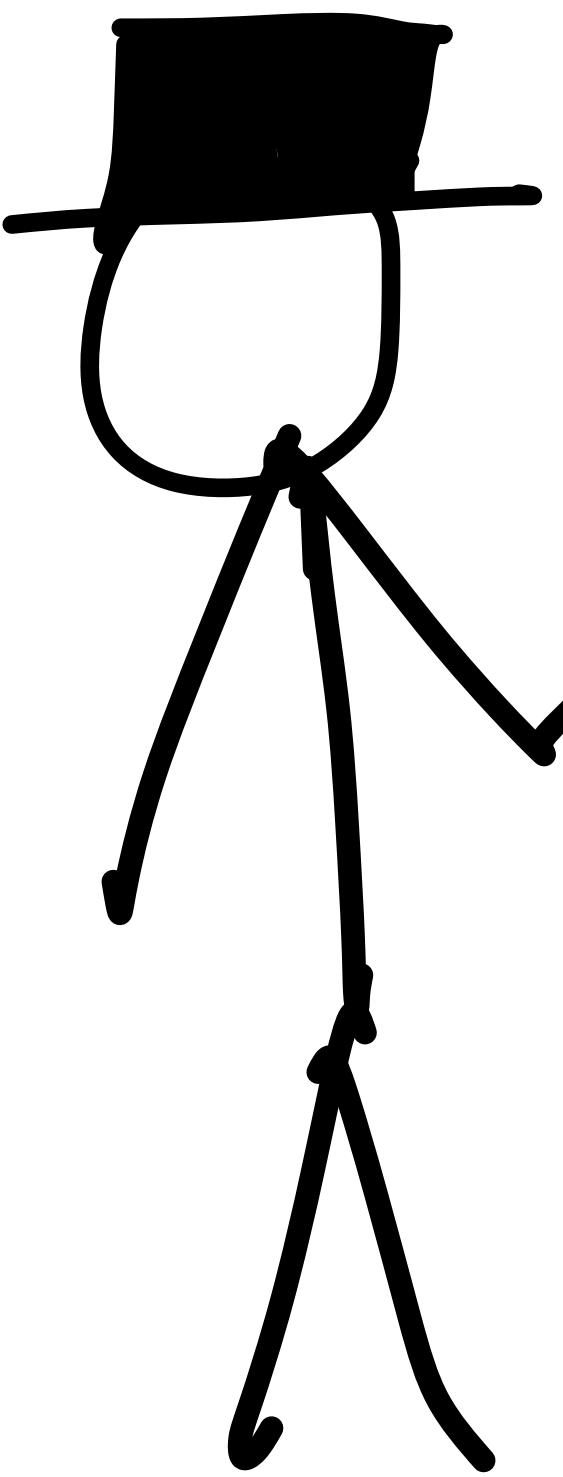
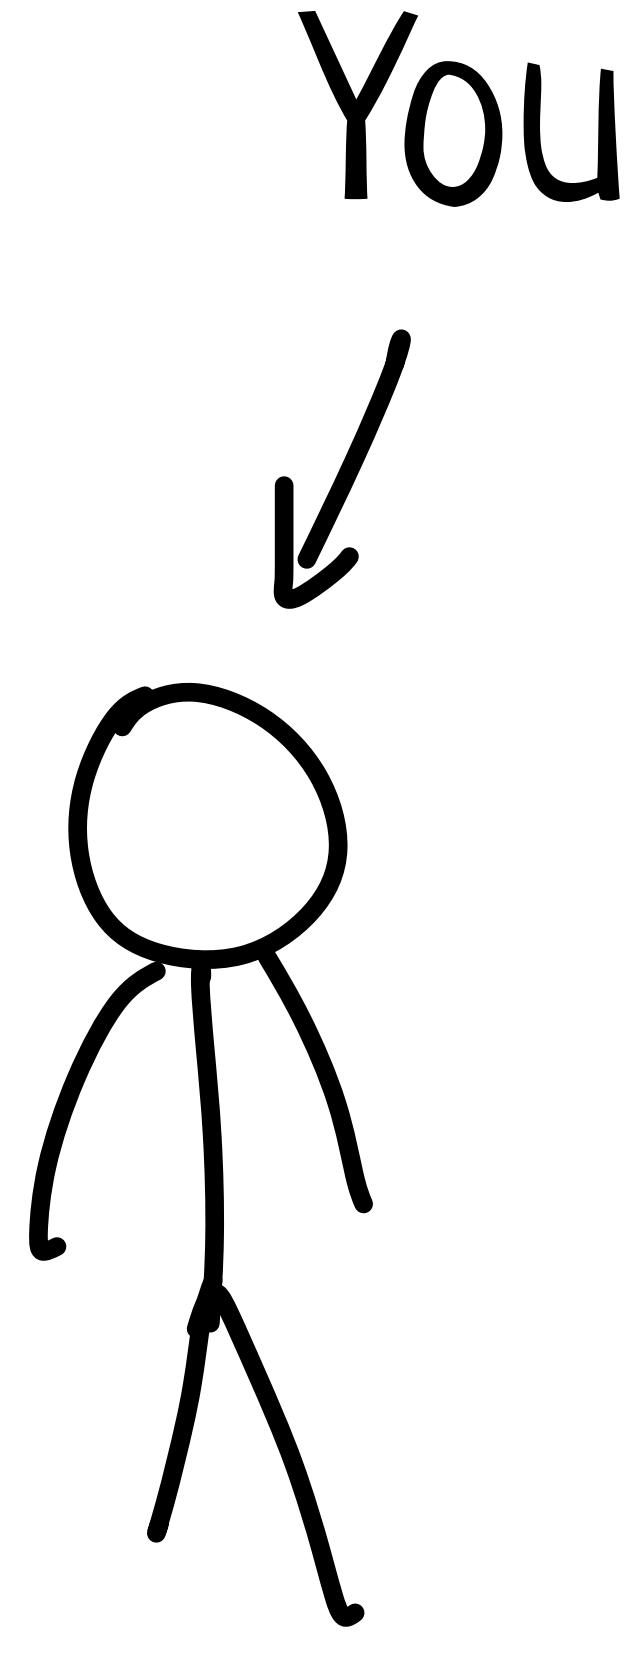


I like to sleep at night



While having a cloud application





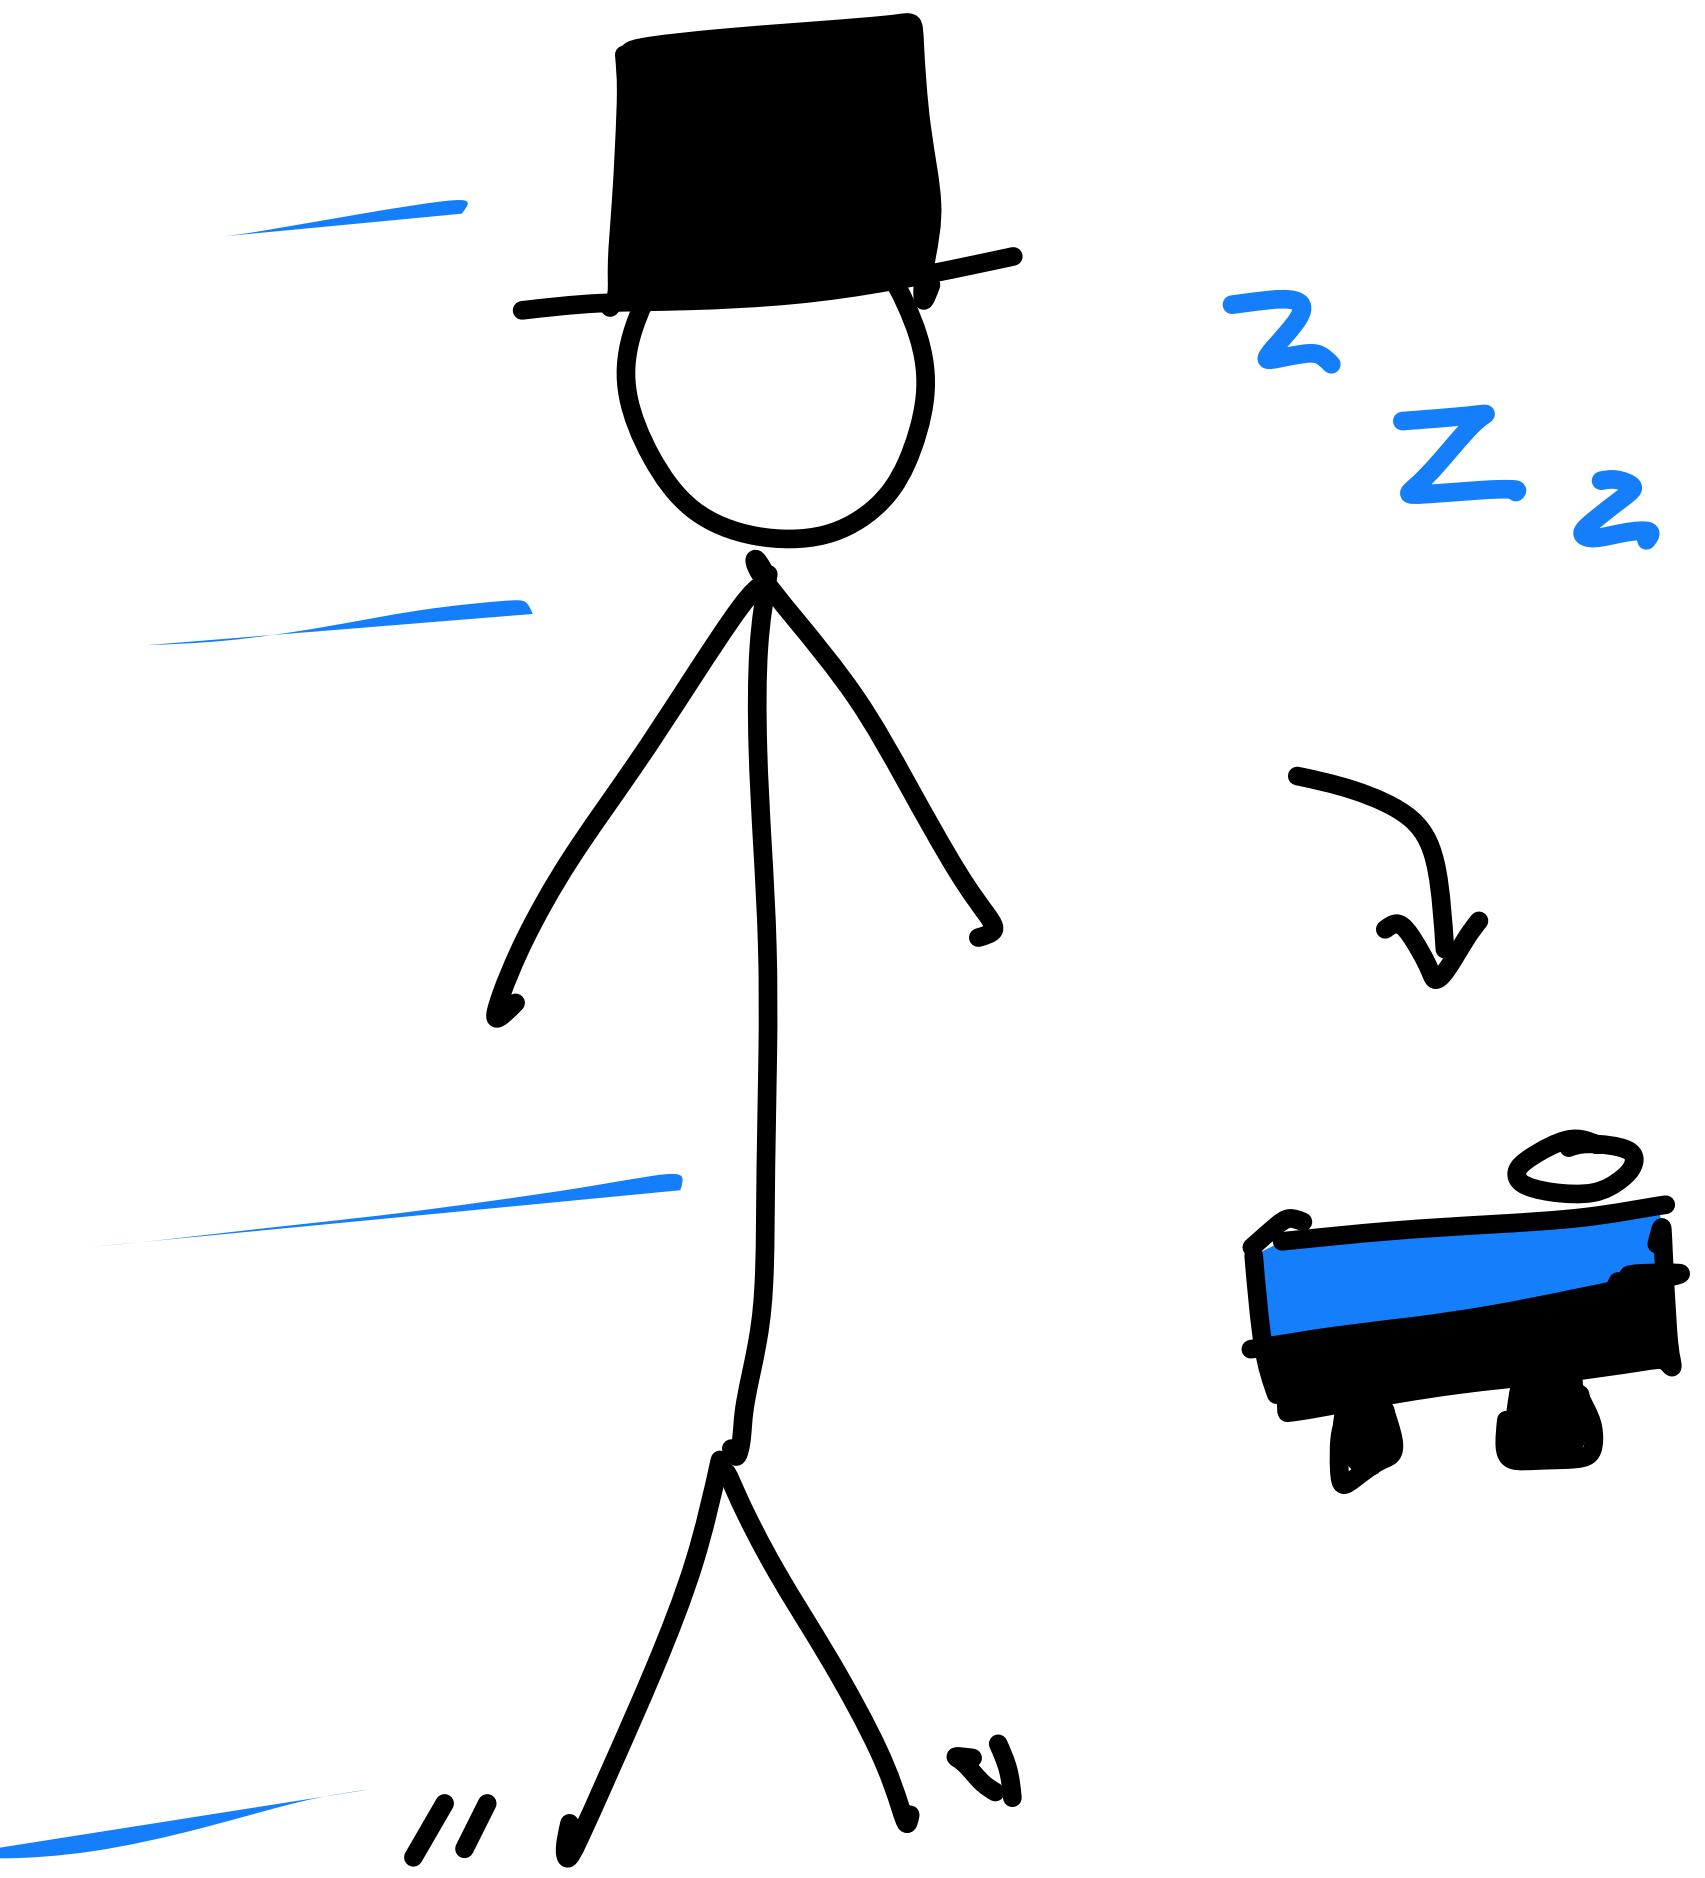
How to sleep at
Night having a cloud
service

One simple hack

To have a cloud service and
sleep at night



Let it burn
and go to
sleep



Thank you for listening

Follow me on social media

DANIEL
SAADA

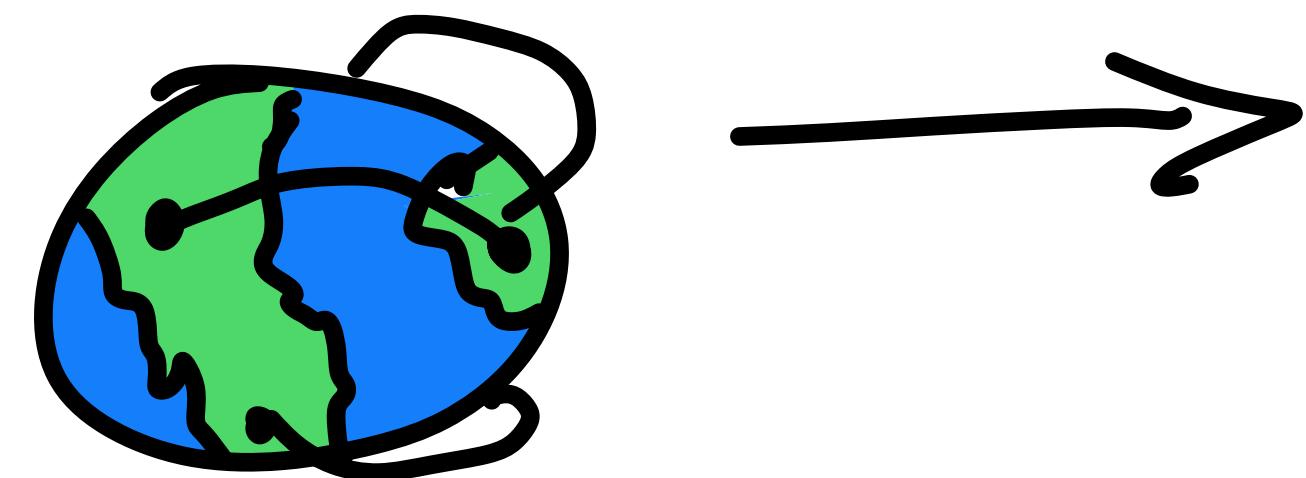
Developer Productivity & Culture

JK

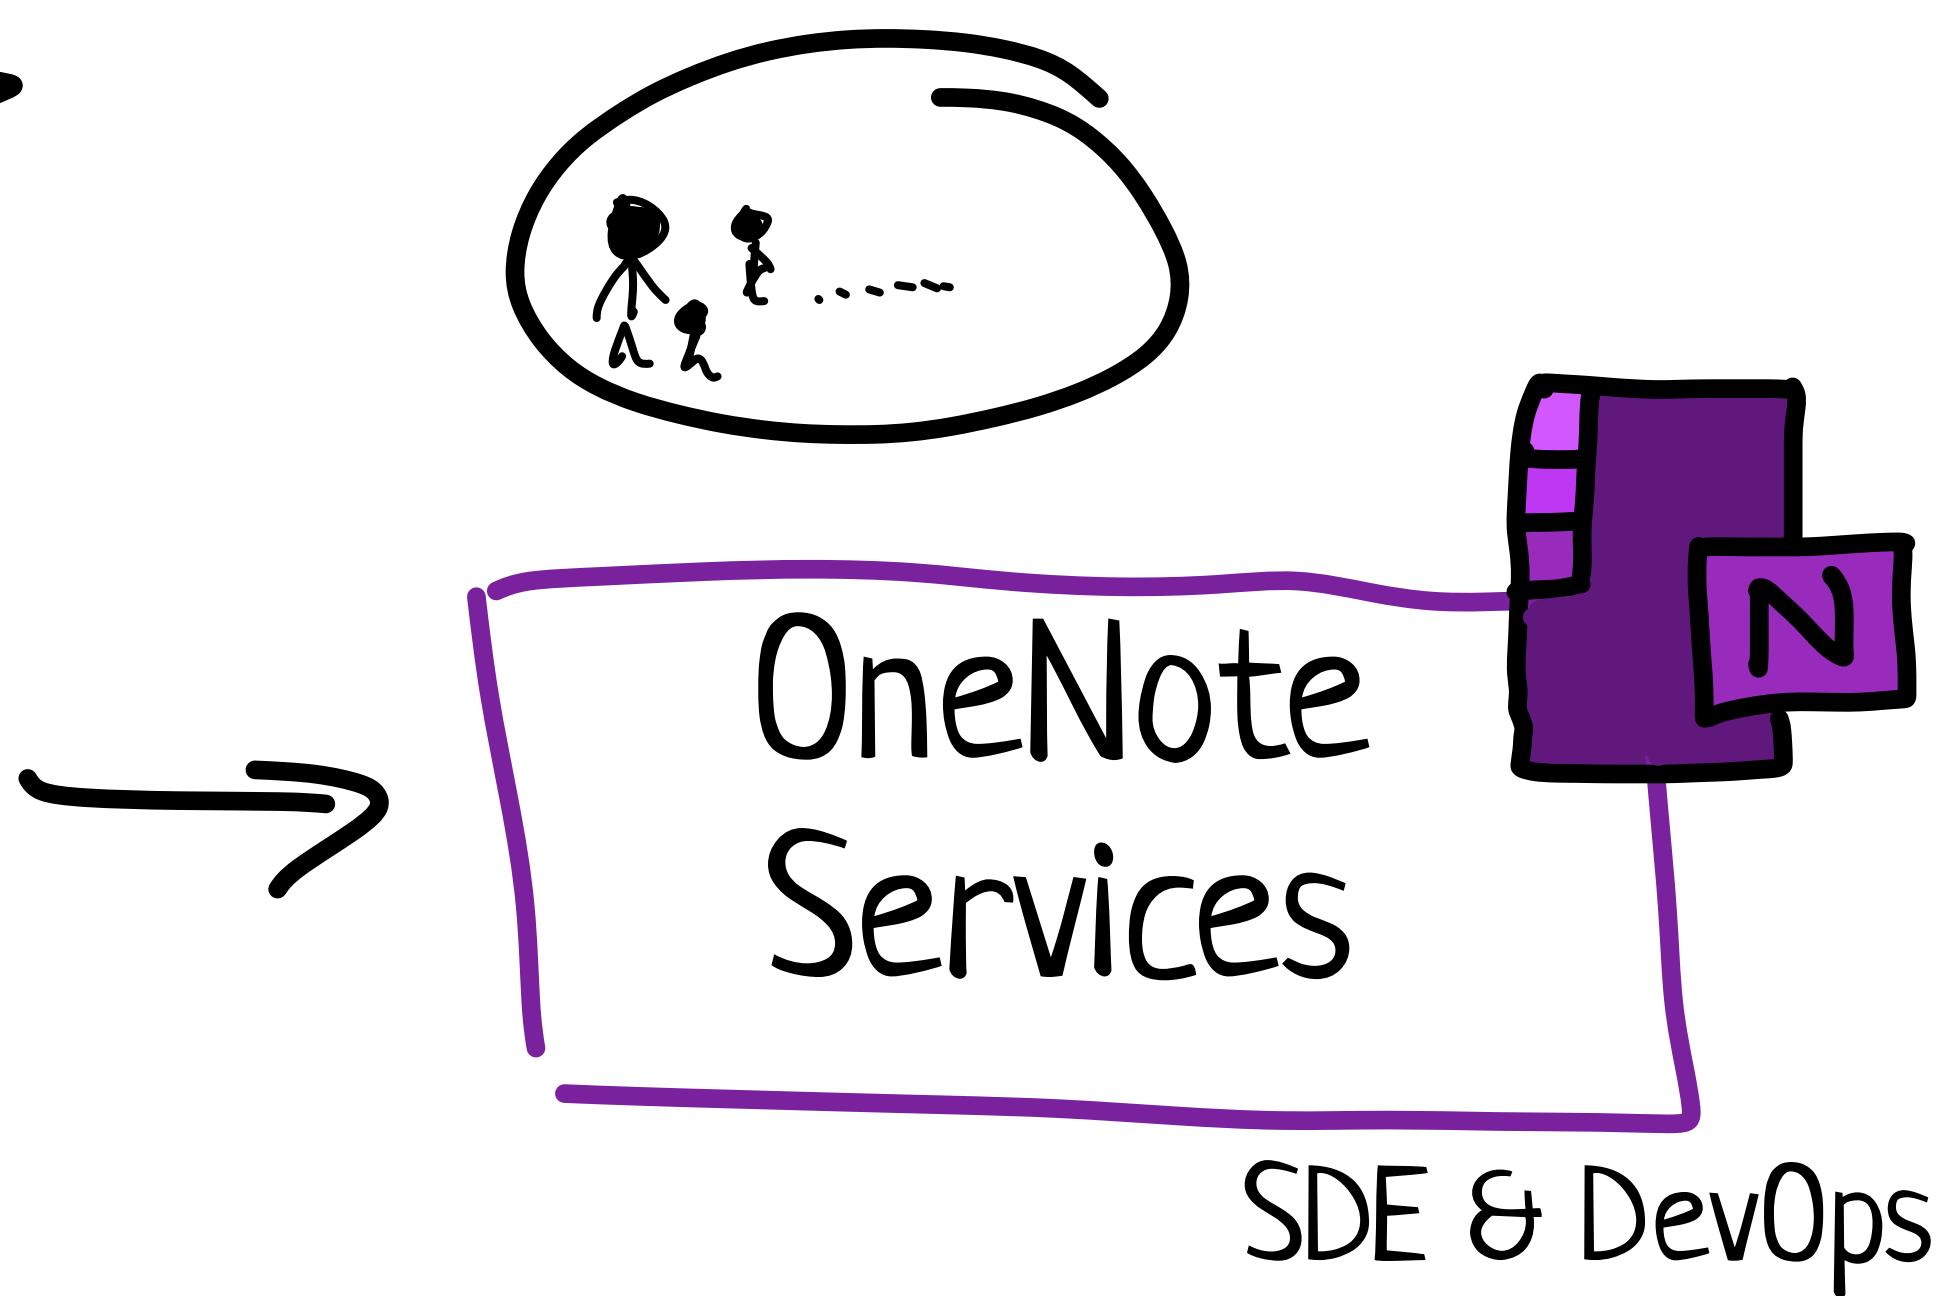
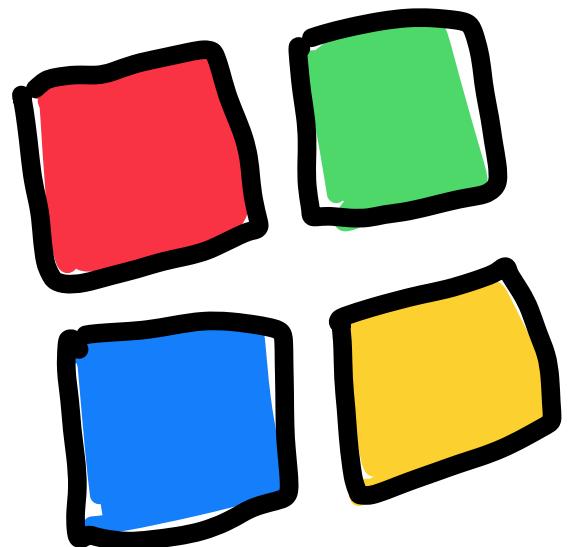


Daniel Sada

Backend services for millions



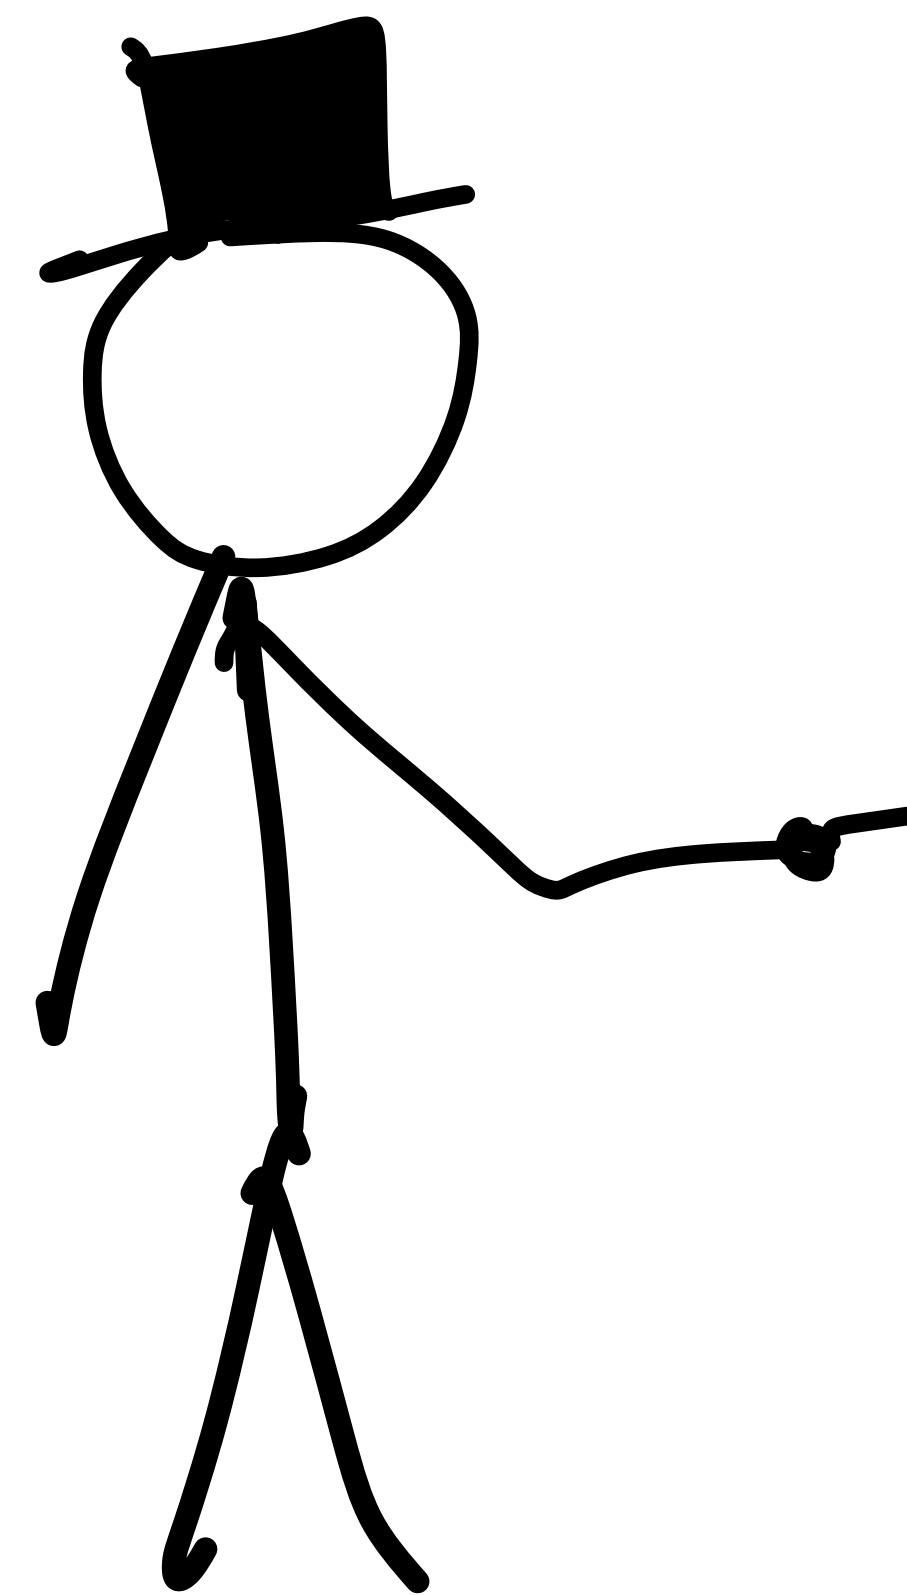
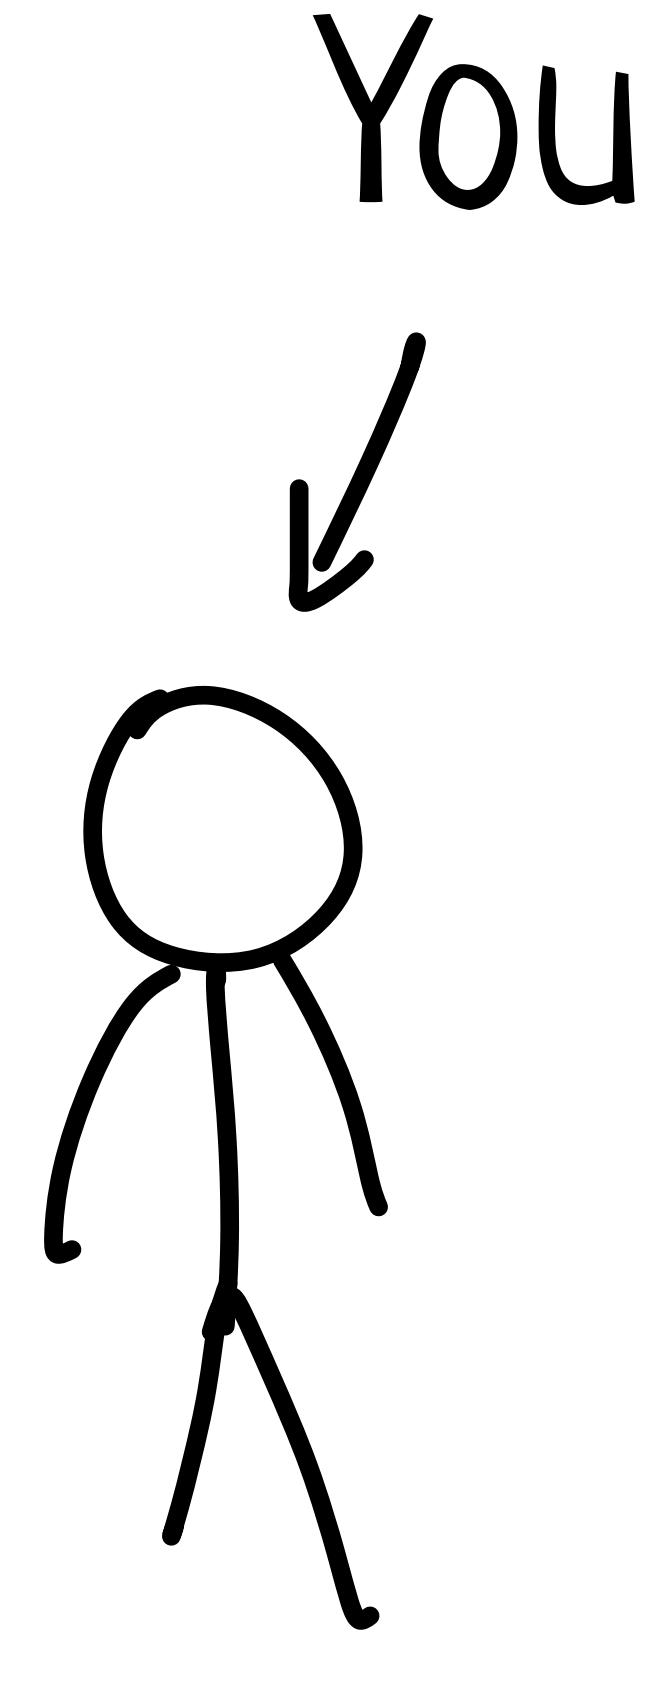
Throughout the world



Snowflake



Developer
Productivity



General concepts

That might apply to you
or not, but have helped
me in the past

We'll see the following levels:

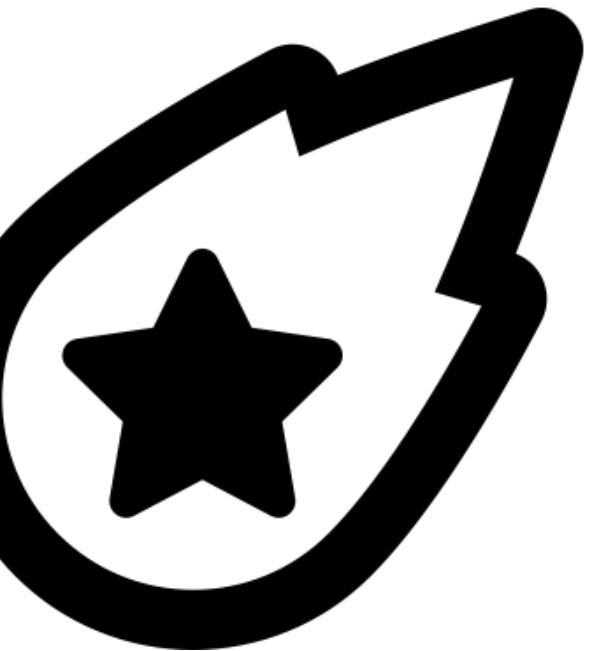
Easy



Medium



Hard



EXPERT



Easy *

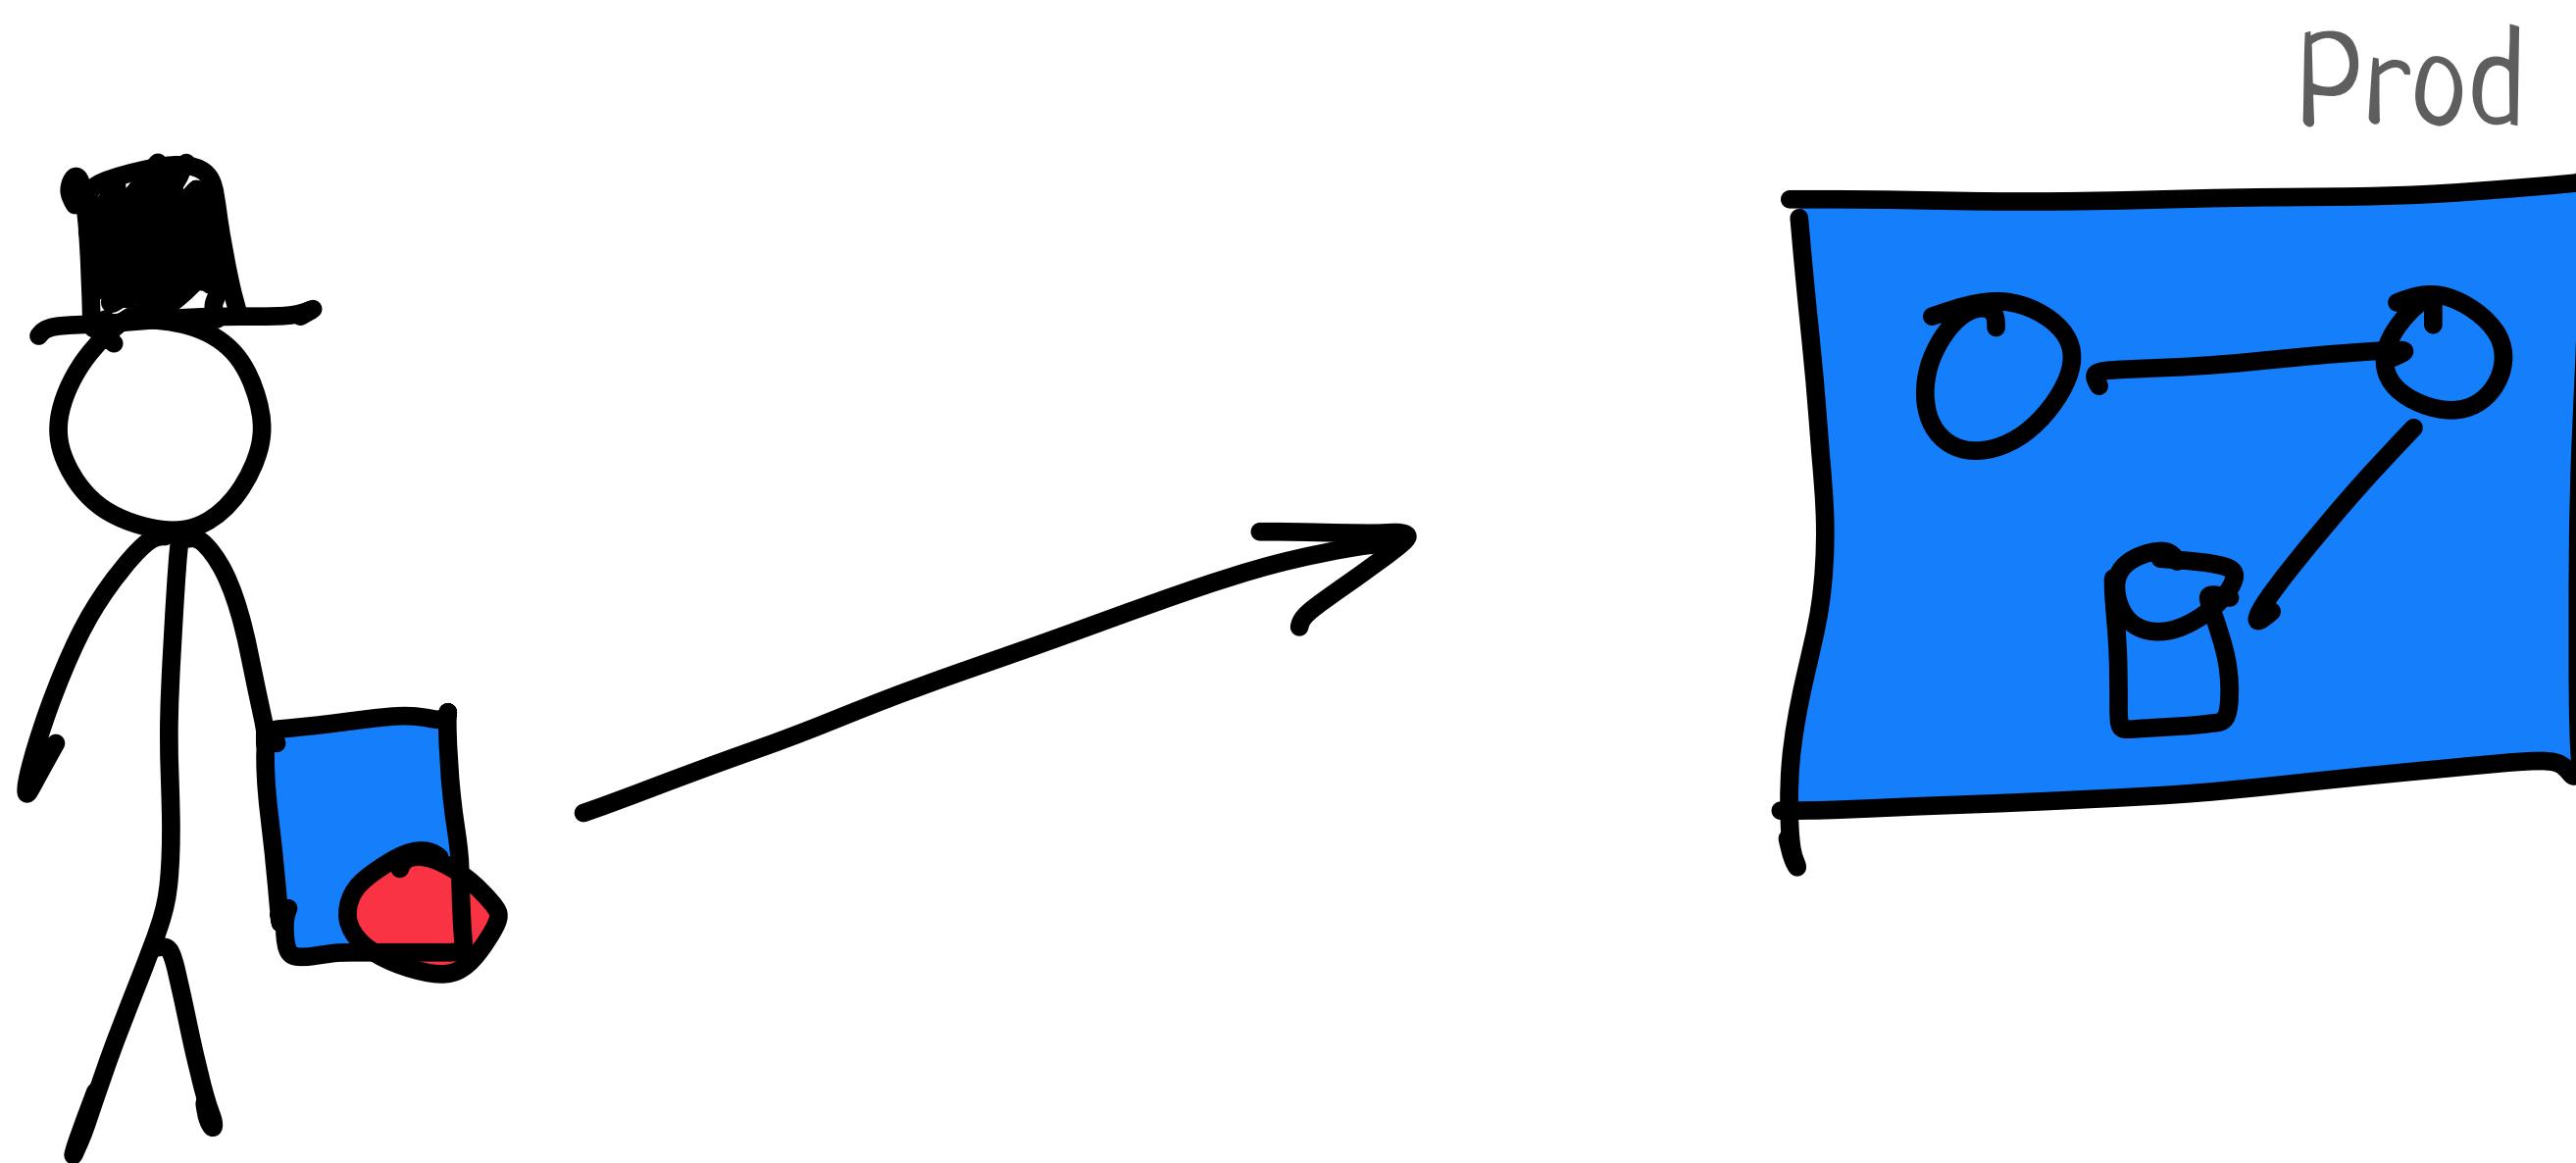


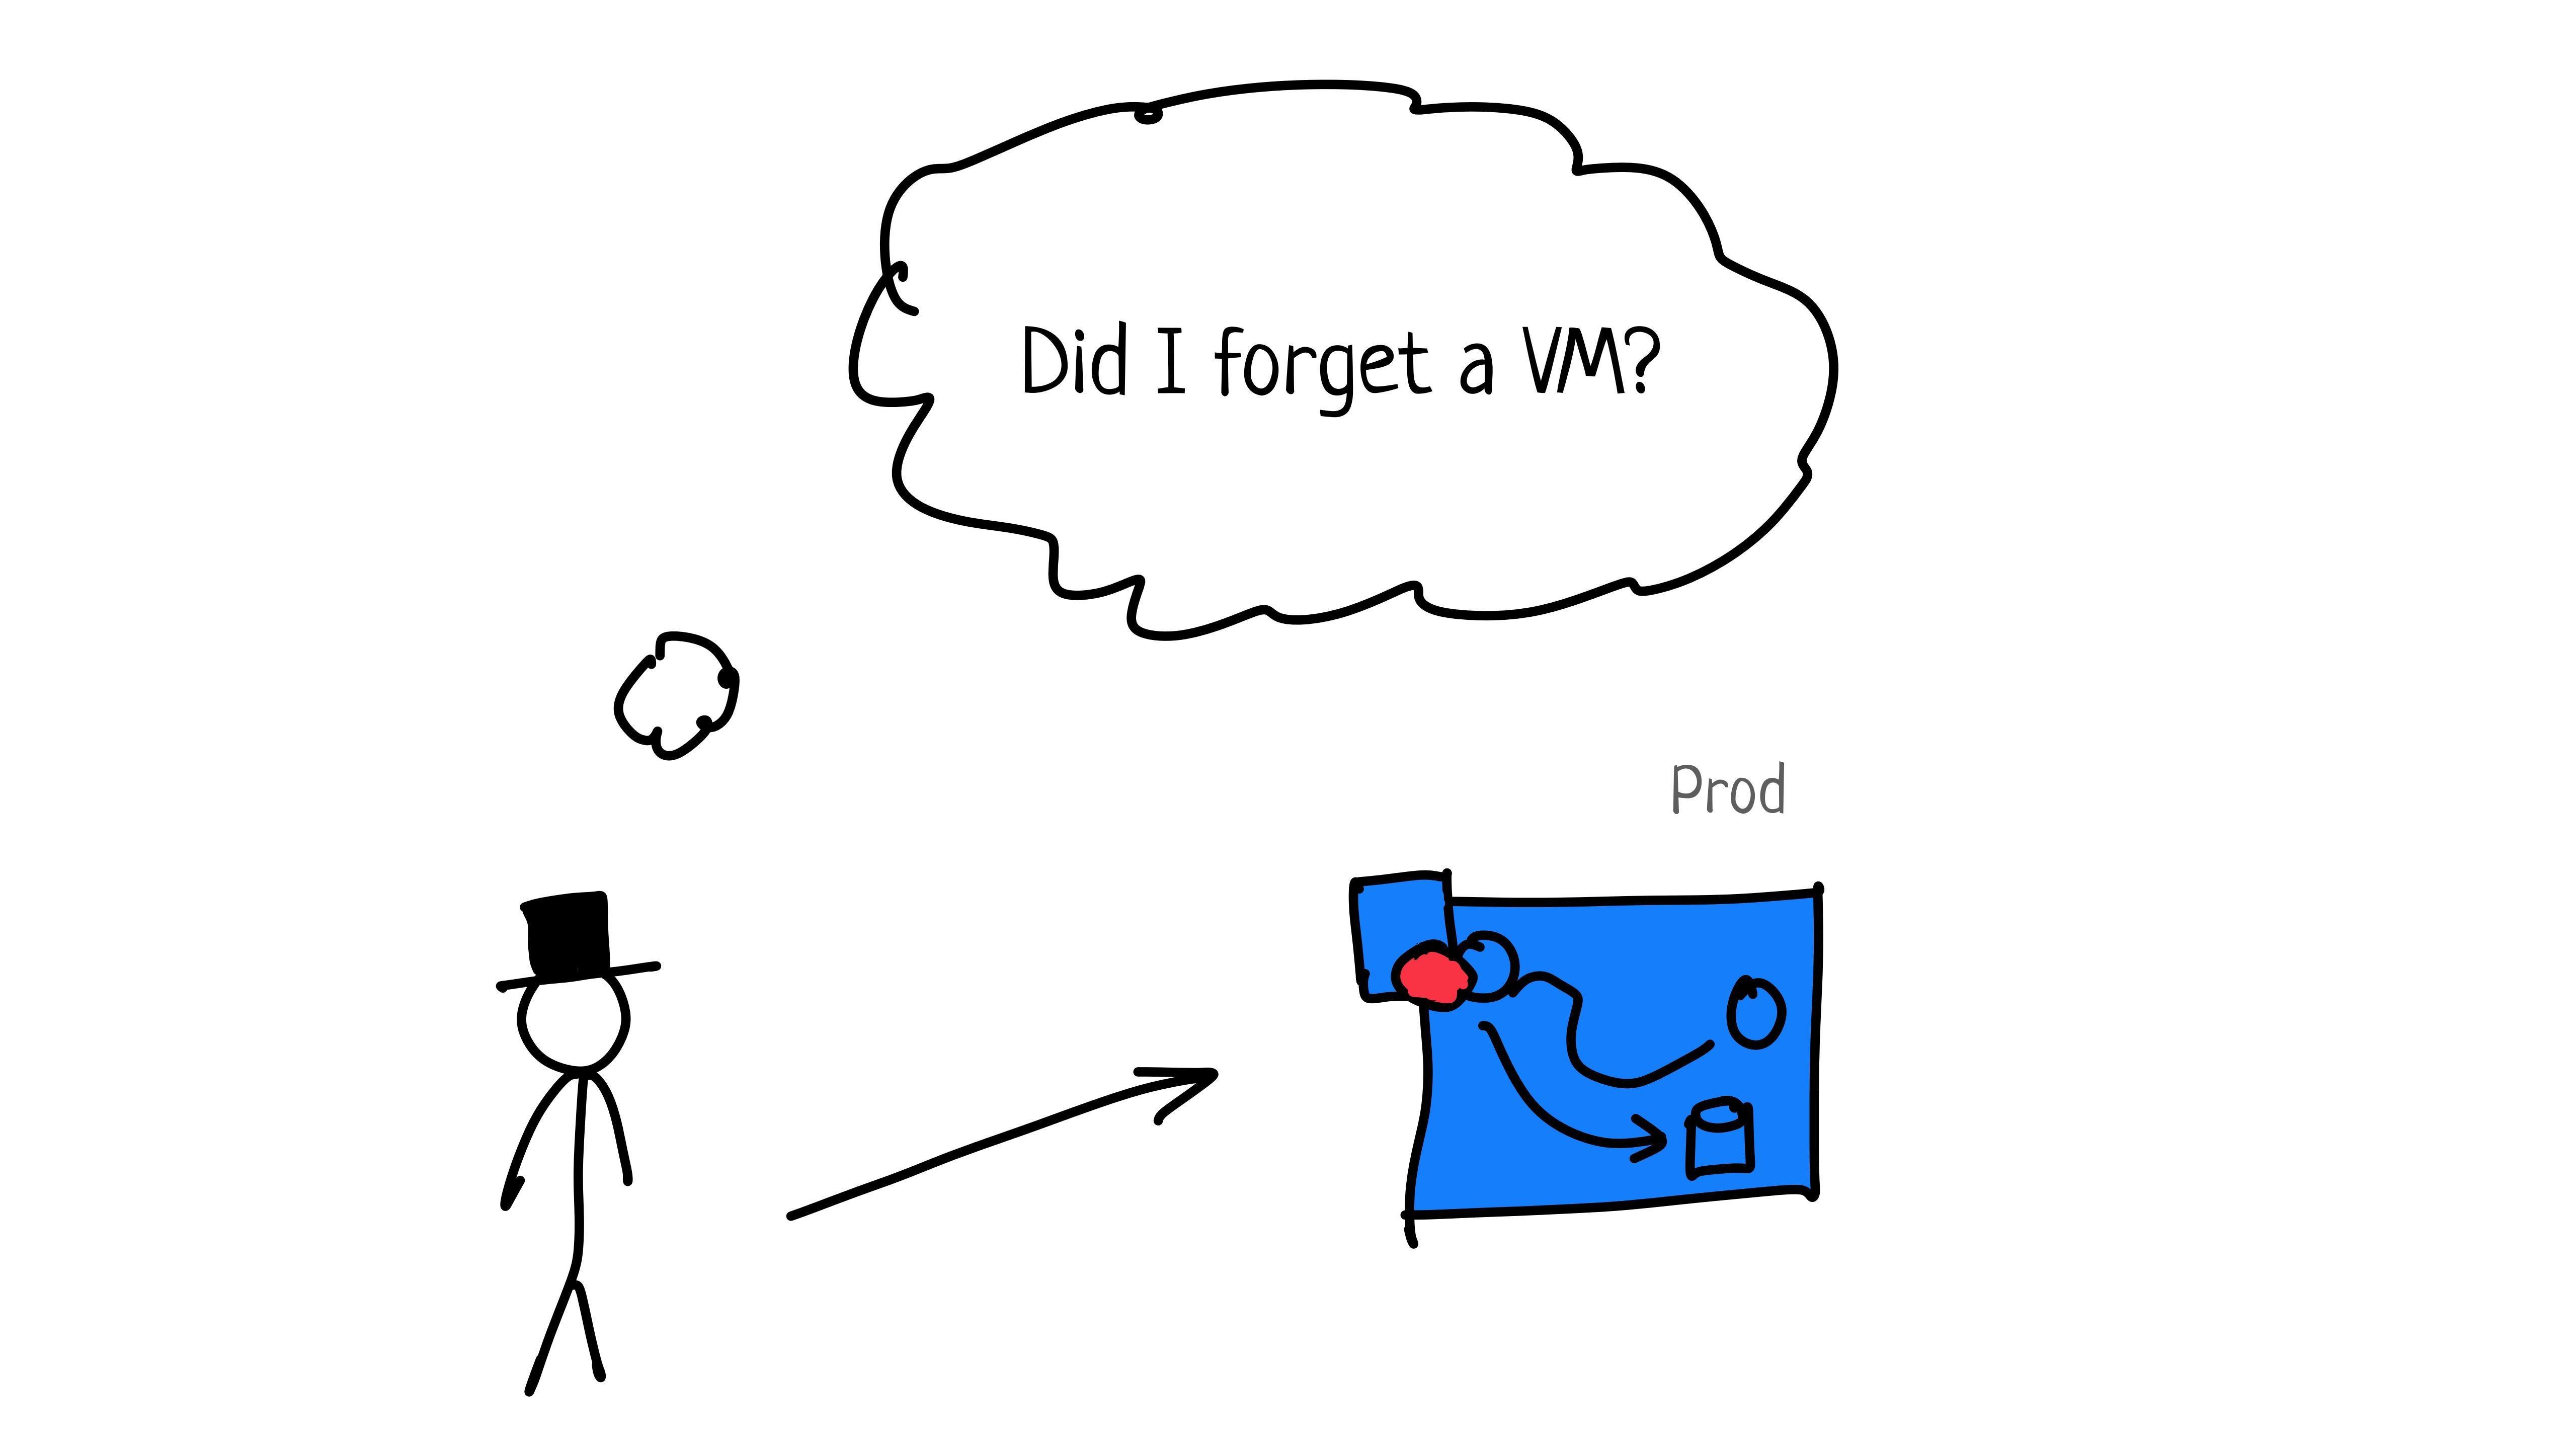
* Not really easy, but simpler than other levels.

IAC

Infrastructure as Code

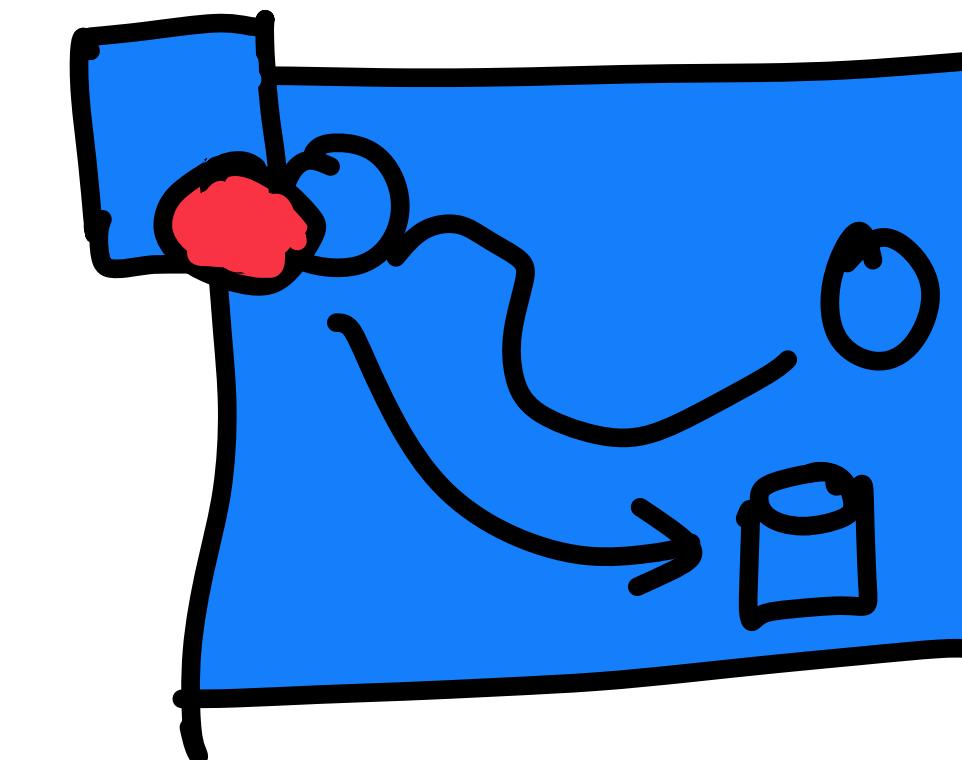
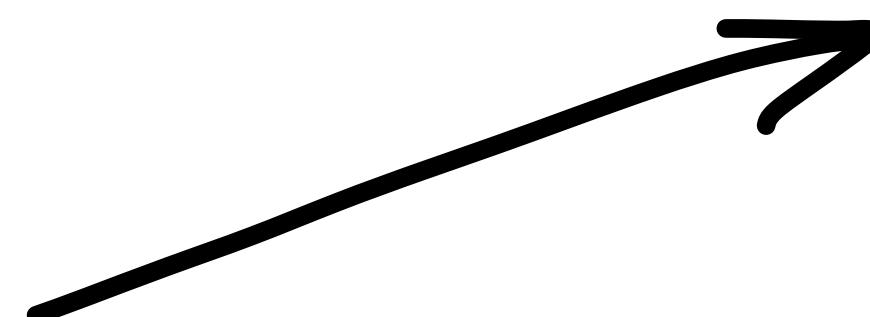
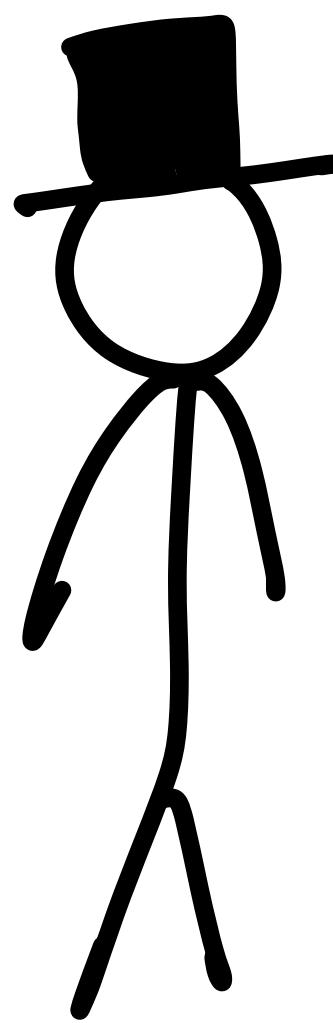
Deploying manually infrastructure in the cloud

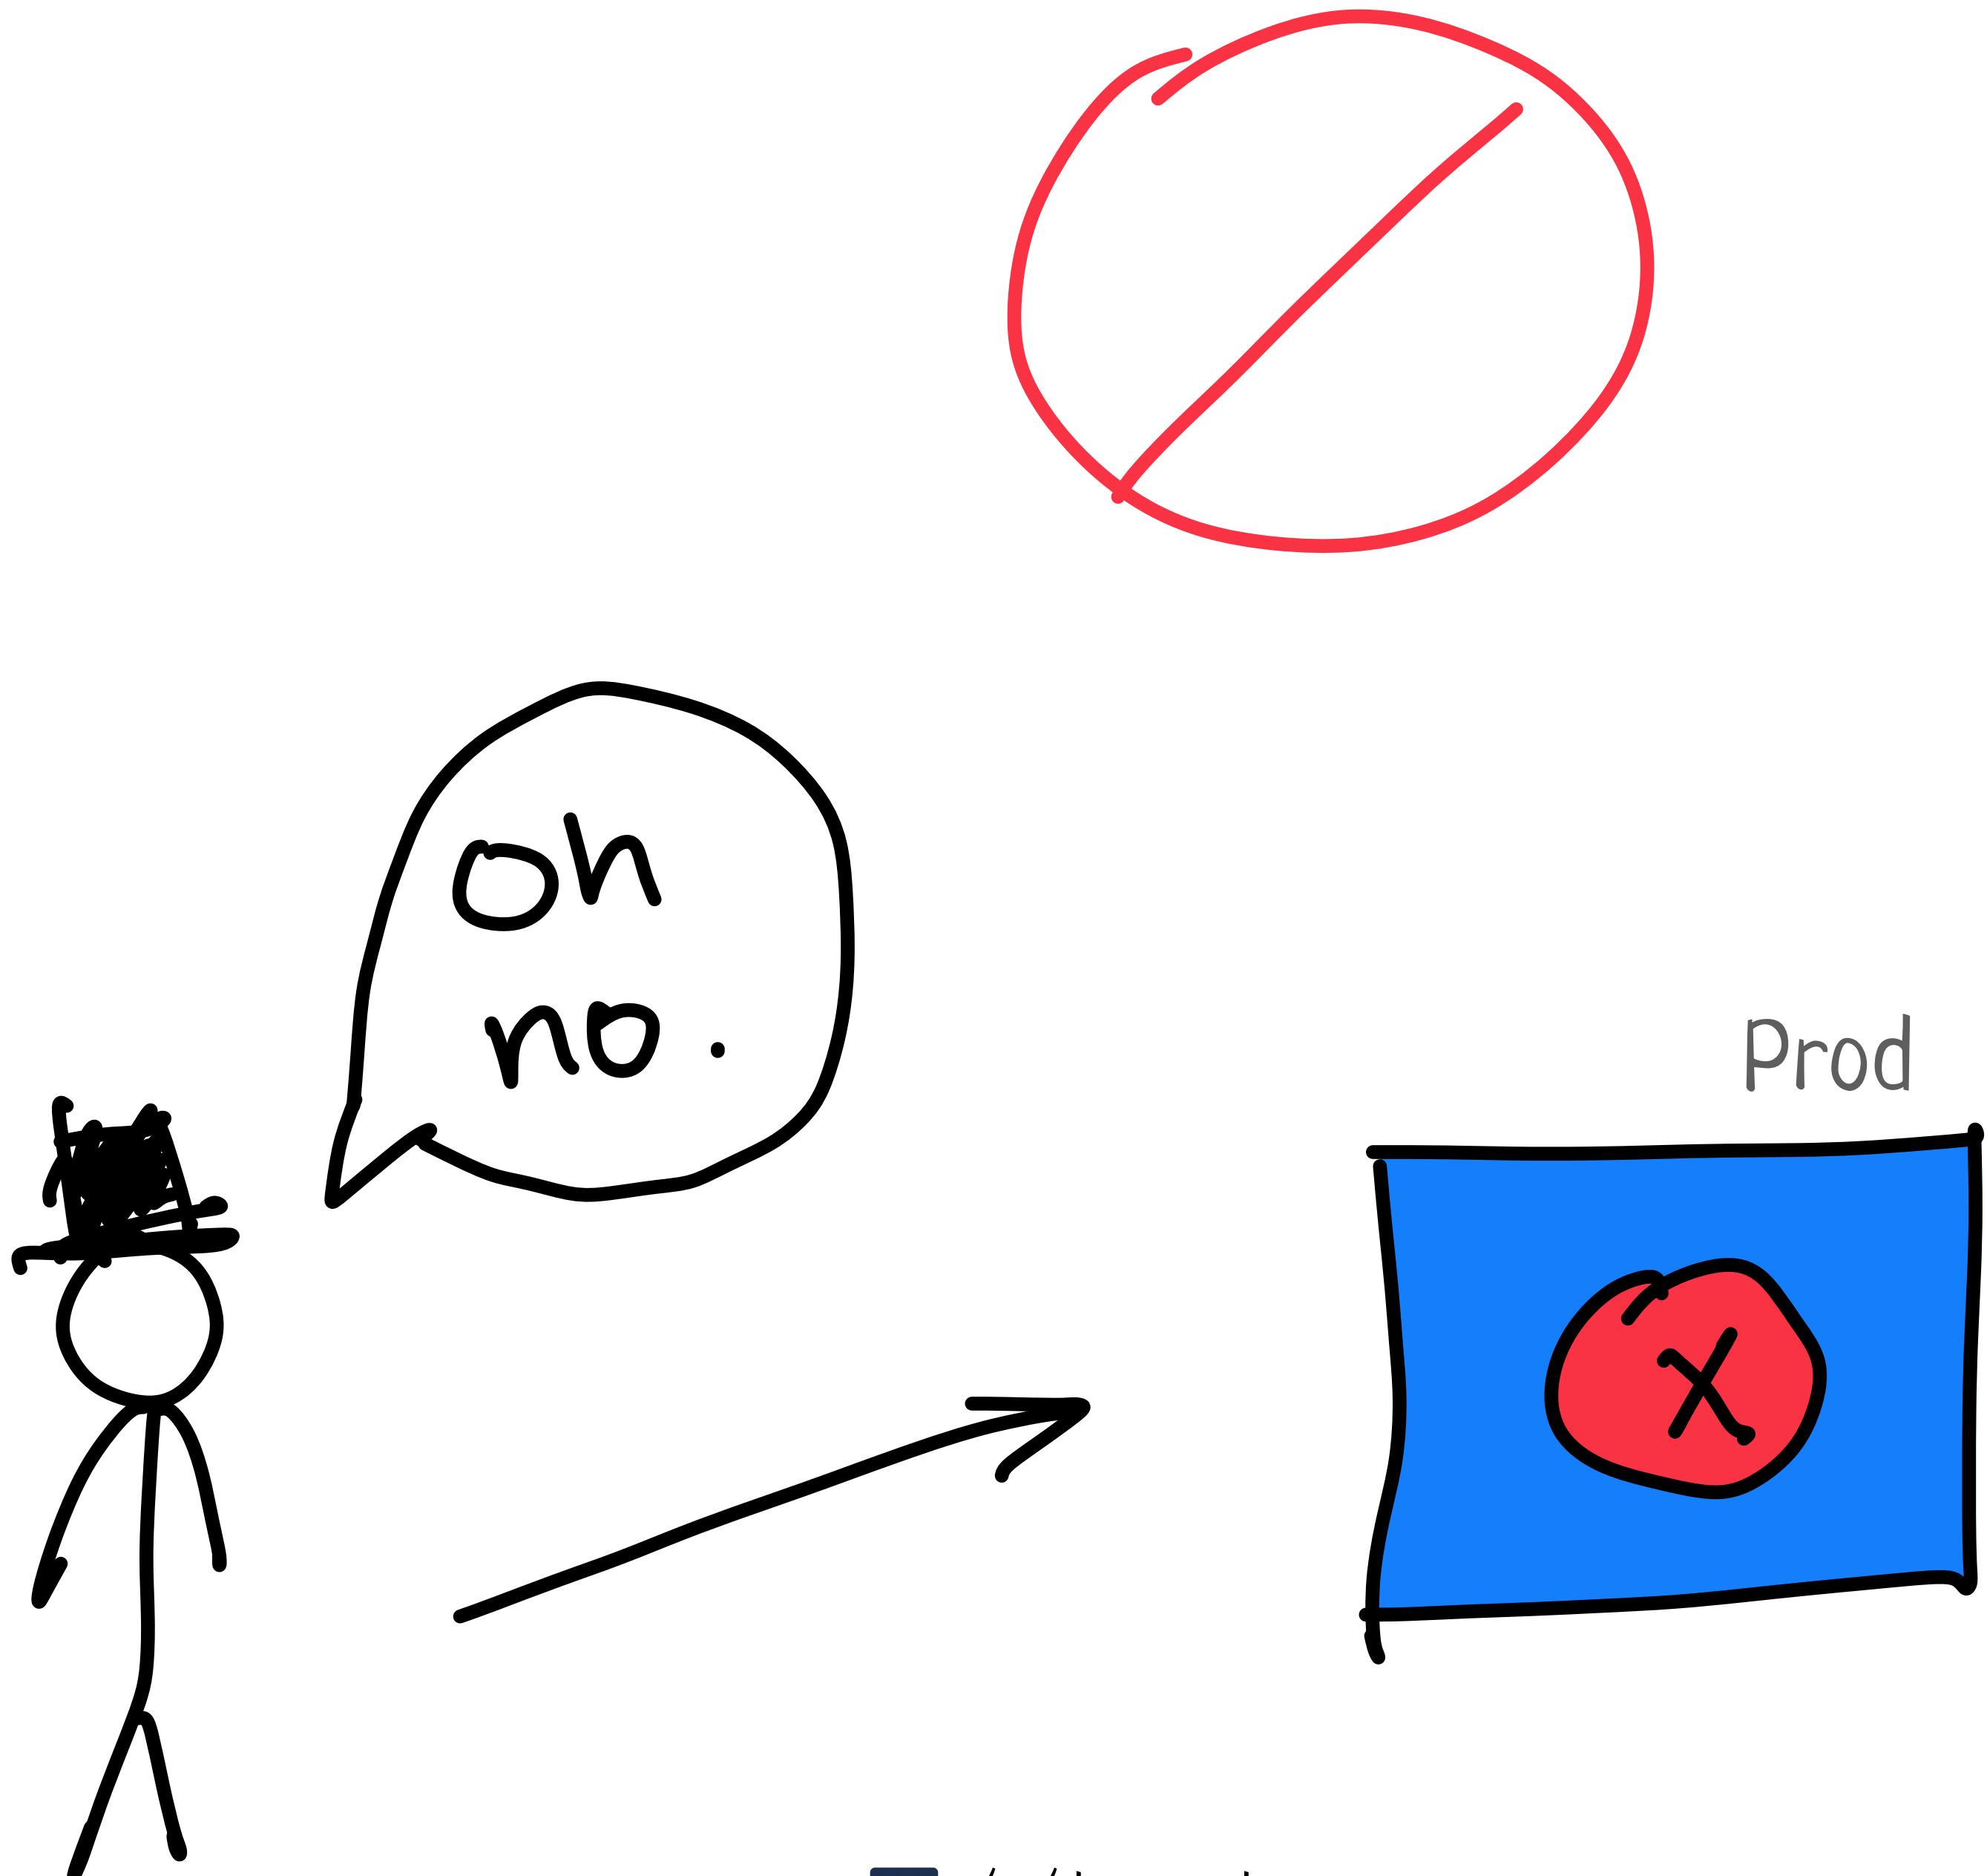




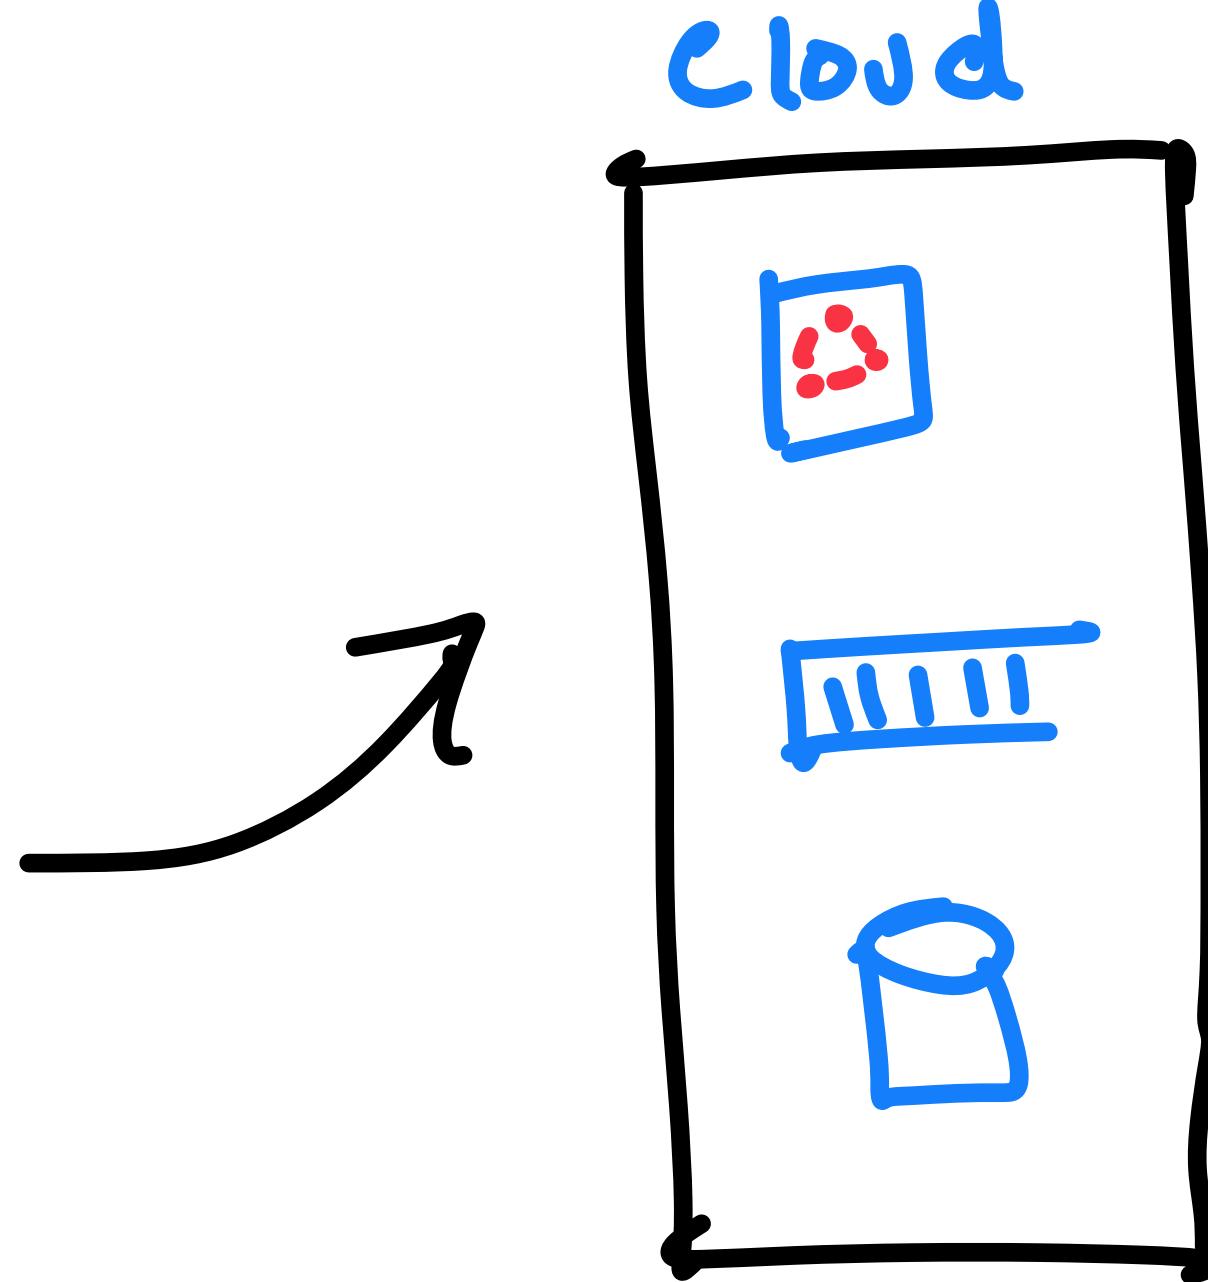
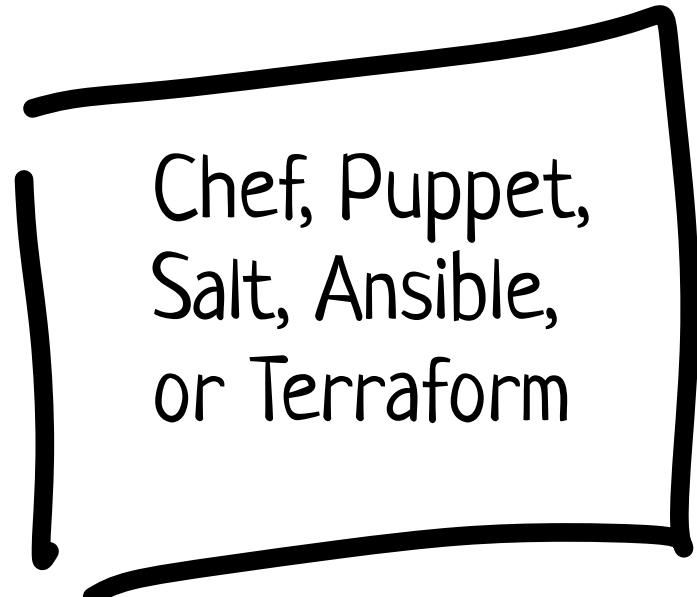
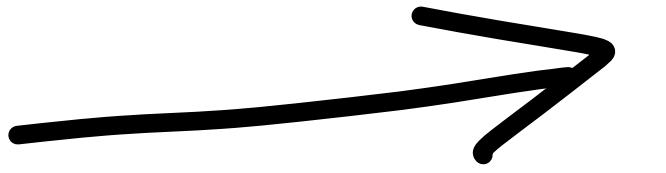
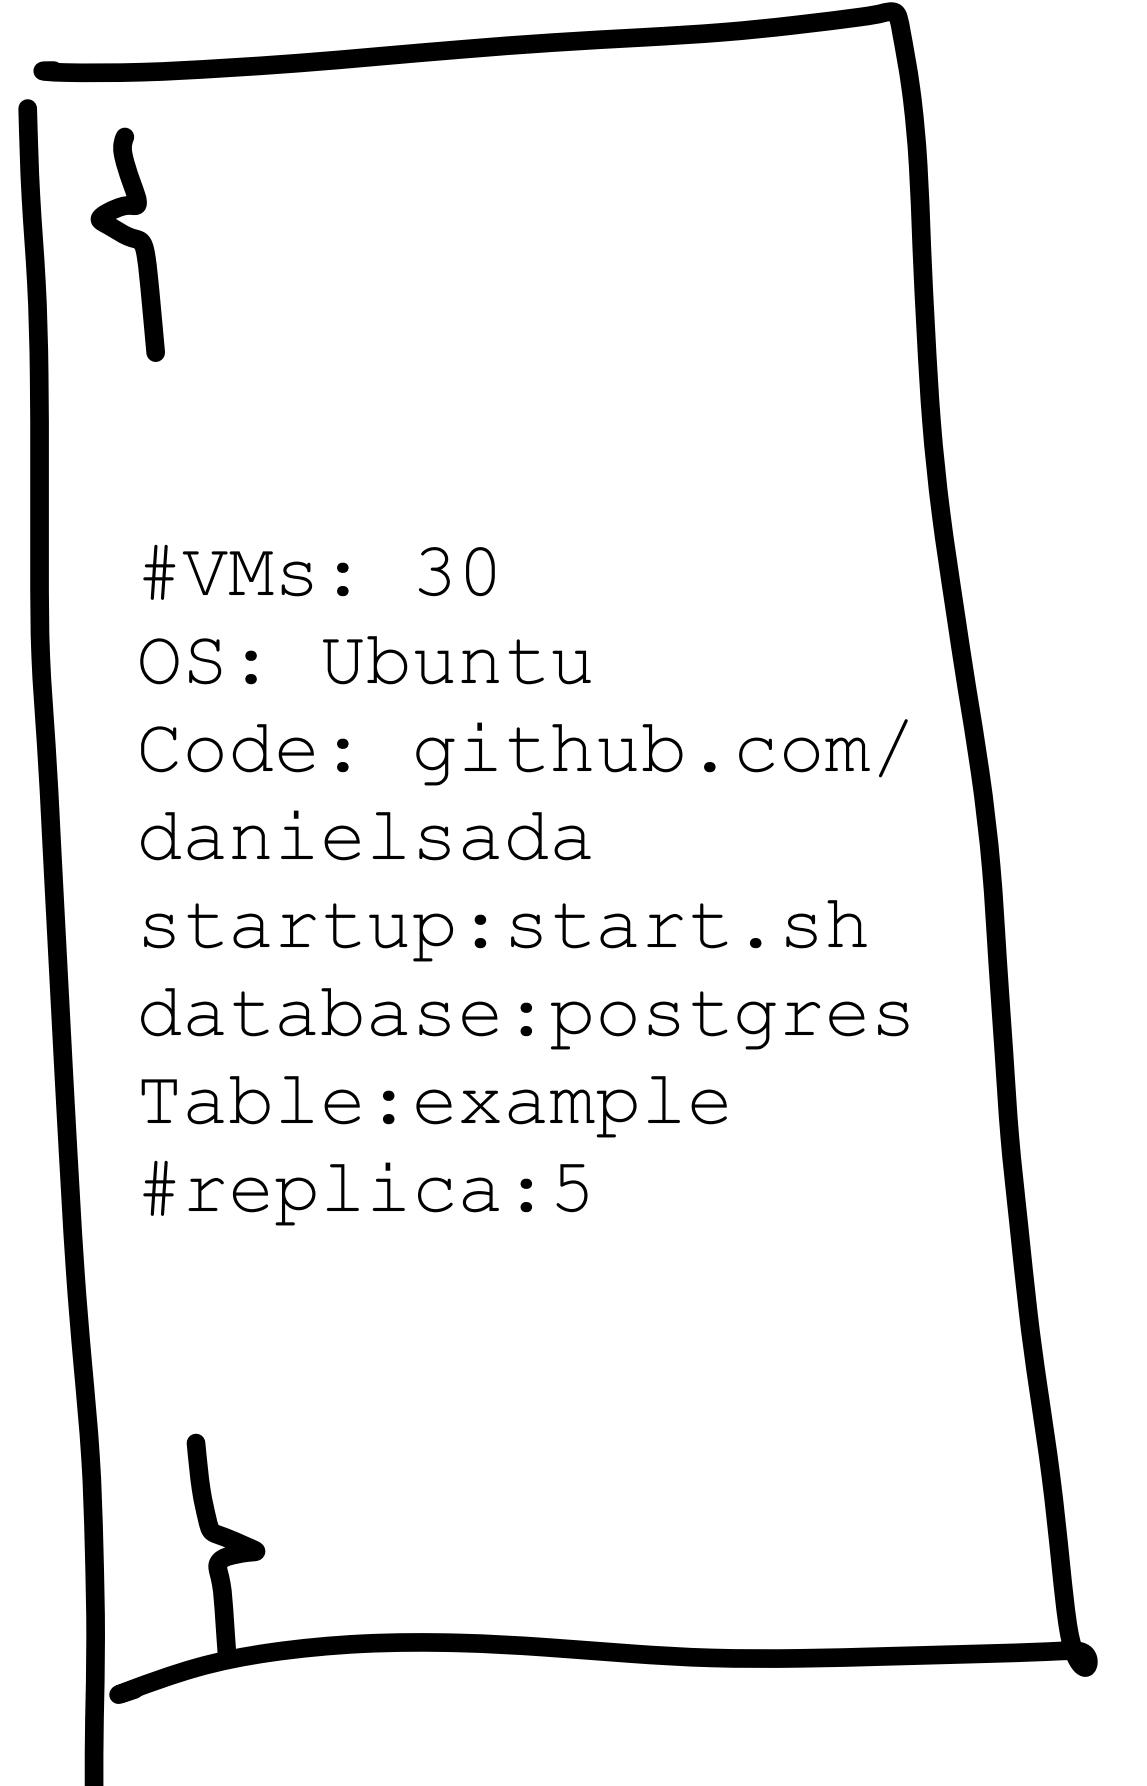
Did I forget a VM?

Prod



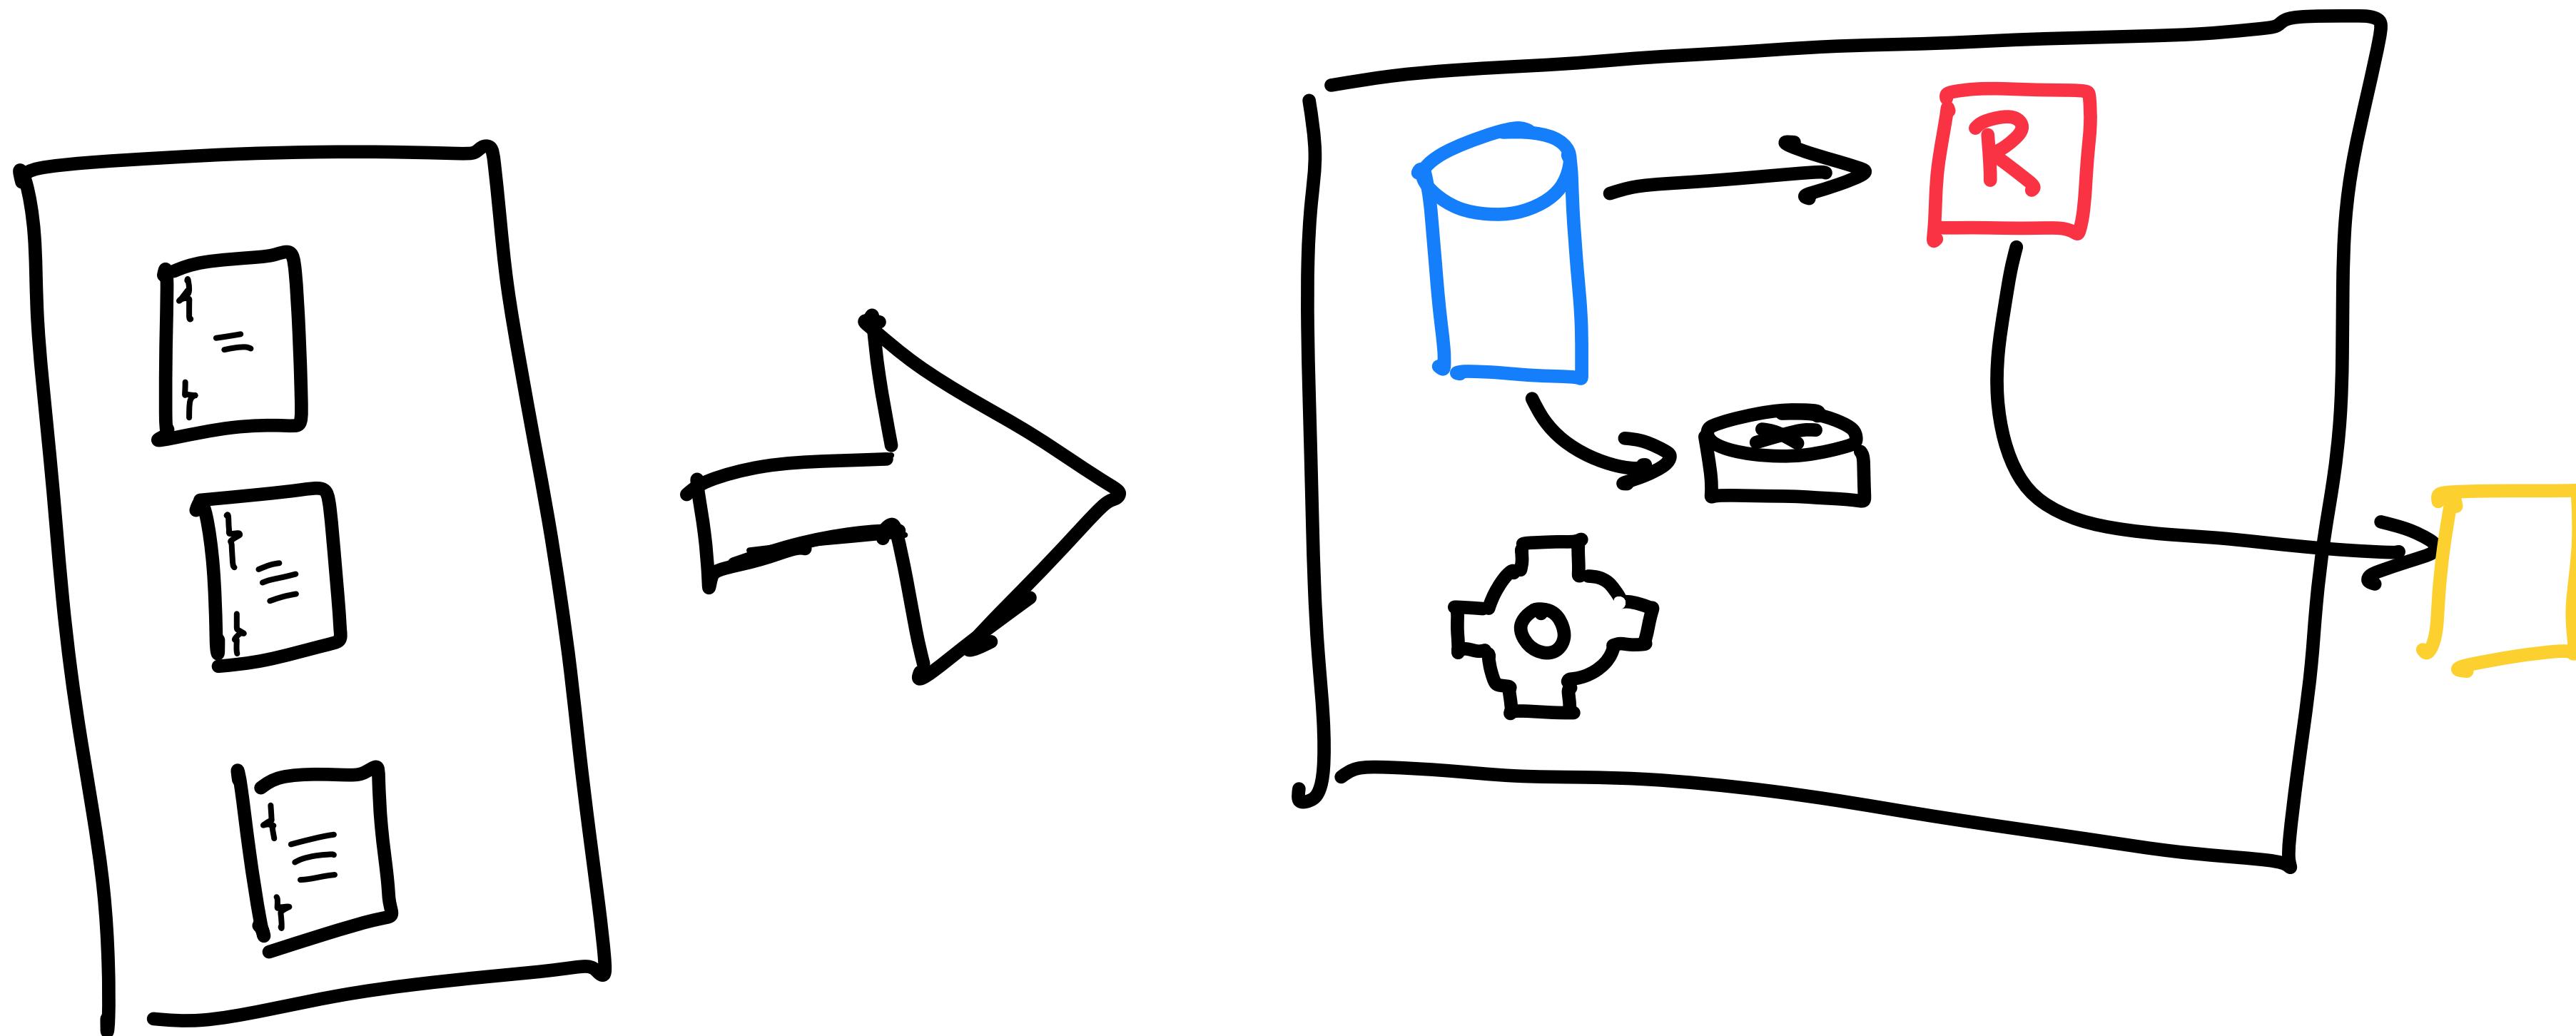


Level 1 – IAC



Script to Deploy

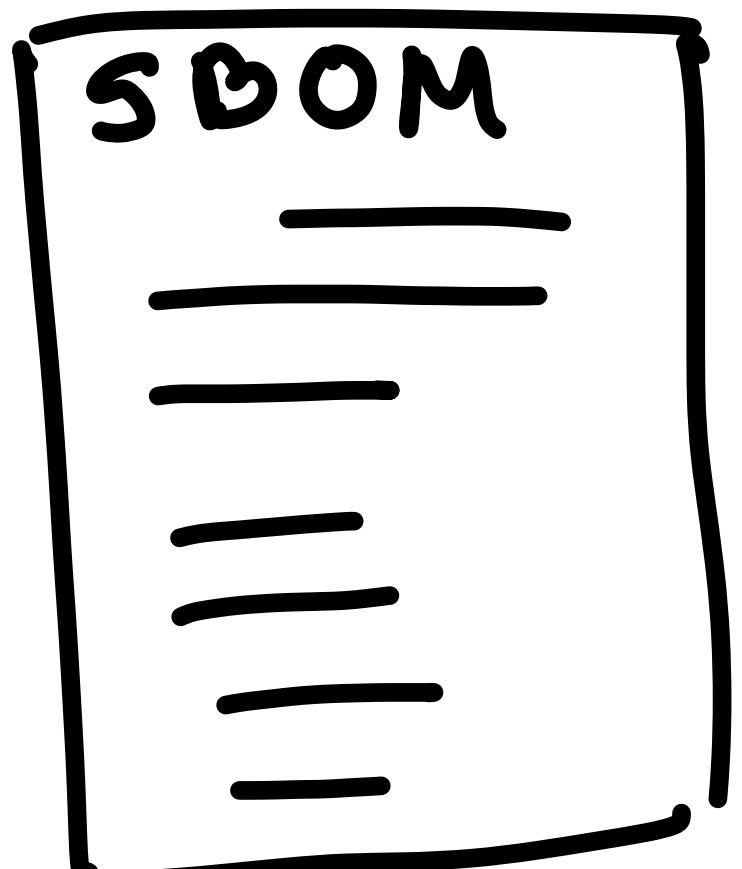
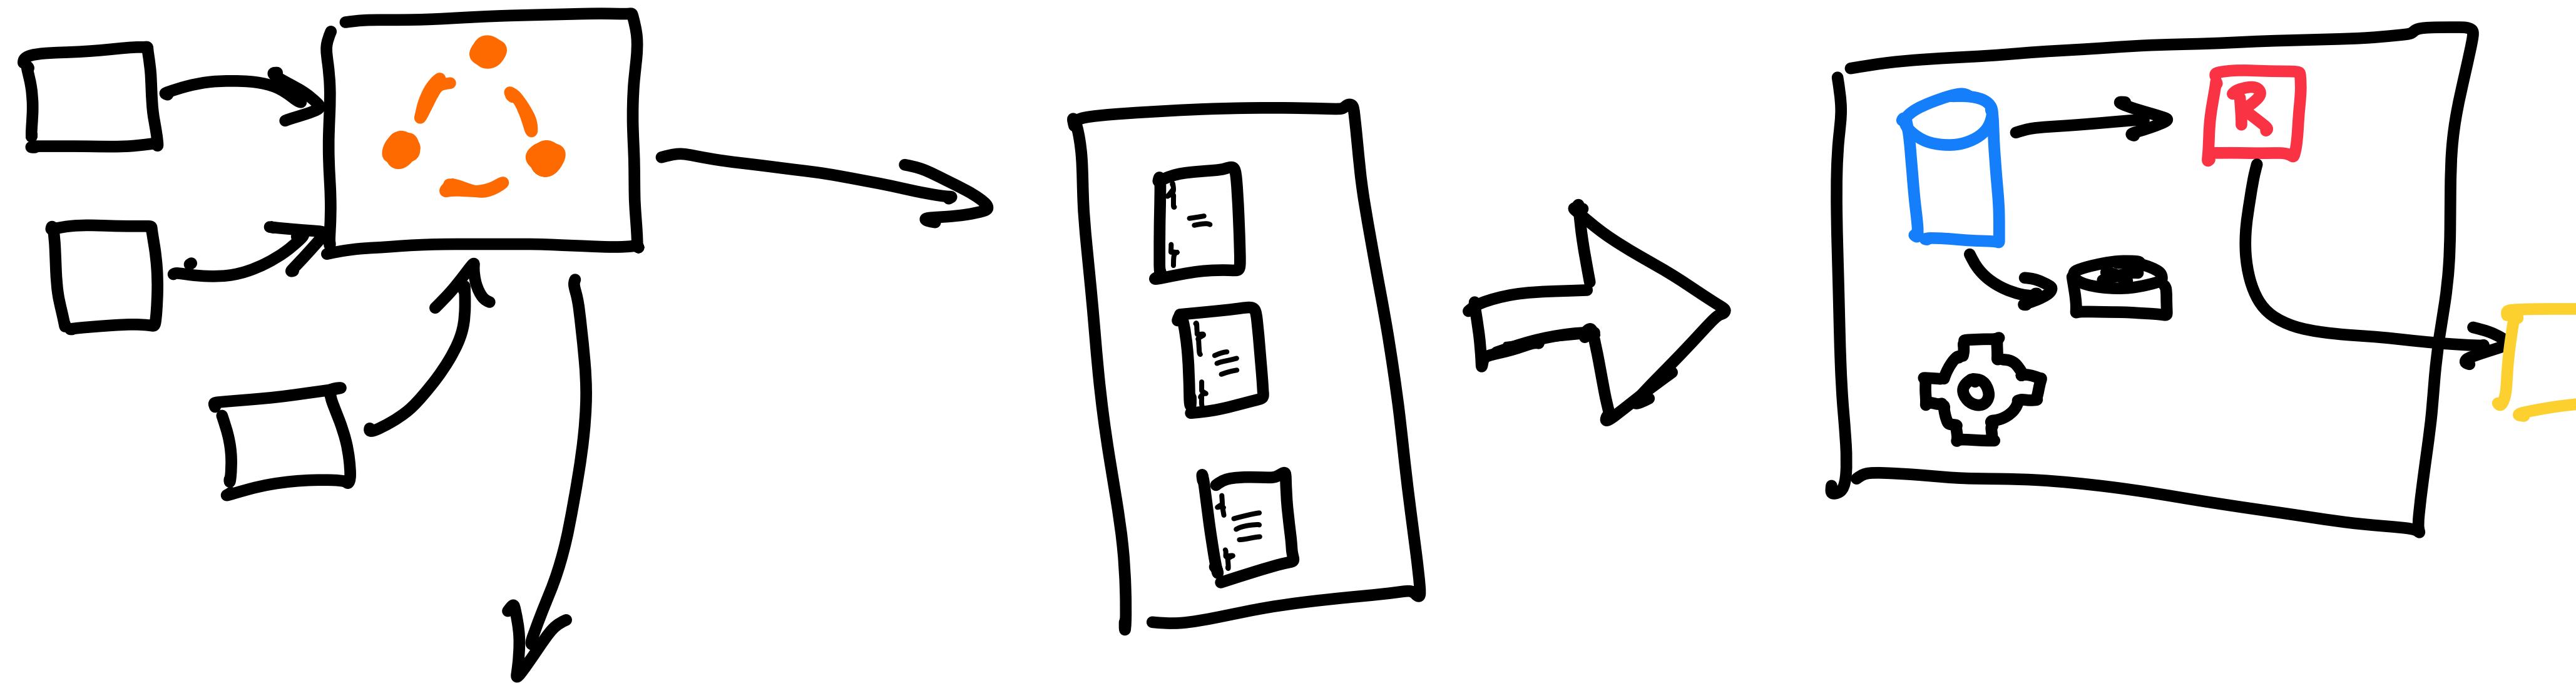
Level 2 – IAC



Set of scripts to describe a system.

Level 3 – IAC

Containers



Custom images + SBOM + Scripts

Software
Bill
Of
Materials

IAC

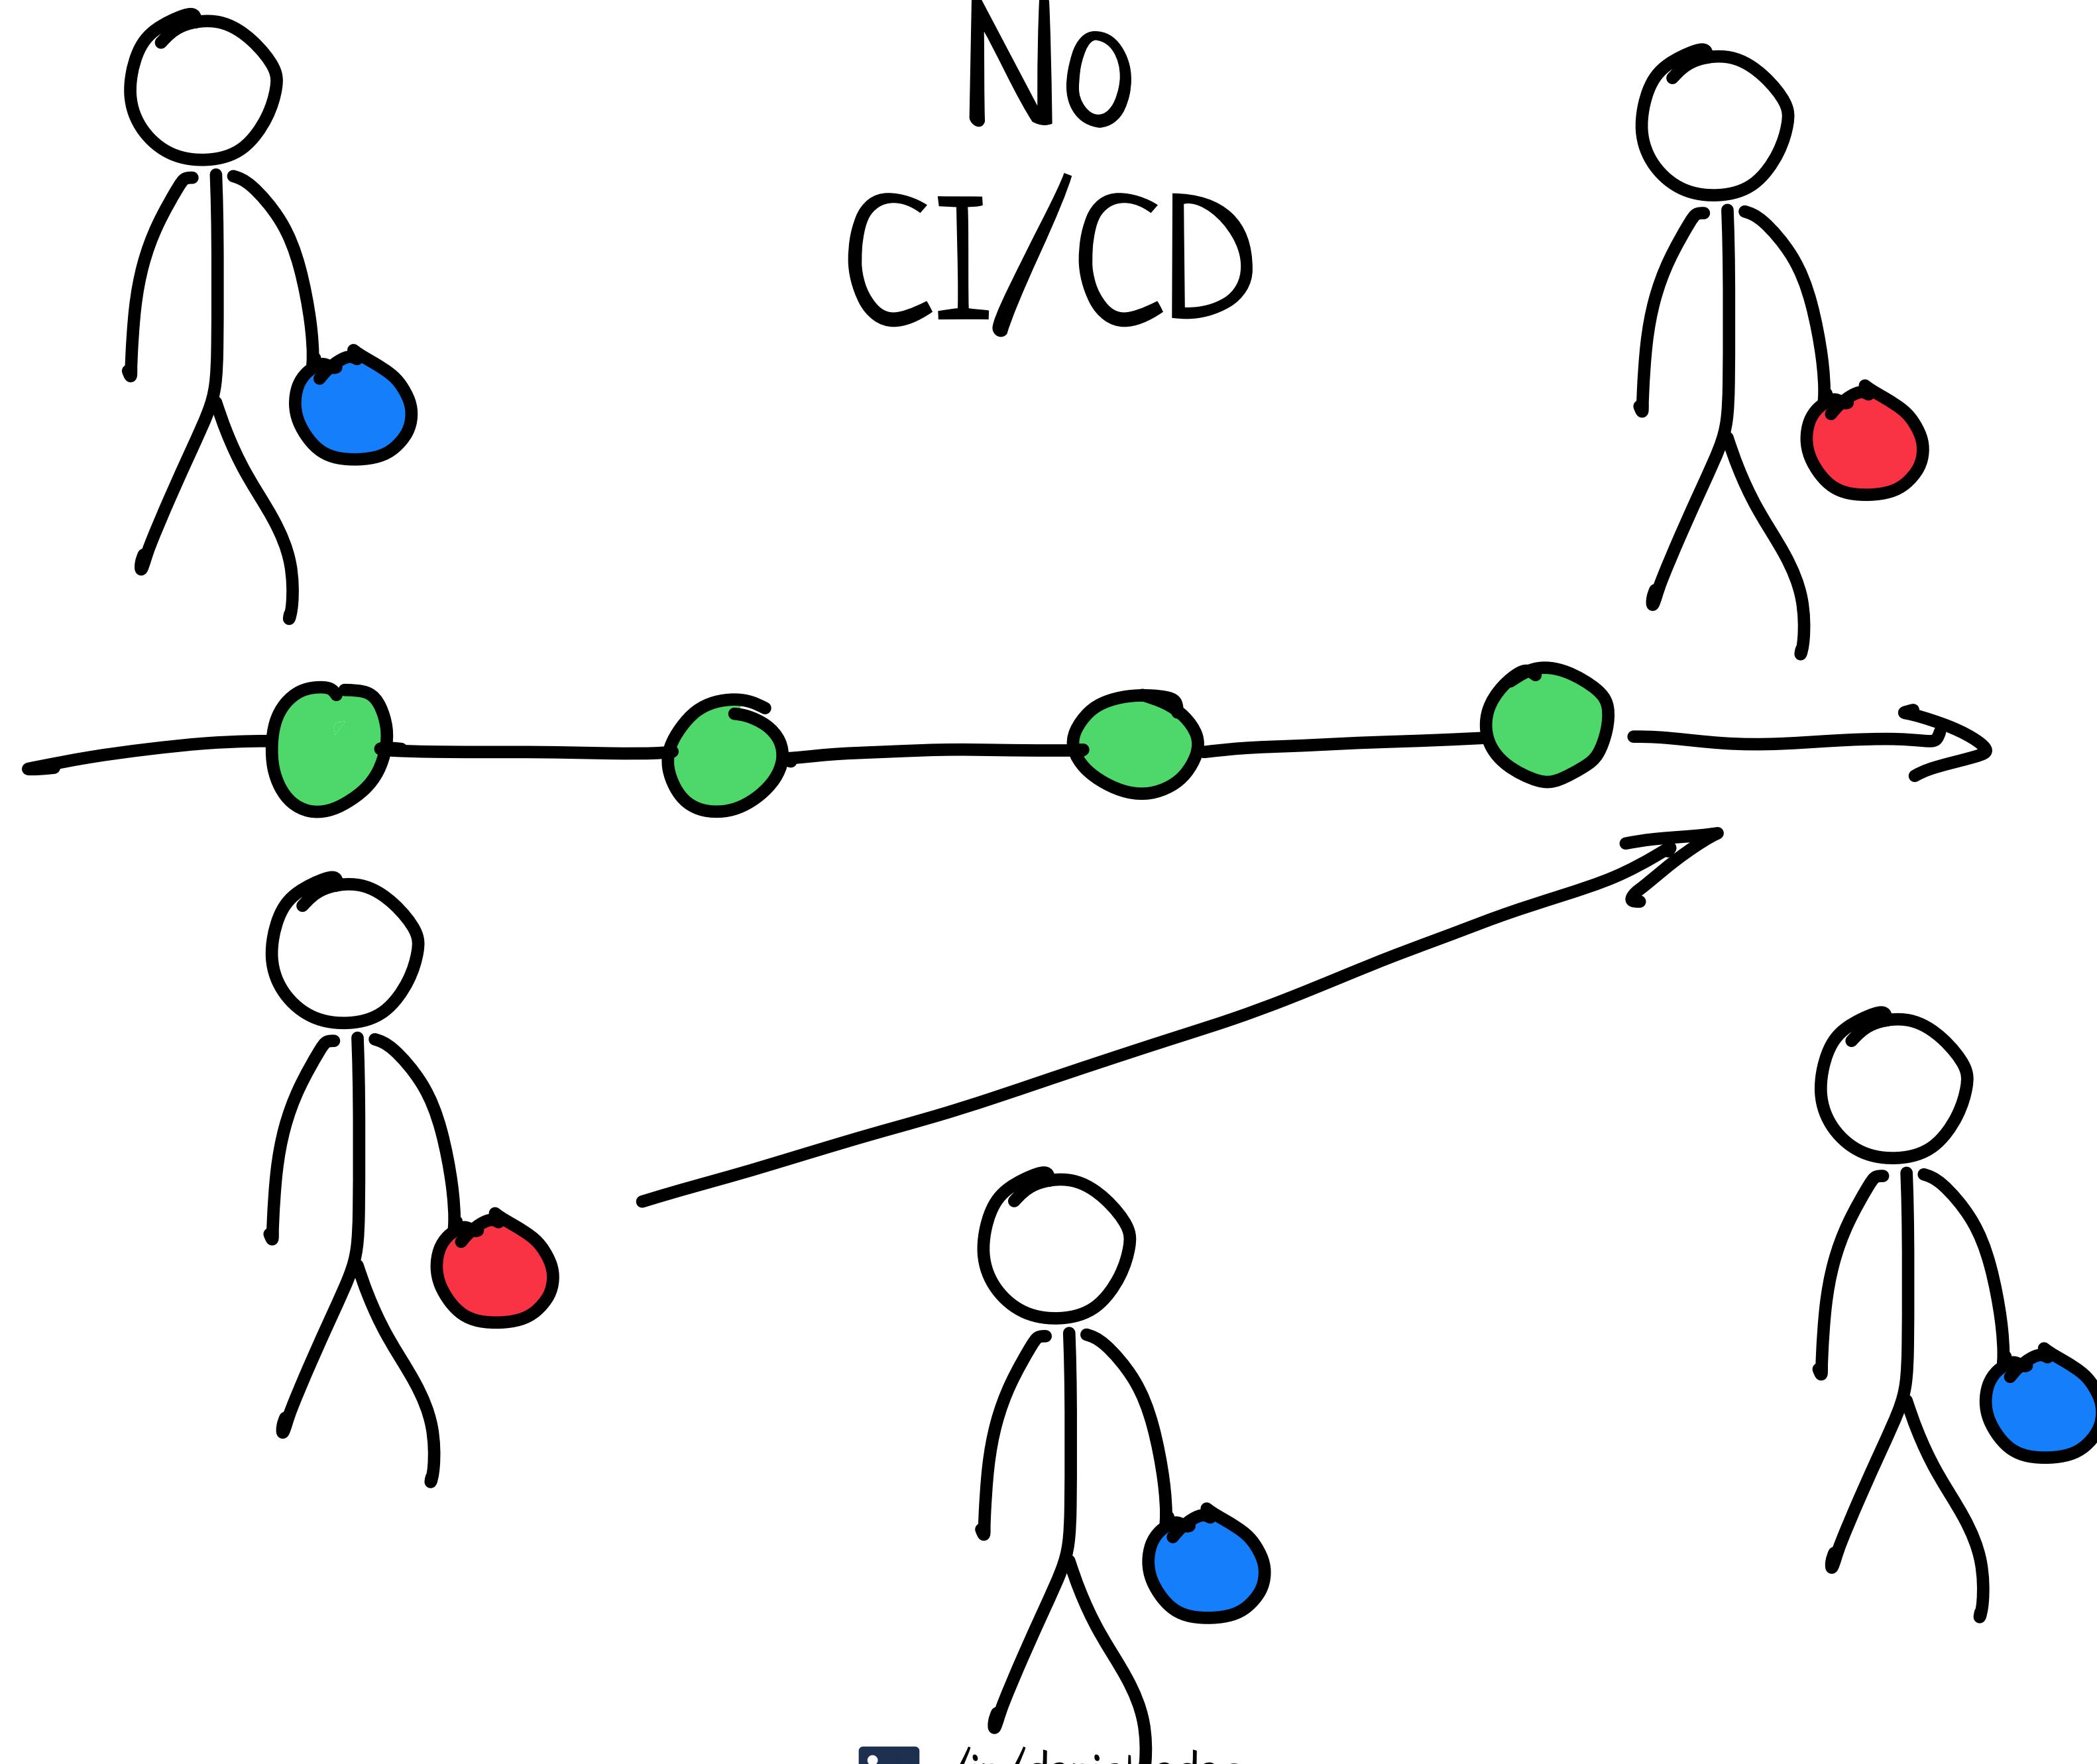
Infrastructure as Code

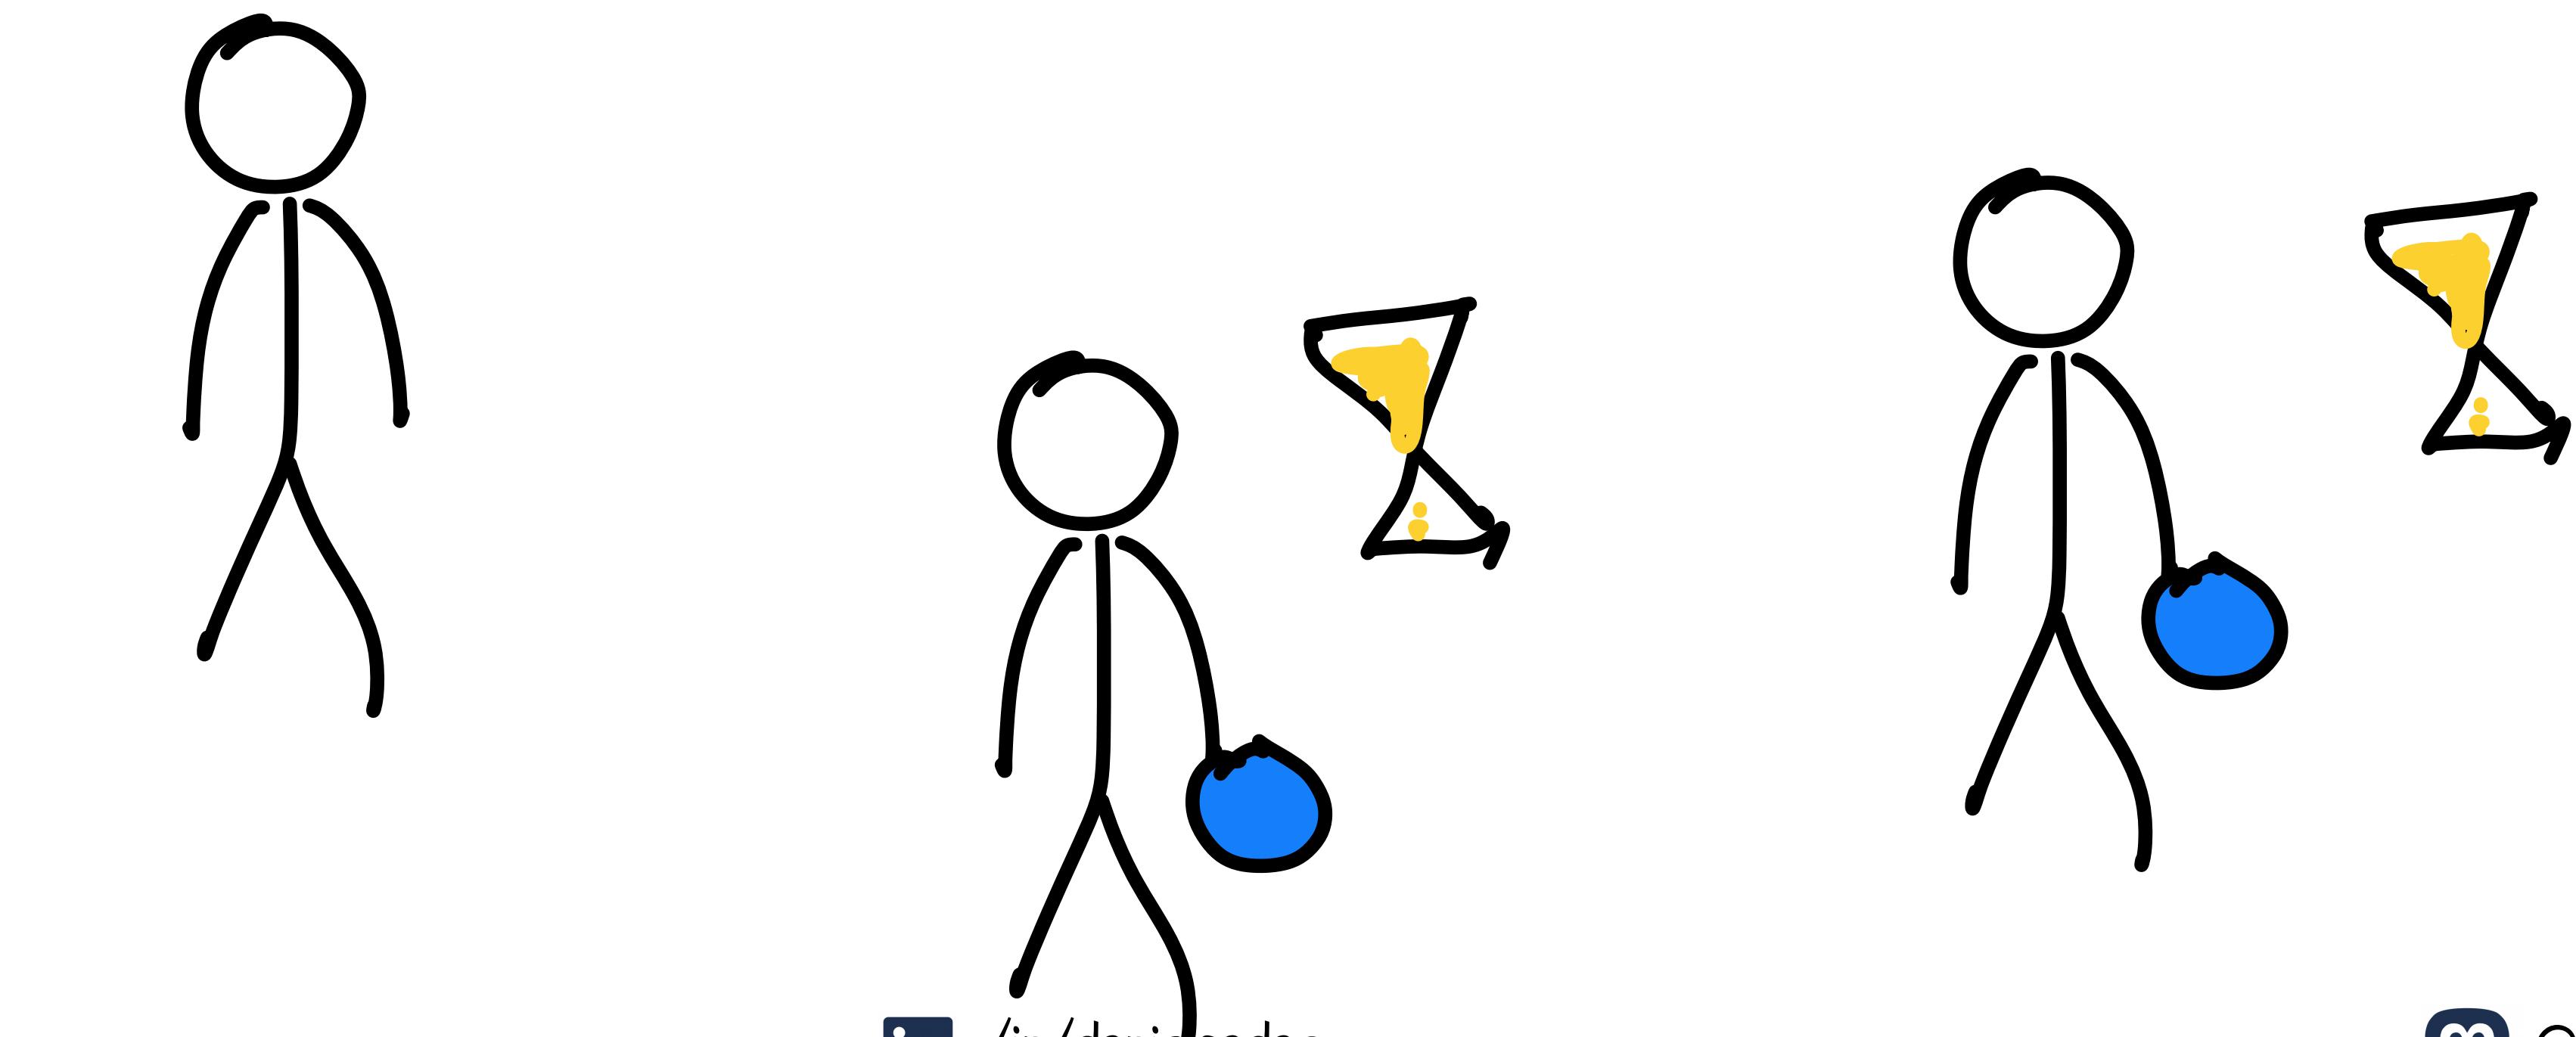
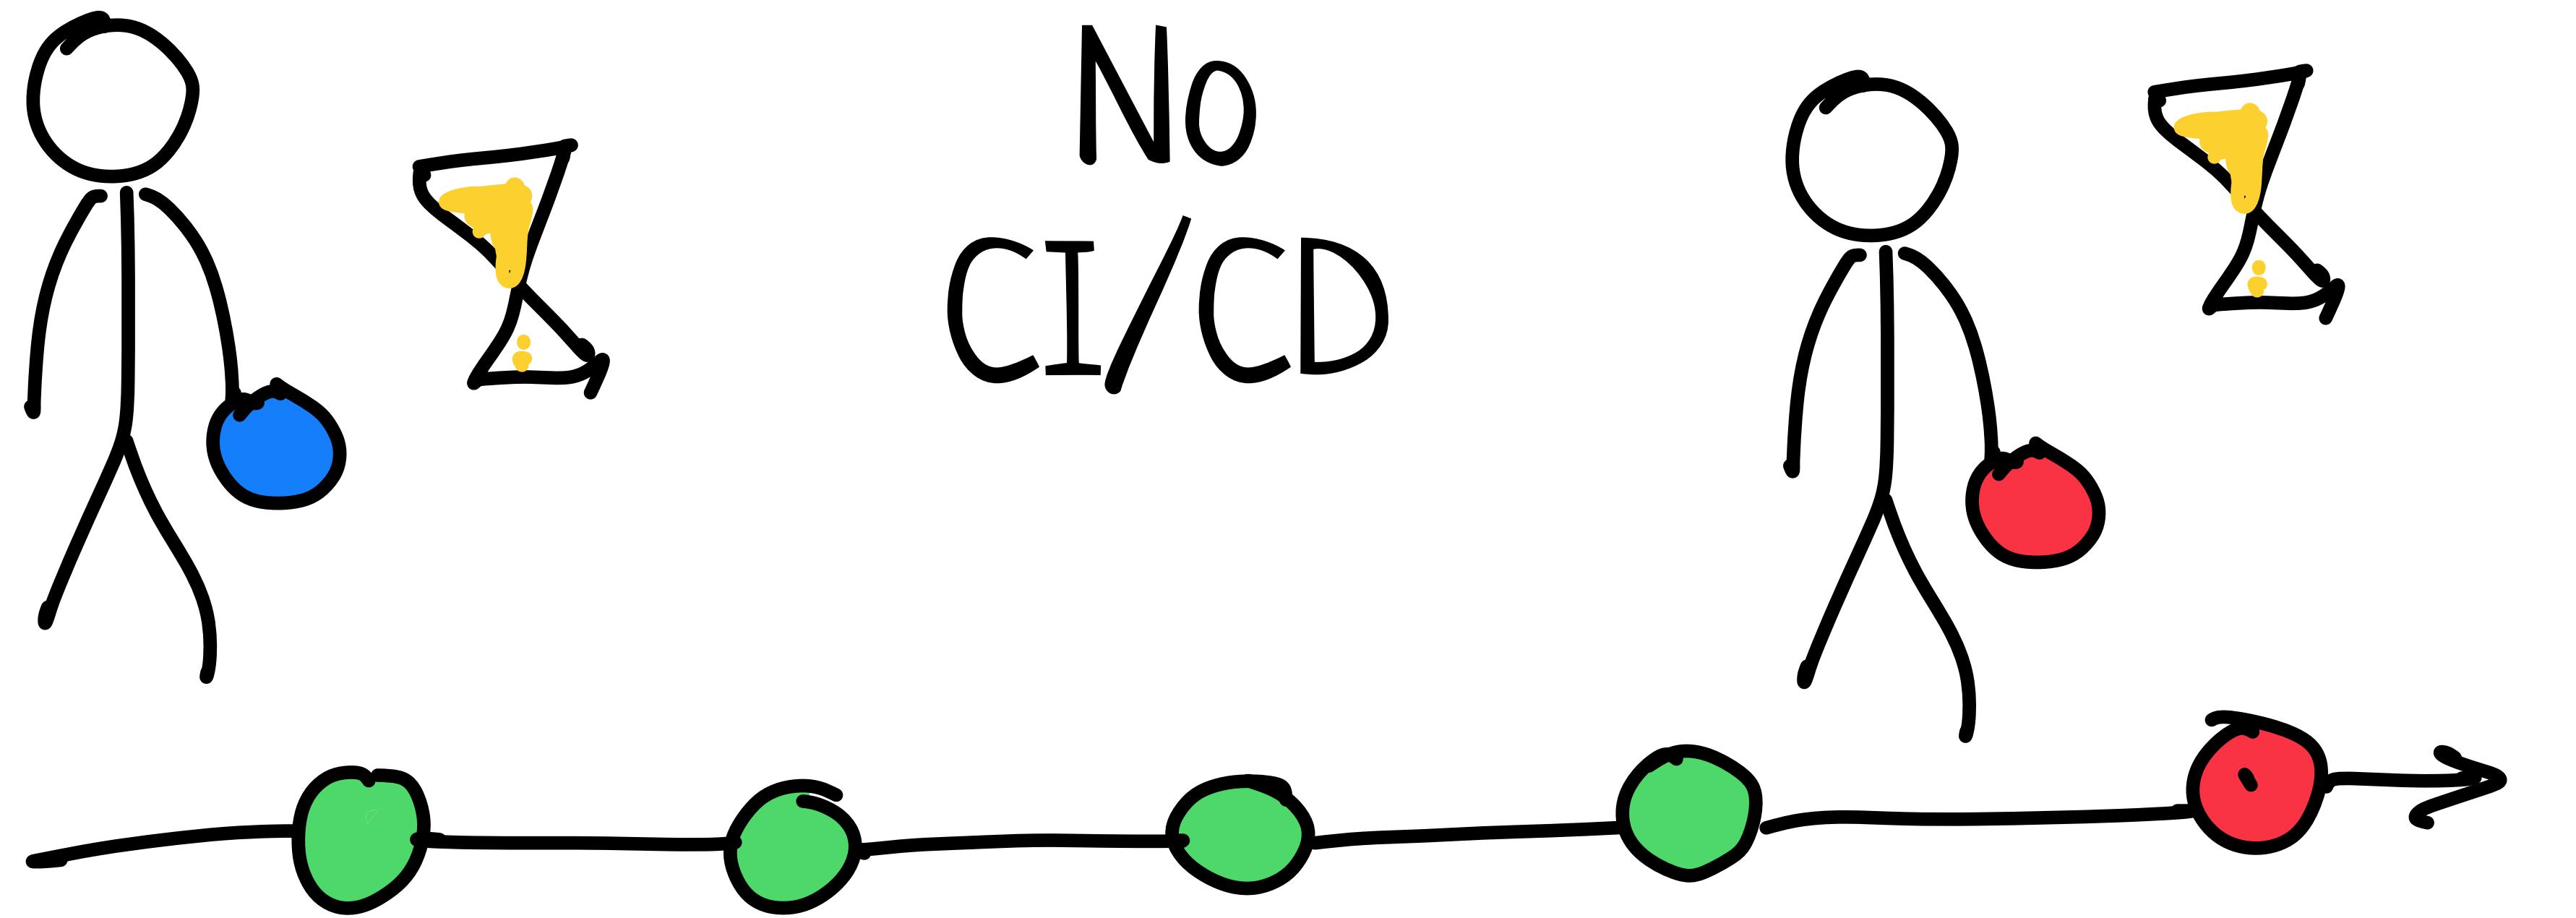
- Chef
- Puppet
- Ansible
- Terraform
- Cloud-specific lang

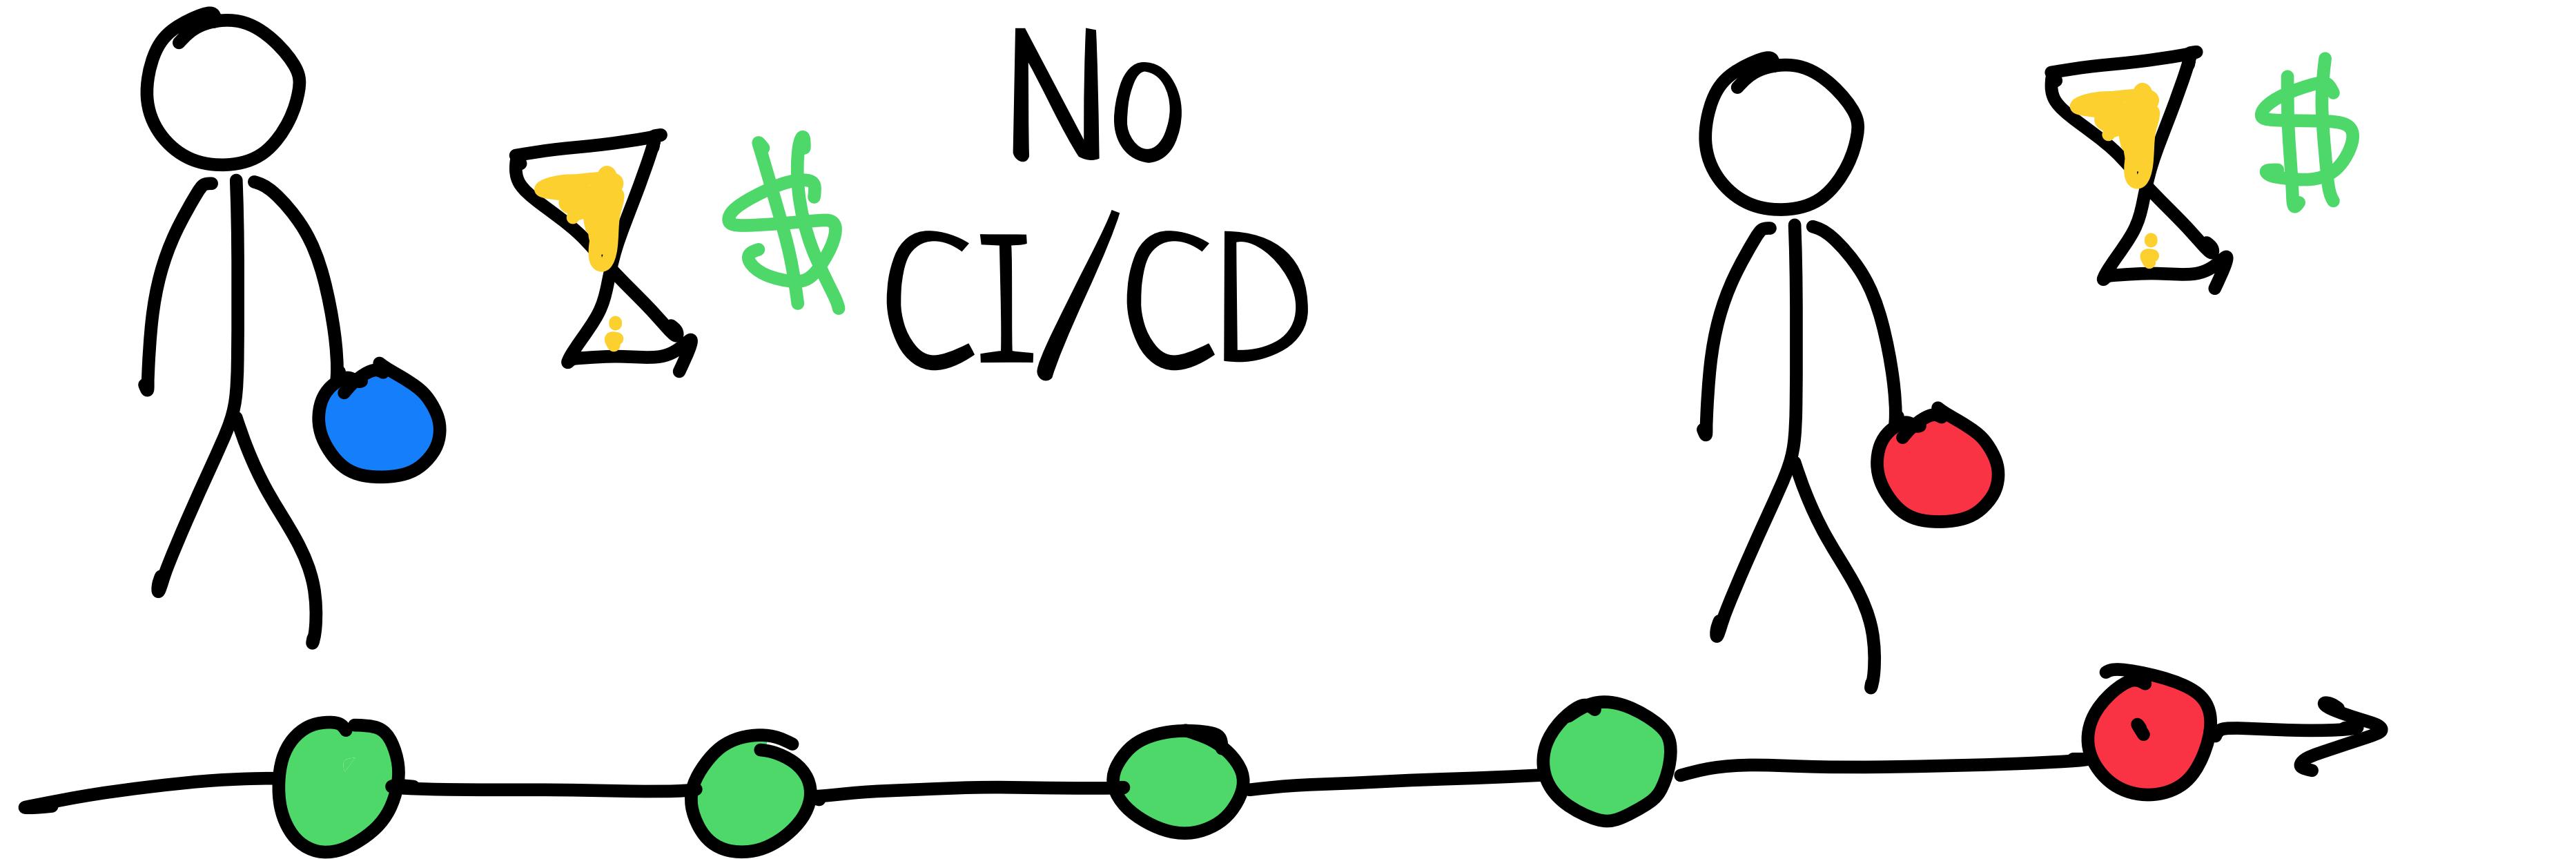
CI/CD

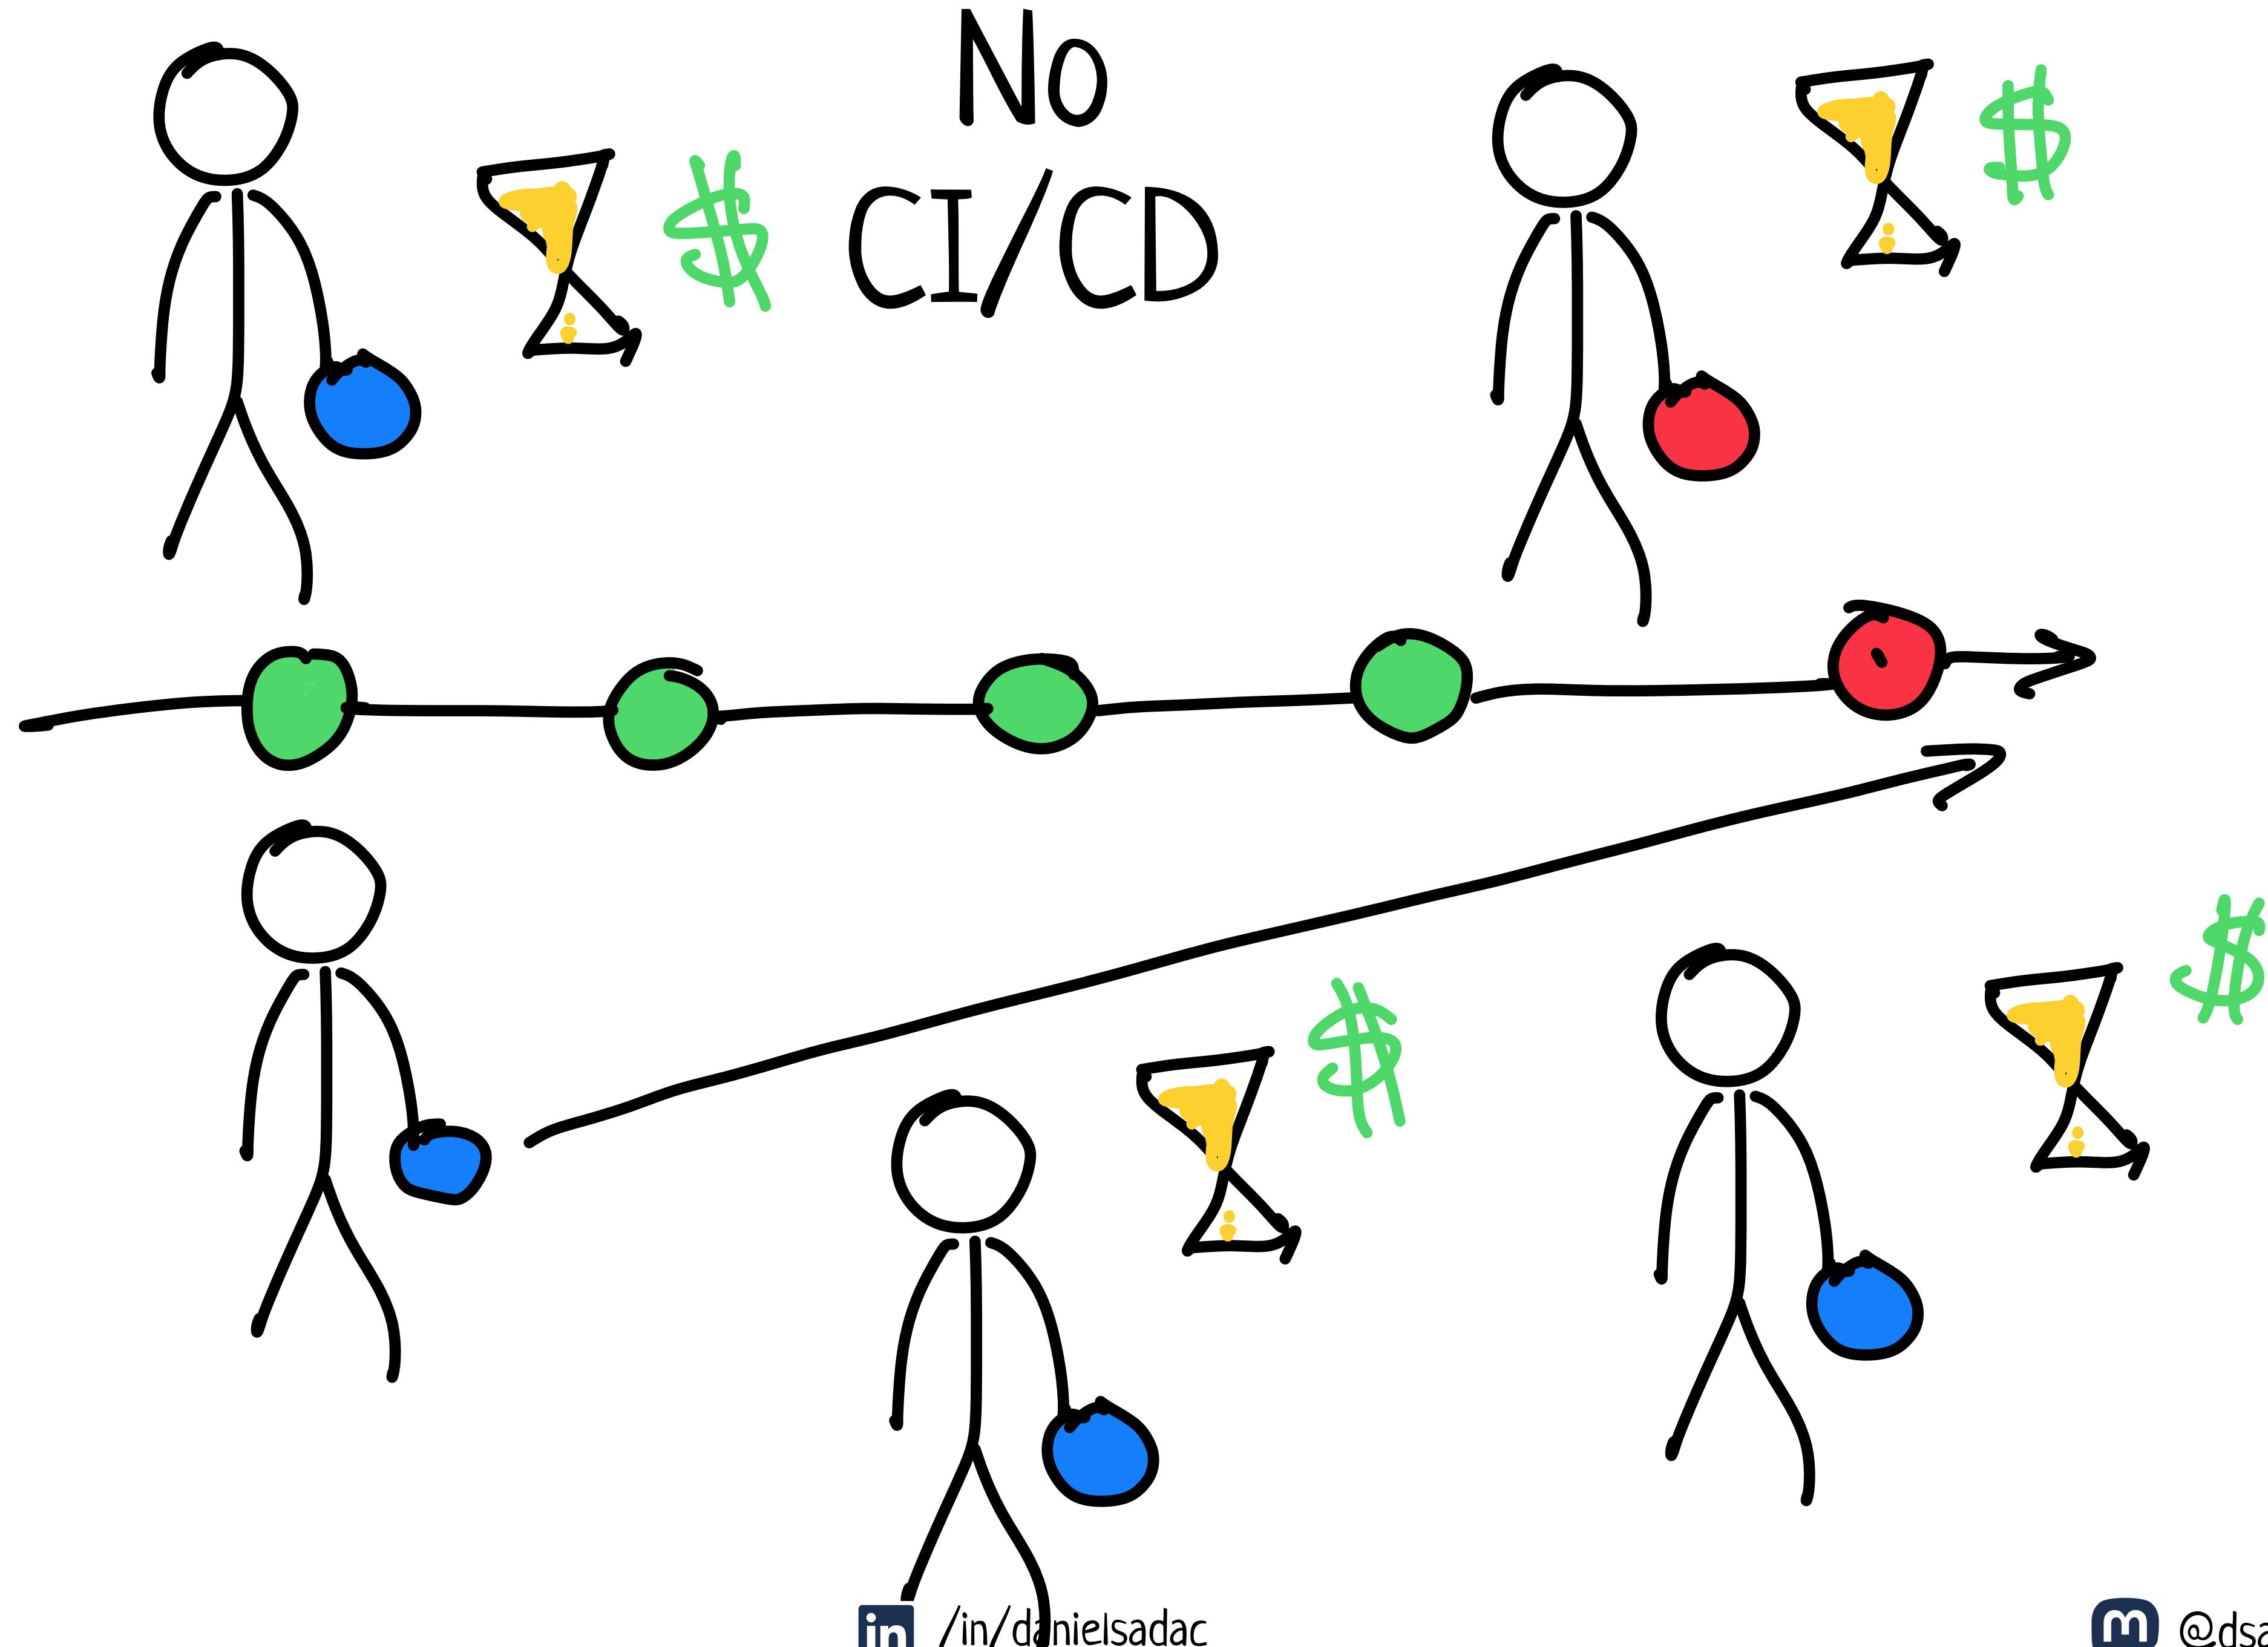
Continuous Integration / Continuous Deployment

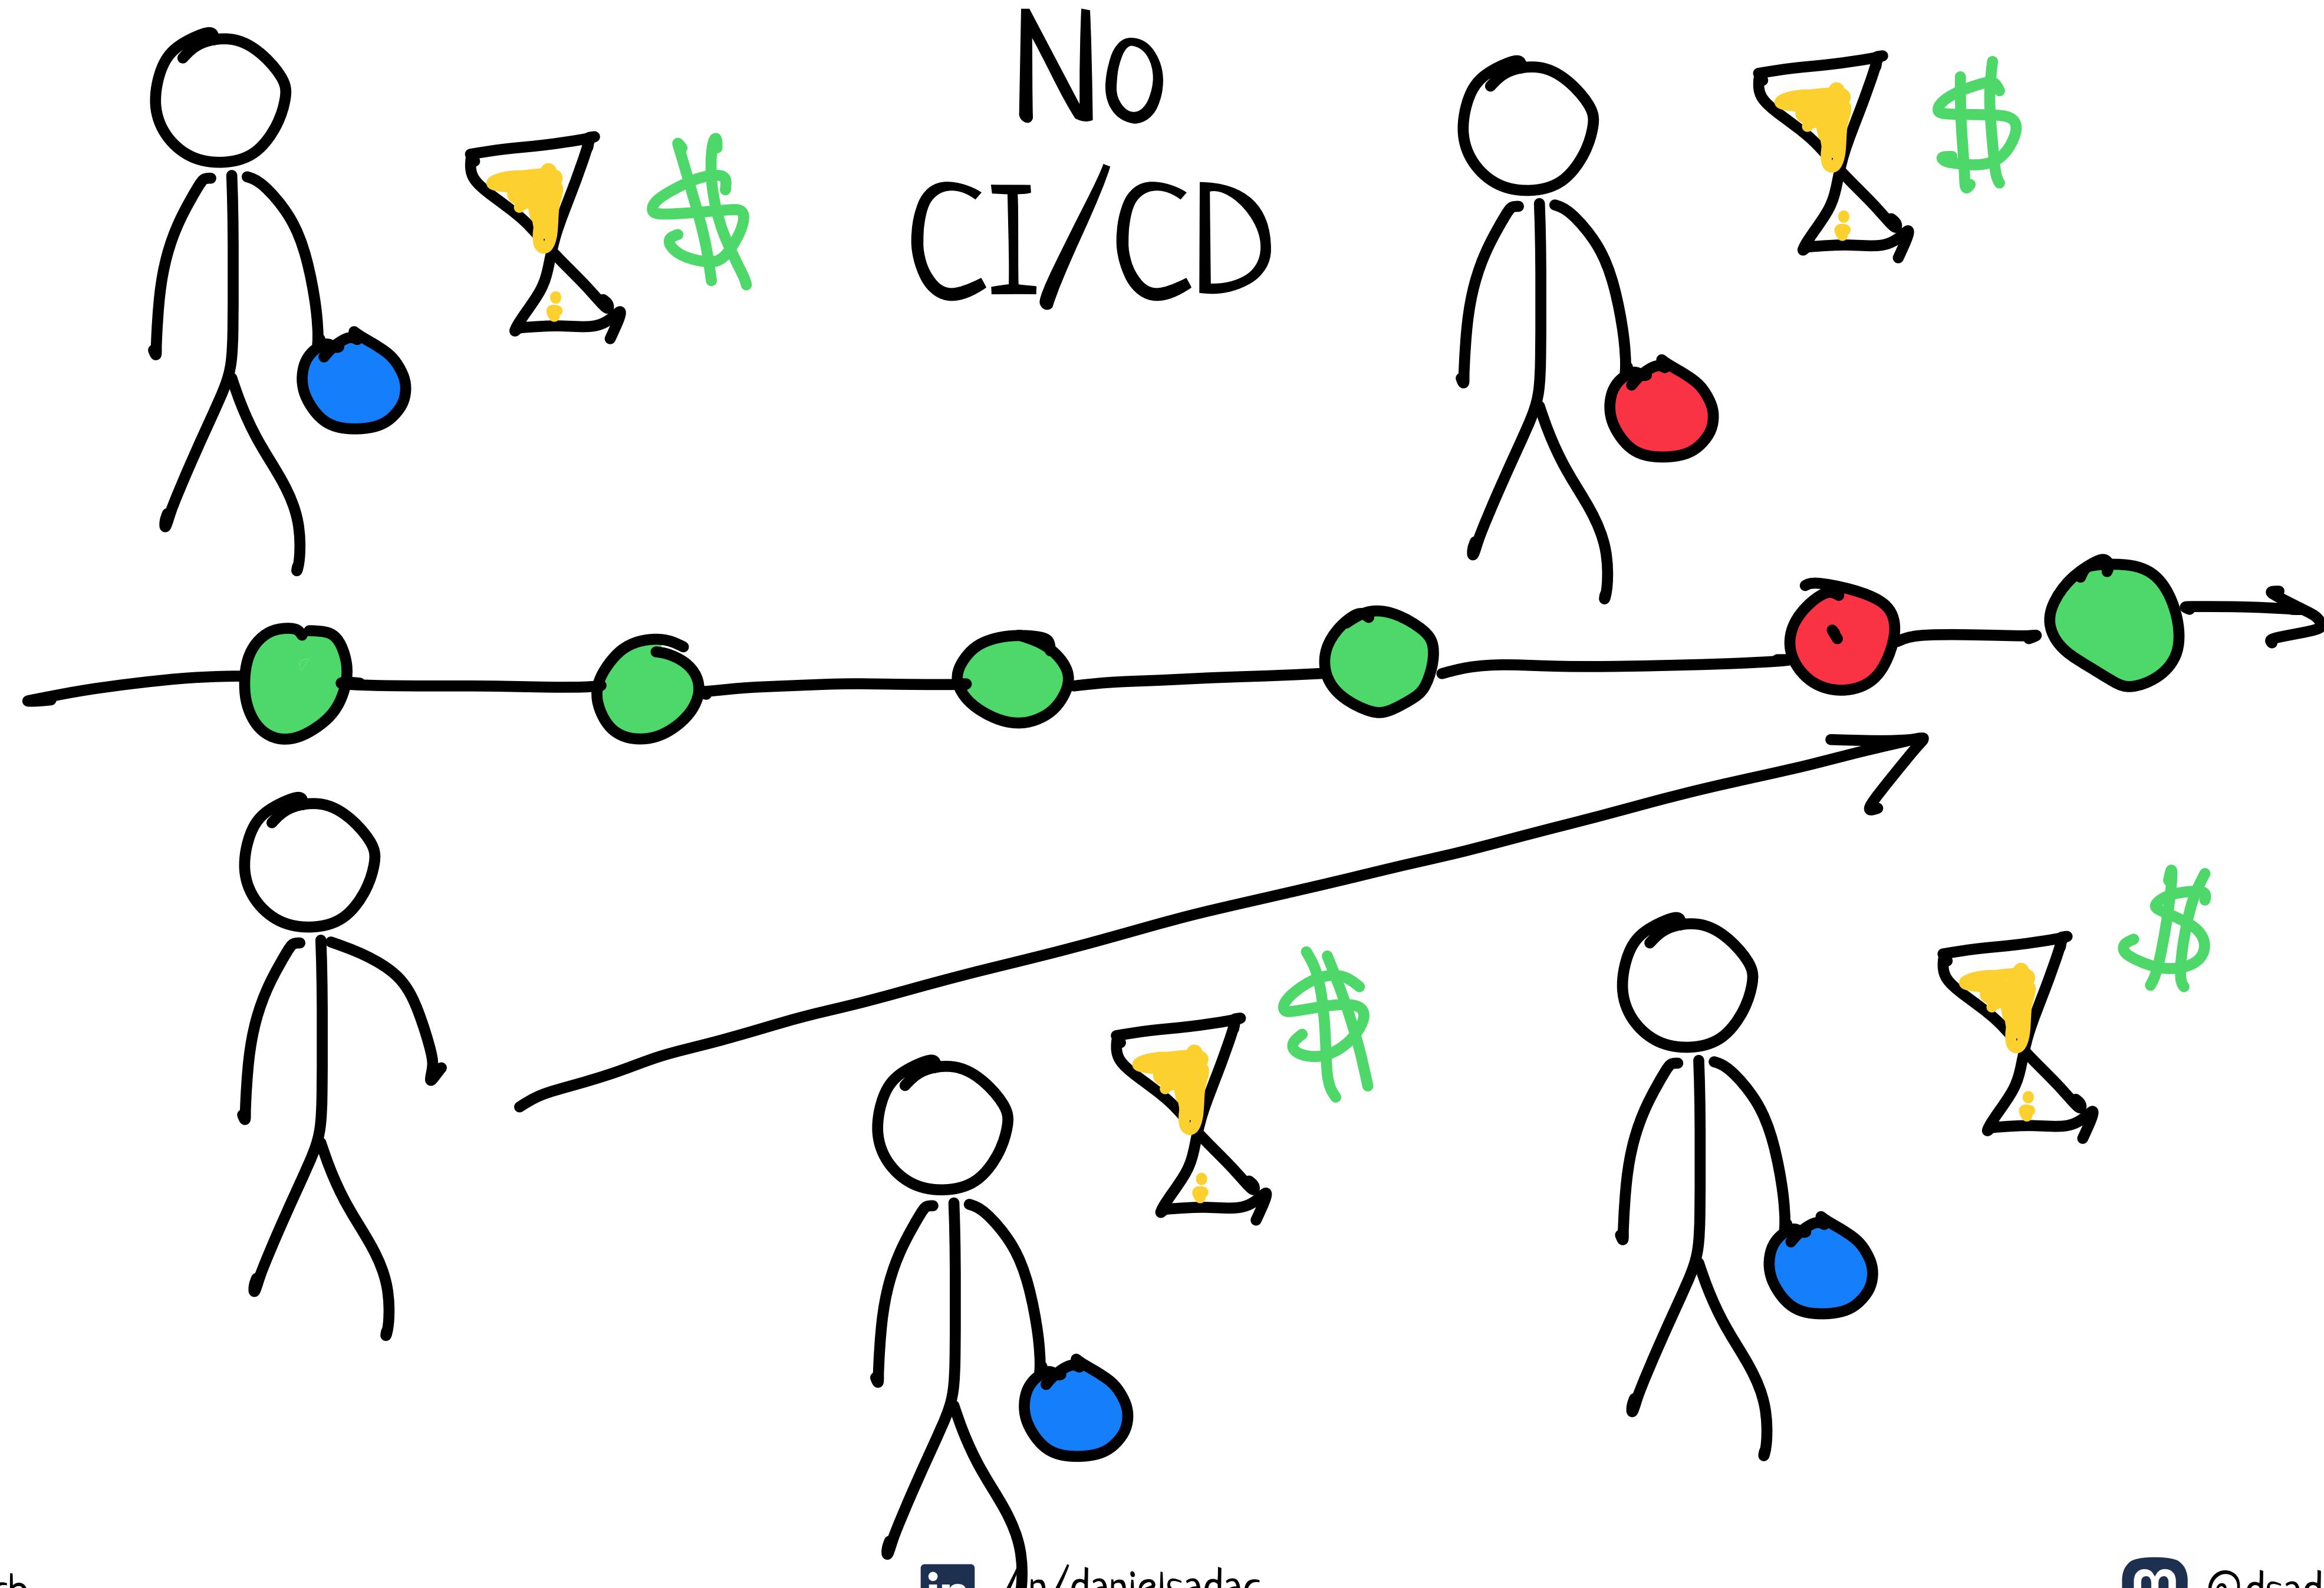
No CI/CD



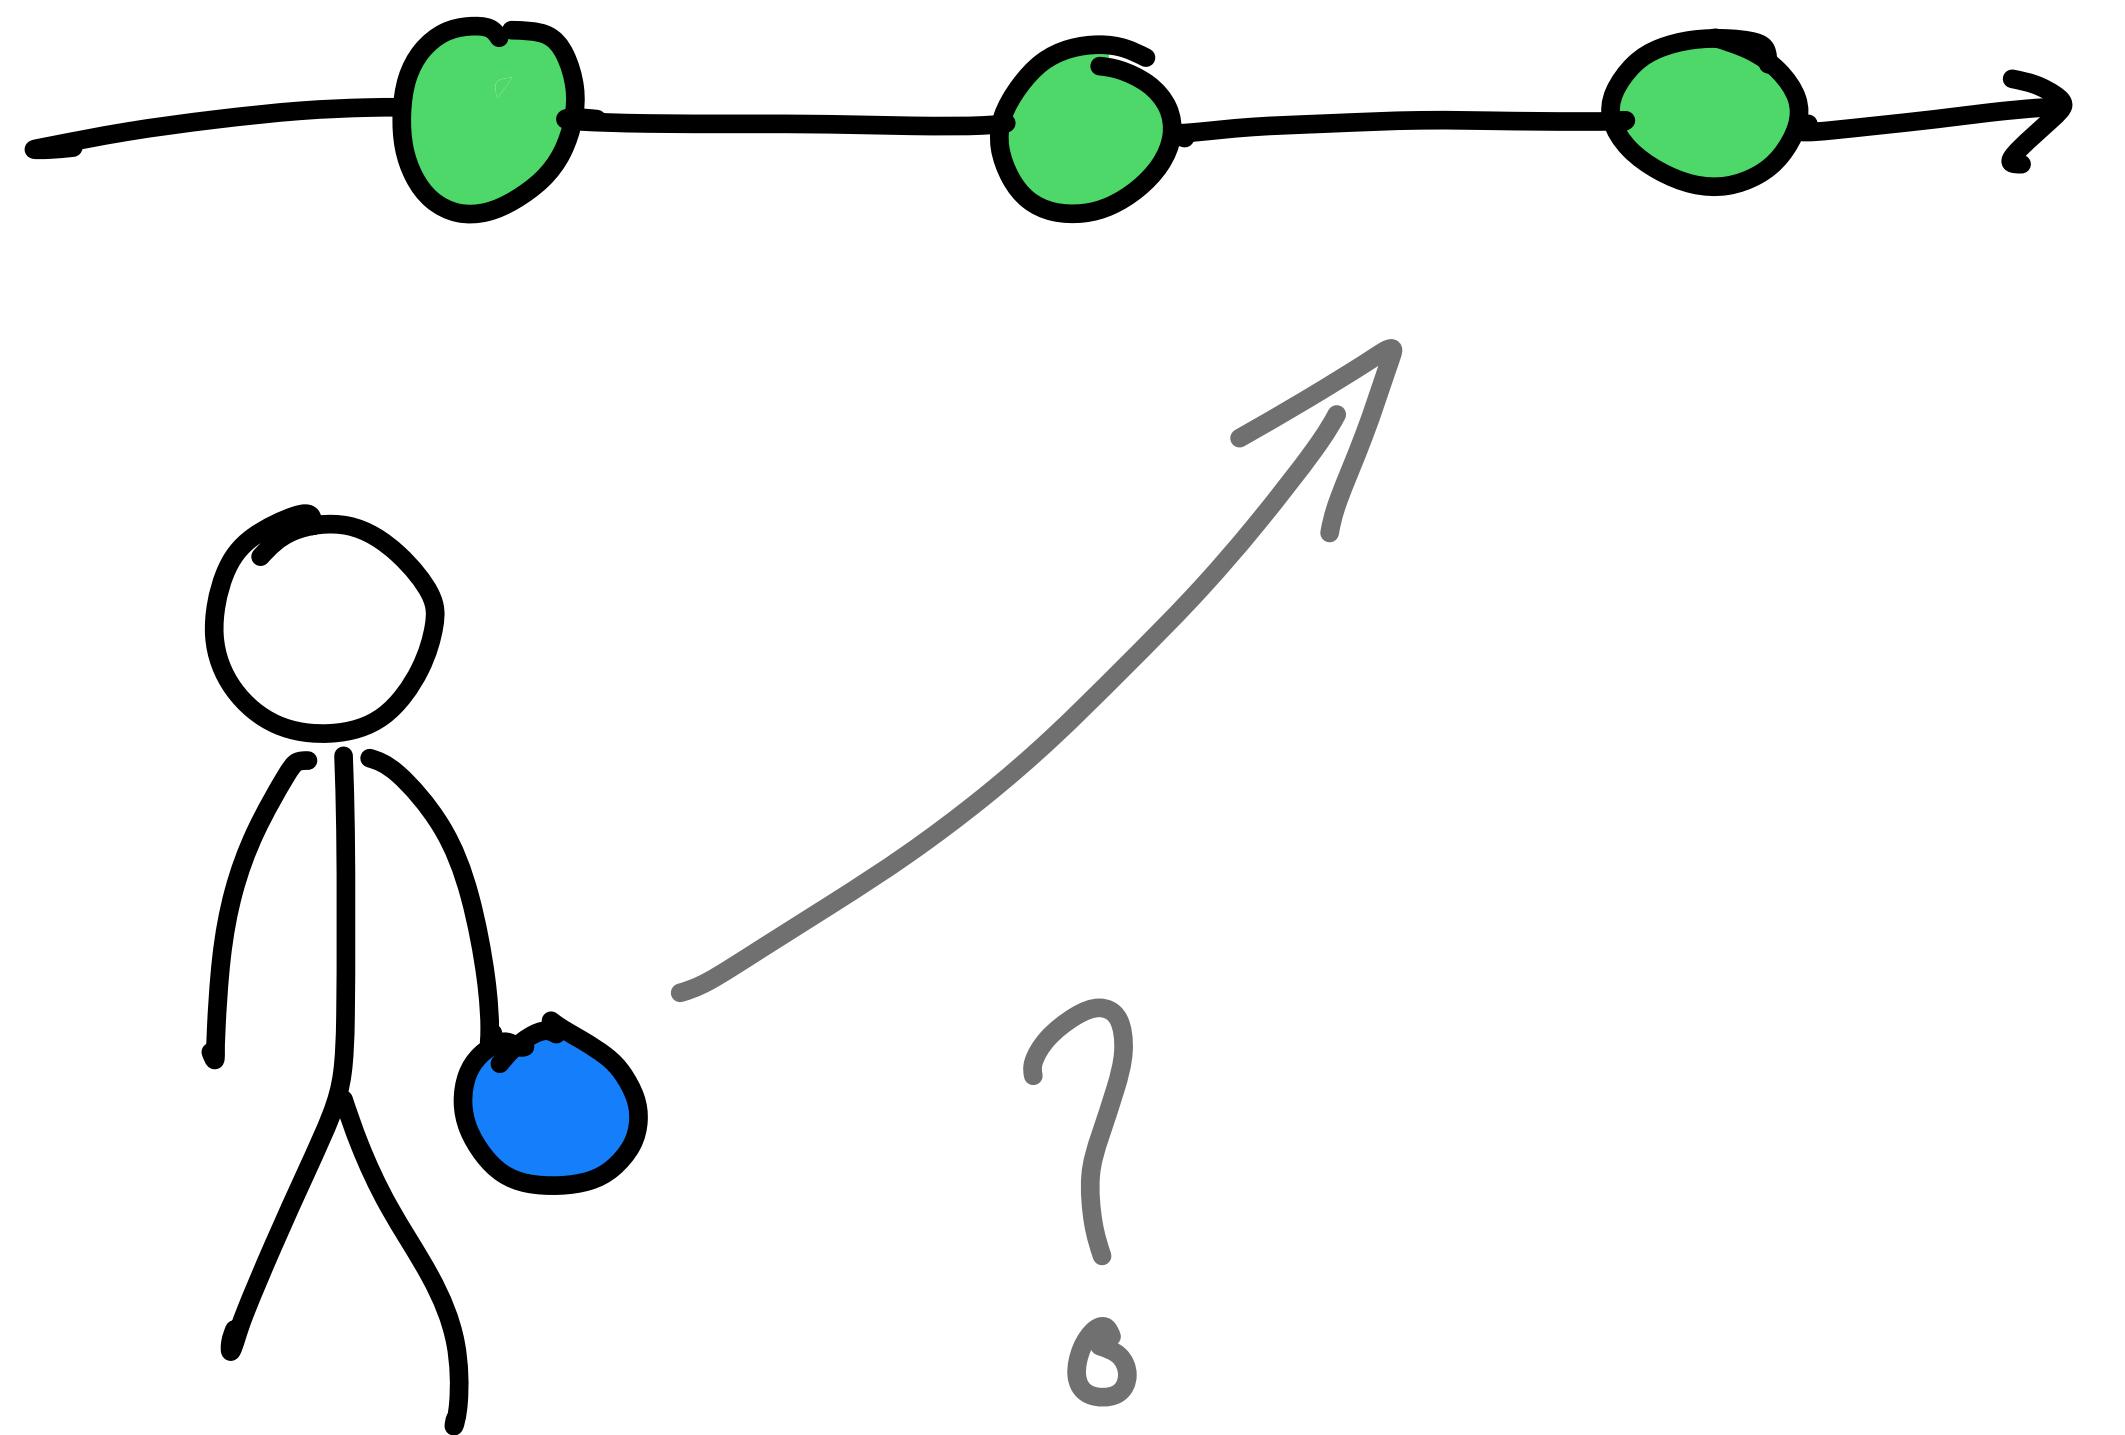




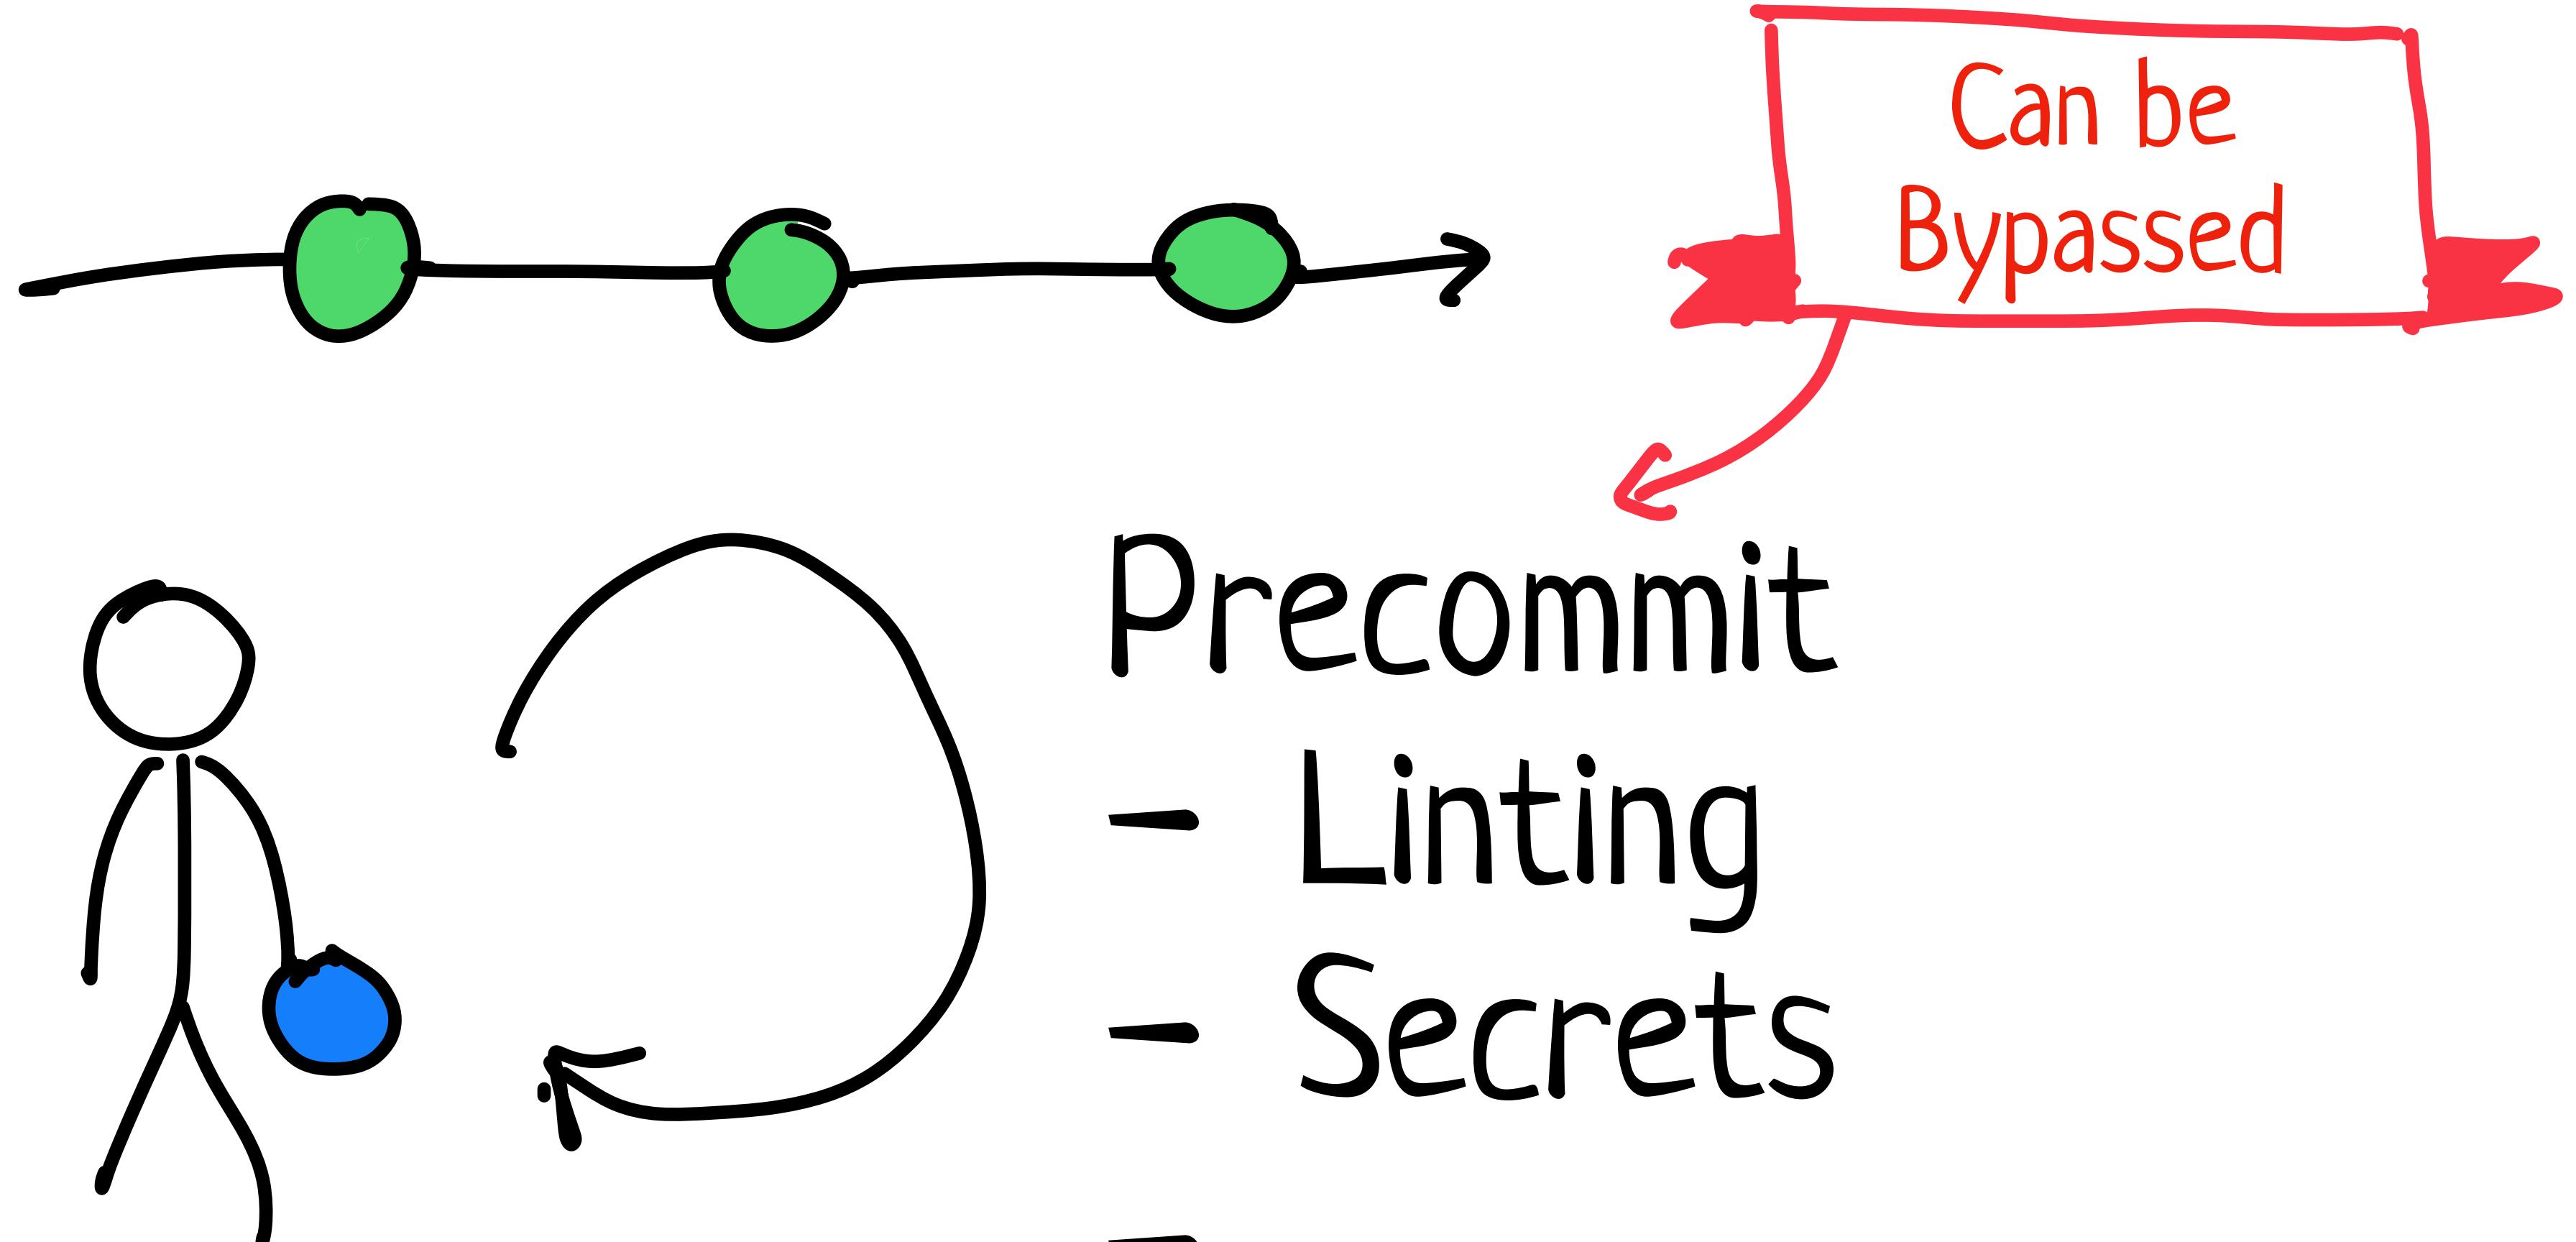




A developer wants to merge

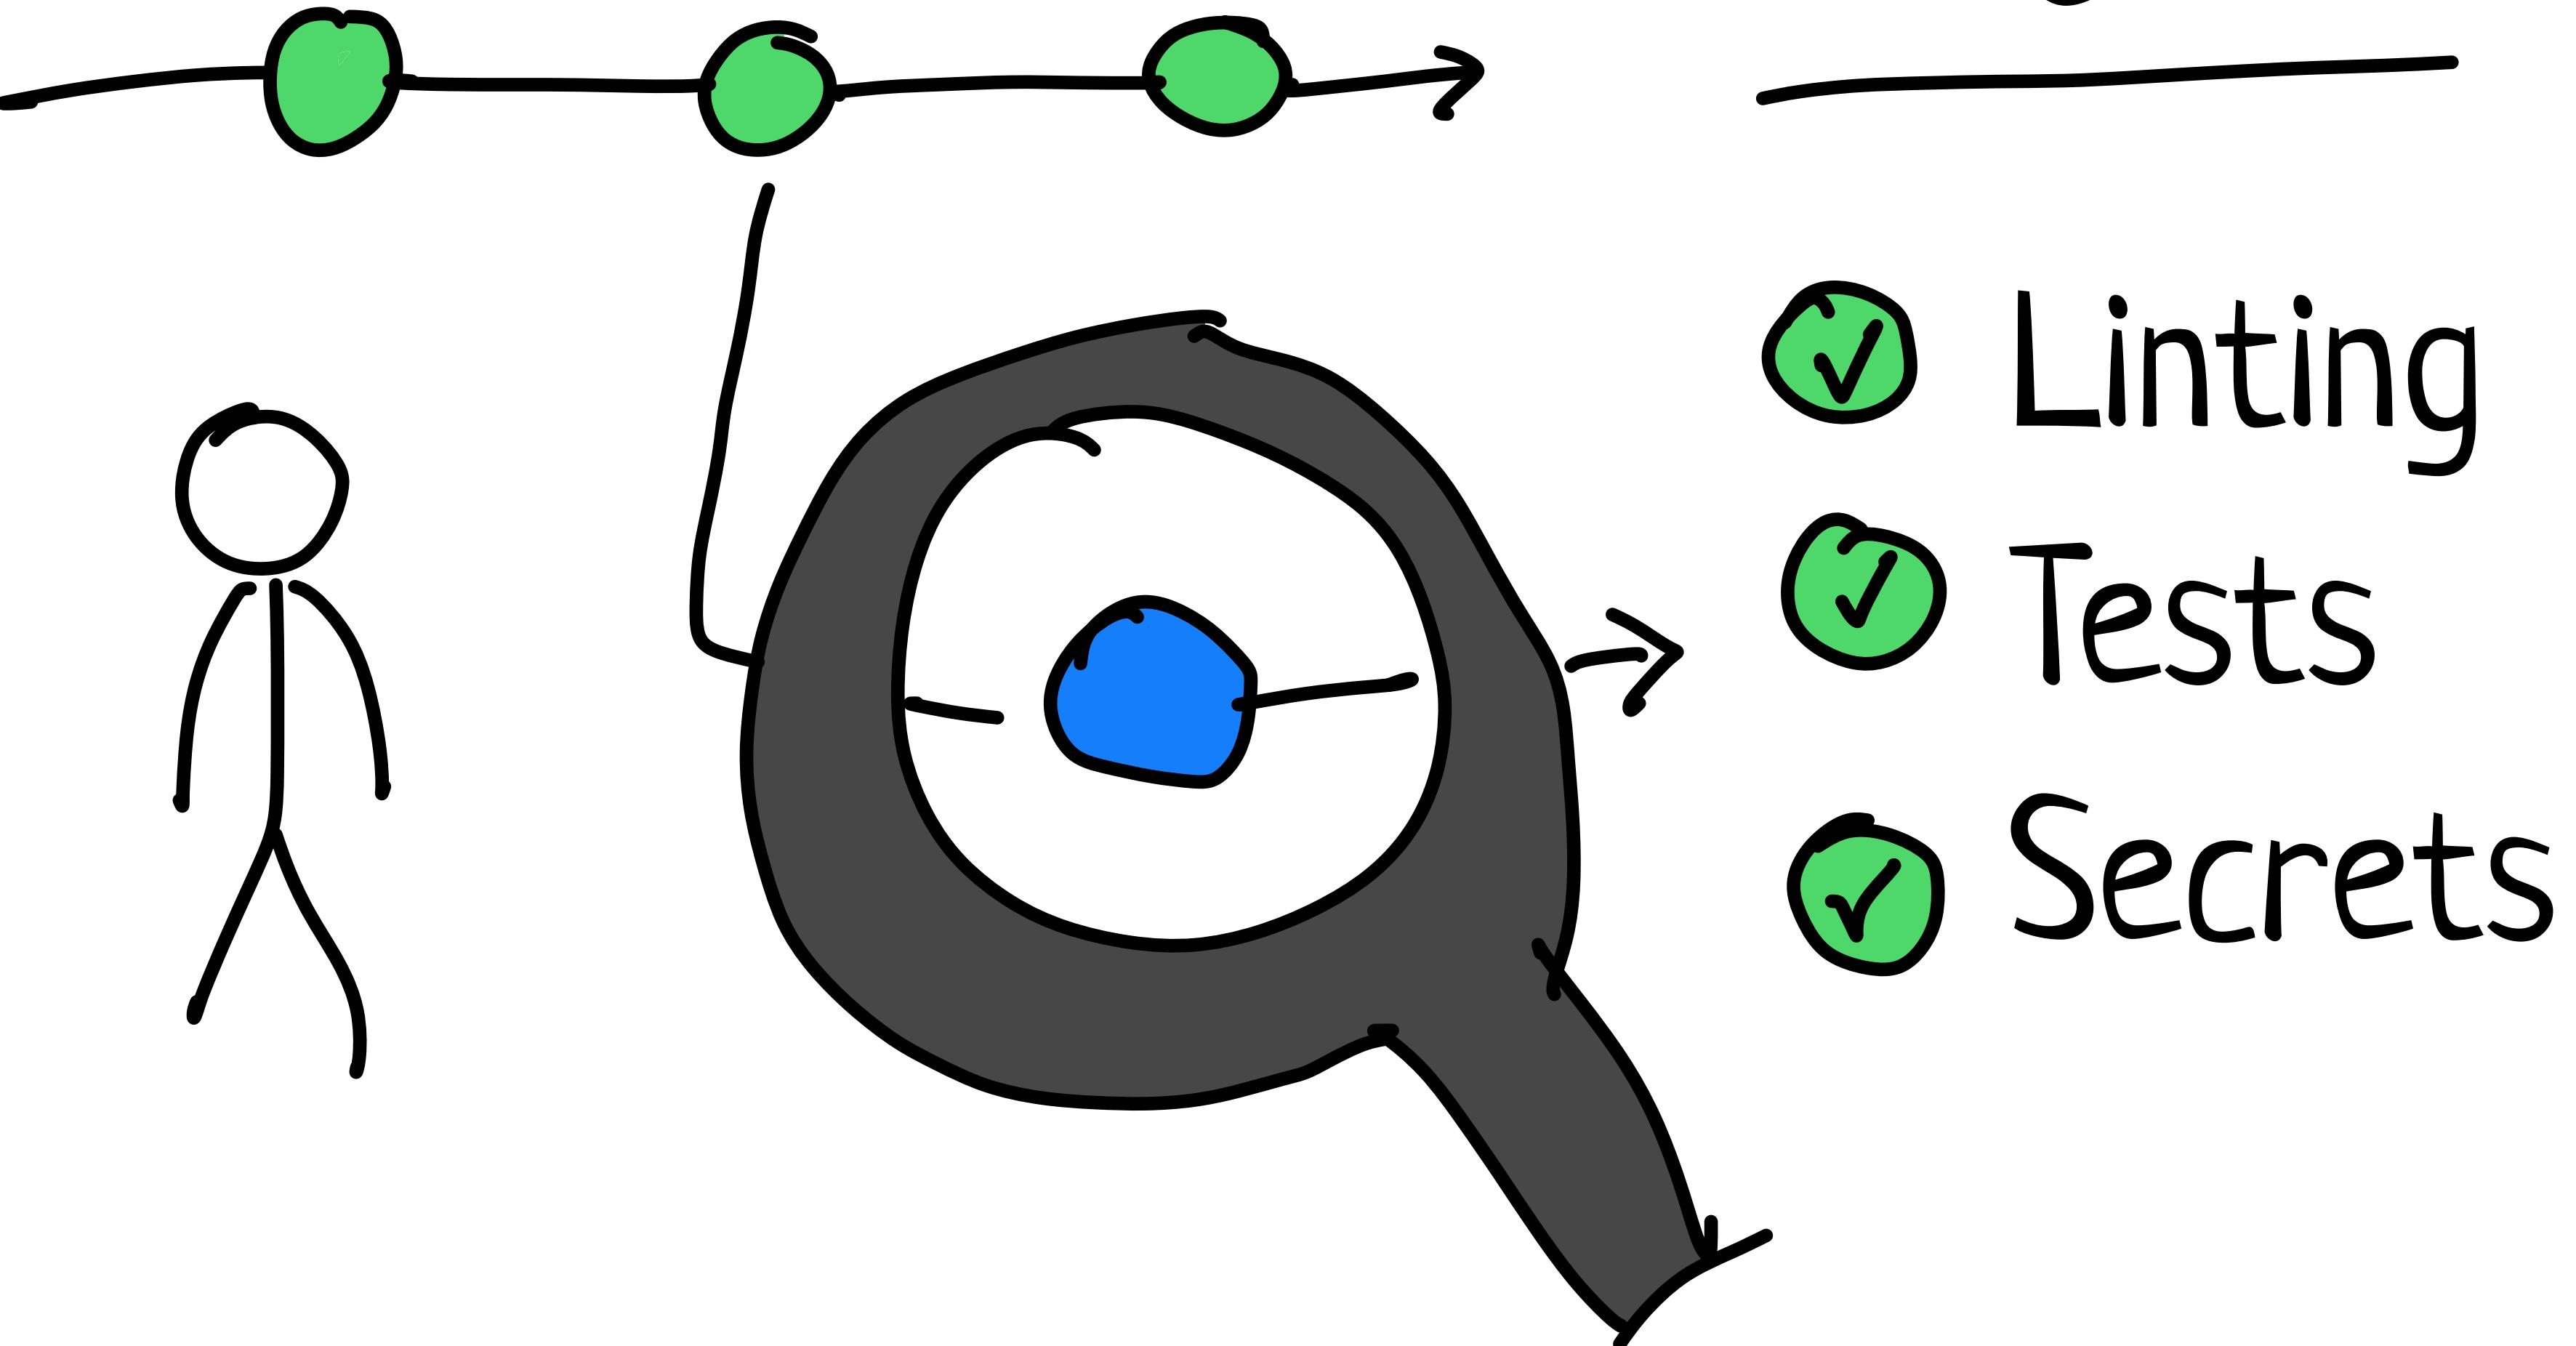


Before commit

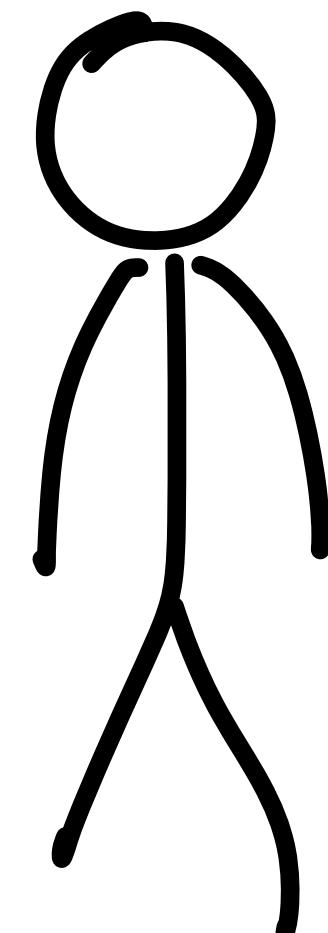
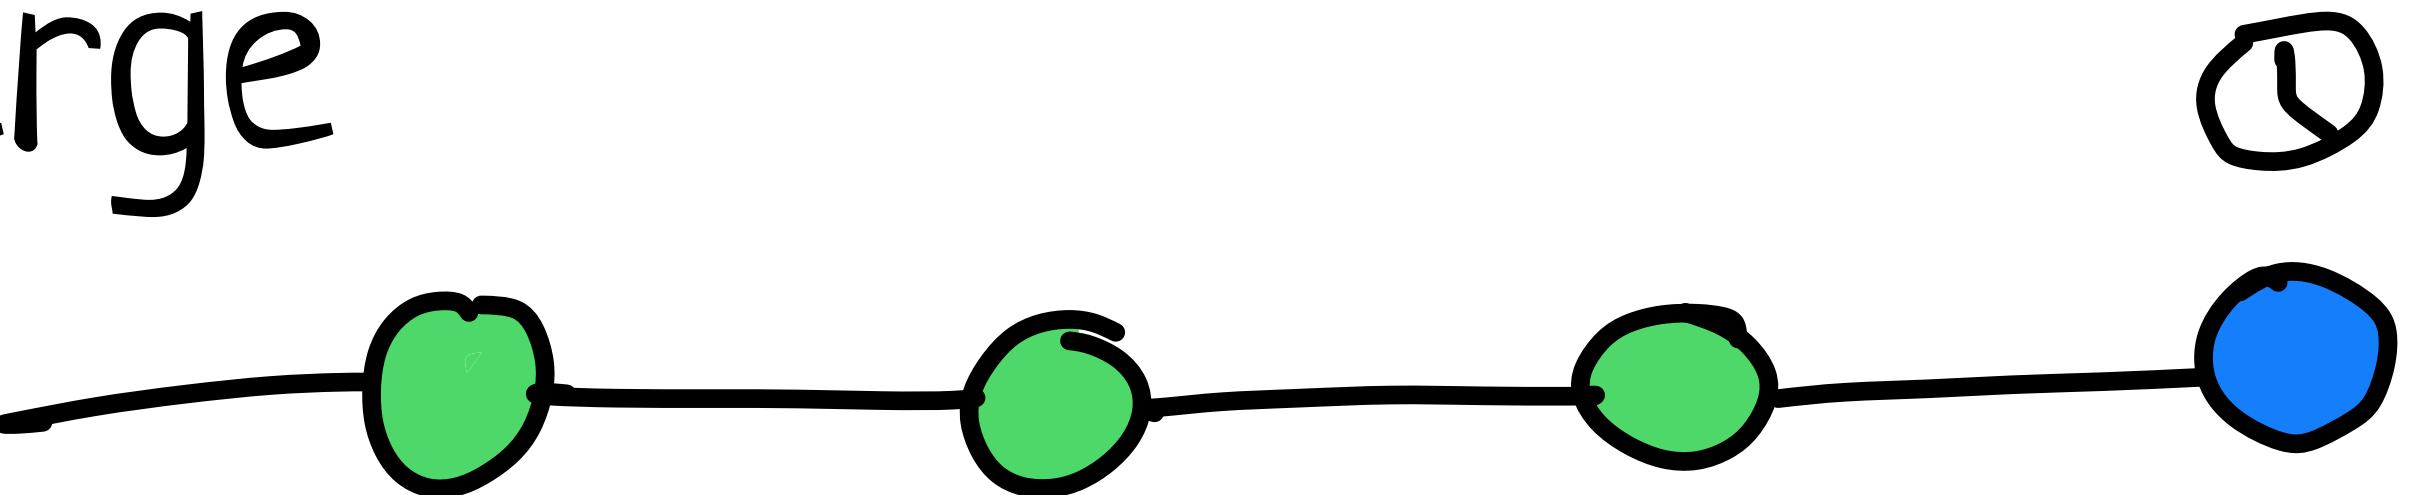


a.k.a git commit checks

While branch is active



After merge



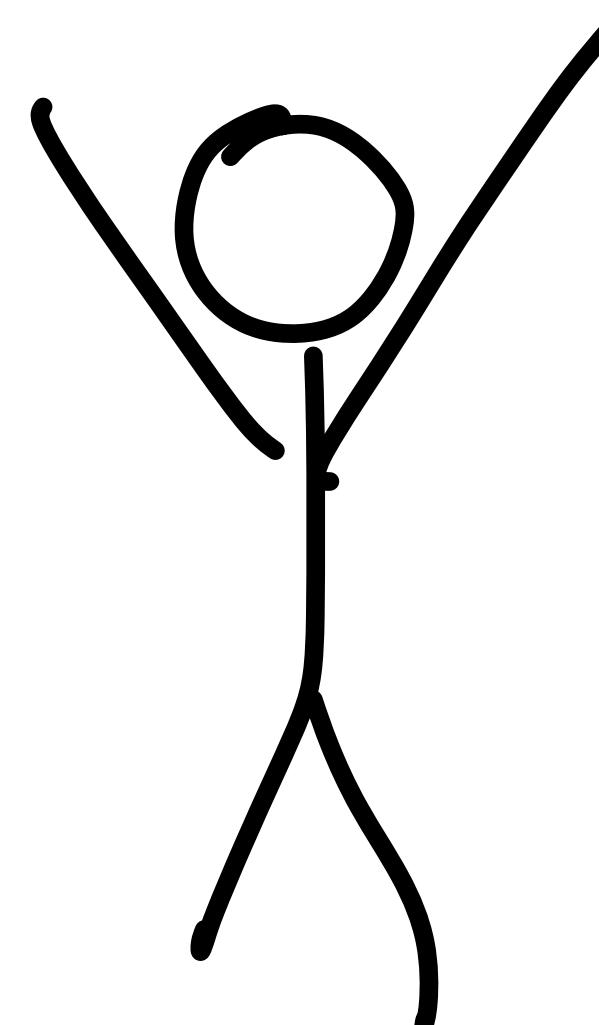
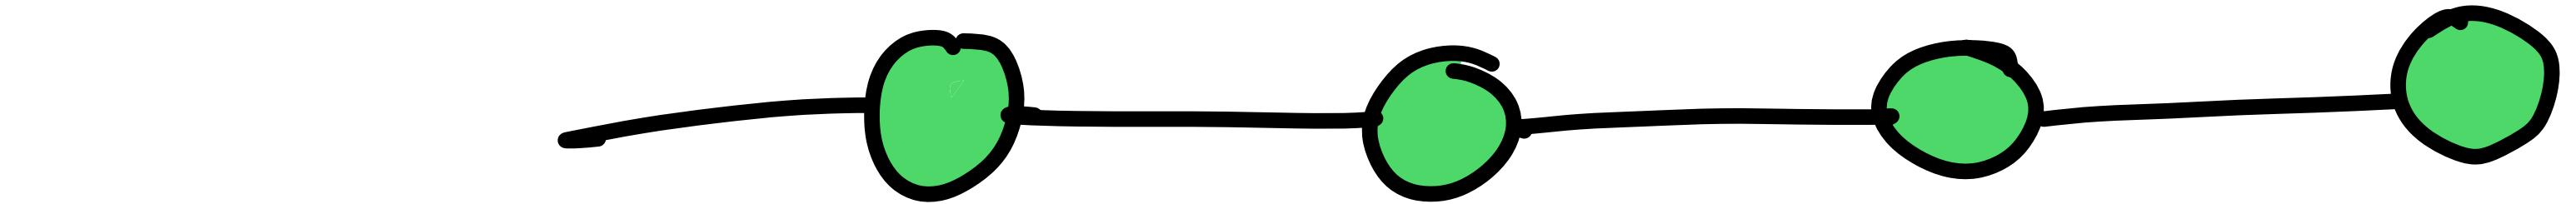
Can vary
after merge!



Build



Tests



Precommit

- Linting
- Secrets
- ...

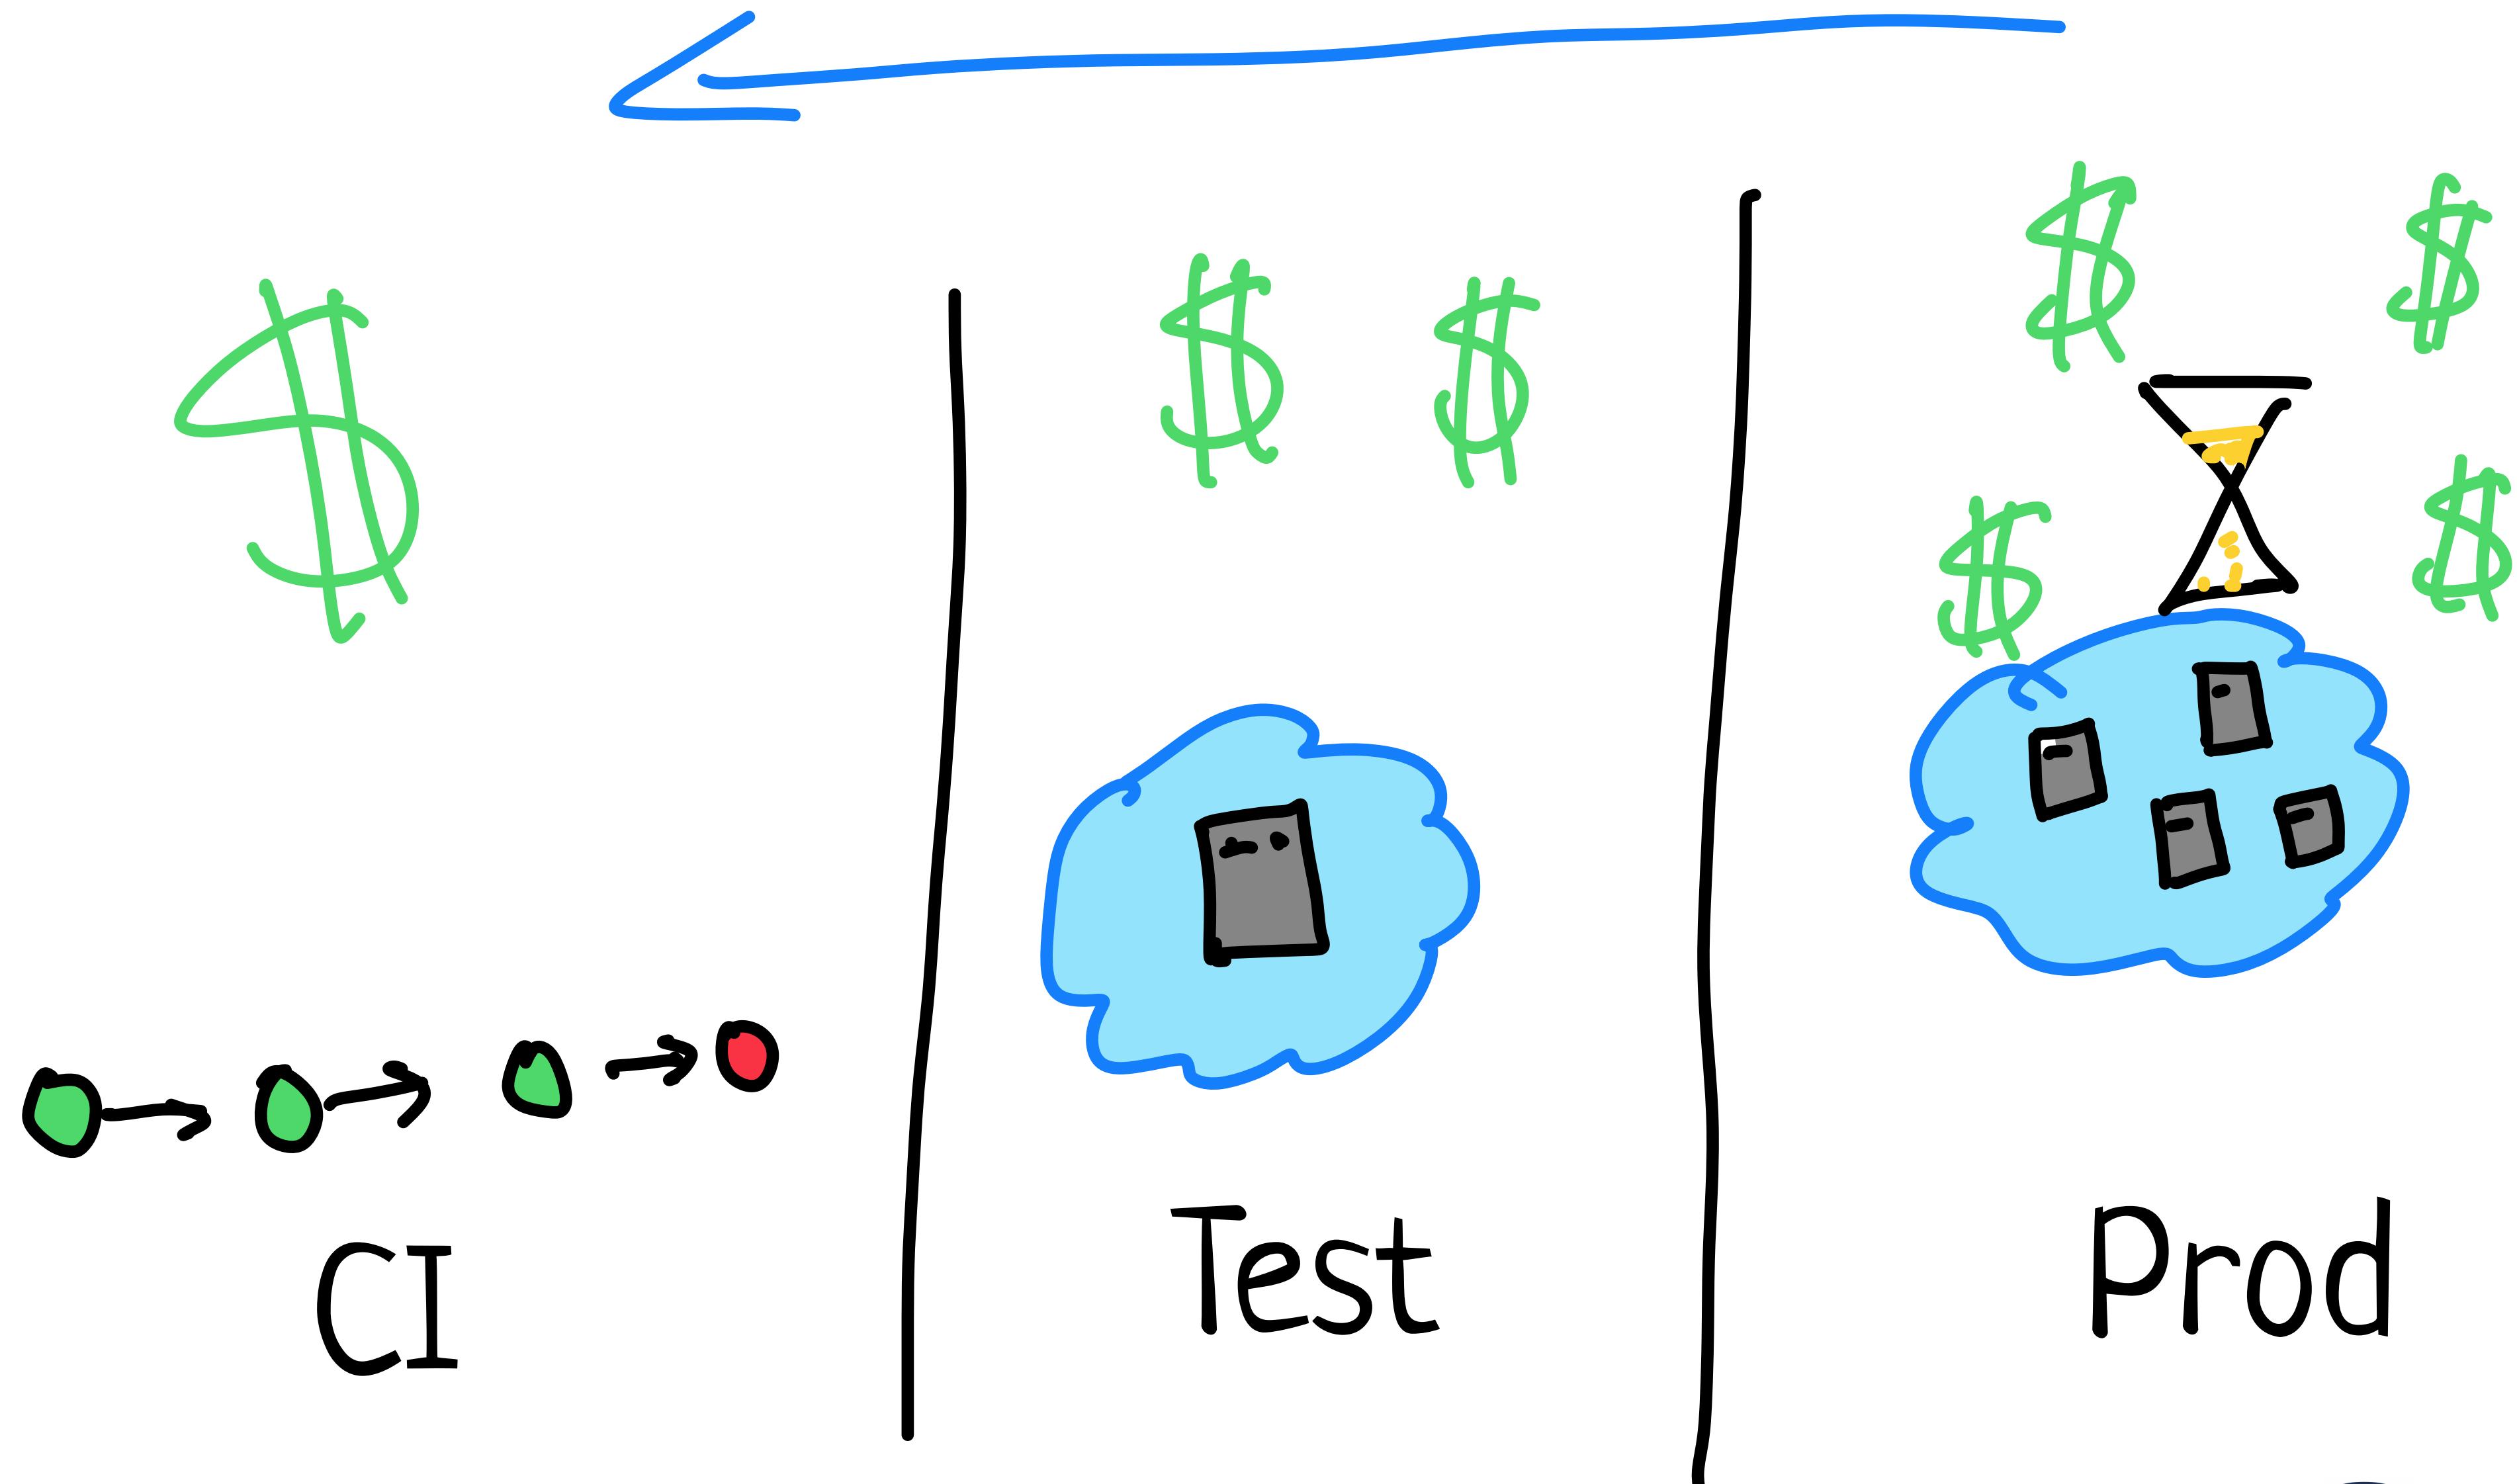
CI

- Build
- Linting
- Tests
- Secrets

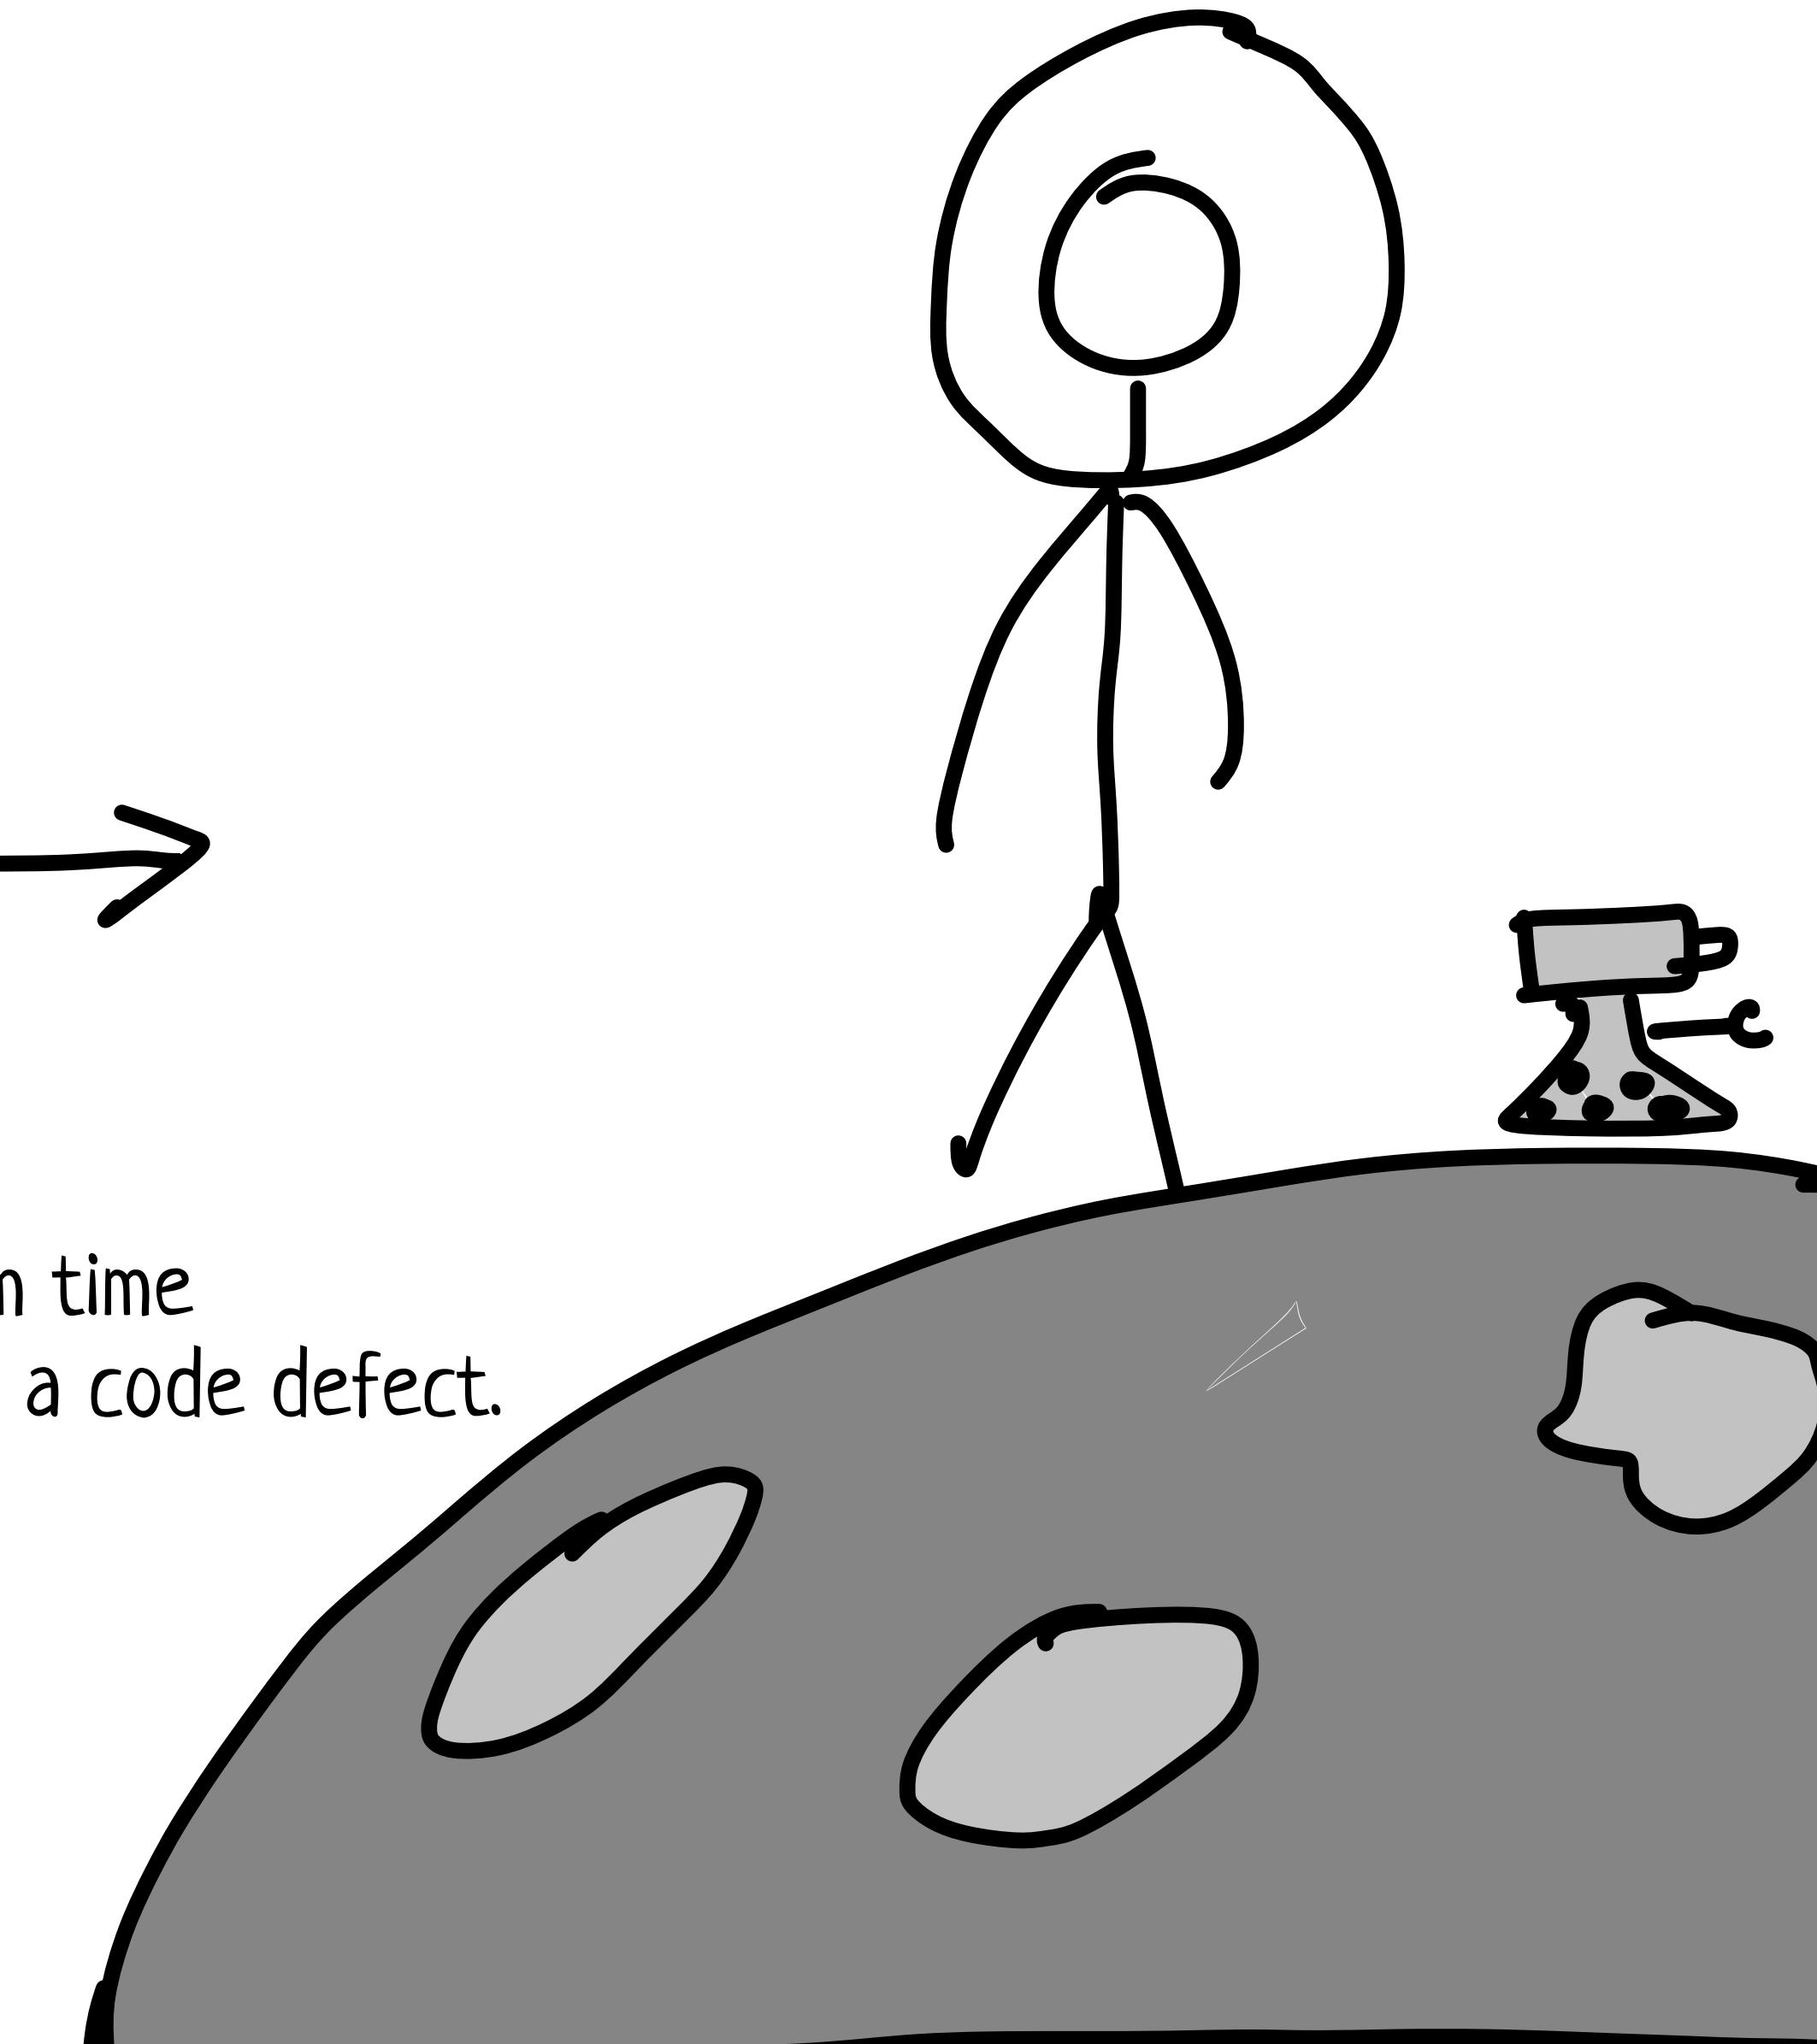
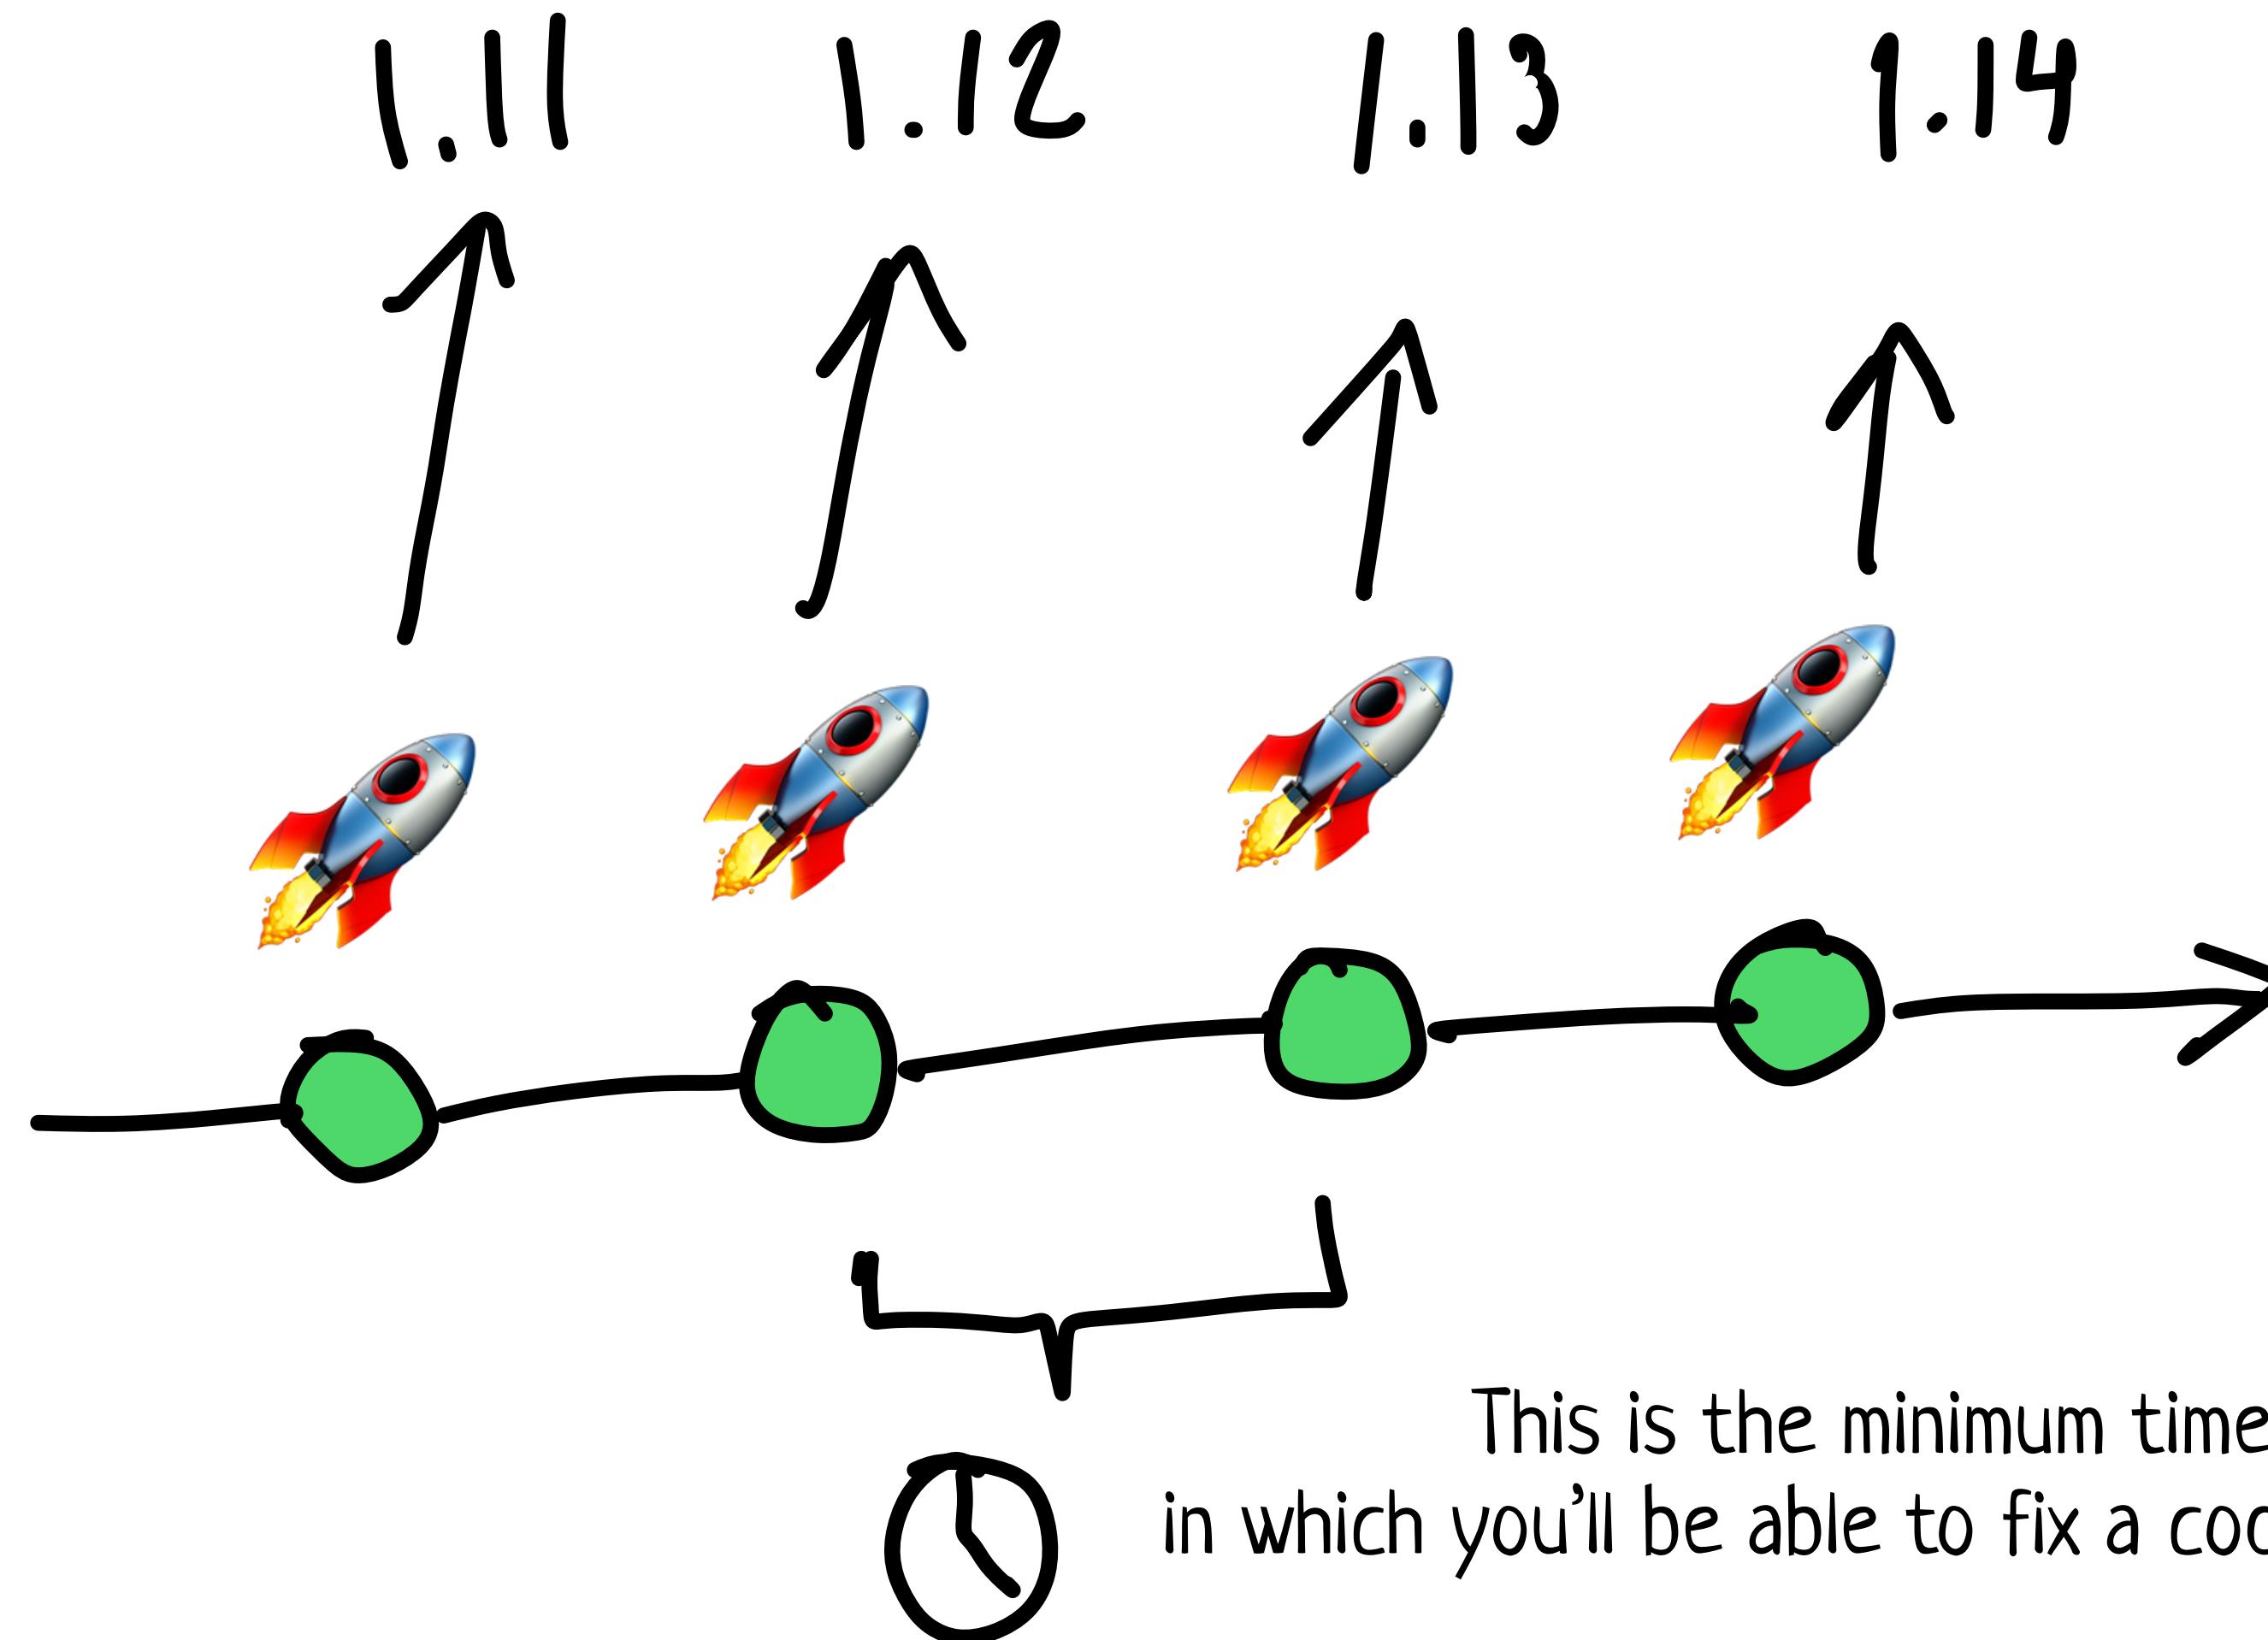
After merge

- Build
- Tests
- Secrets

Failures are easier to catch, the earlier they come.



Continuous deployment

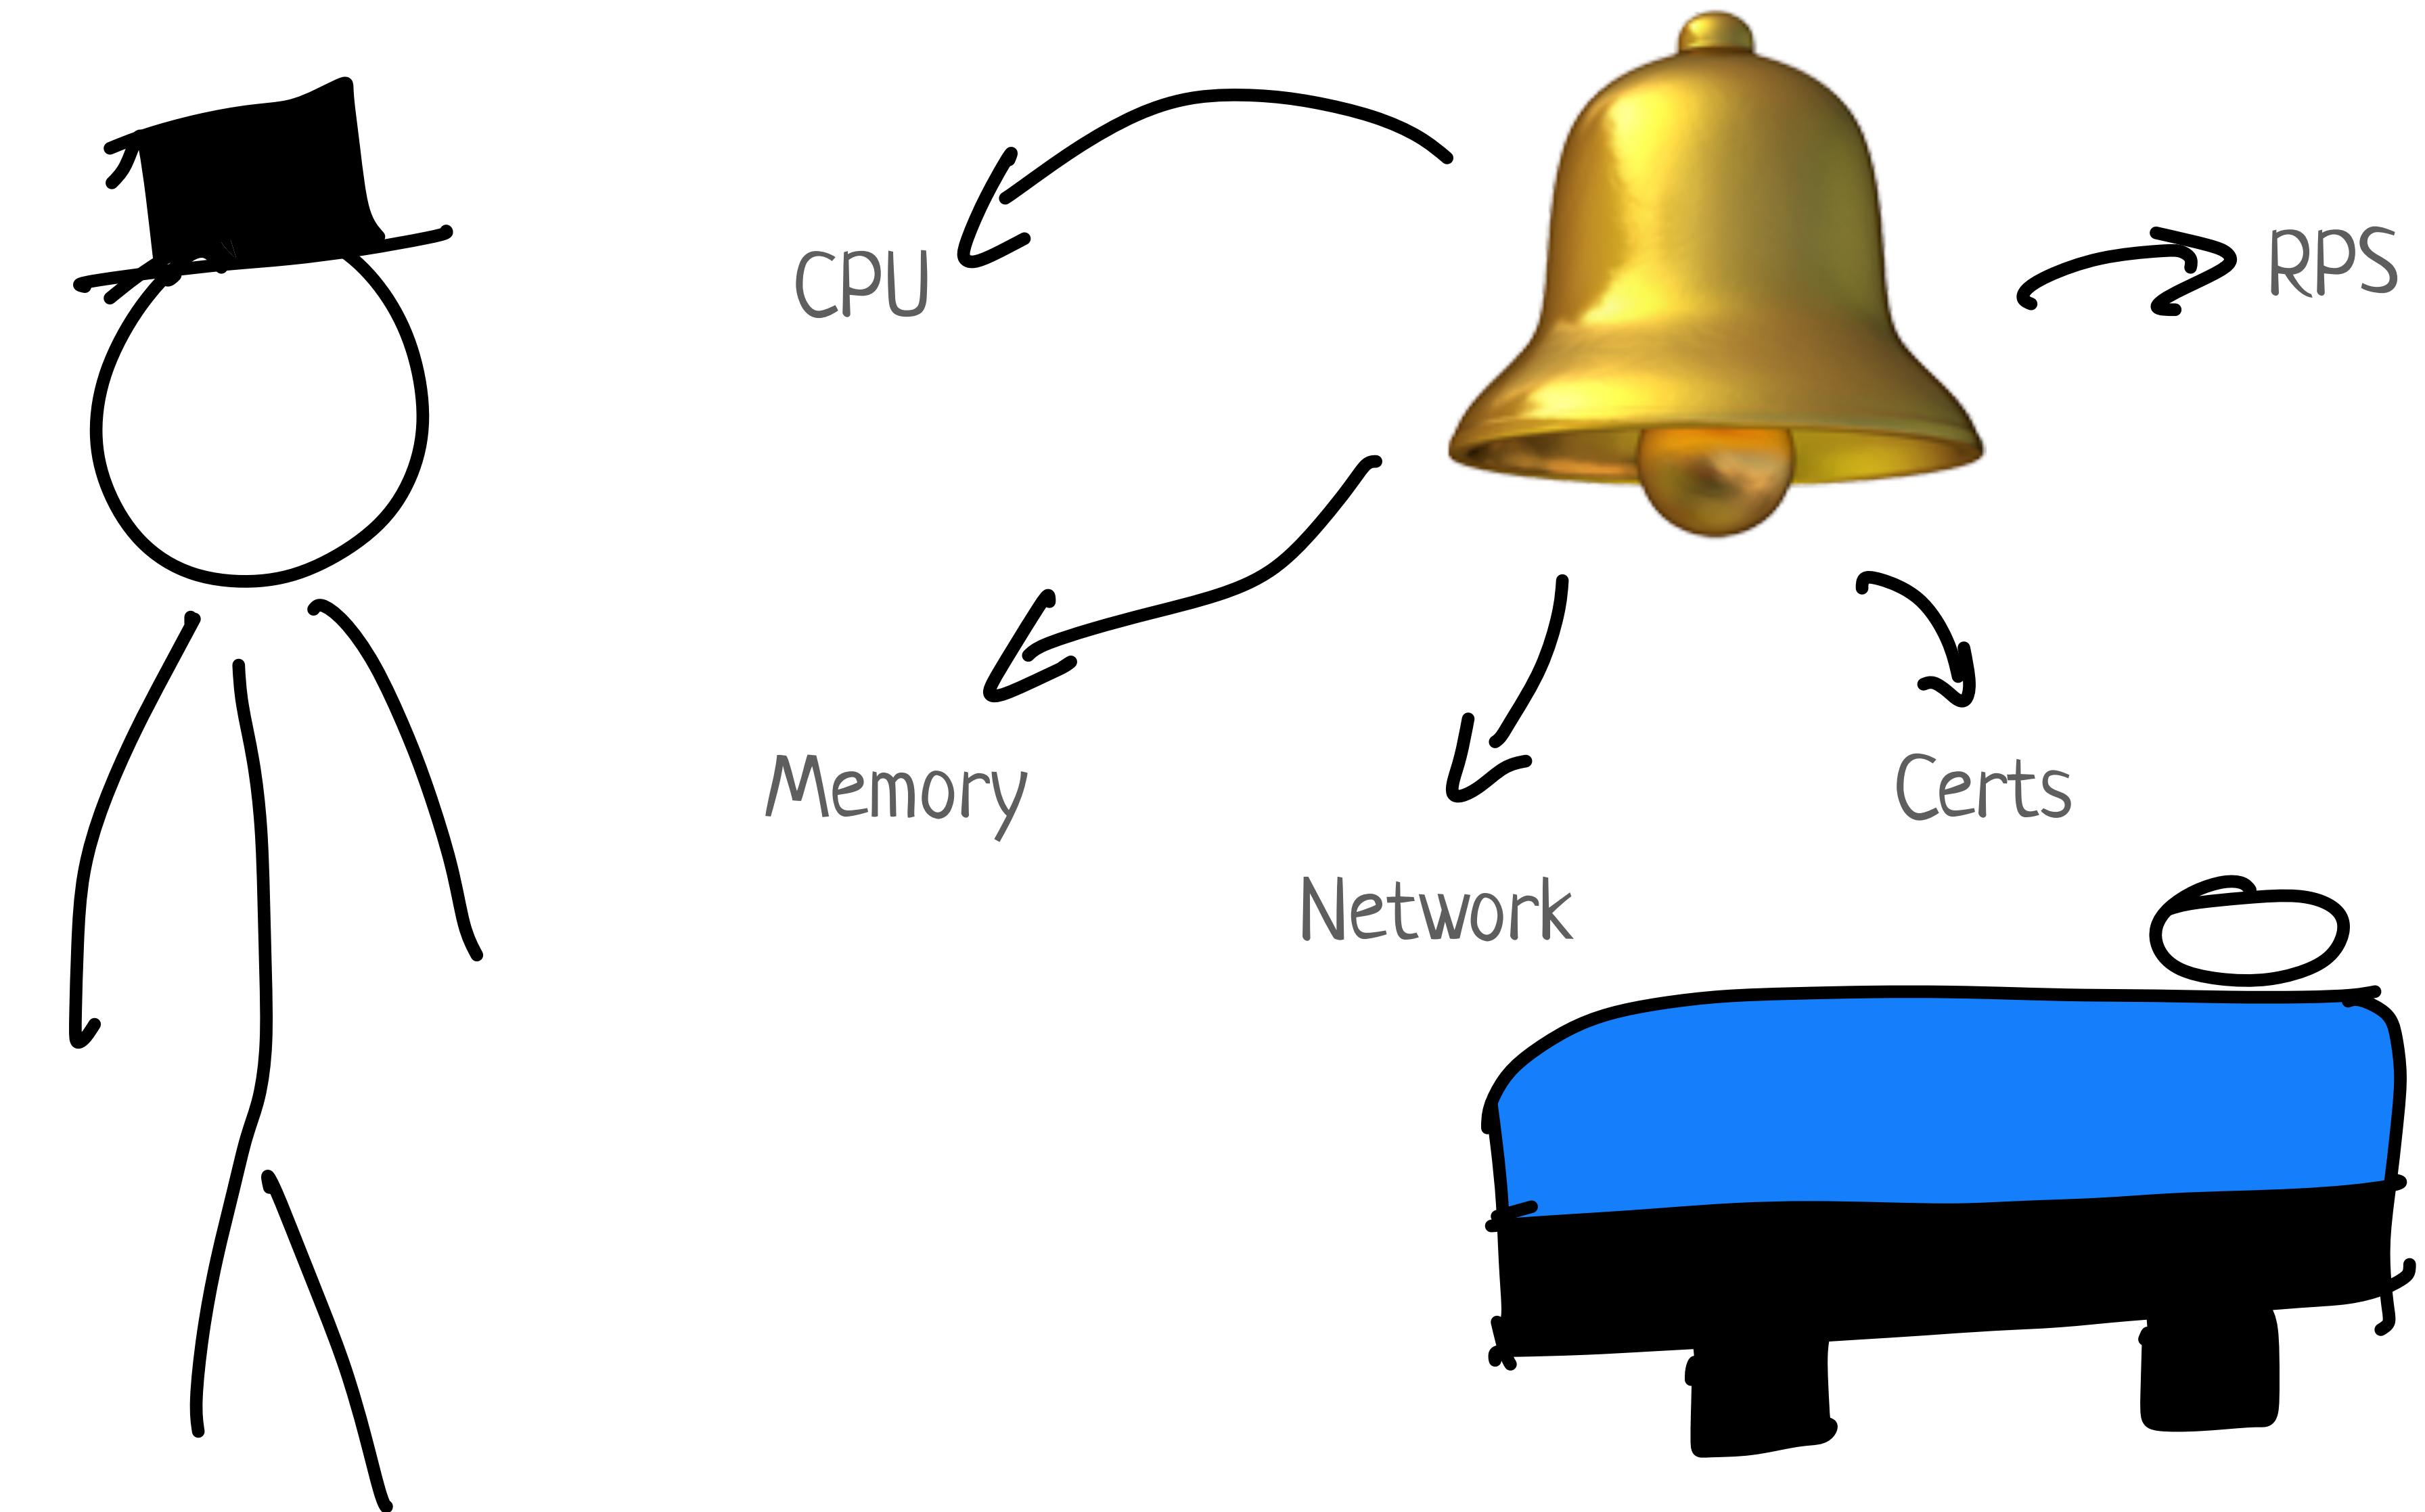


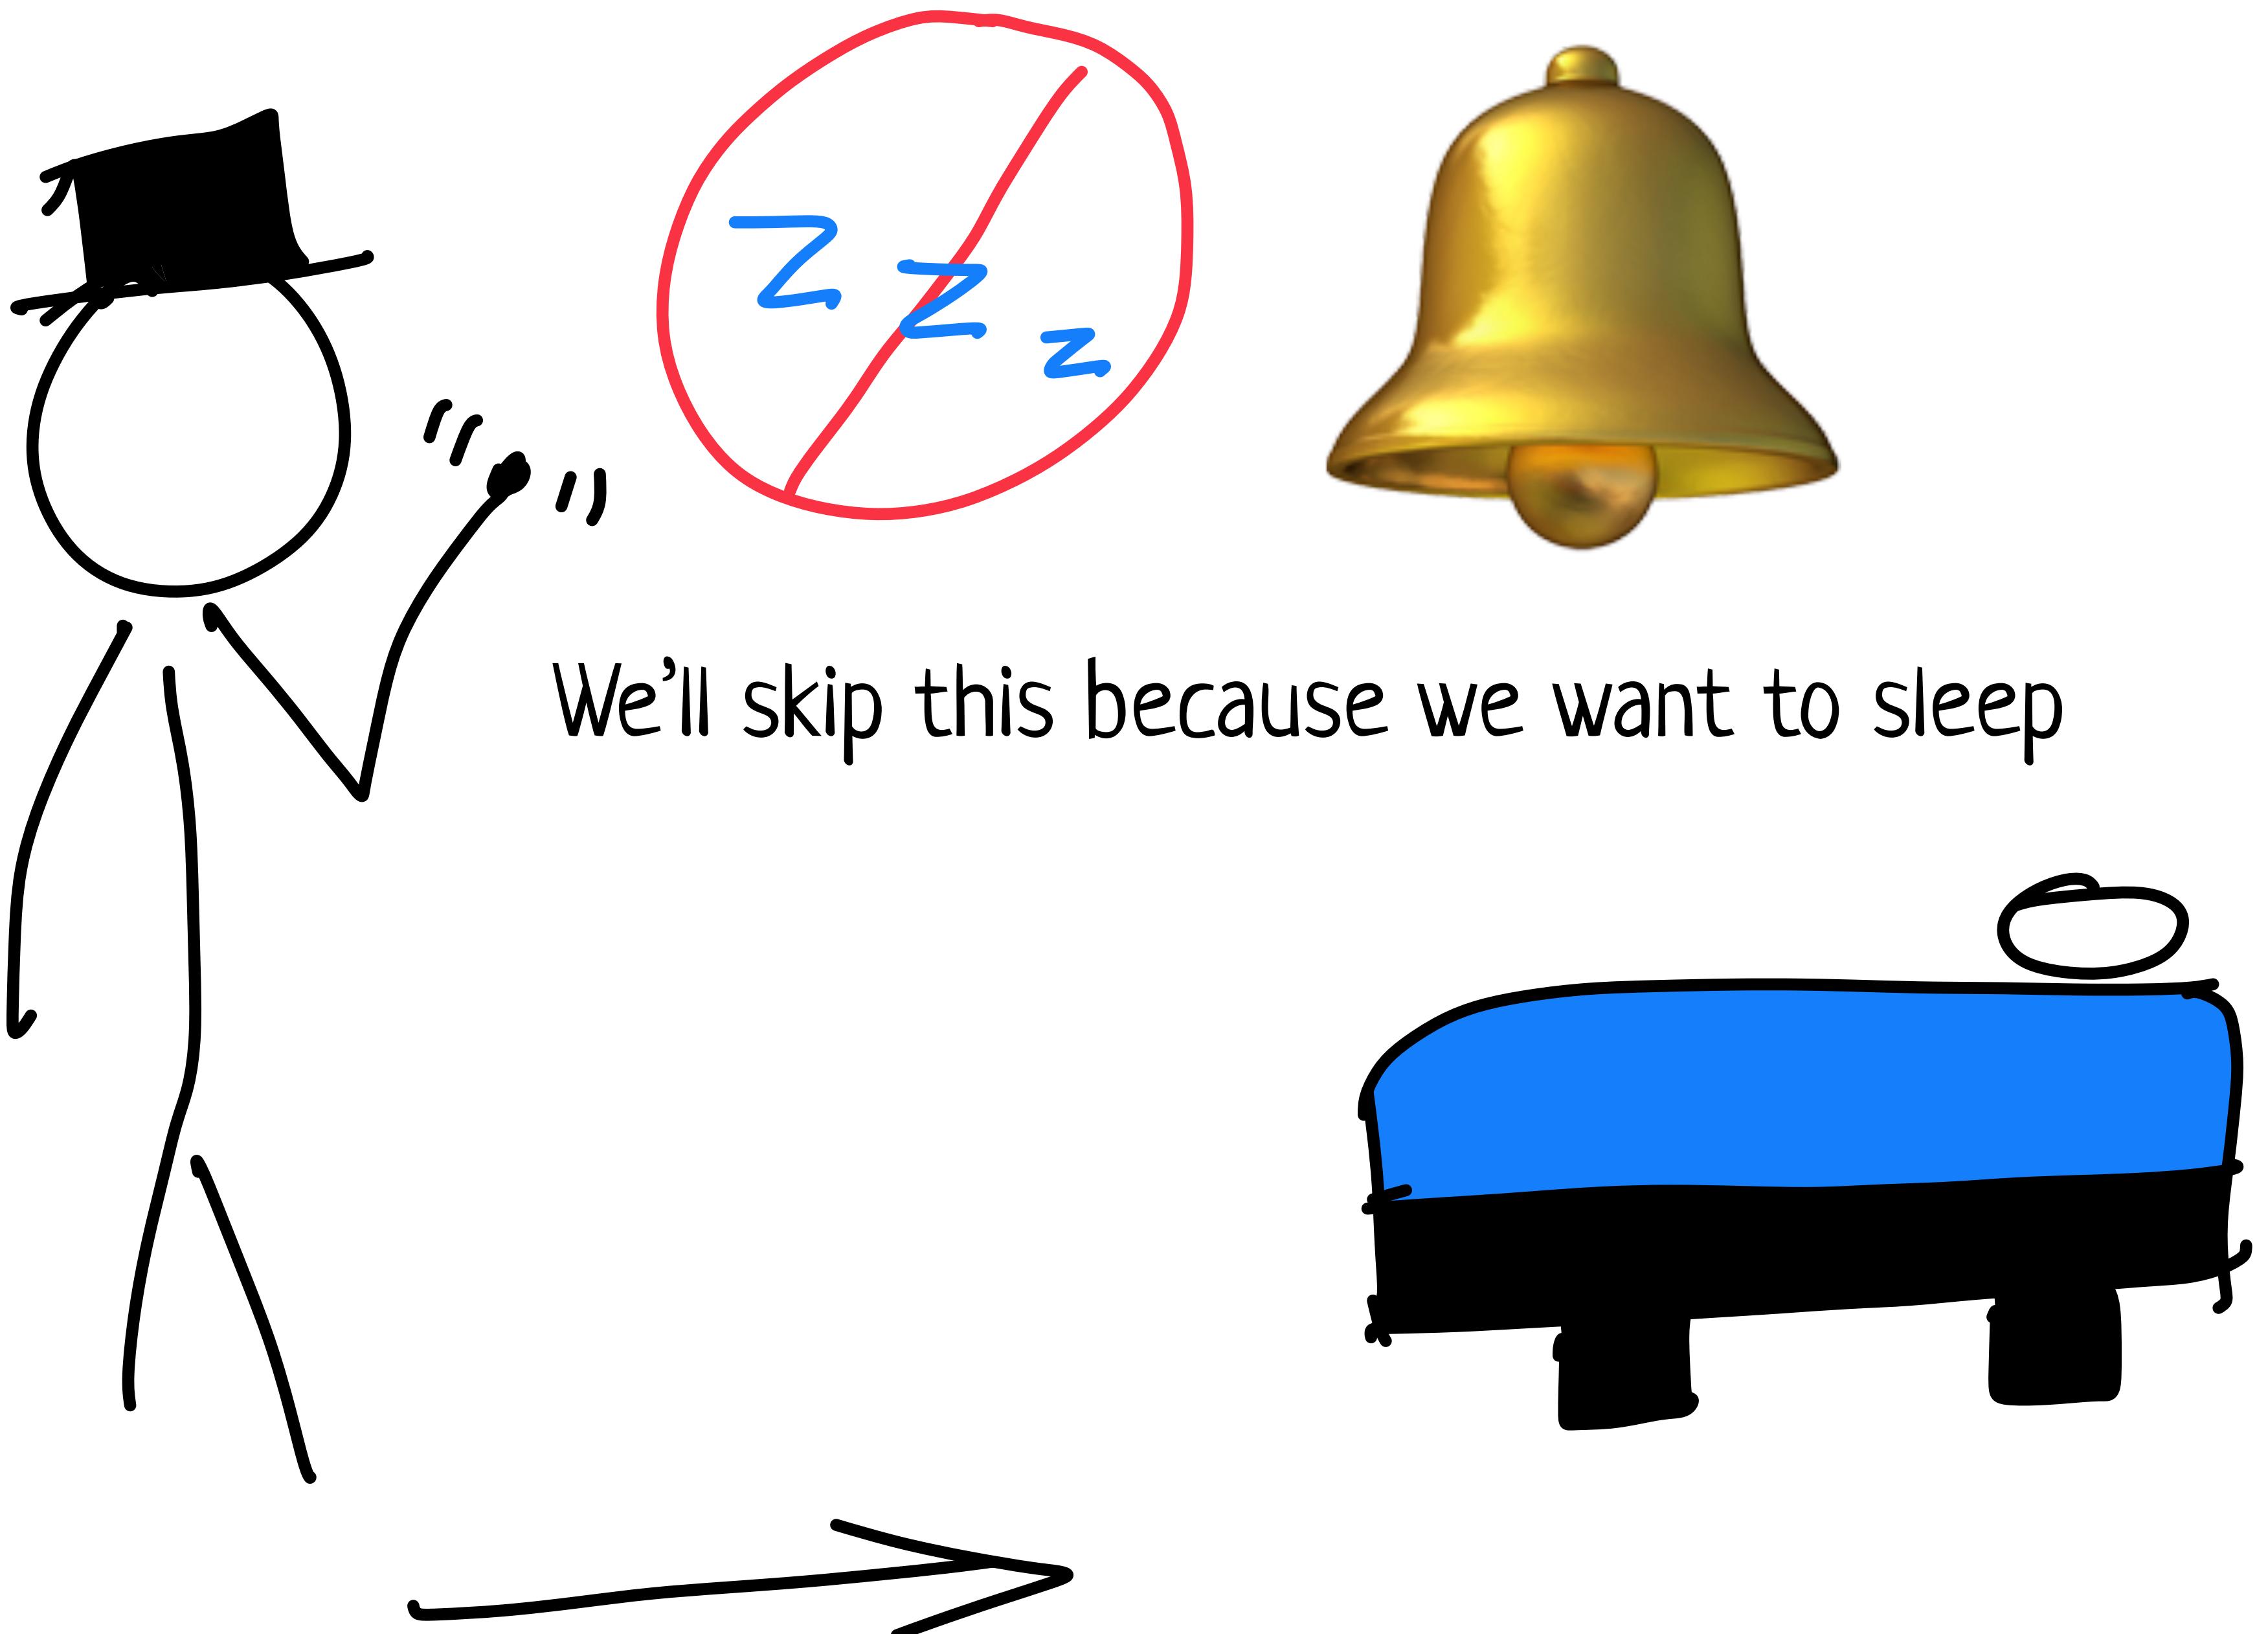
CI/CD

Continuous Integration
/ Continuous Deployment

- Github Actions
- Jenkins
- CircleCI
- TravisCI
- Etc...

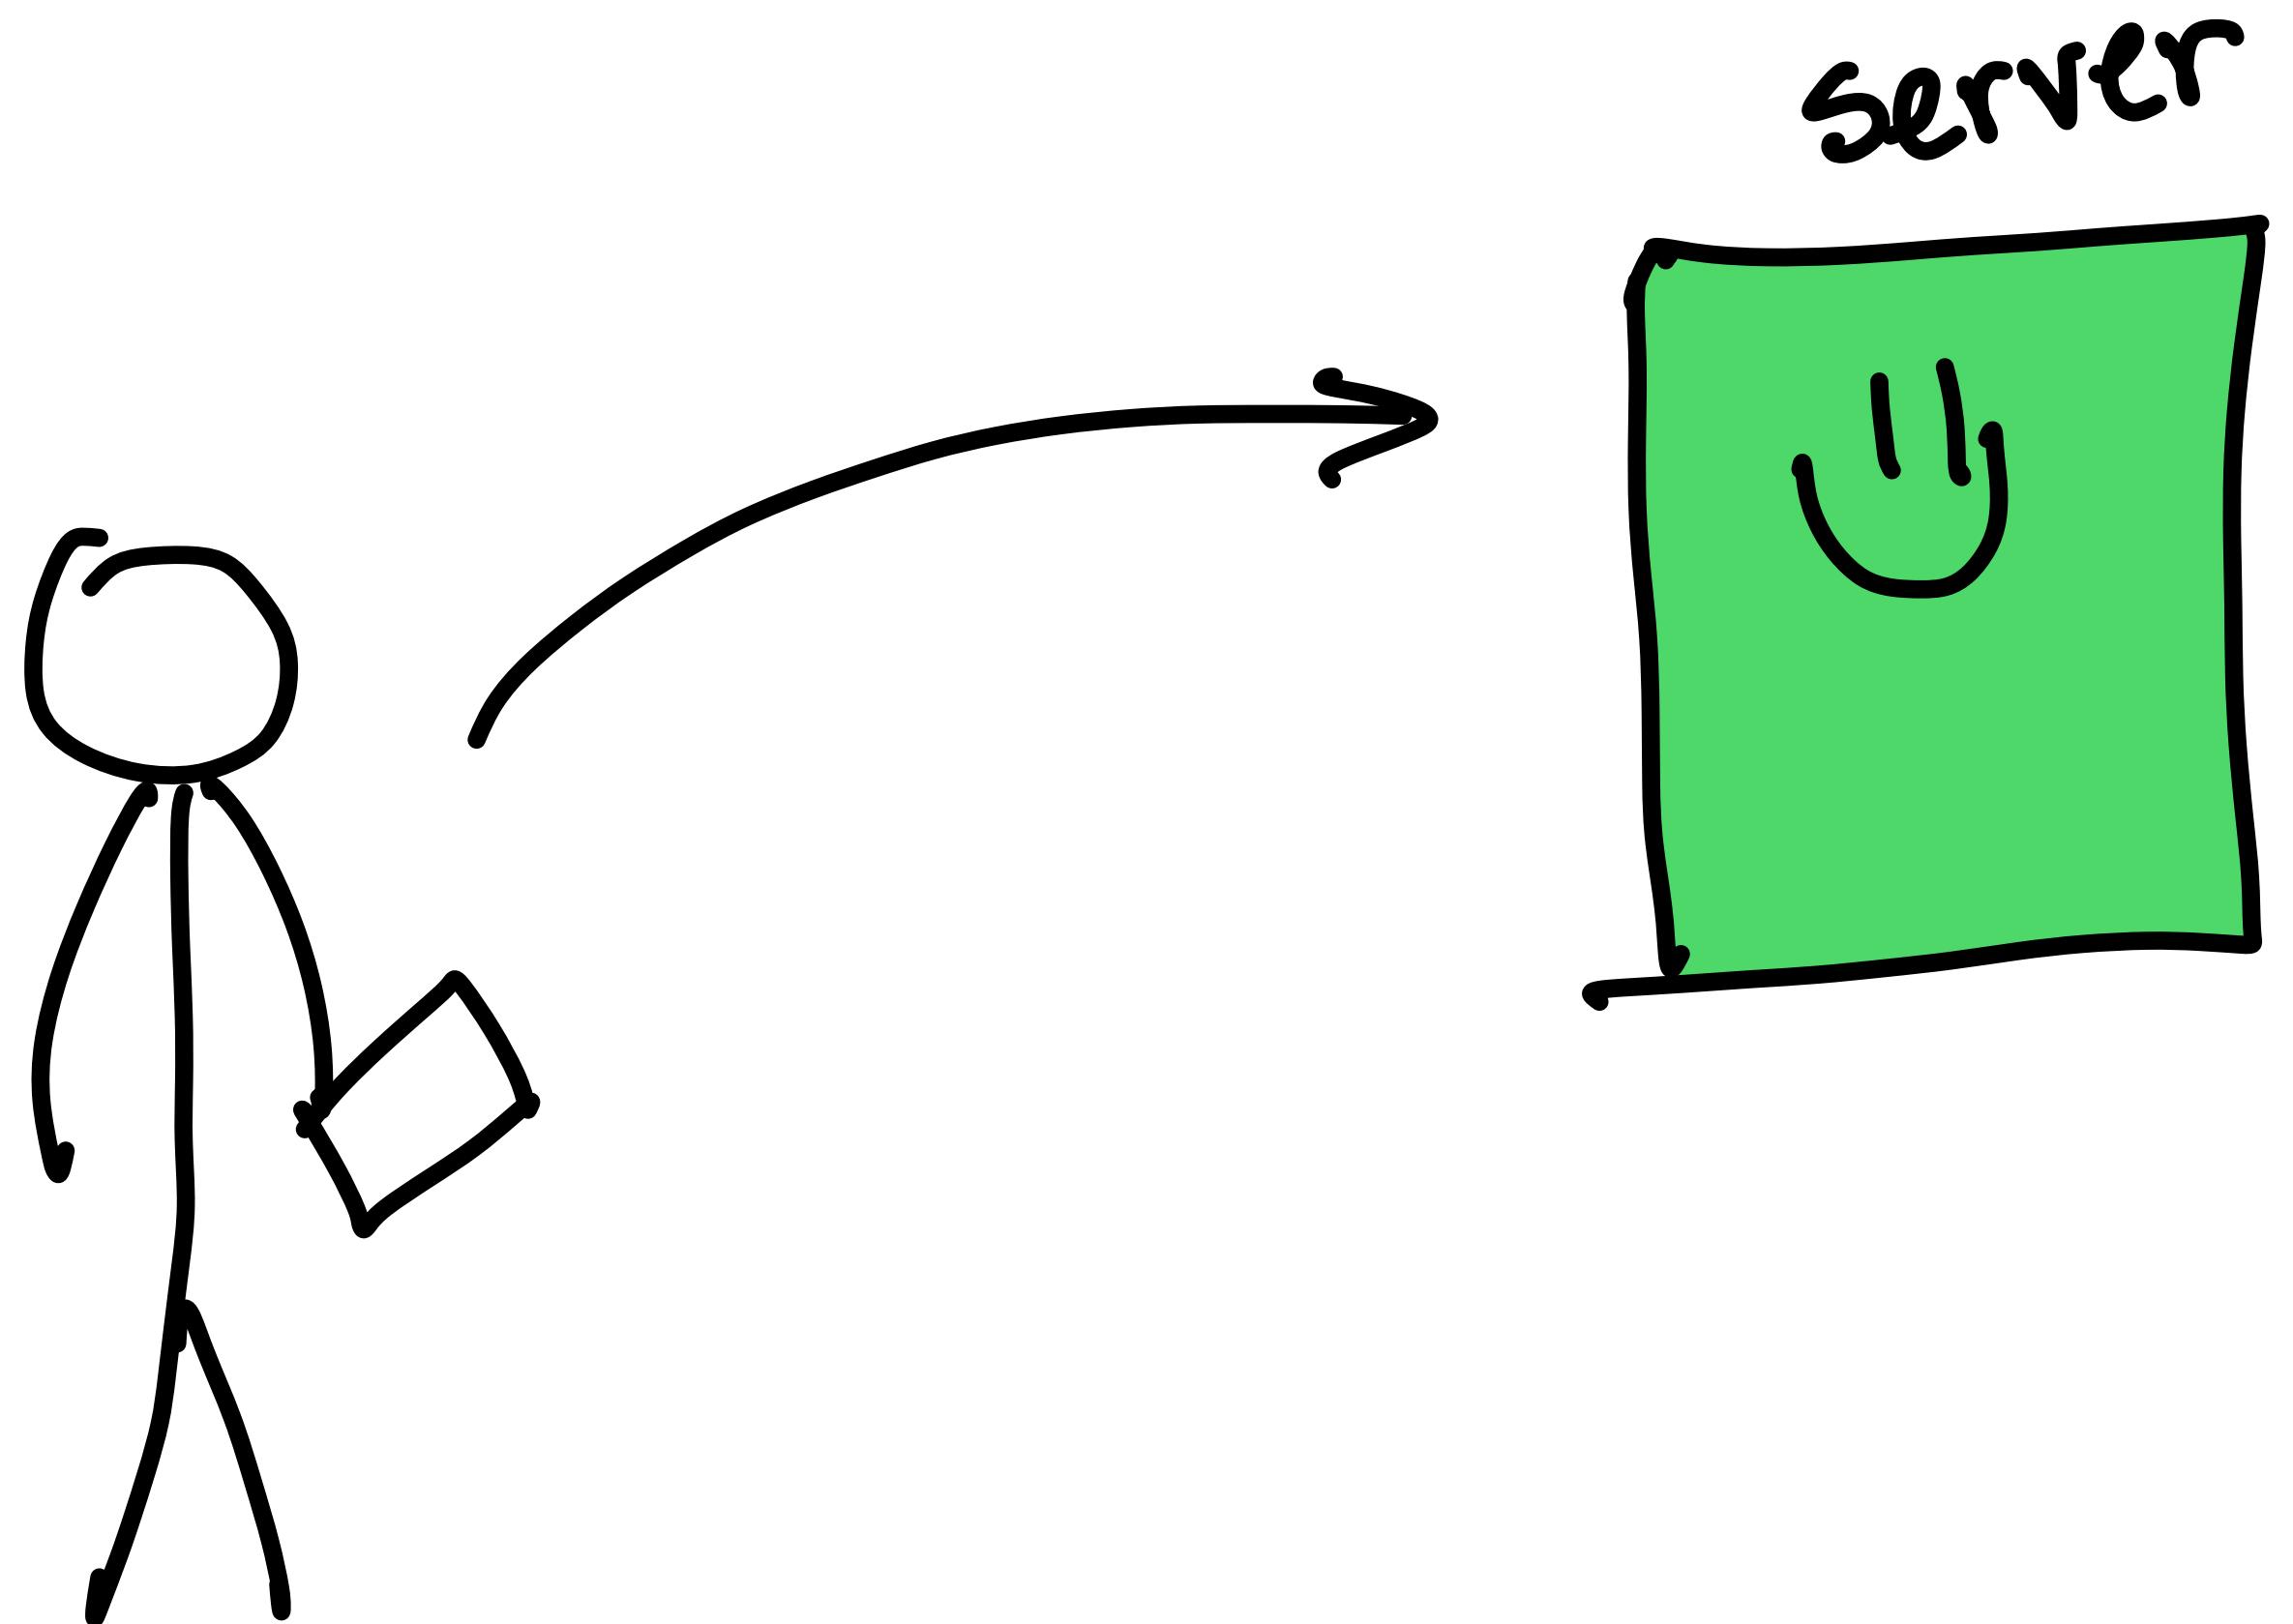
Alerts



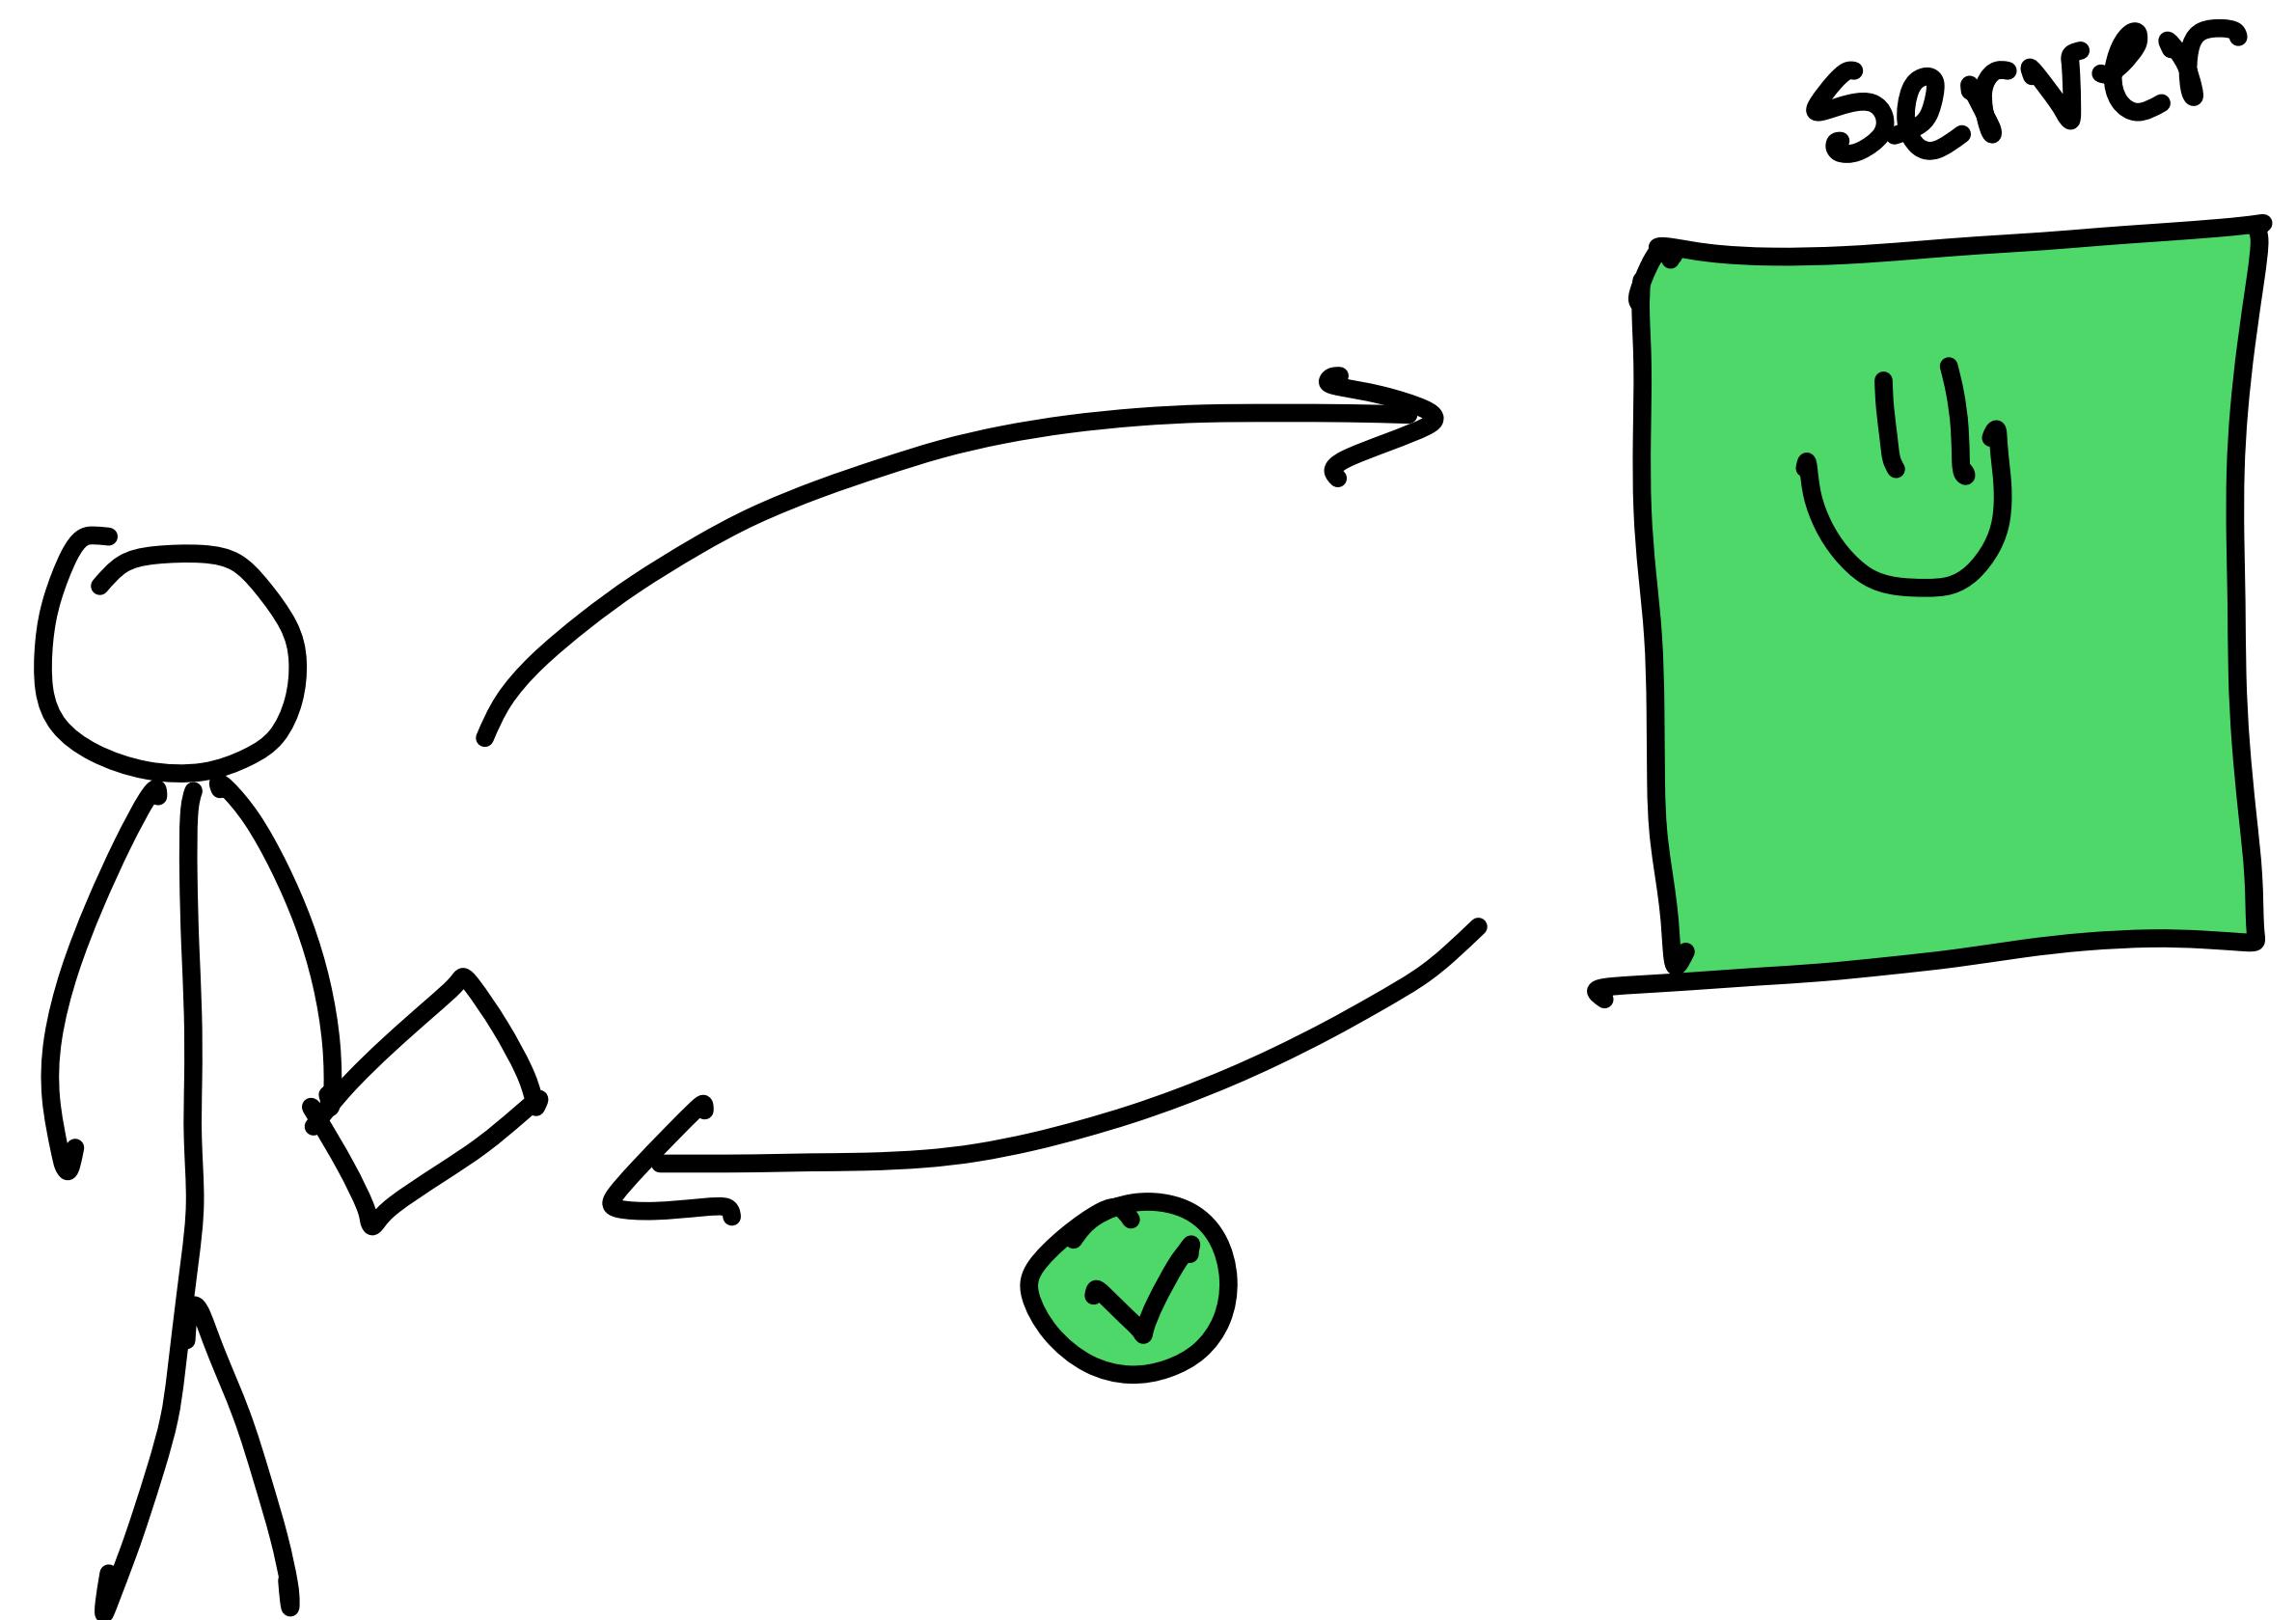


Load Balancer

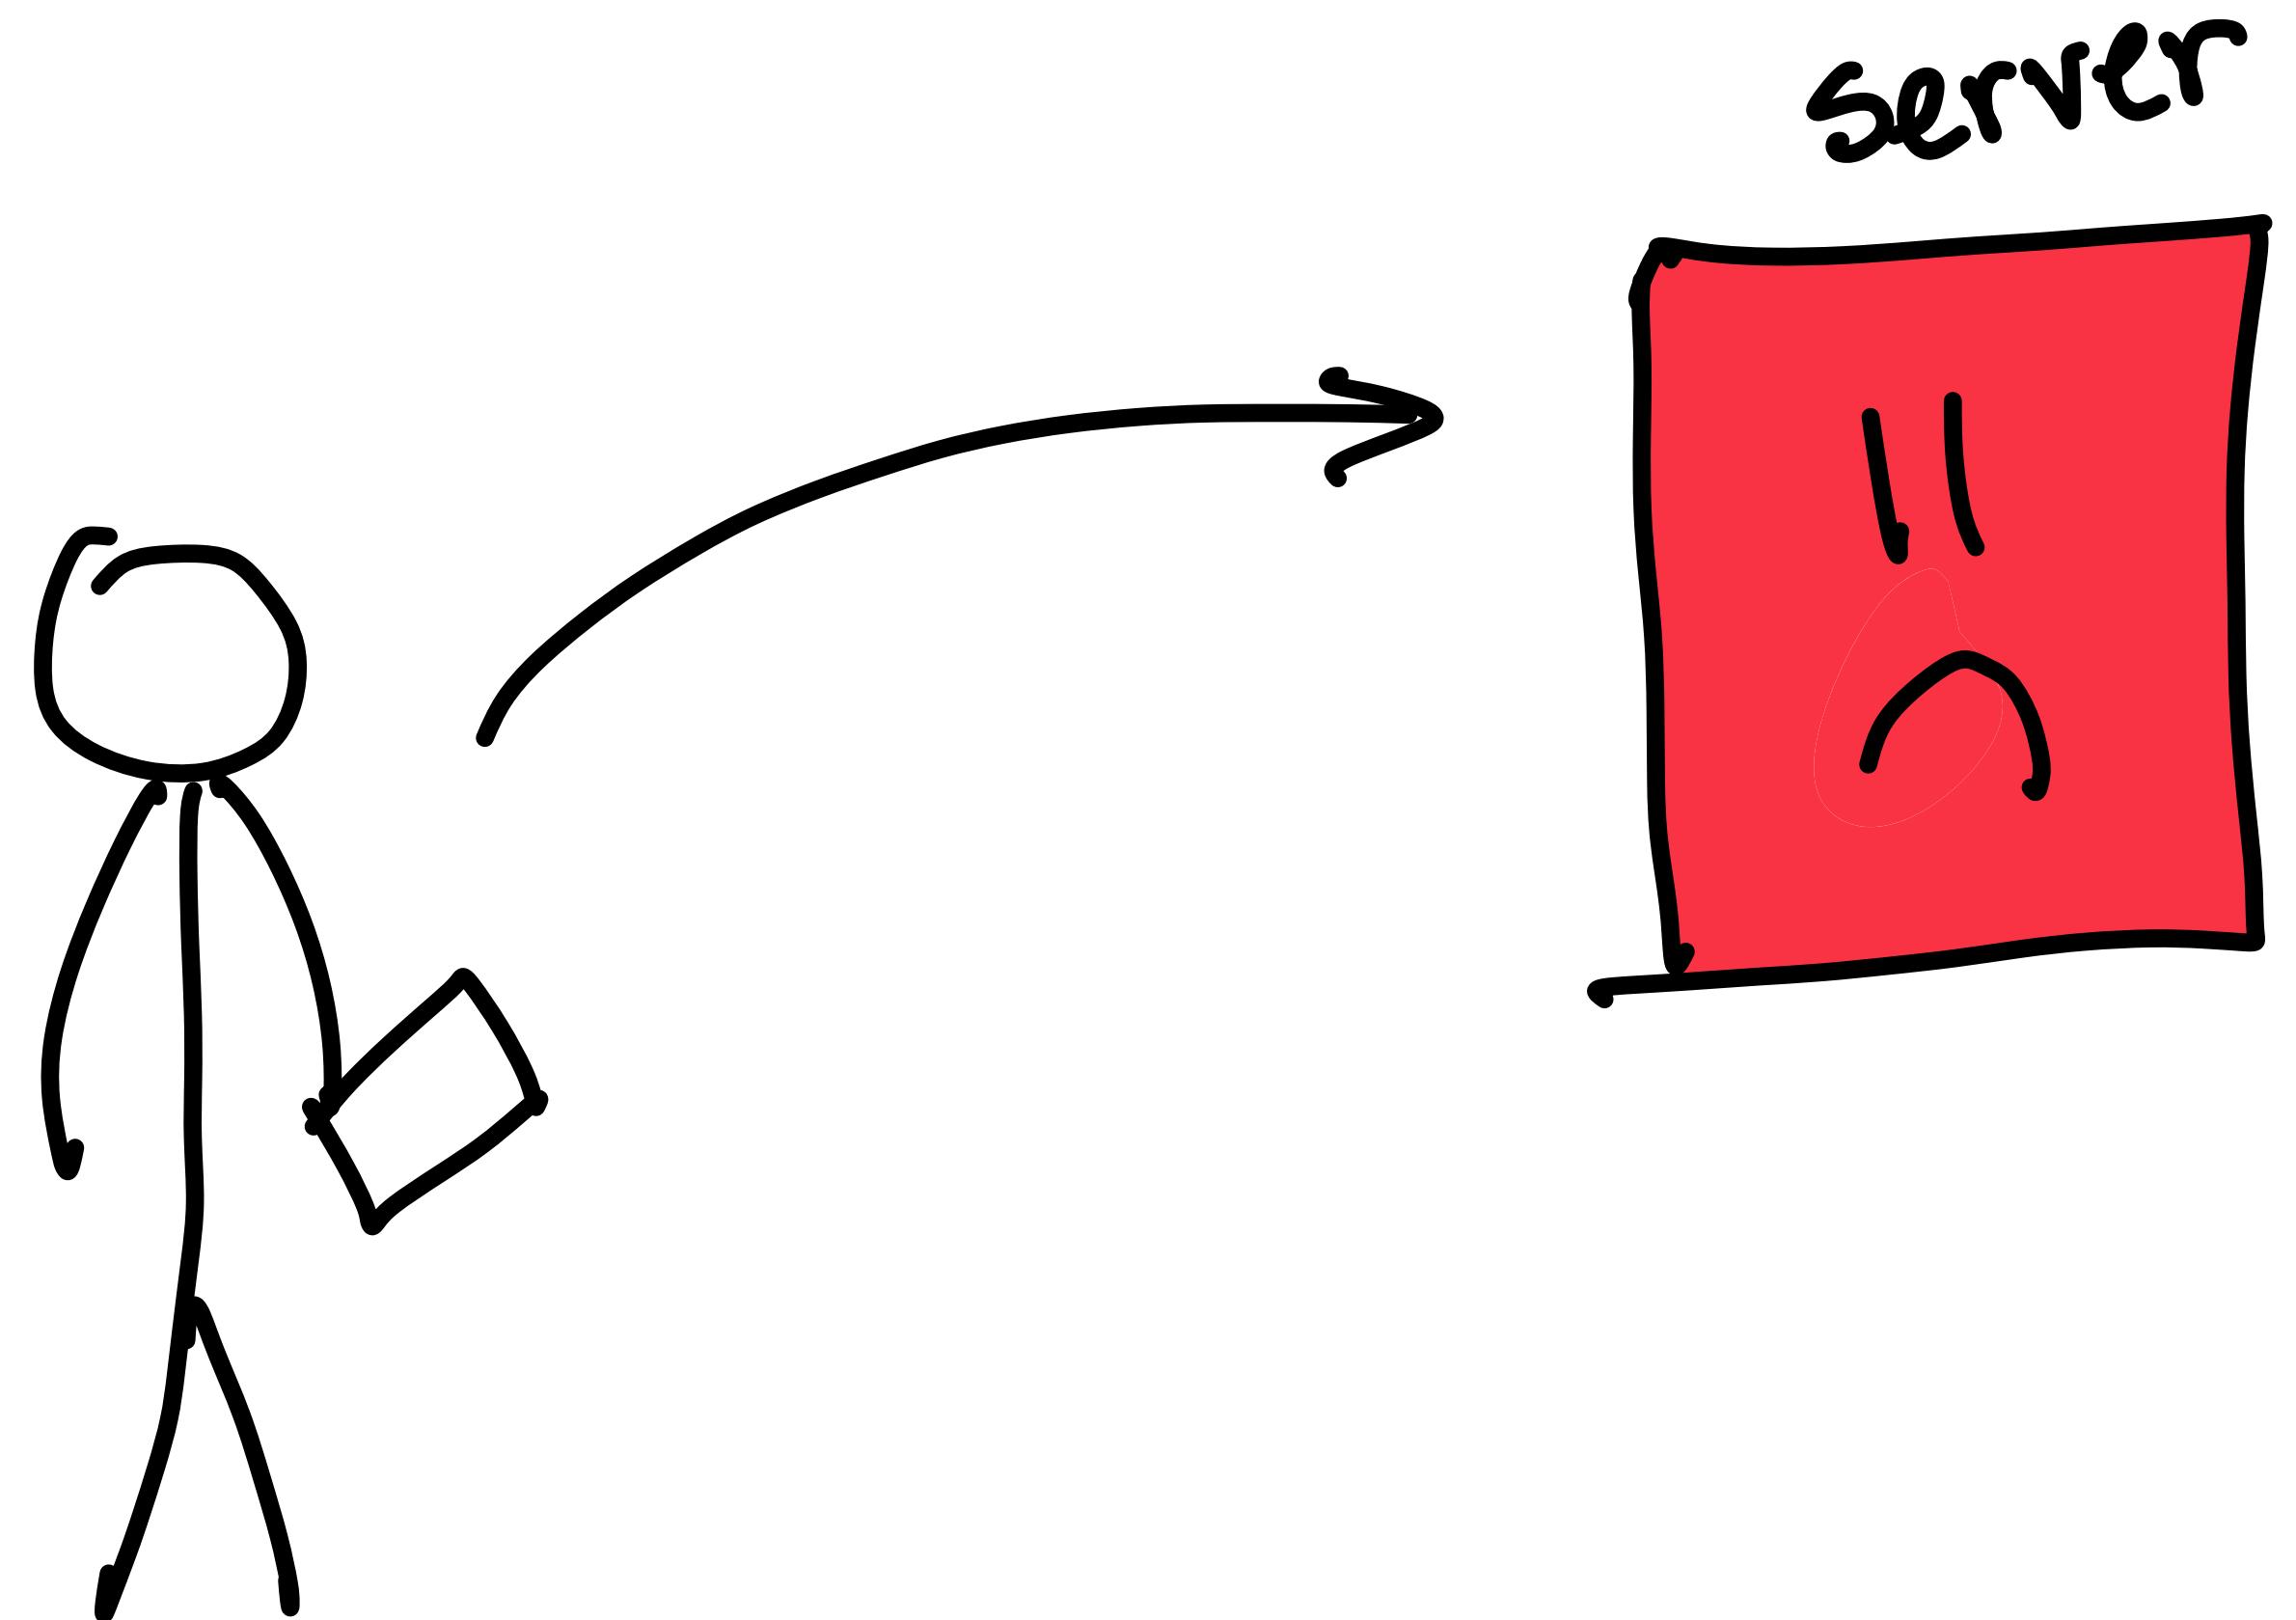
If we have only one server and no load balancing



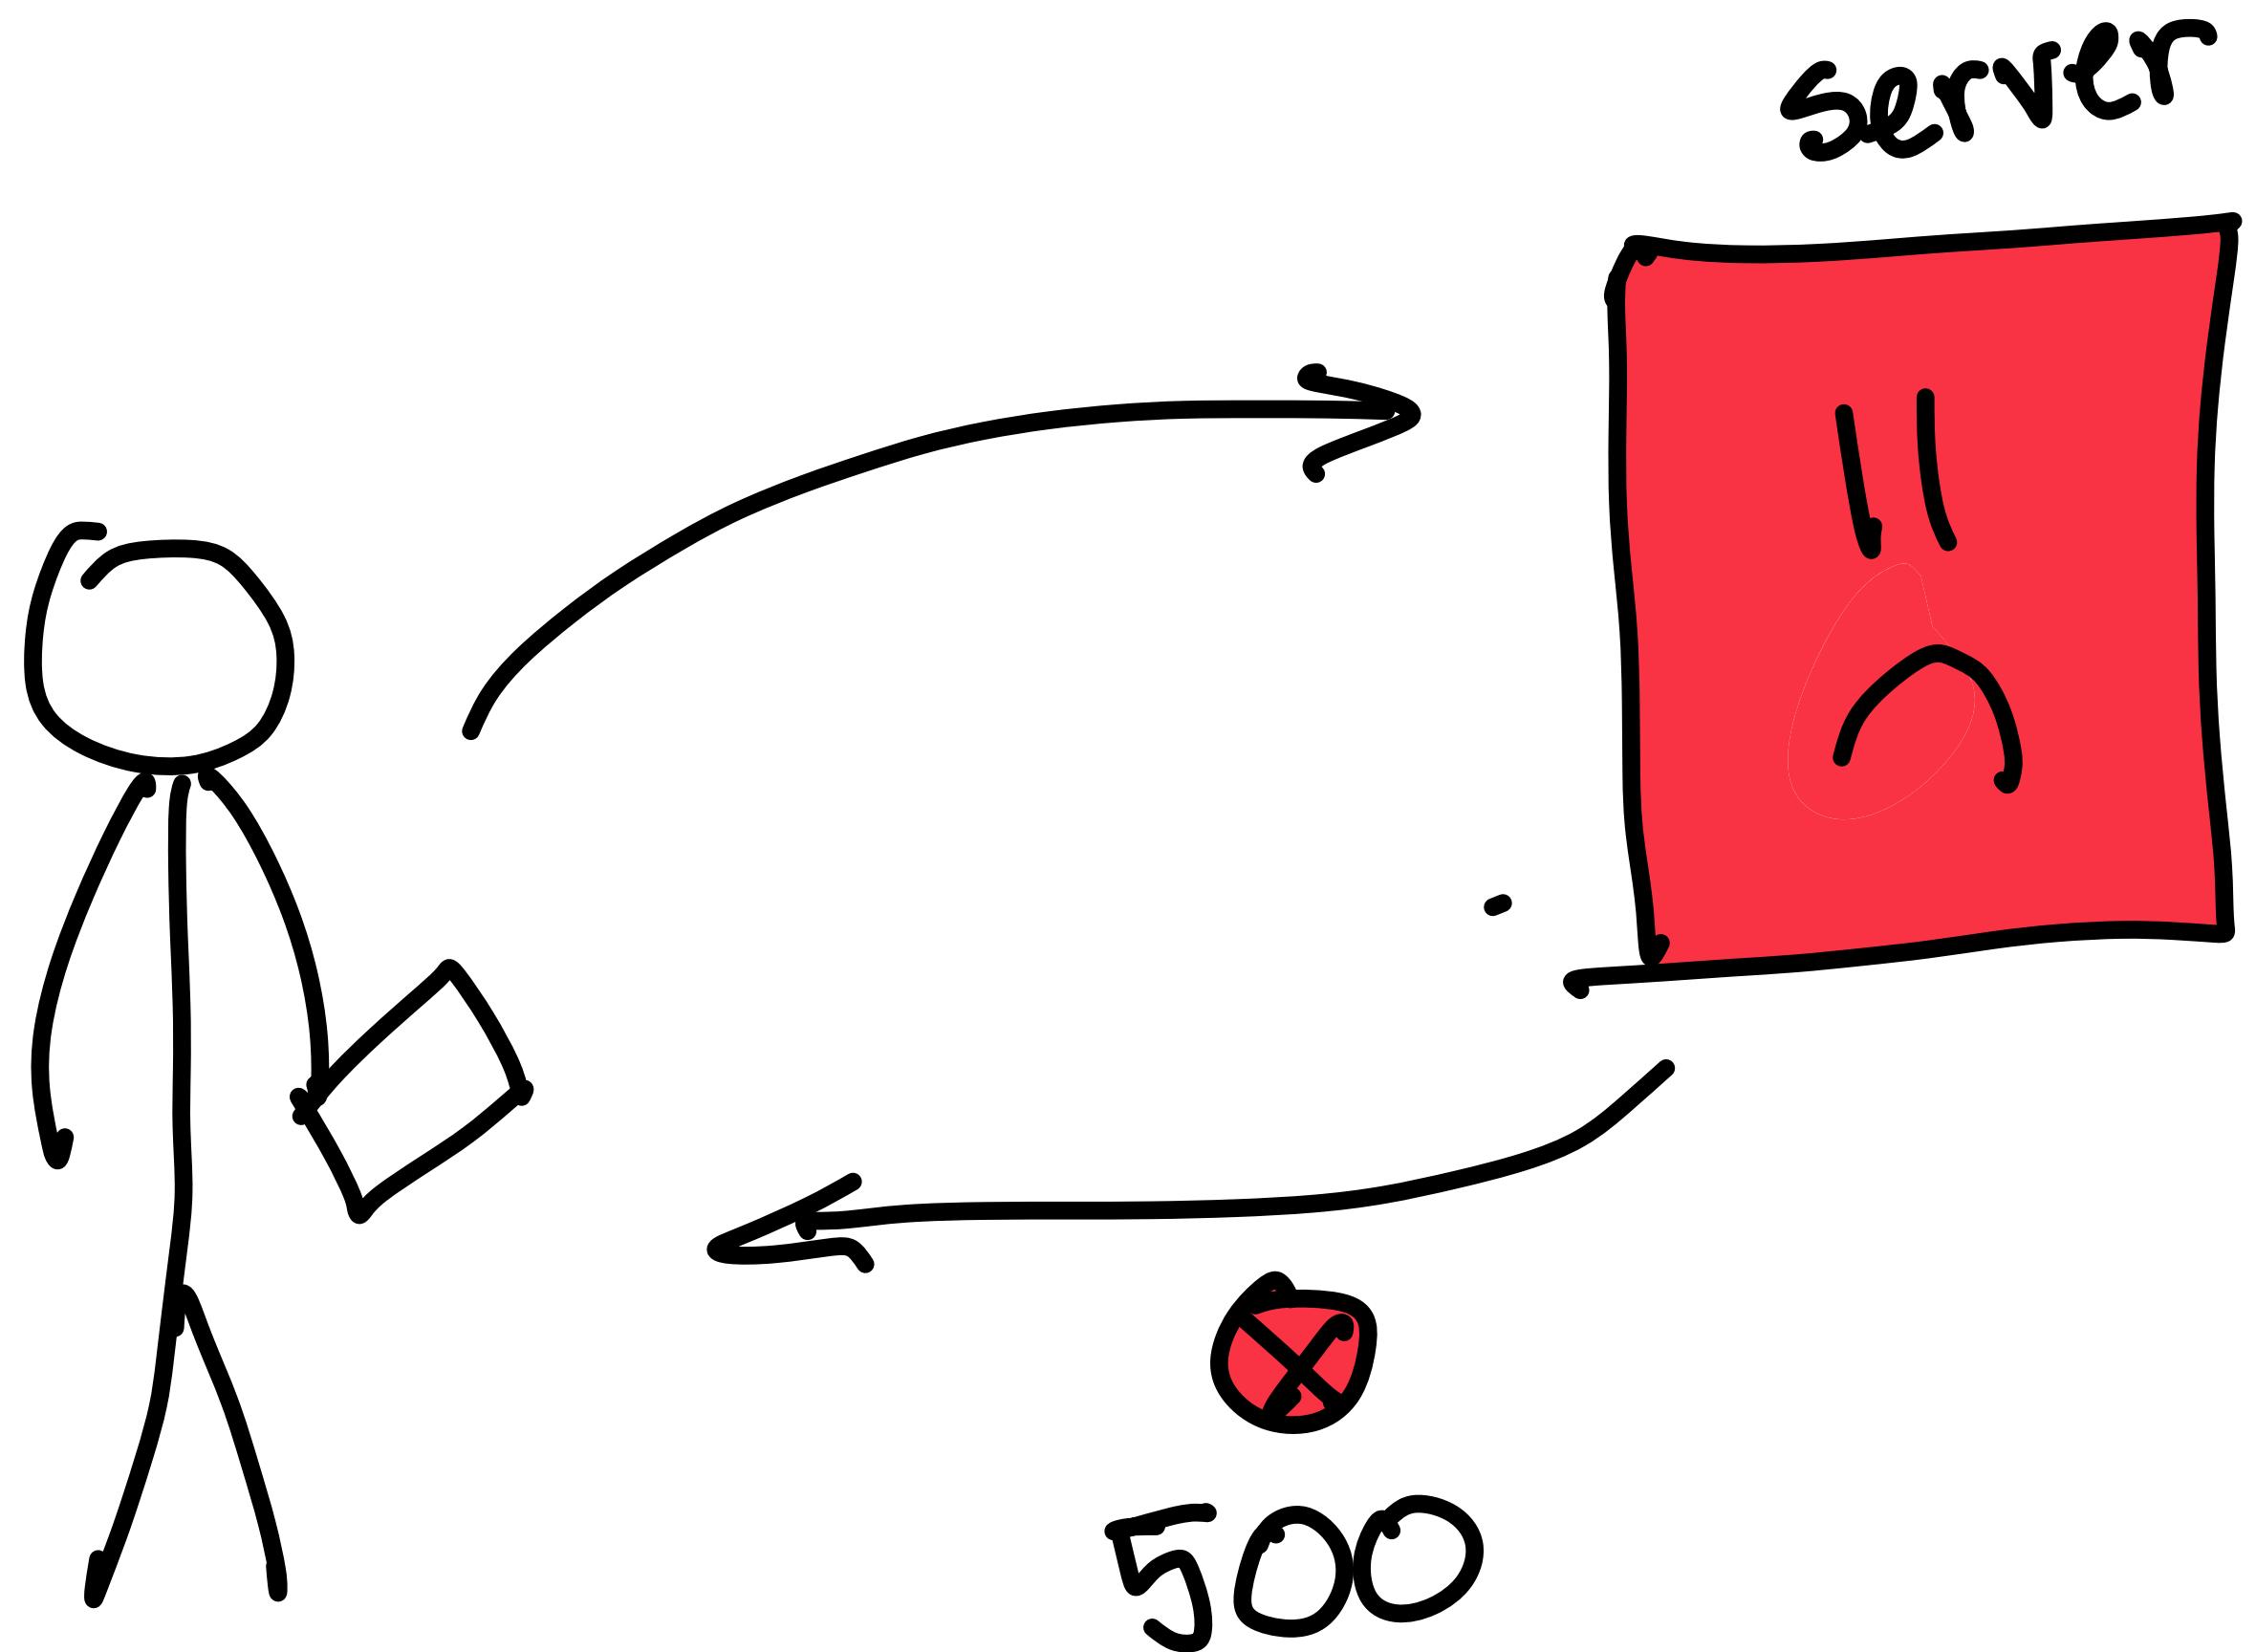
If we have only one server and no load balancing



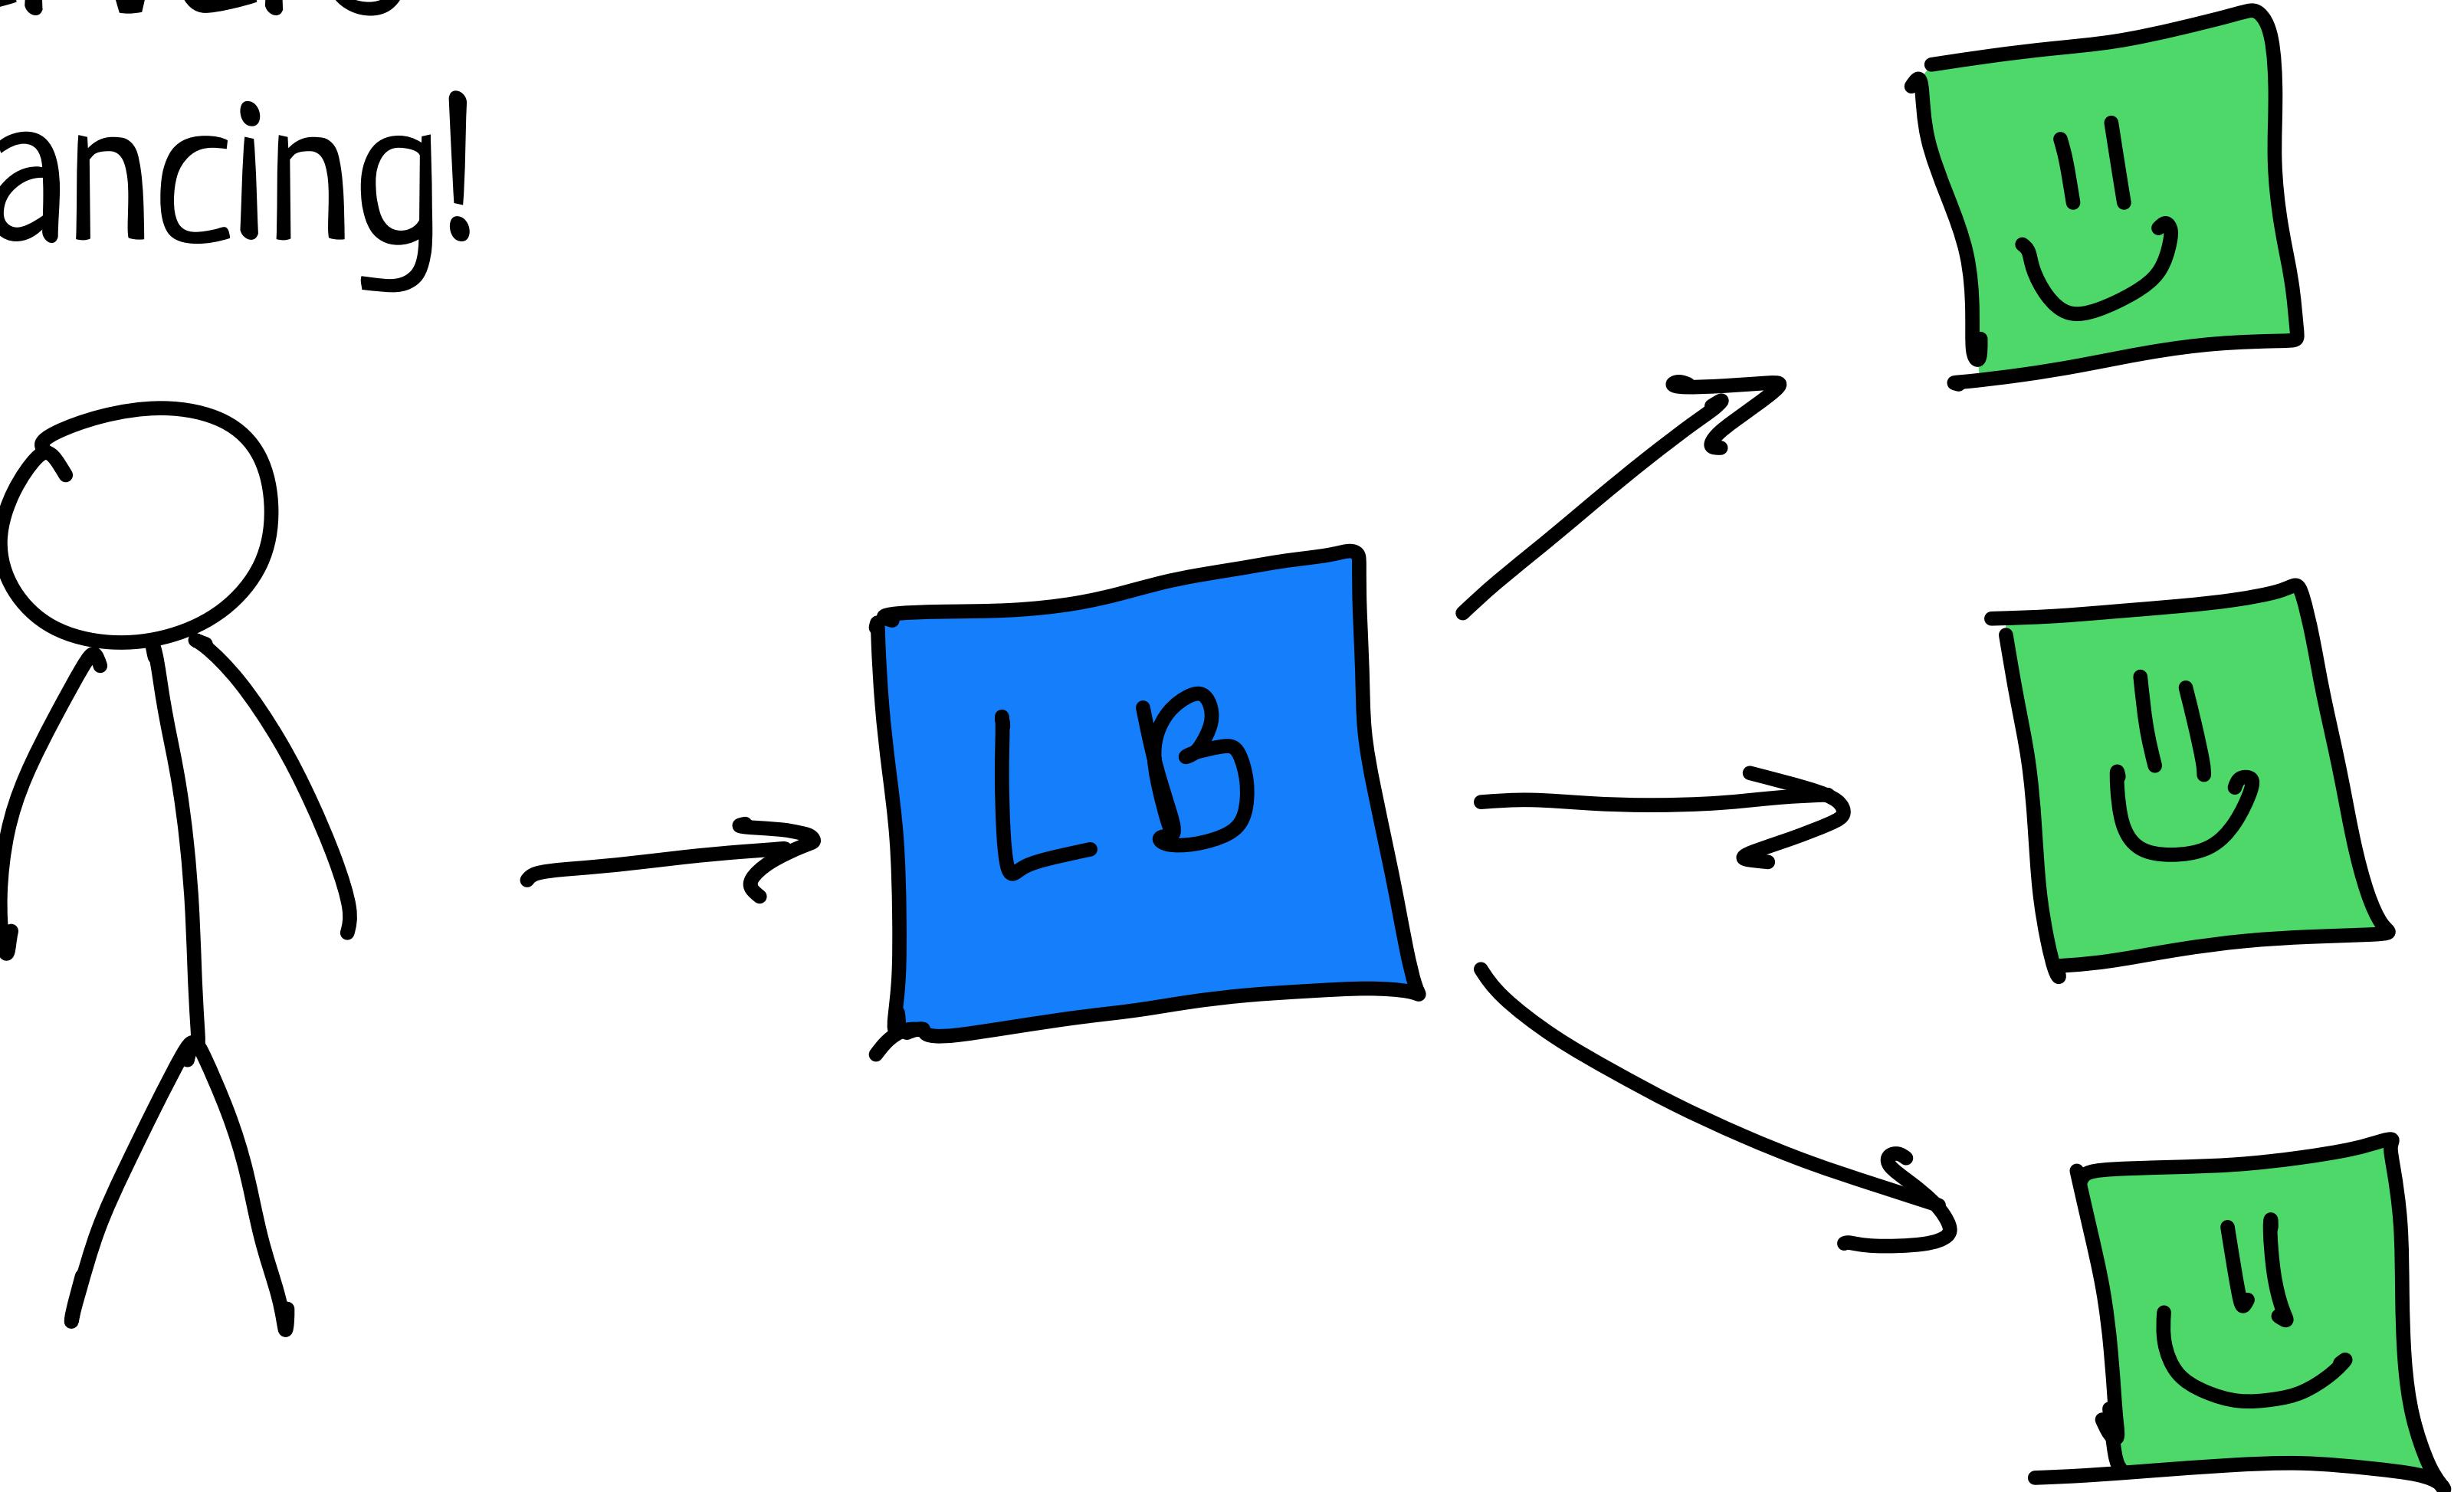
If we have only one server and no load balancing



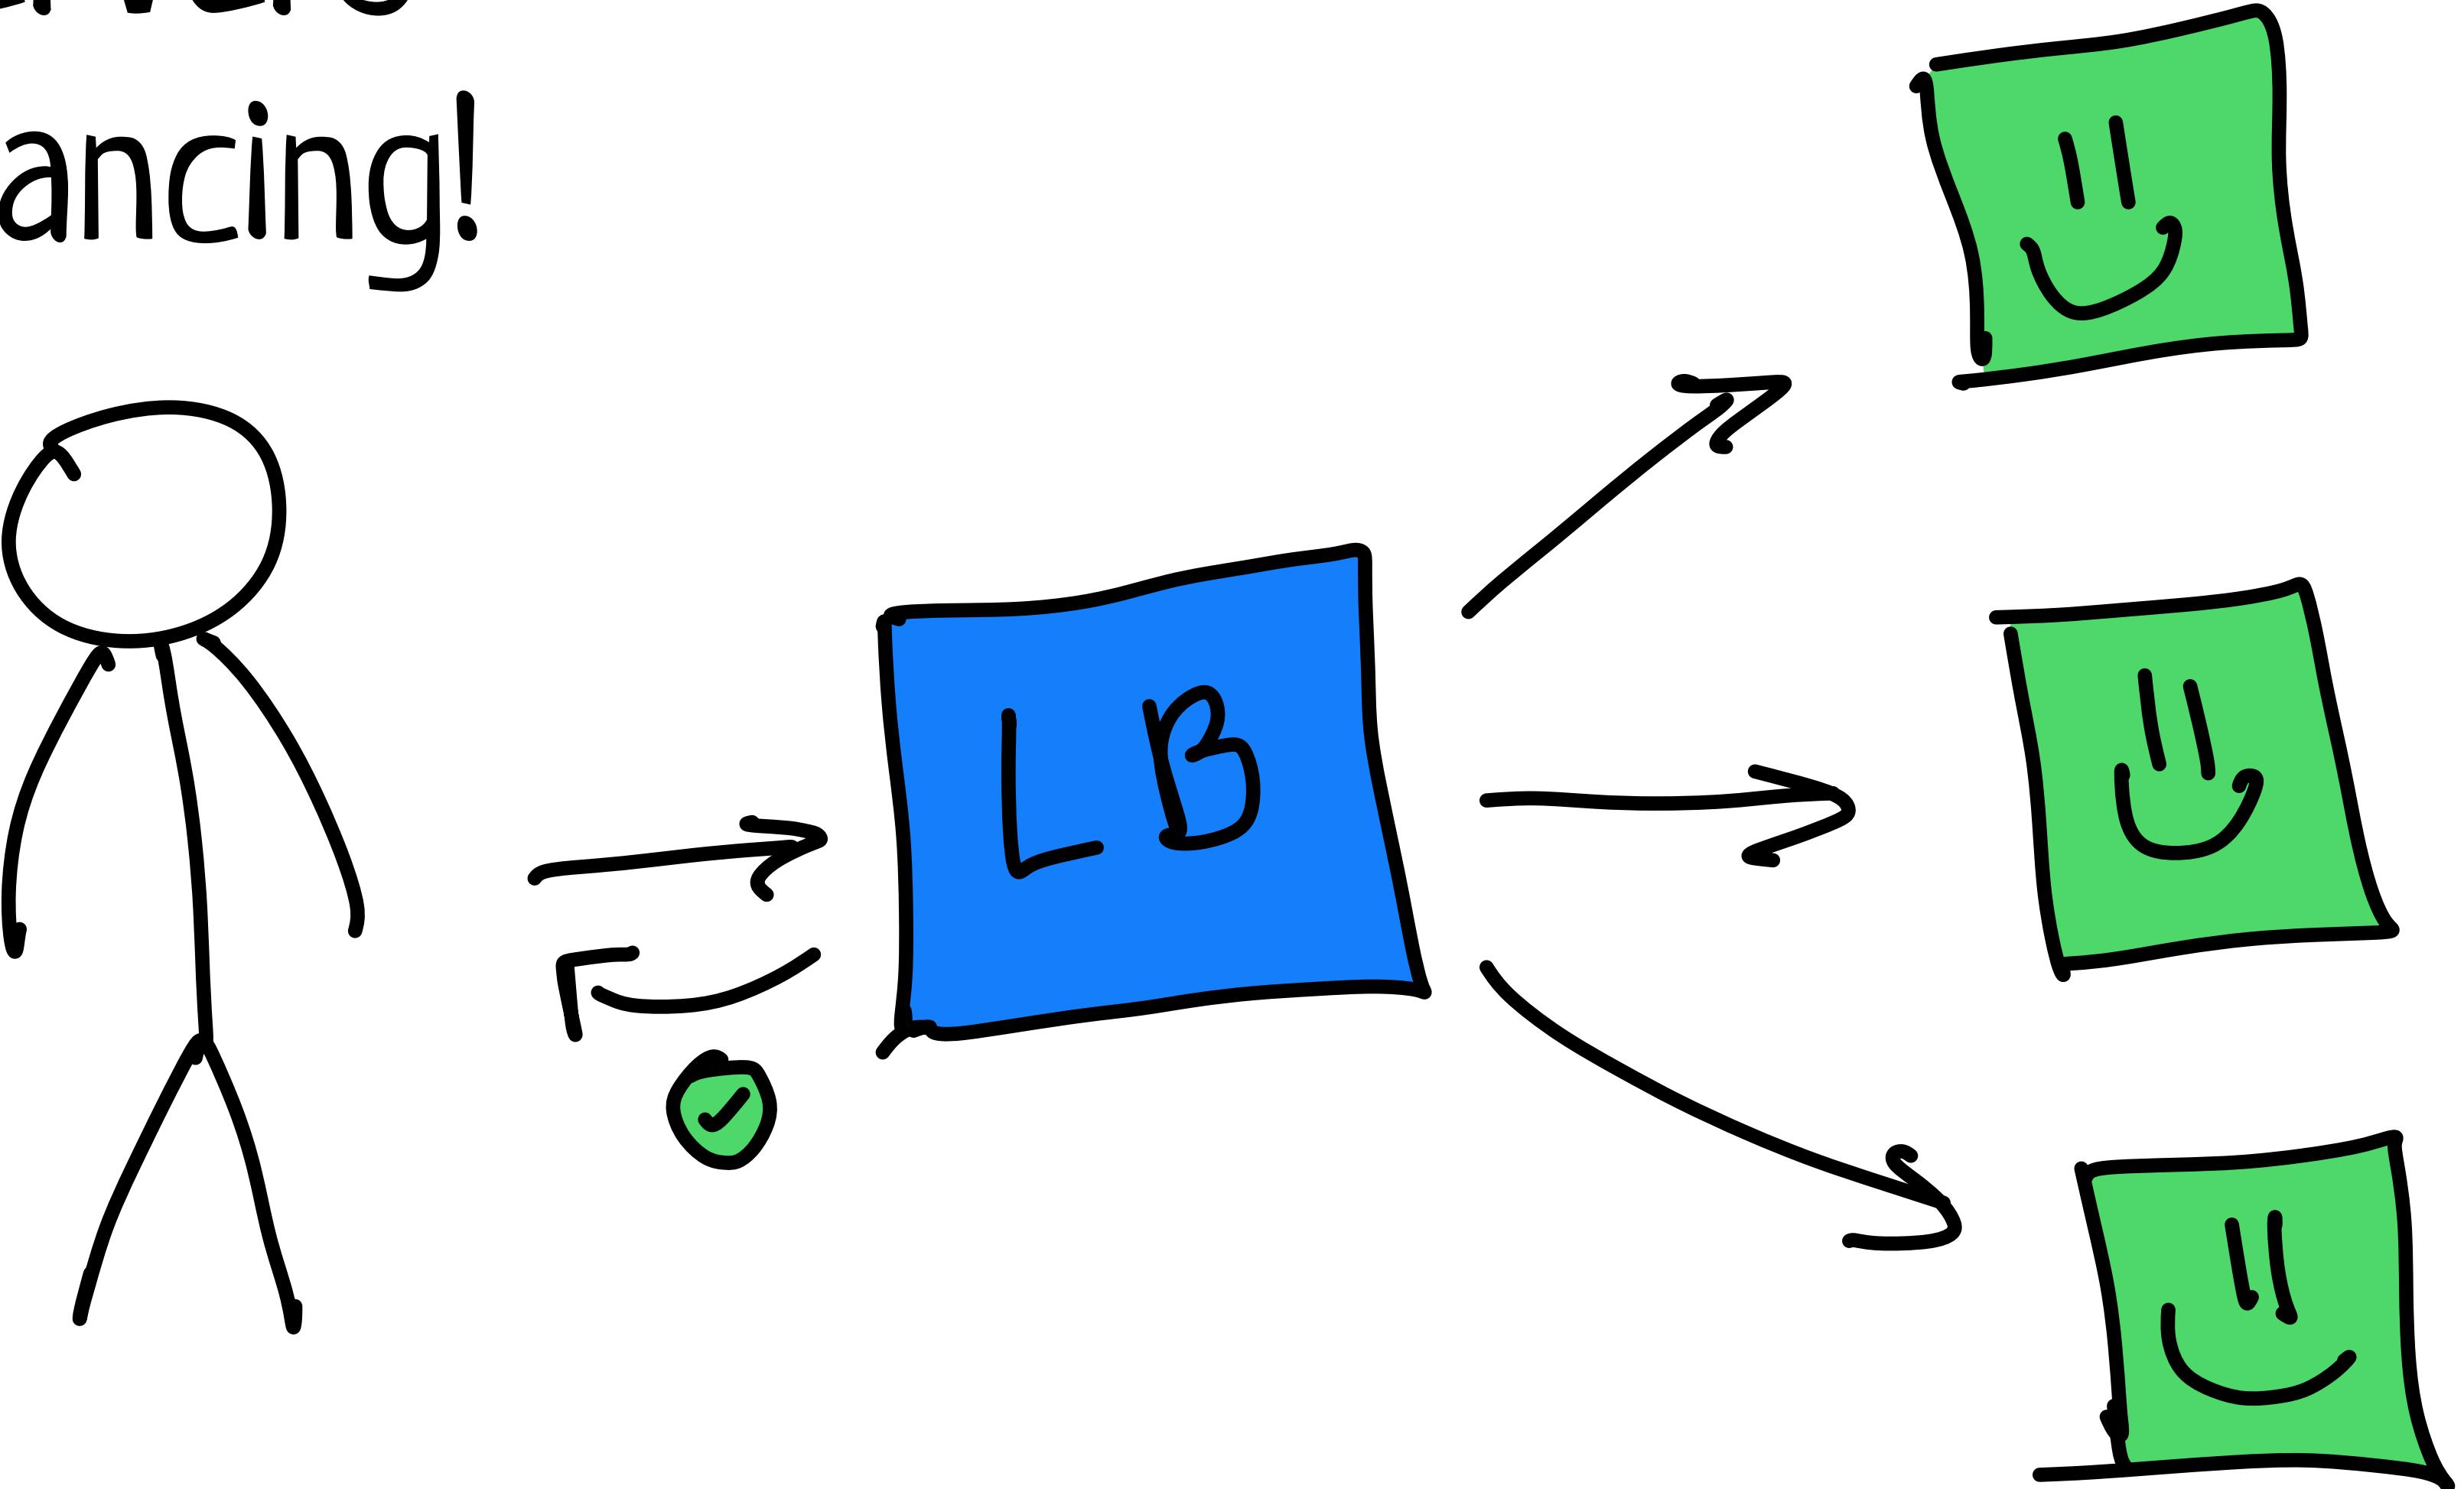
If we have only one server and no load balancing



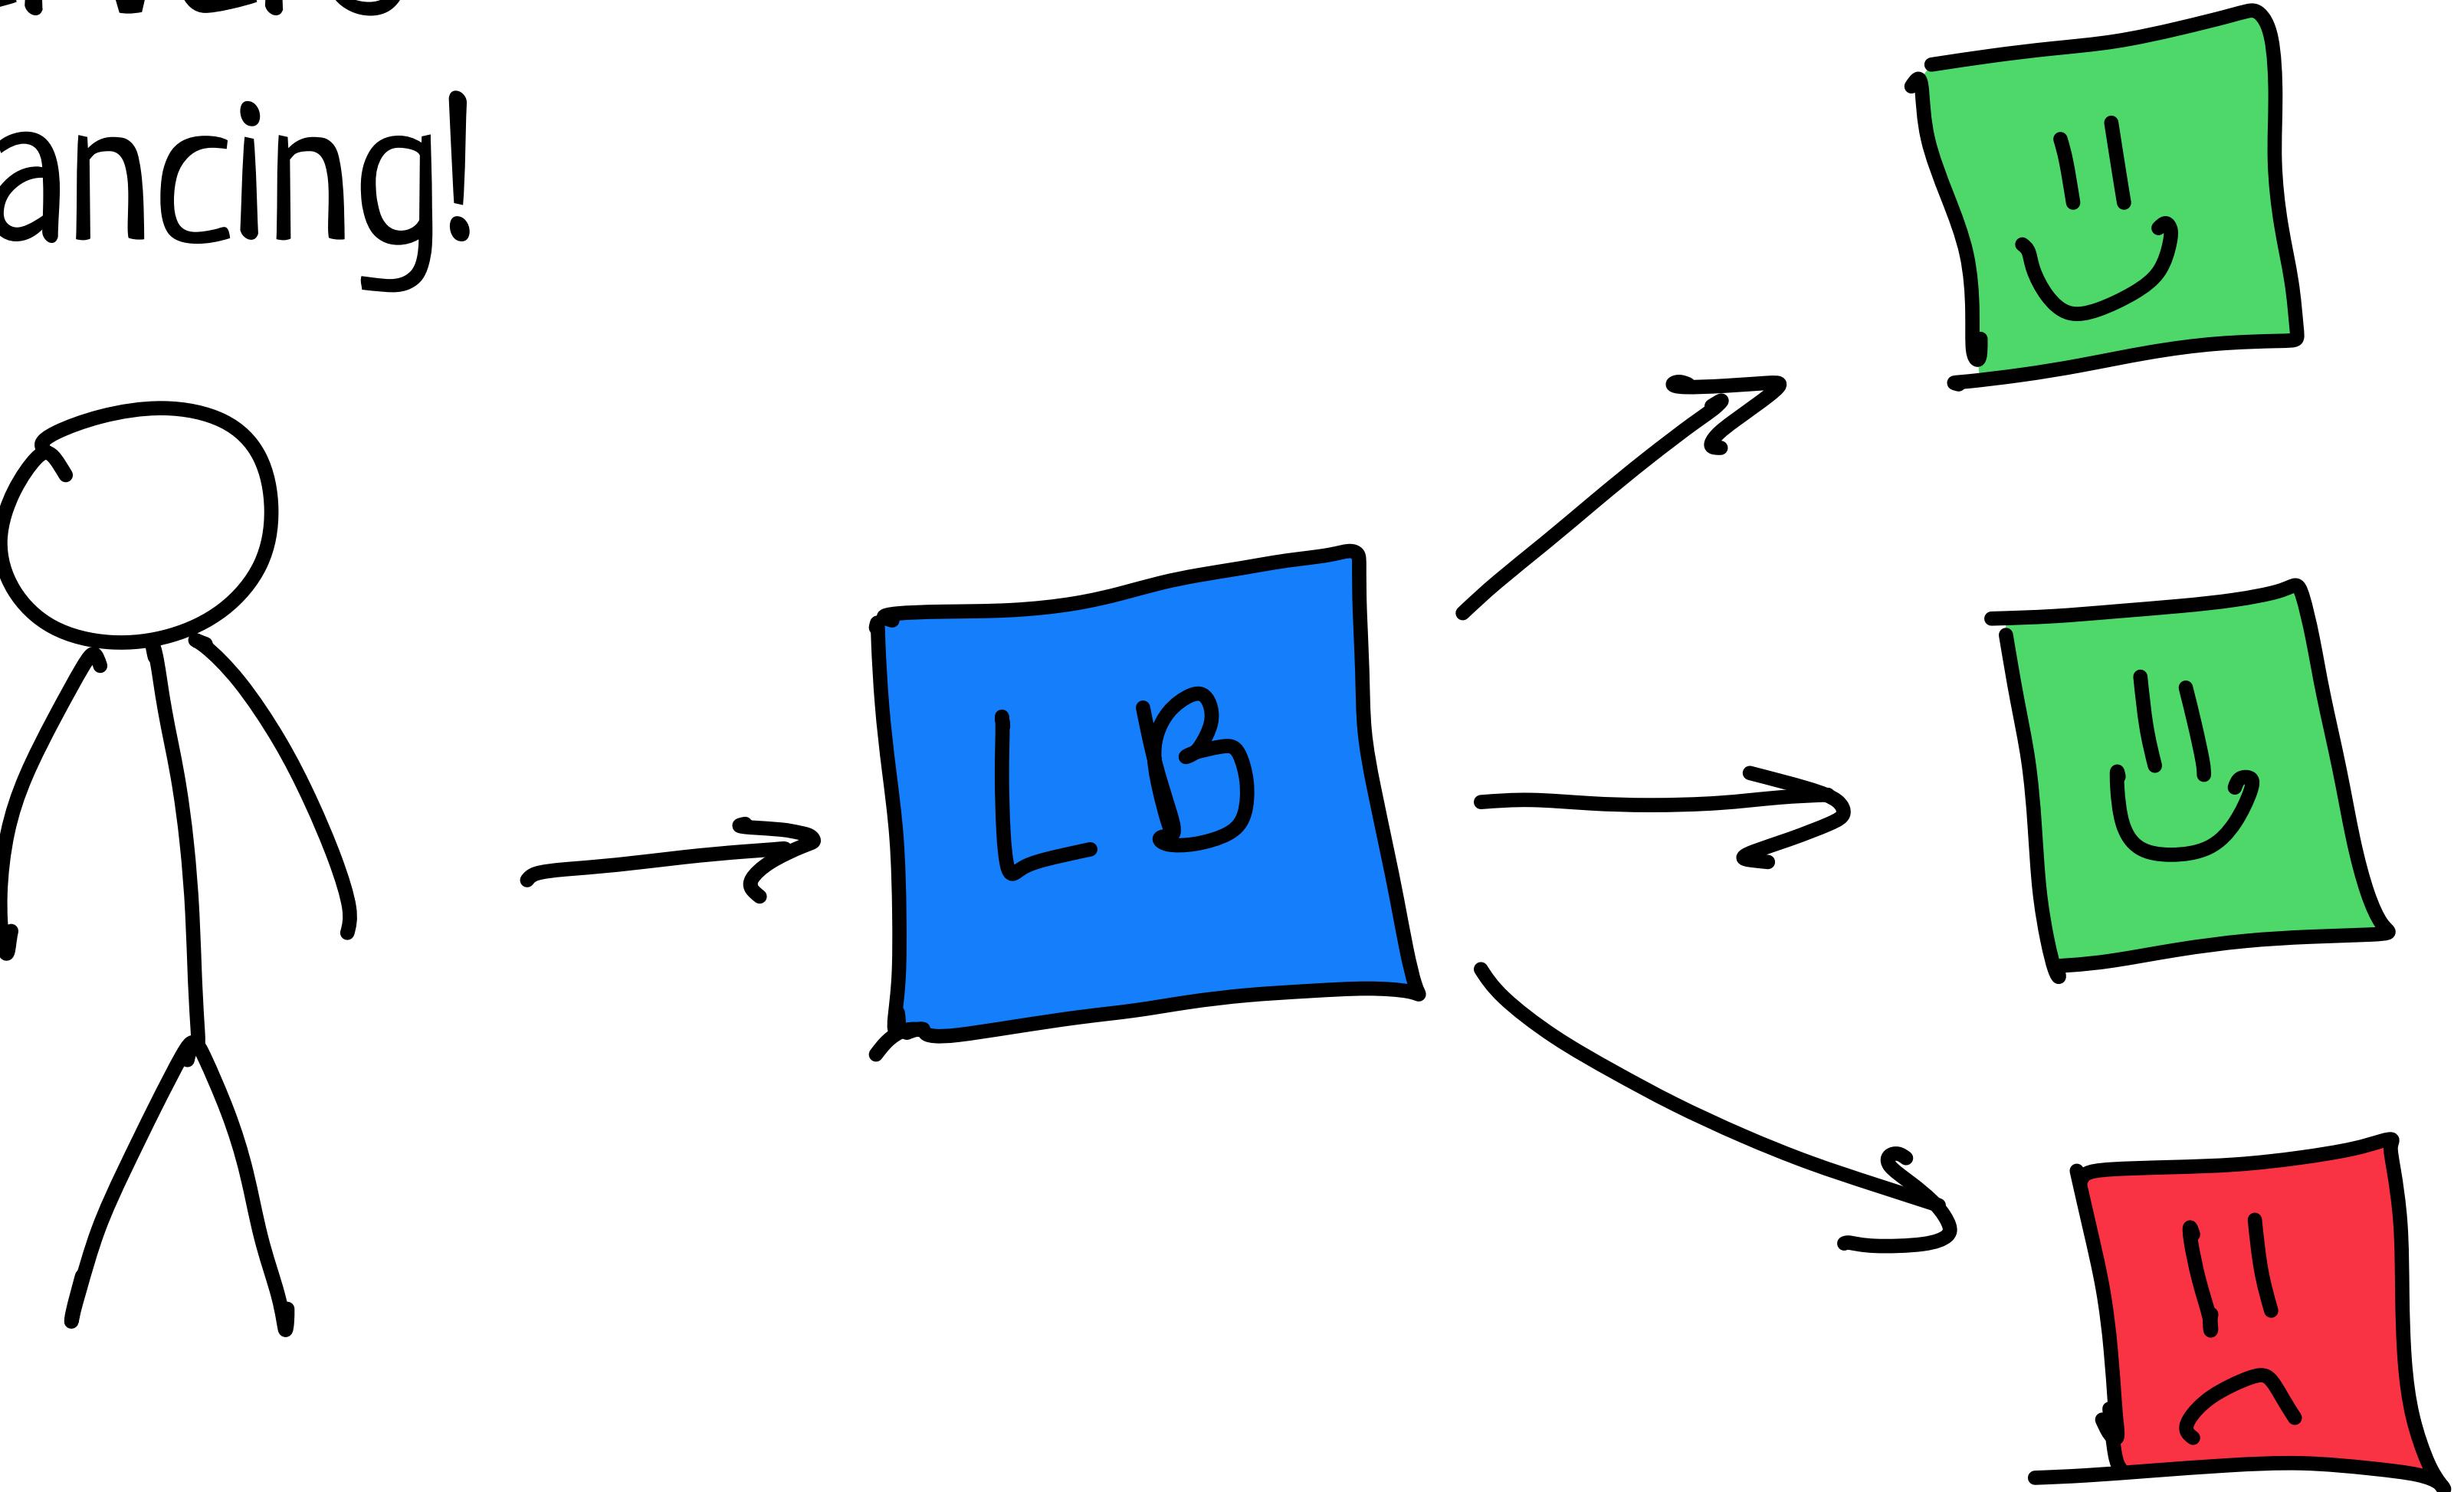
More servers More balancing!



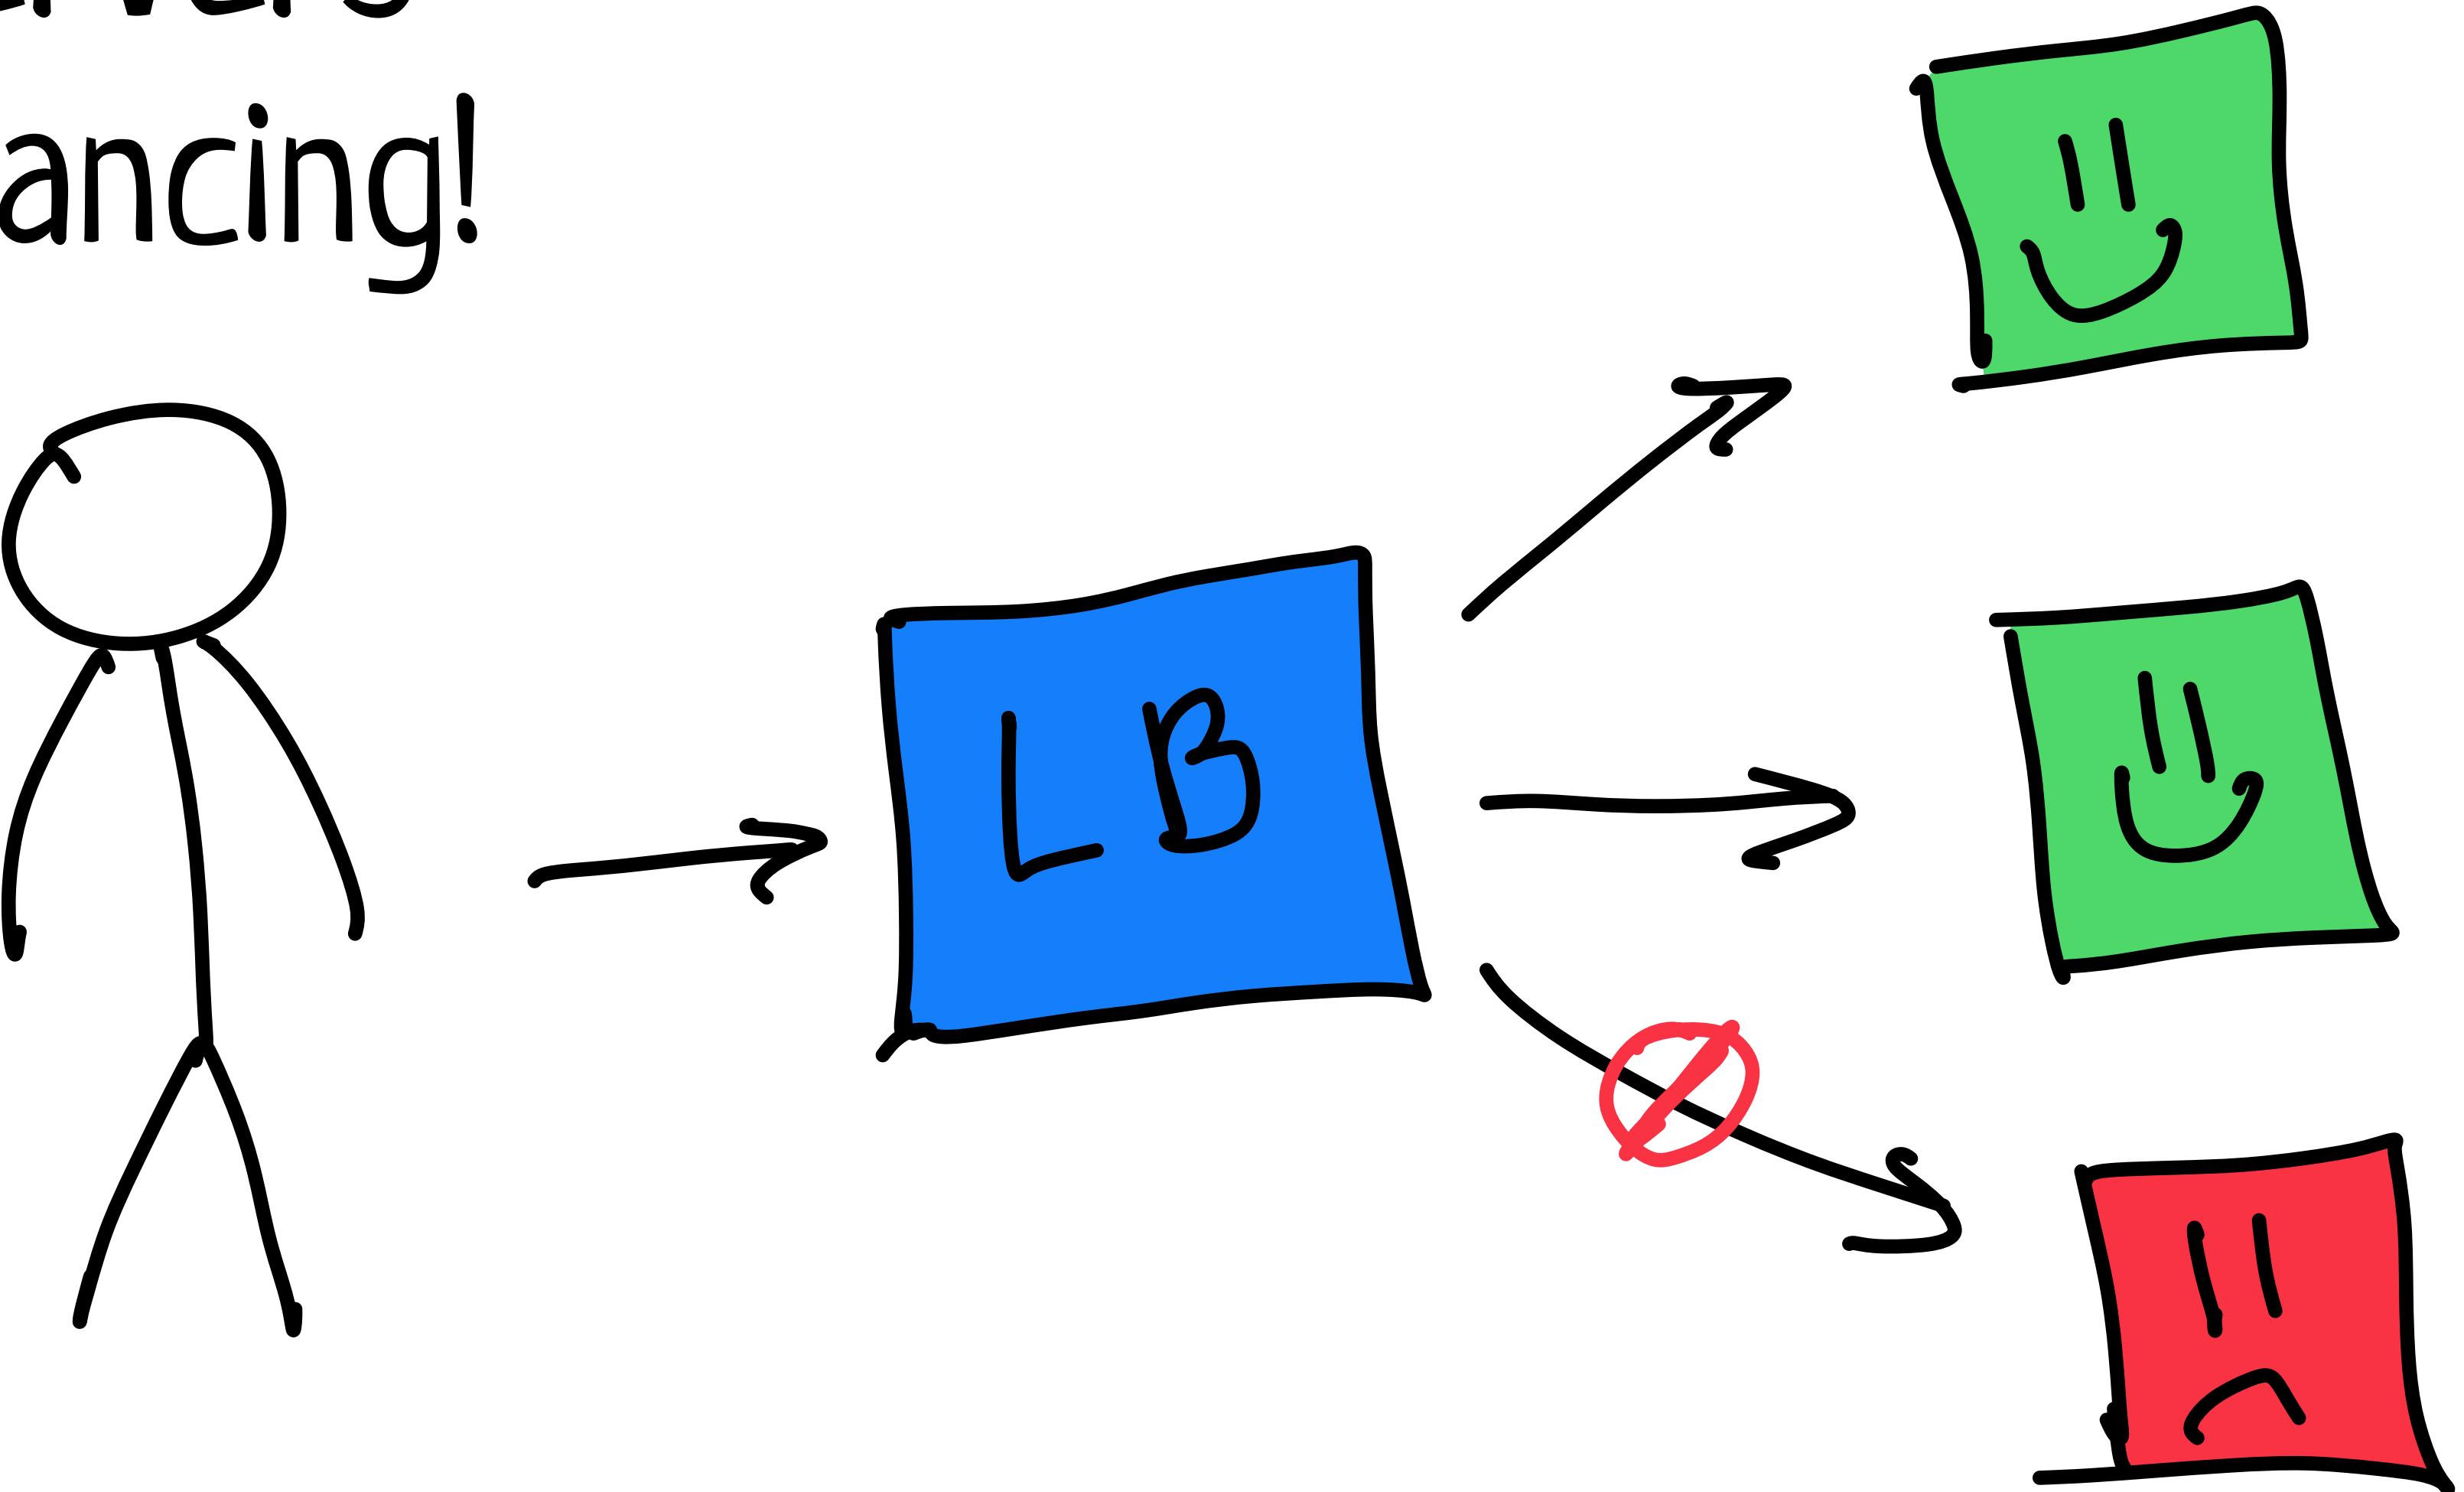
More servers More balancing!



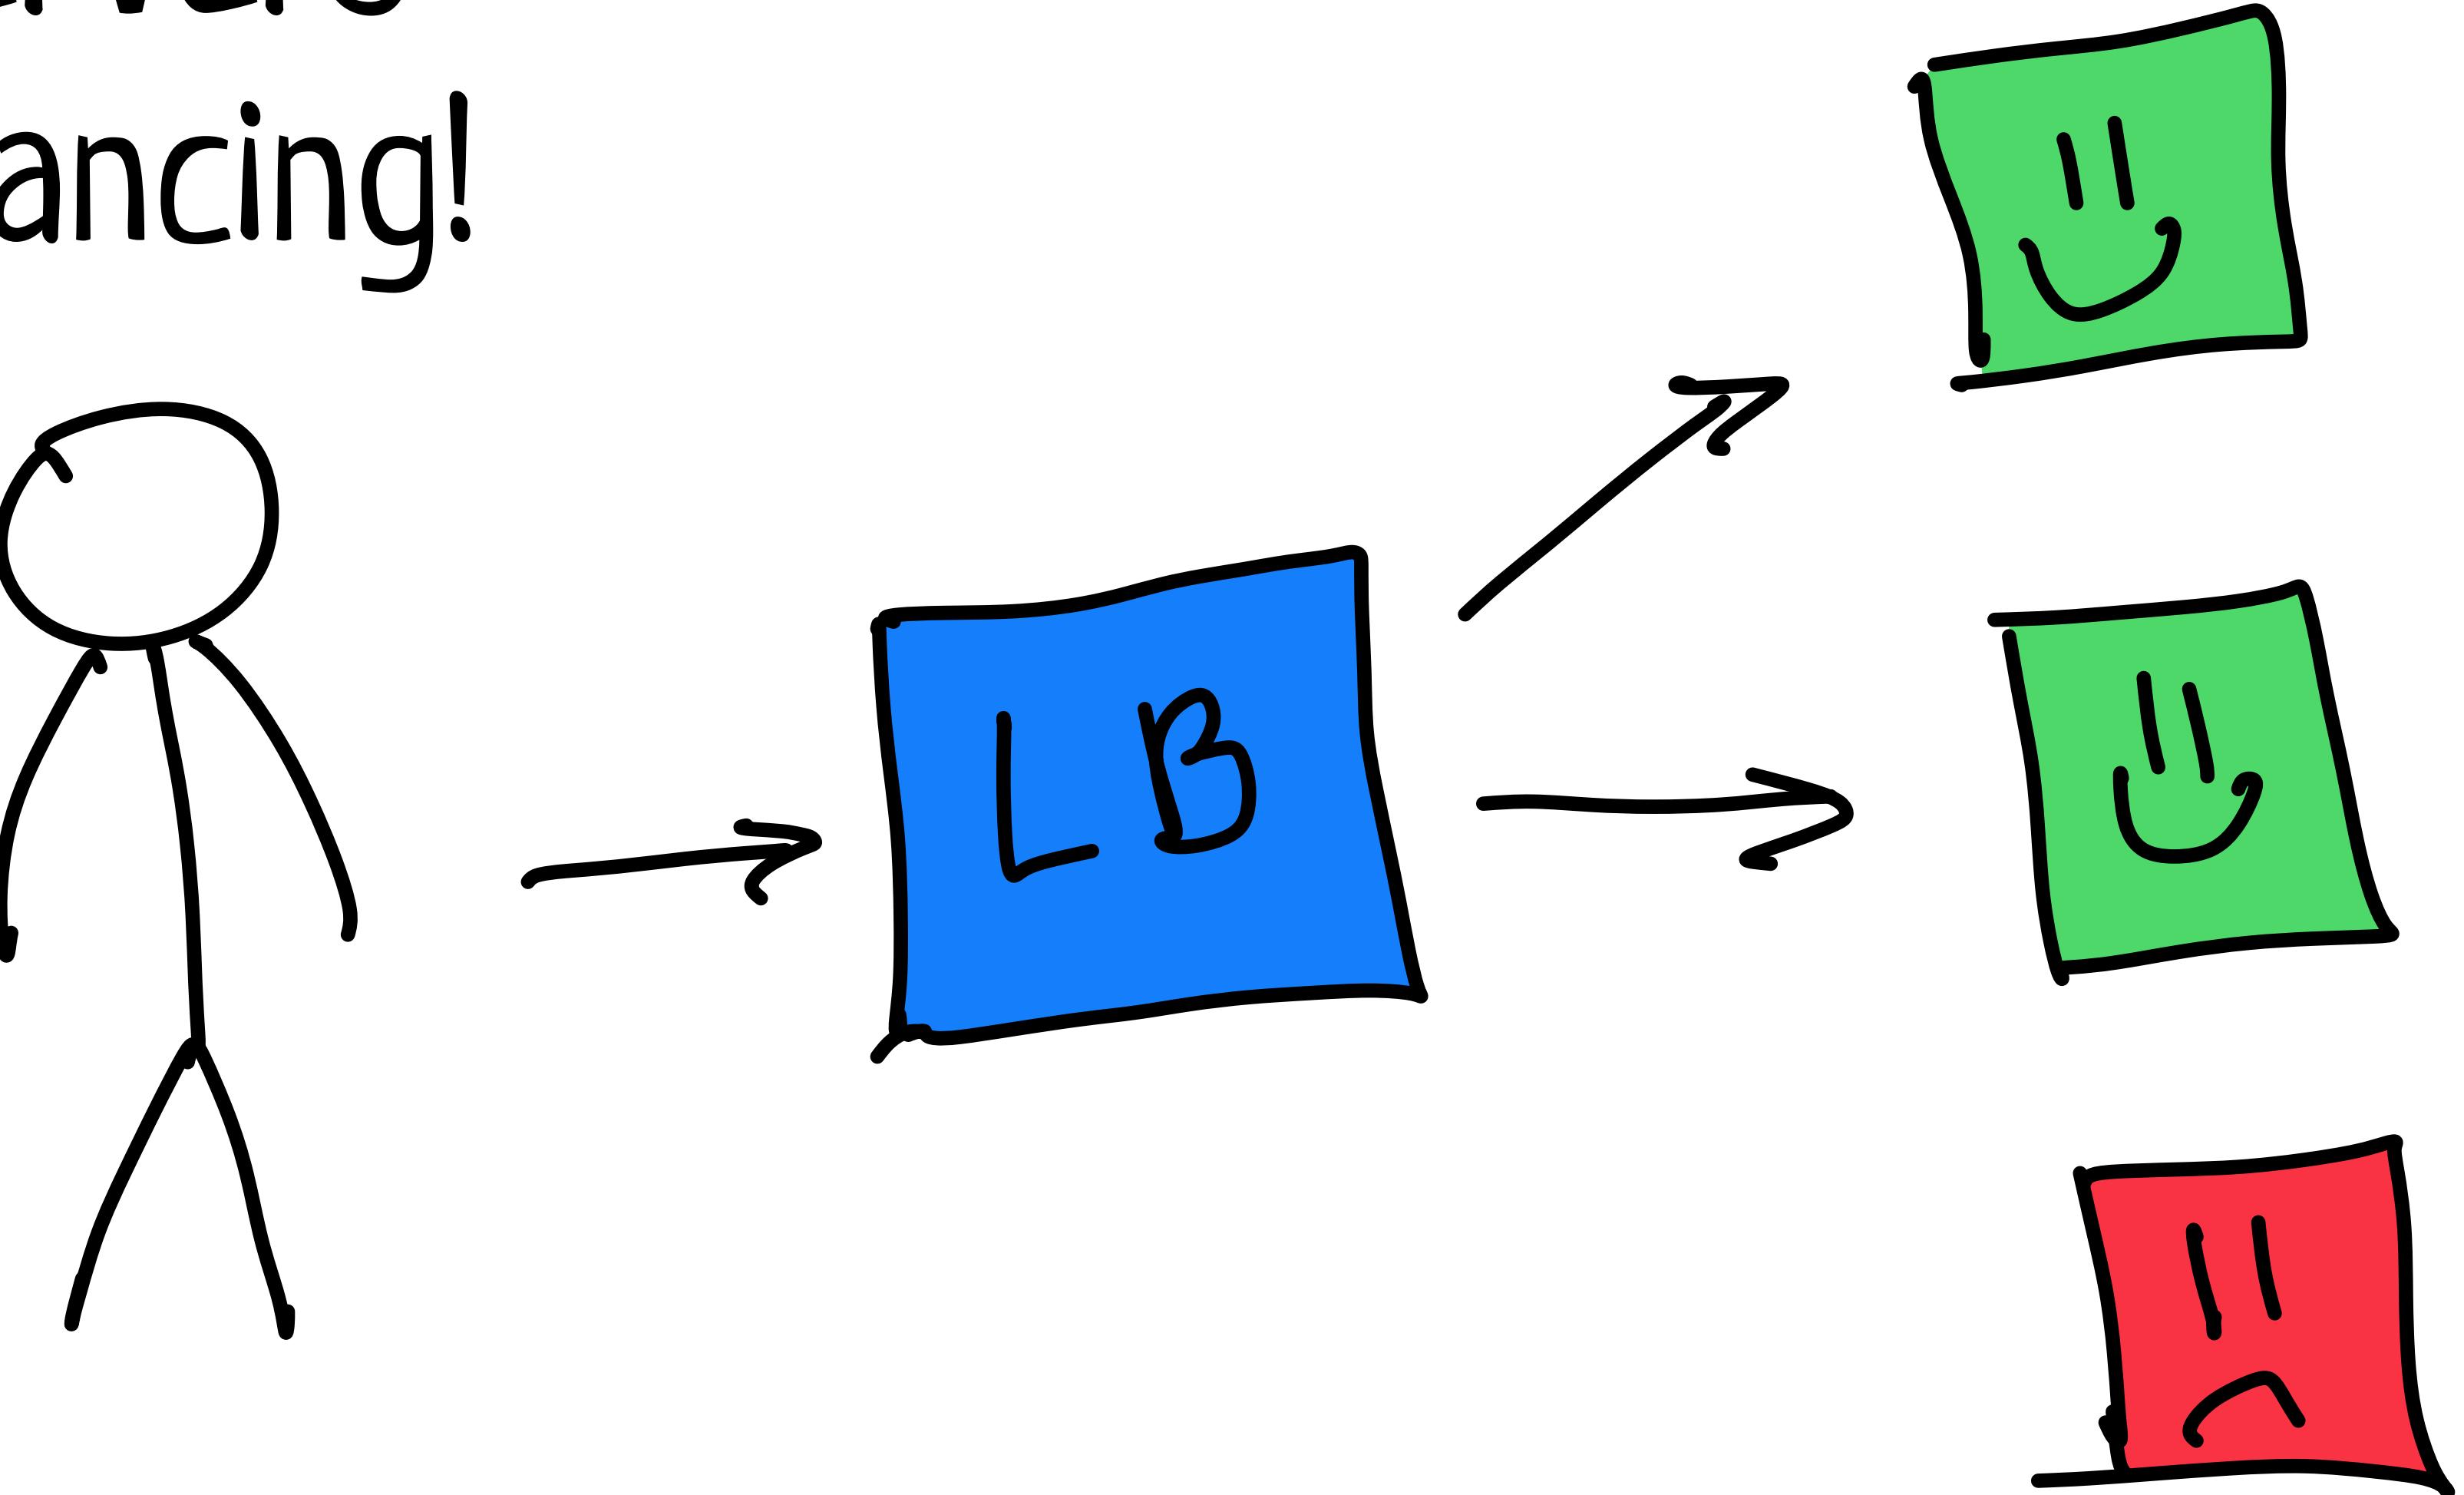
More servers More balancing!



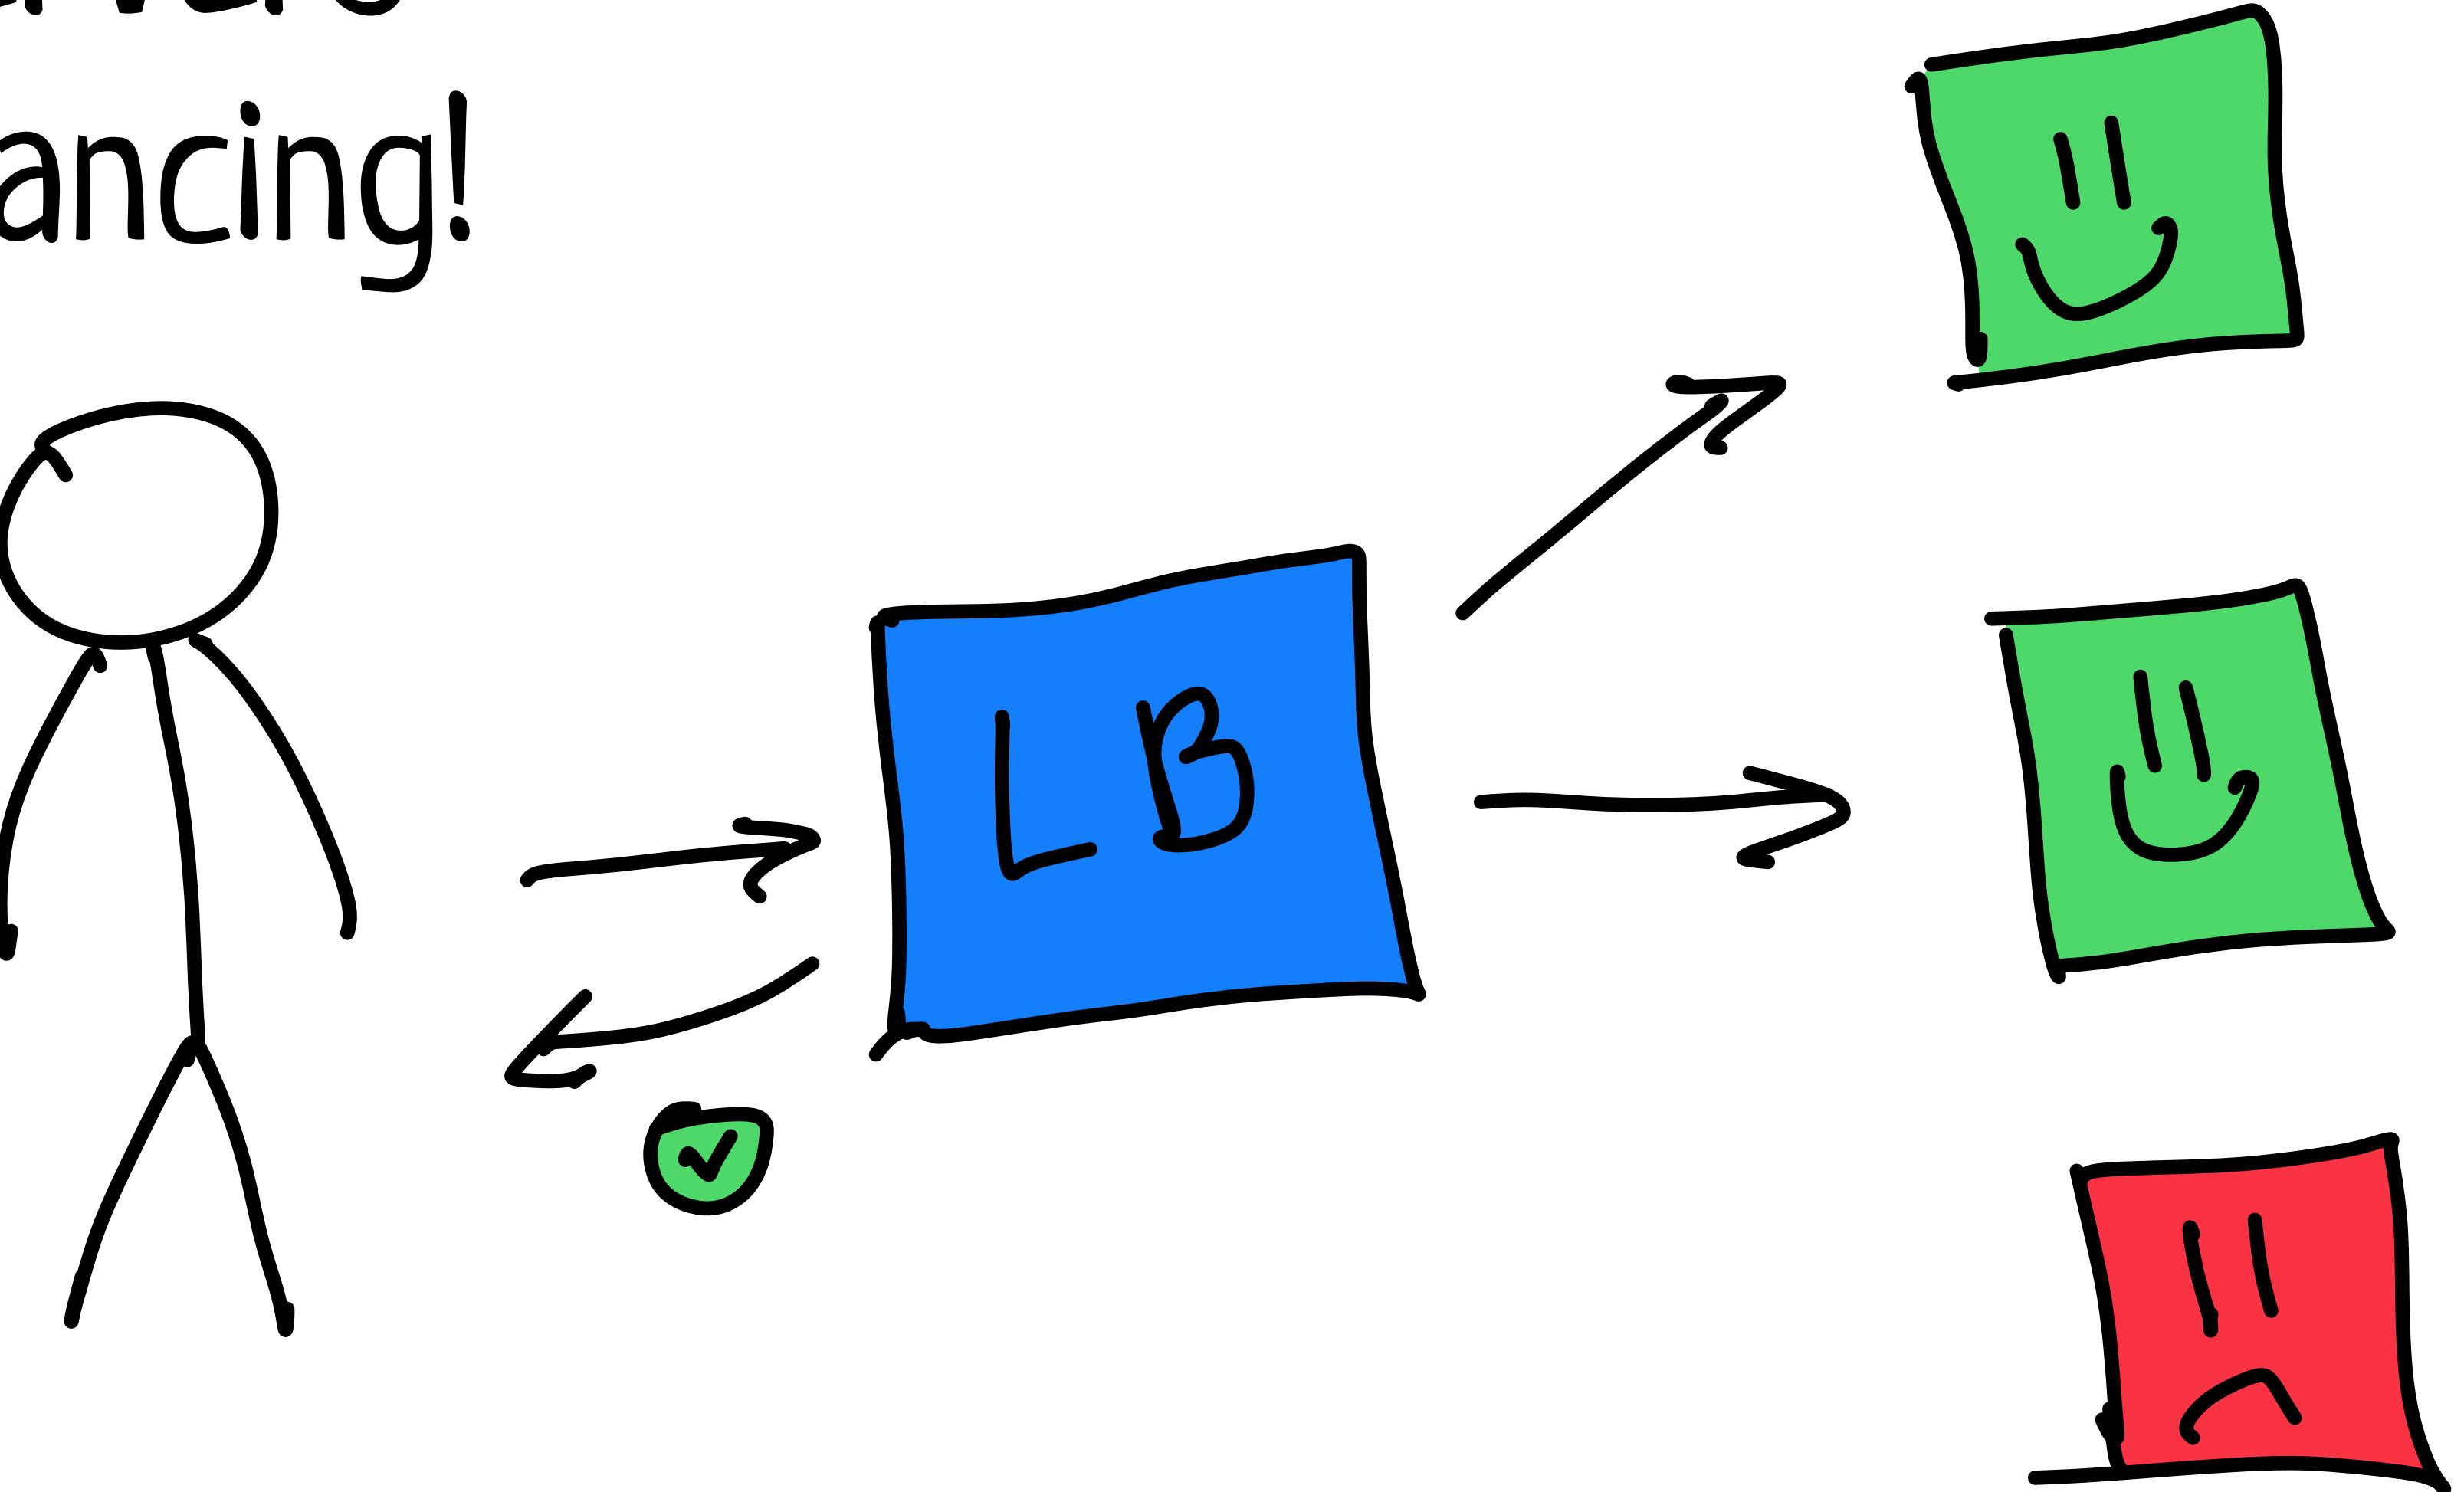
More servers More balancing!



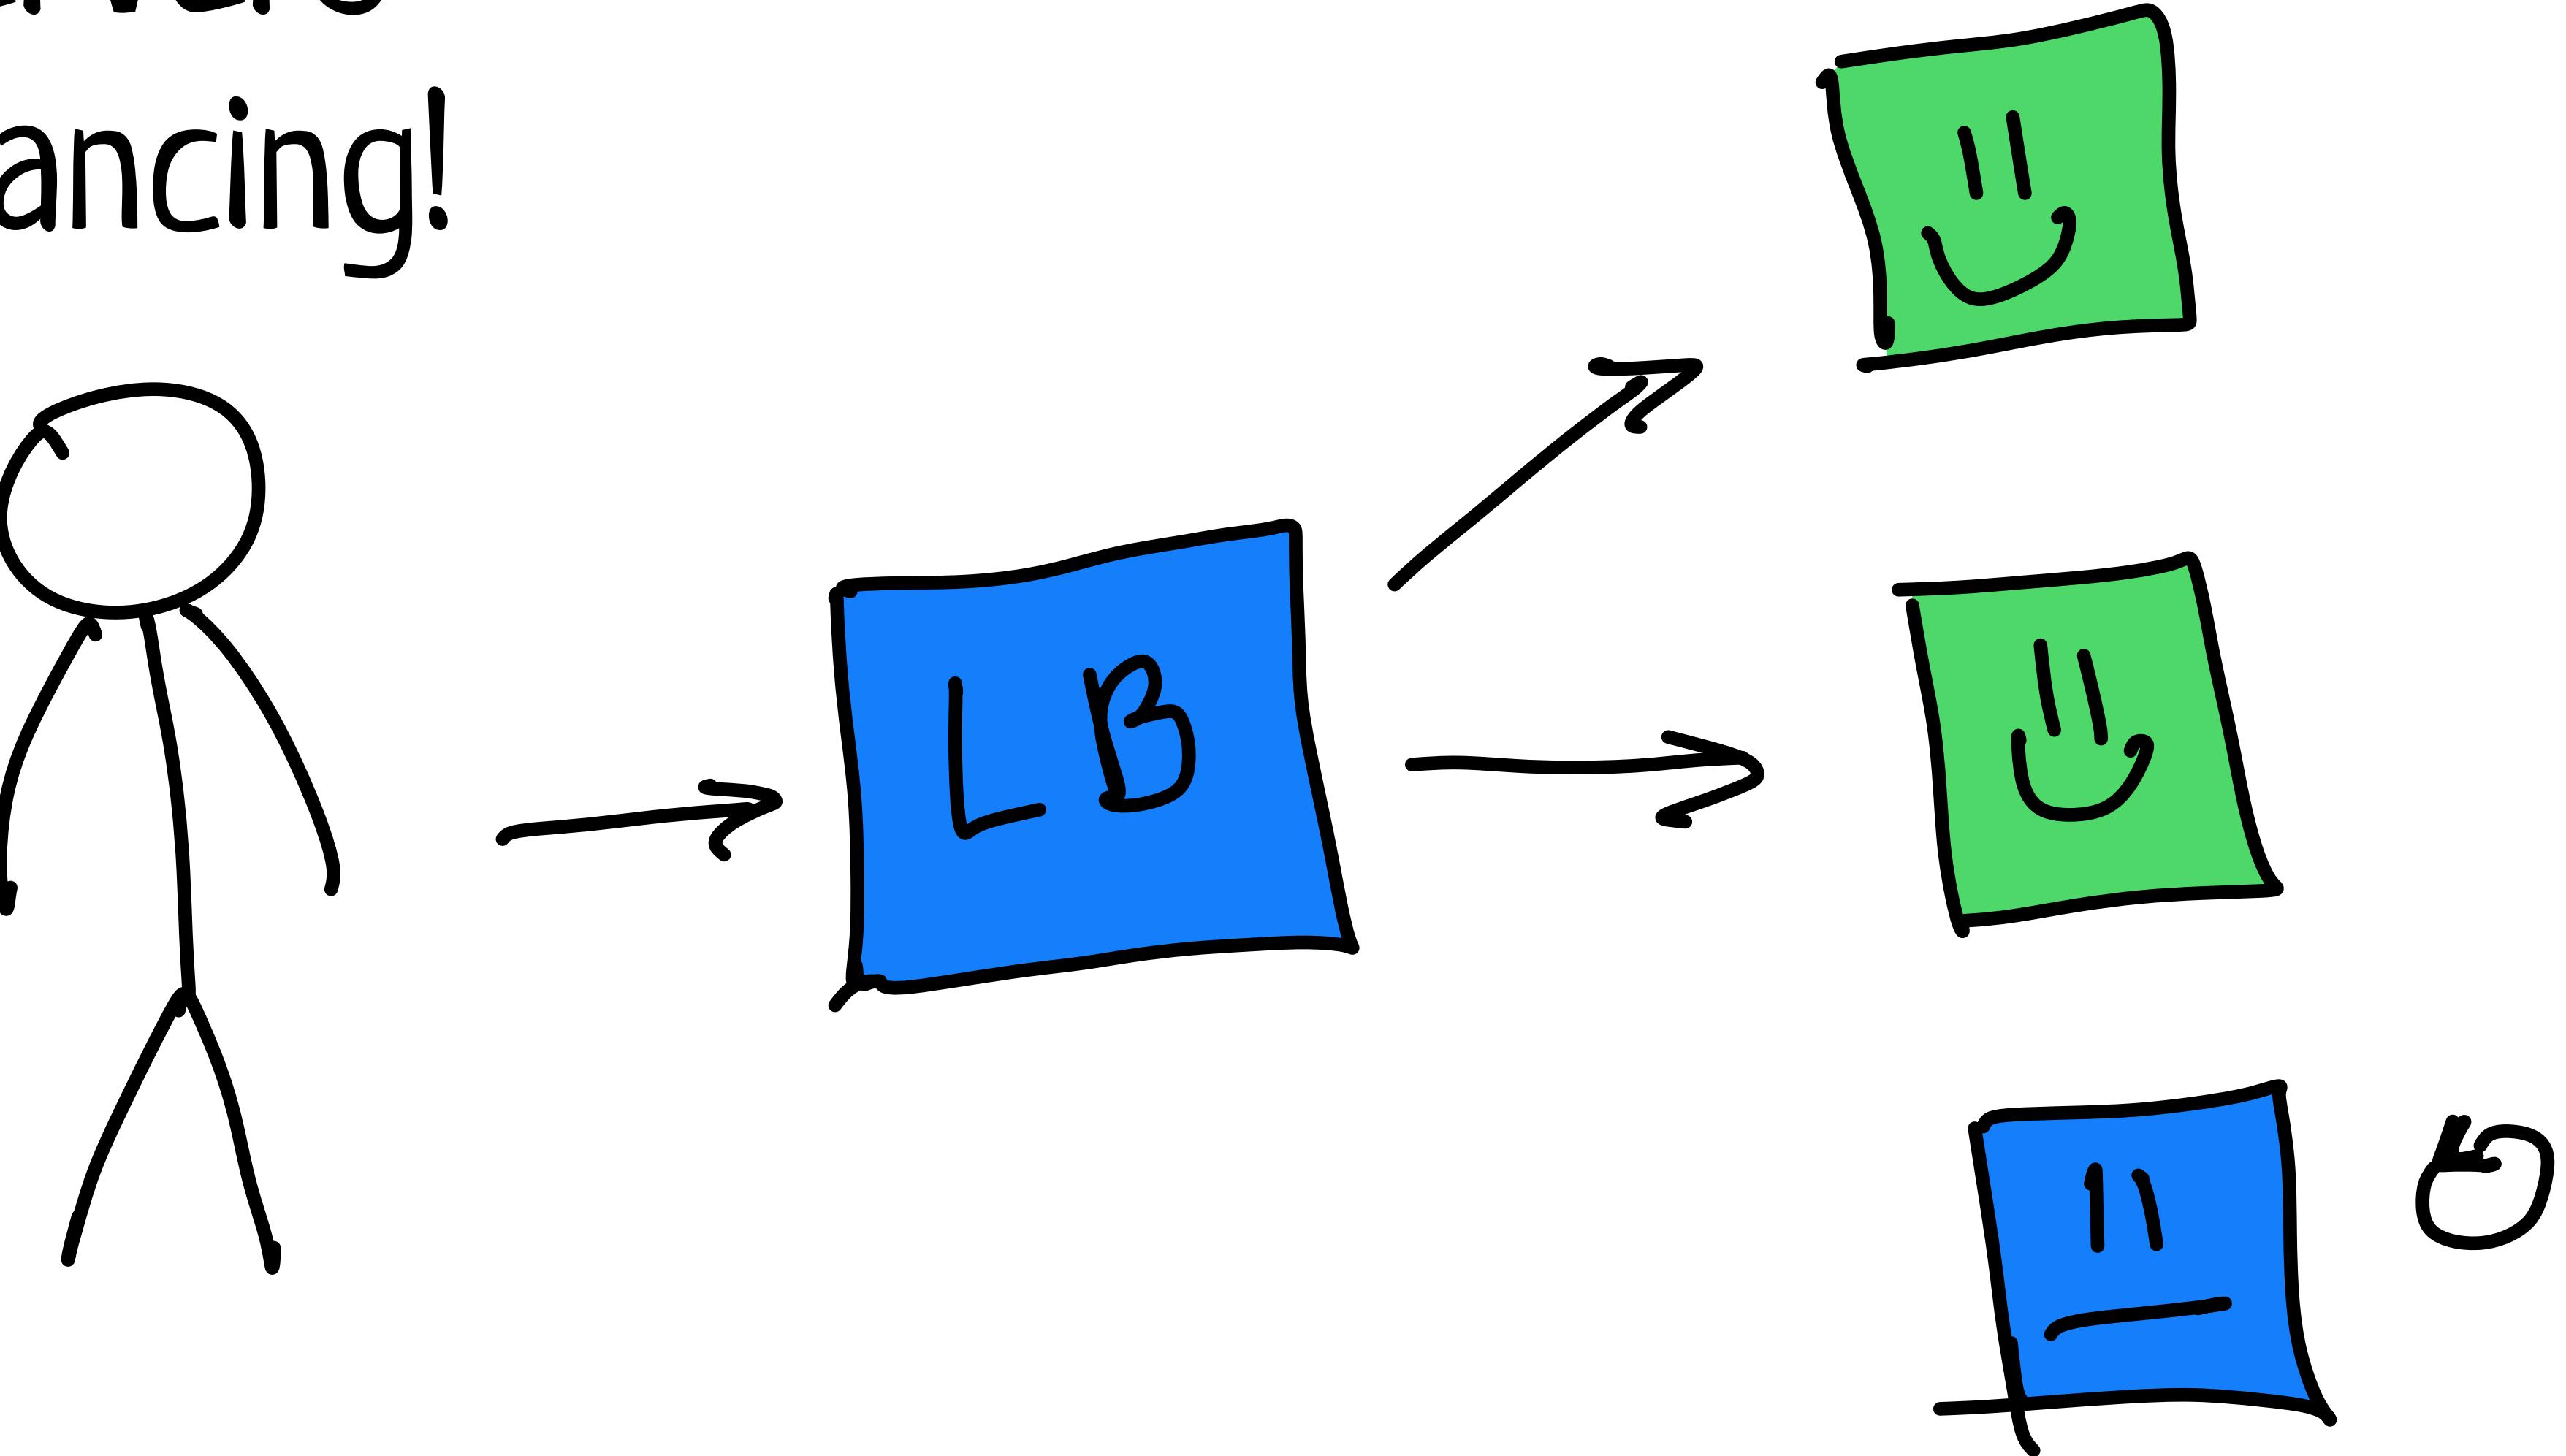
More servers More balancing!



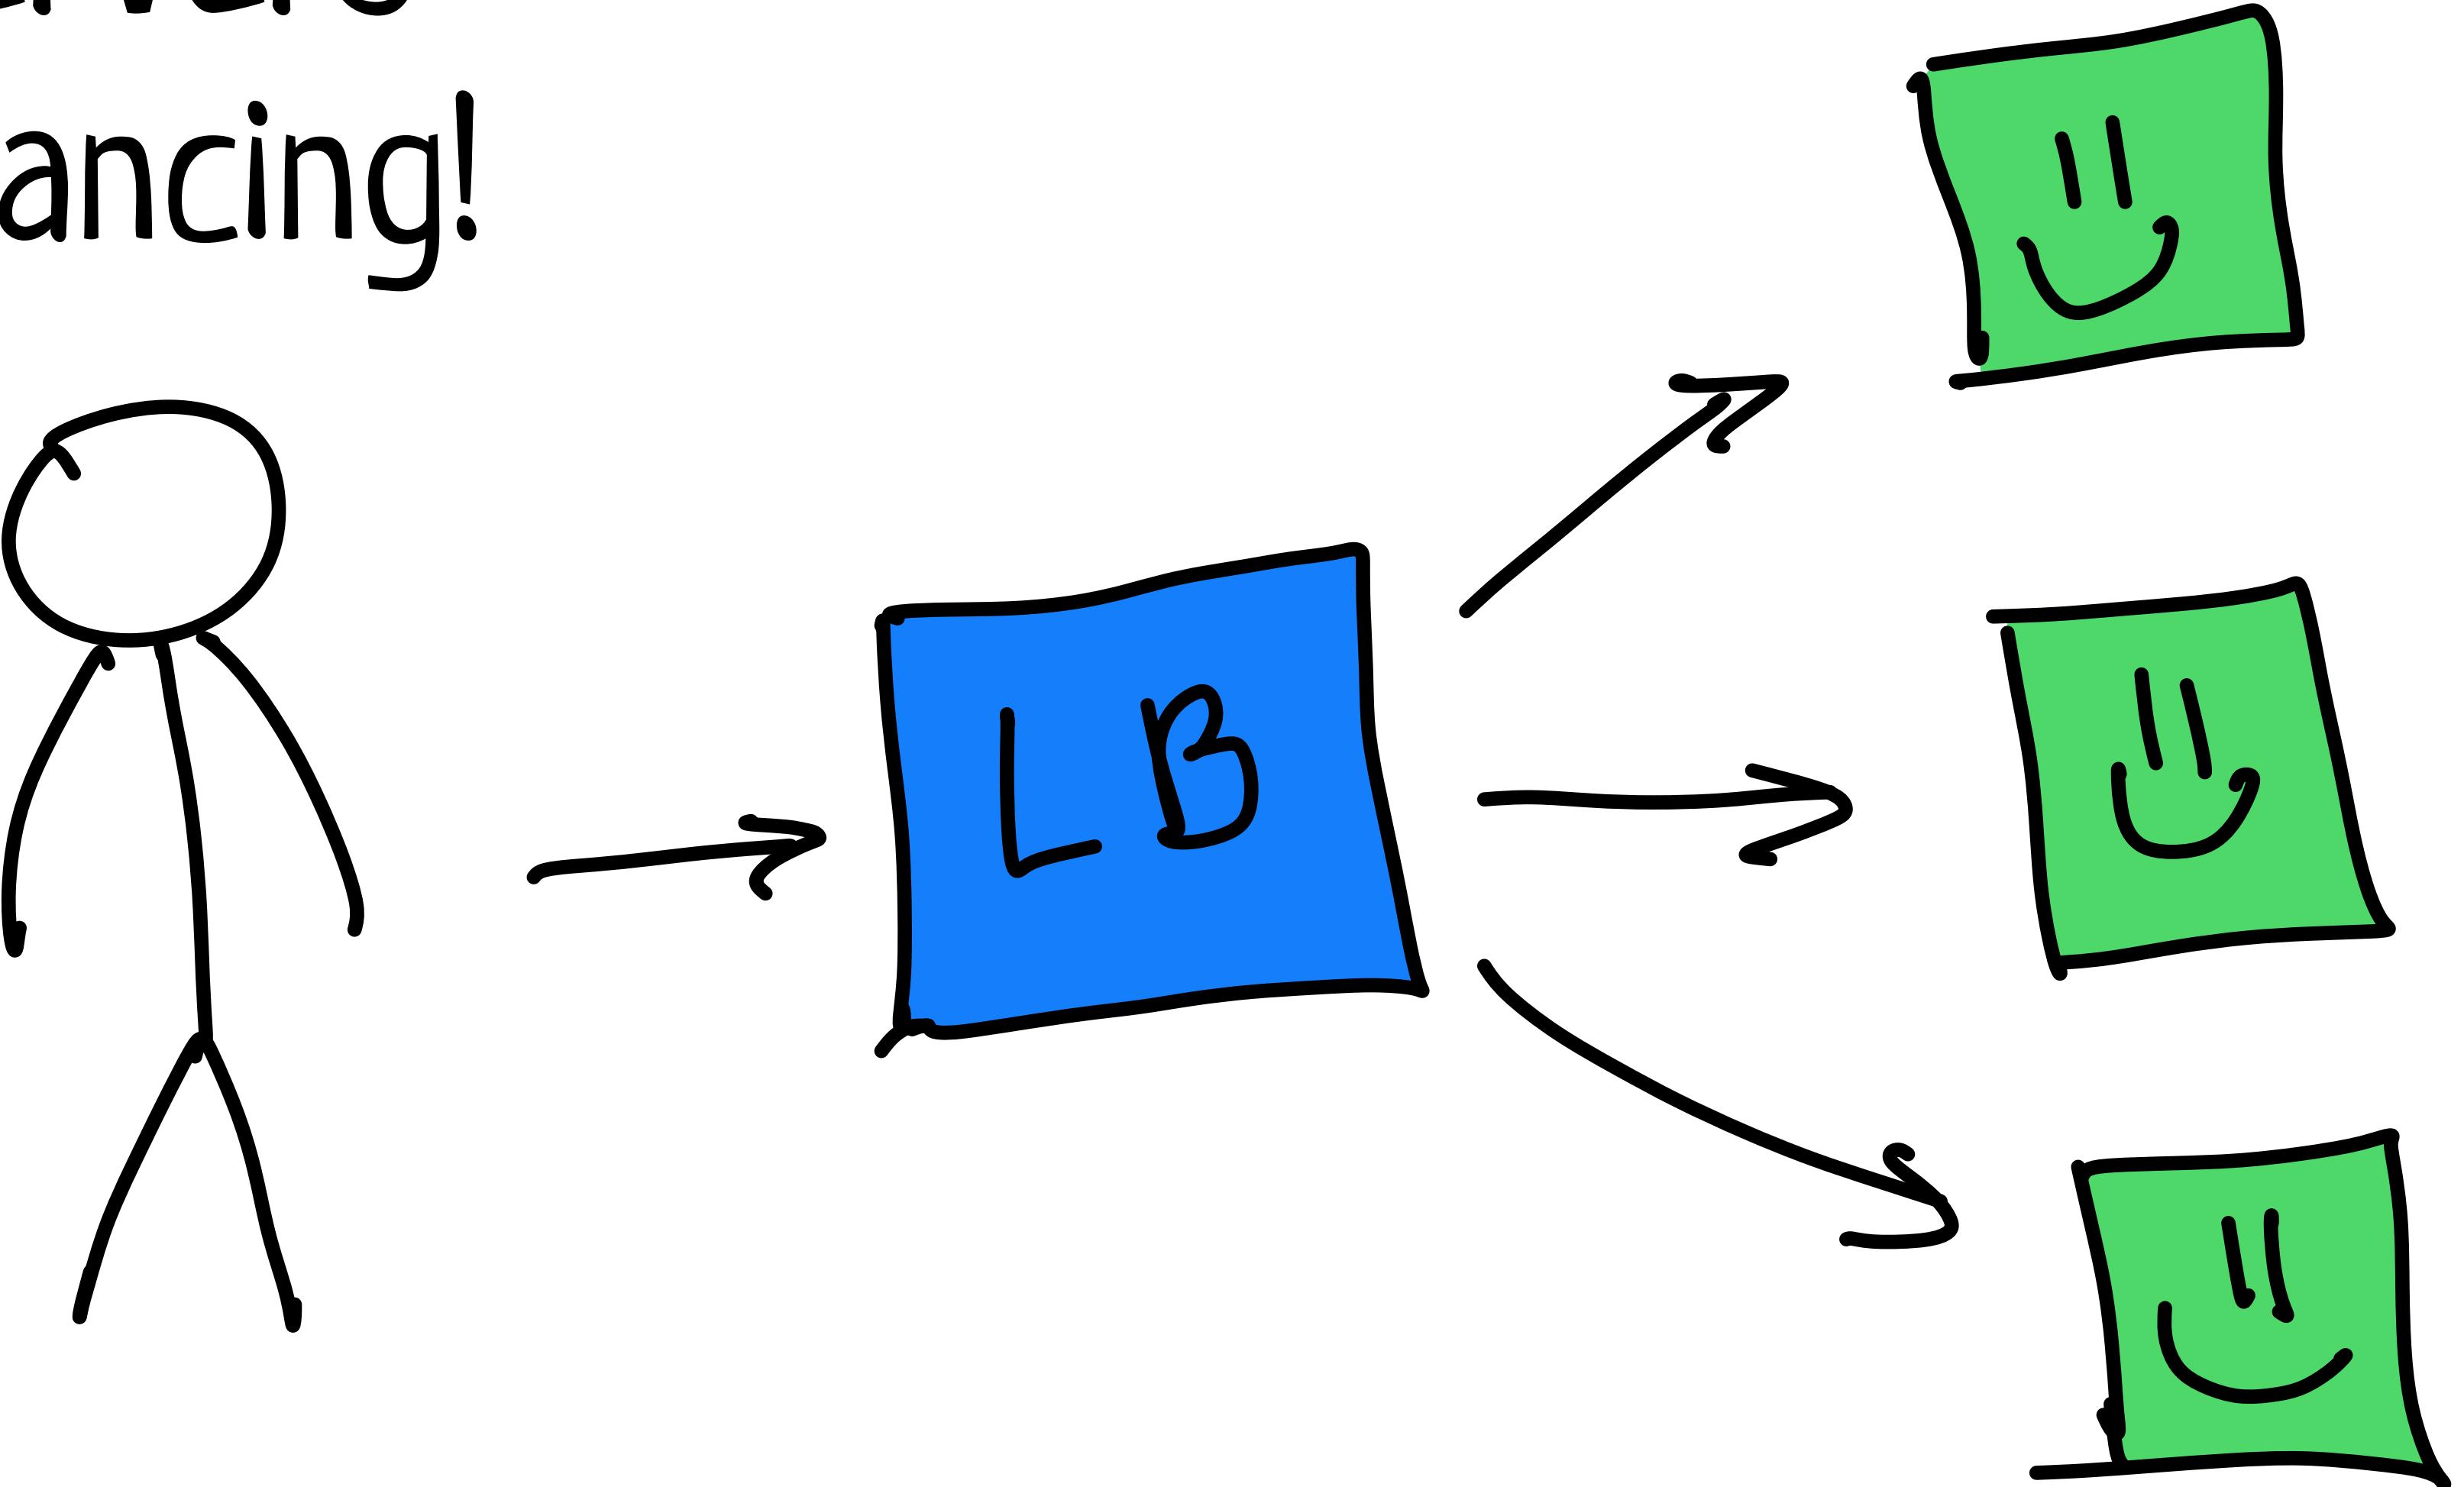
More servers More balancing!



More servers More balancing!



More servers More balancing!



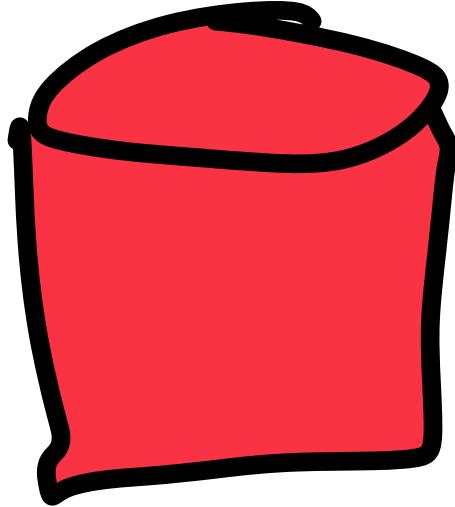
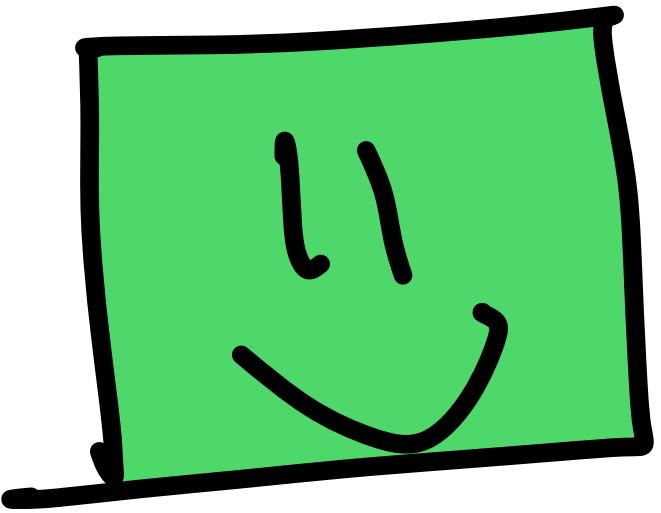
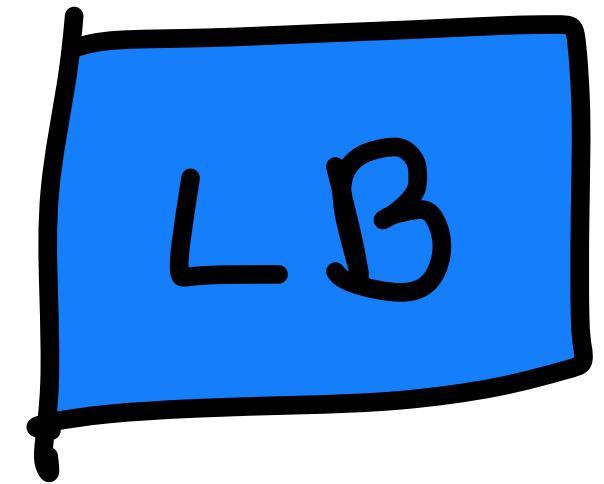
Load Balancer

- NGINX
- Azure Traffic Manager
- AWS Elastic Load Balancing
- GC Load Balancing

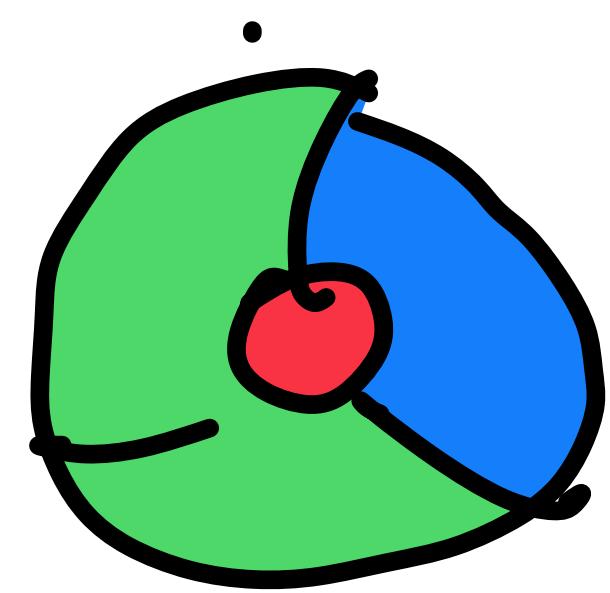
UUID

Universal Unique Identifier/Atomic Unique Identifiers

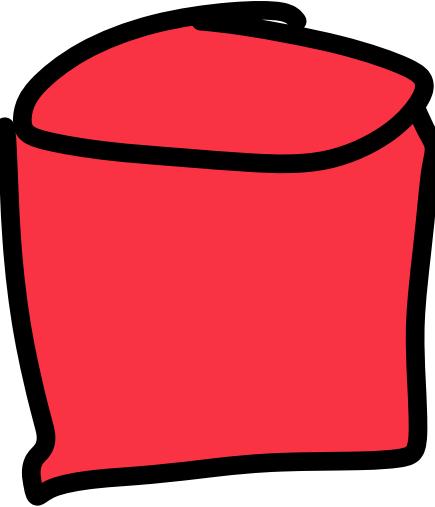
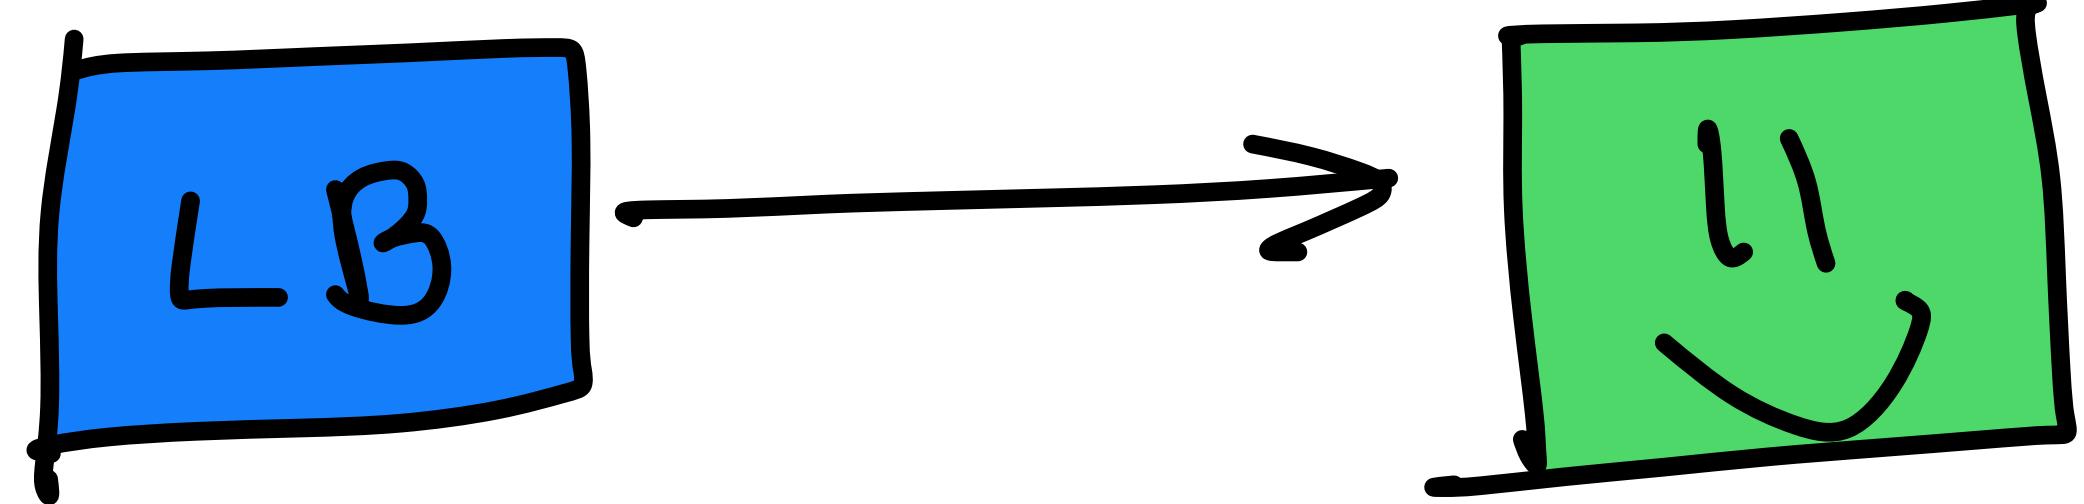
No UUIDs
Per transaction



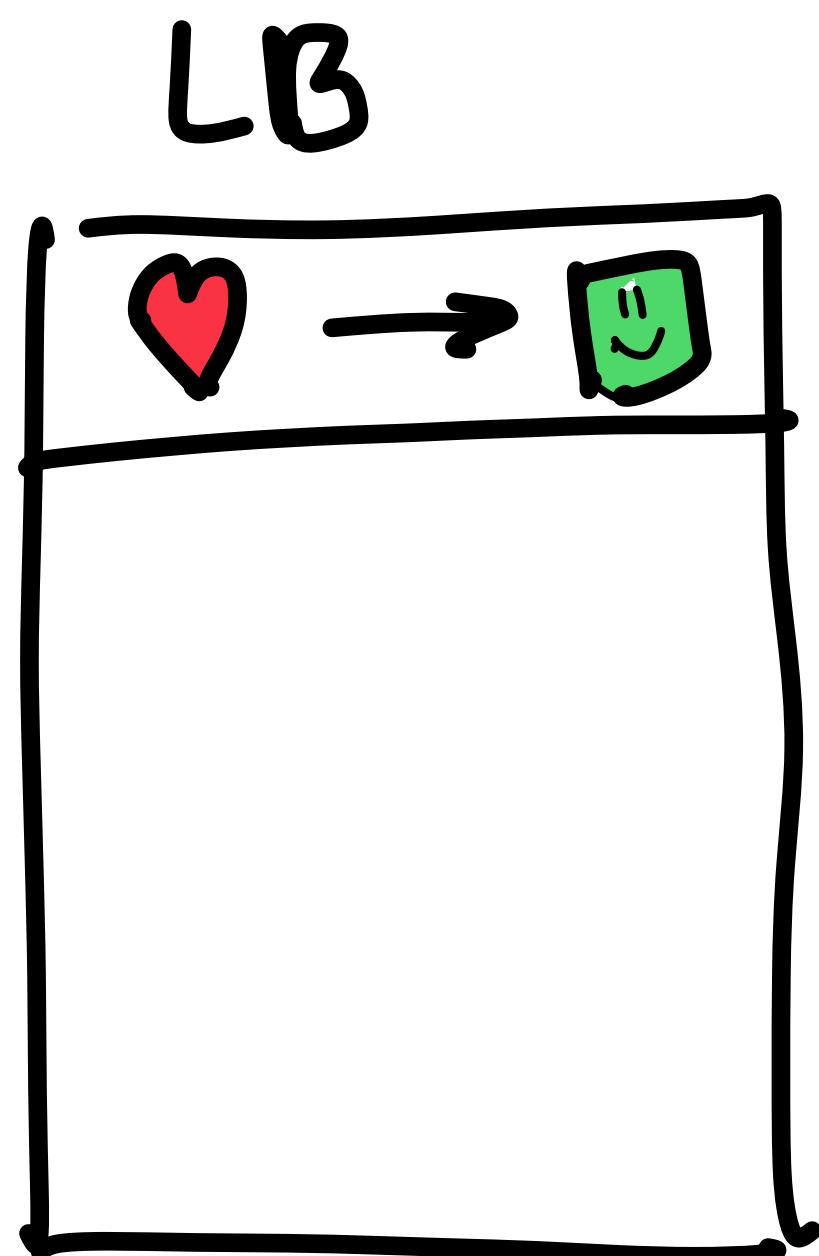
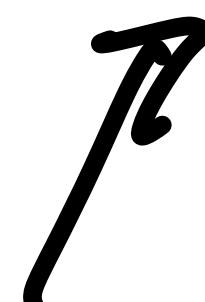
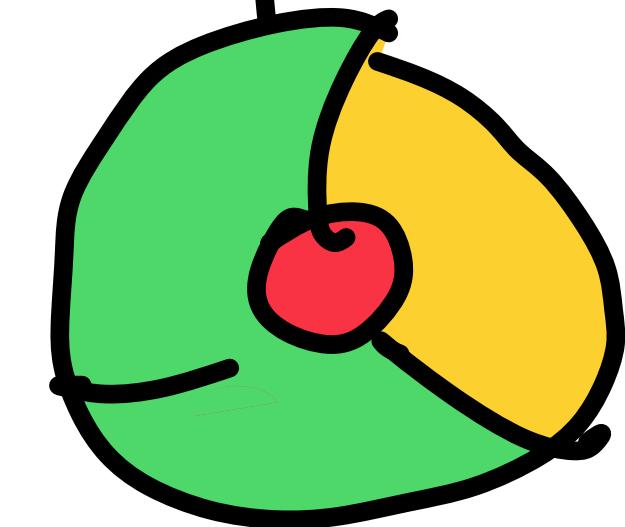
Logs



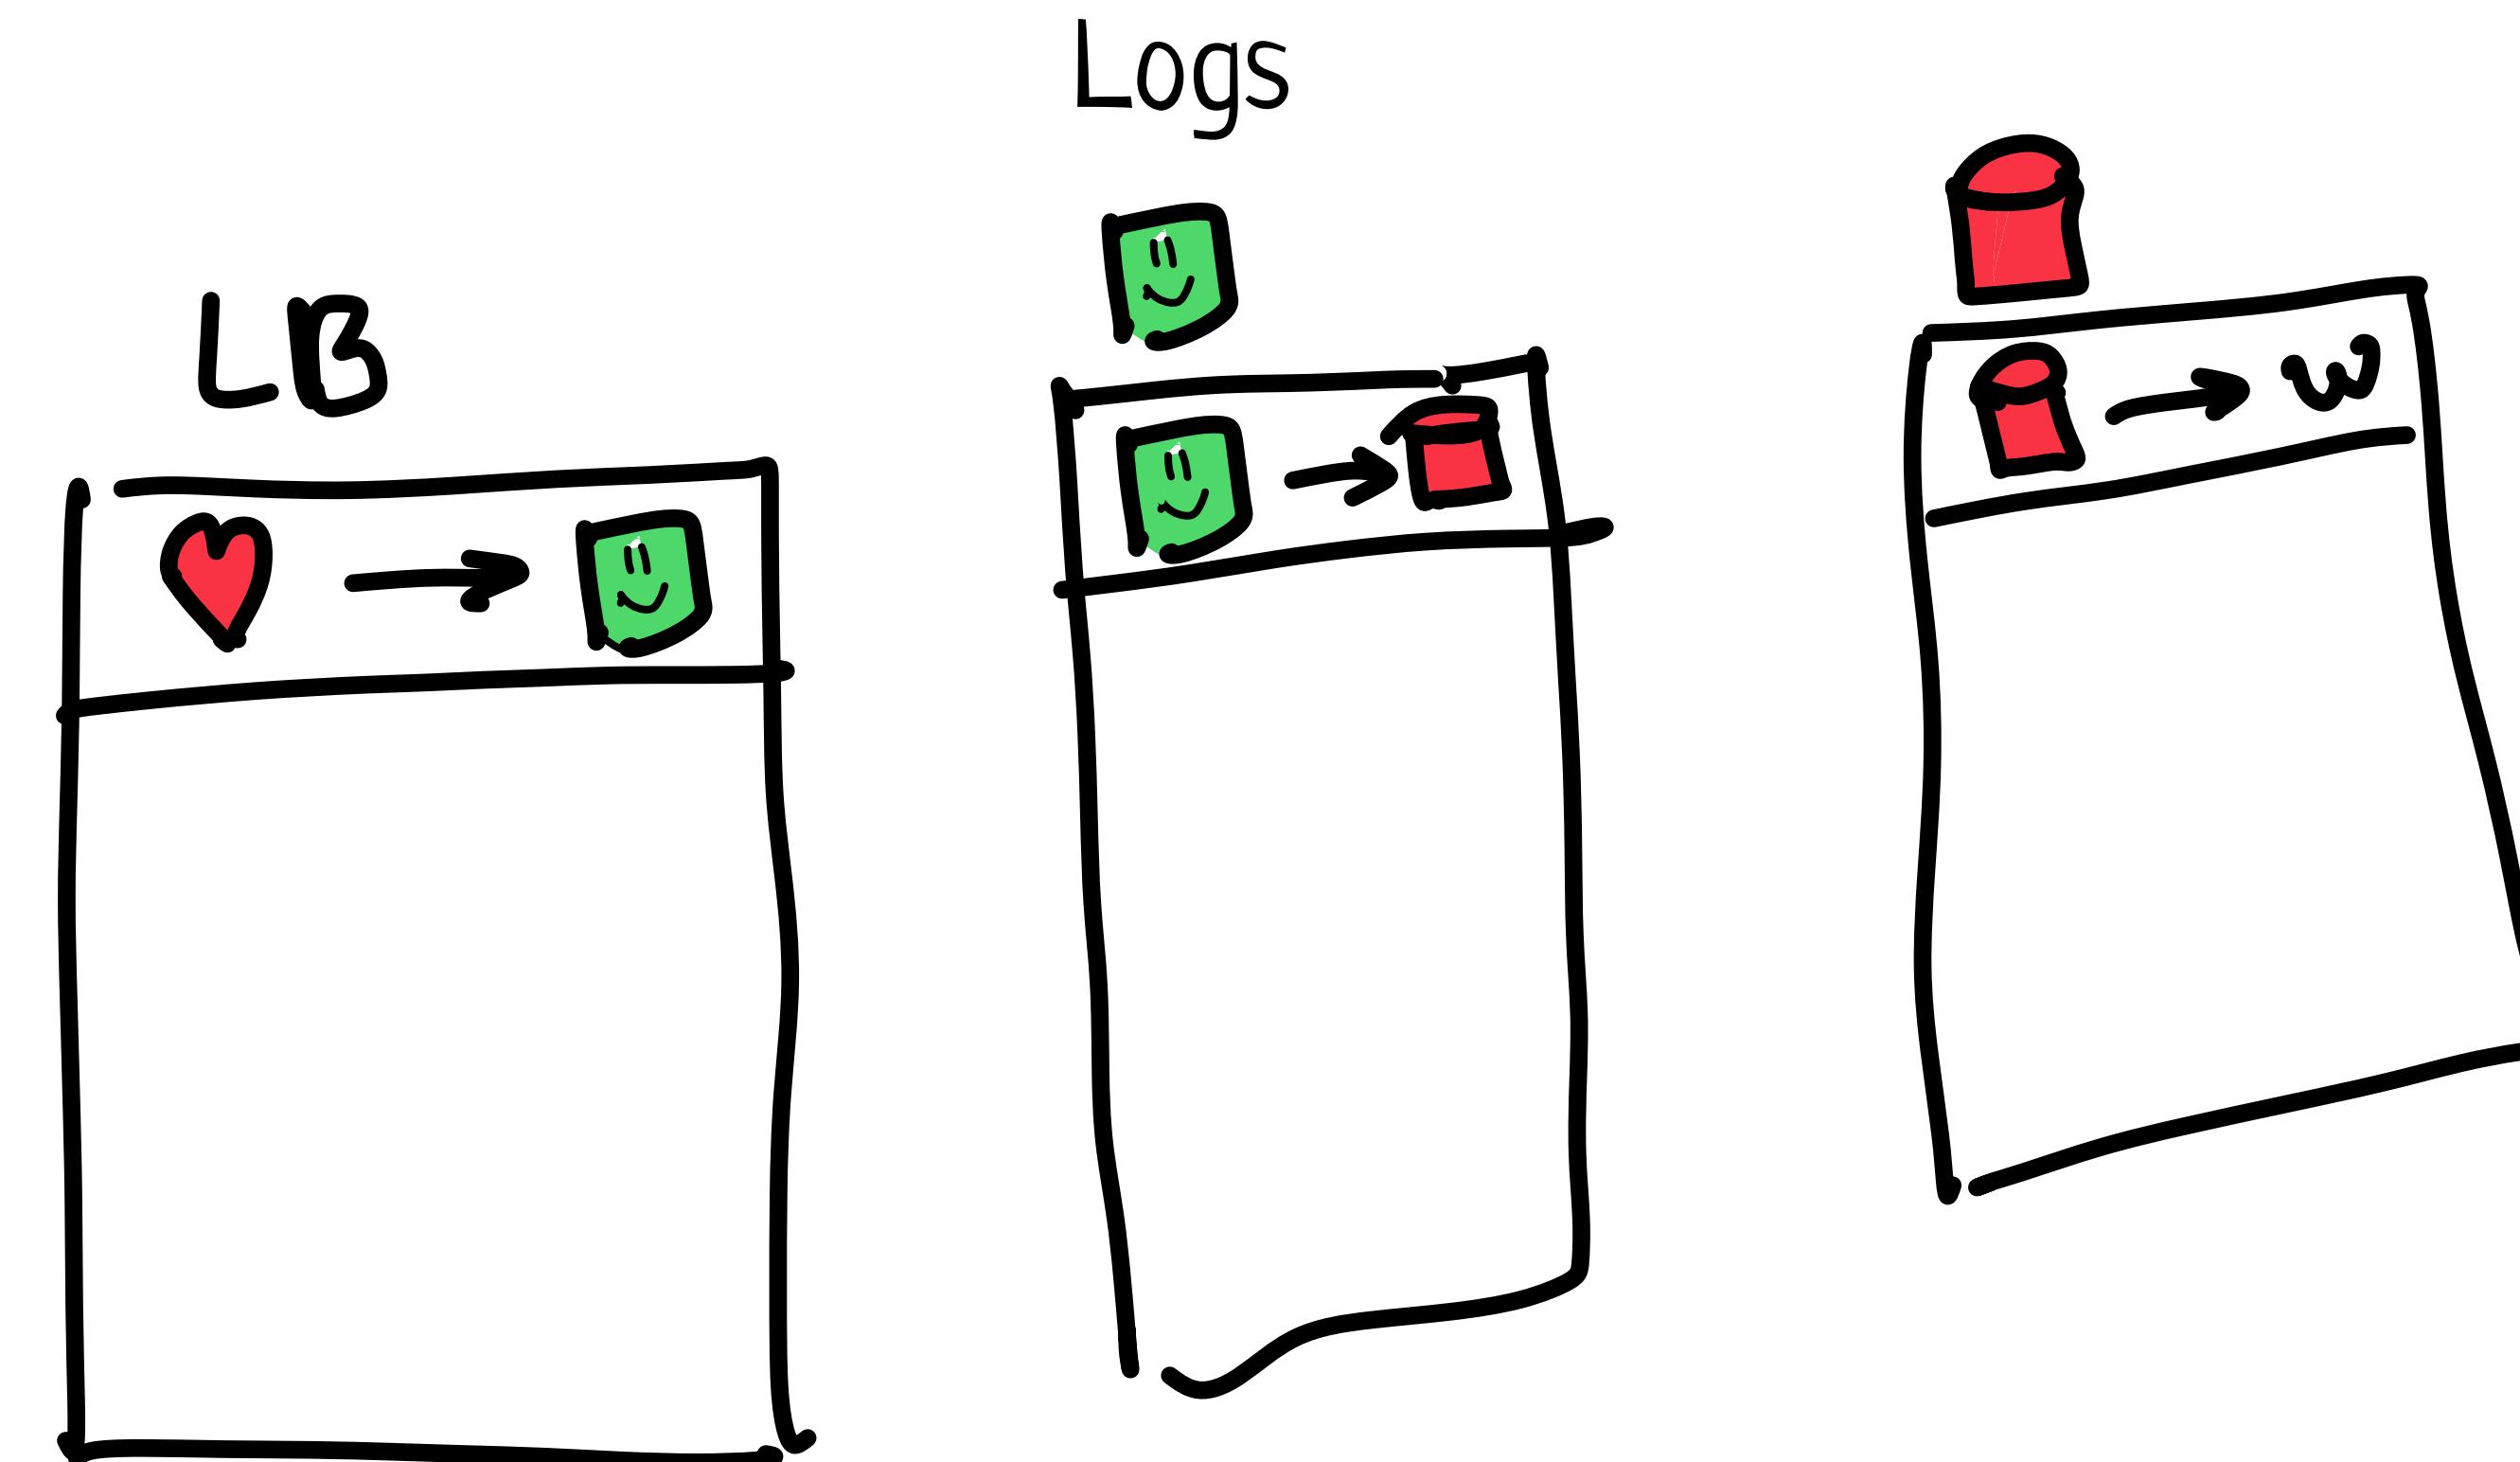
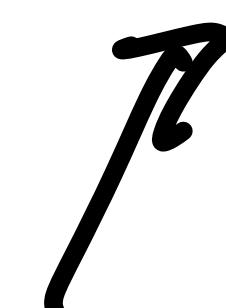
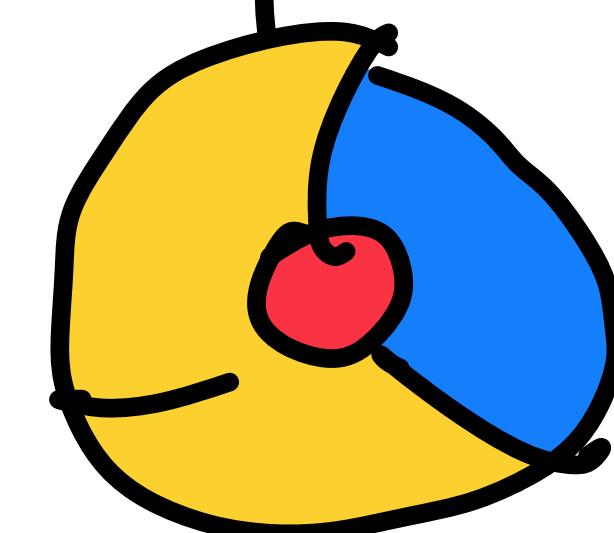
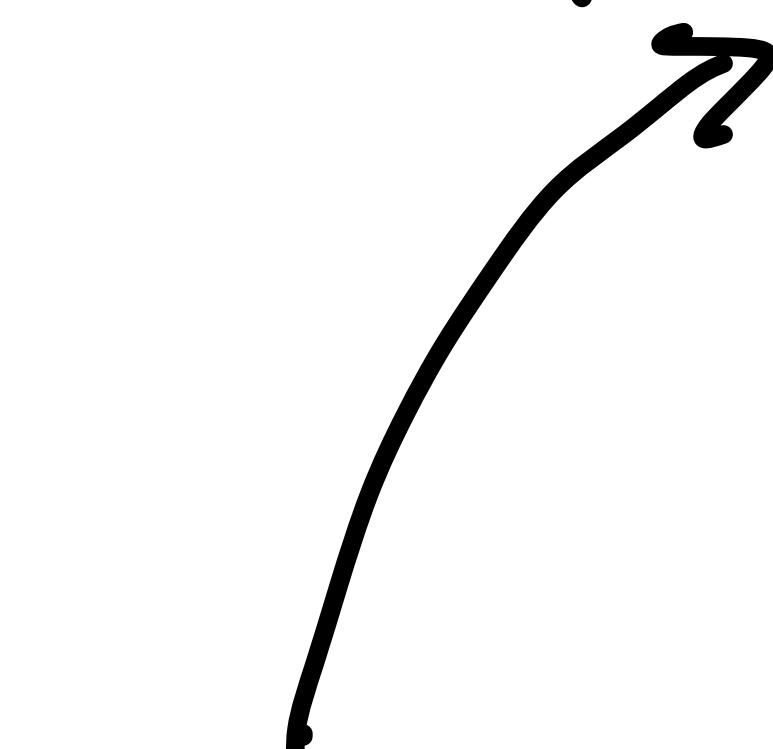
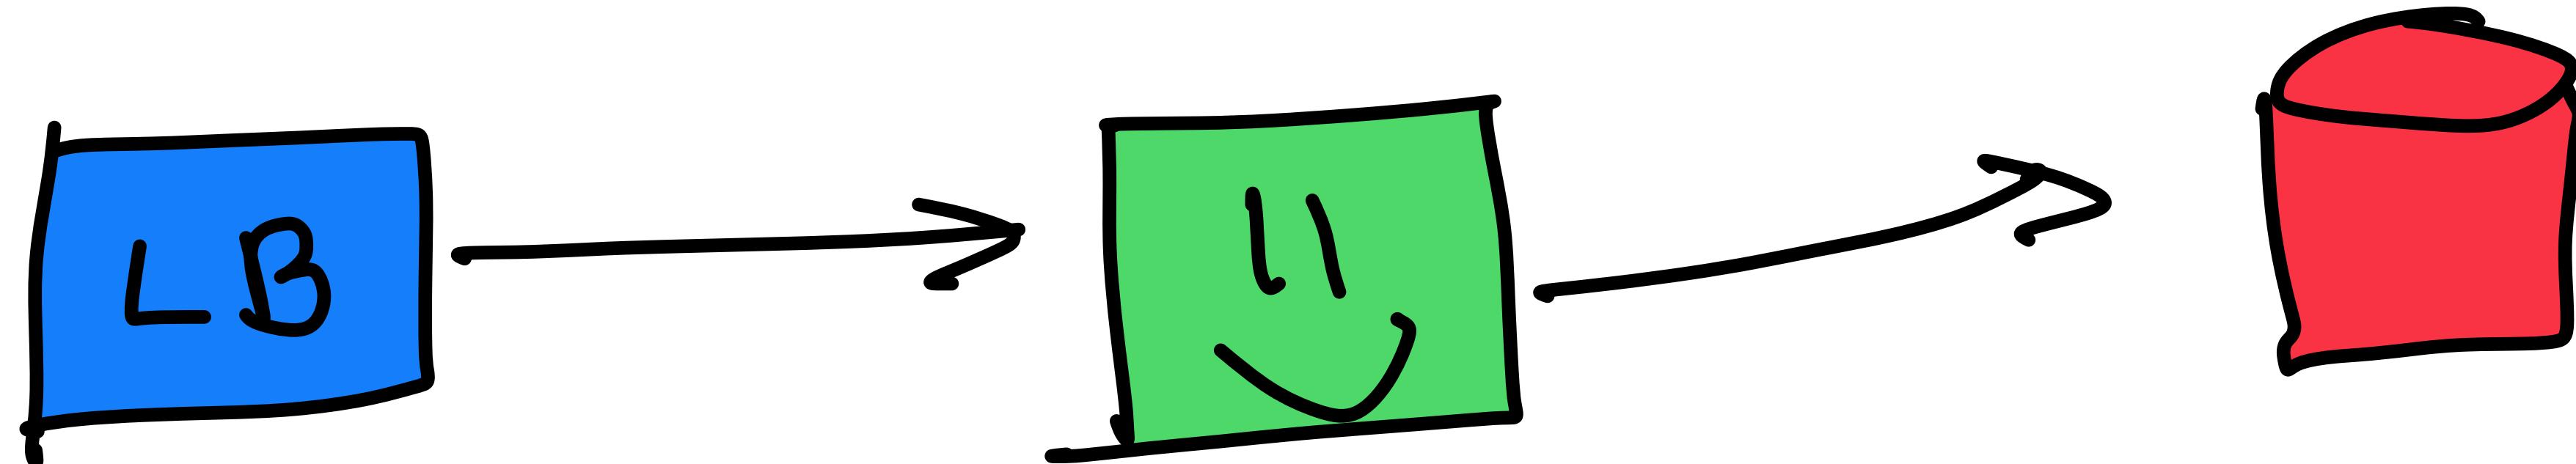
No UUIDs
Per transaction



Logs

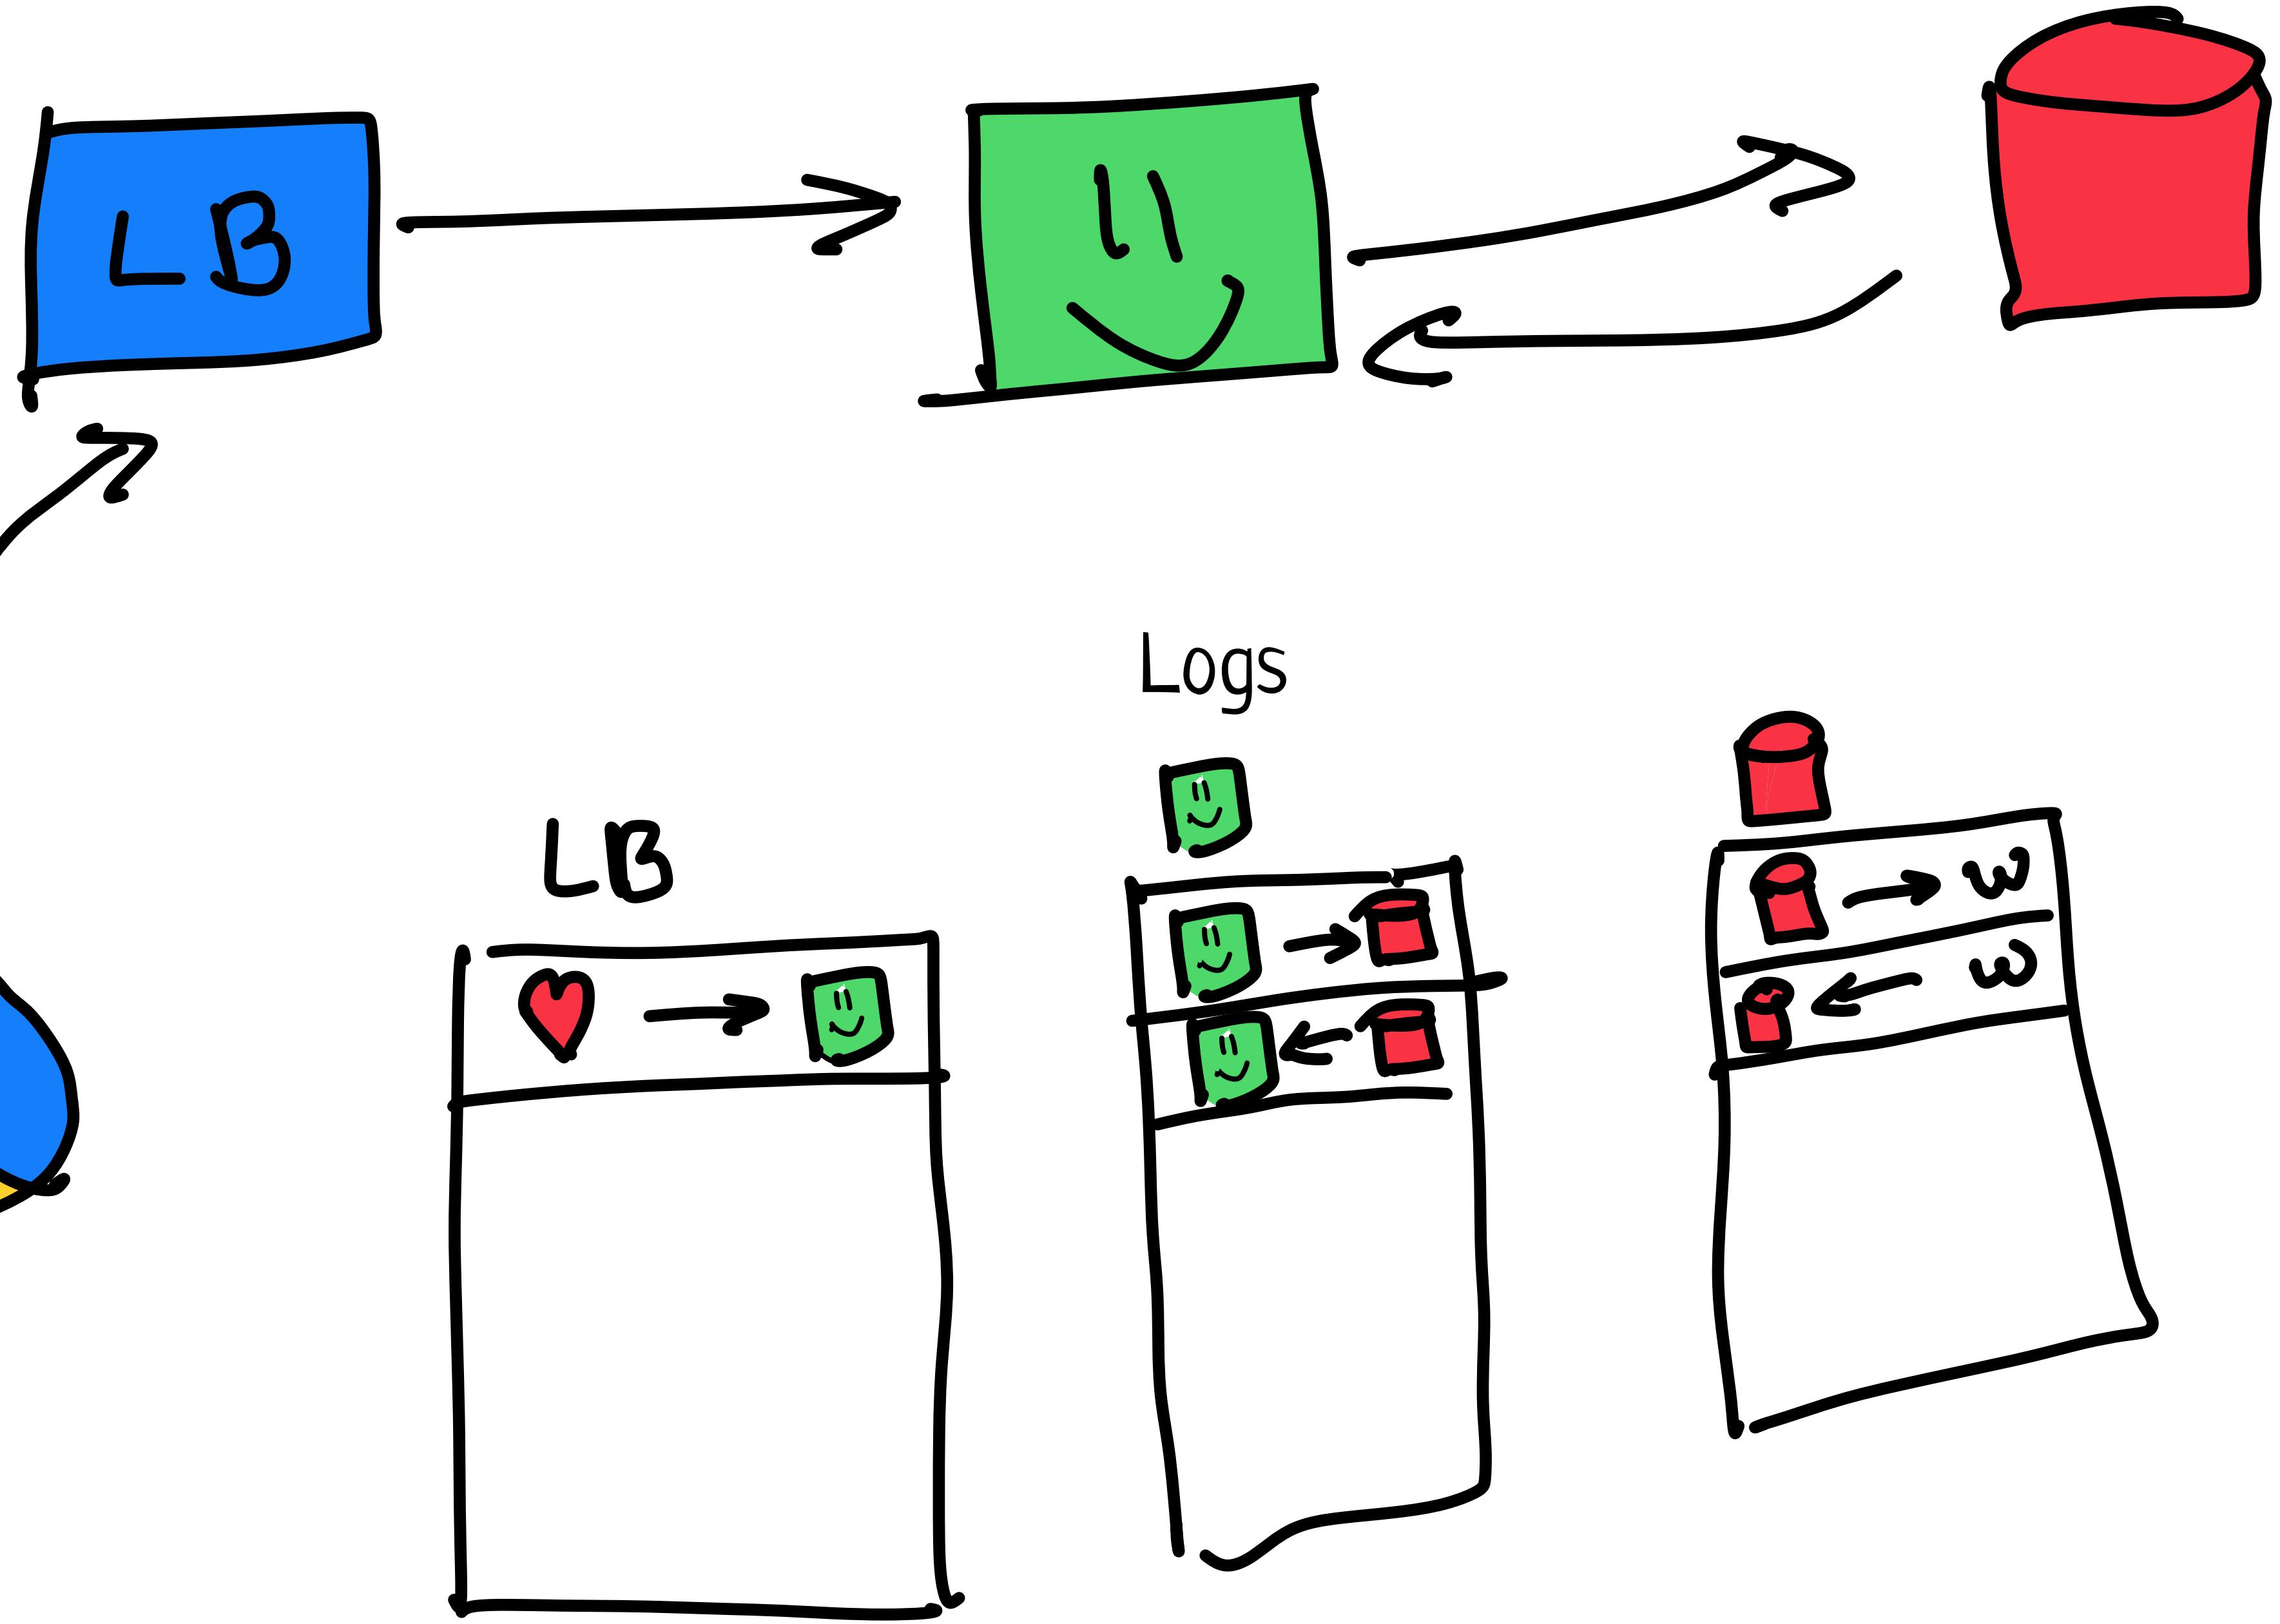


No UUIDs
Per transaction

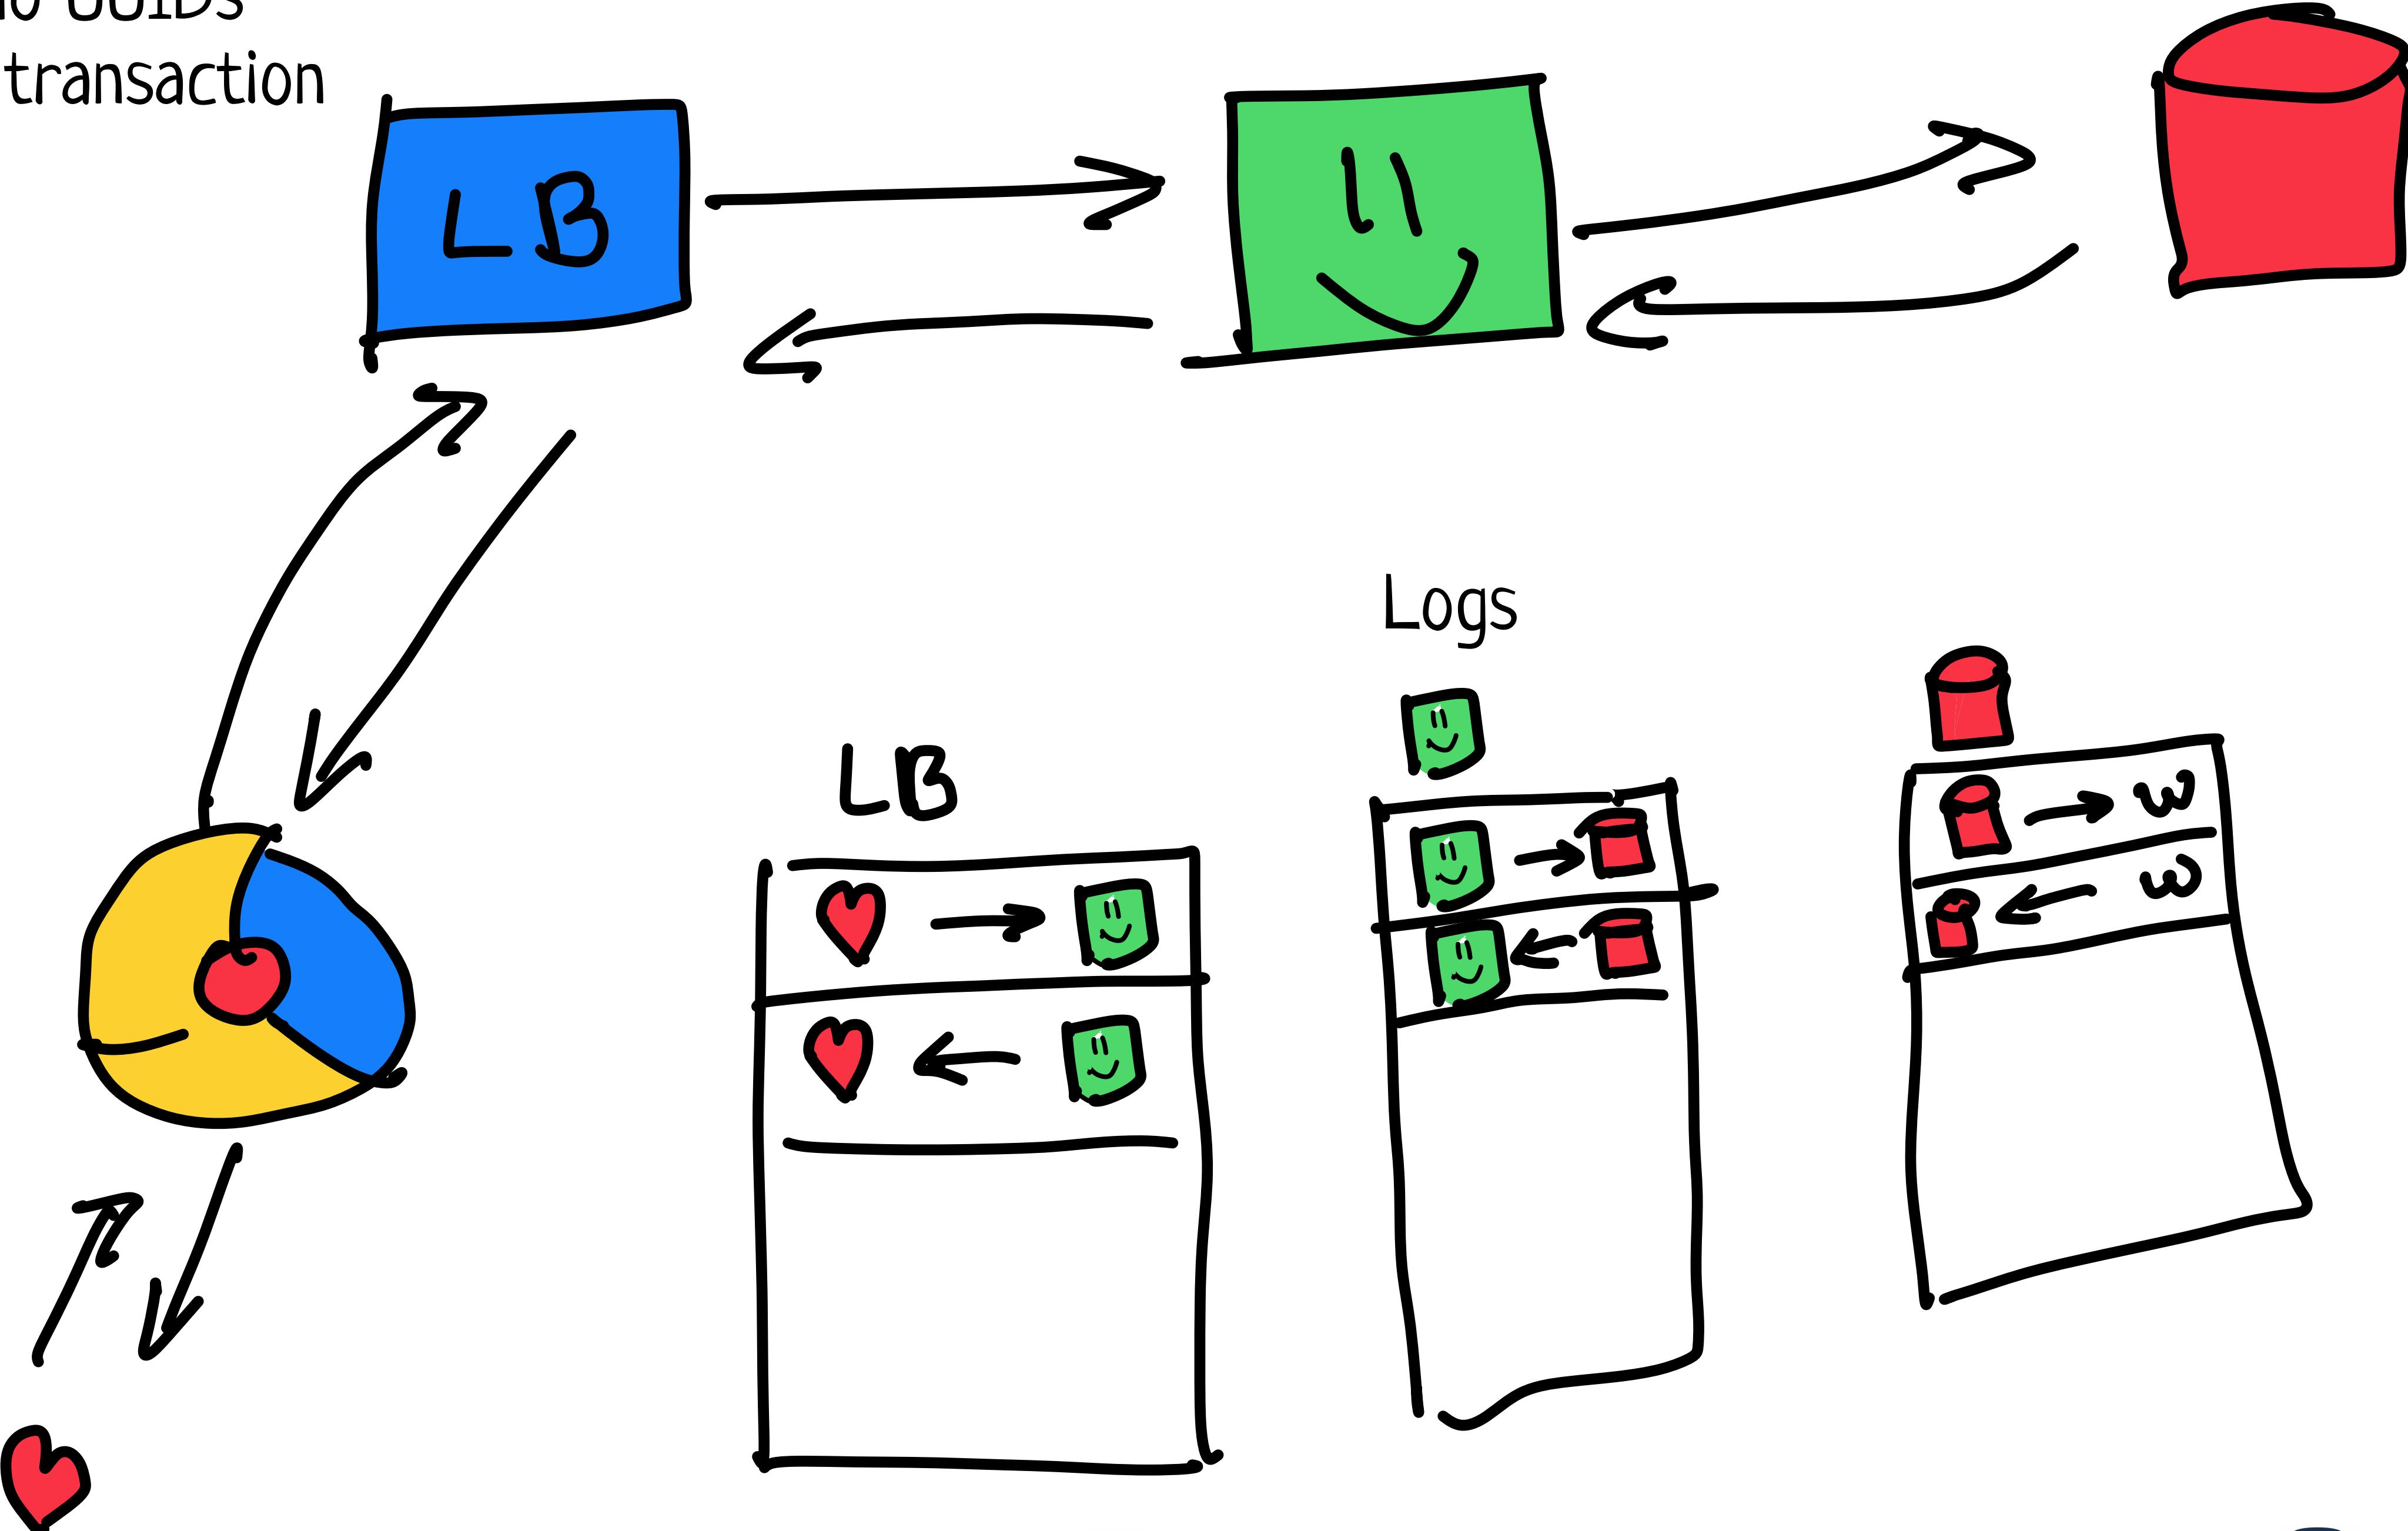


Logs

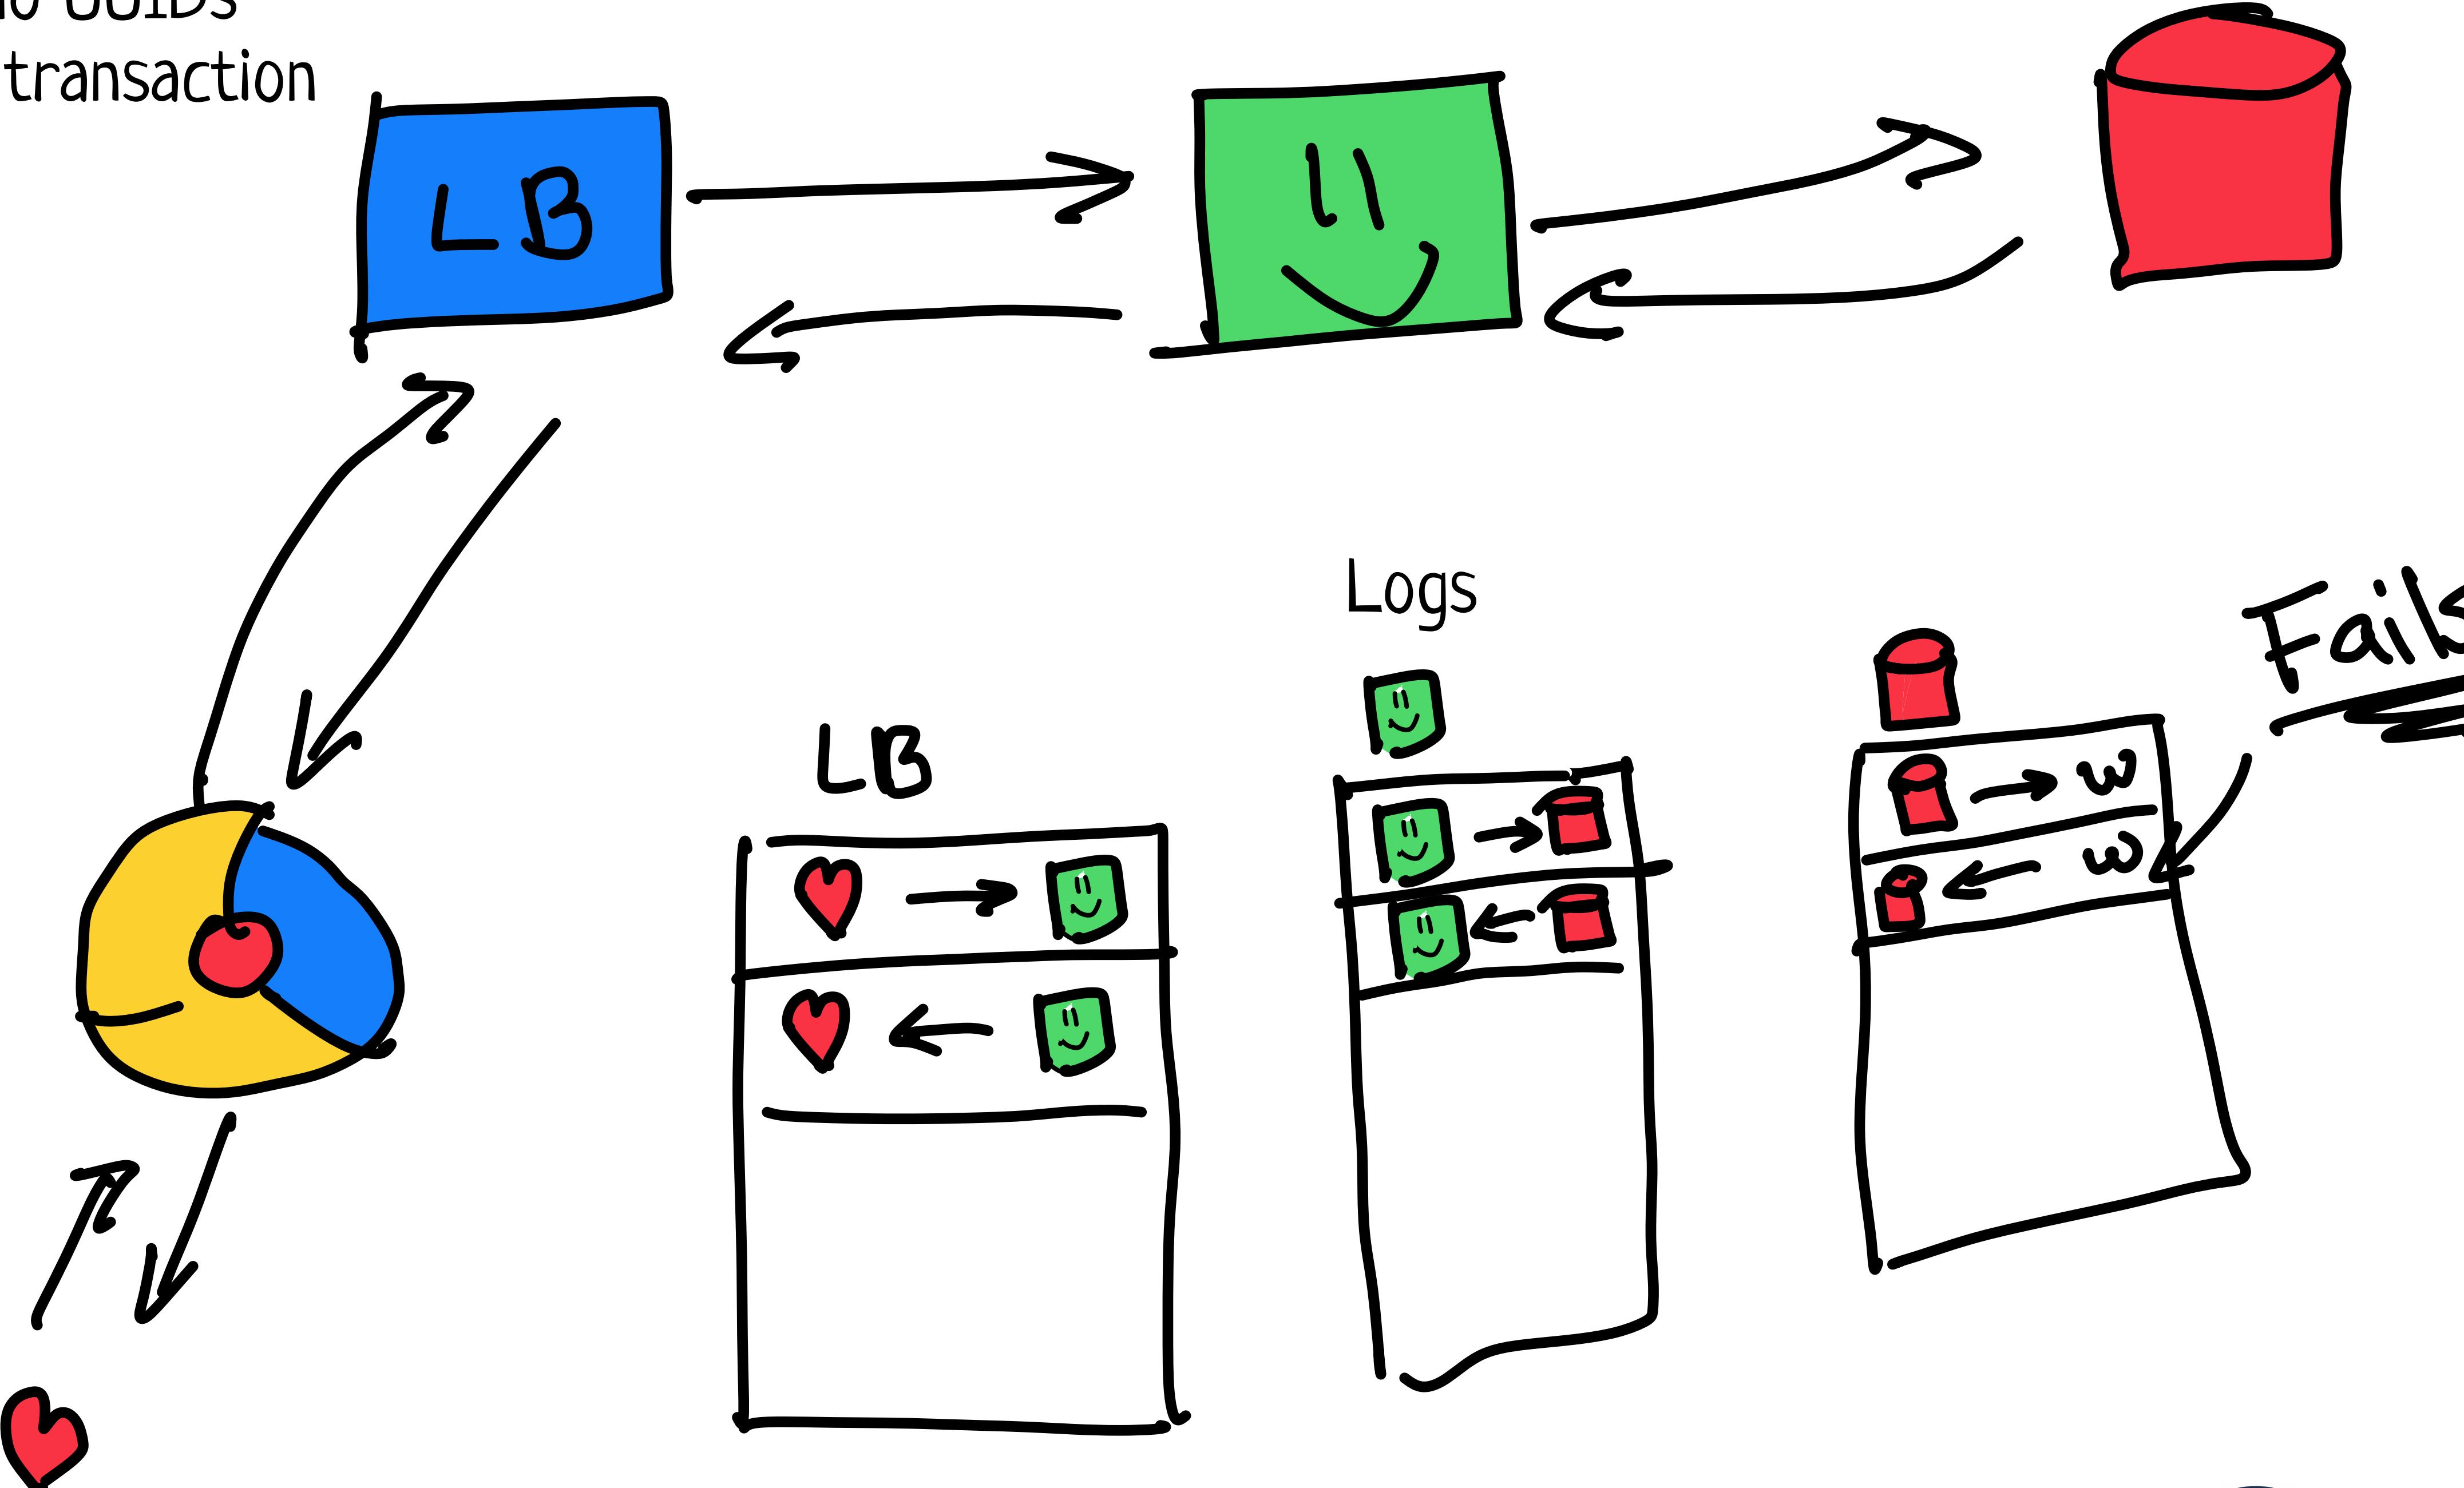
No UUIDs
Per transaction

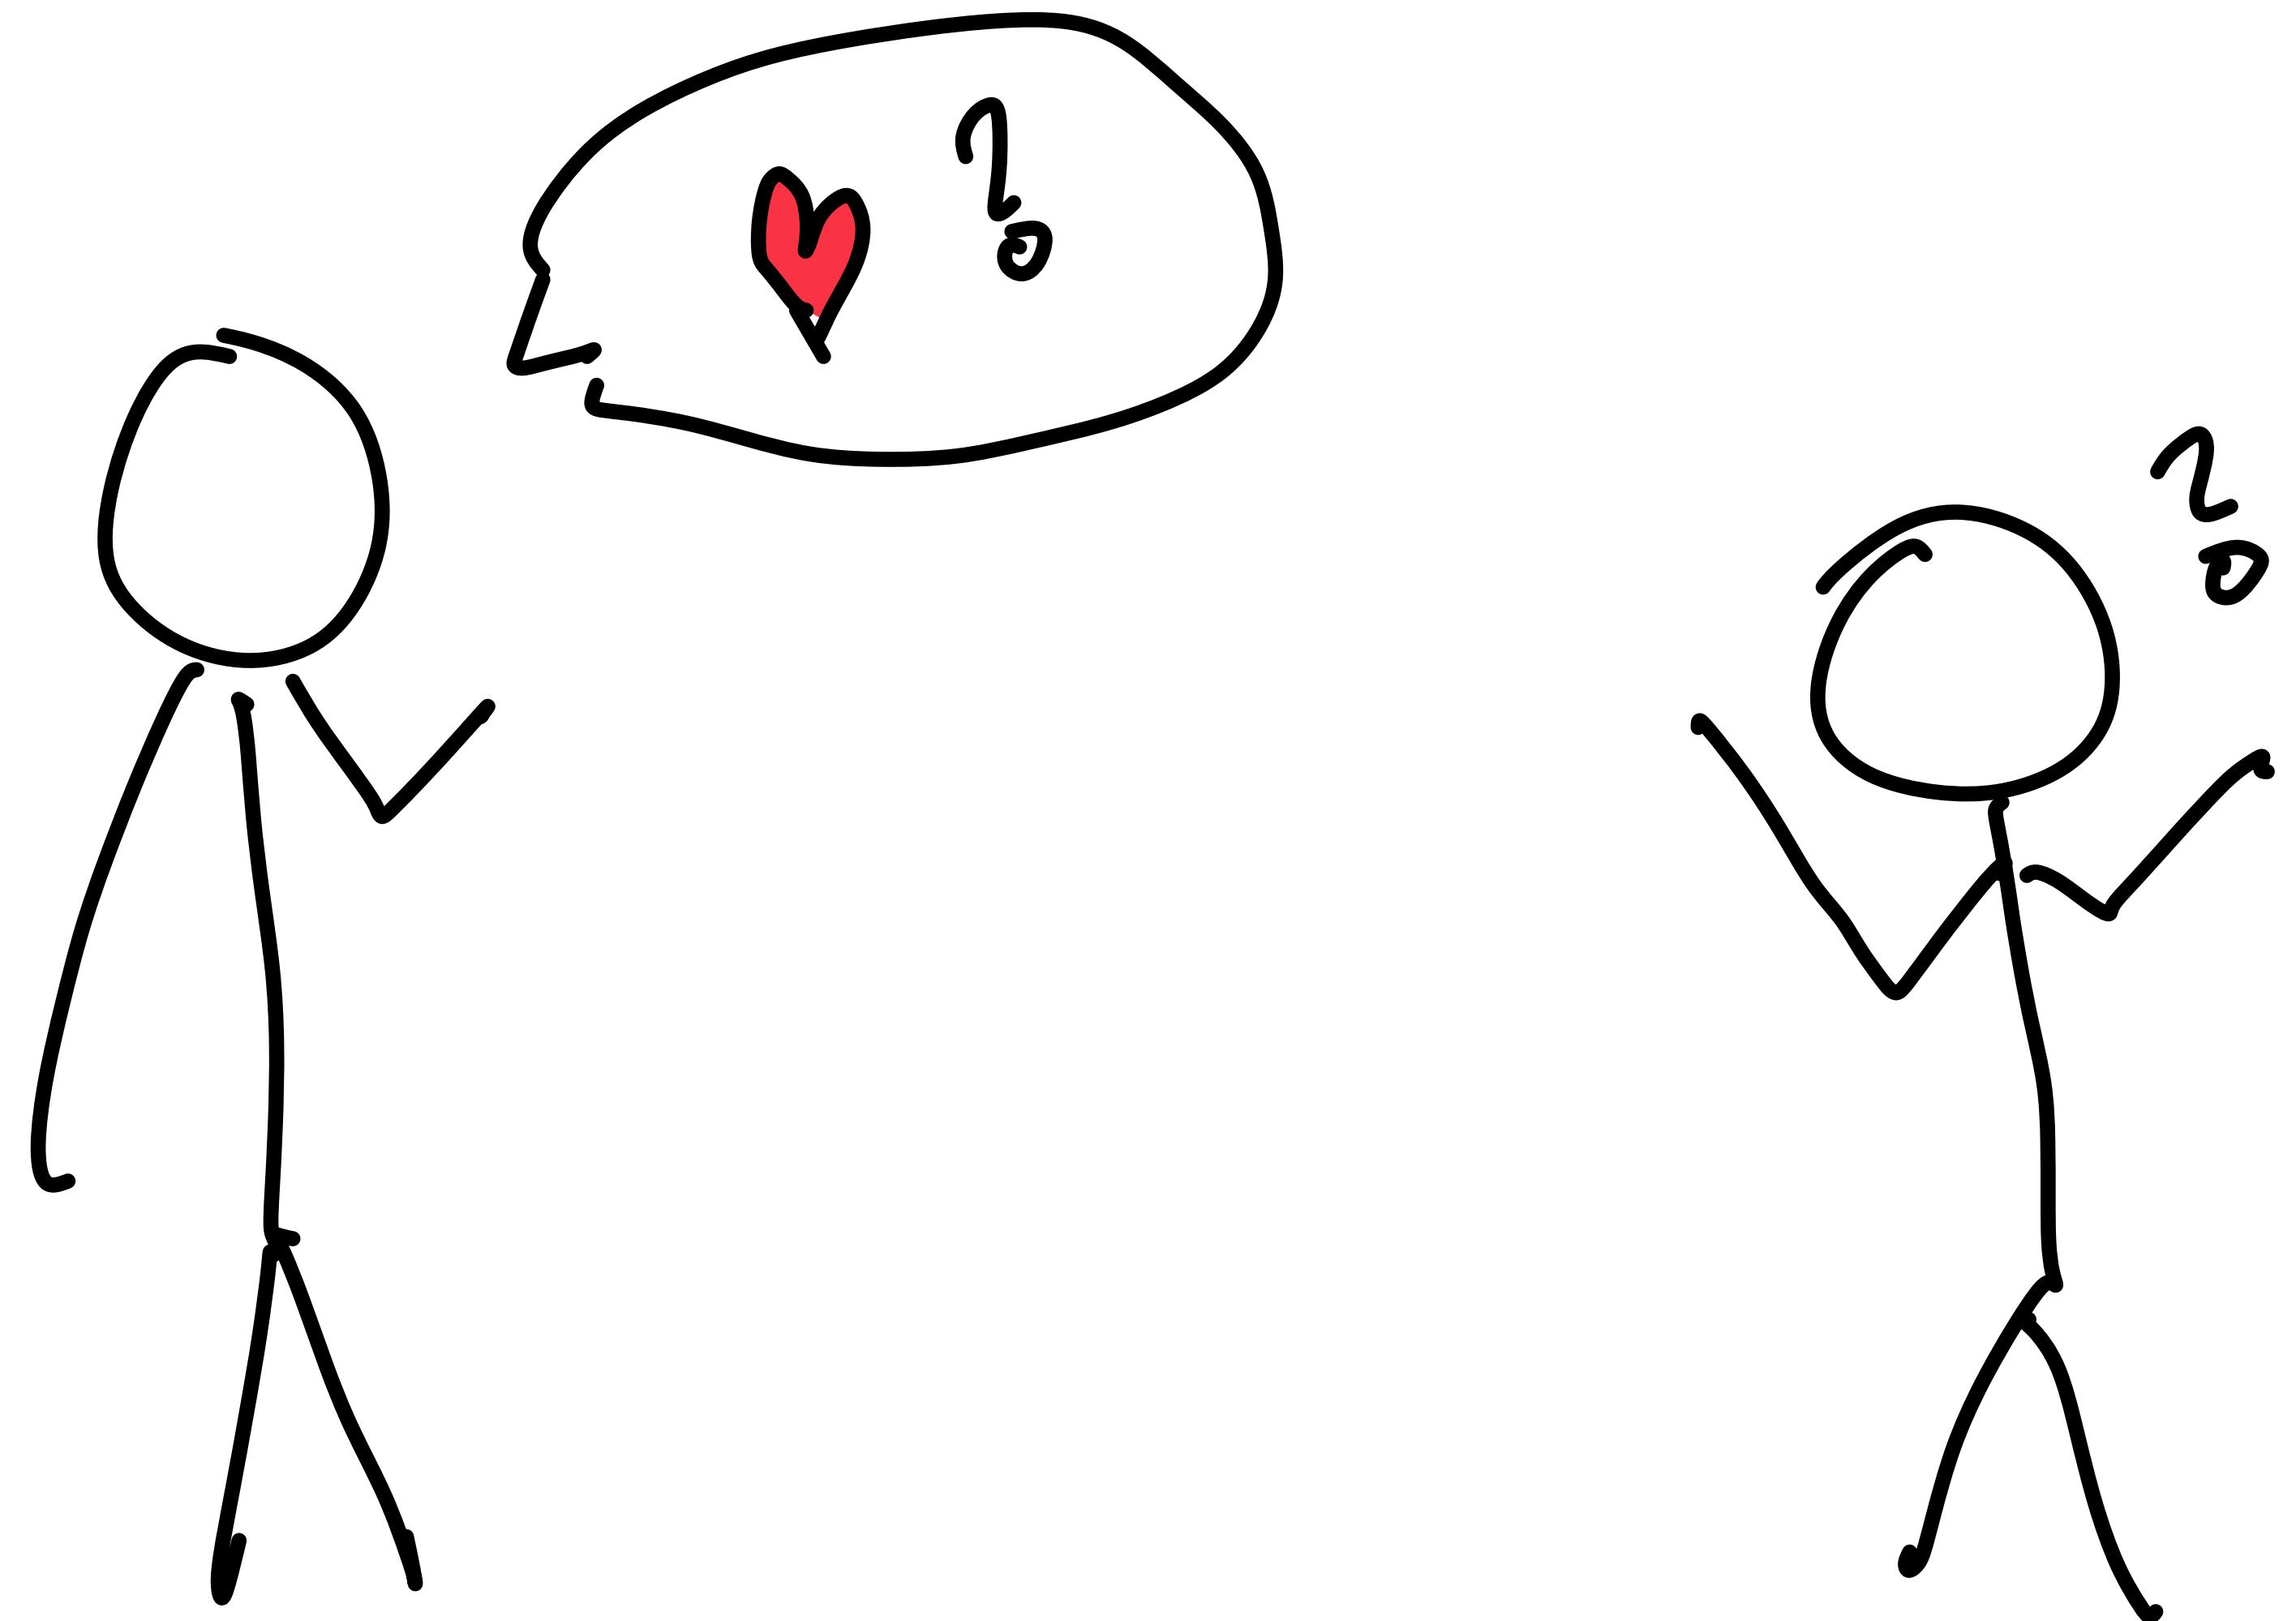


No UUIDs
Per transaction



No UUIDs
Per transaction



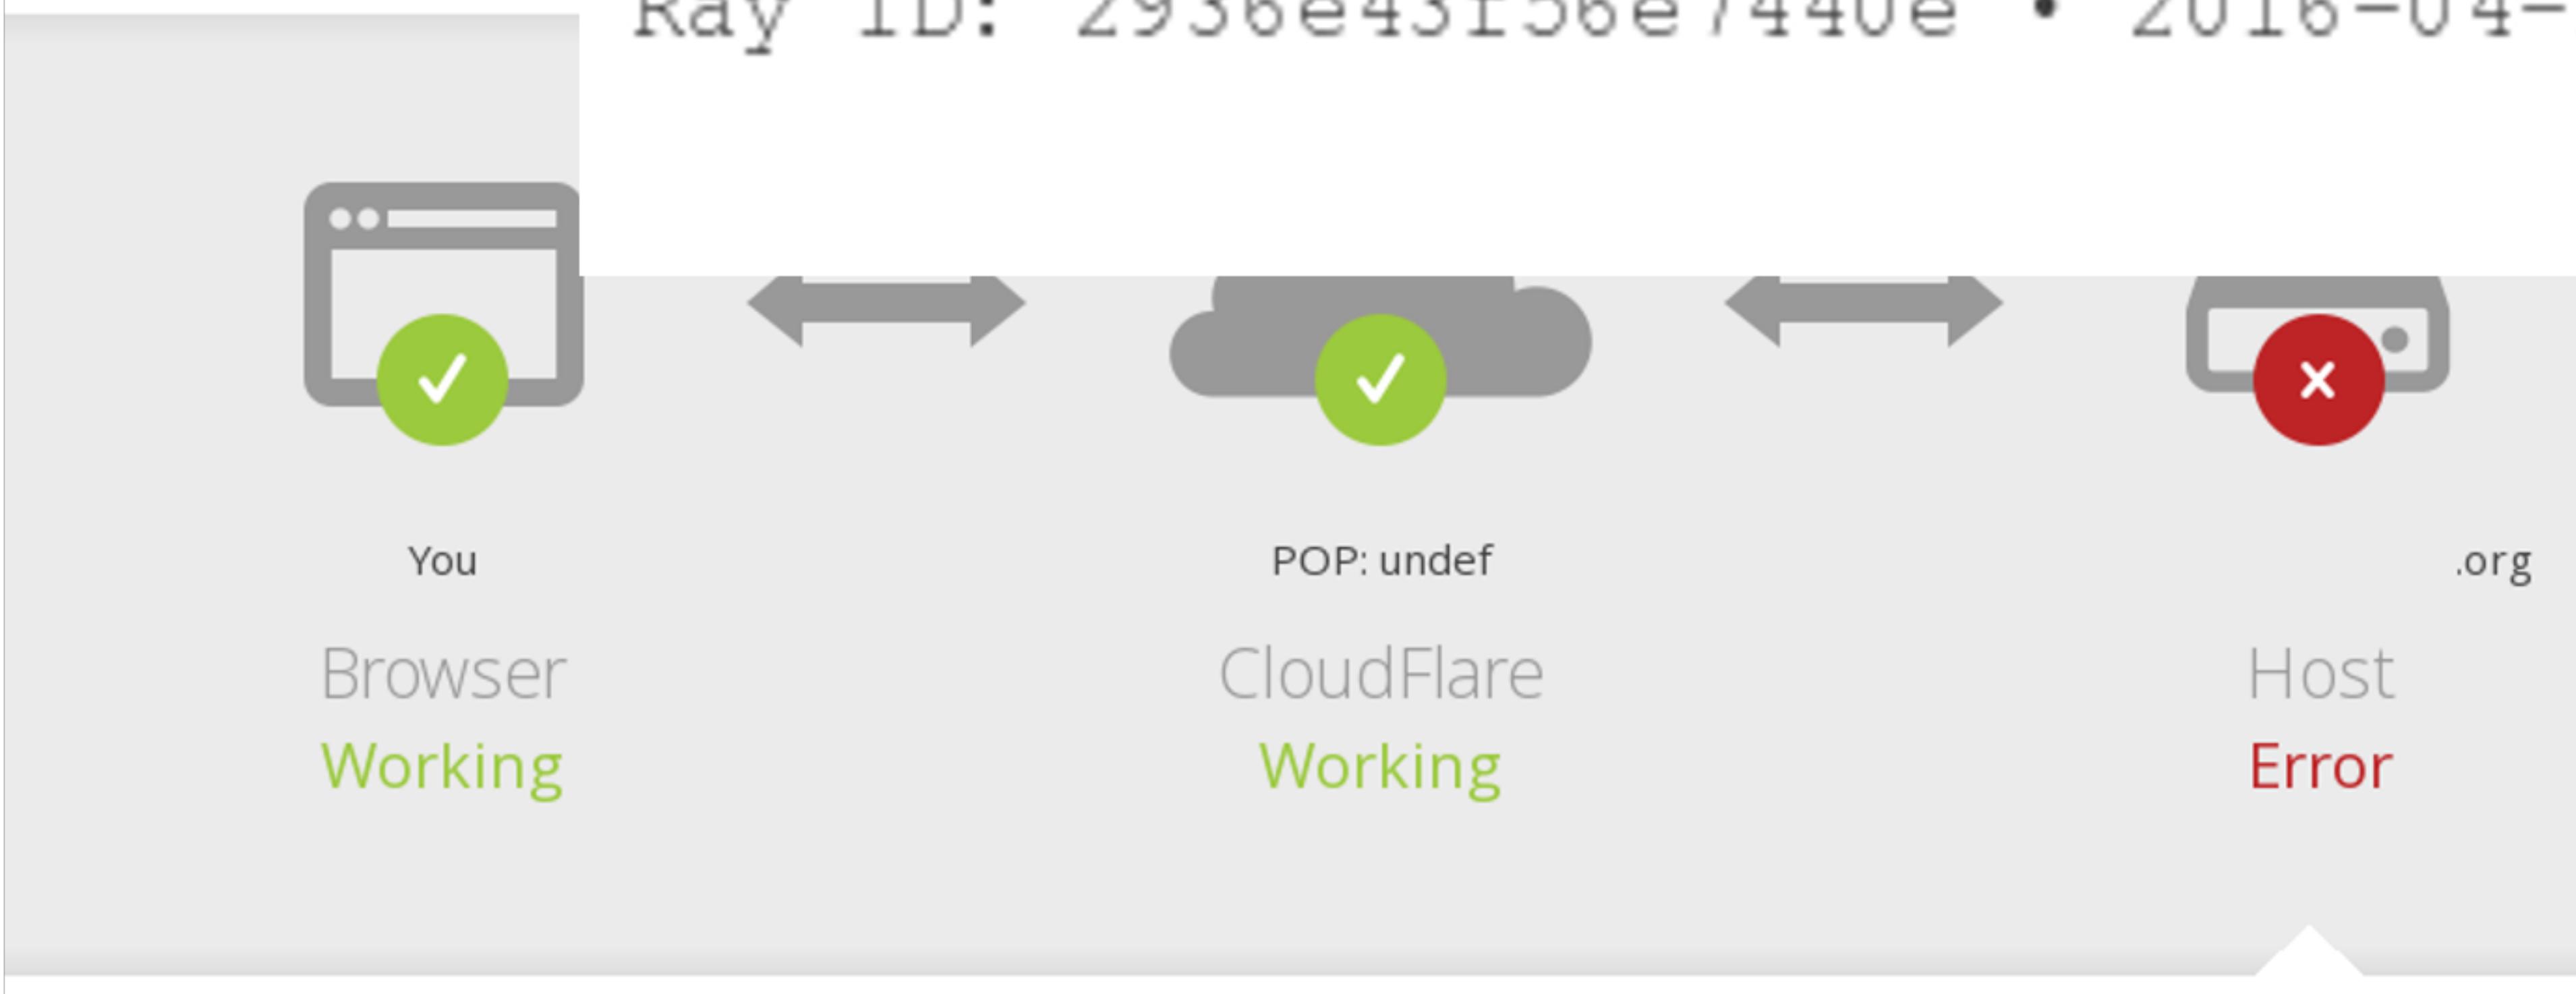


Error 502

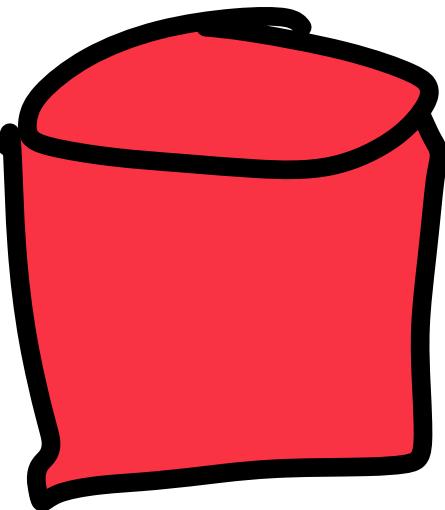
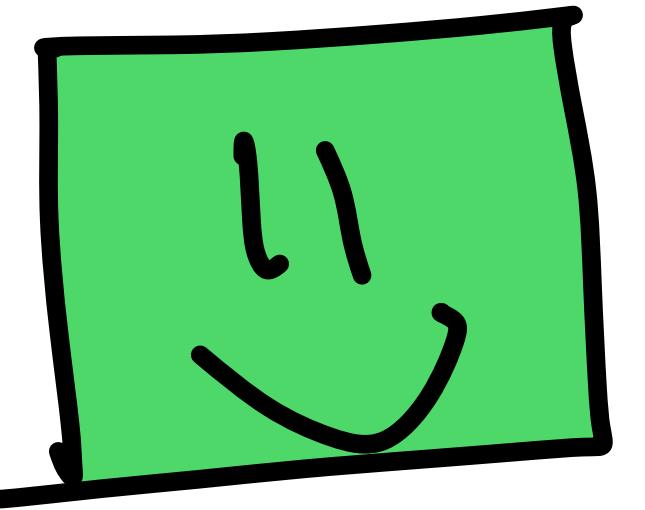
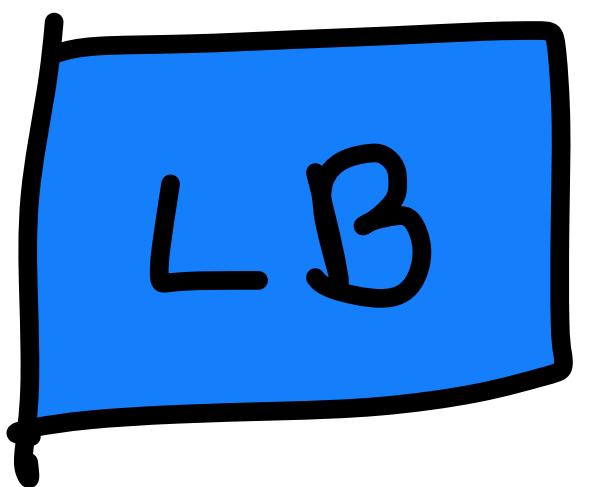
Ray ID: 2936e43f56e7440e • 2016-04-14 11:41:57 UTC

Bad gateway

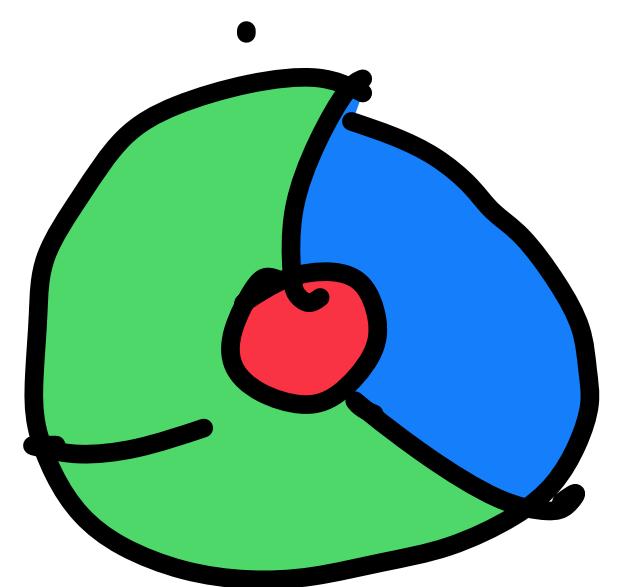
Ray ID: 2936e43f56e7440e • 2016-04-14 11:41:5



Atomic Uuids

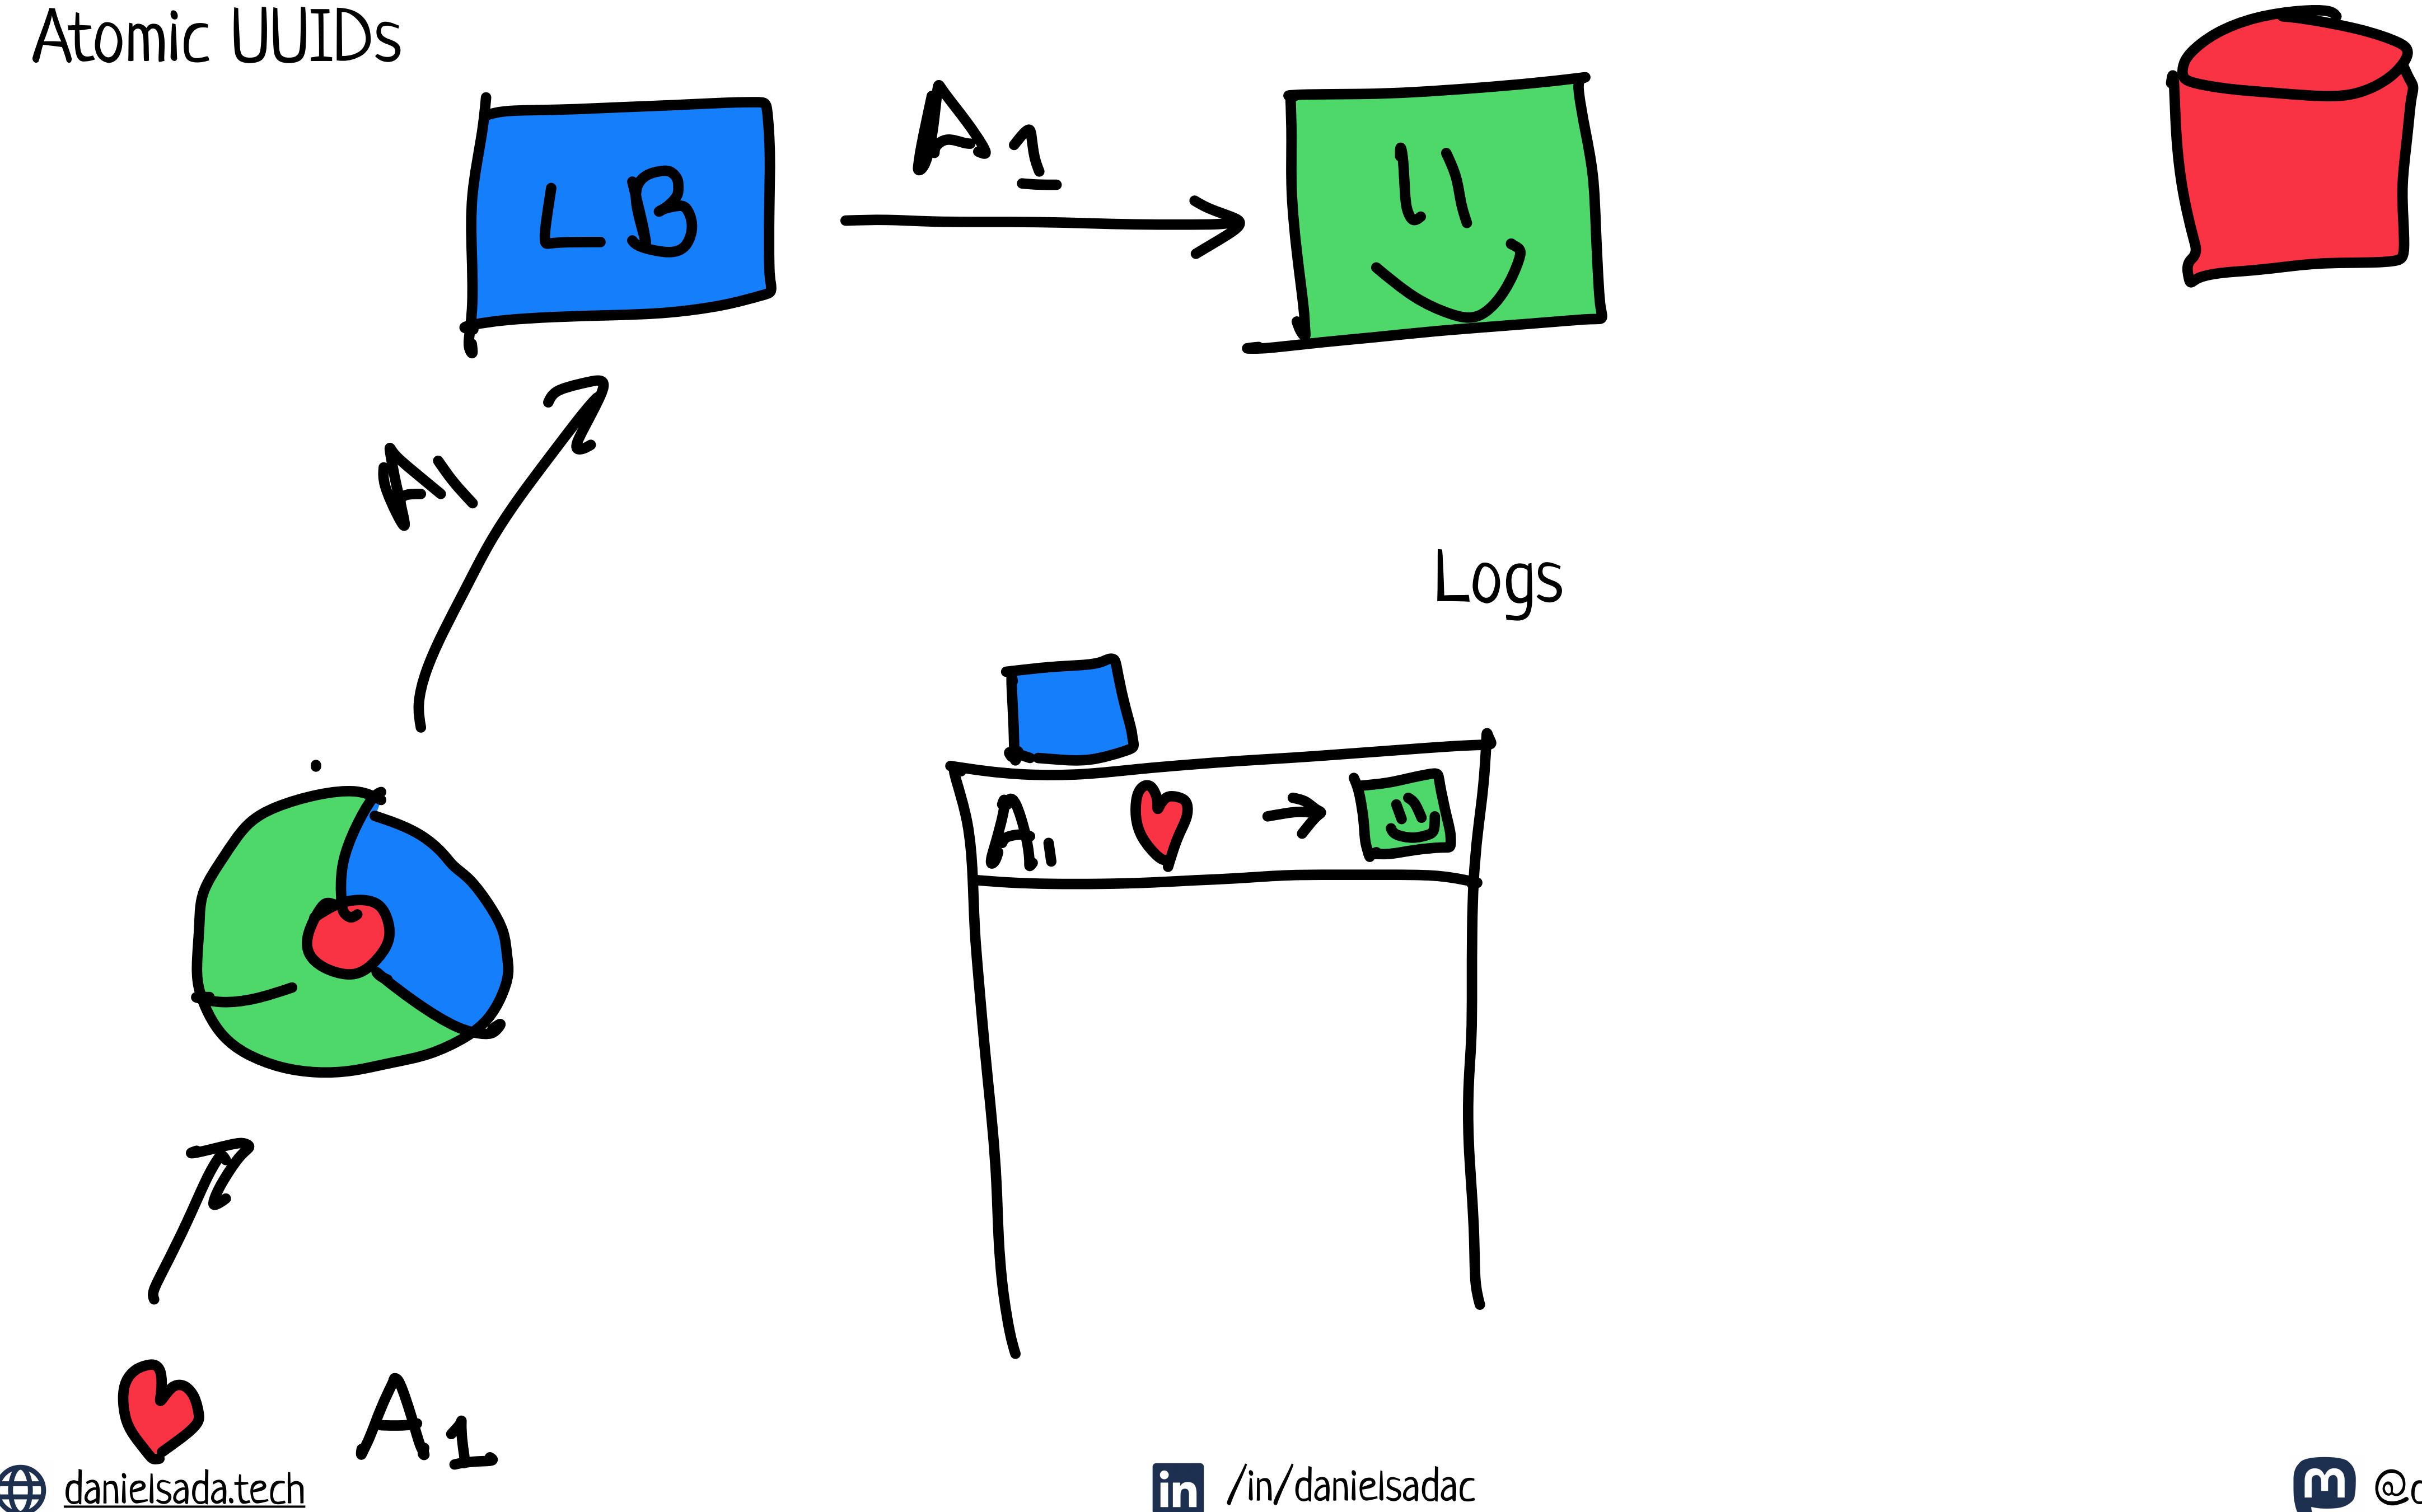


Logs

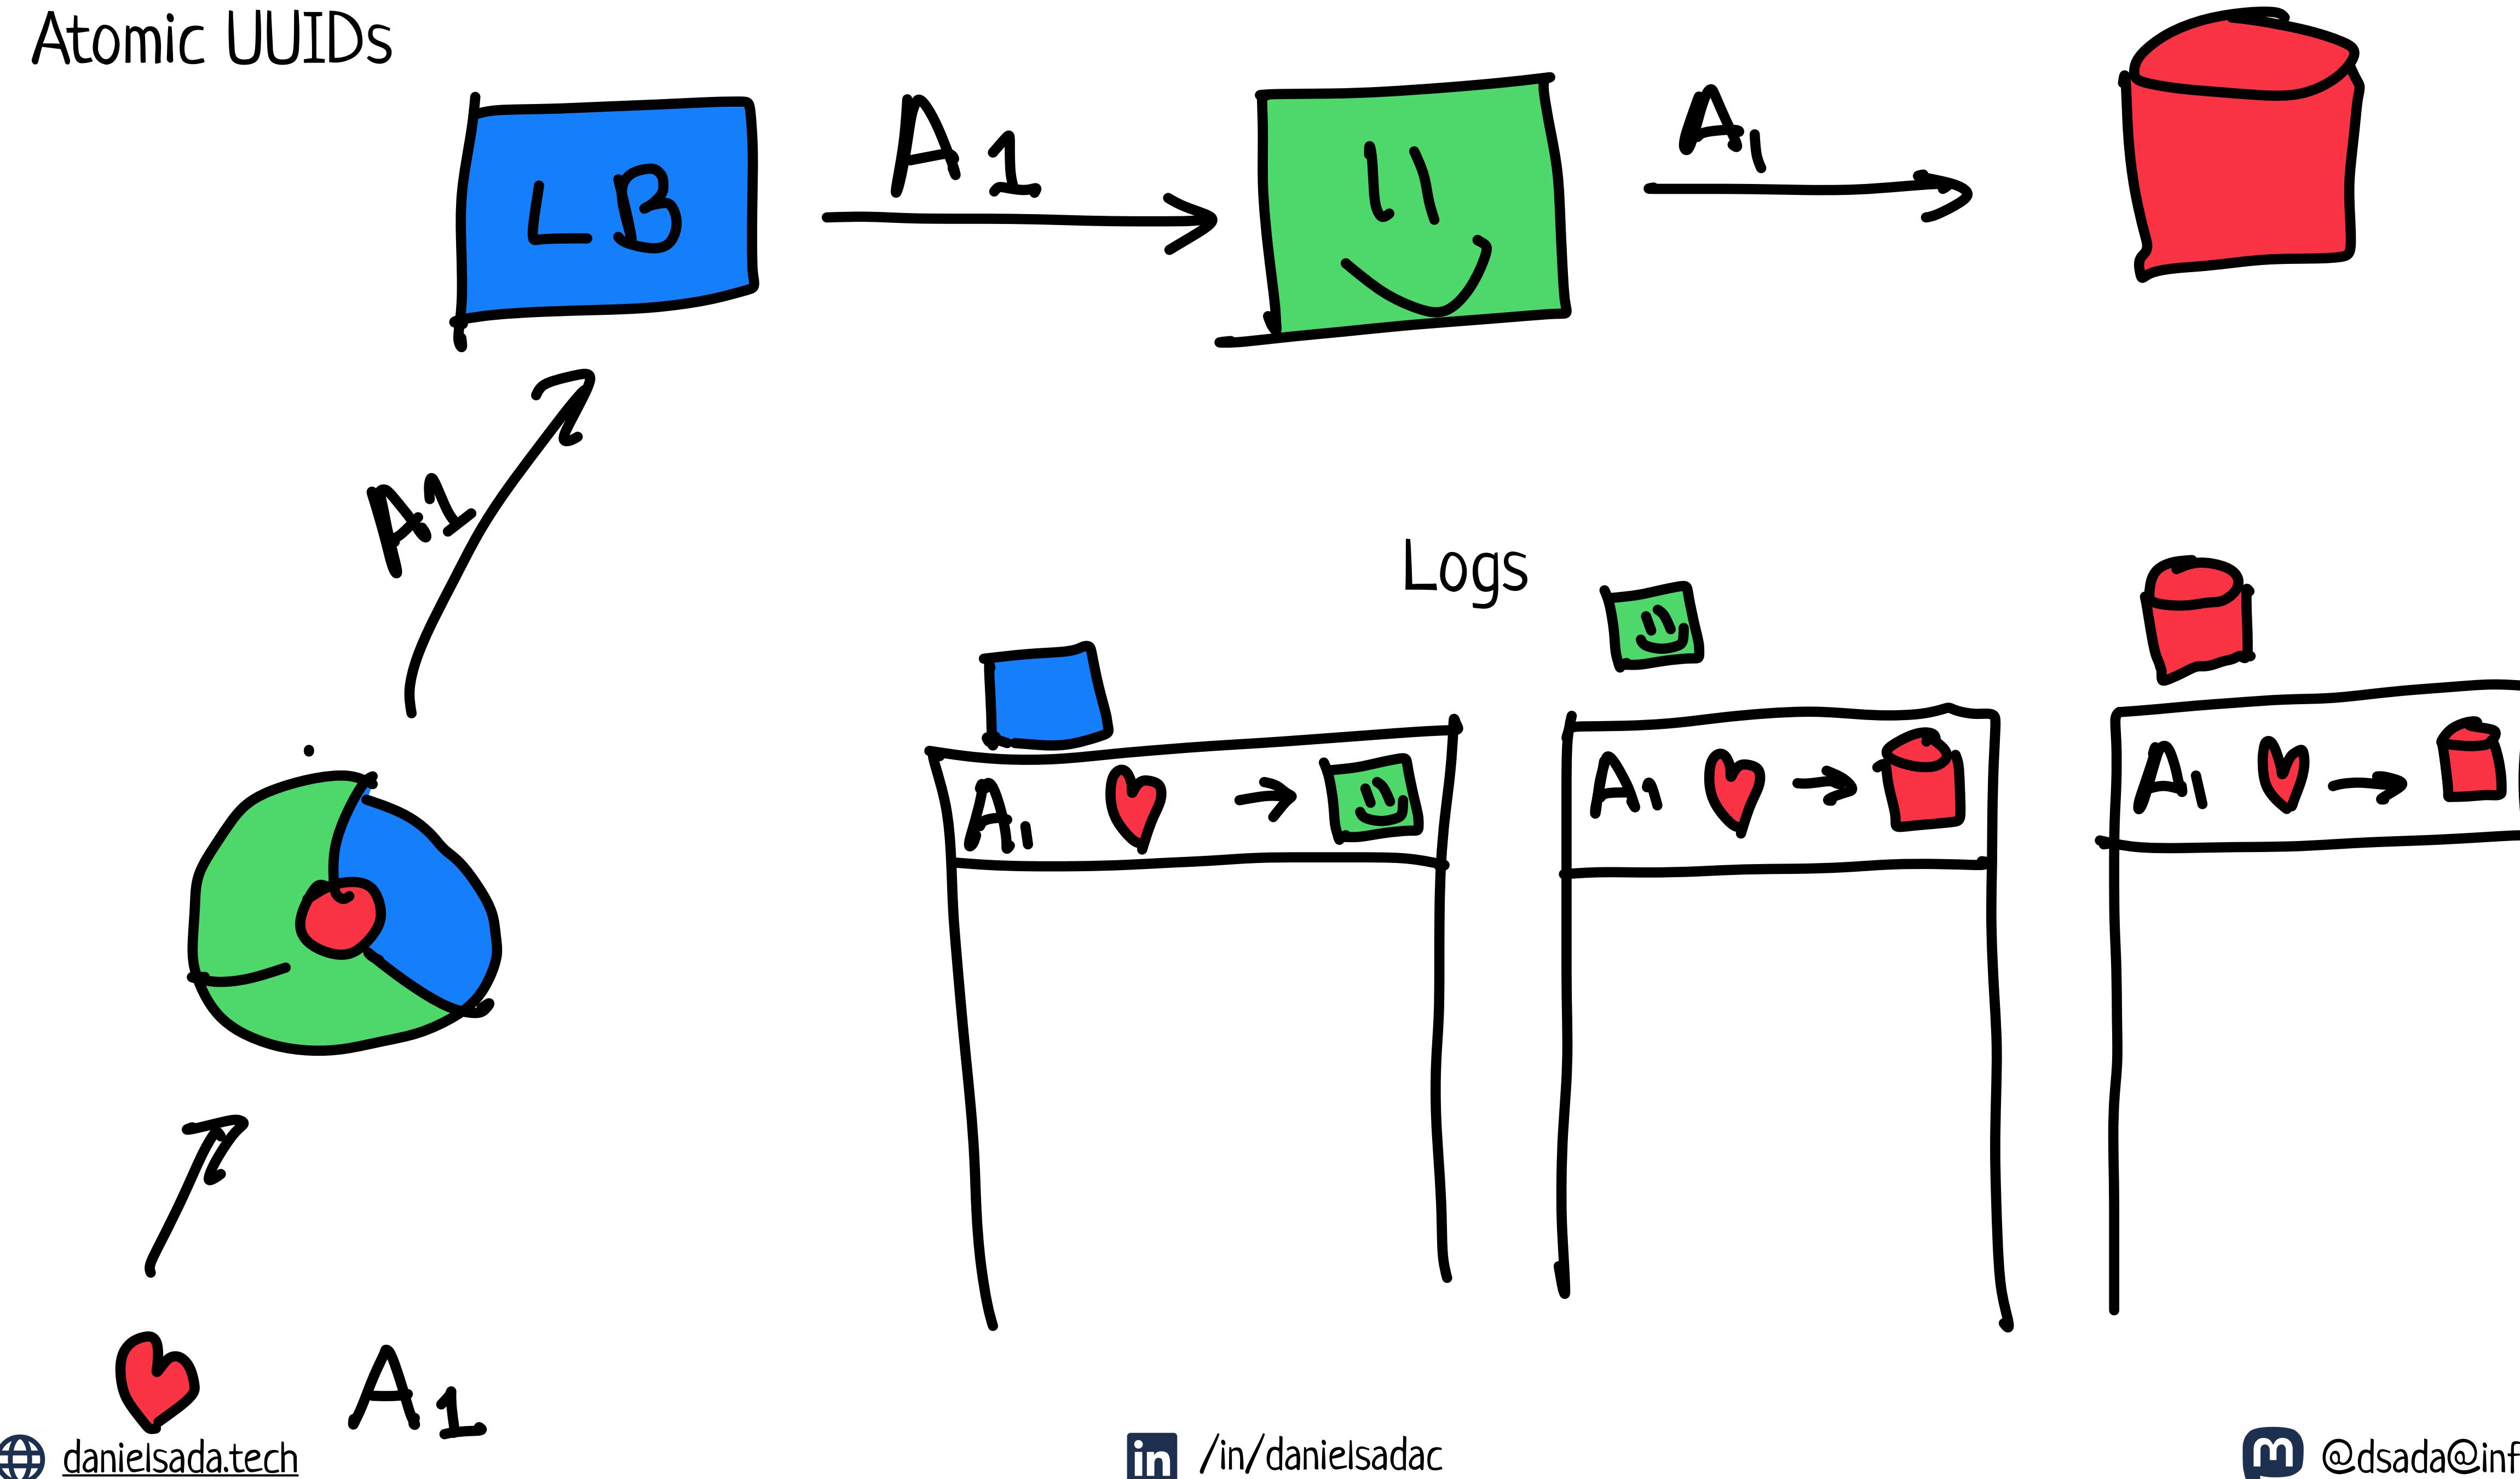


A₁

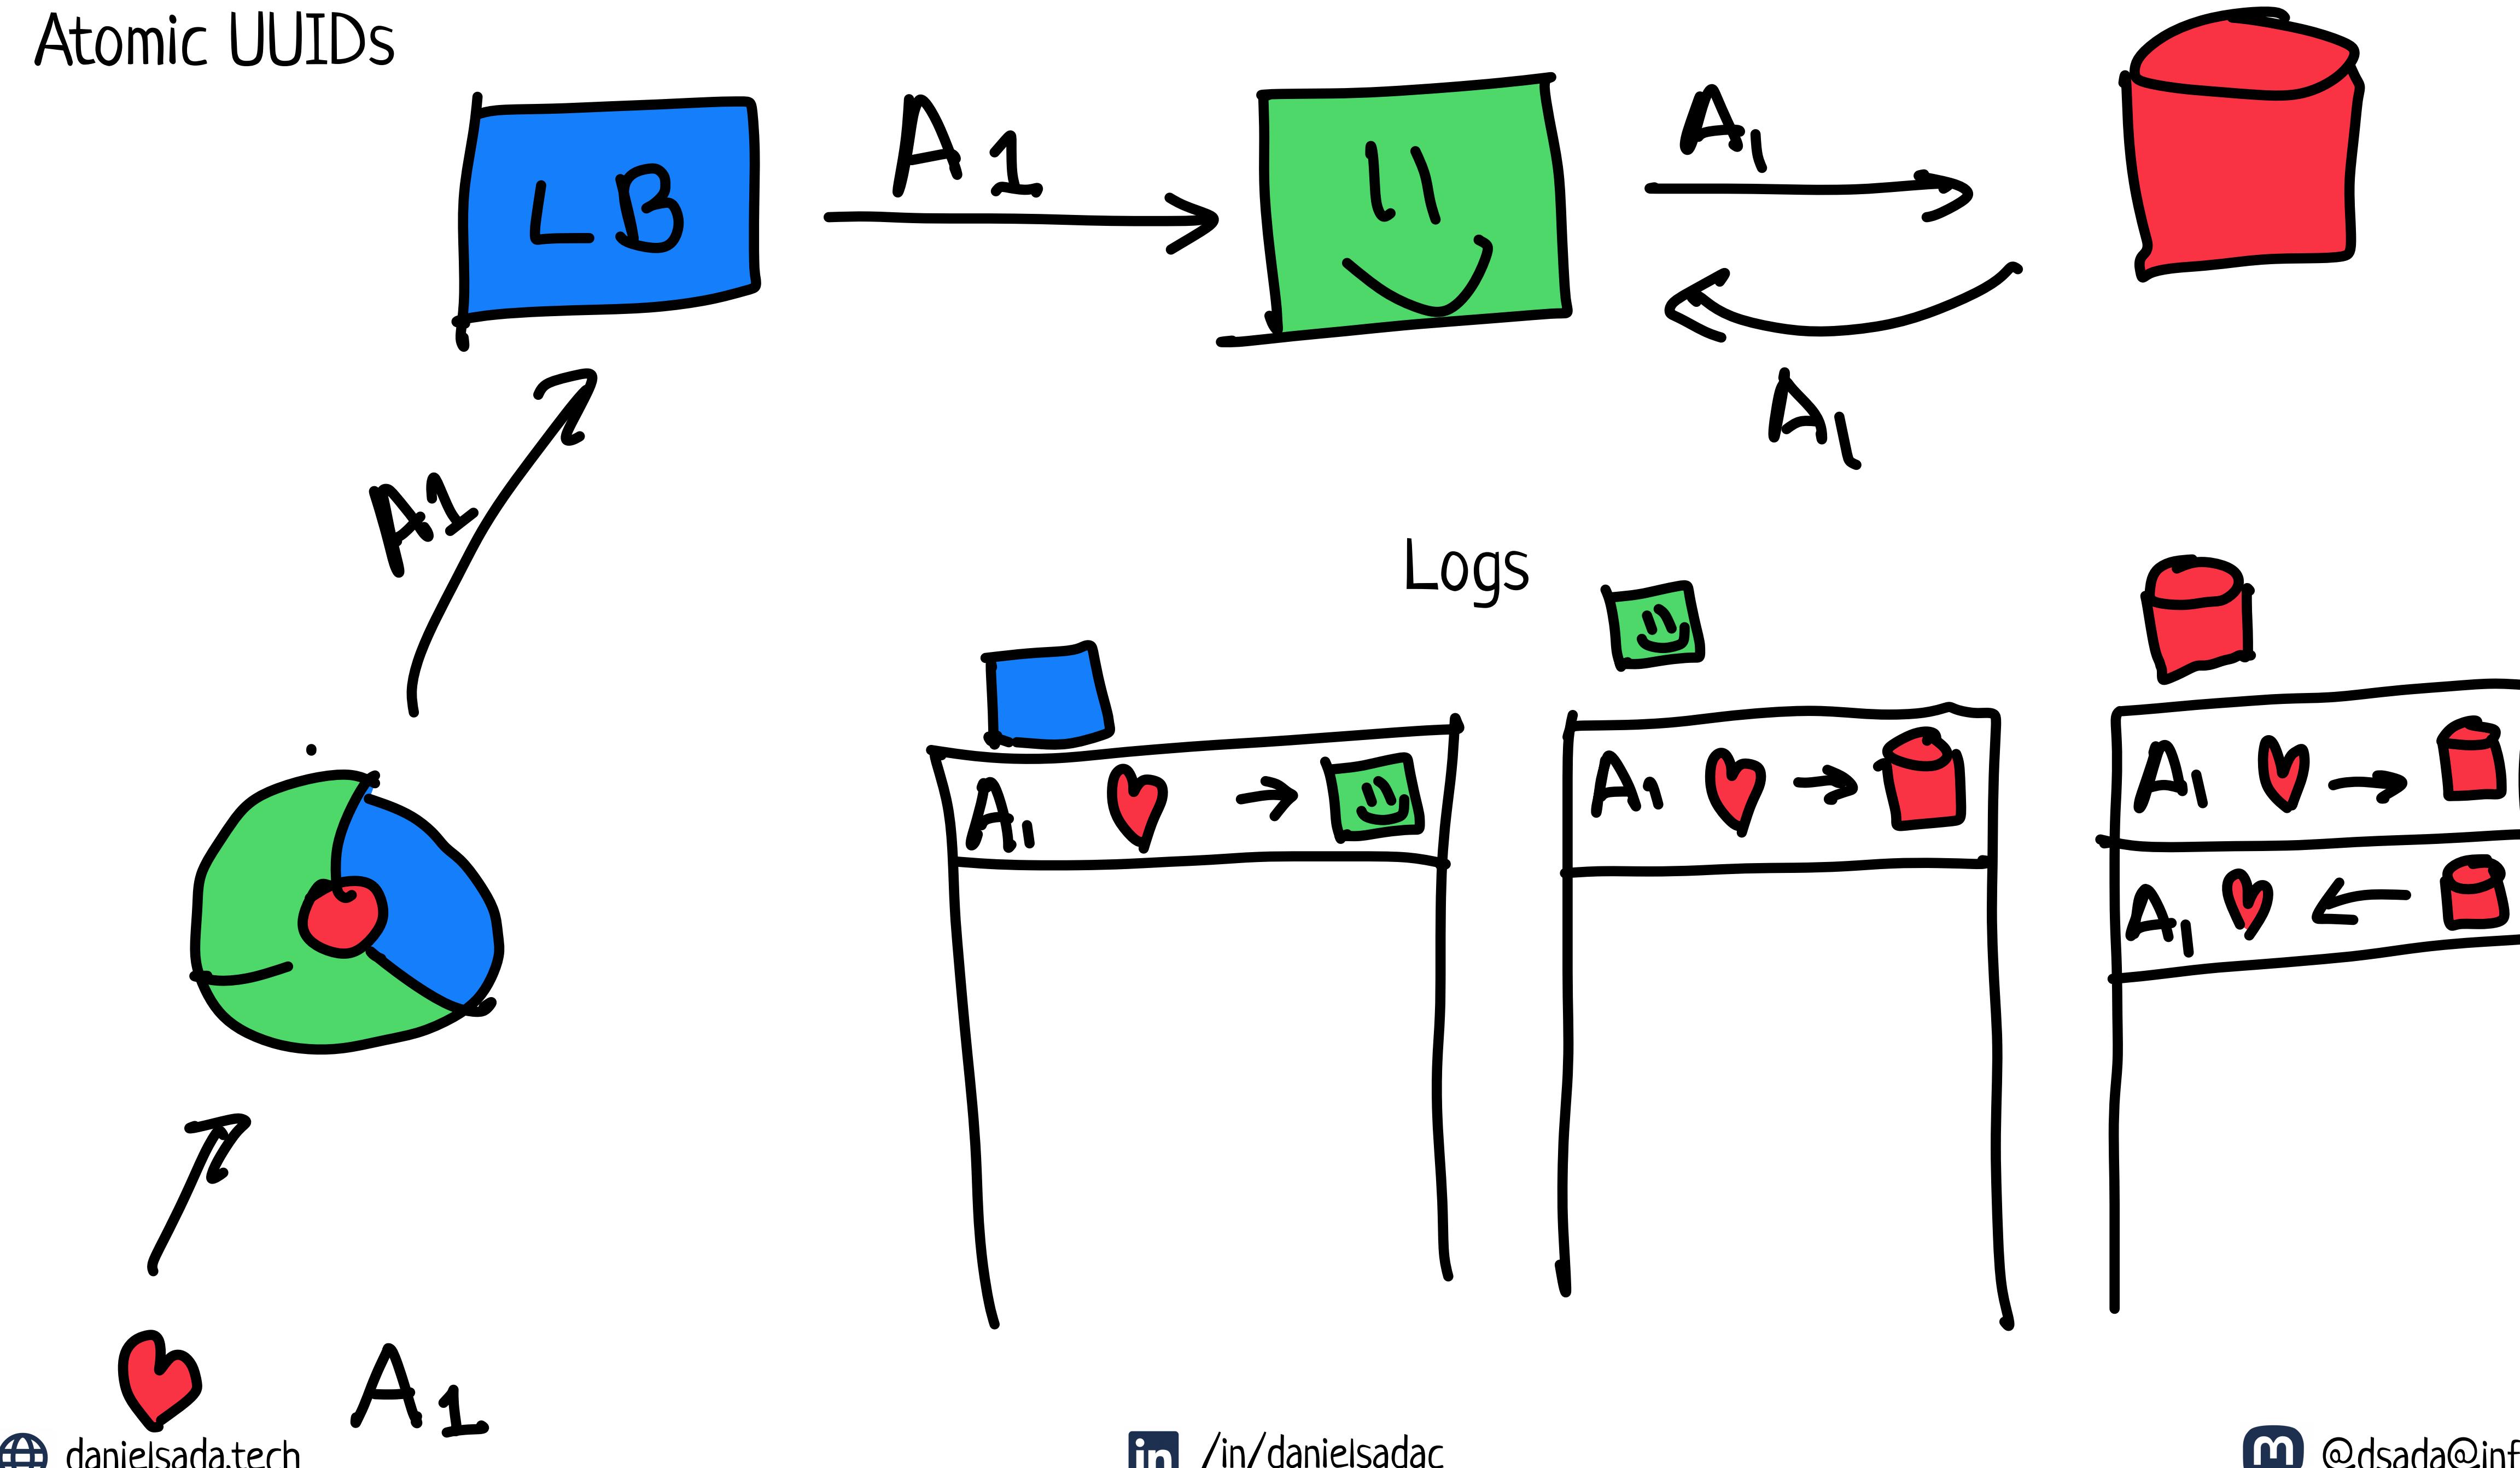
Atomic Uuids



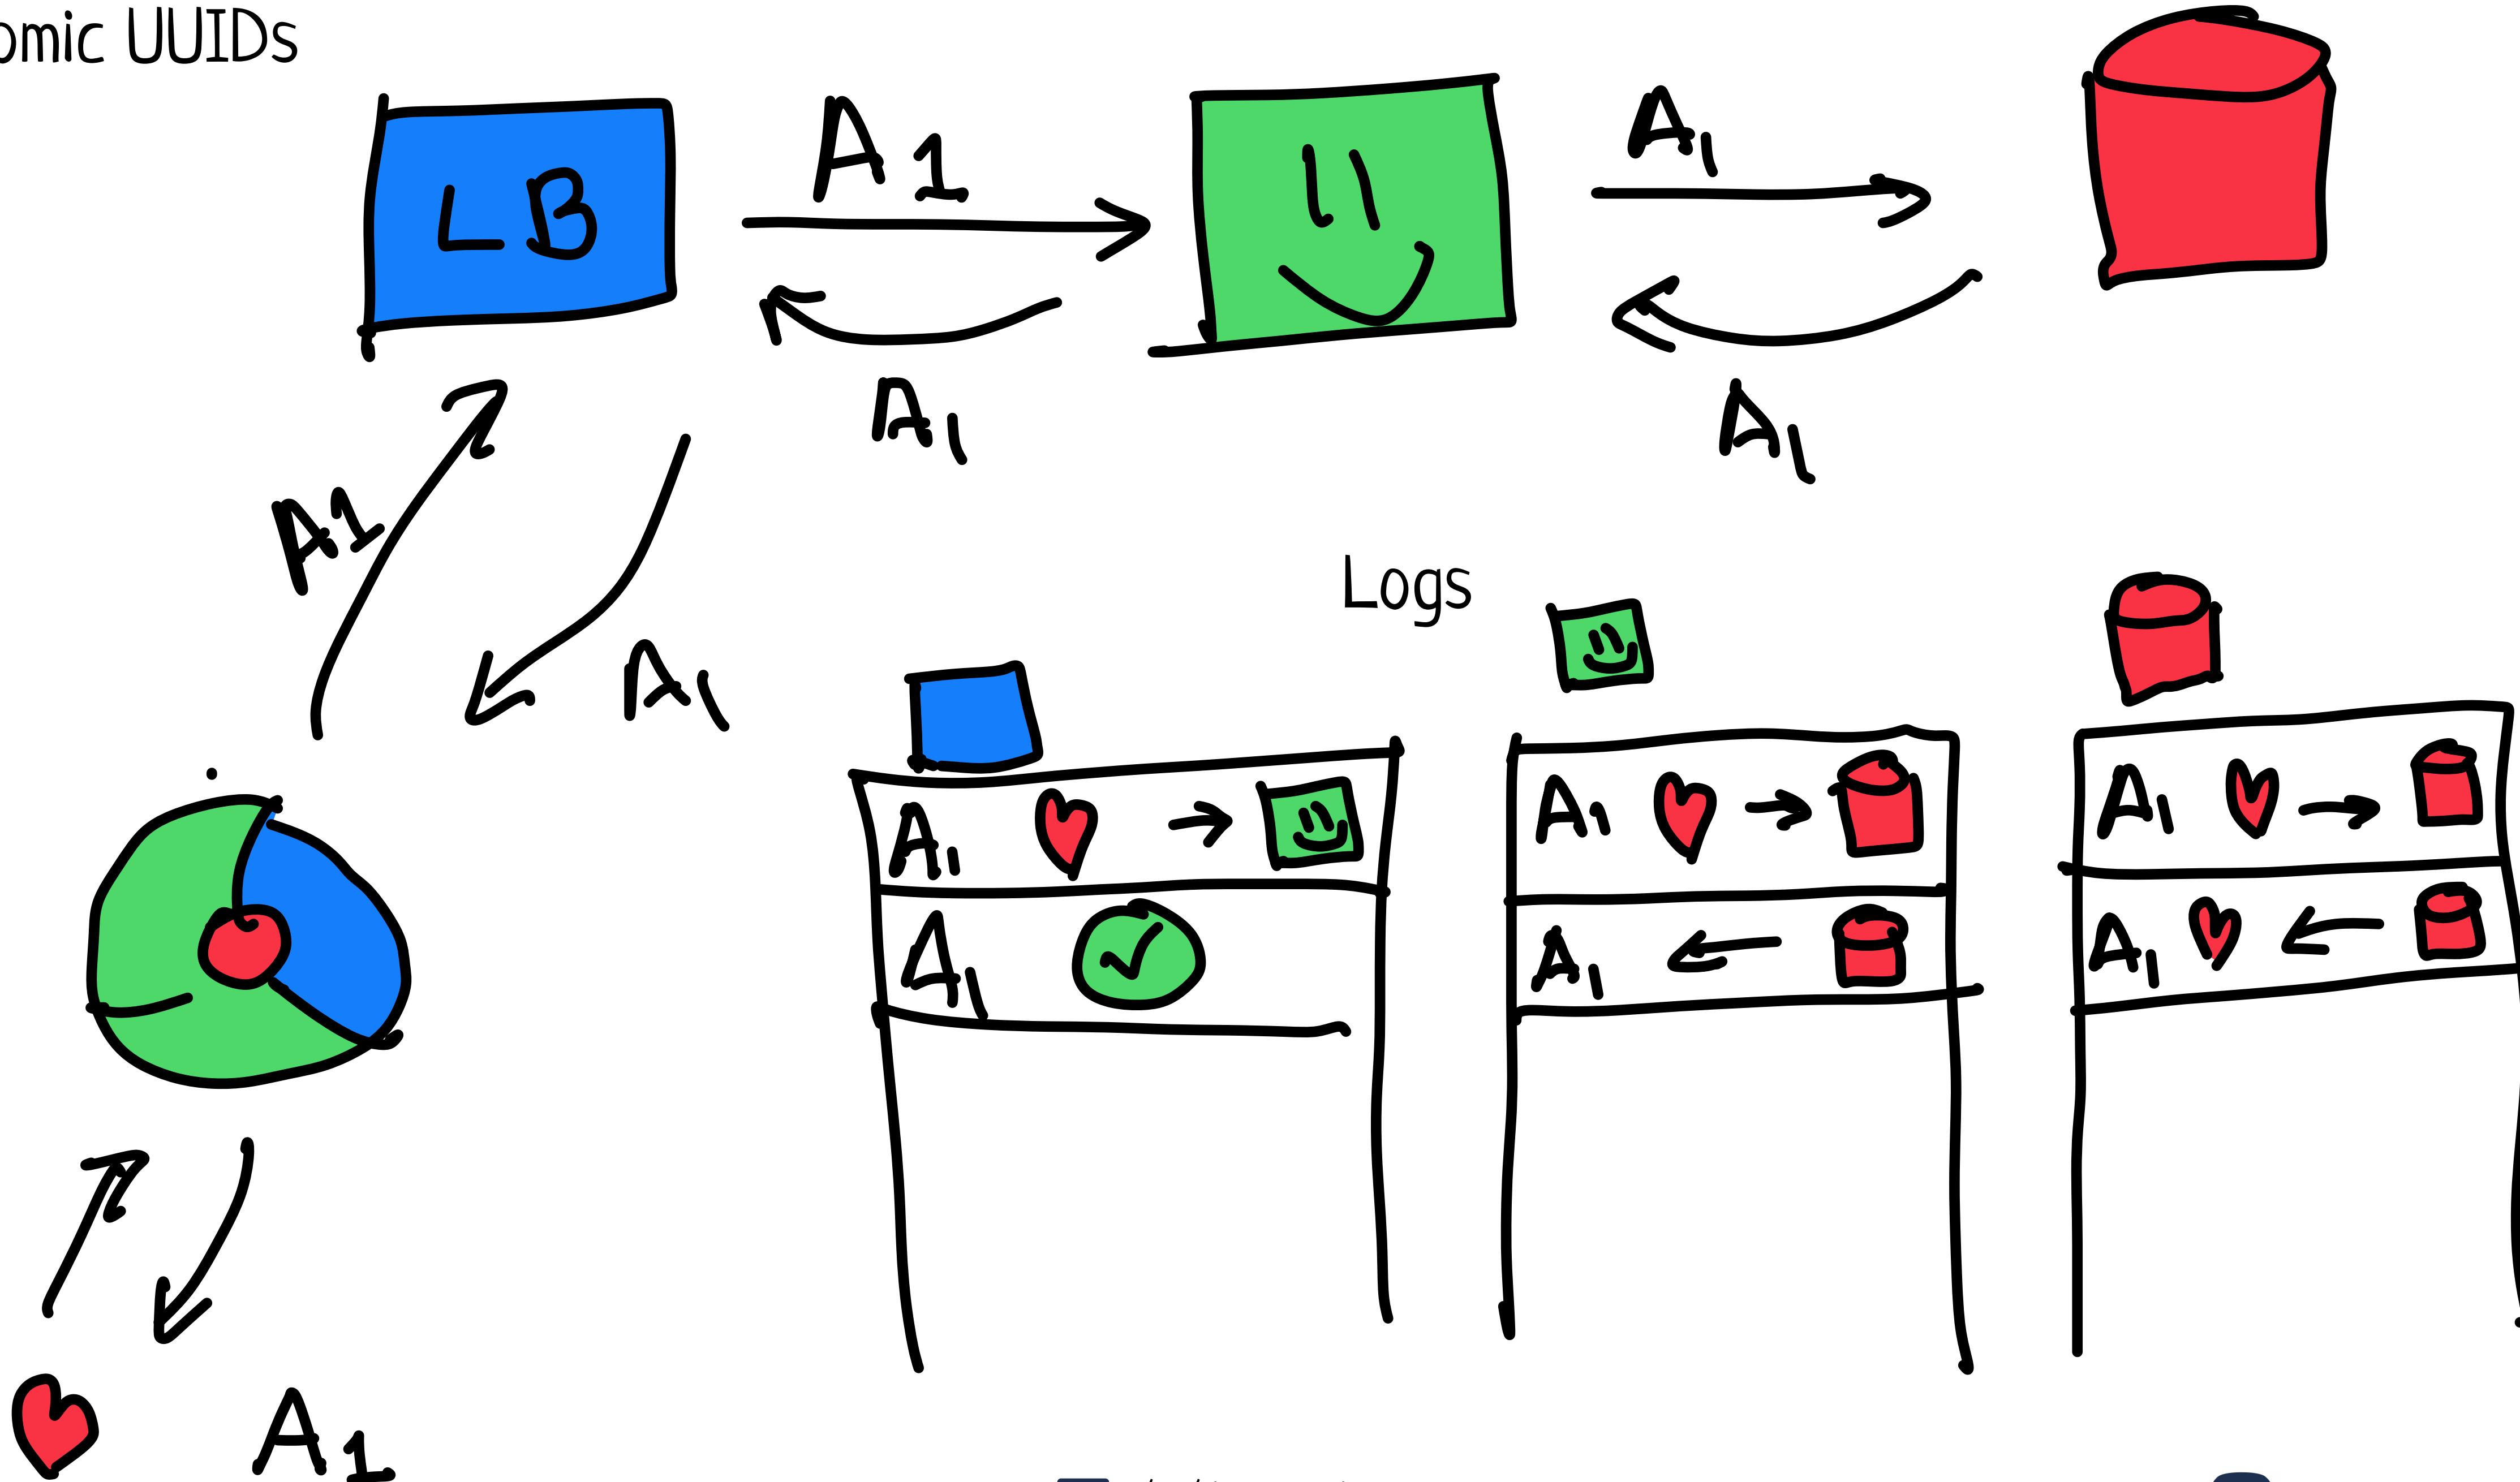
Atomic Uuids



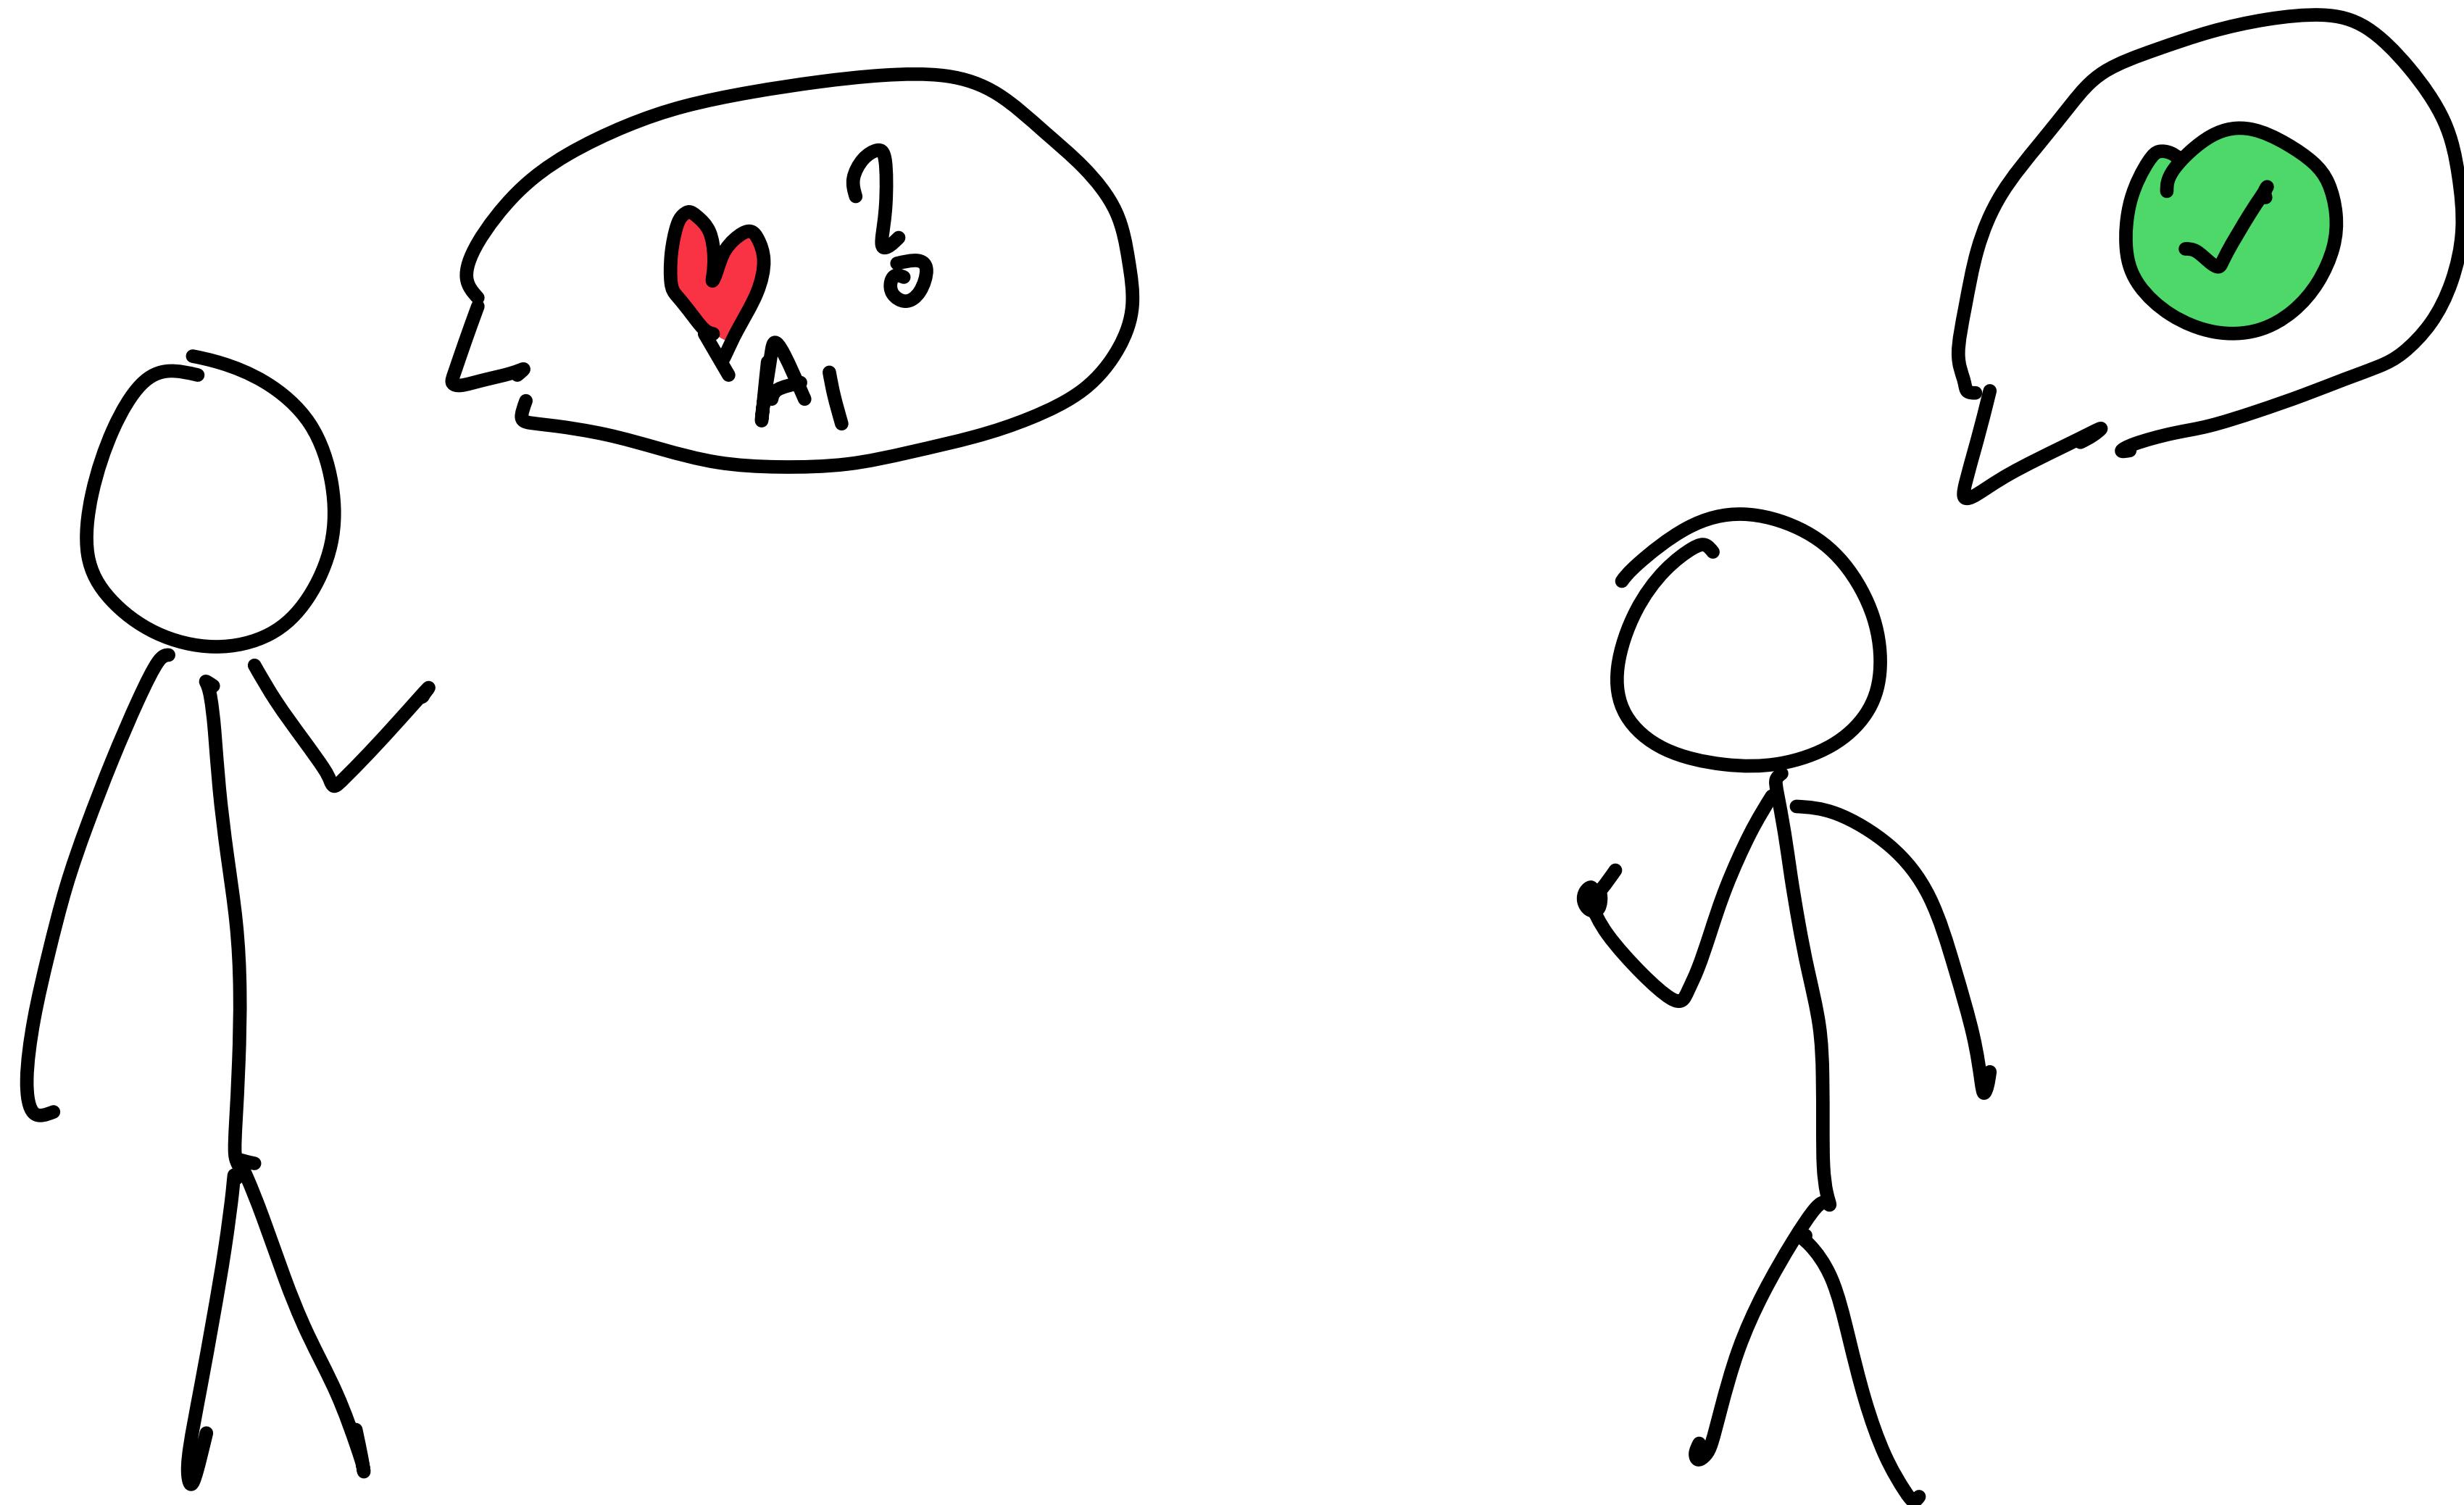
Atomic UUIDs



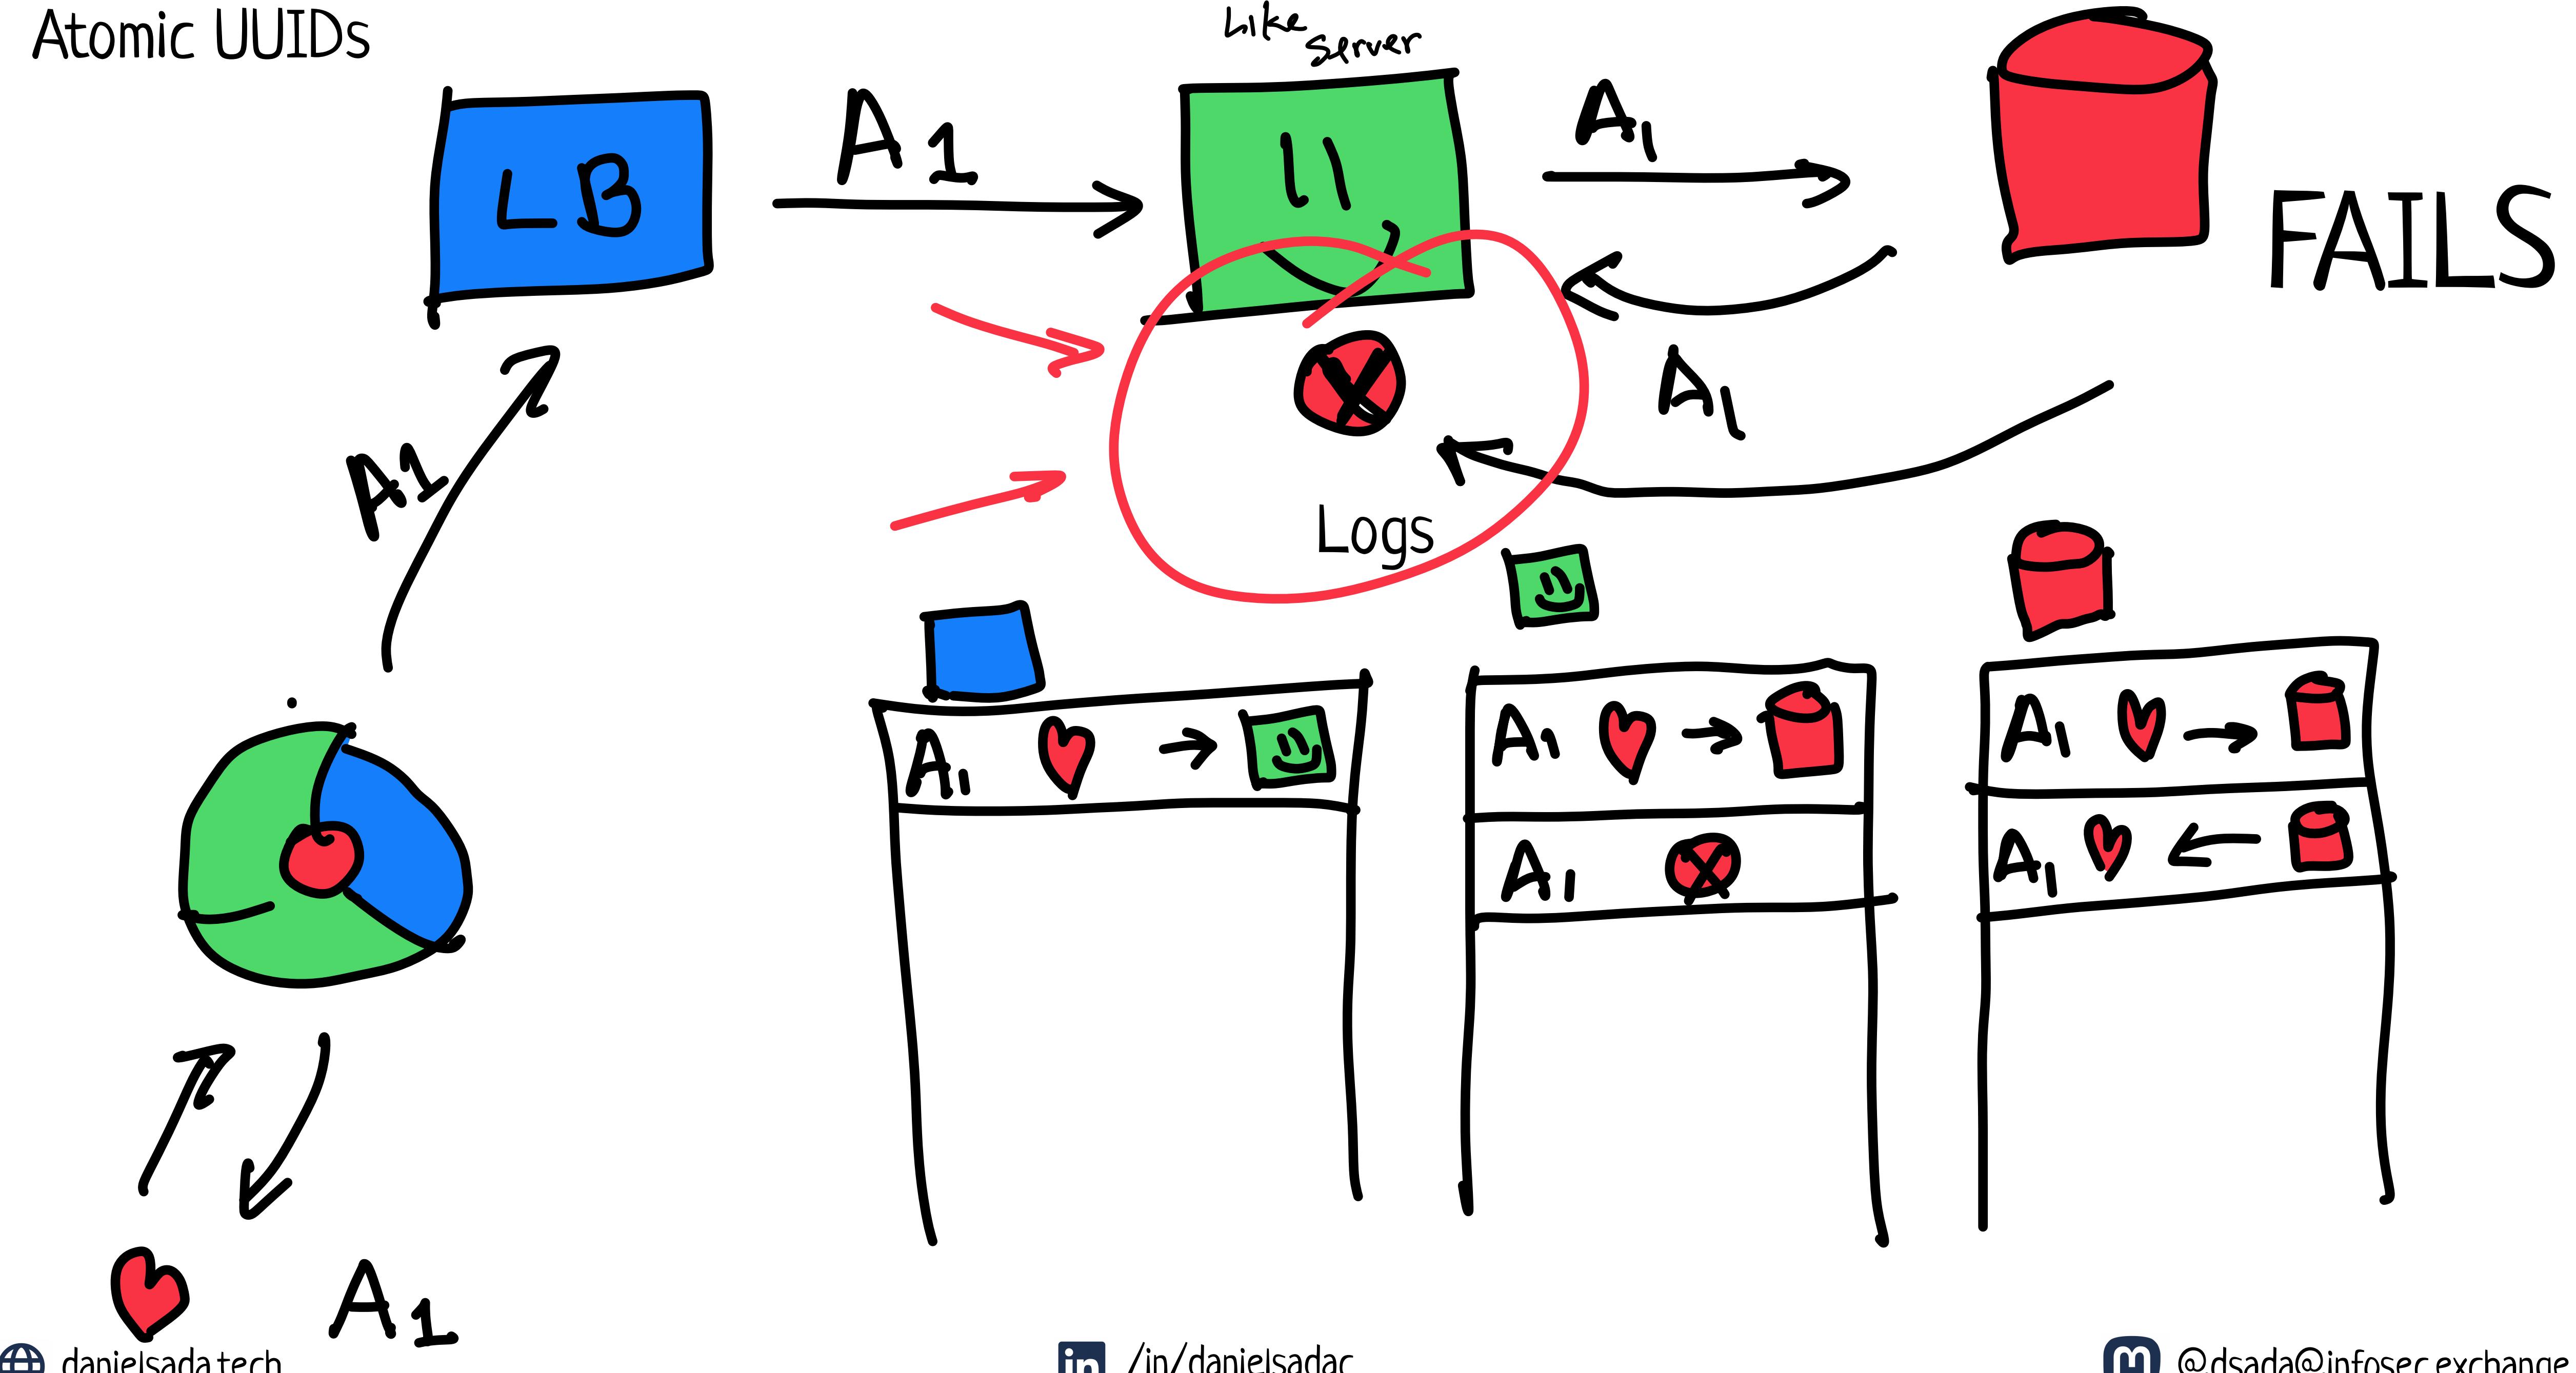
Atomic UUIDs



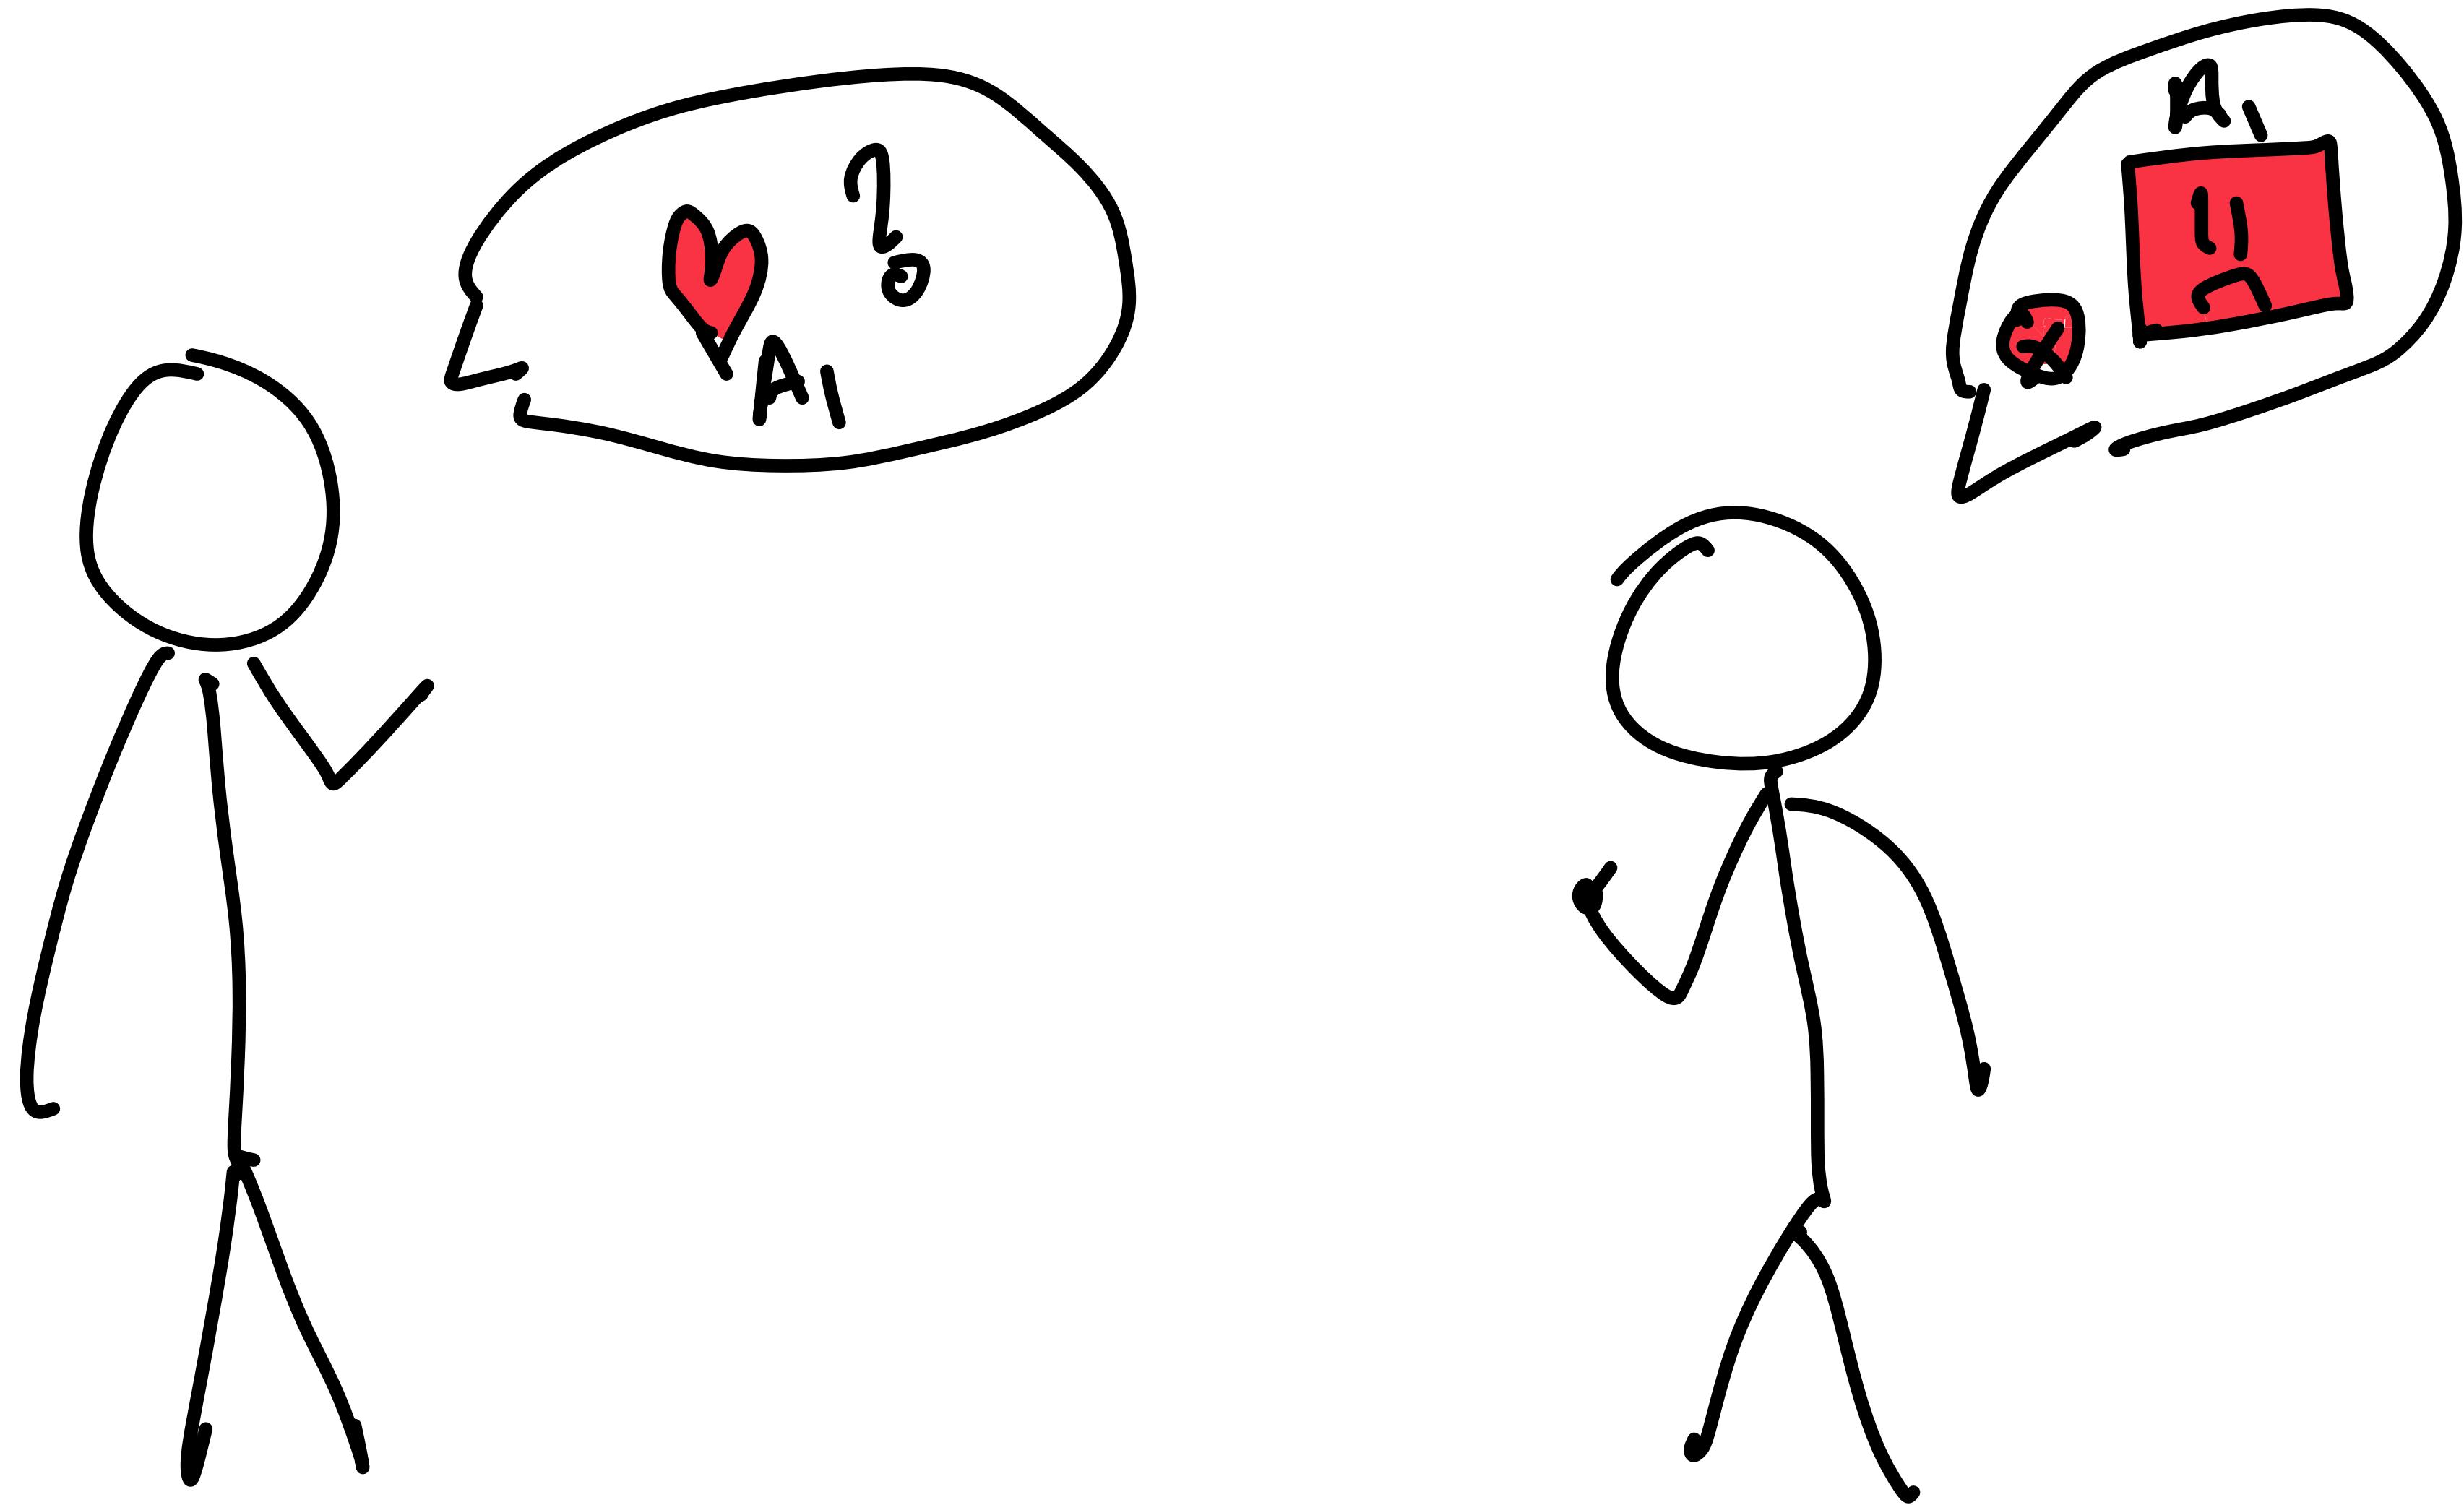
Atomic Uuids



Atomic UUIDs



Atomic Uuids



UUID

Universal Unique Identifier

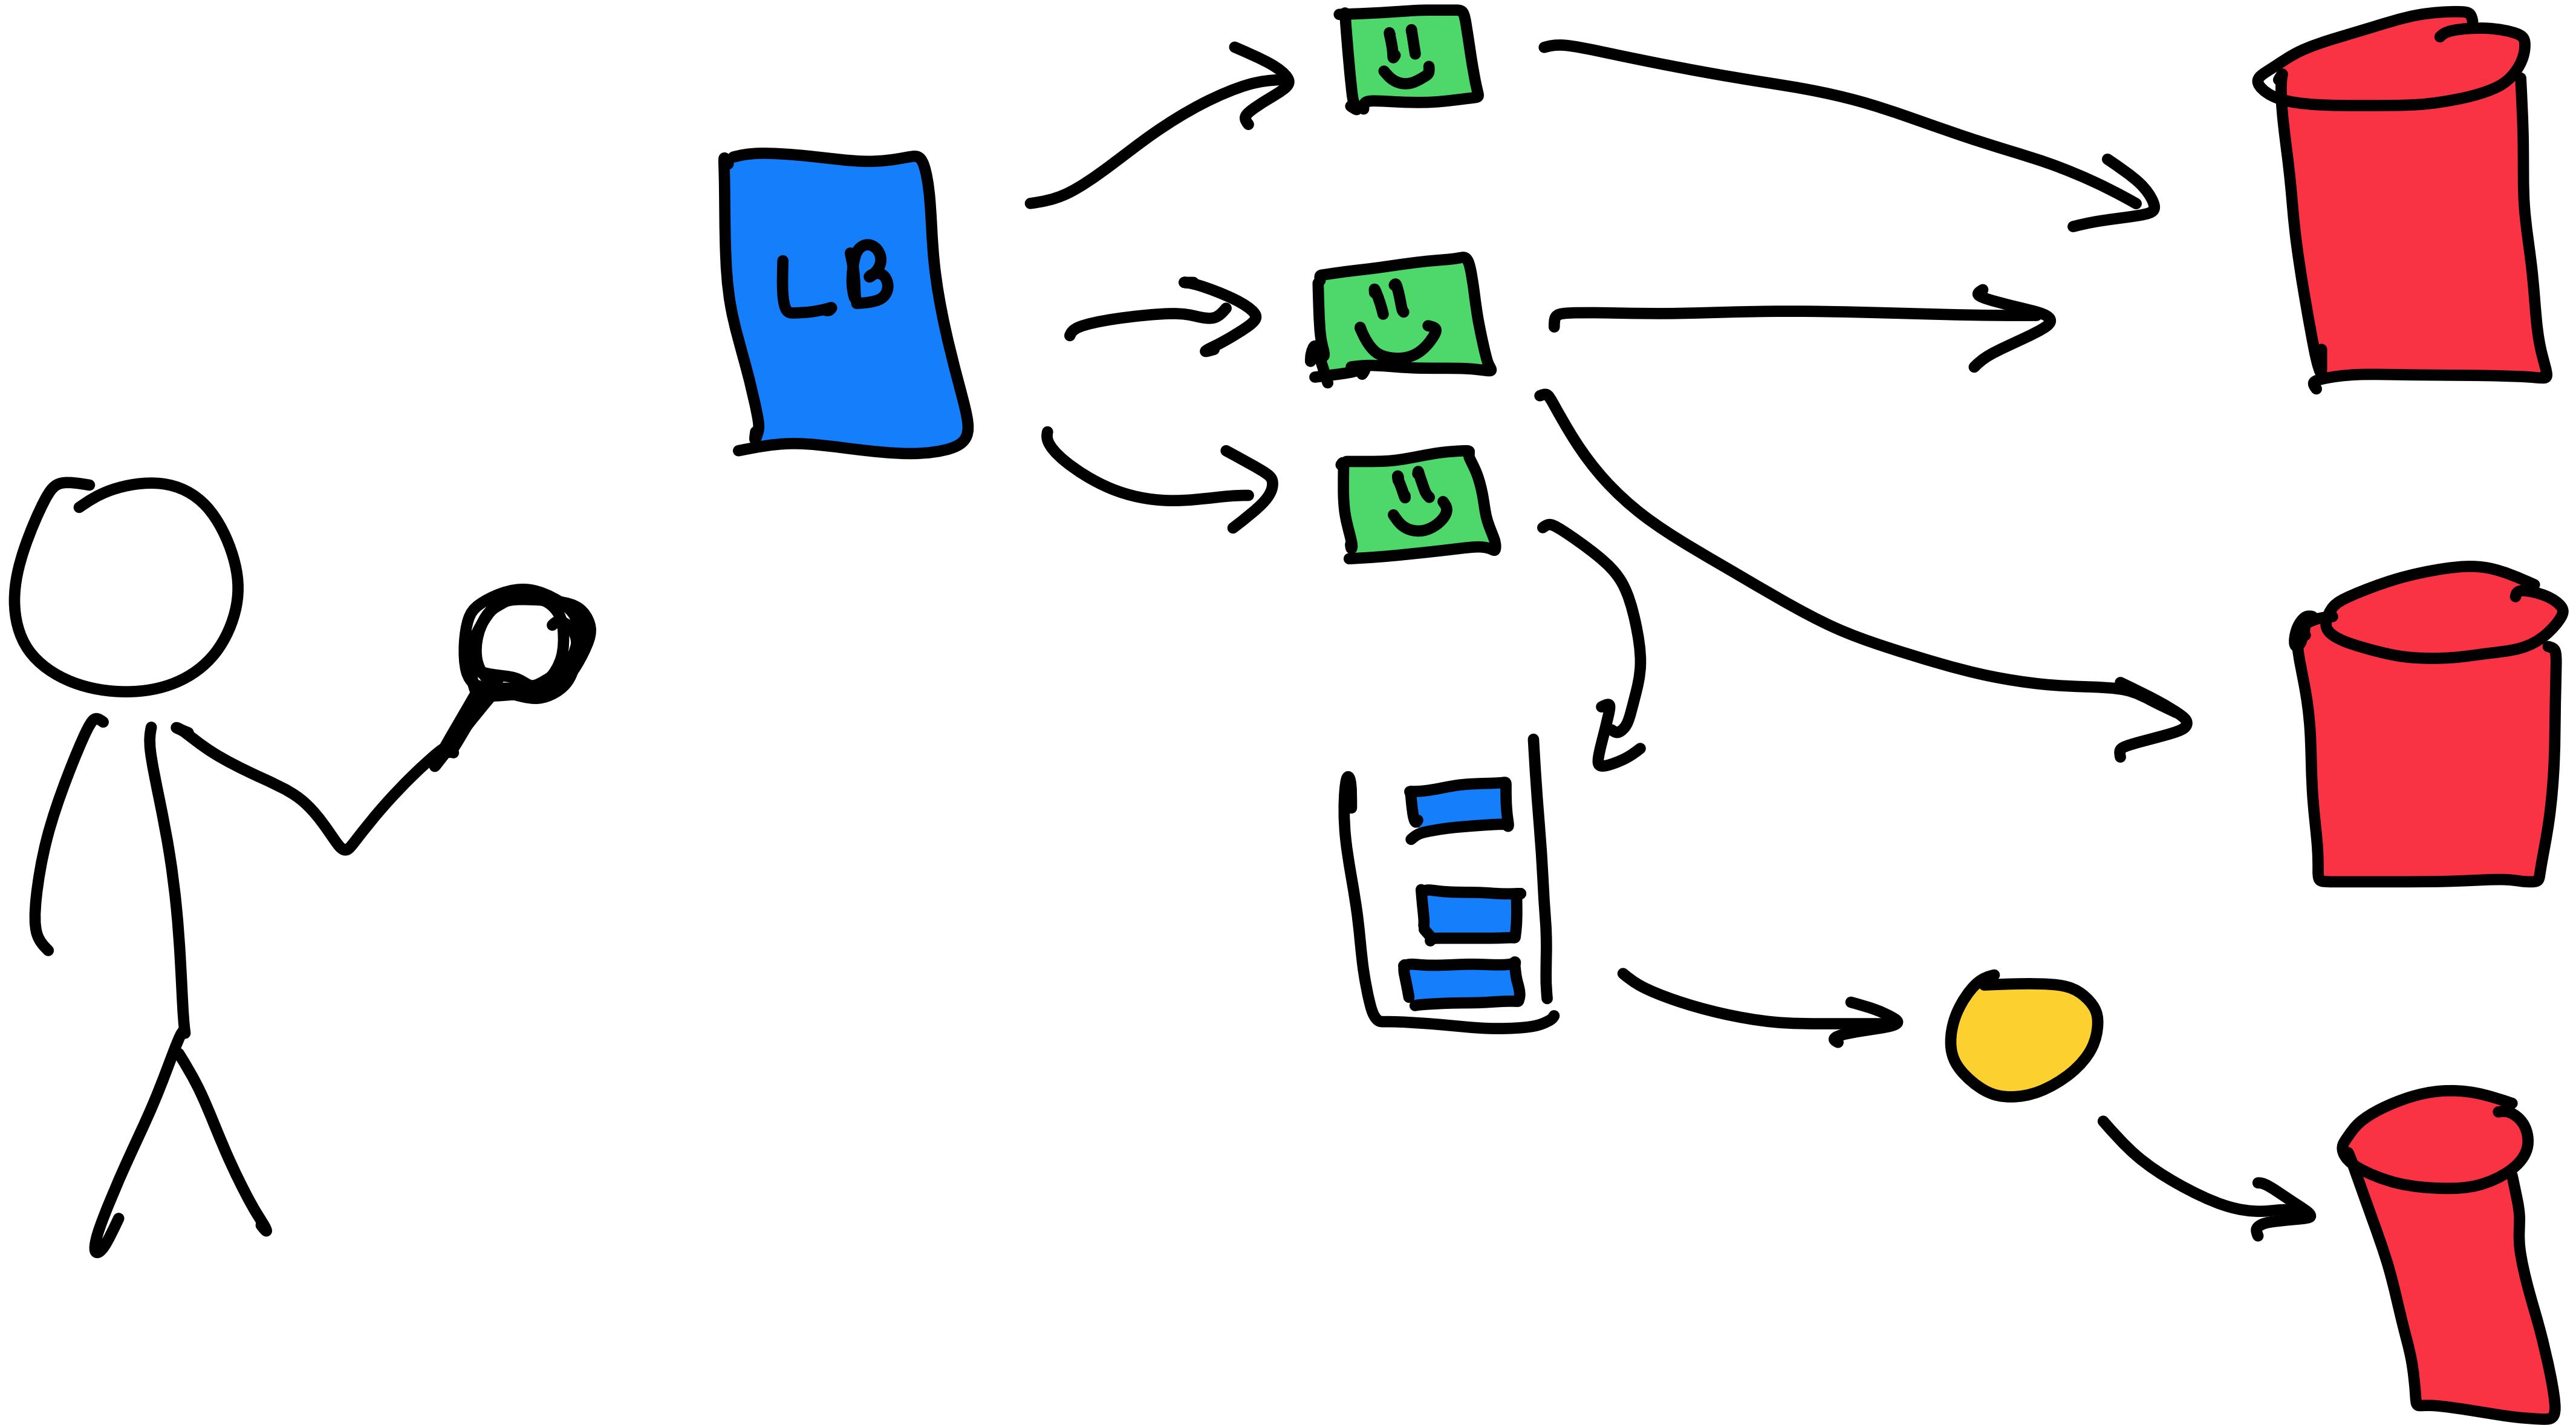
- Implement it yourself
- Avoid generating your own IDs
- You can generate them client side.

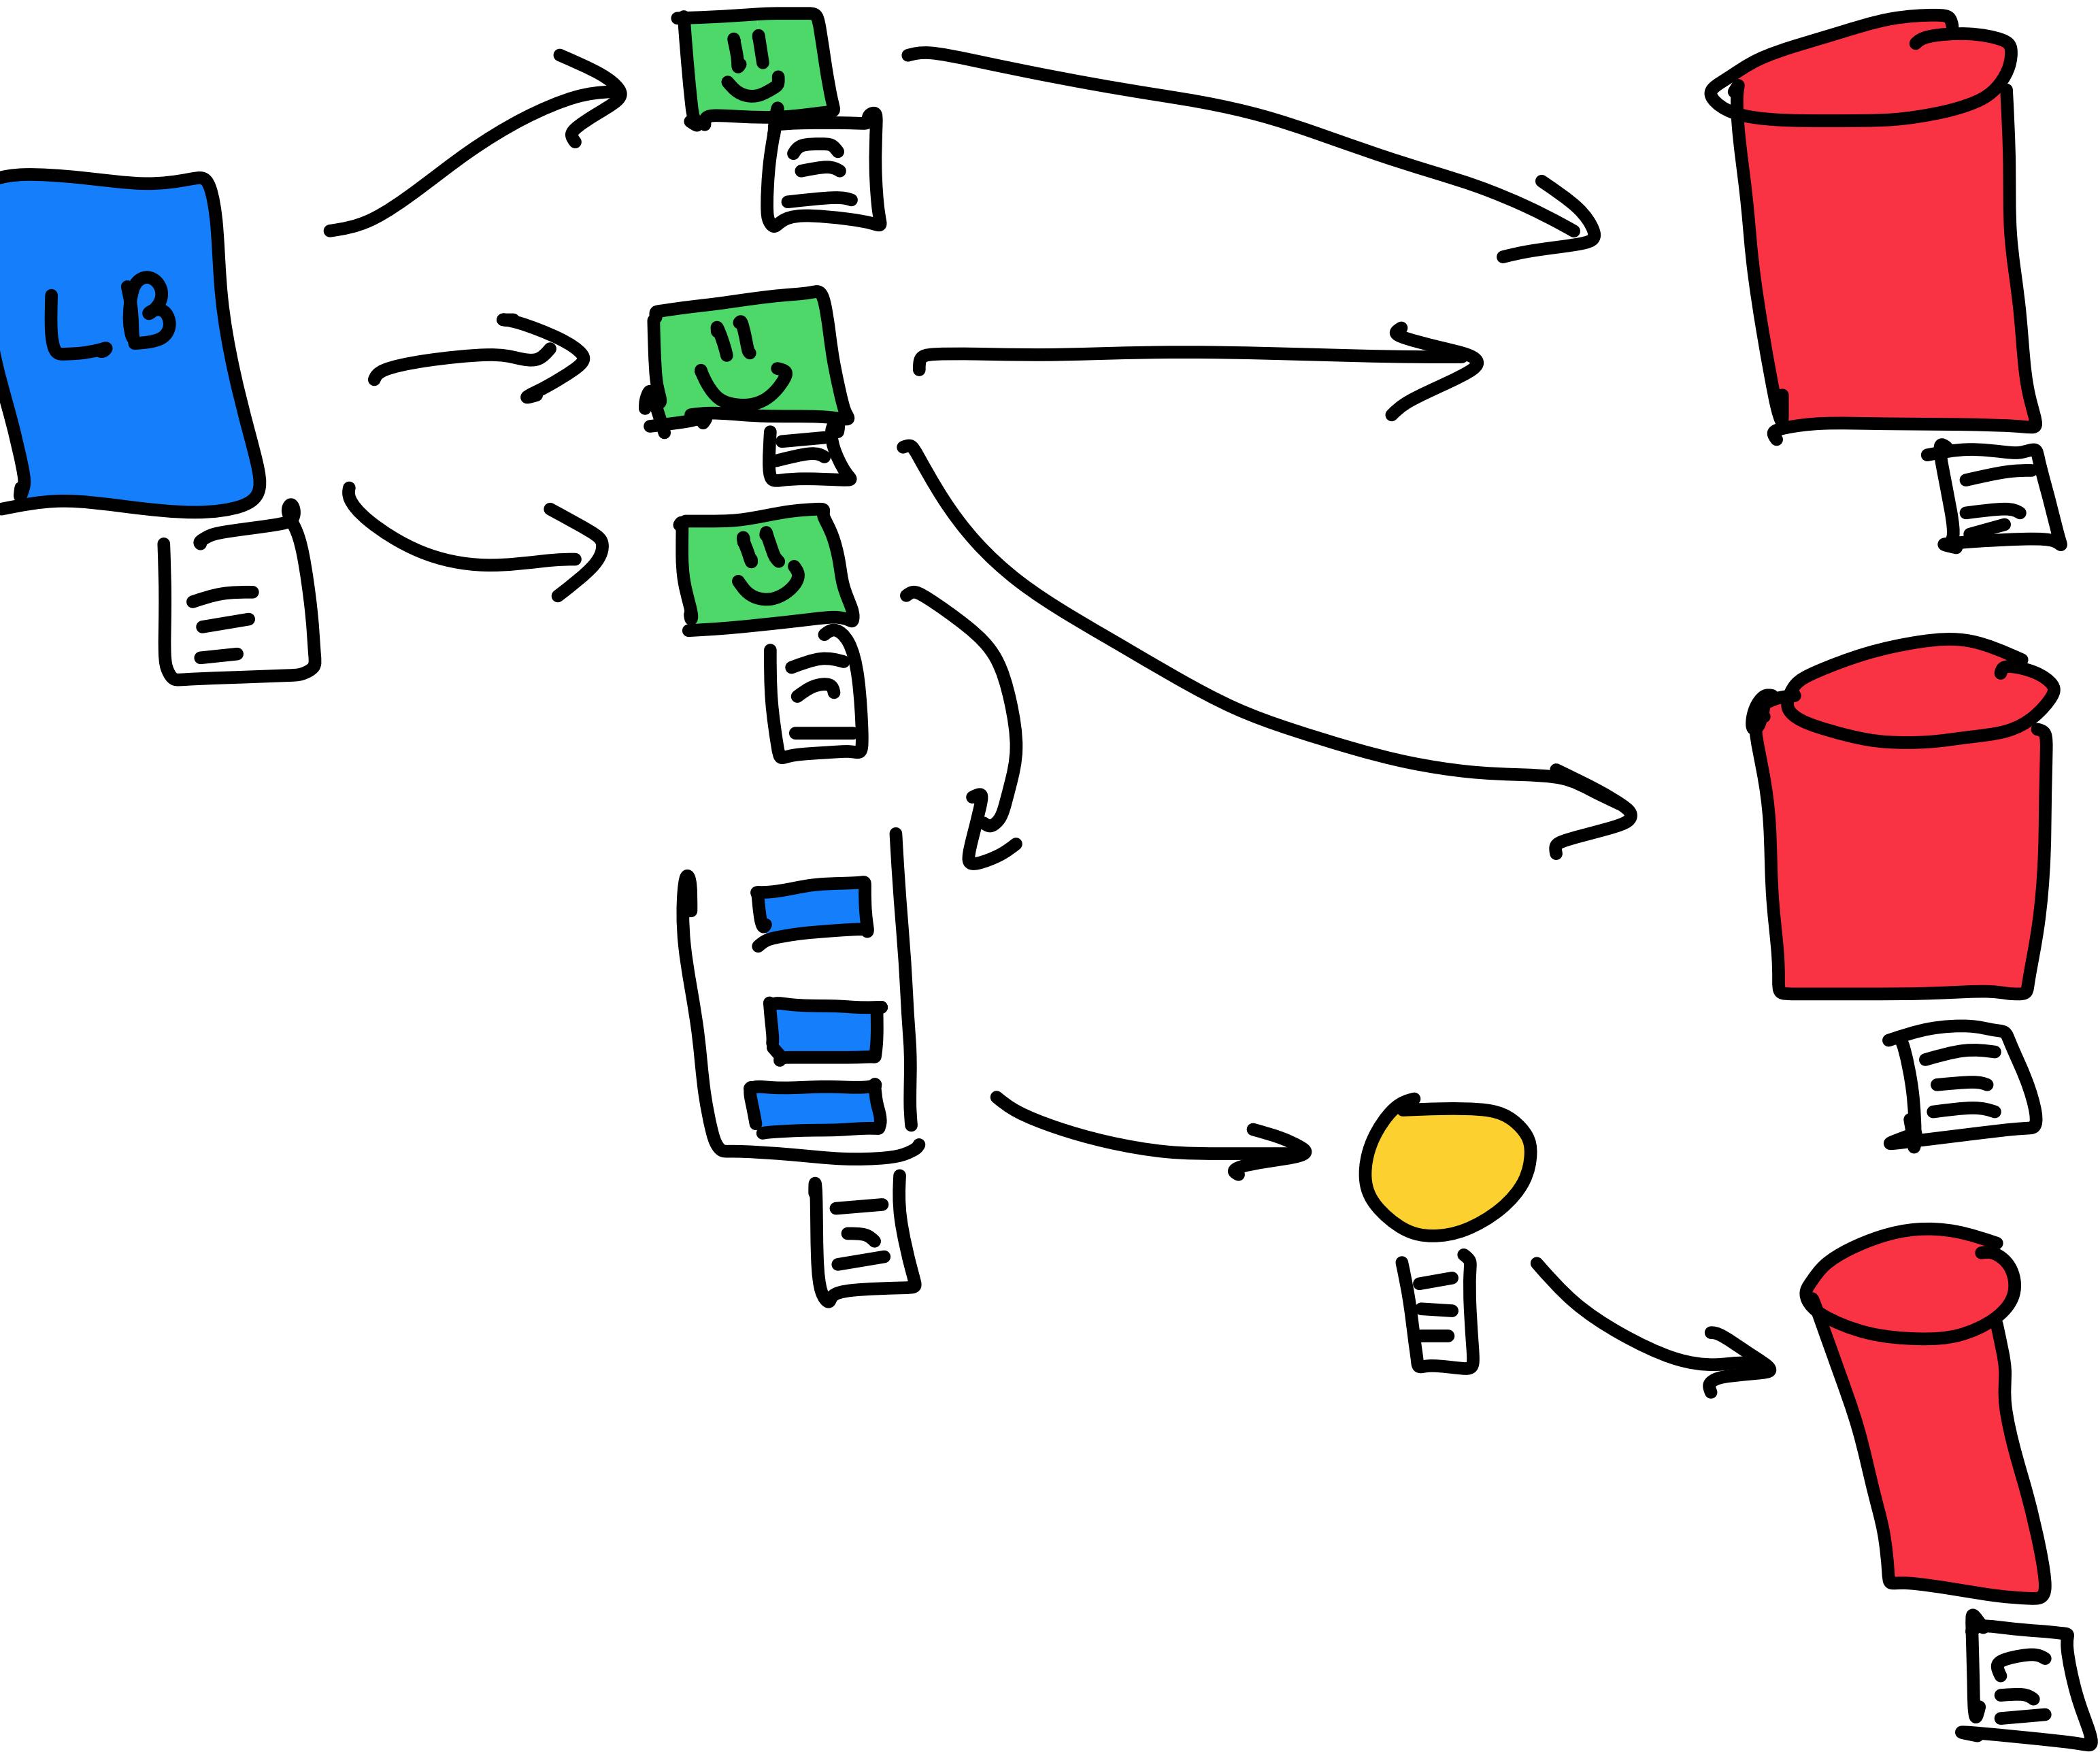
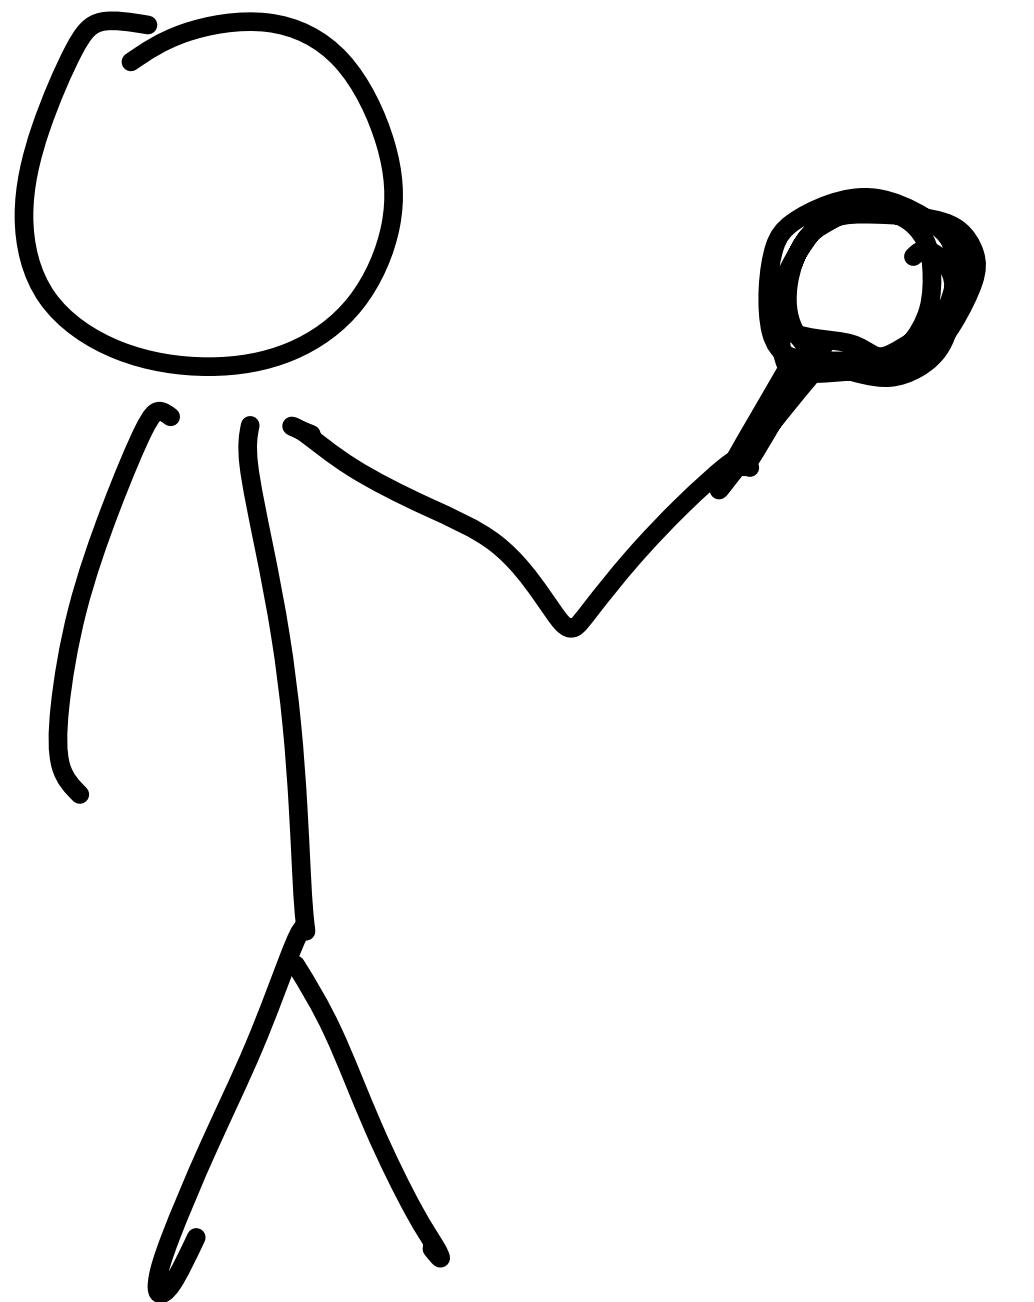
Medium

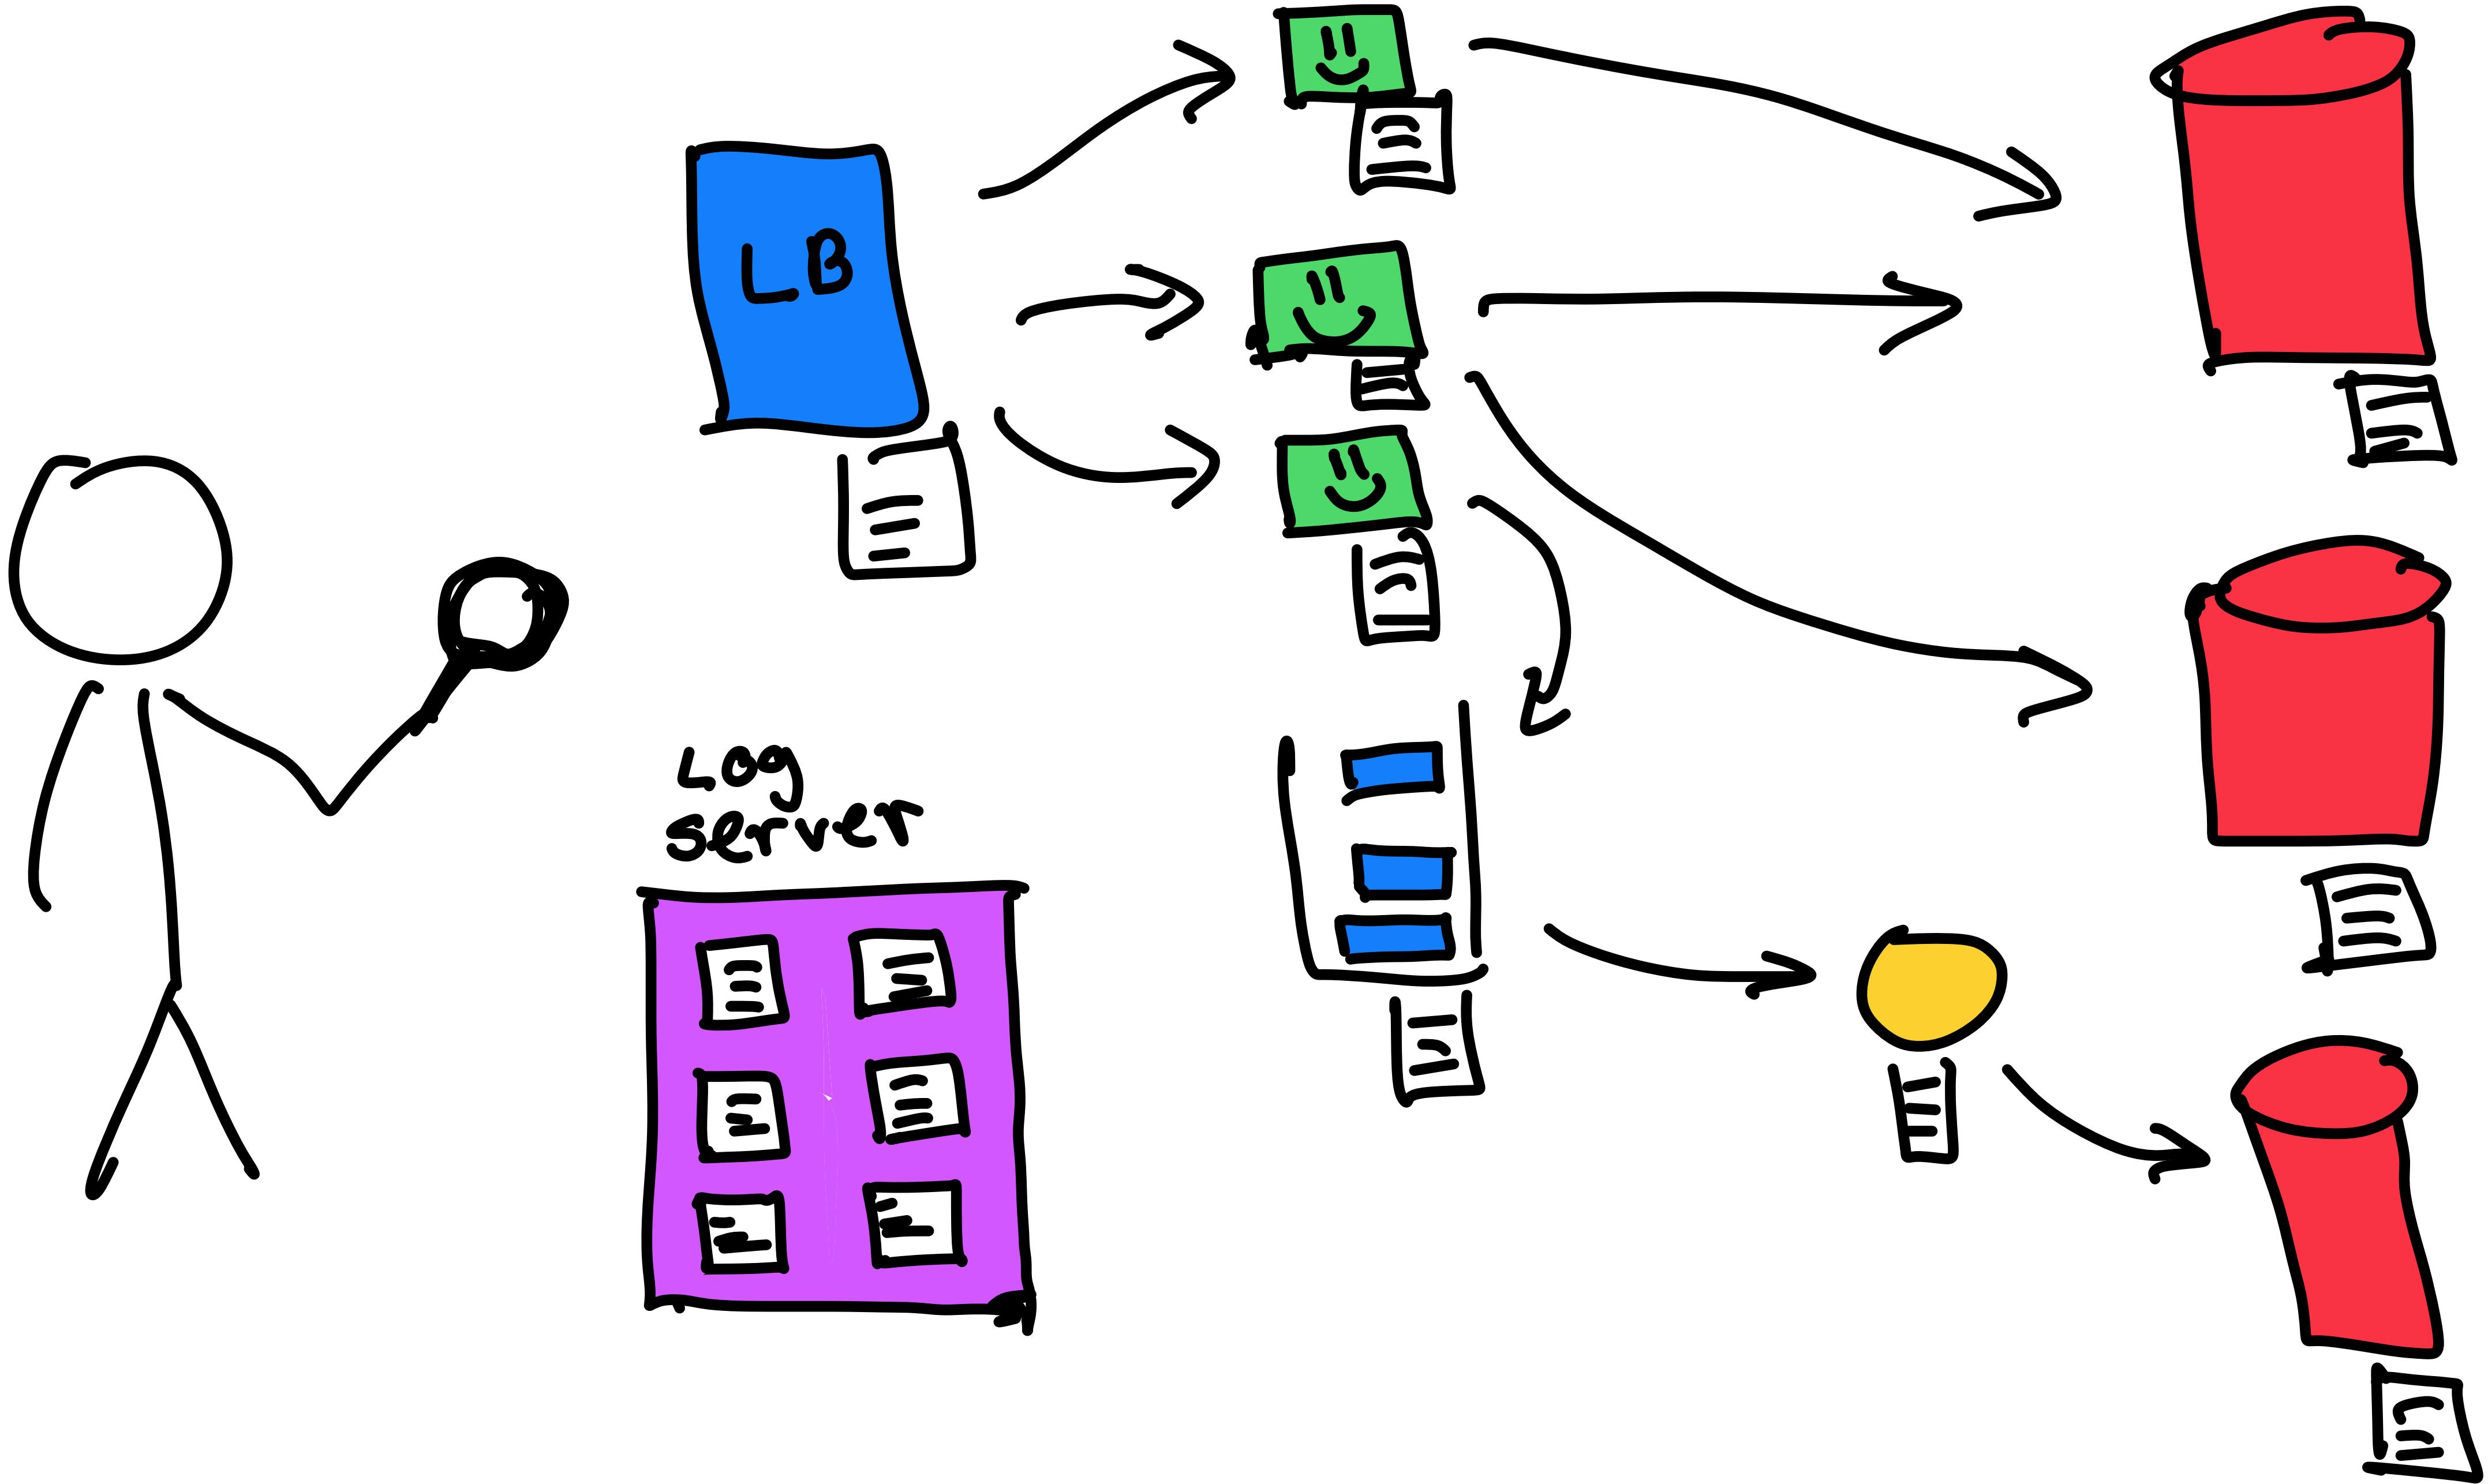


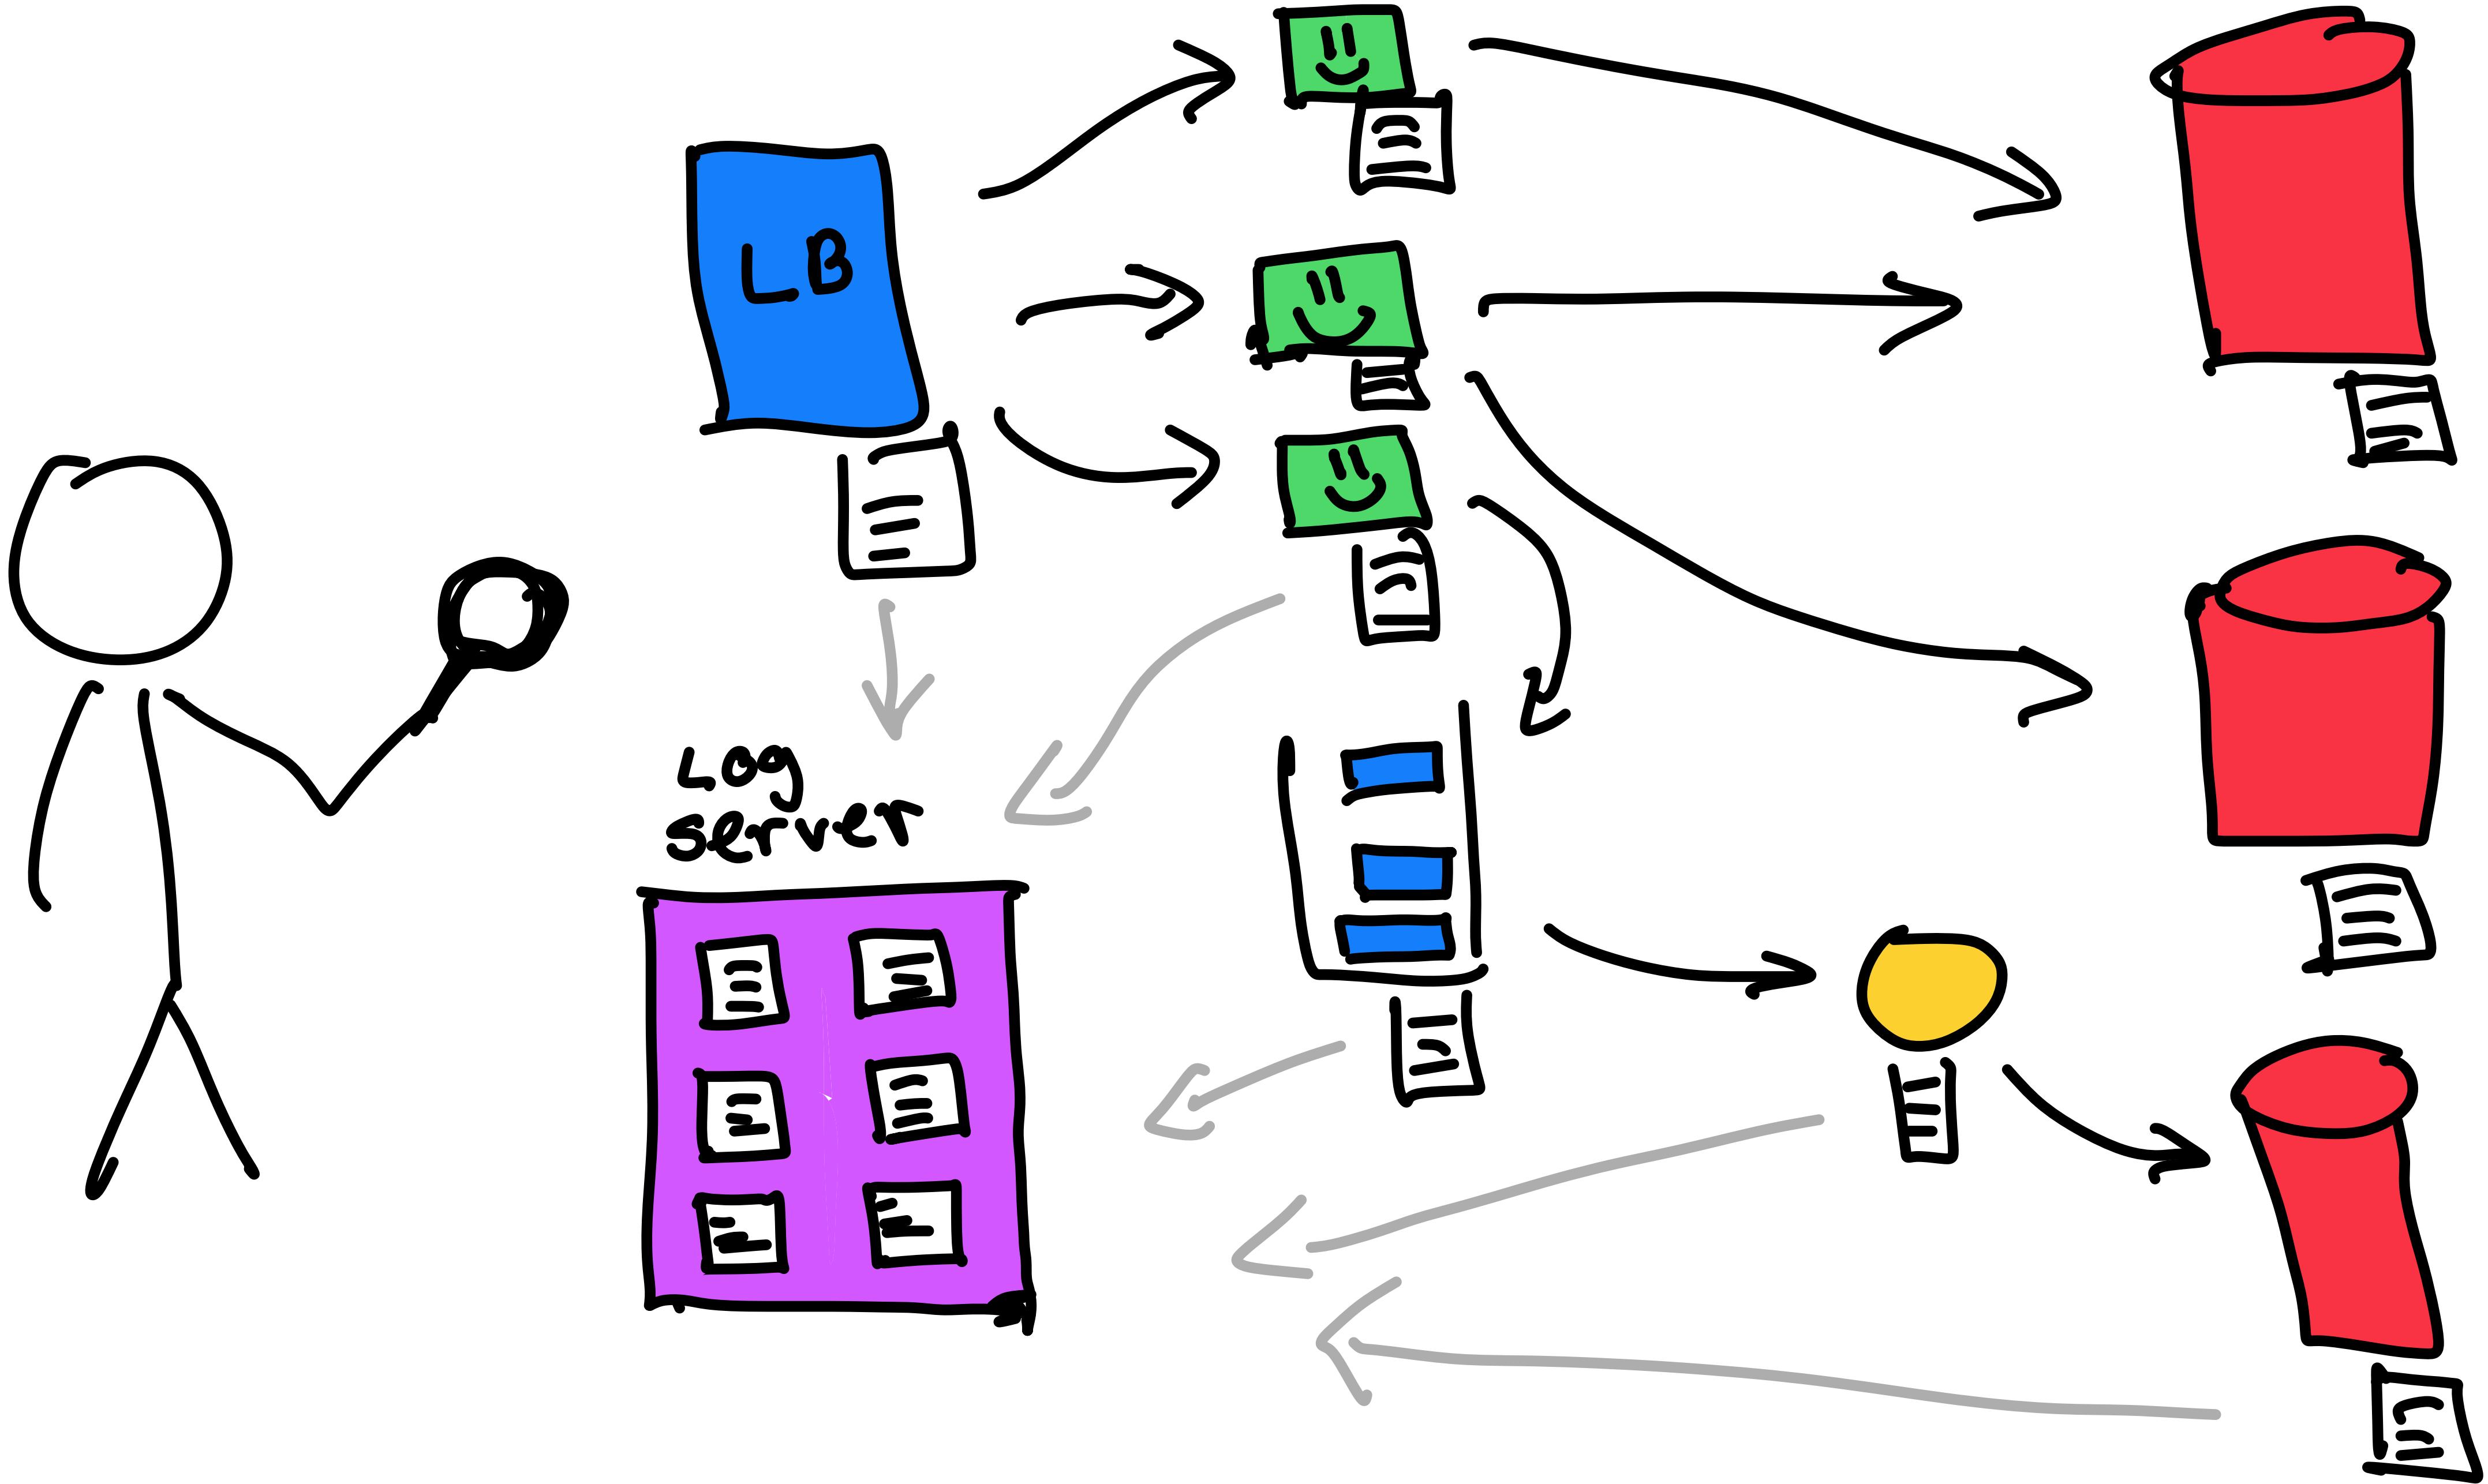
Starts getting spicy

Centralized Logging

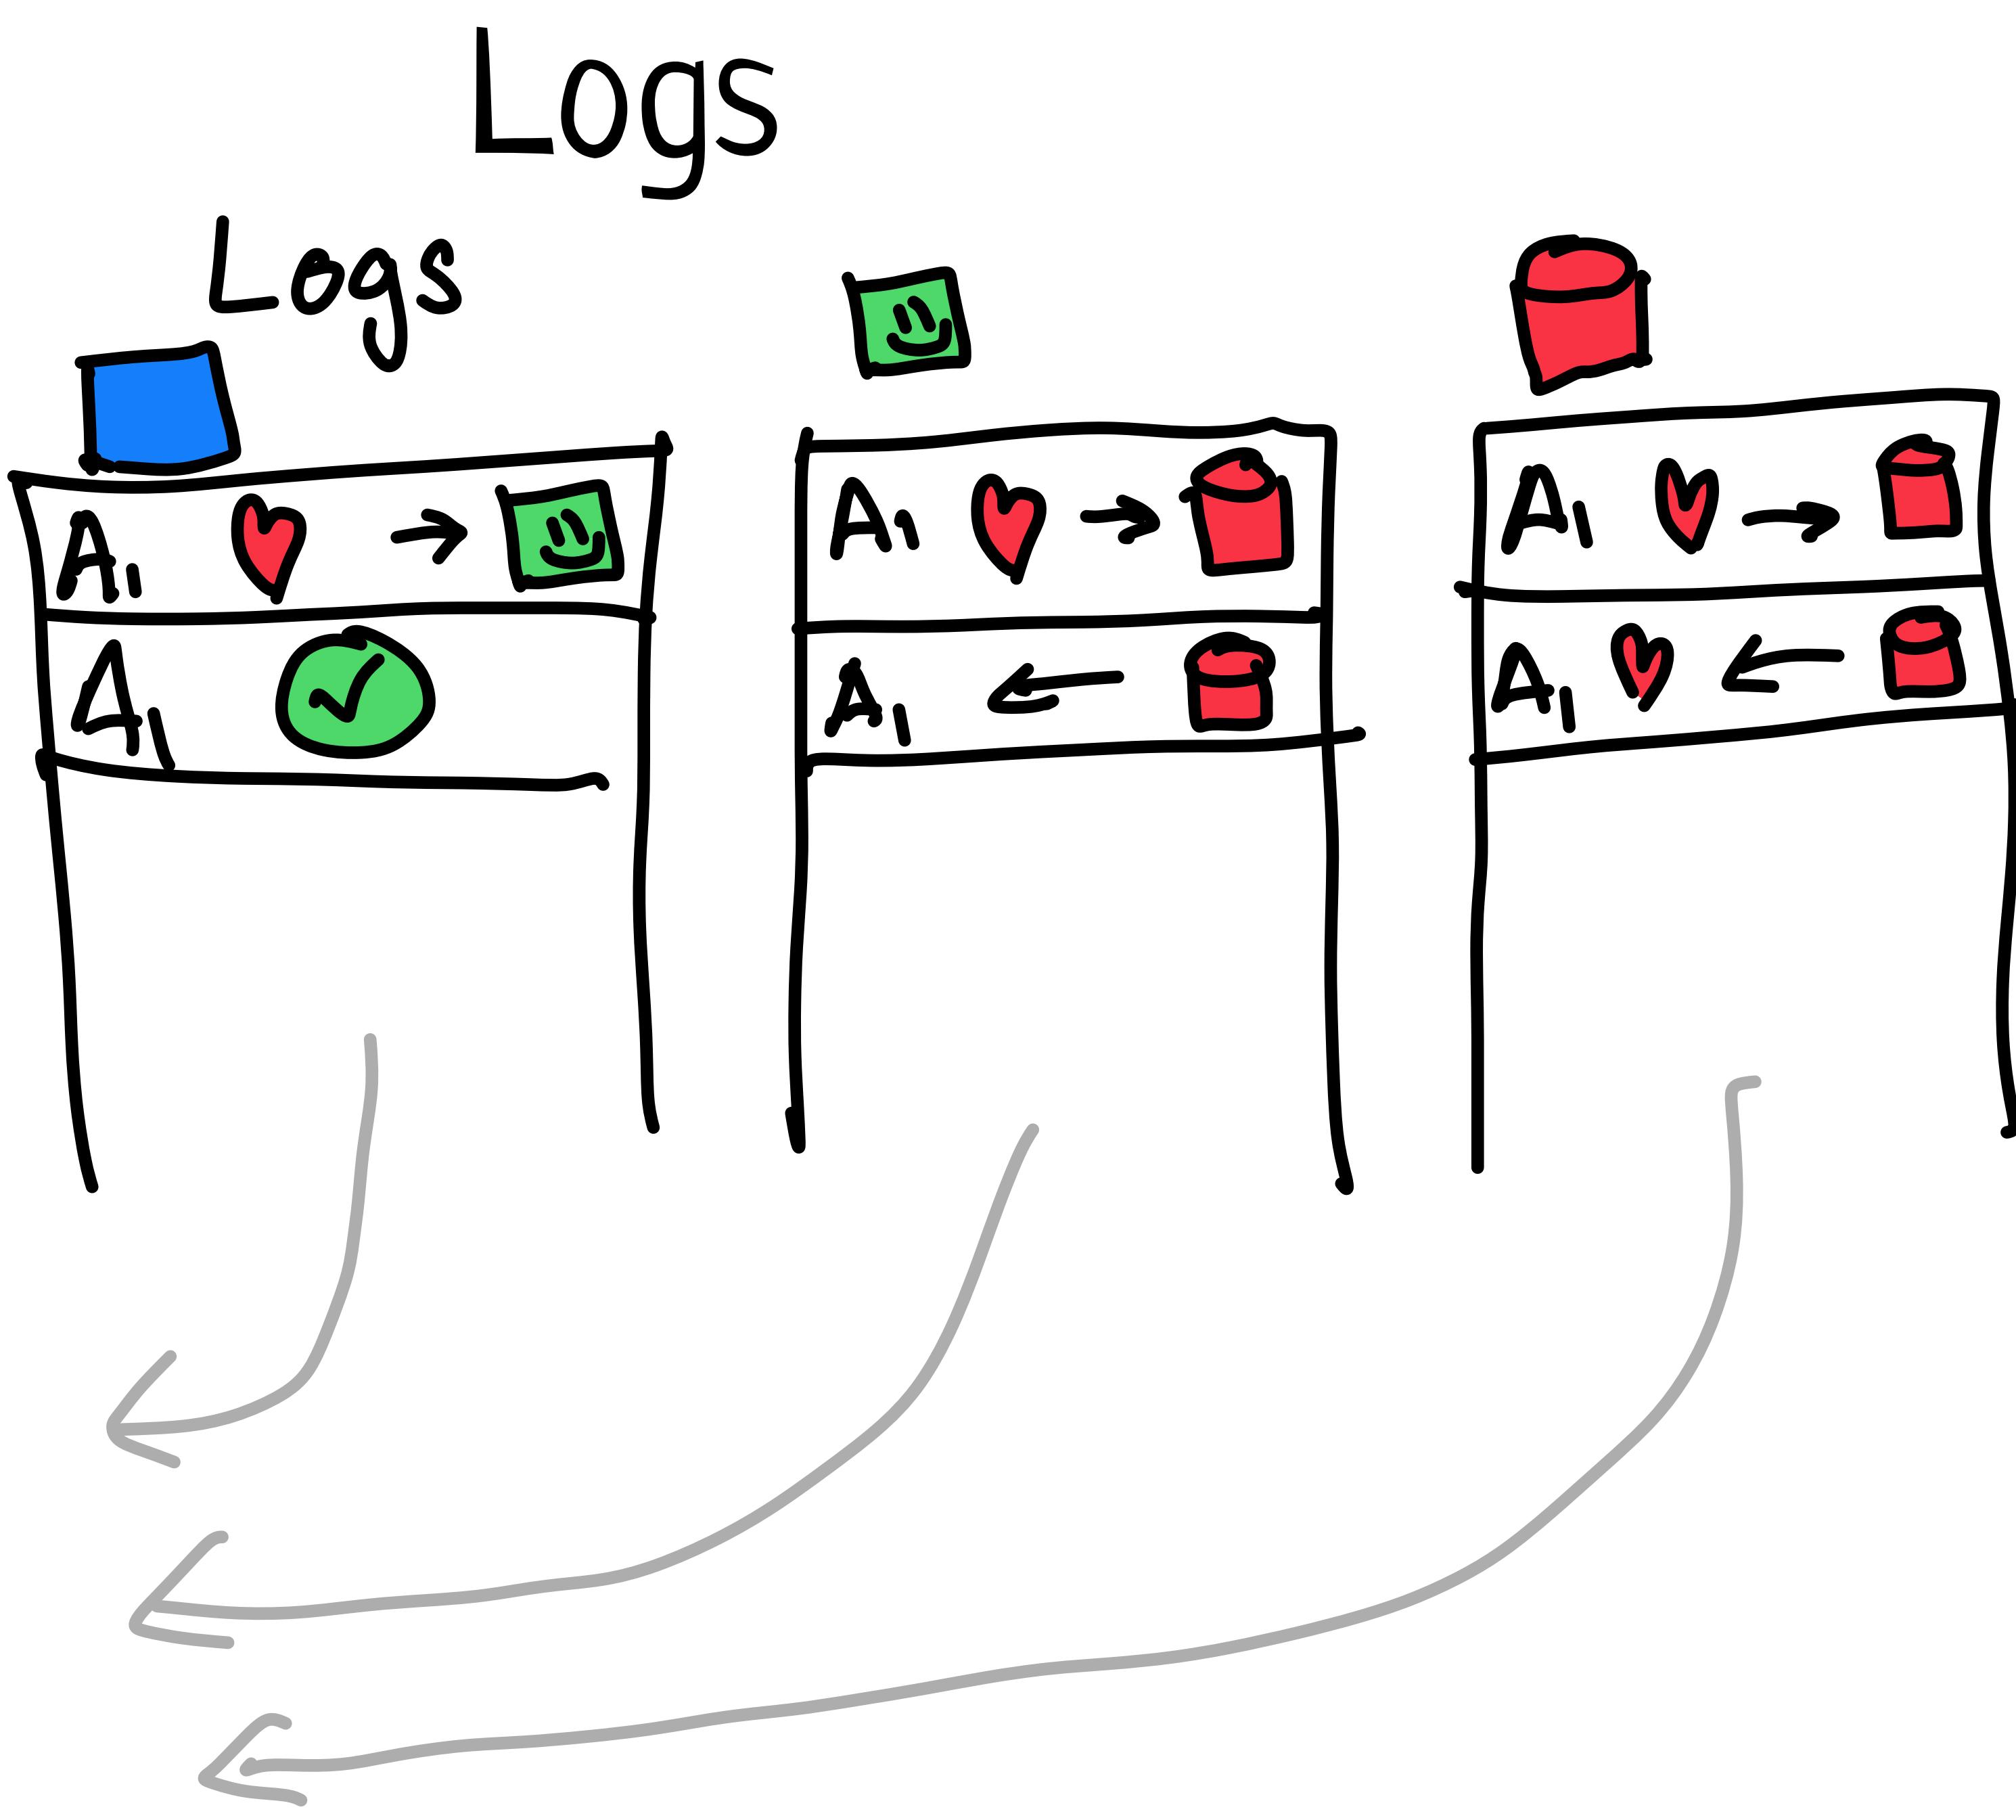
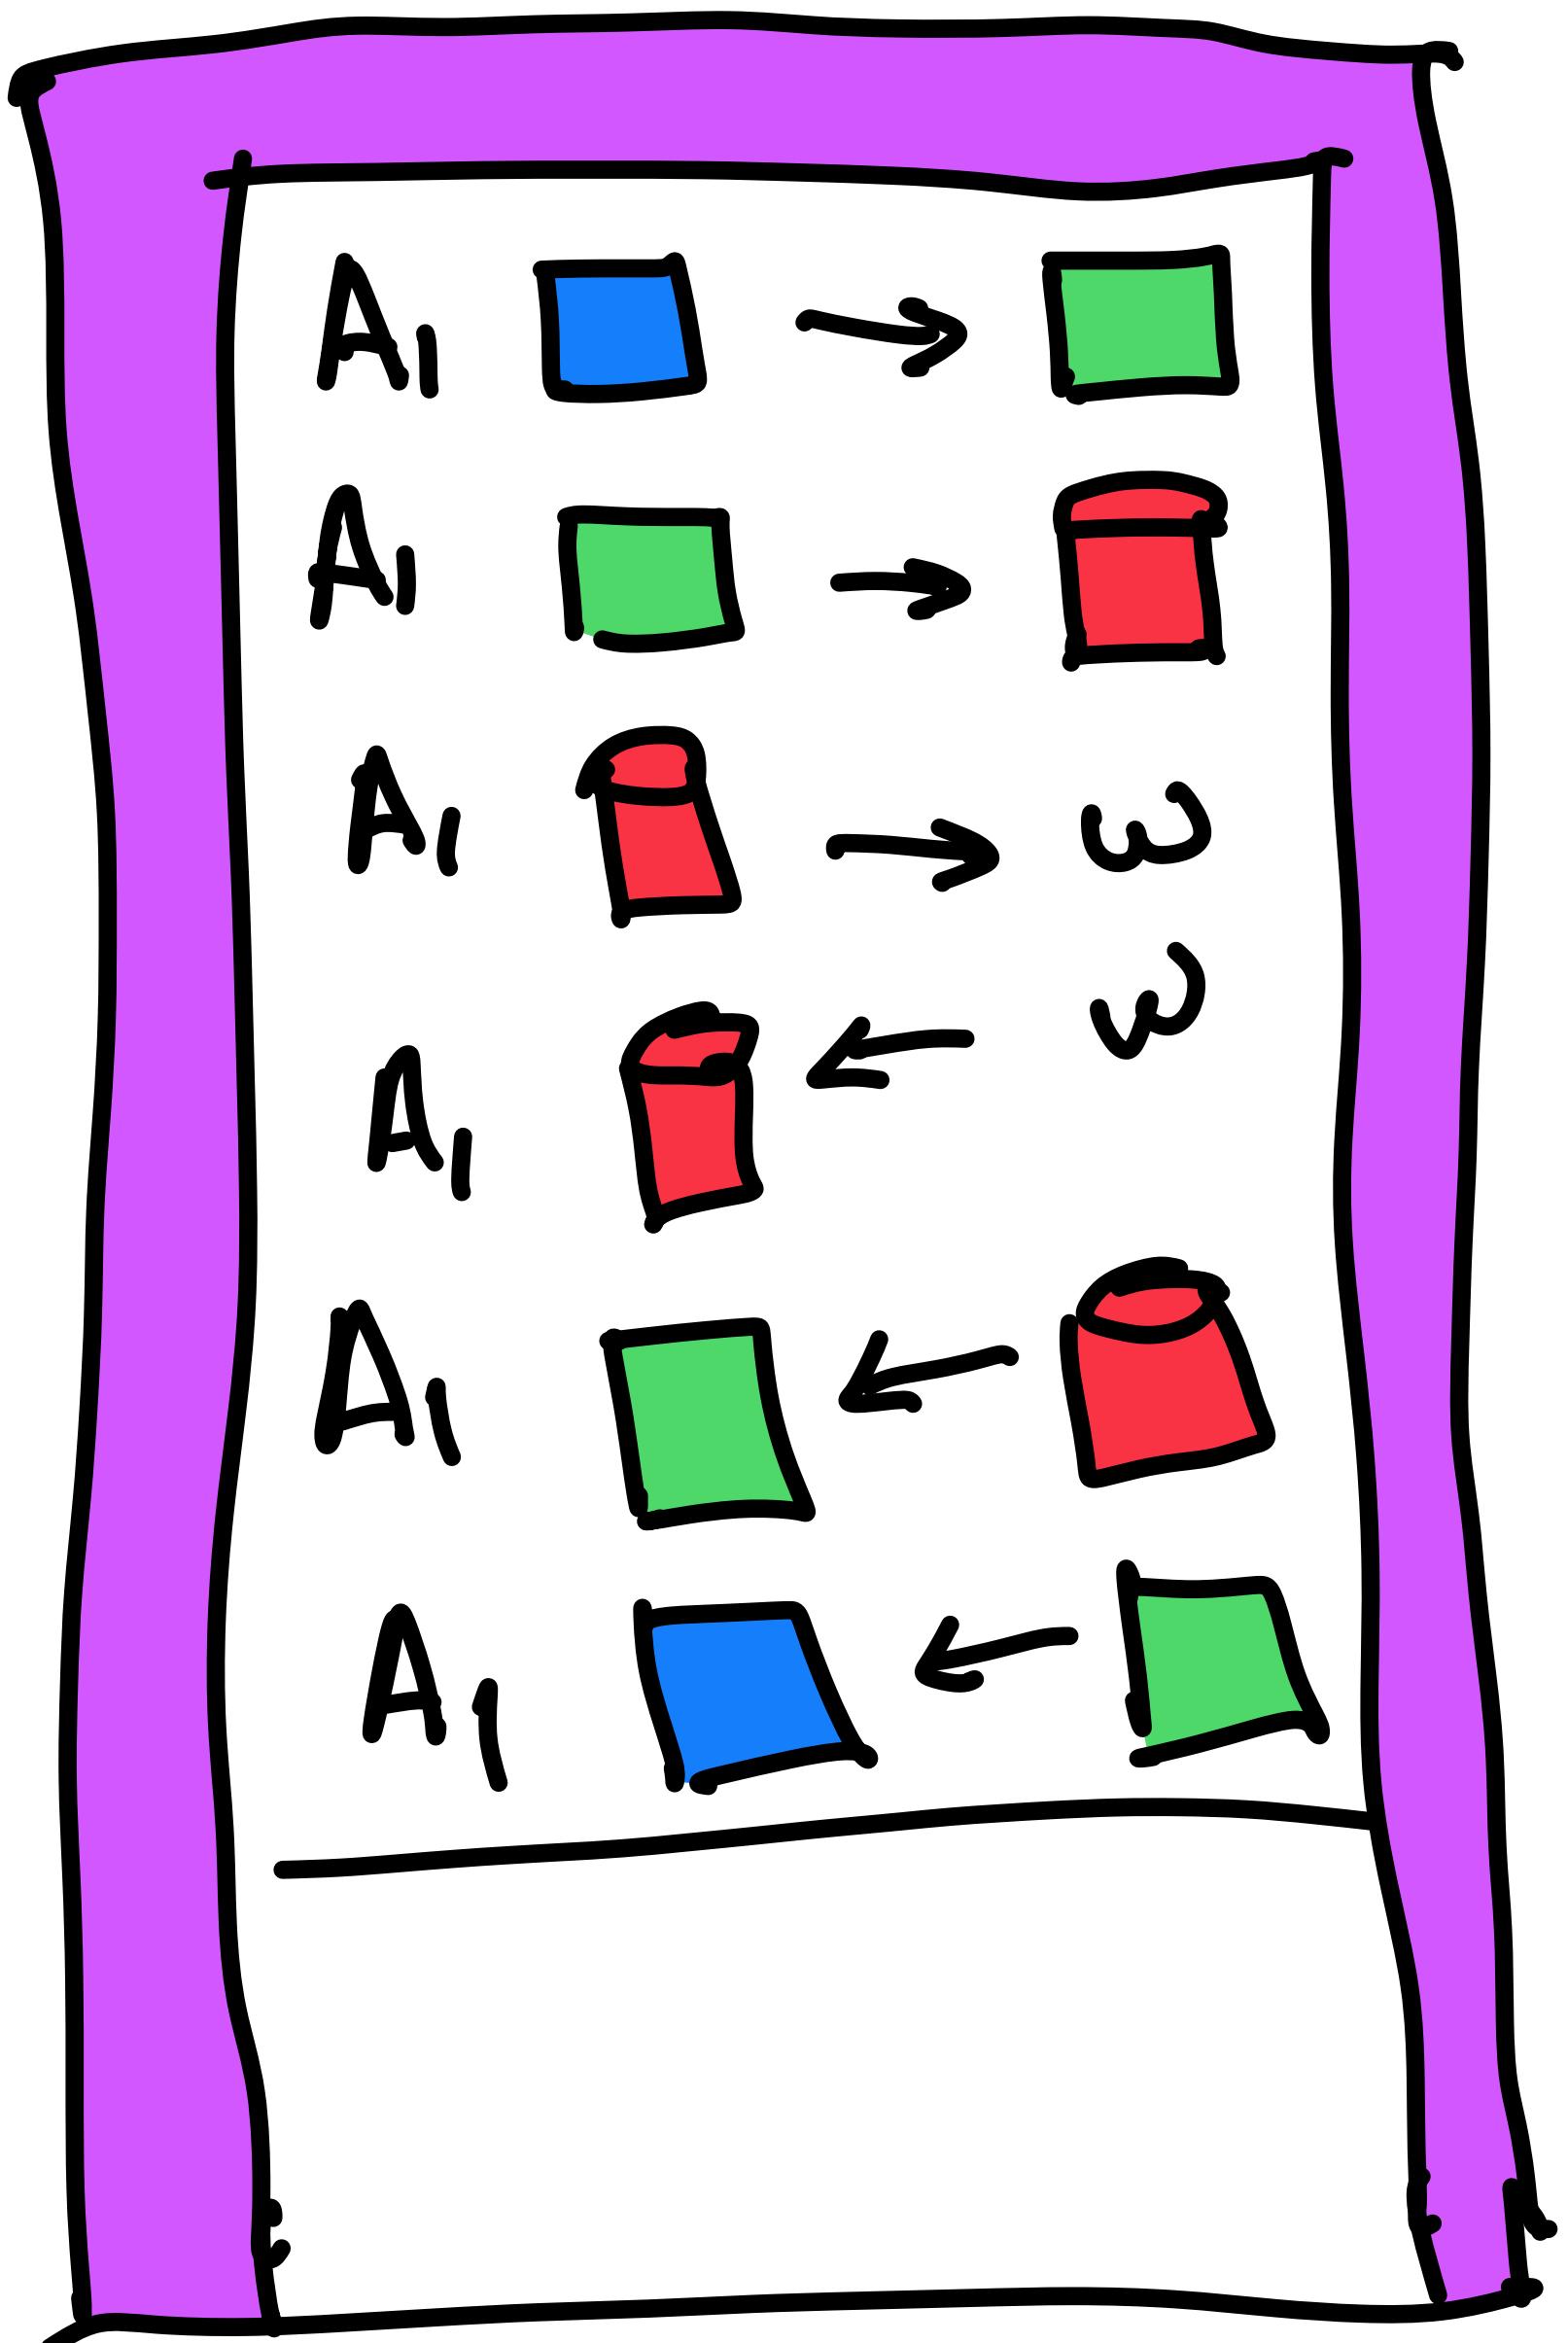








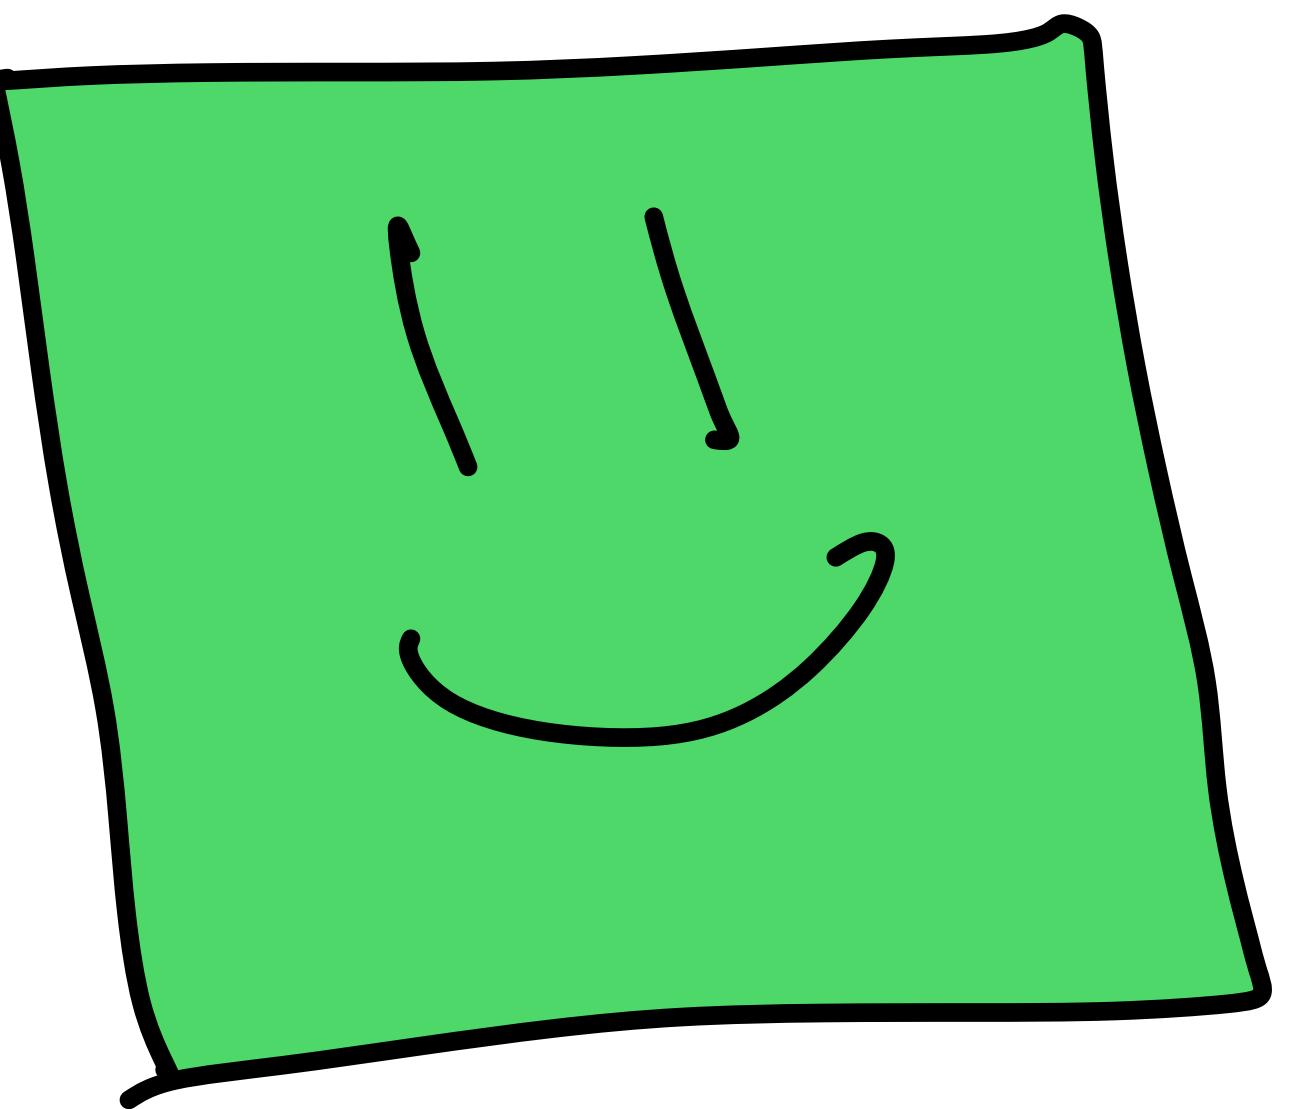
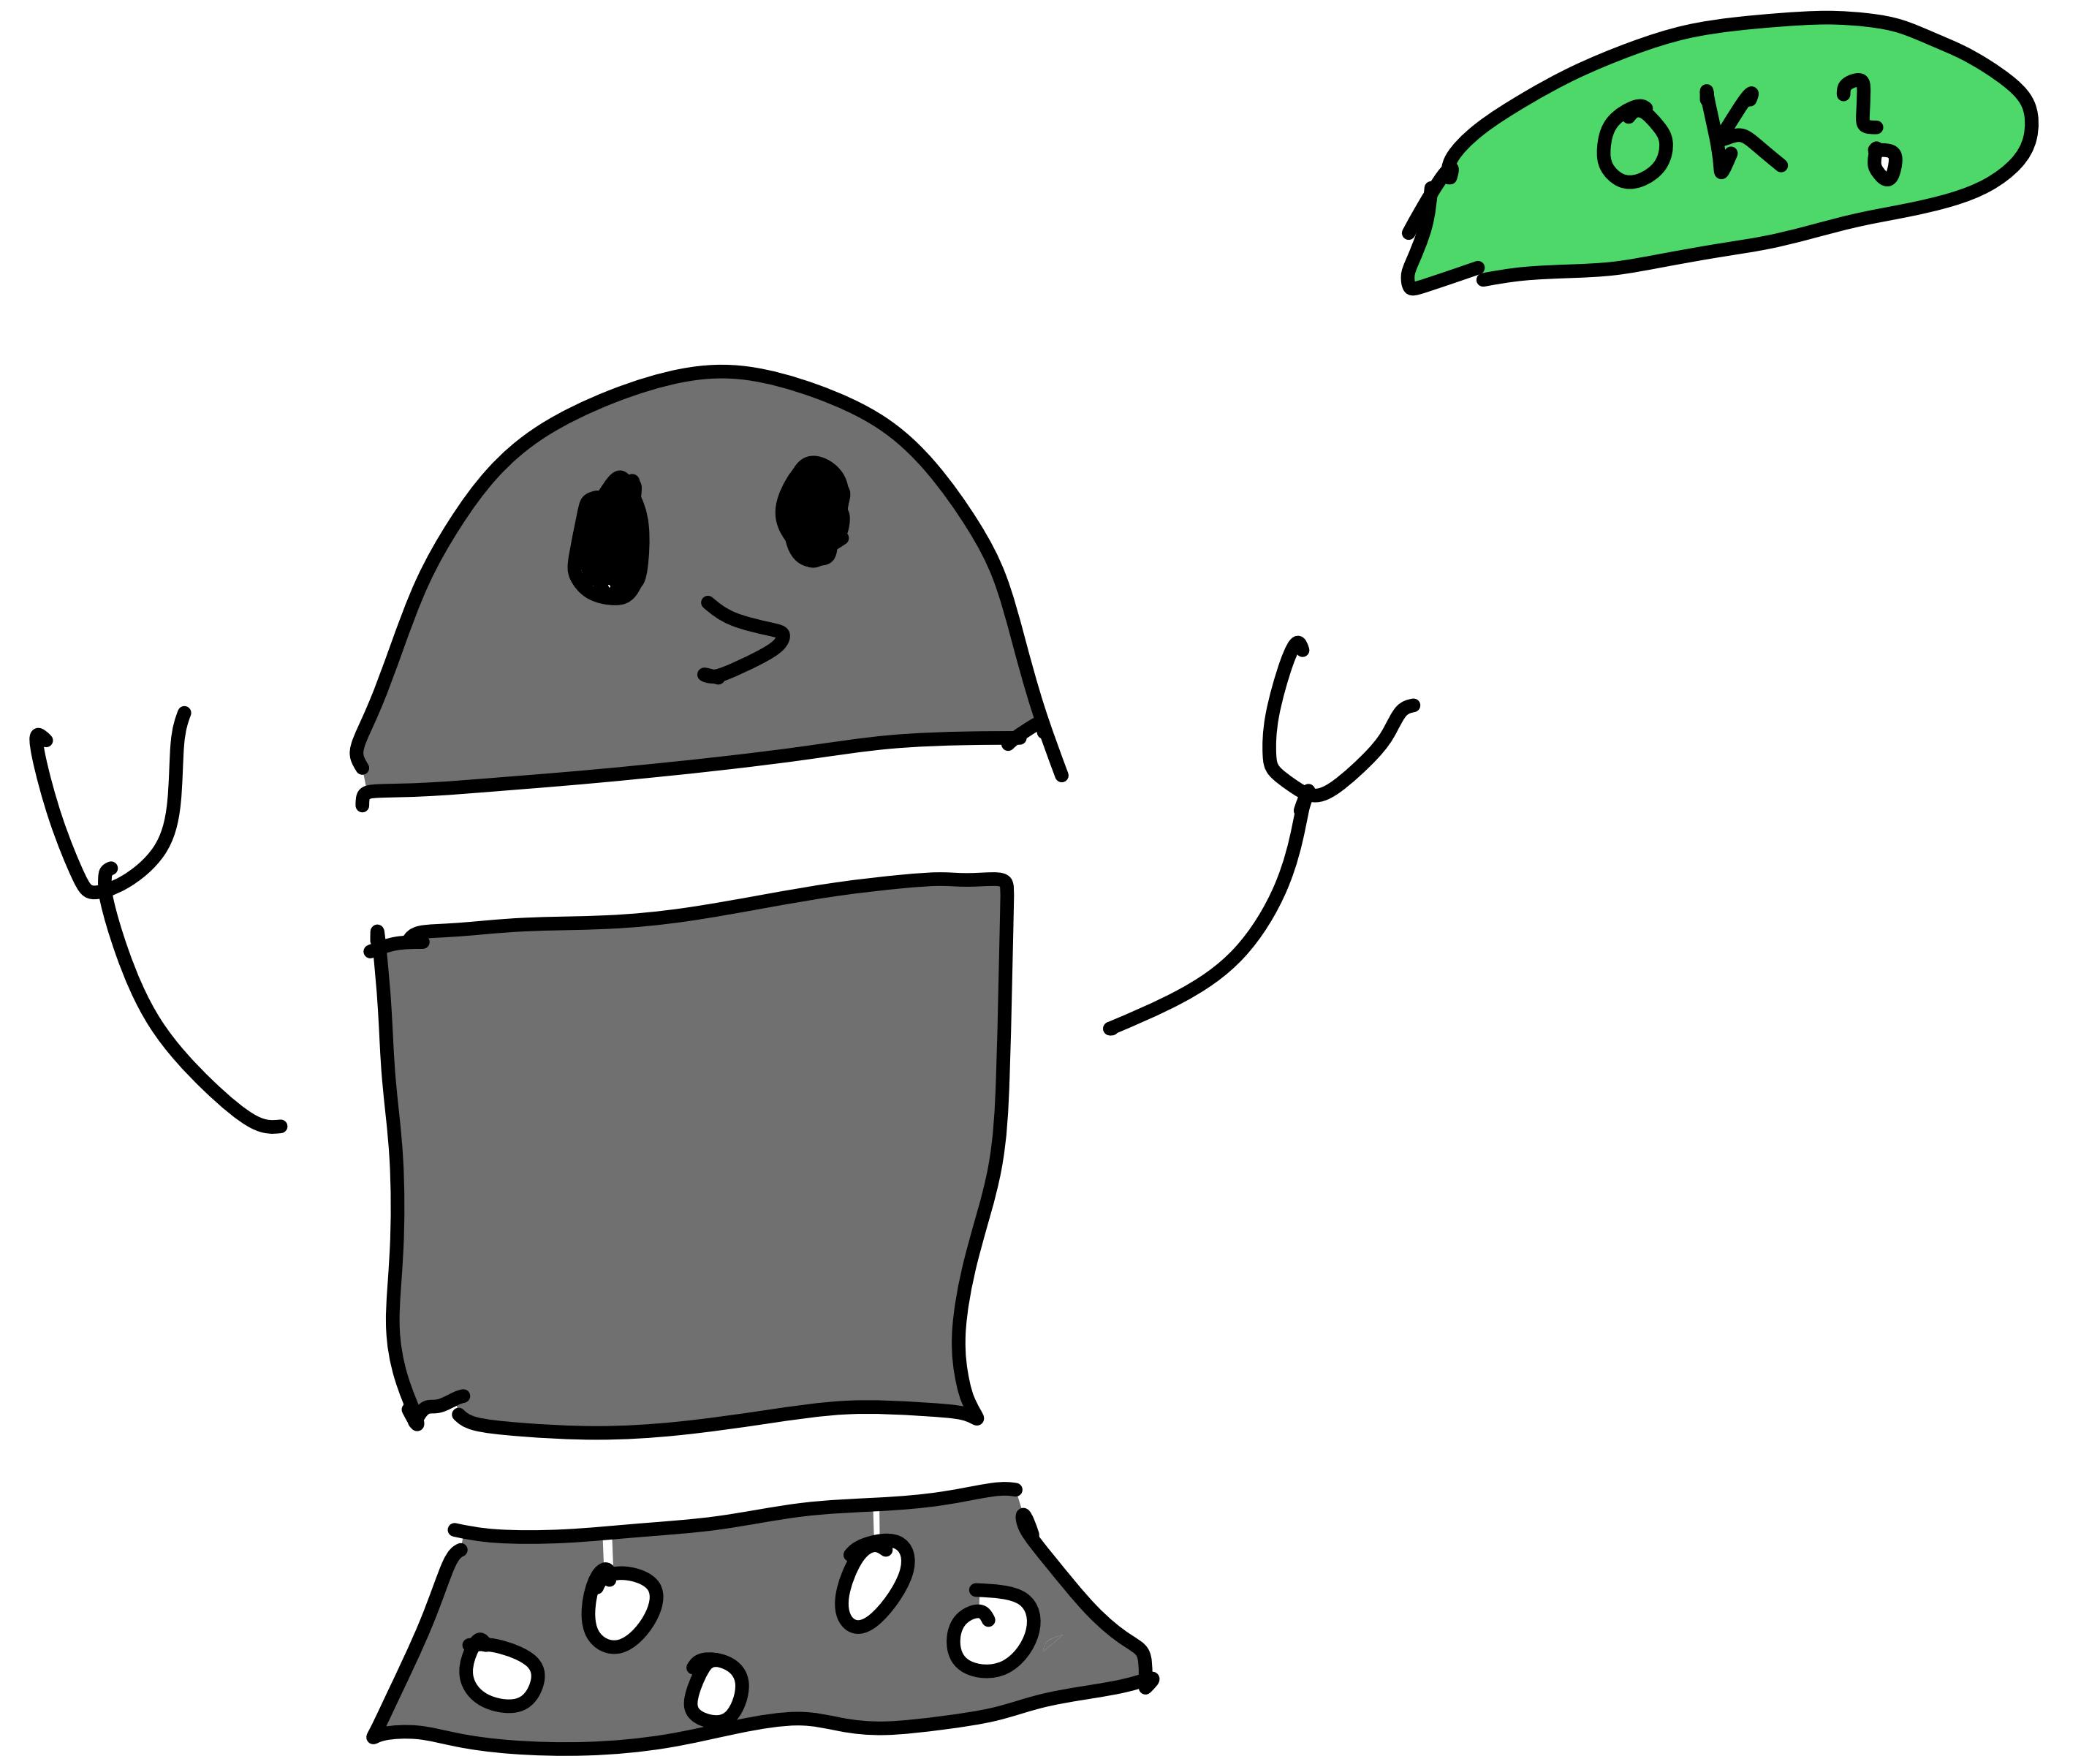
Atomic Uuids example

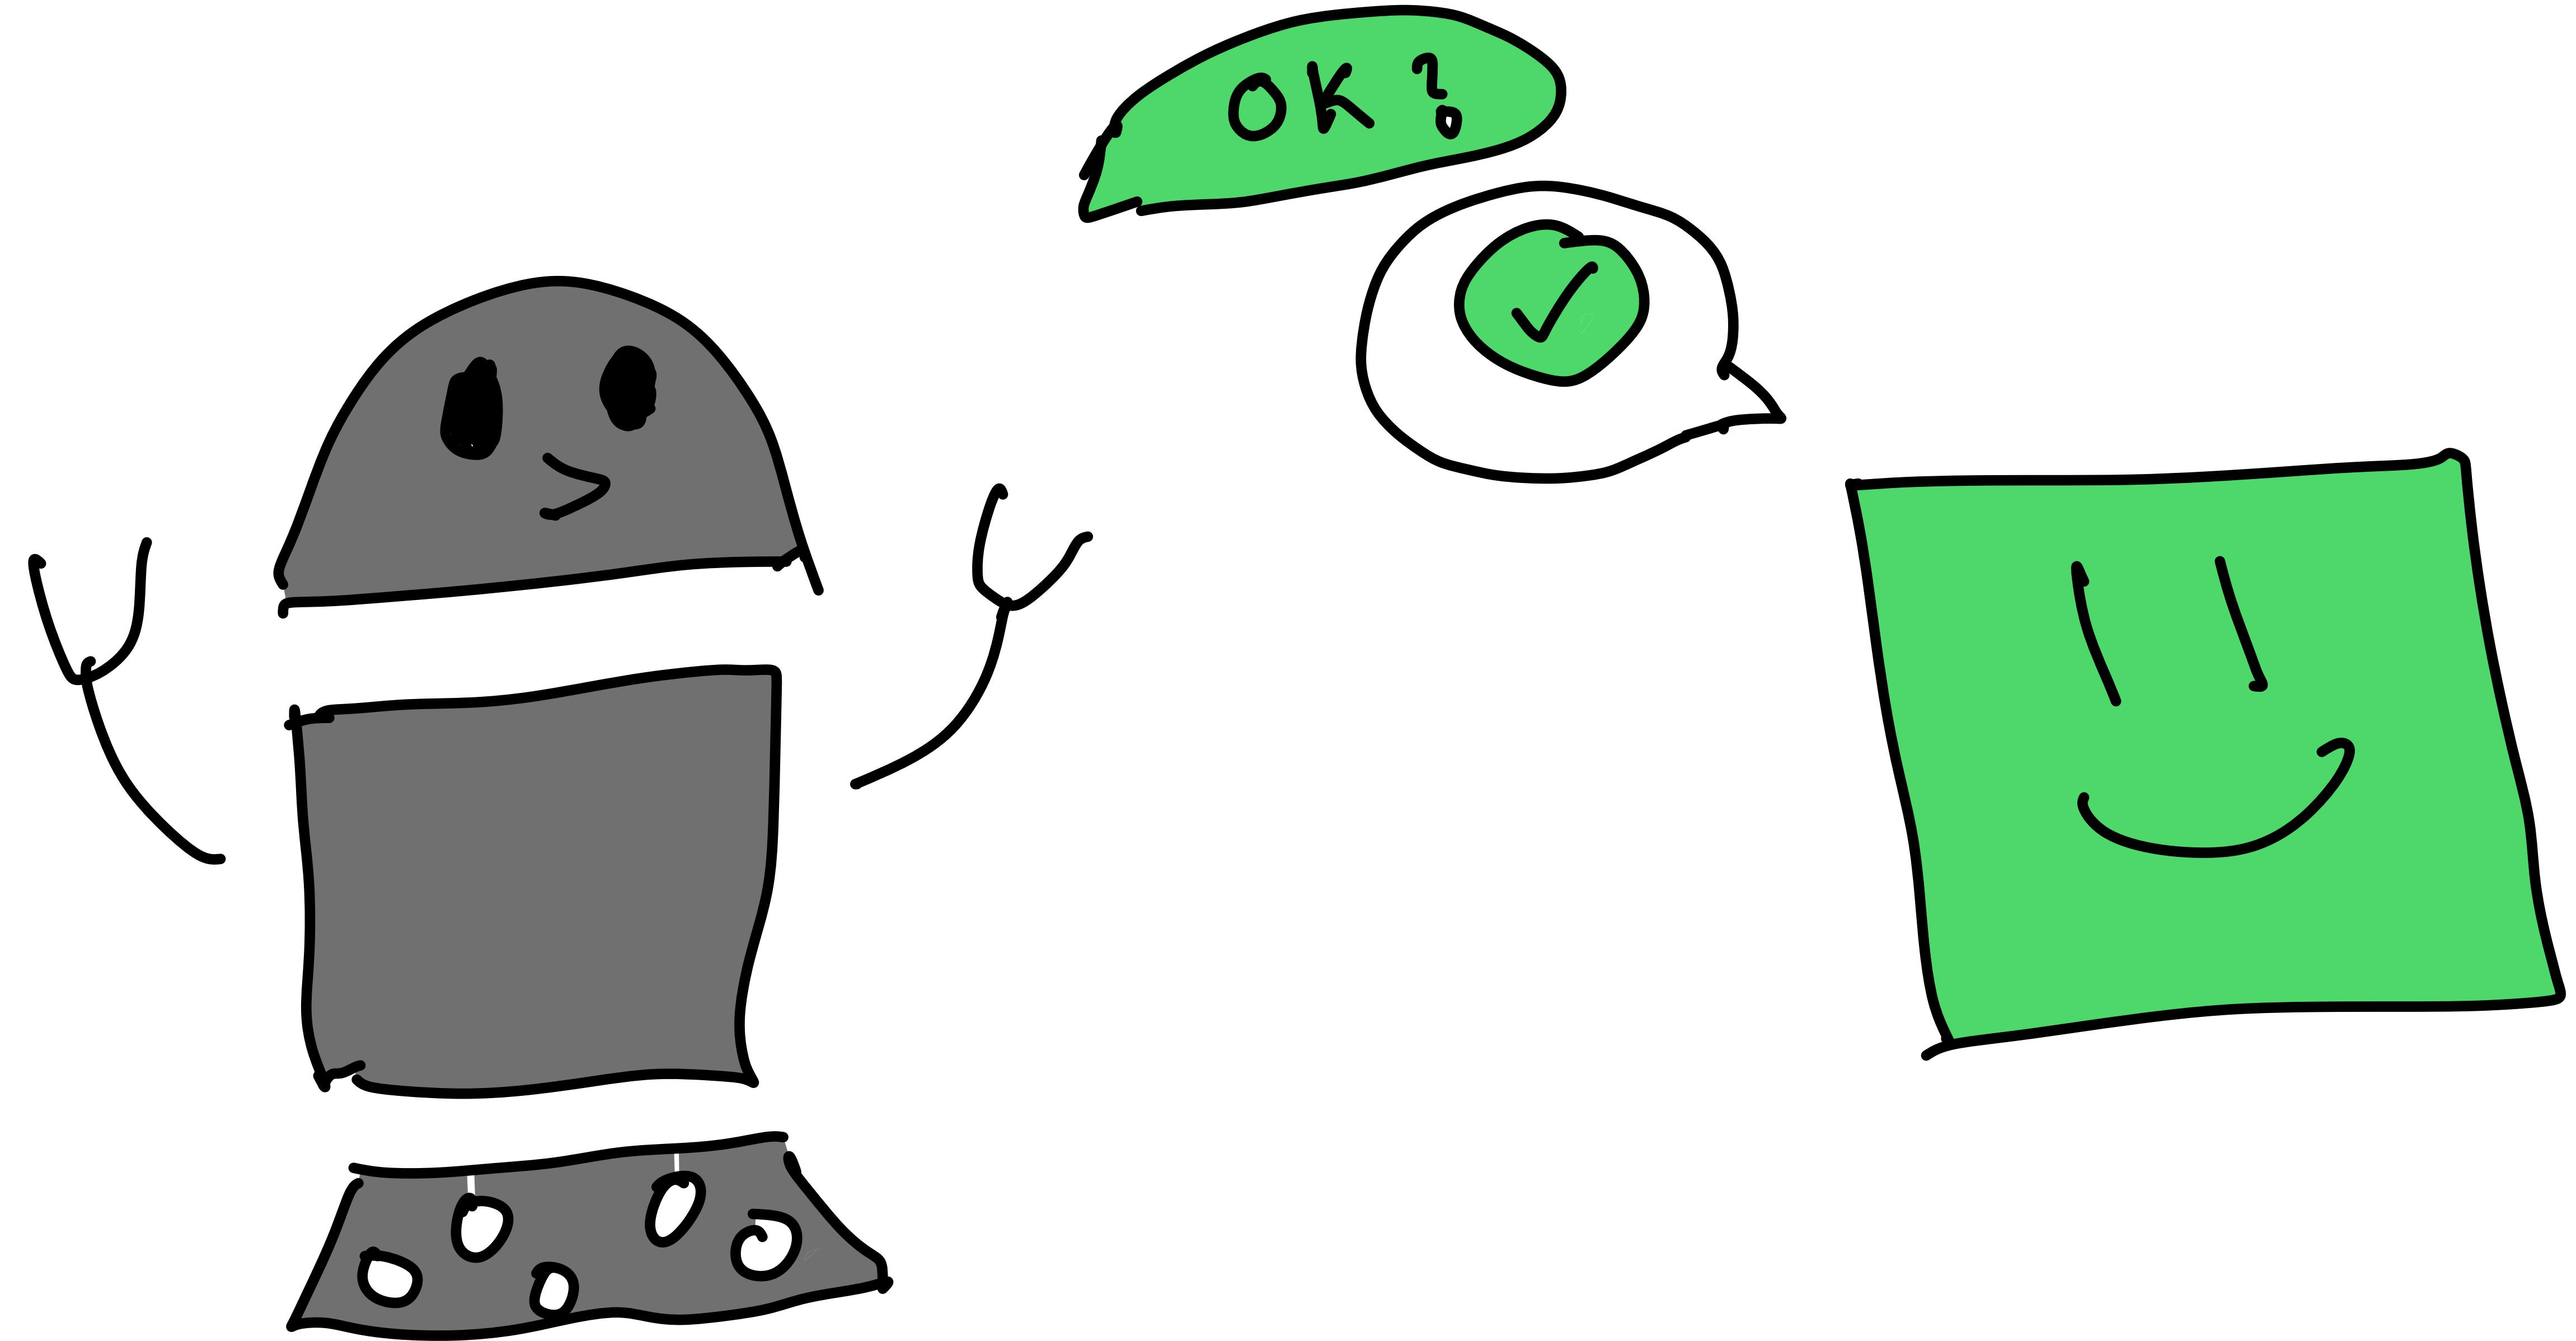


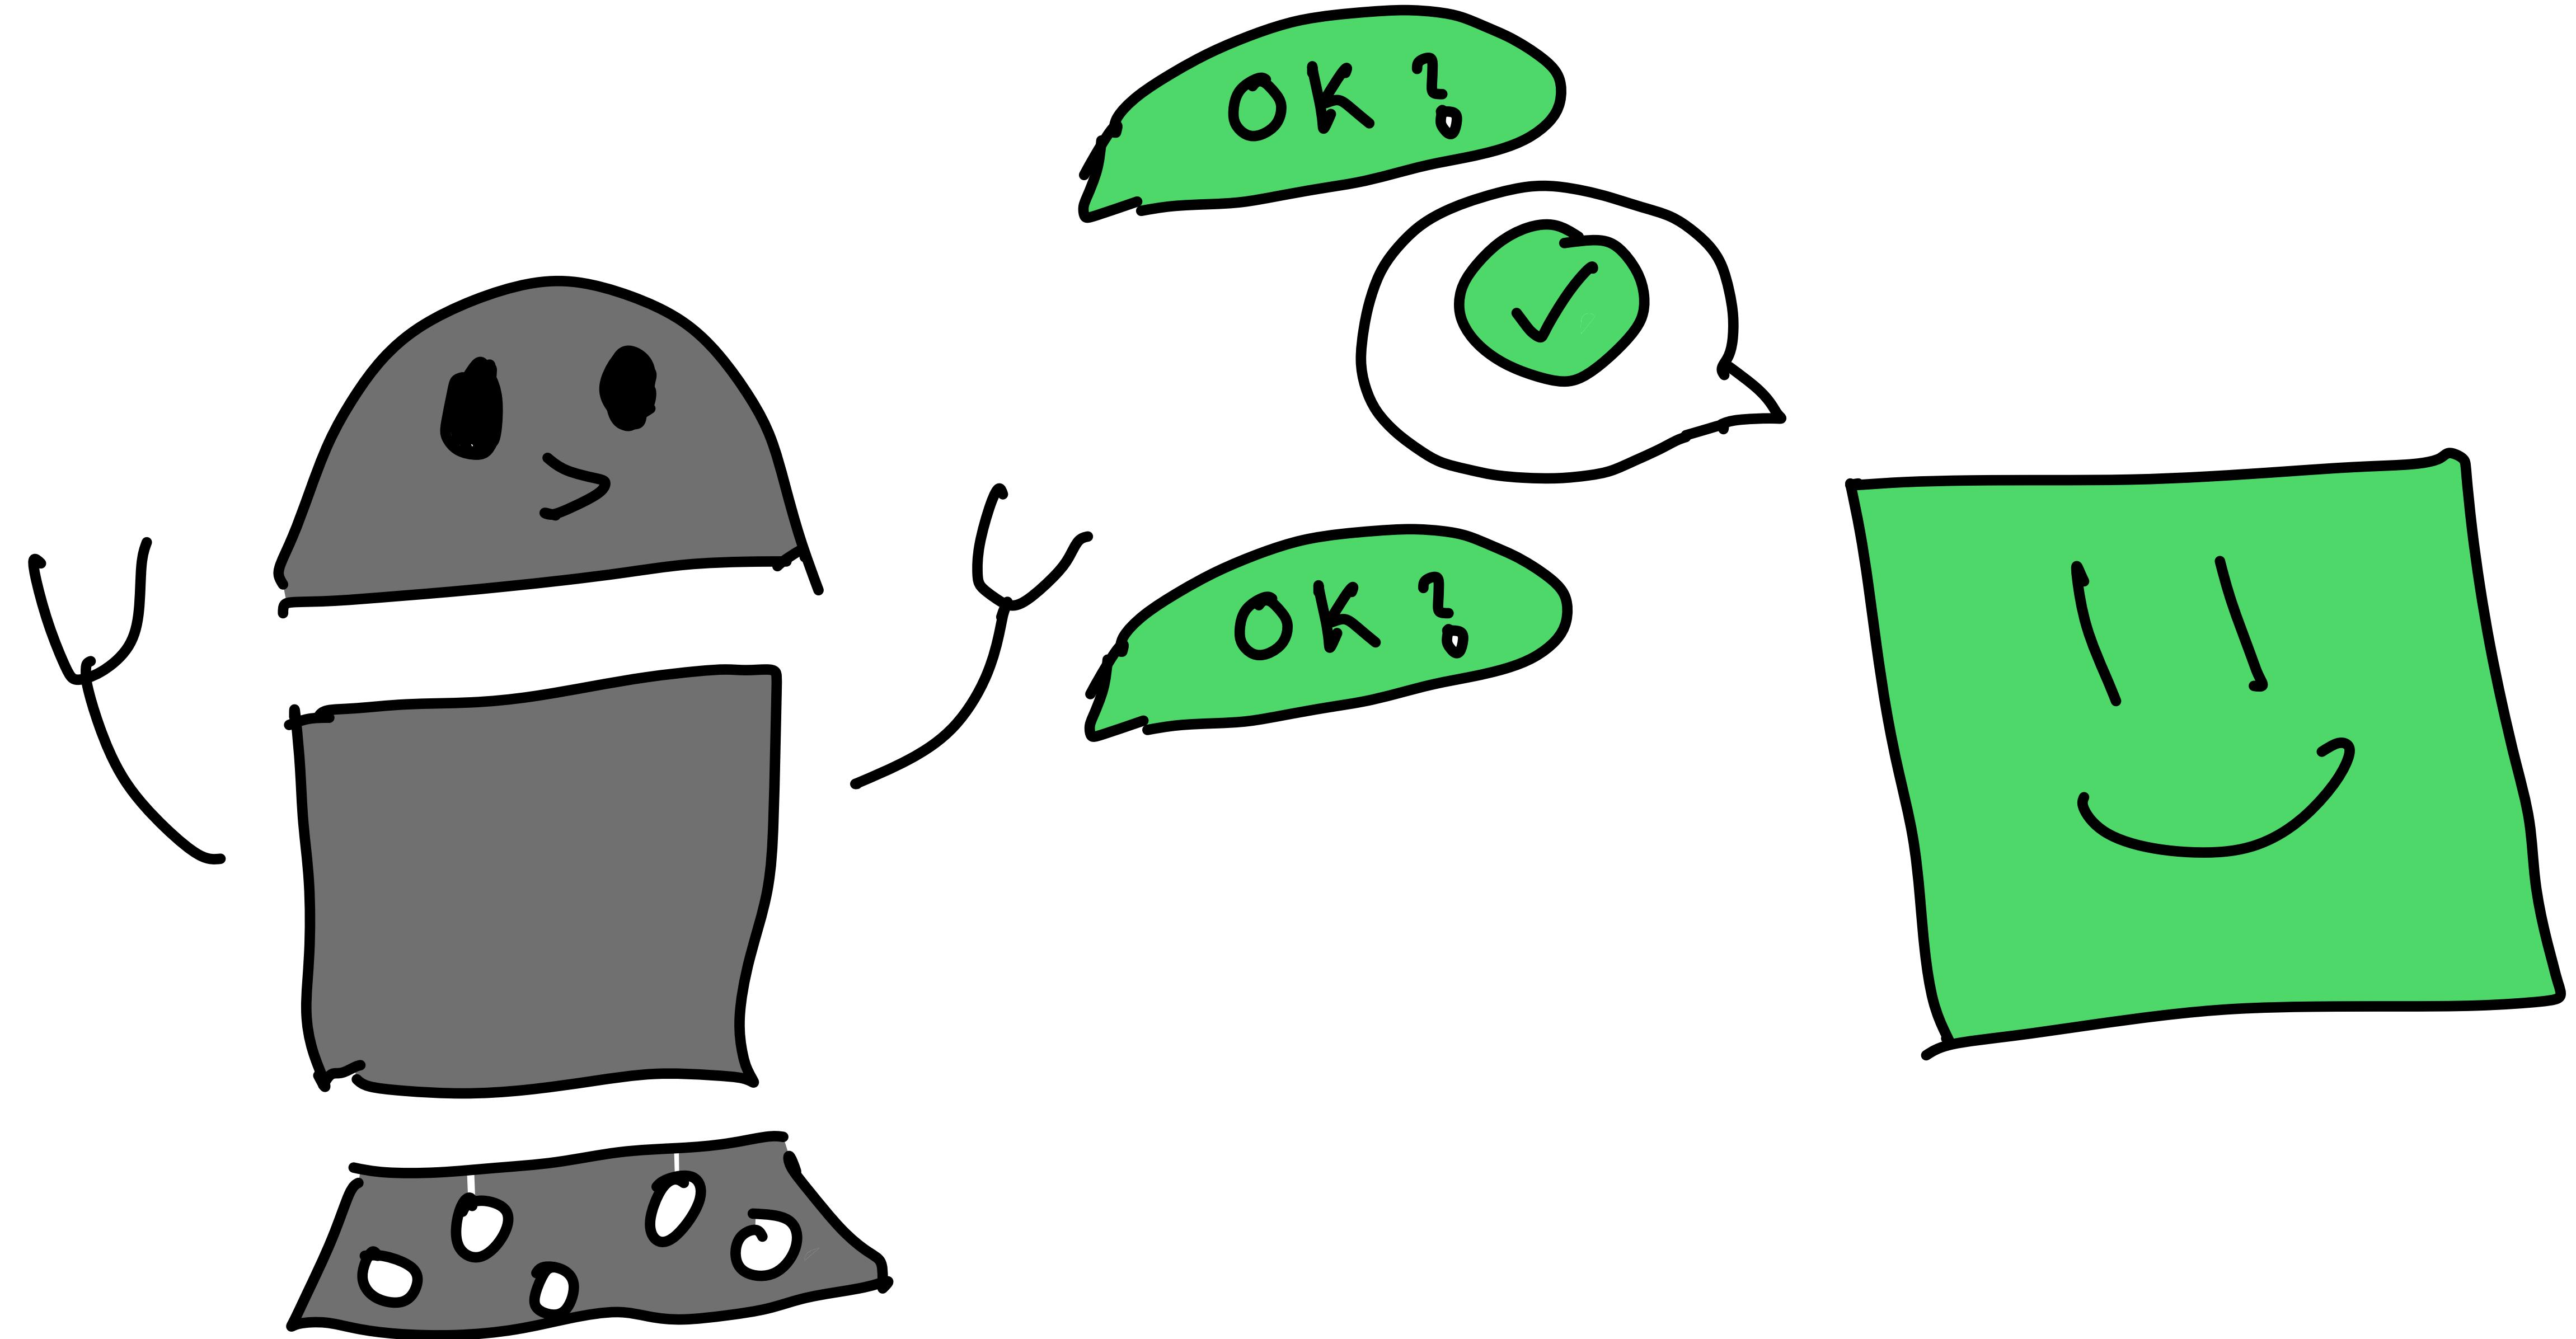
Centralized Logging

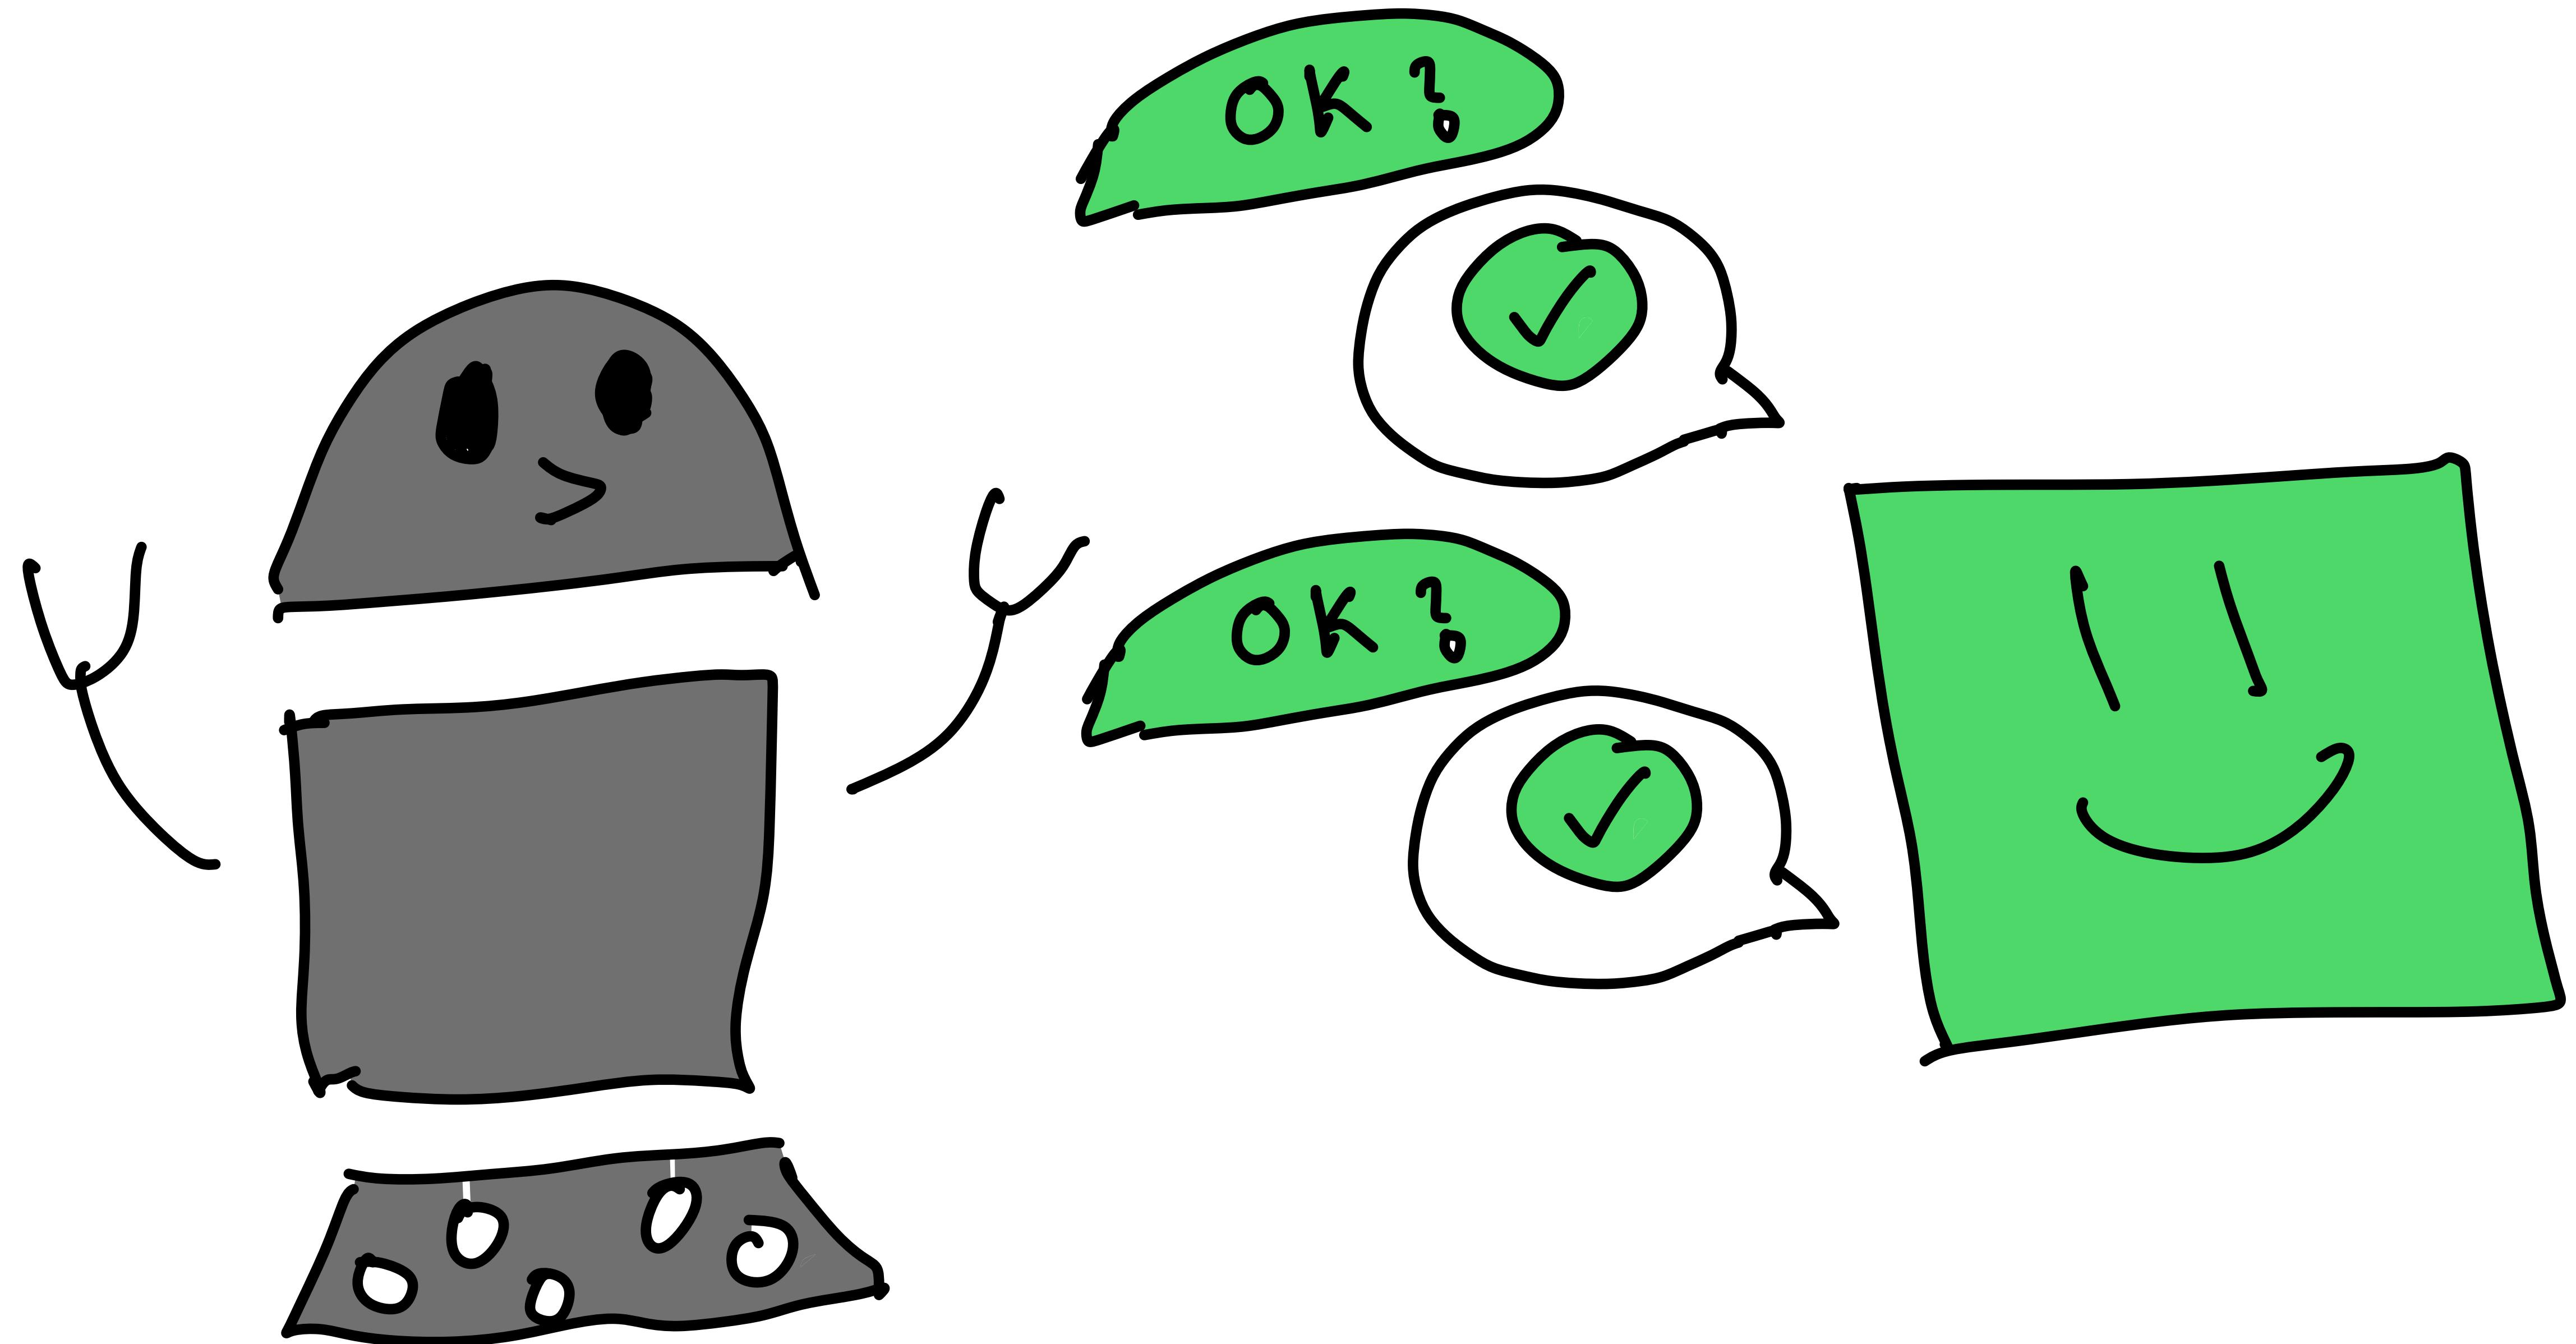
- ELK stack
(Elastic Search, Logstash, Kibana)
The “E” in ELK is an
Sponsor in this conference,
Go talk to elastic!
- Grafana
- Splunk
- Prometheus

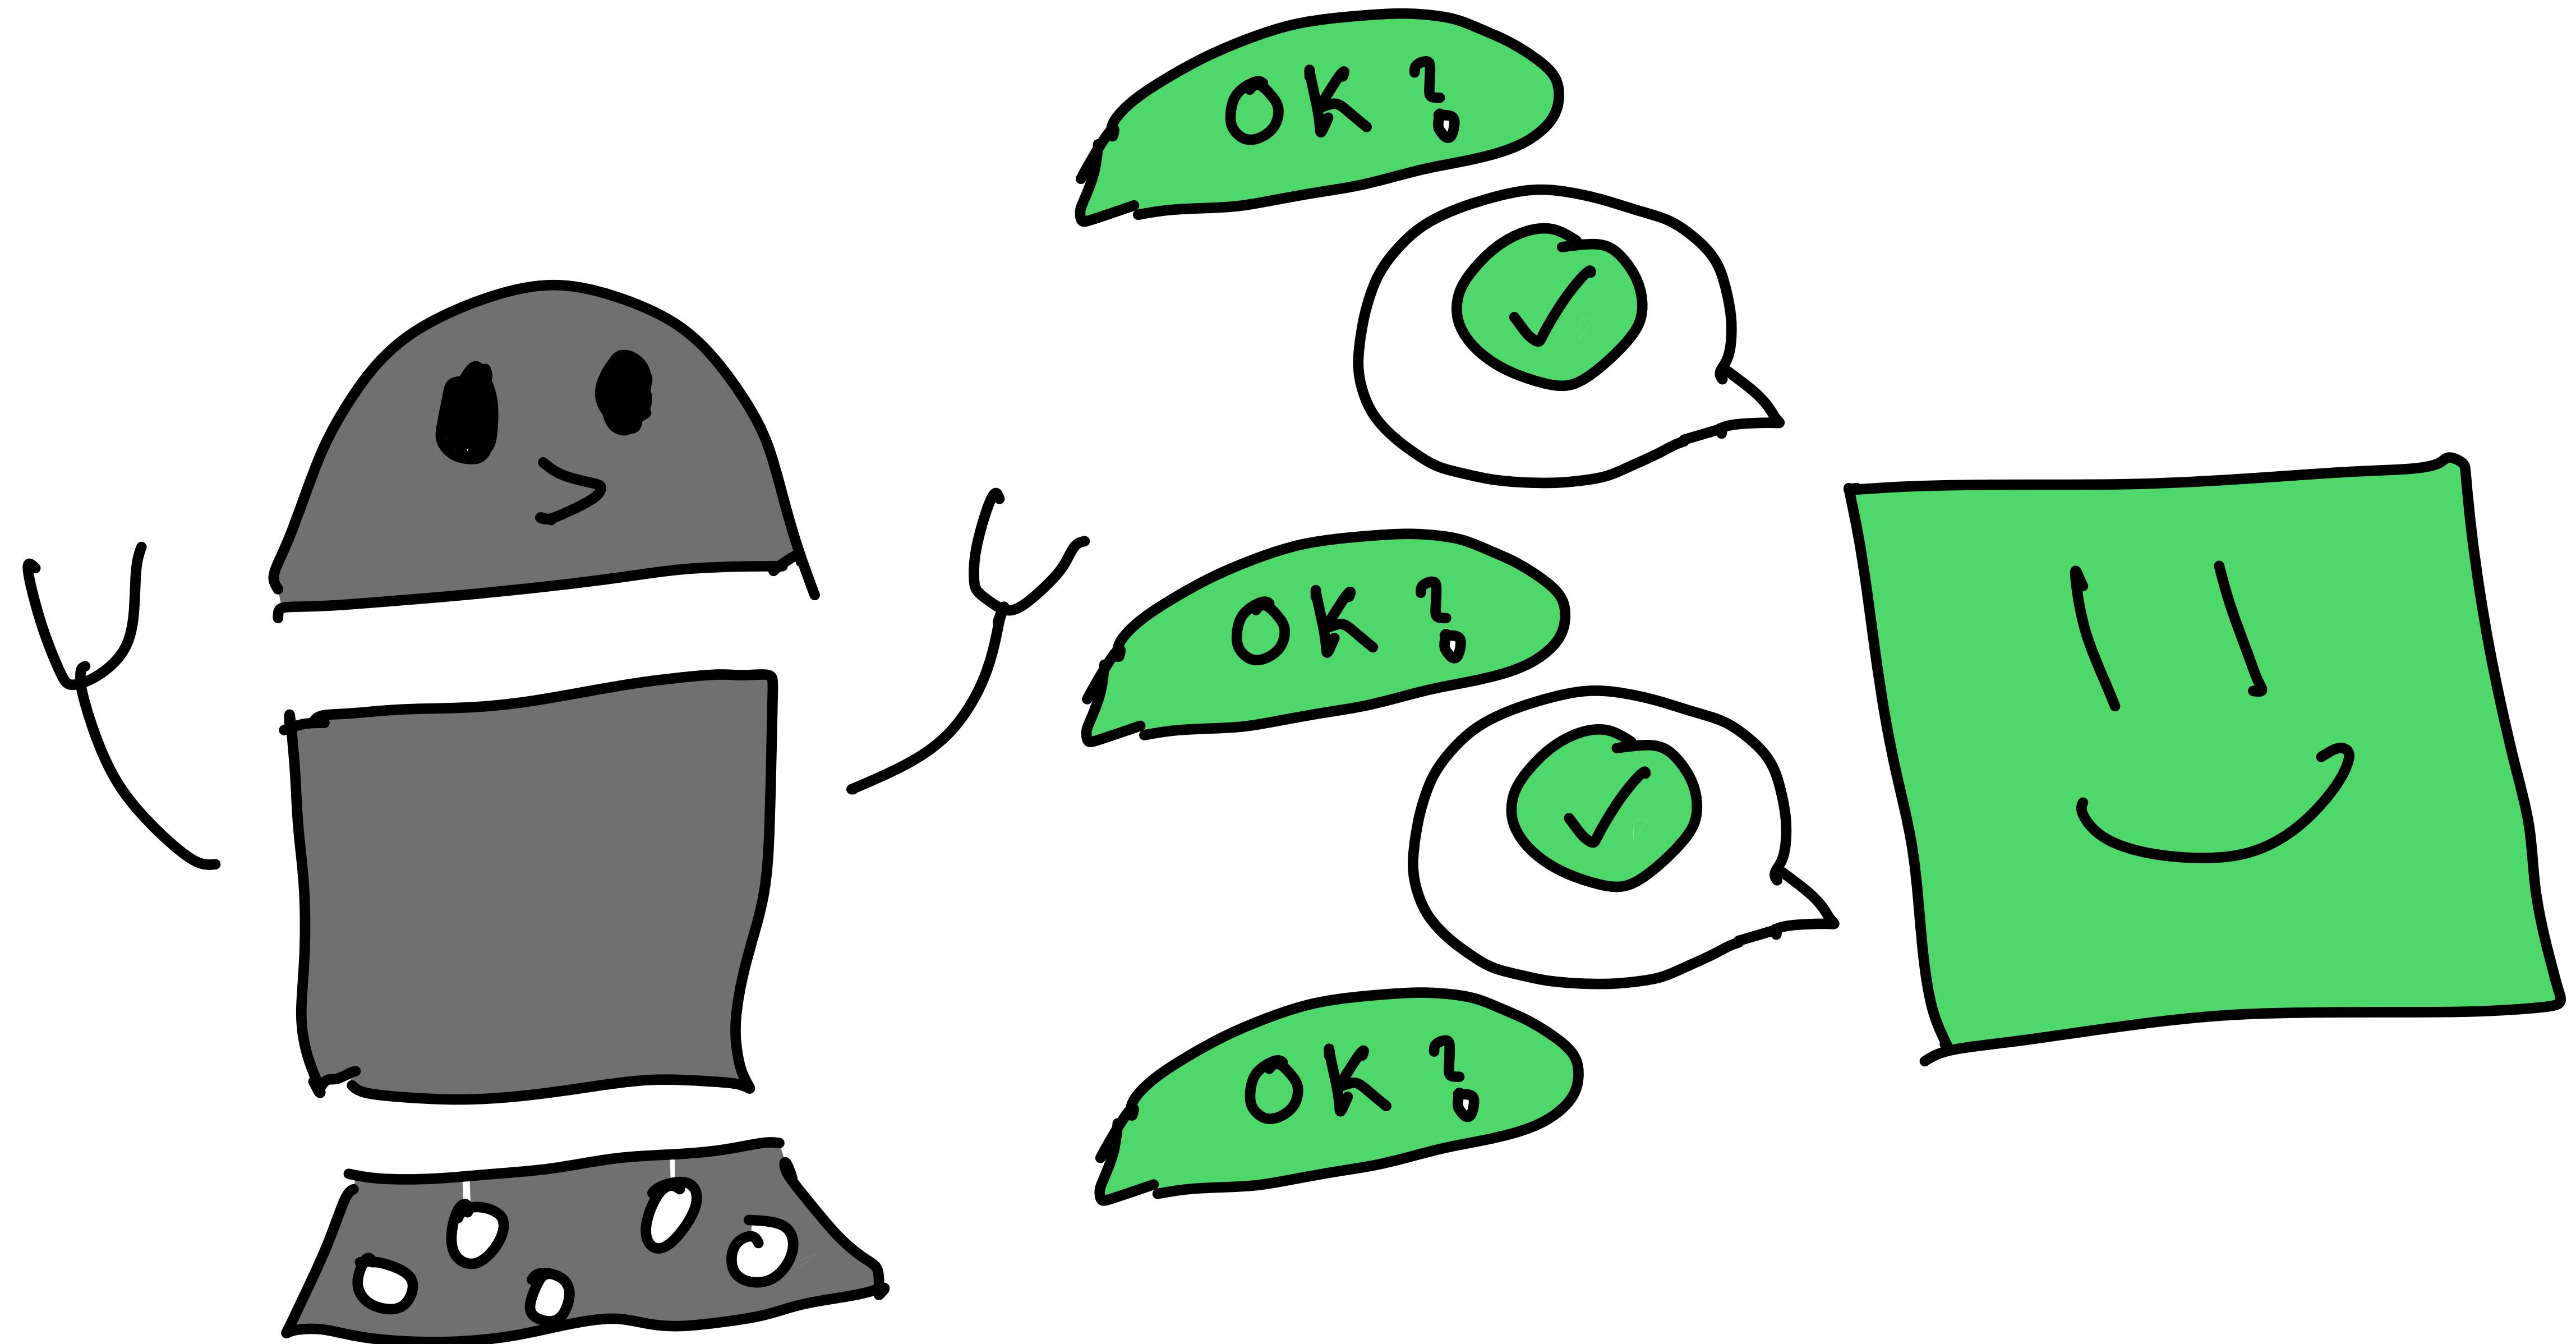
Health Checks / Automation

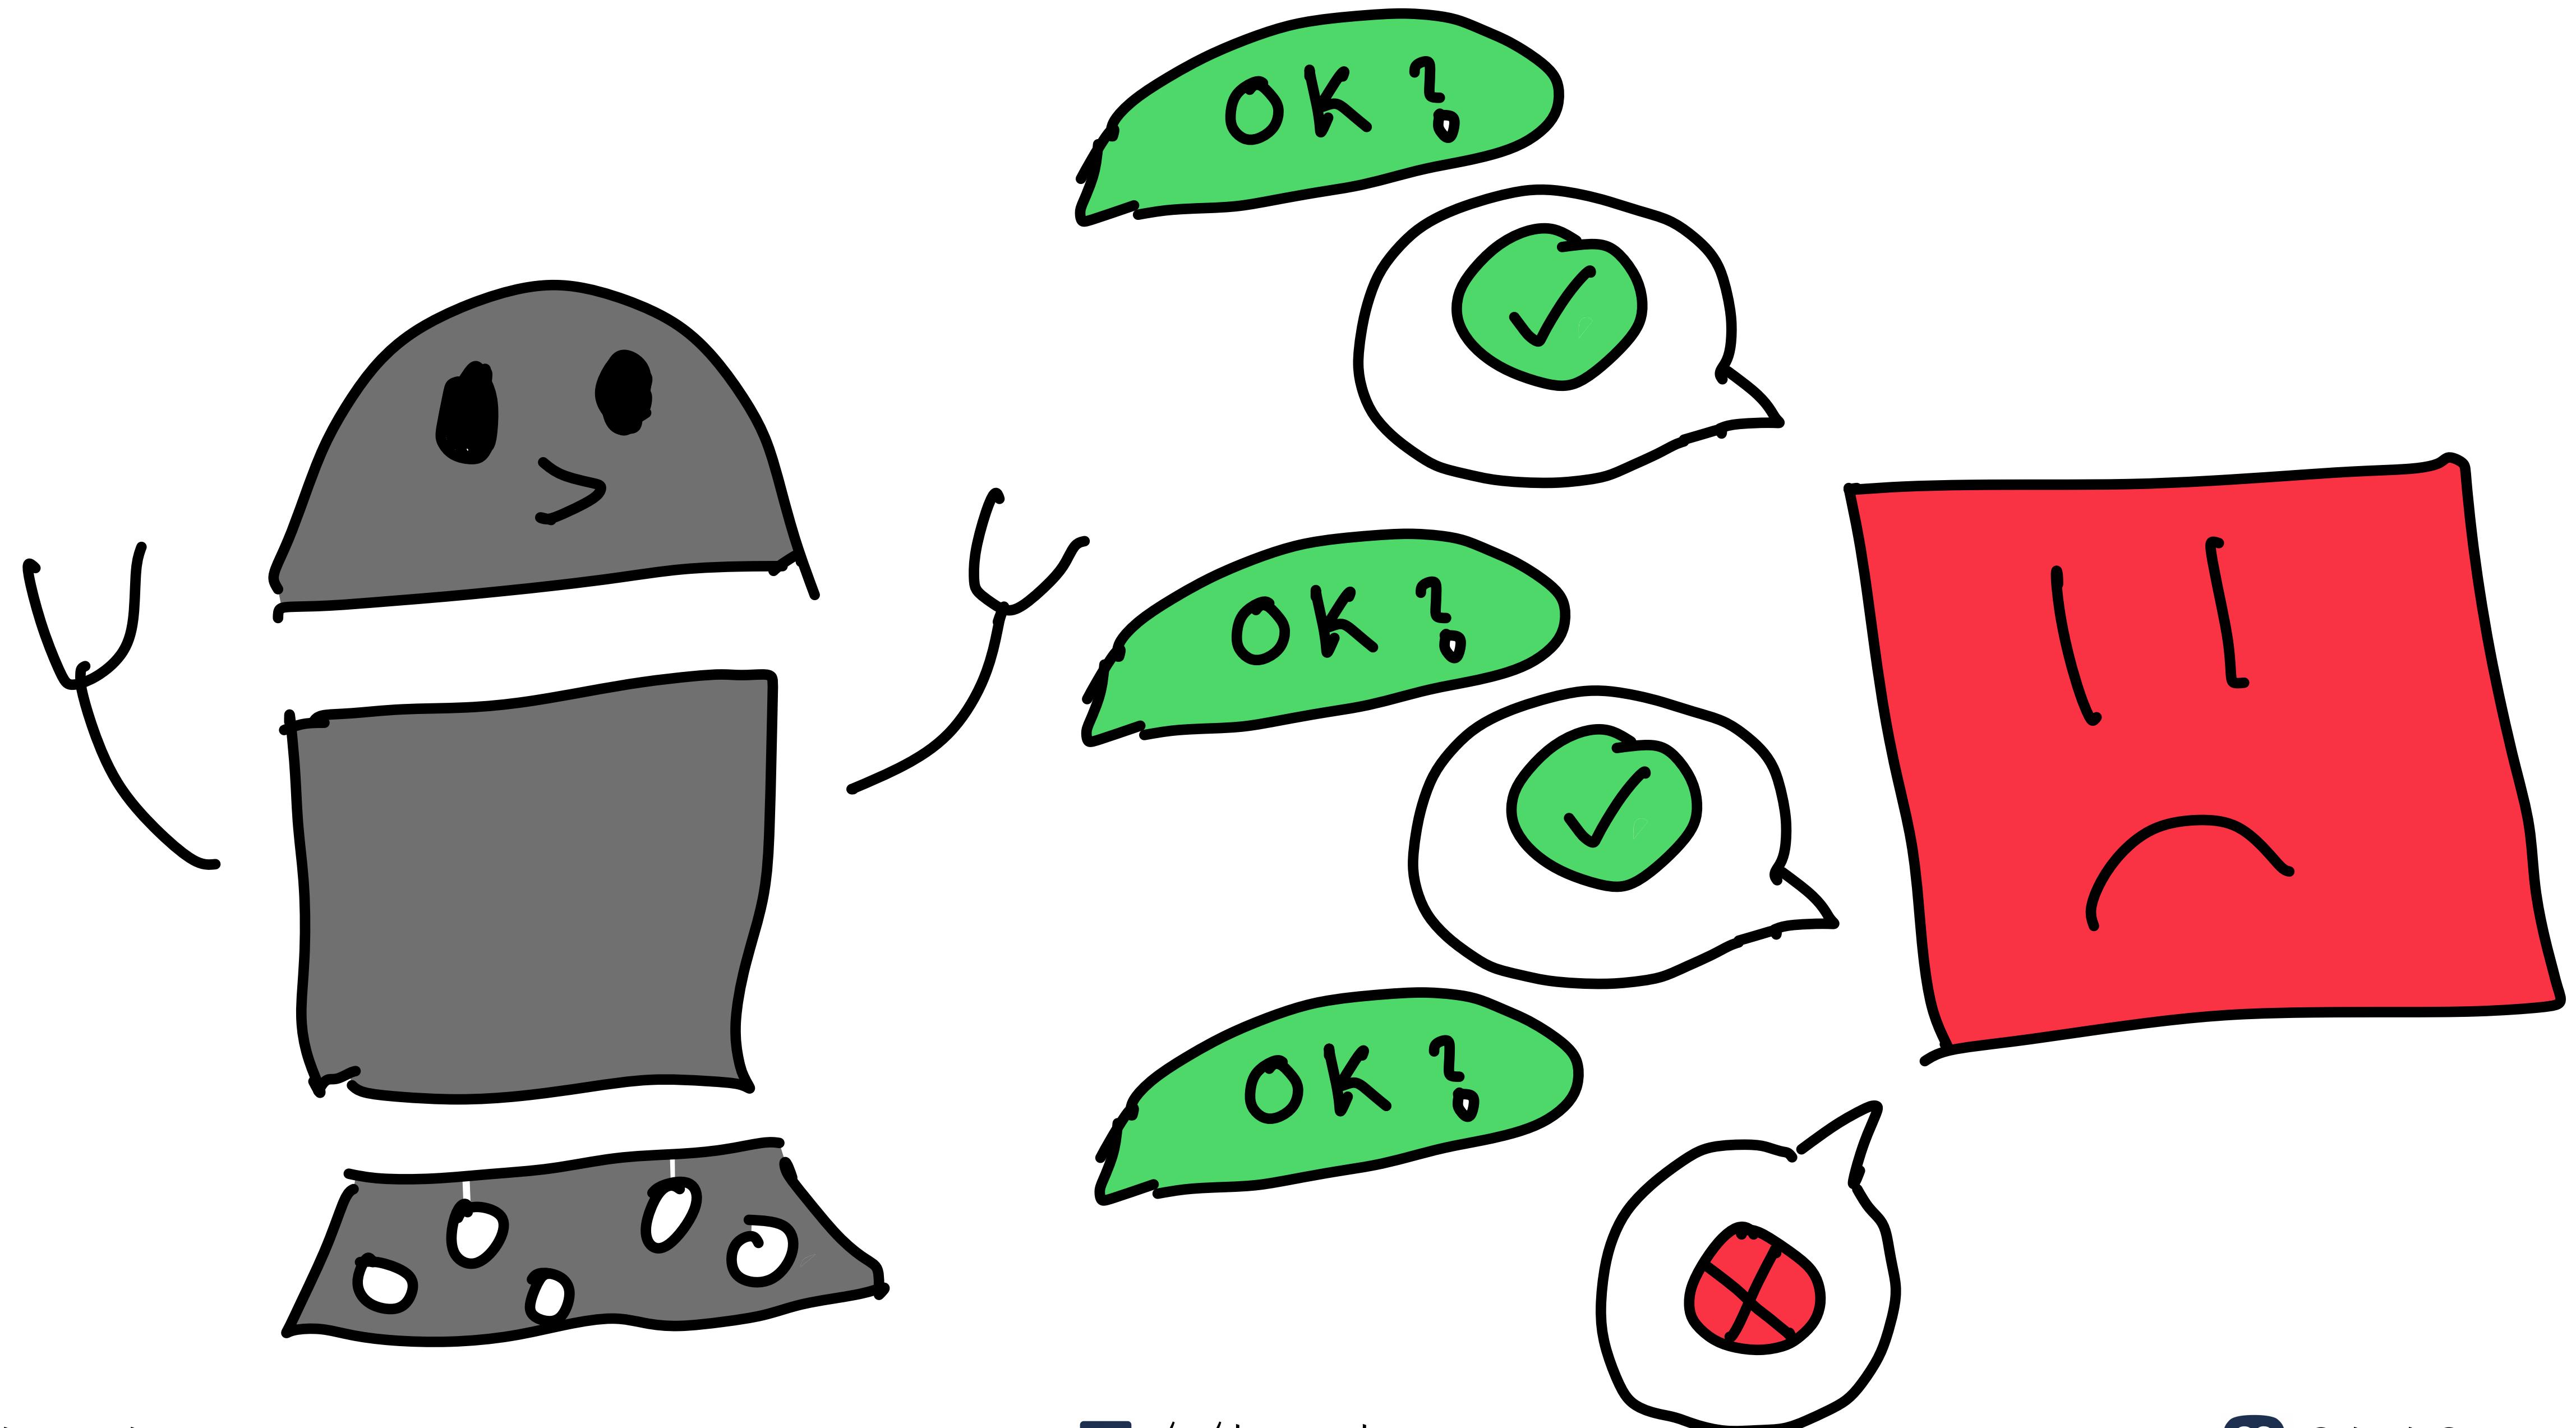


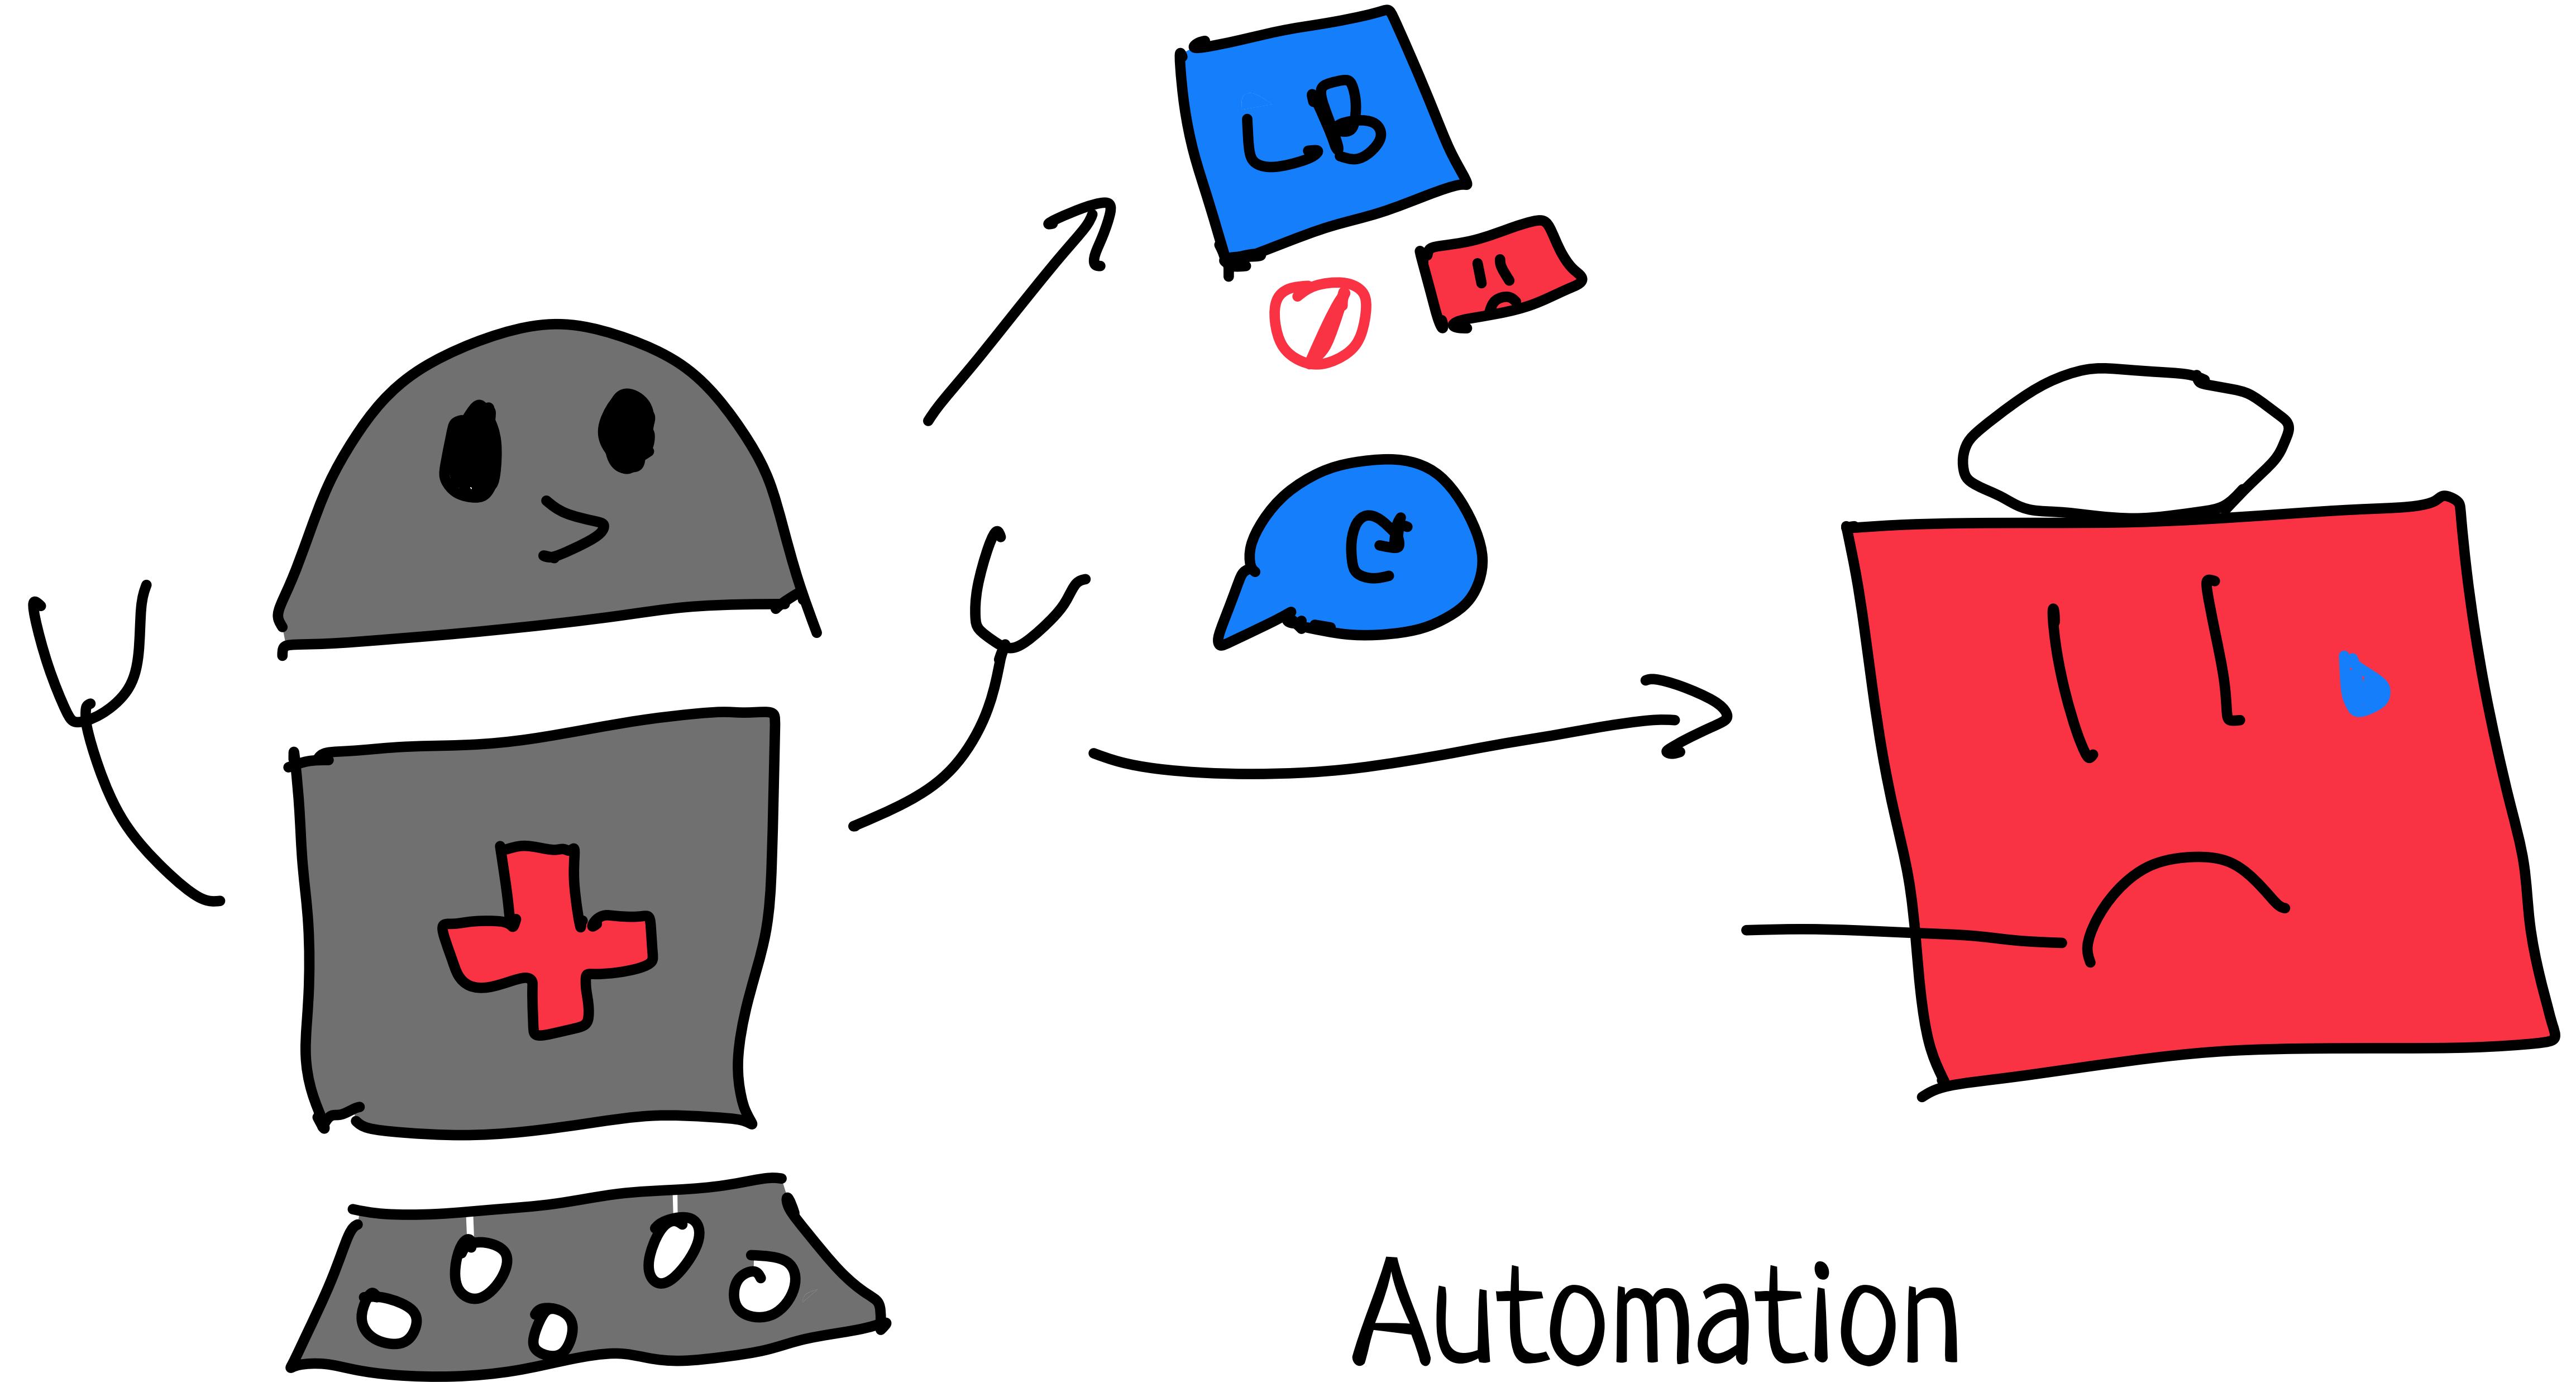




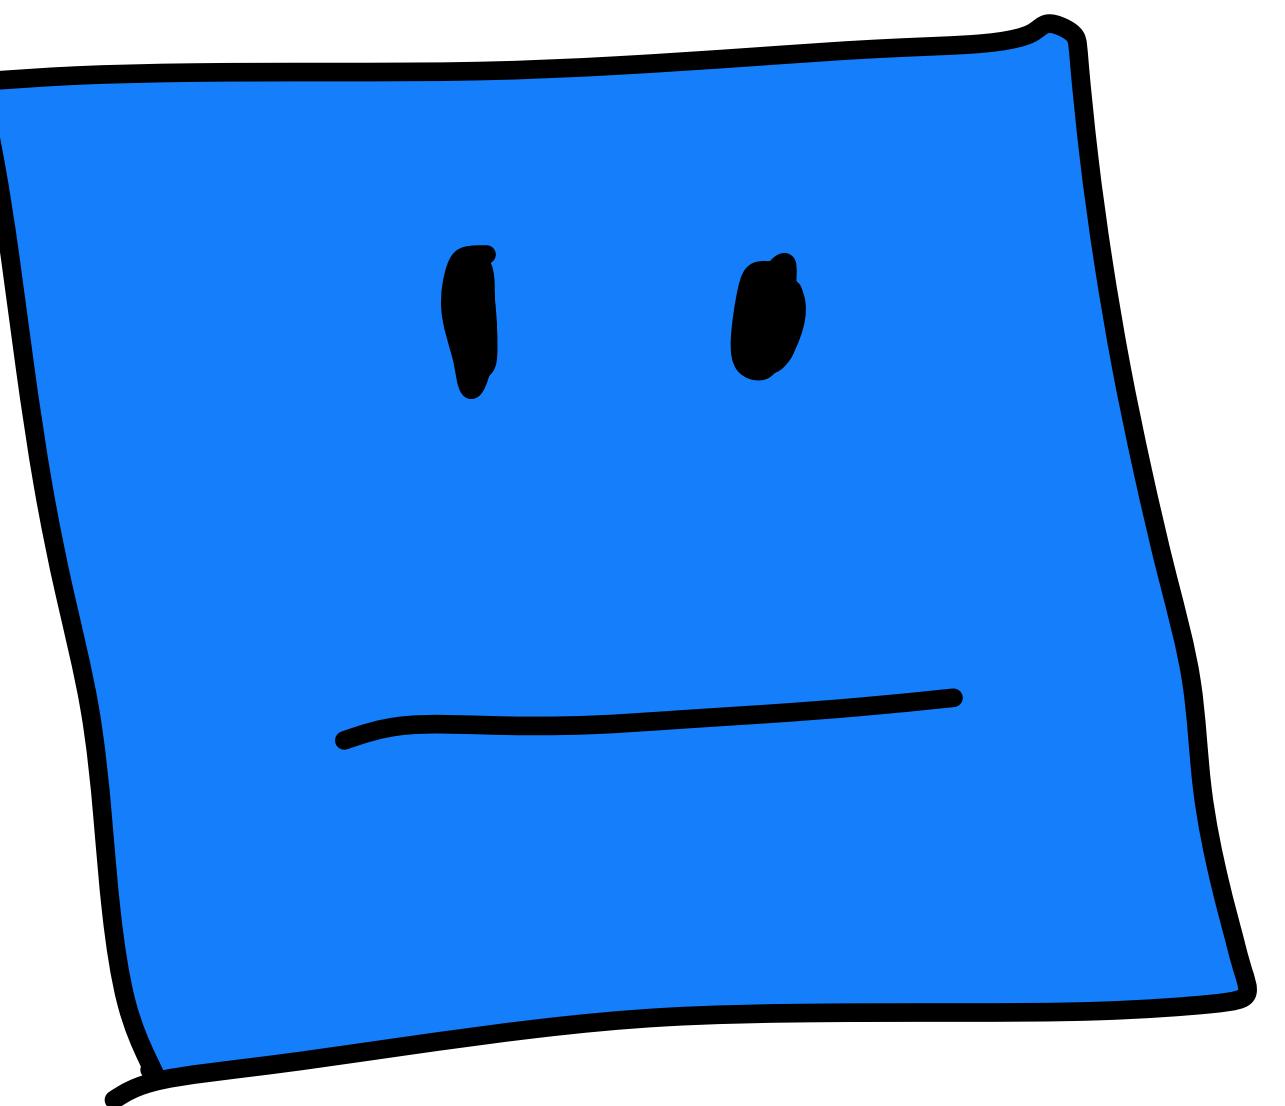
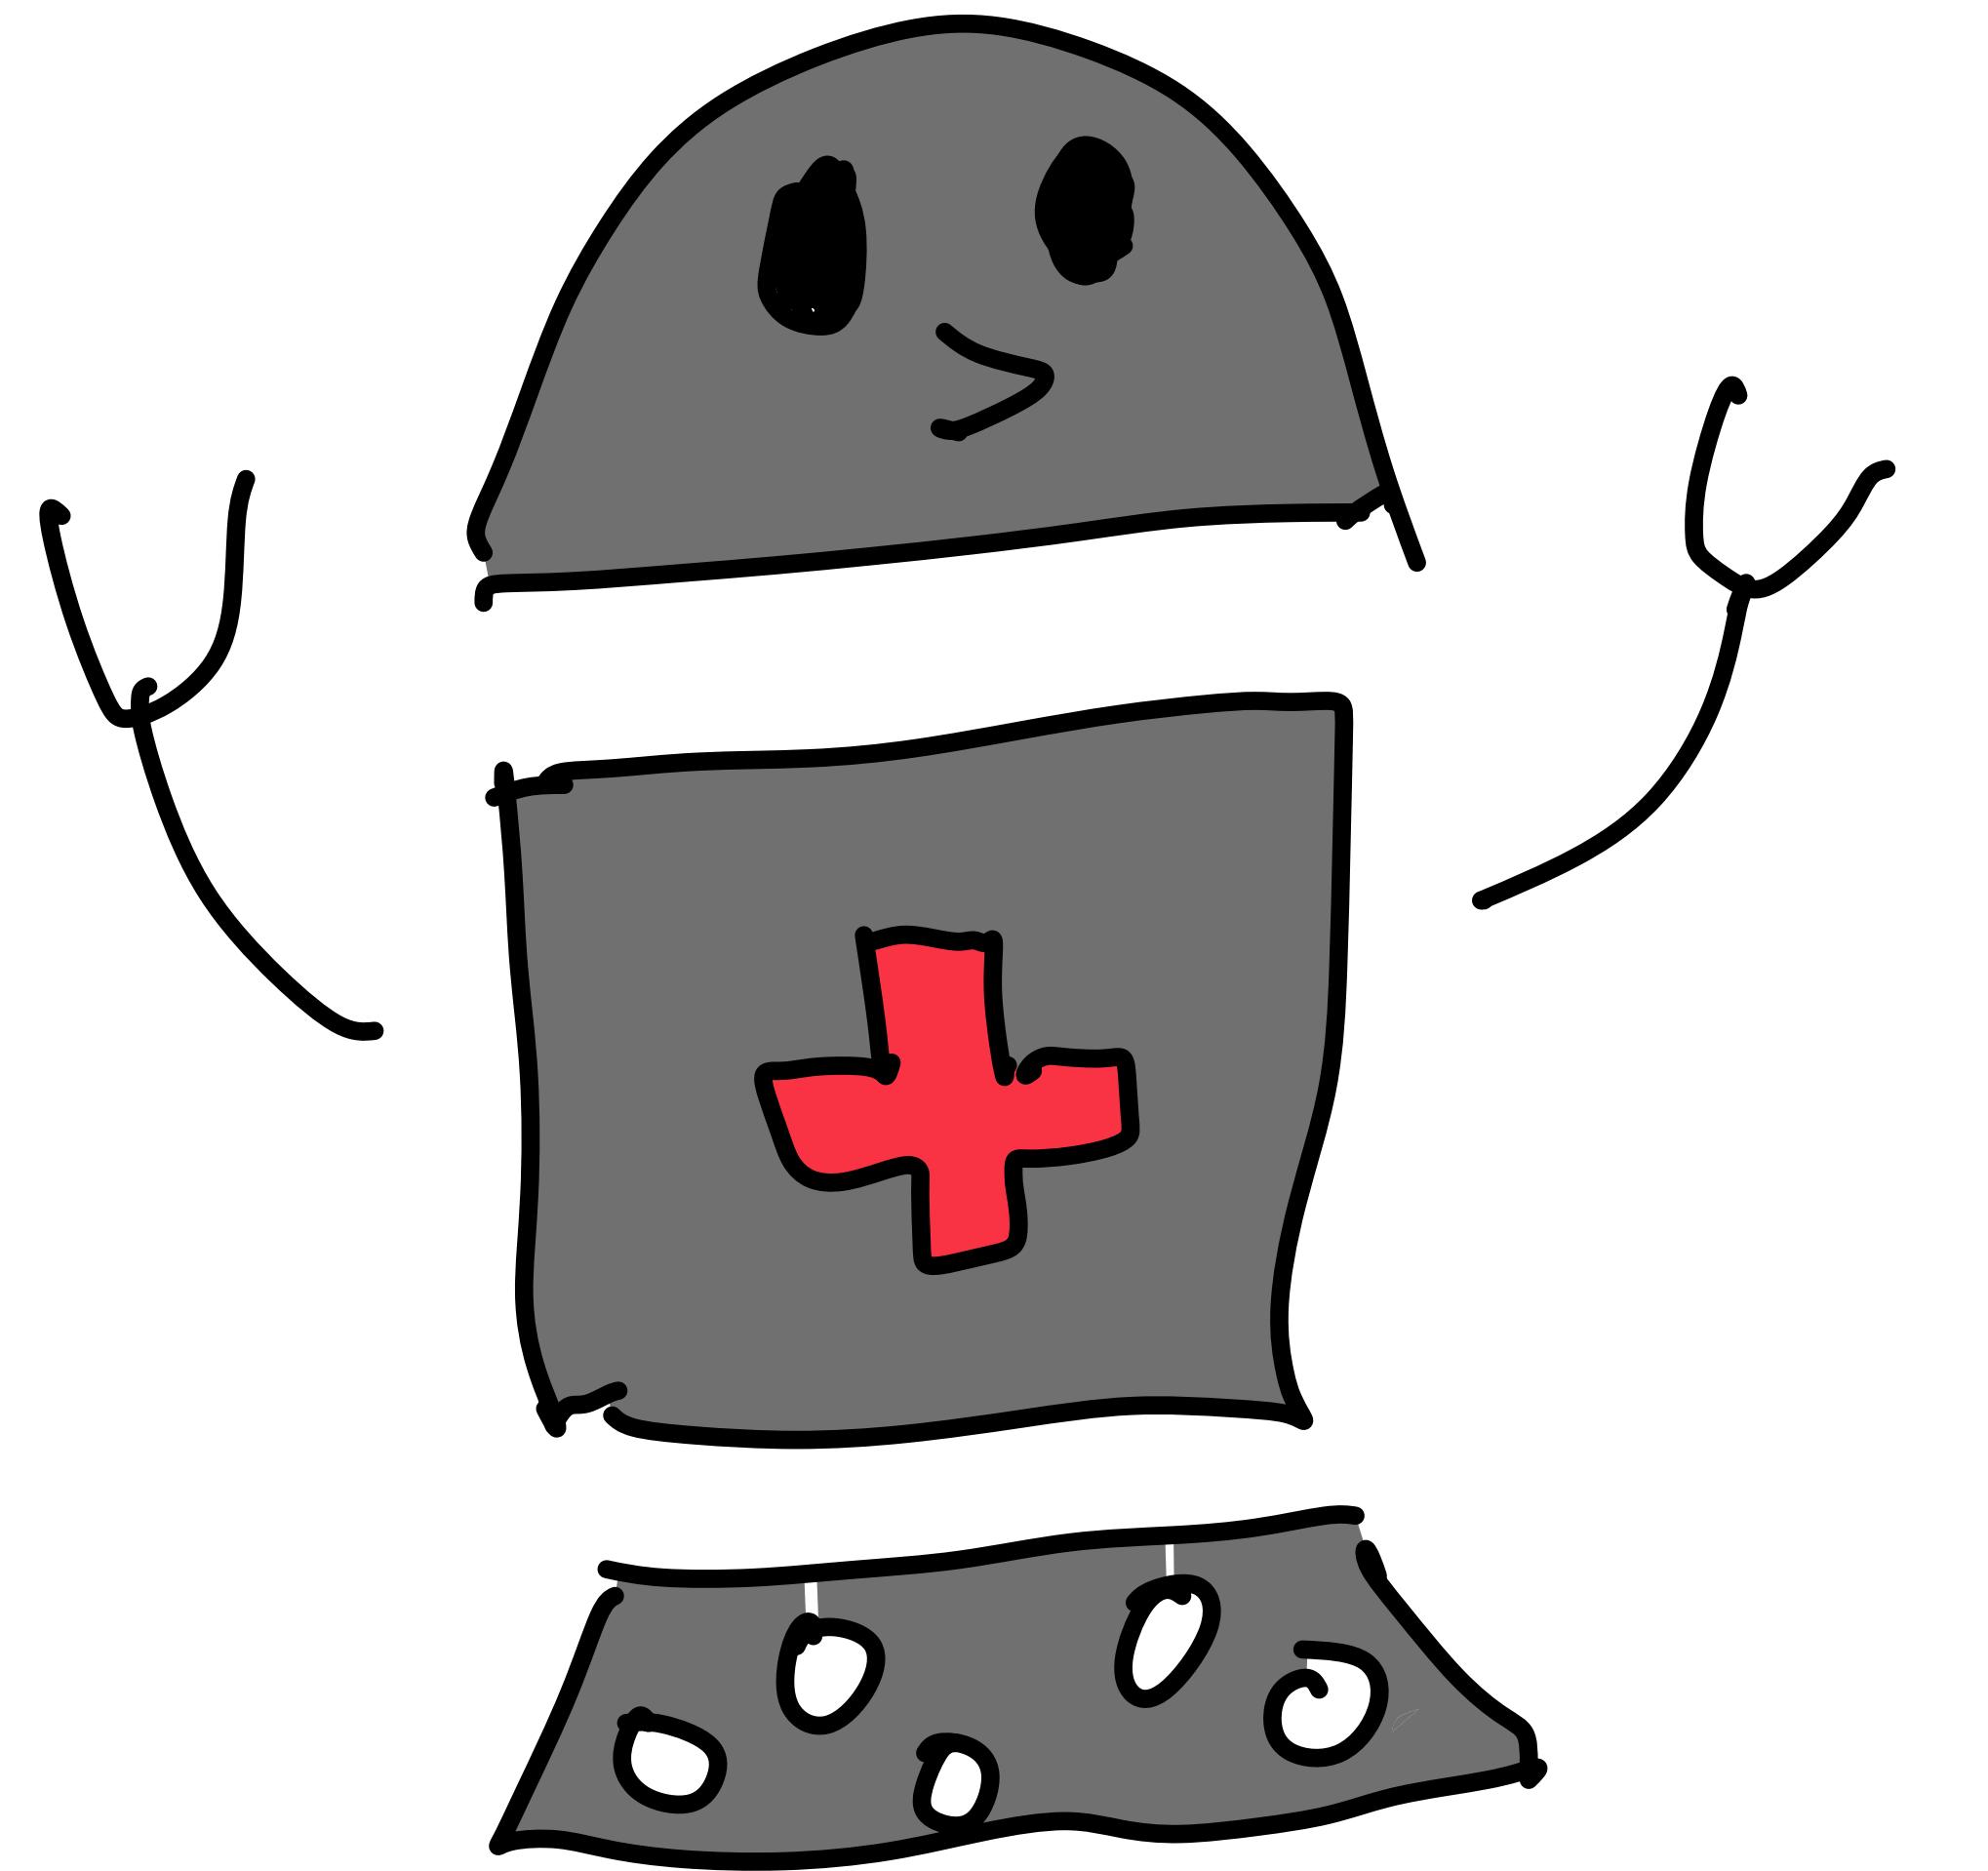


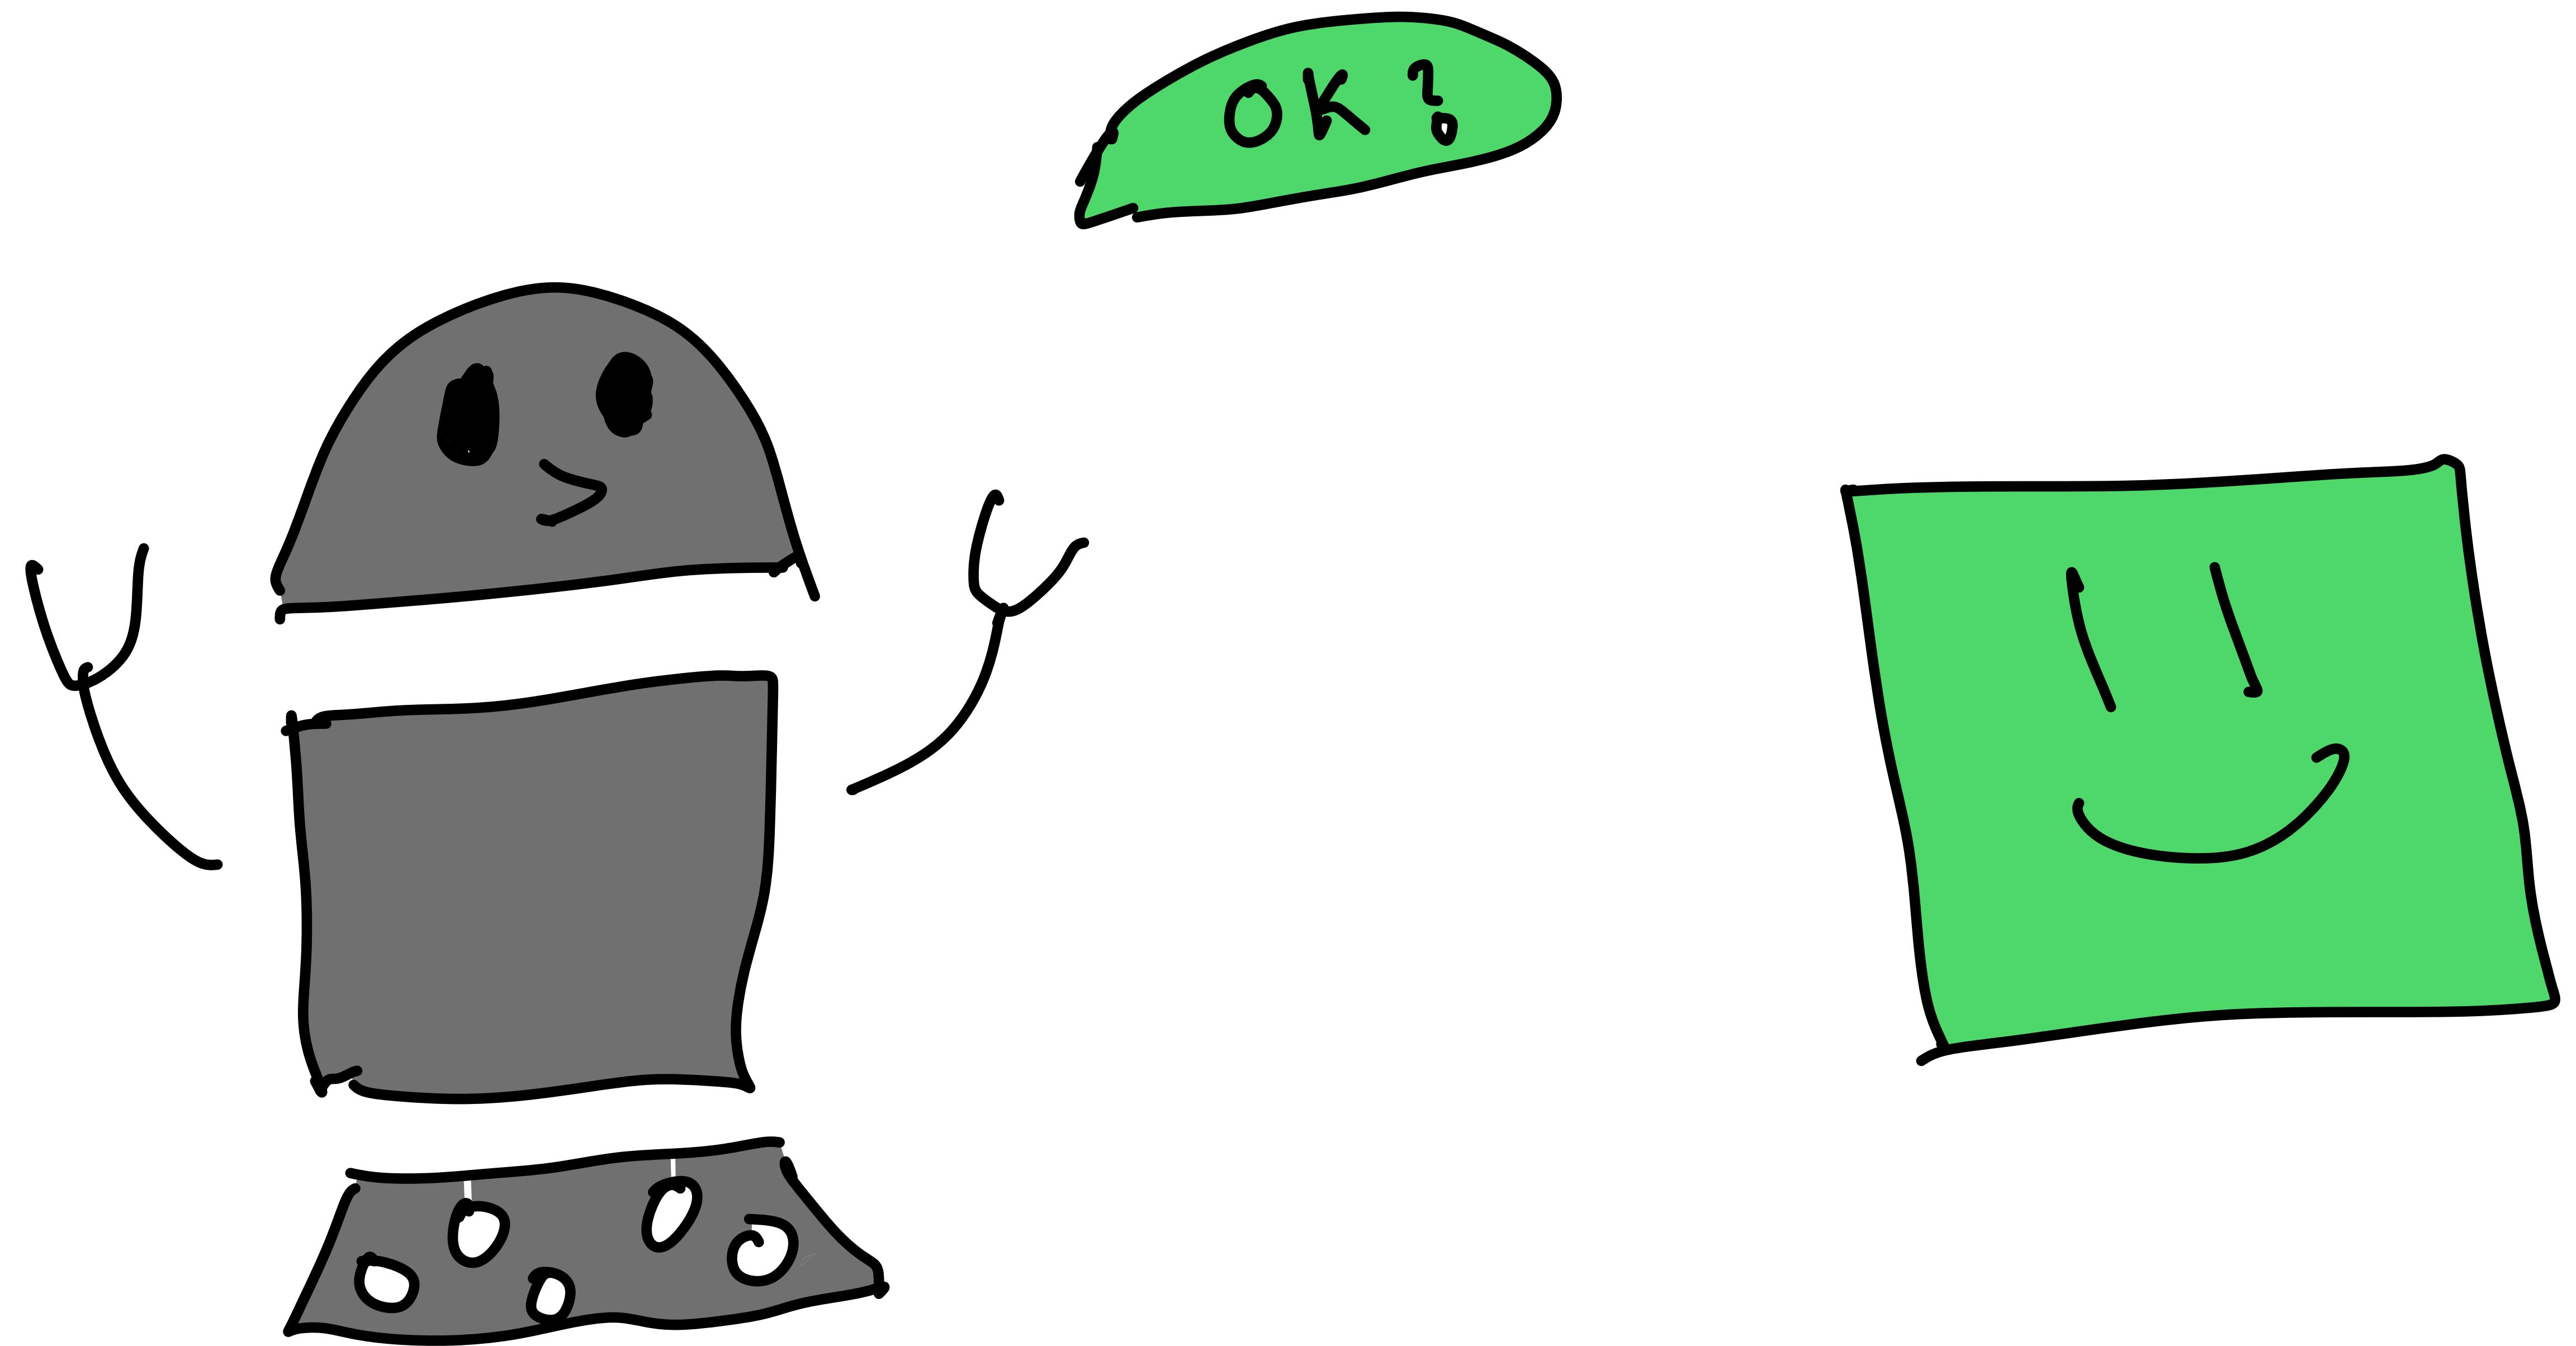


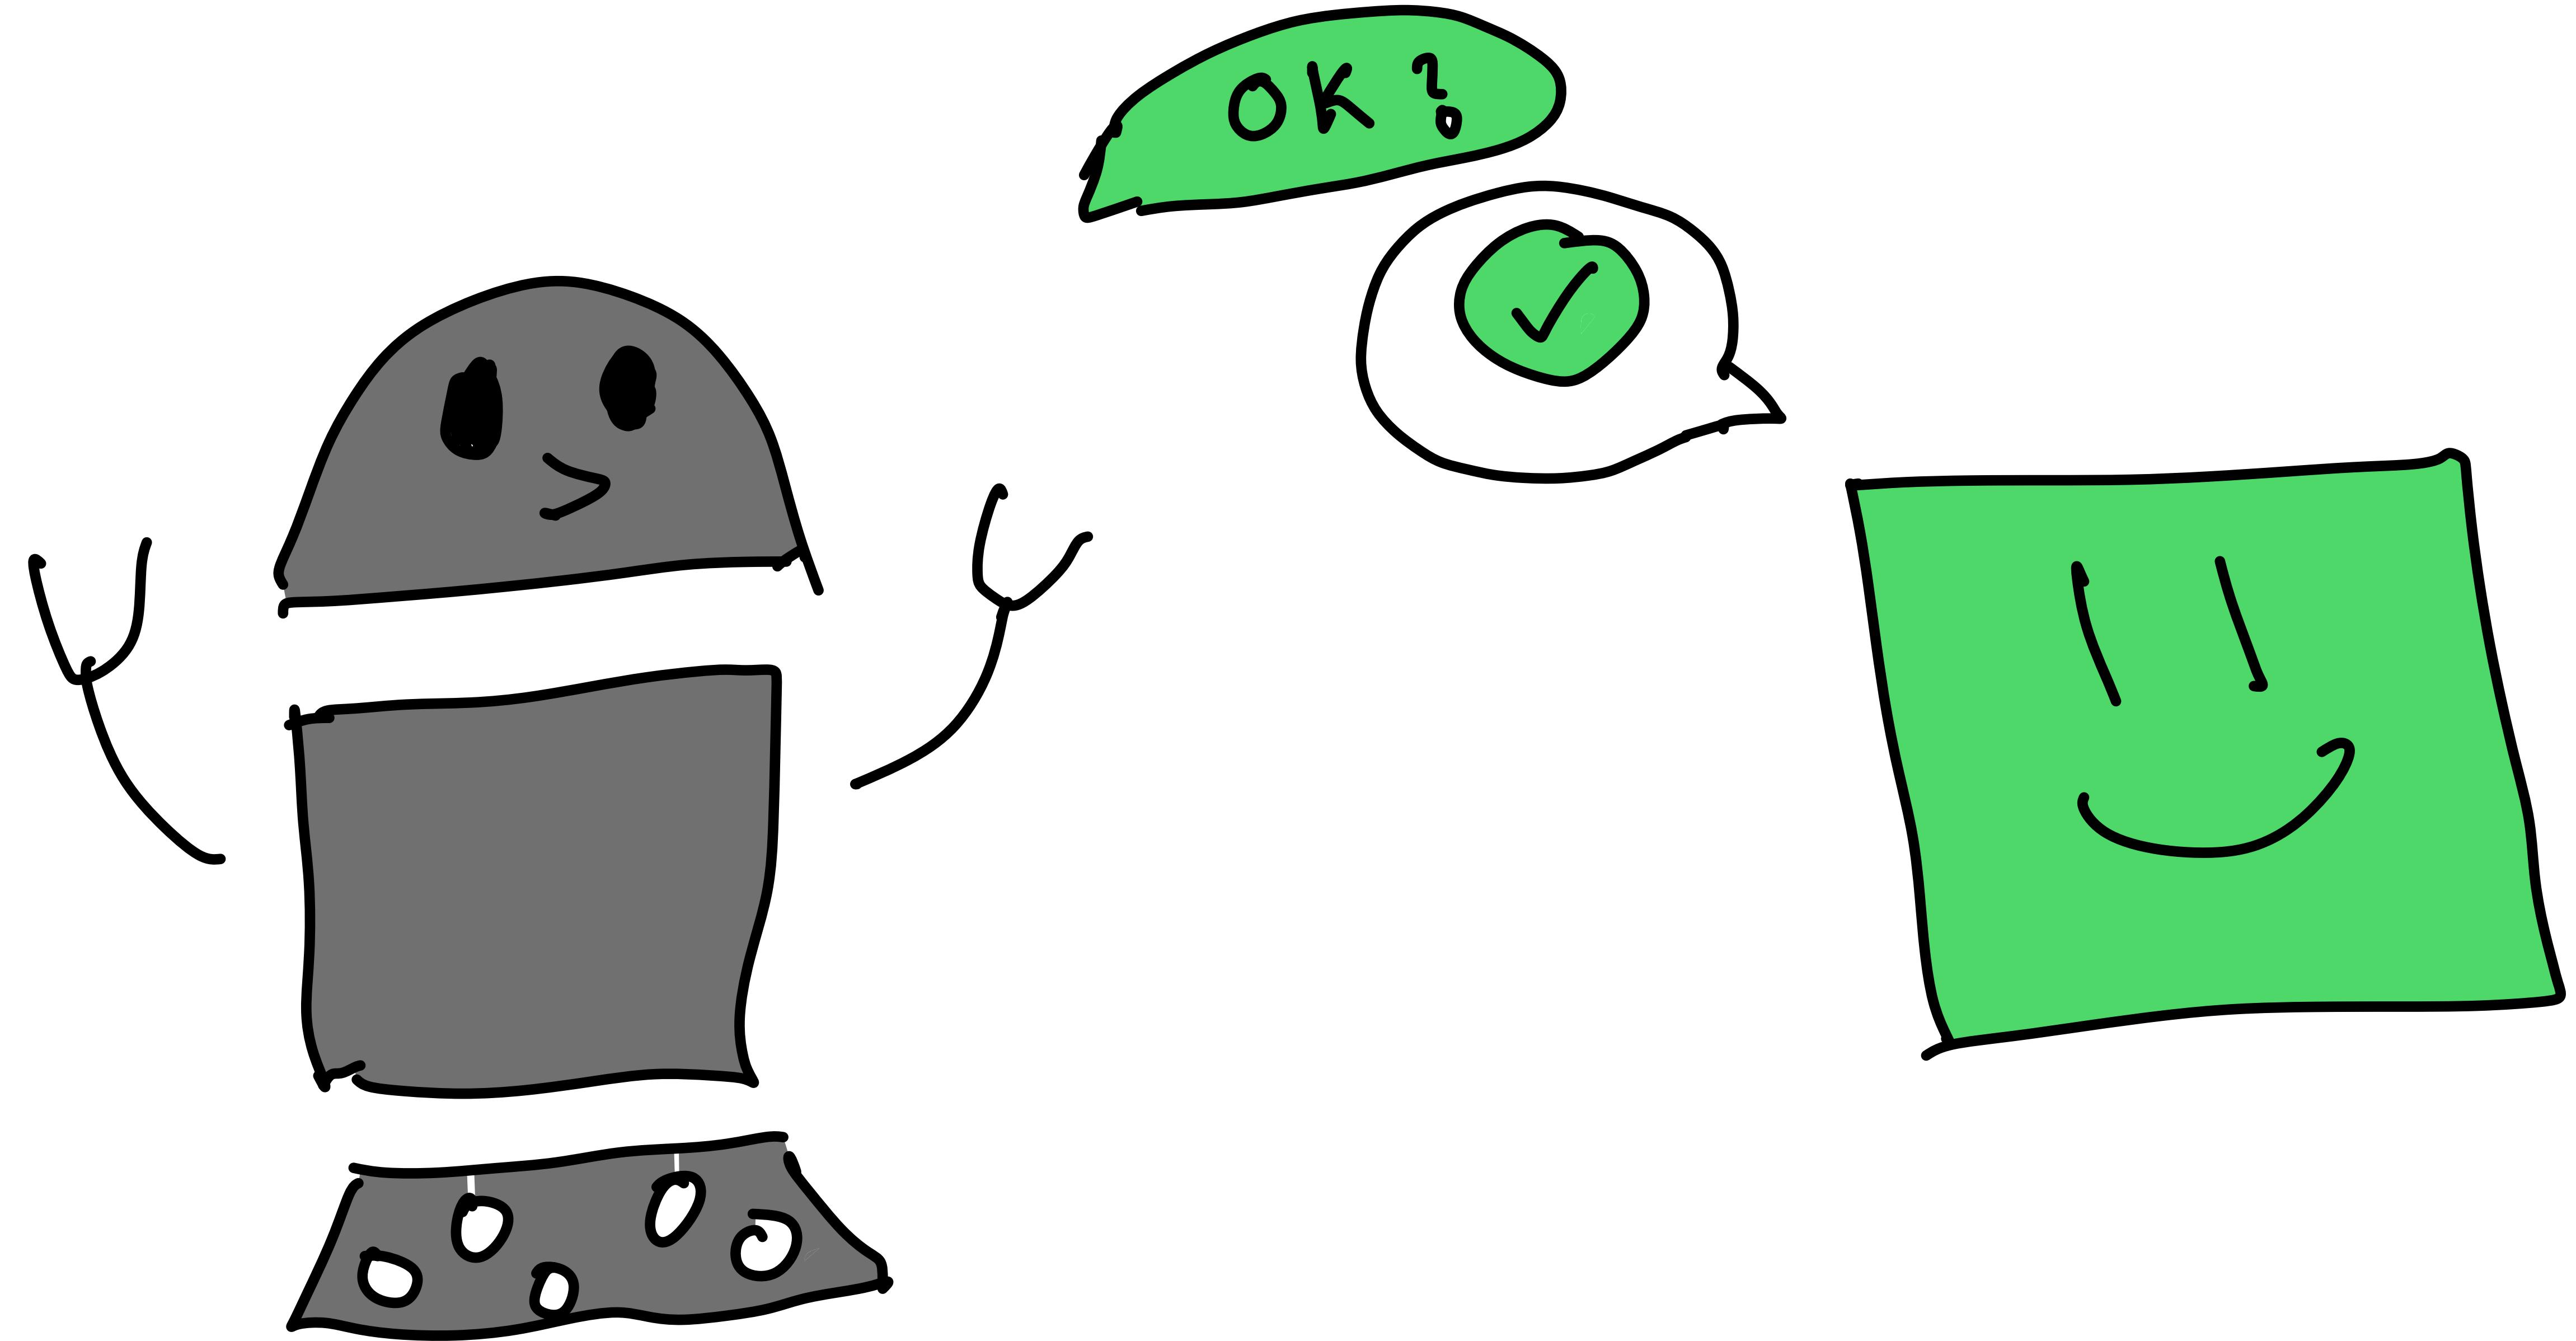




Automation



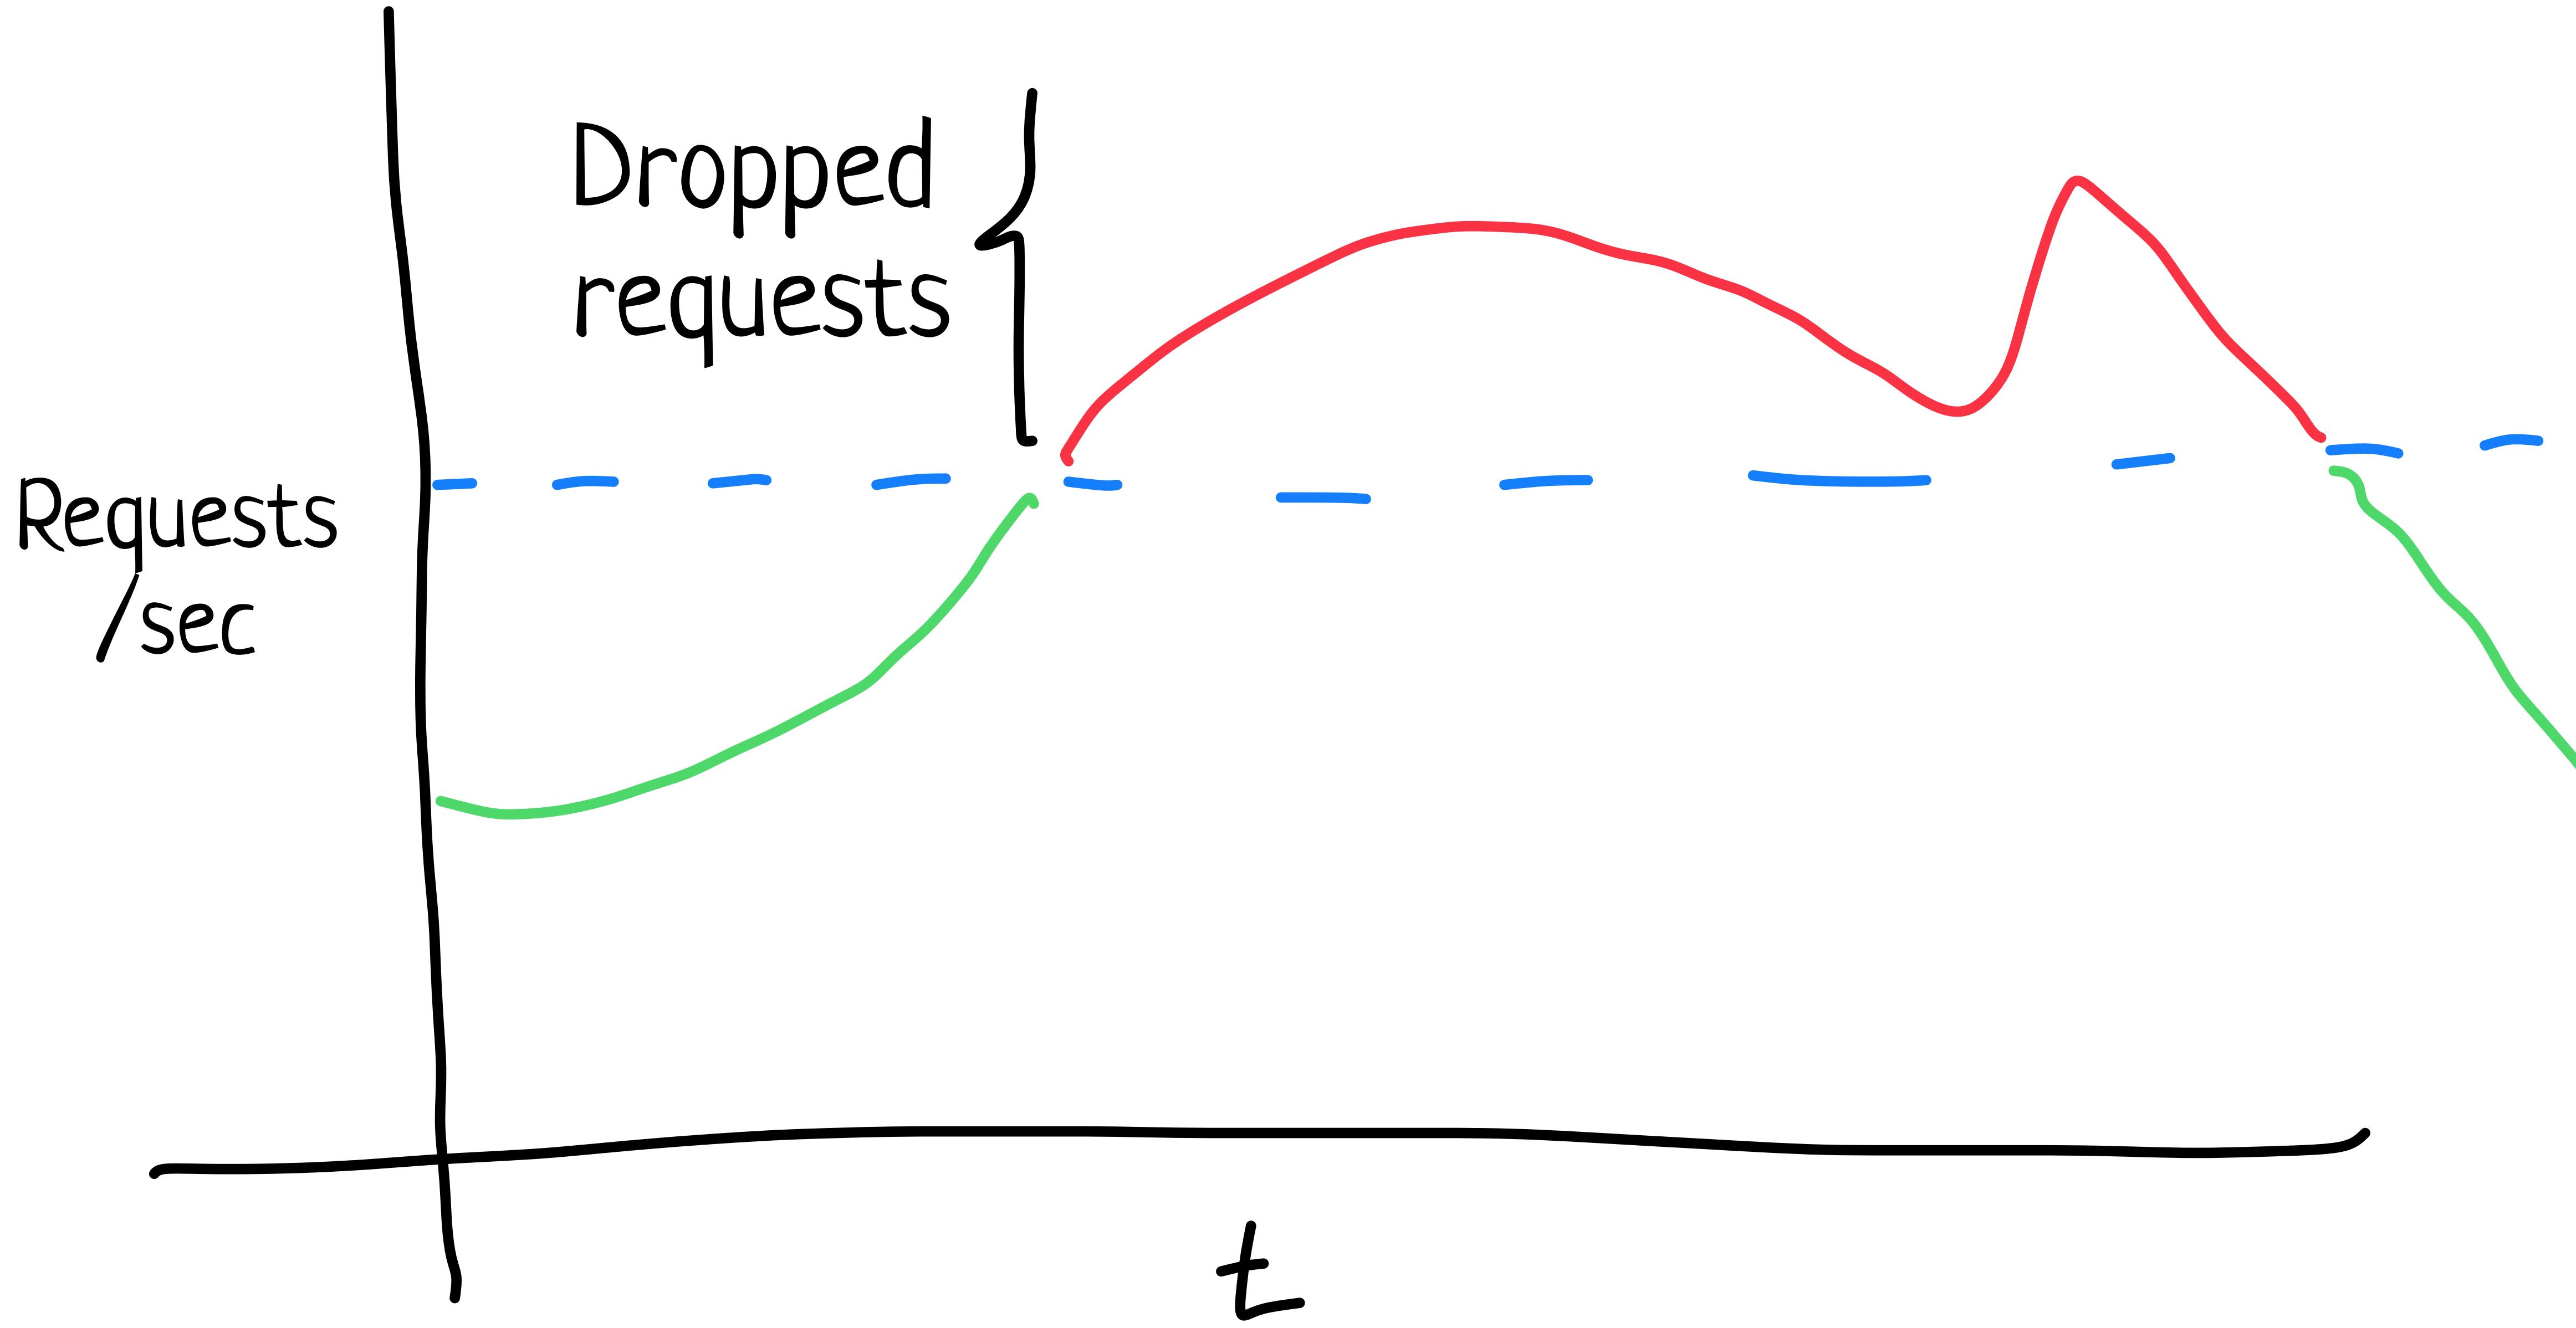


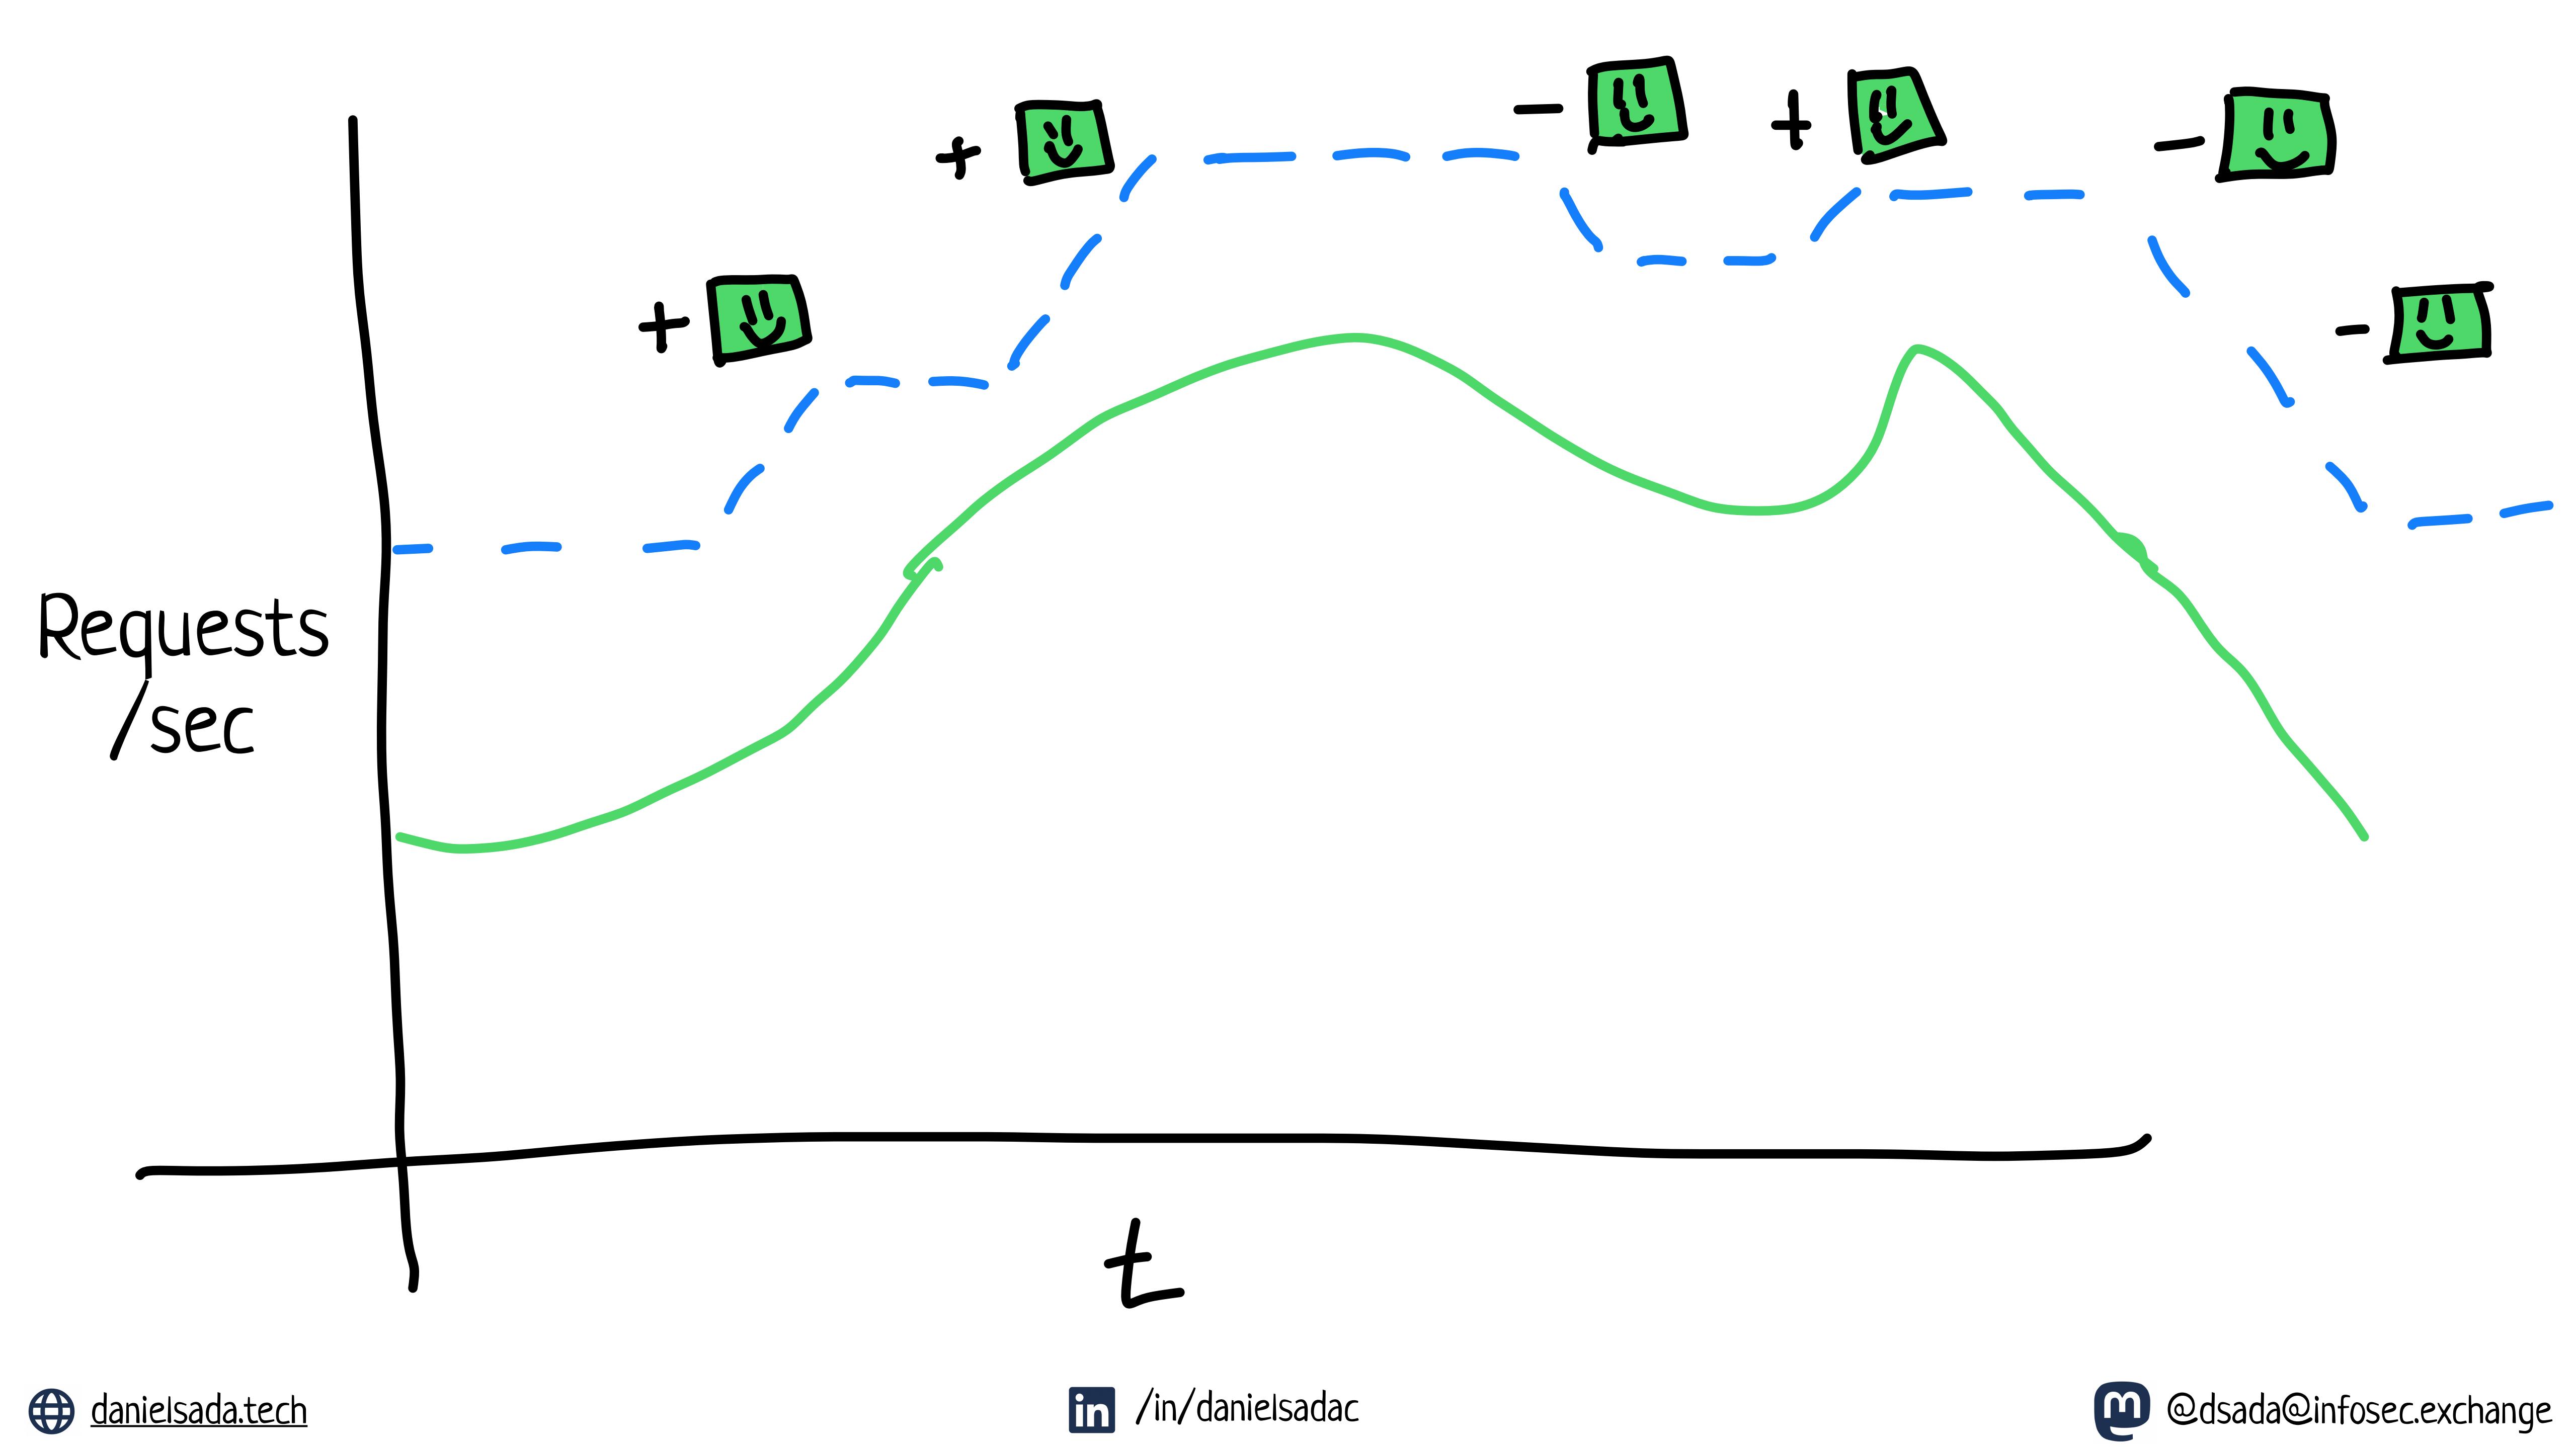


Health Checks / Automation

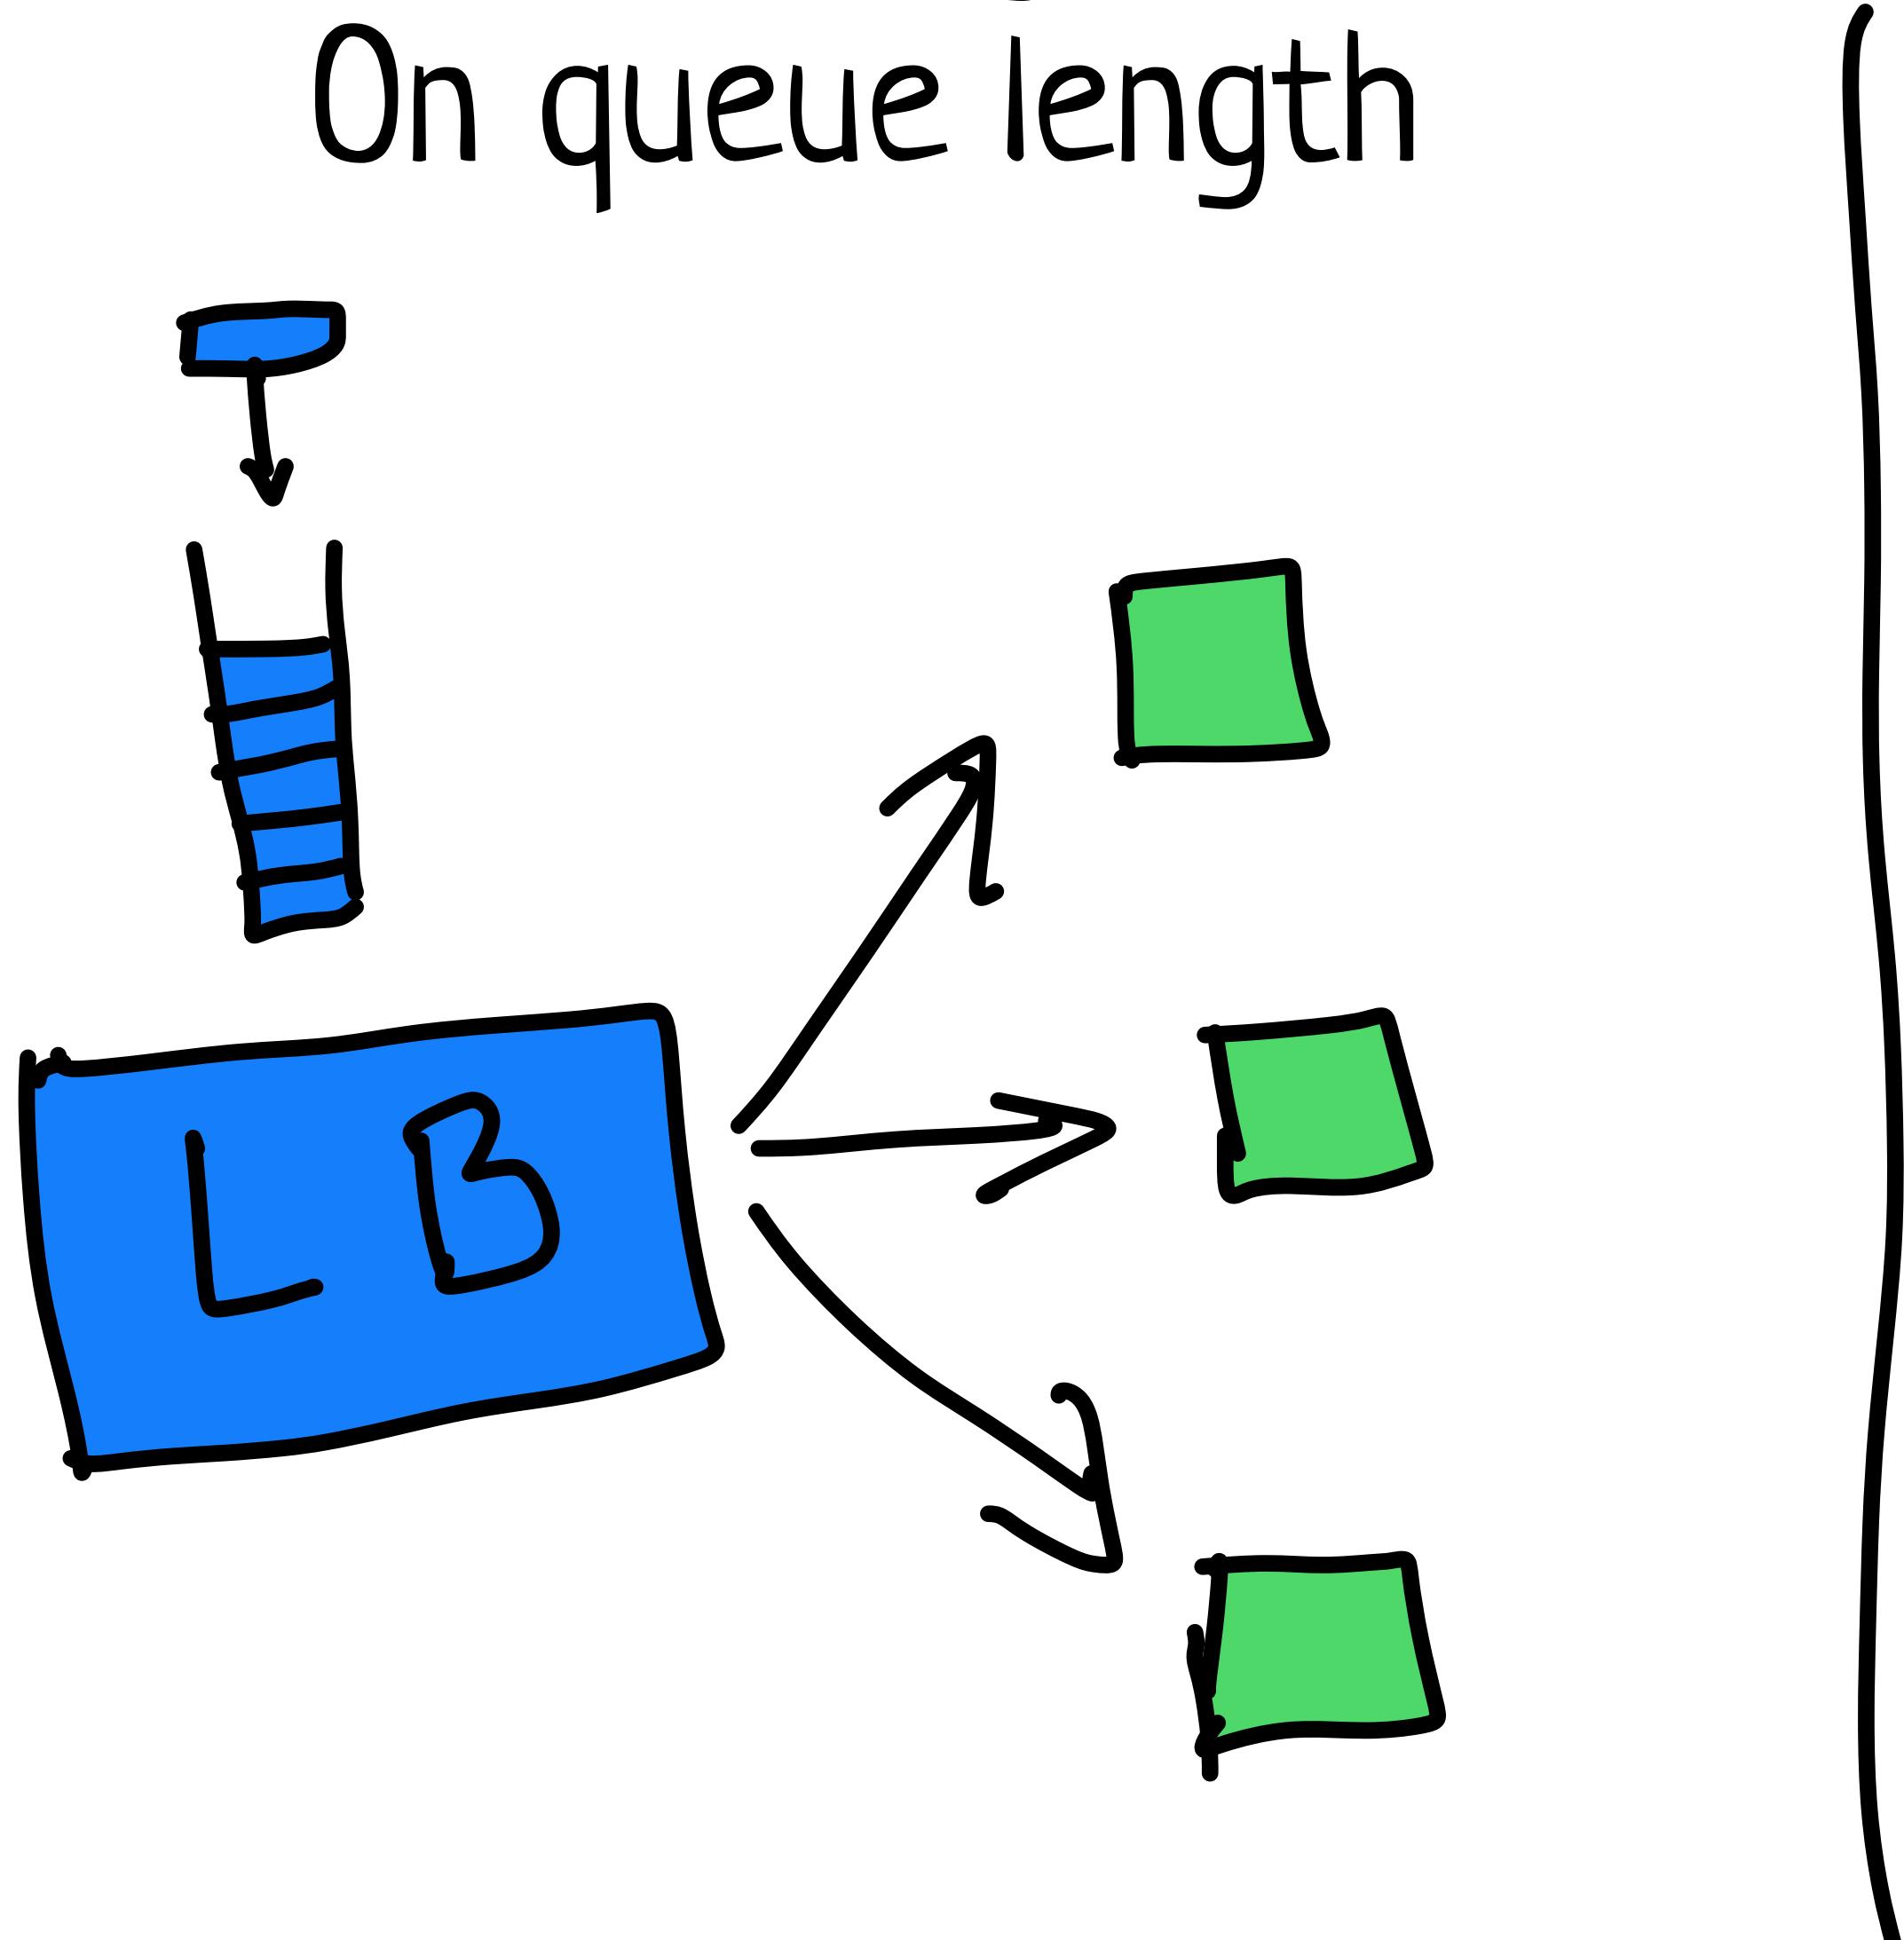
- Mostly using orchestrators
- On other environments, you can use your load balancer or expose endpoints.
- Cloud specific offerings.

Autoscaling

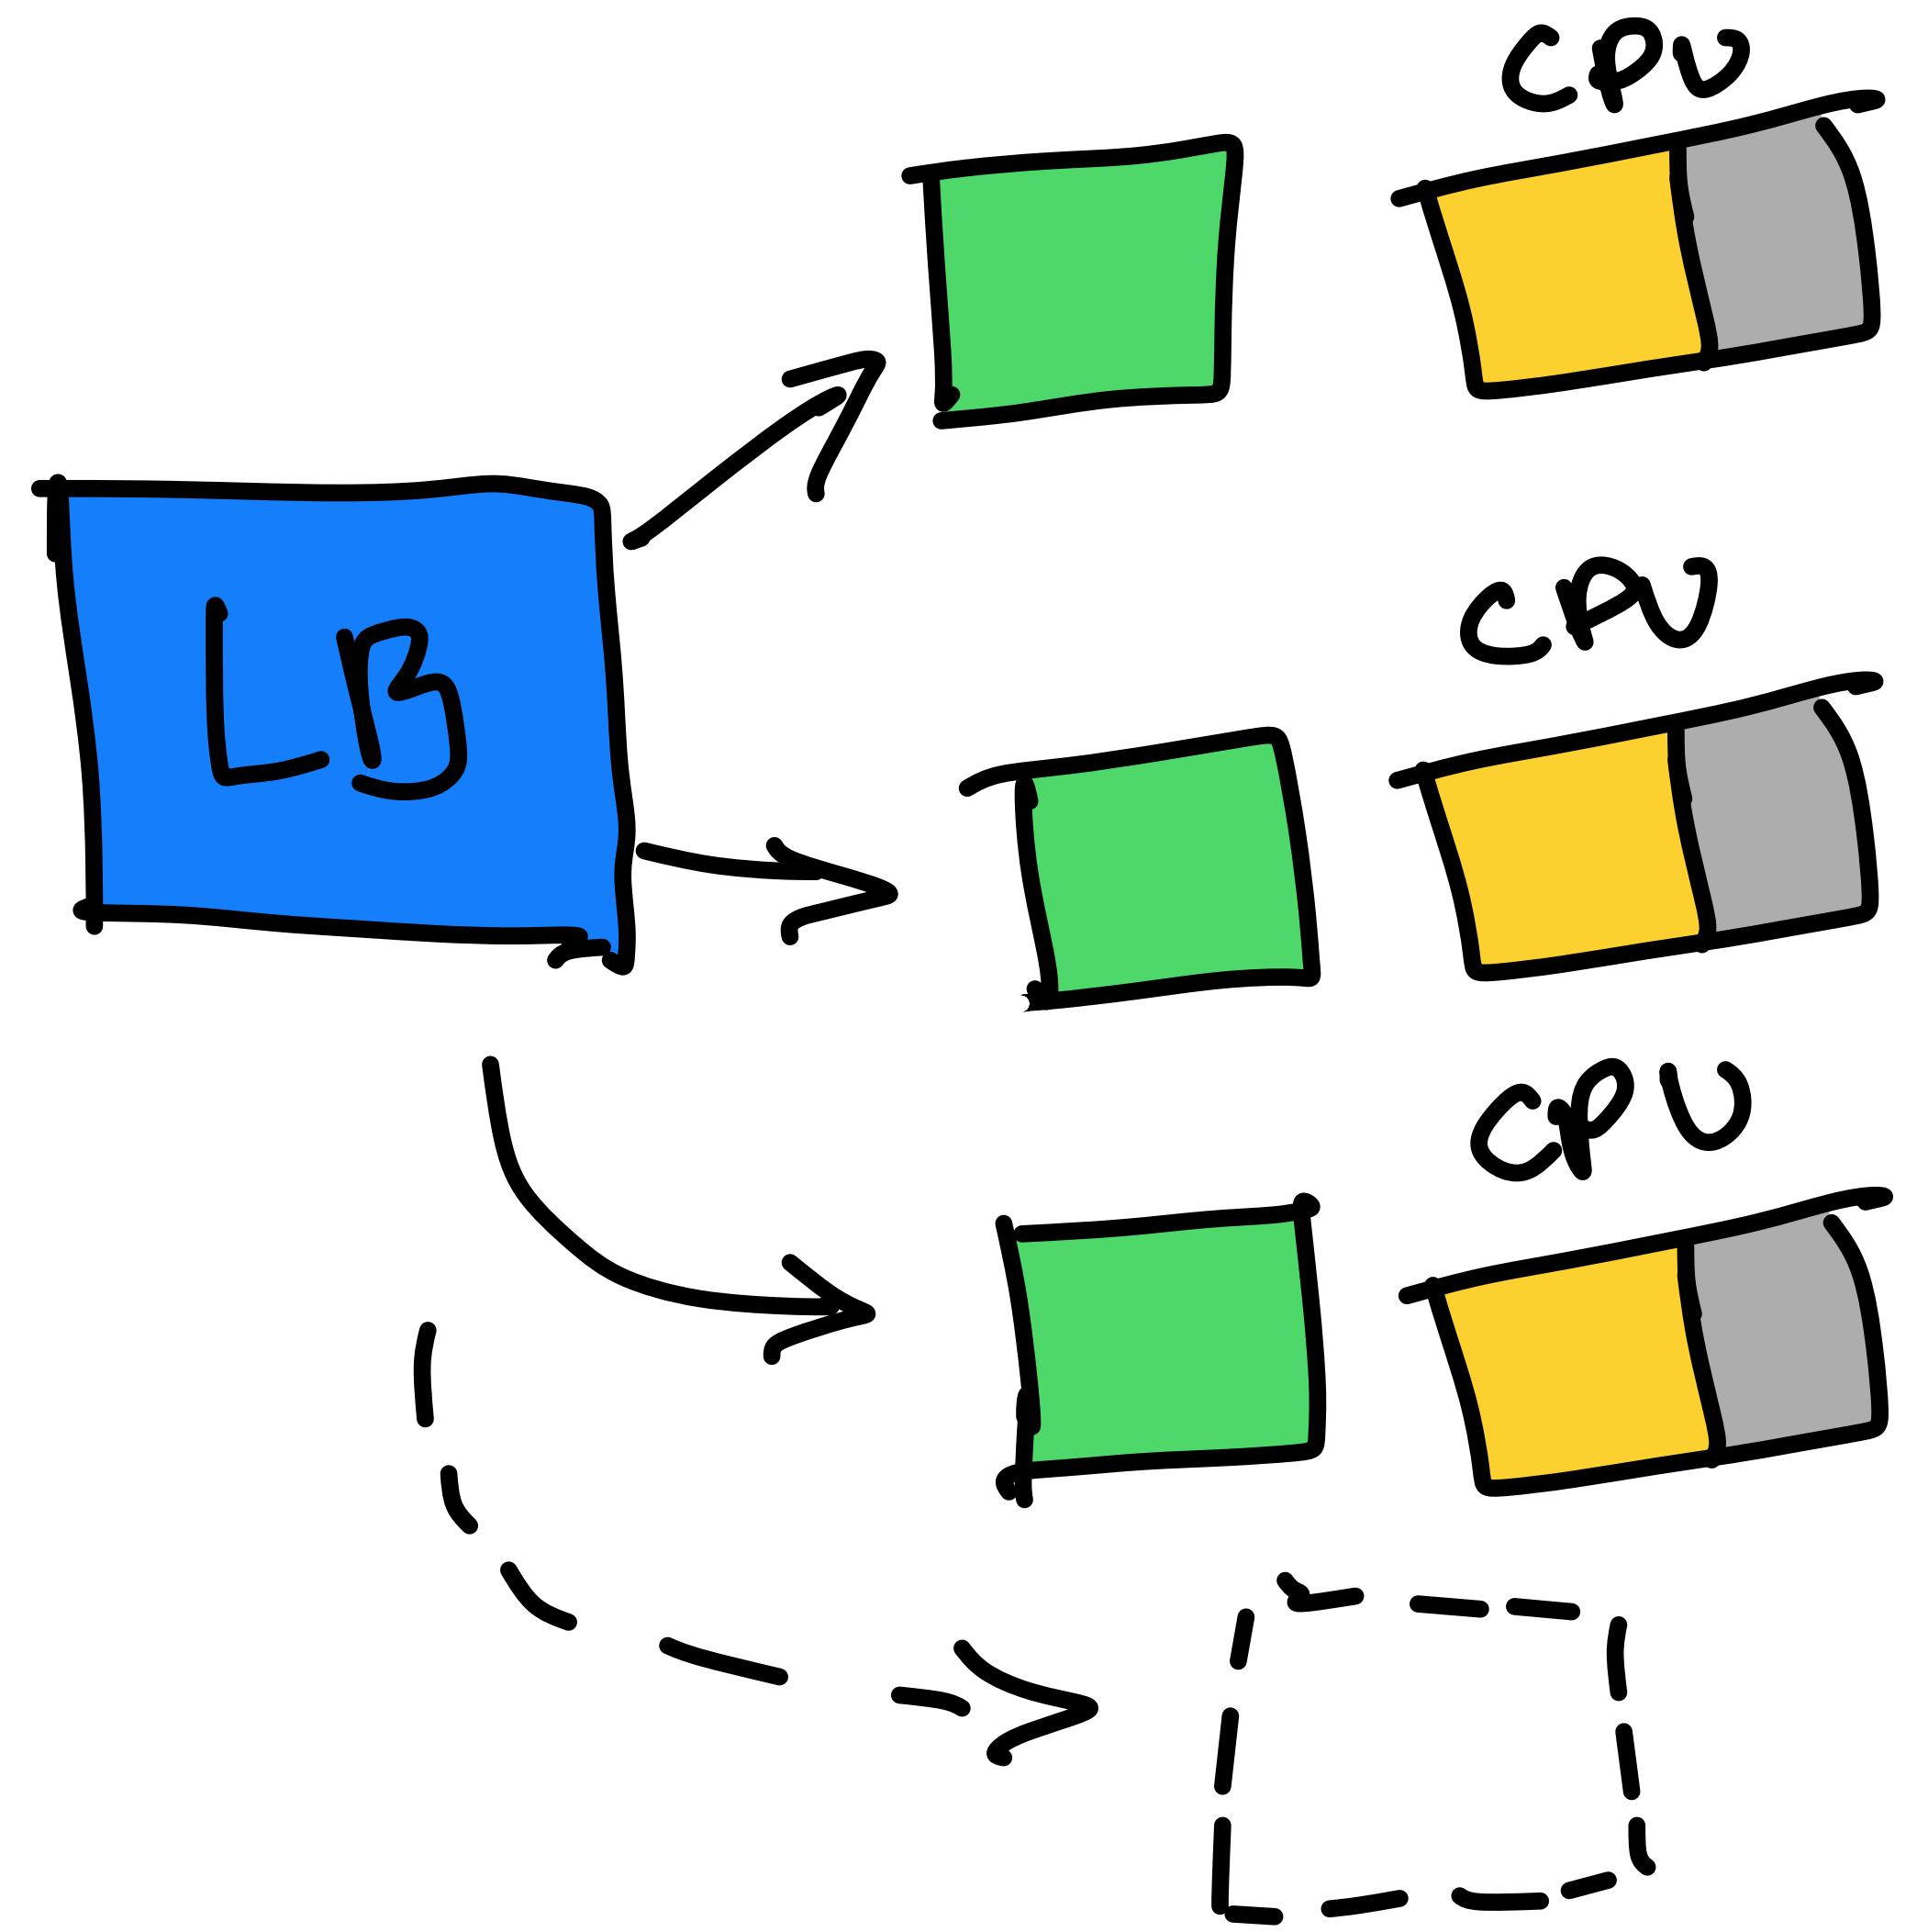




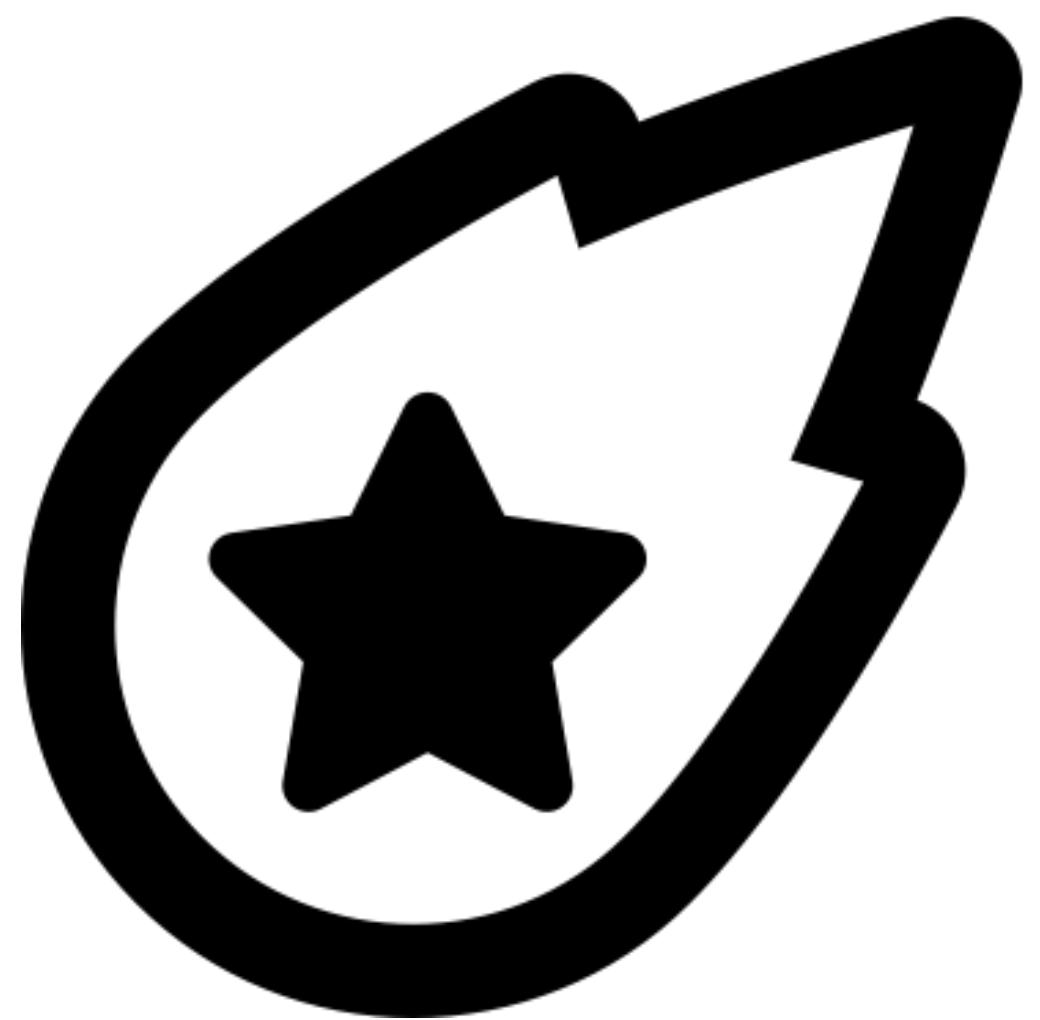
Autoscaling based On queue length



Autoscaling based On resource usage

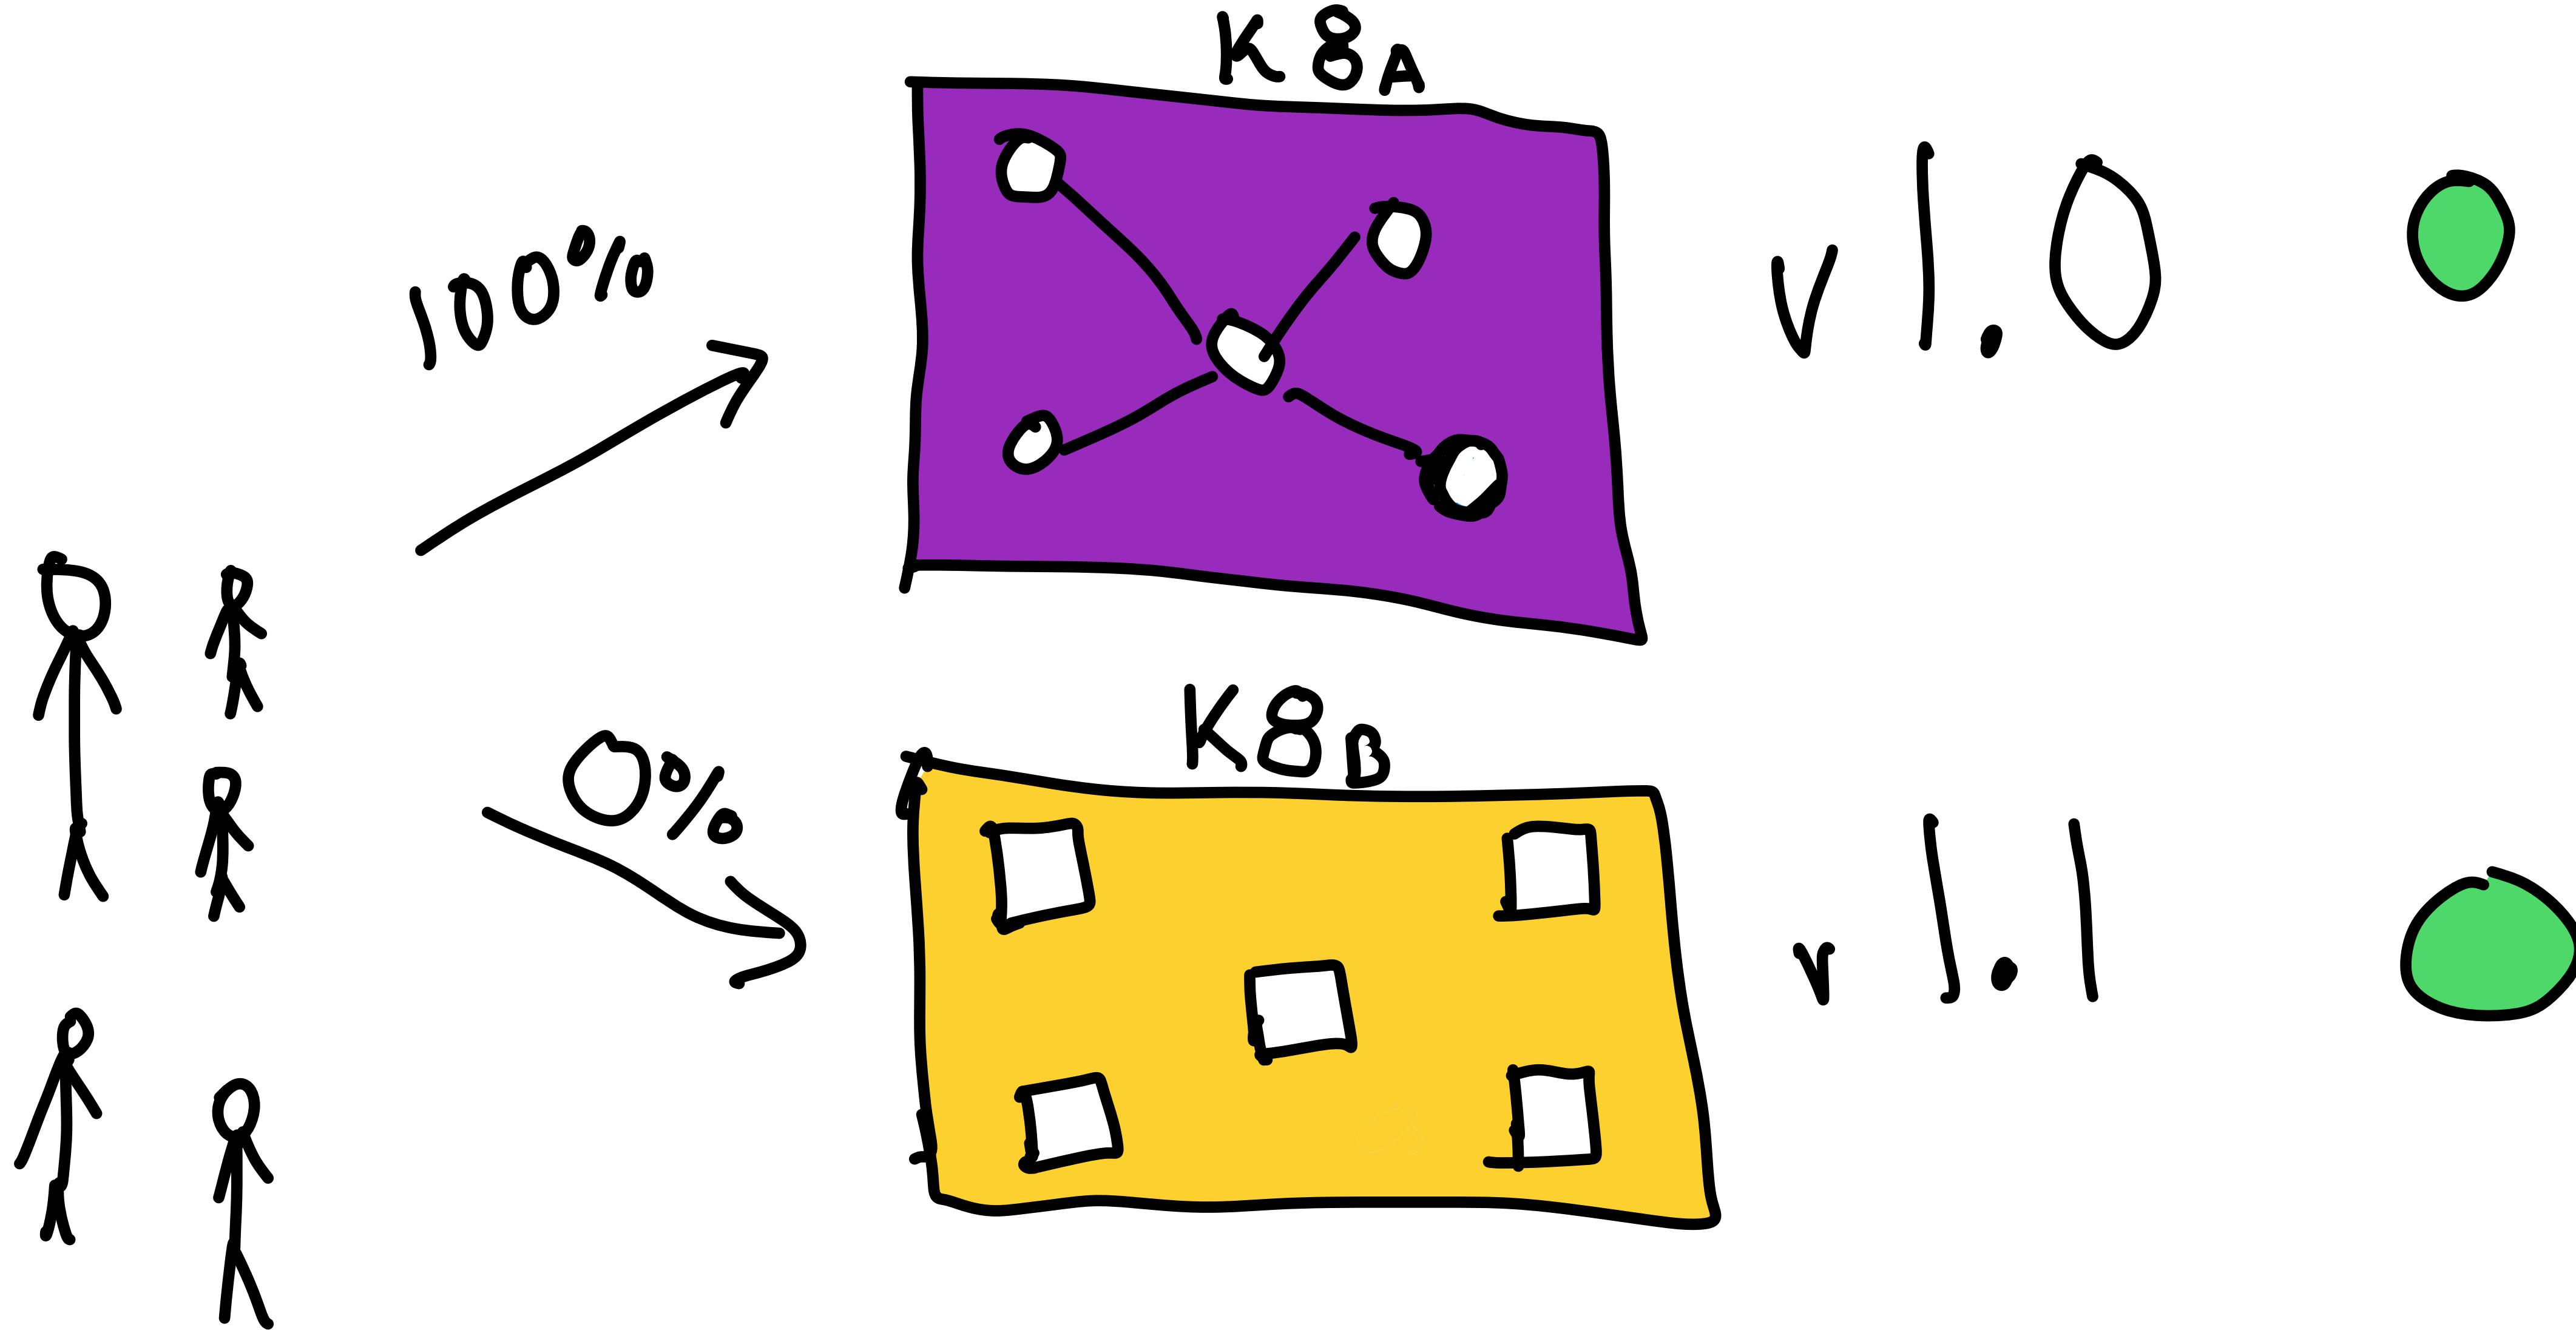


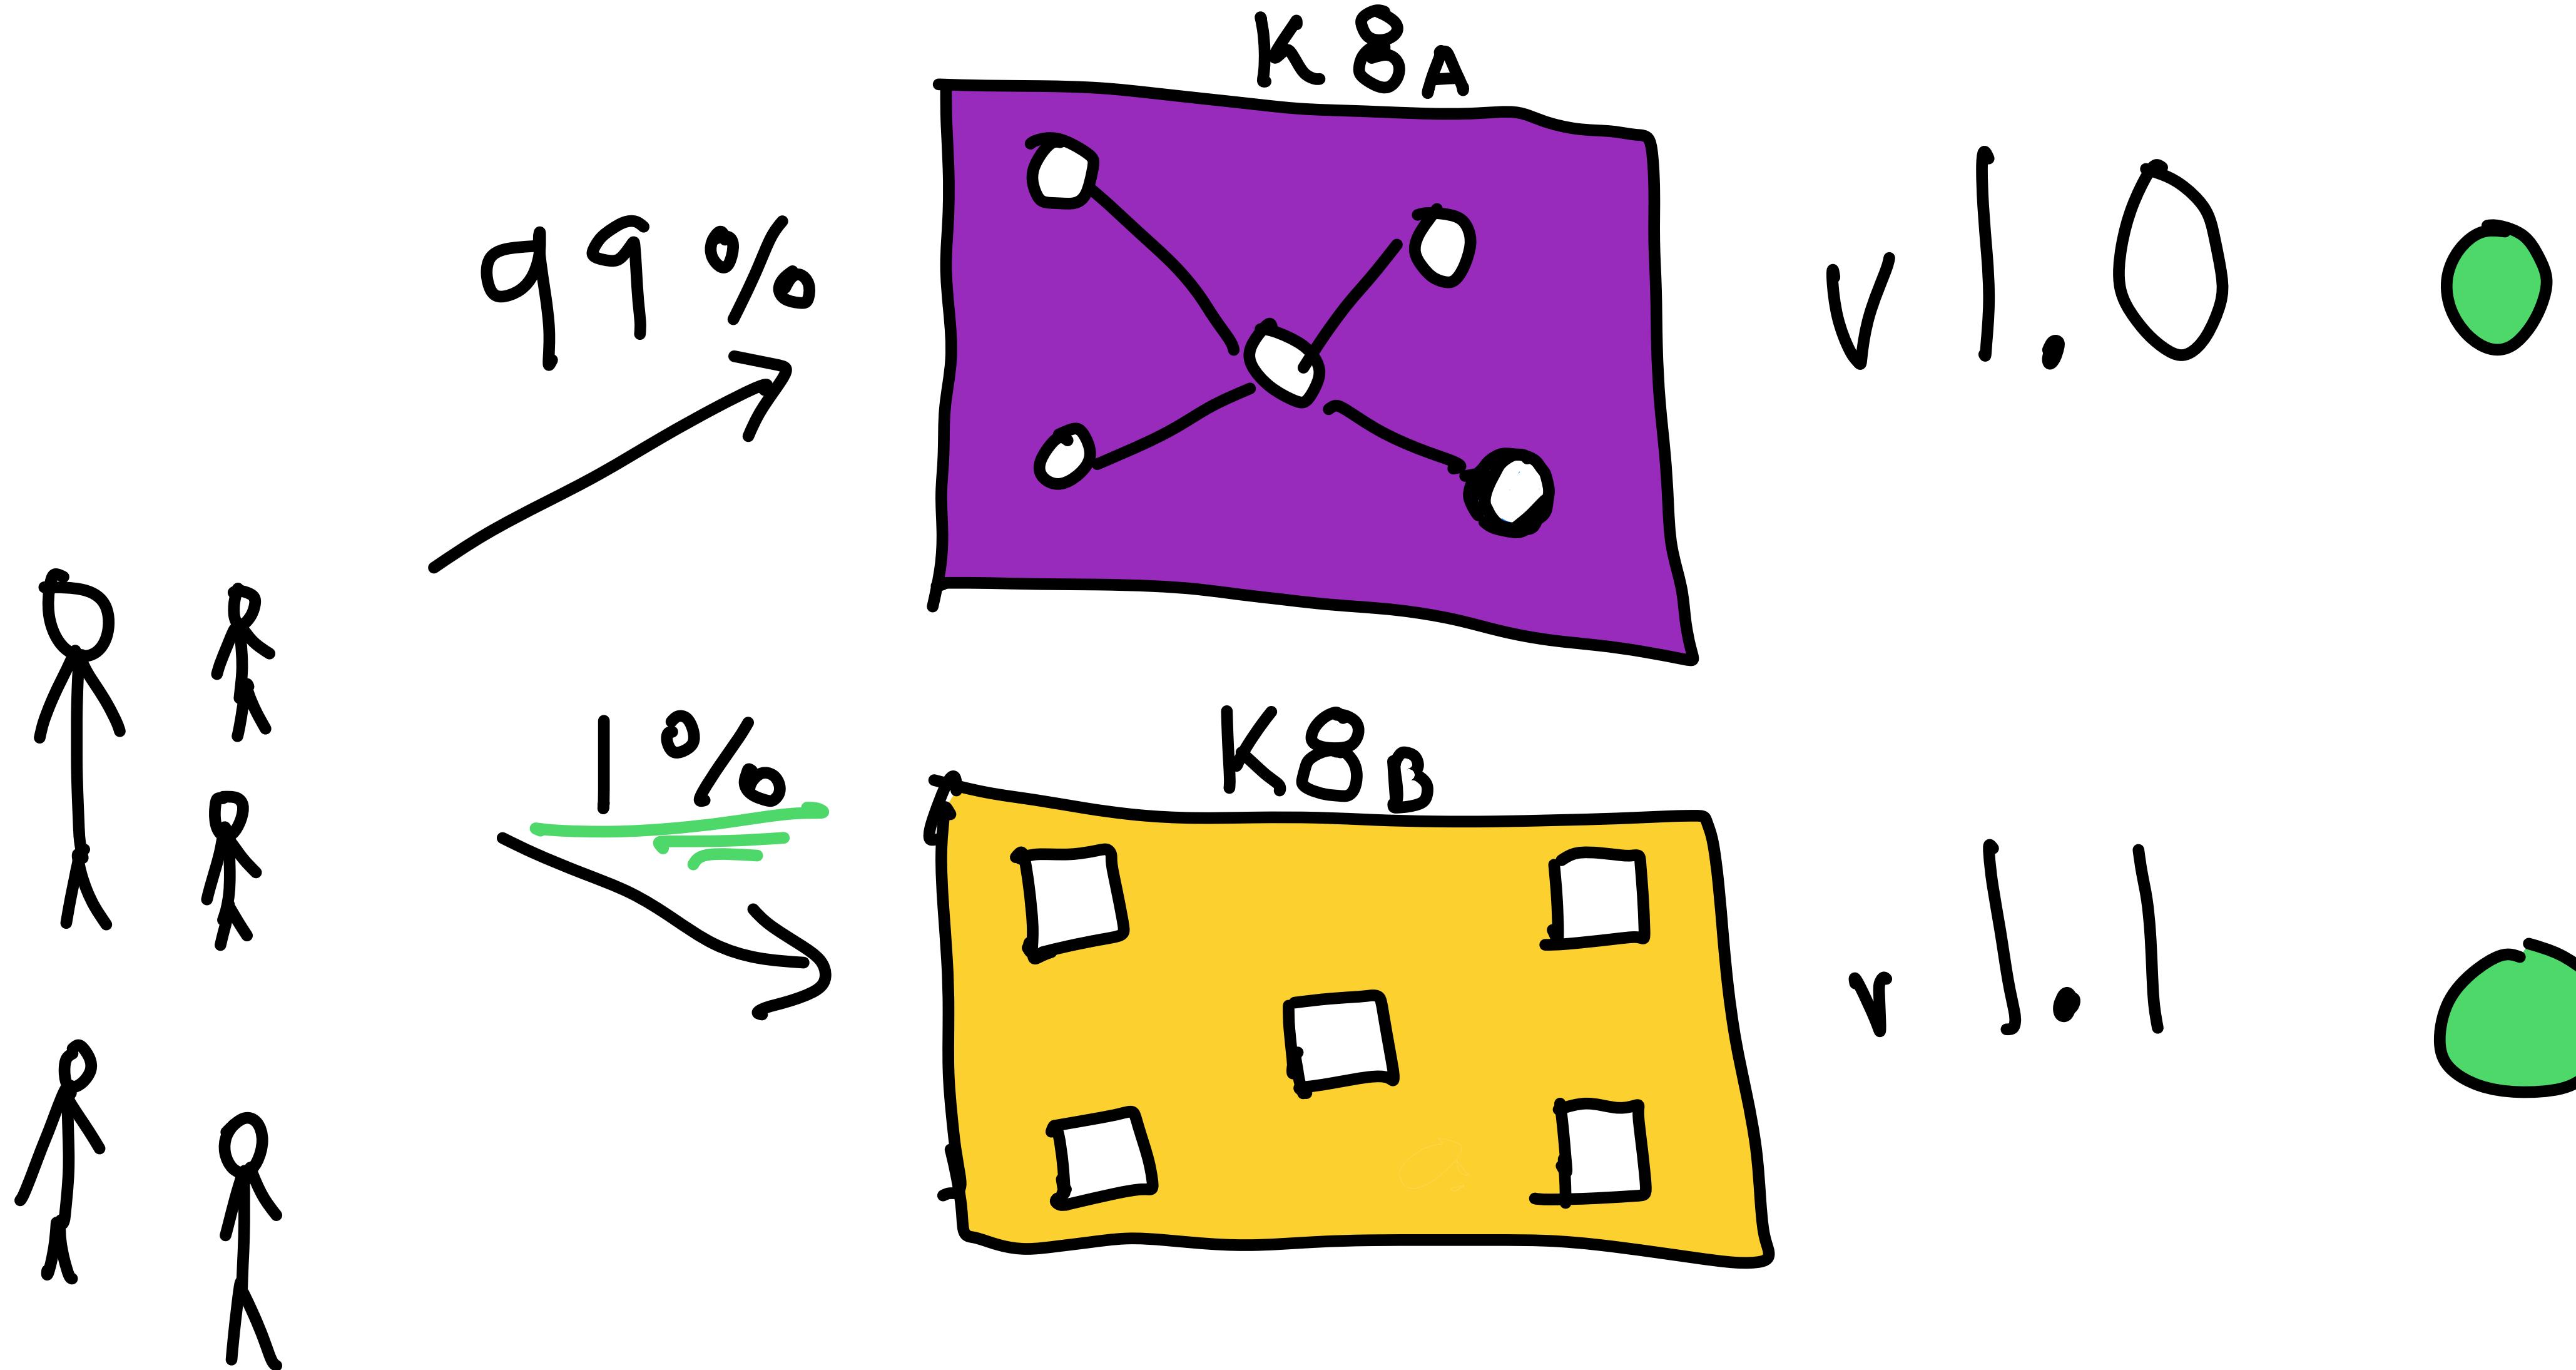
Hard

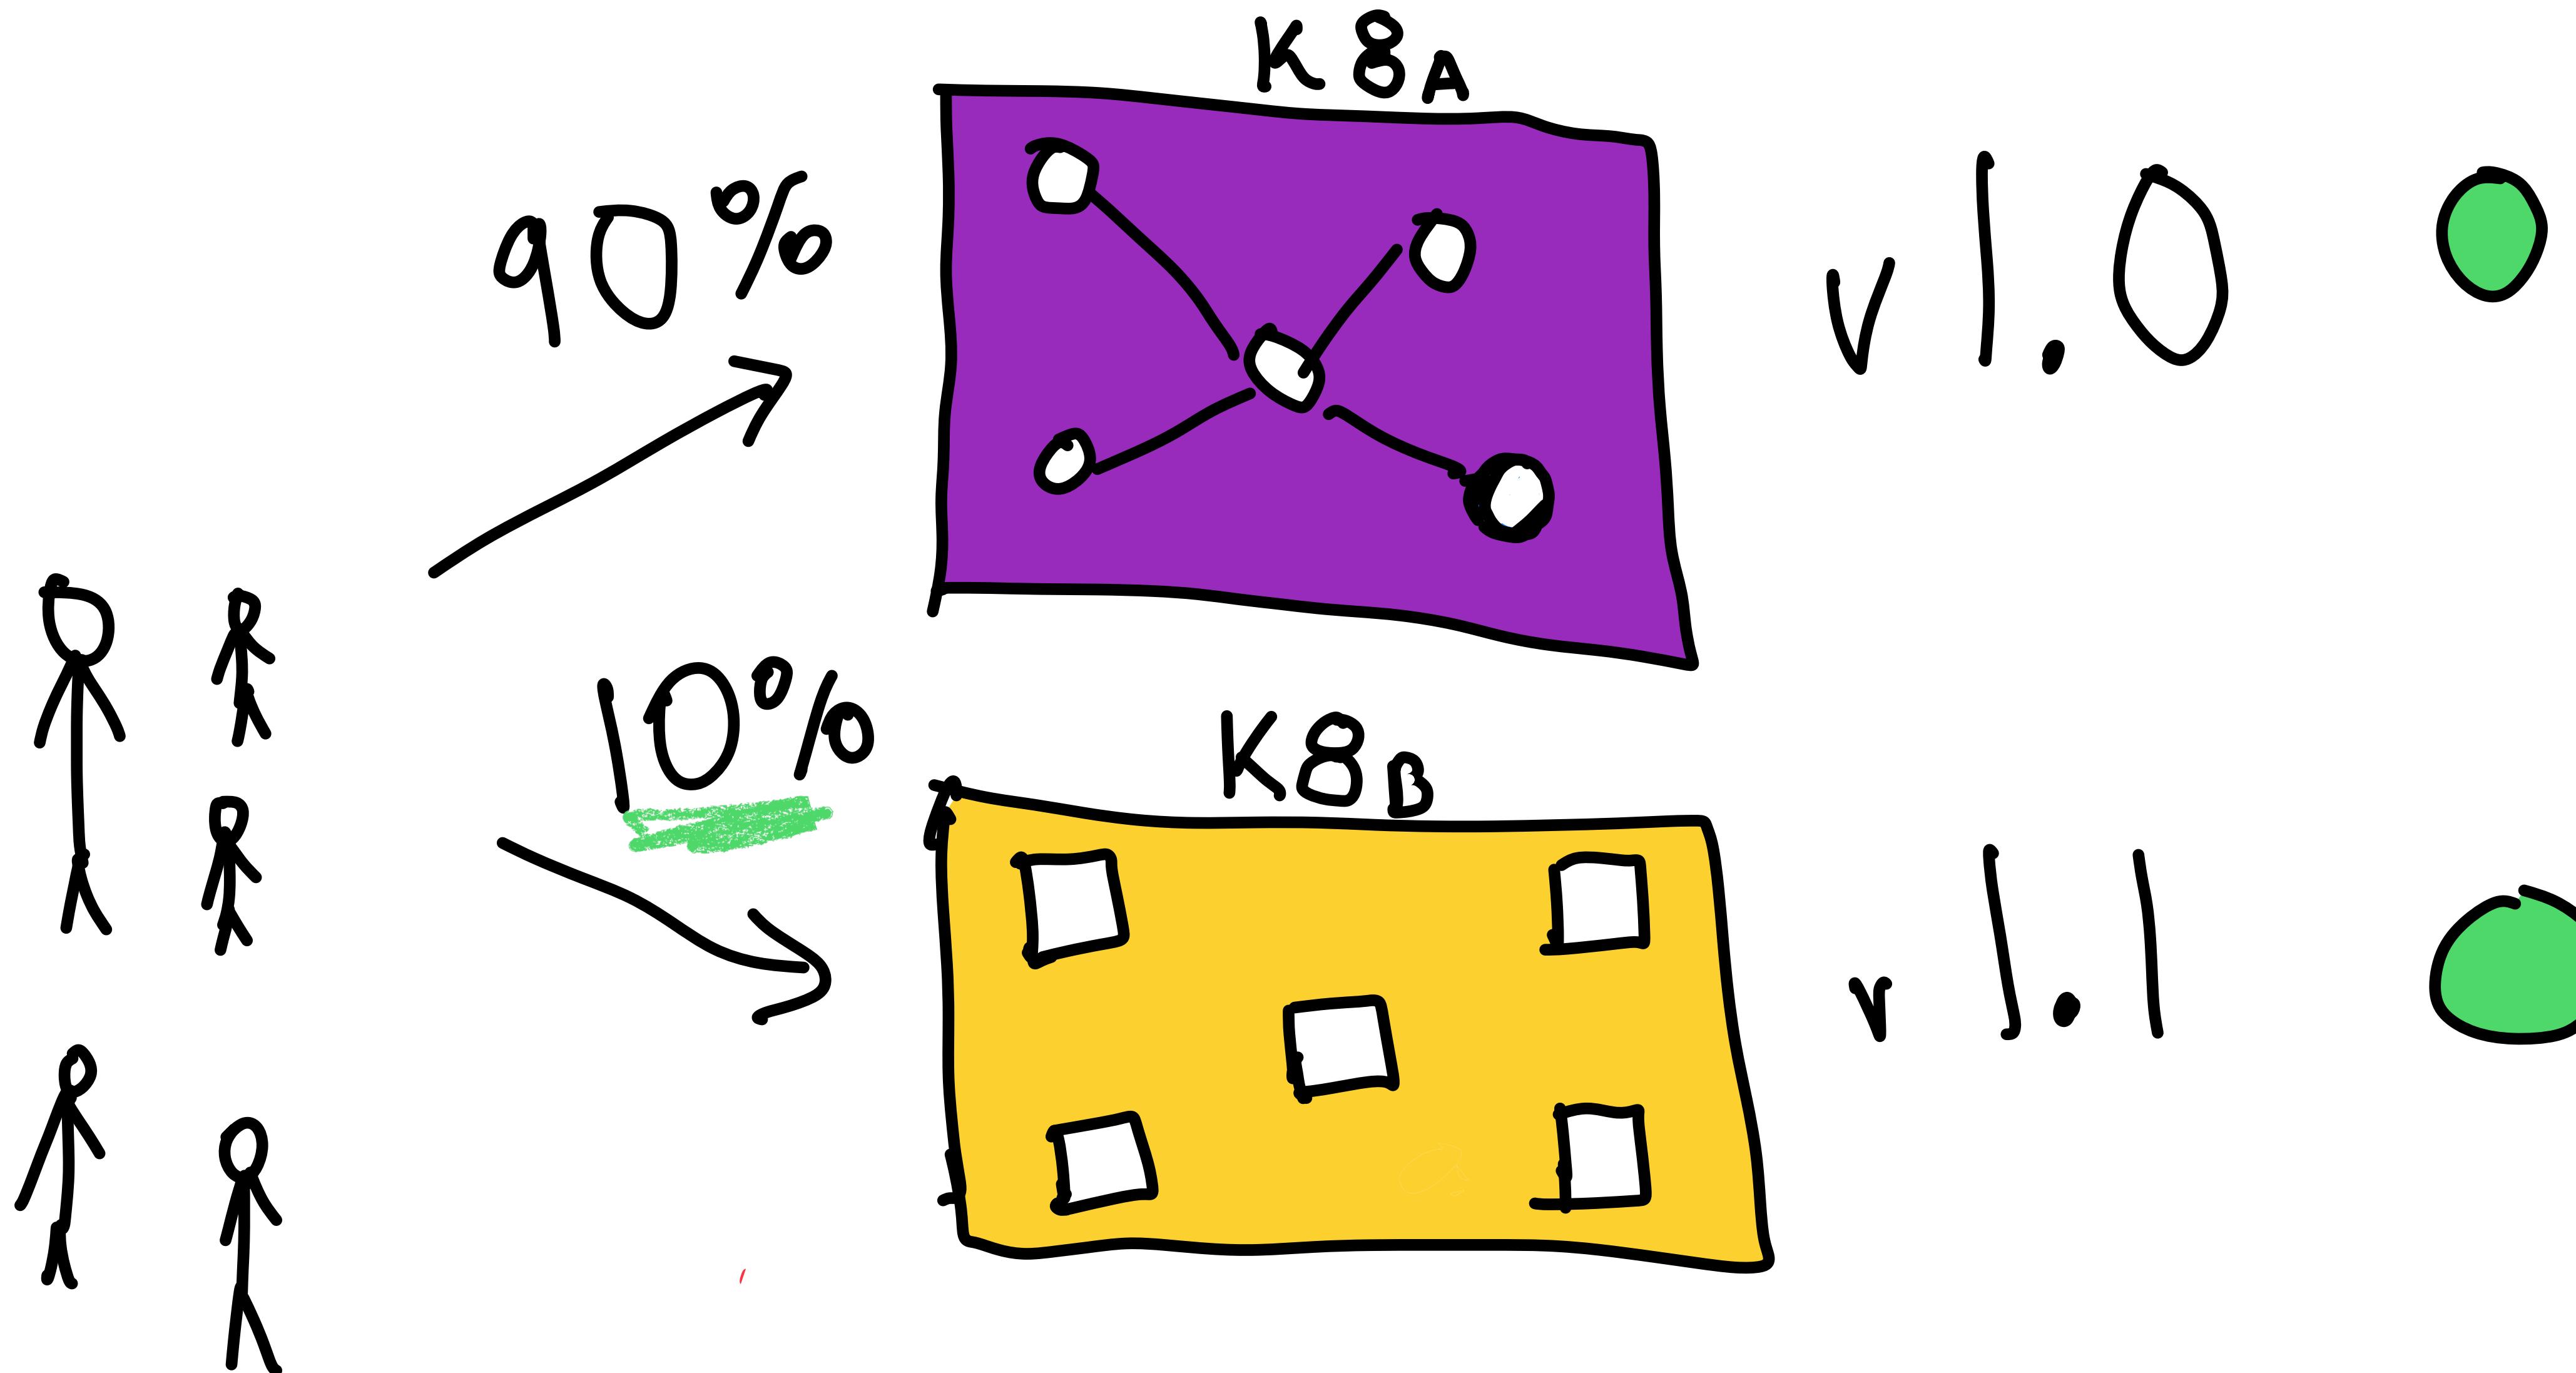


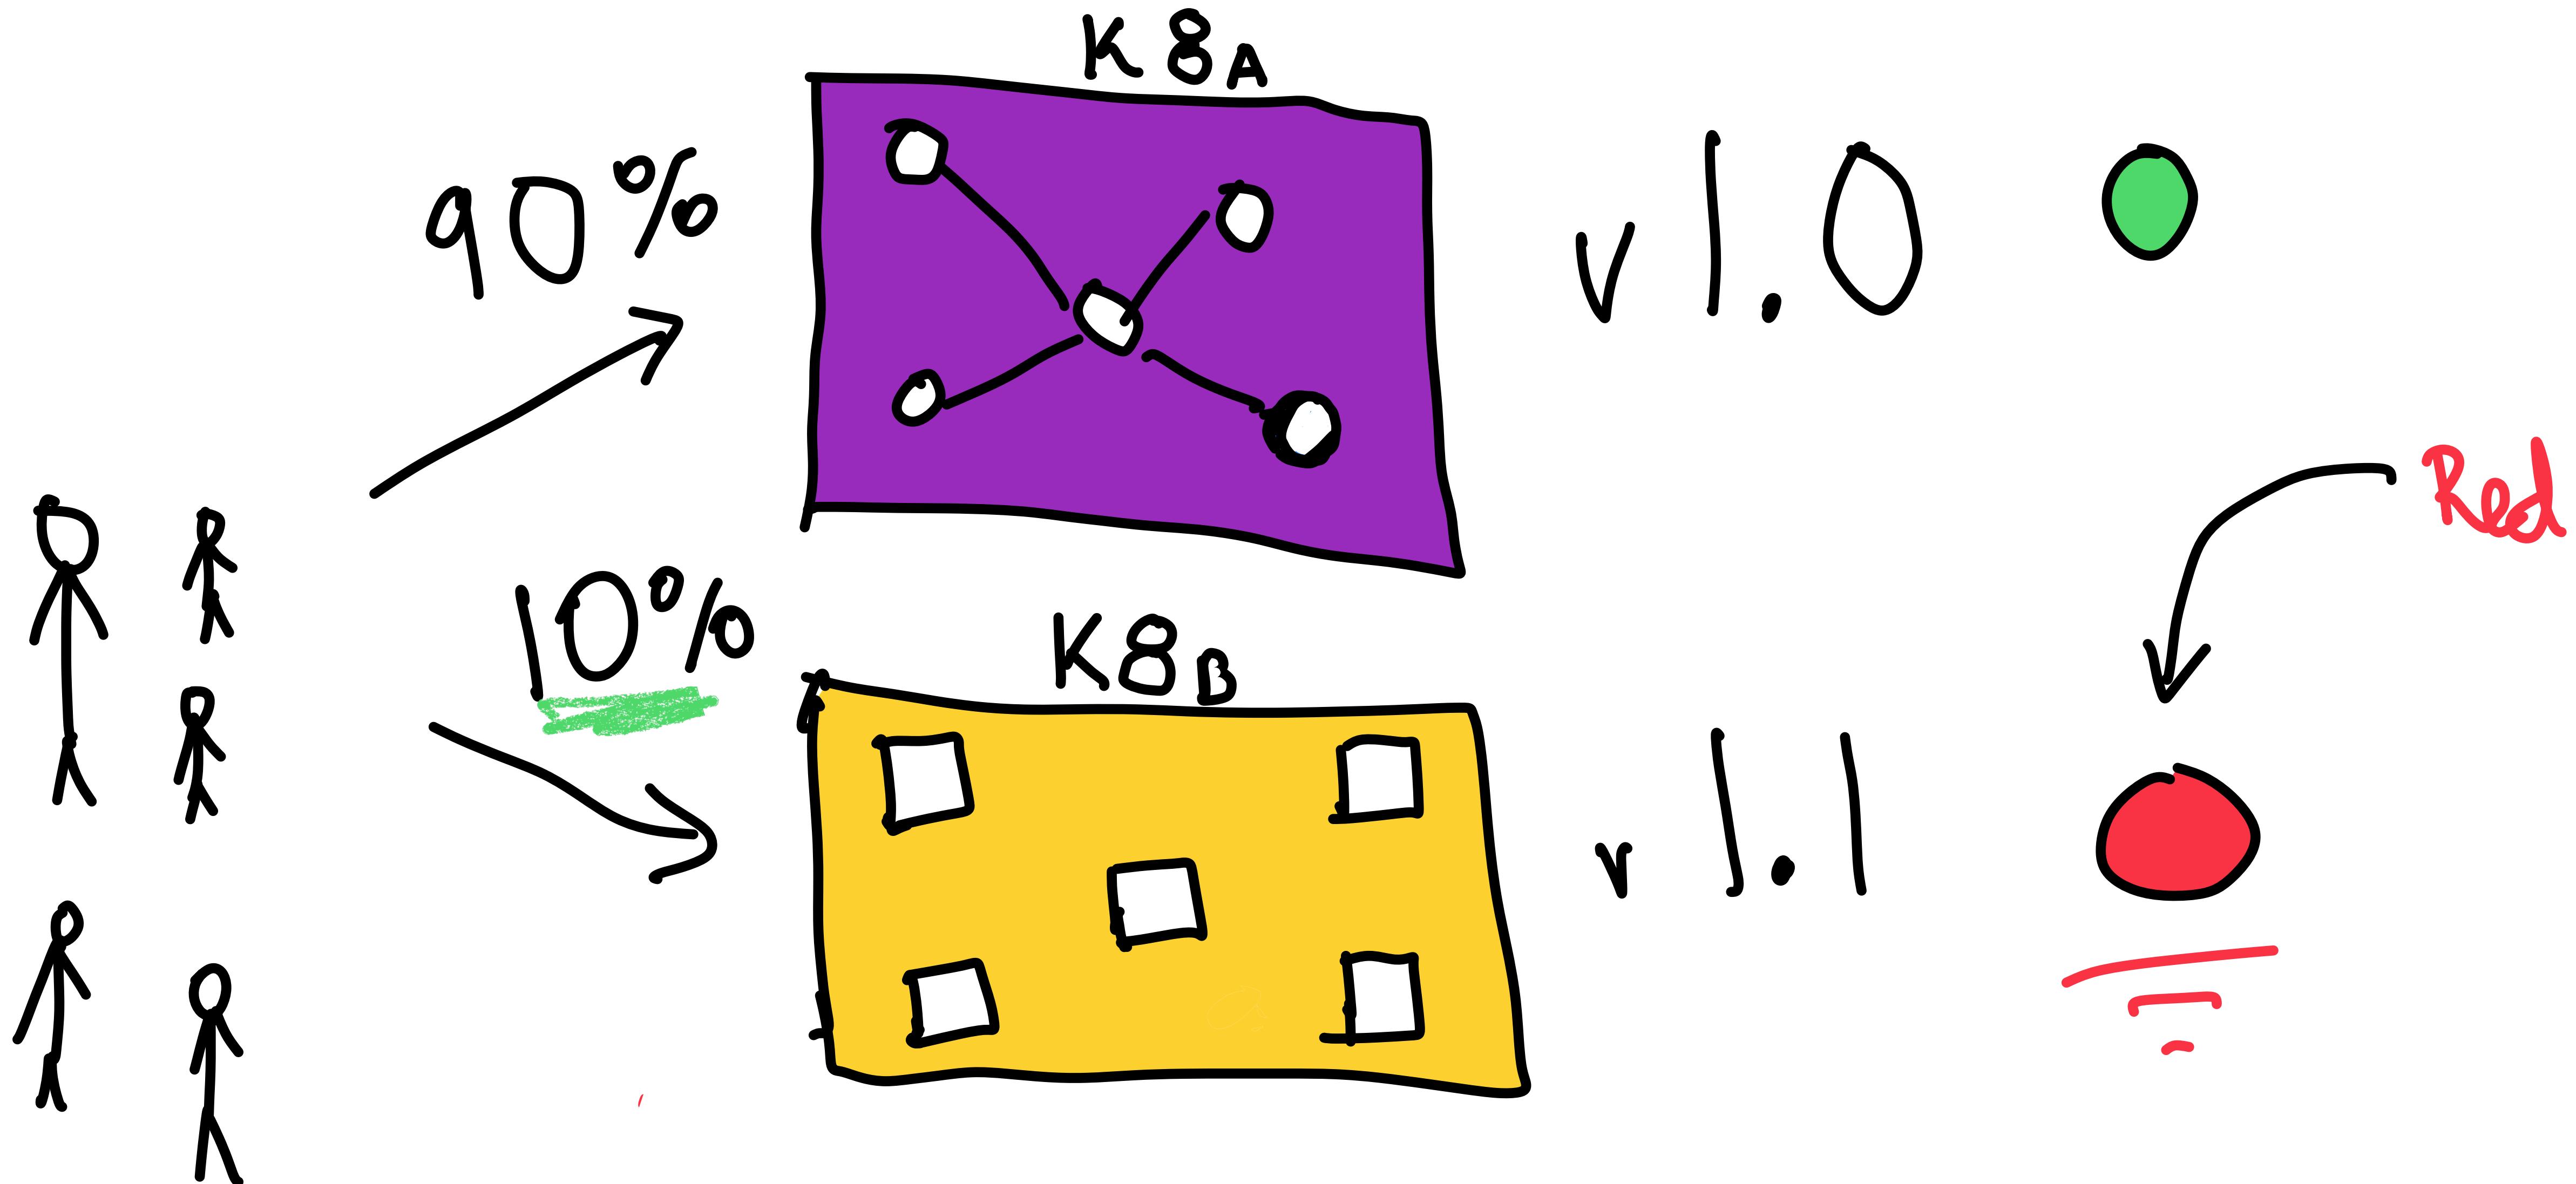
This is expensive and requires fine-tuning

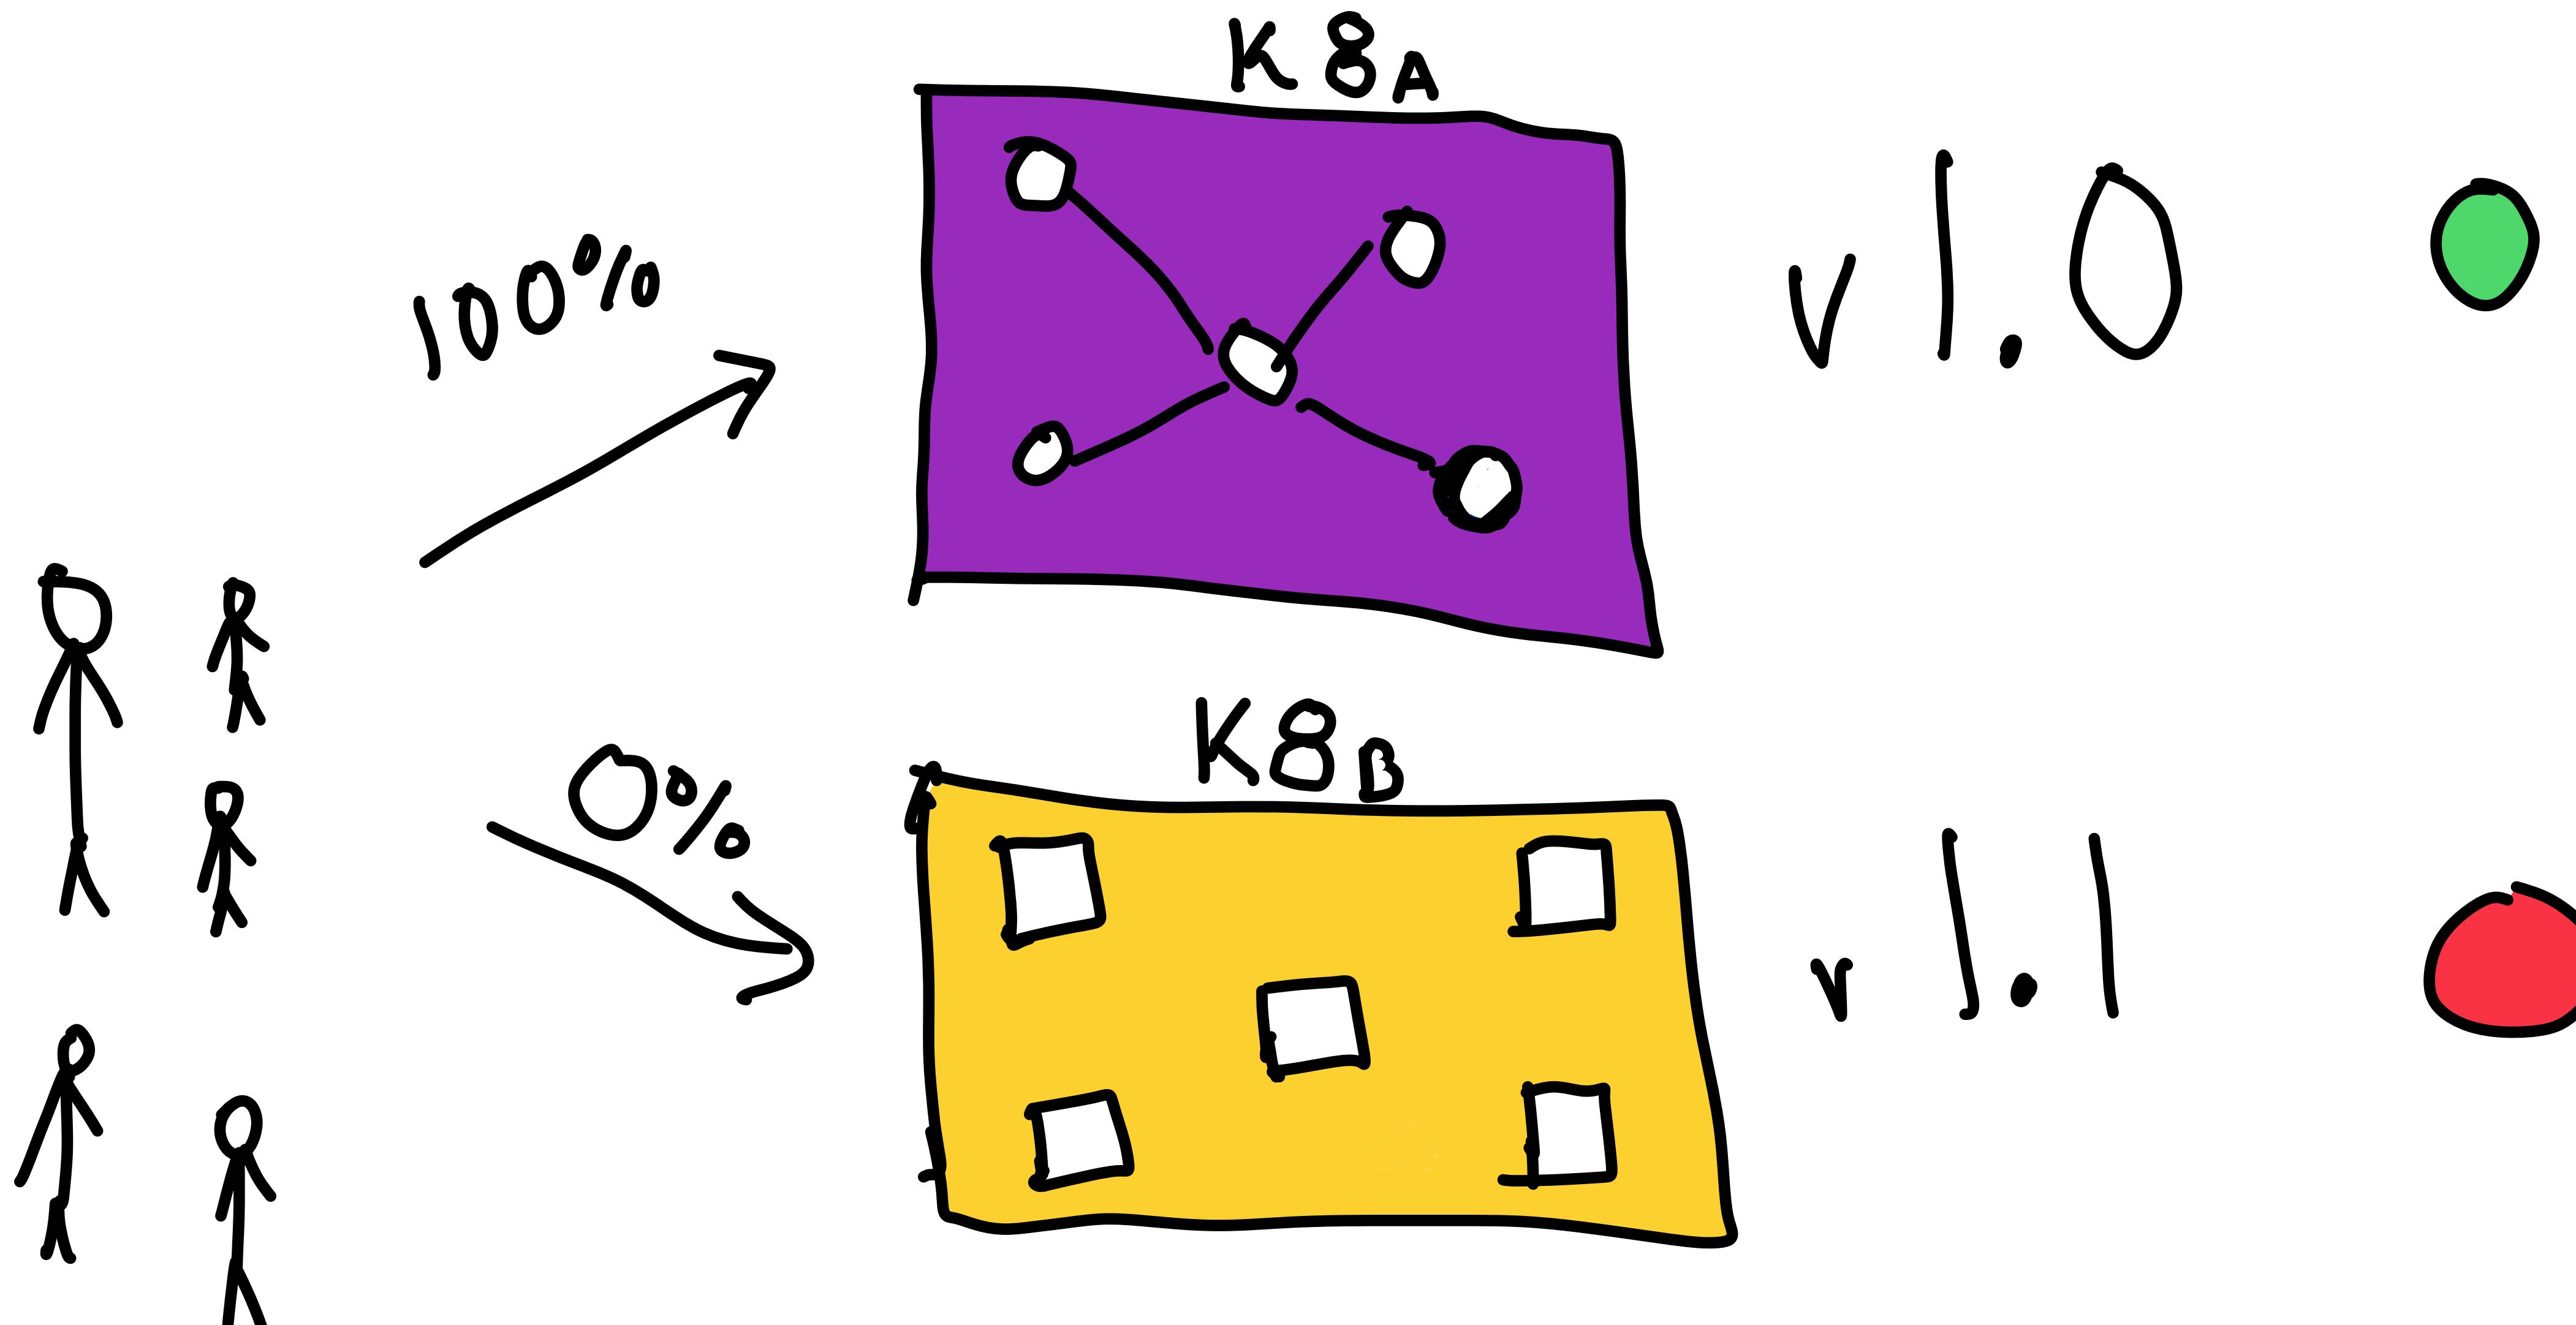
Green-Blue Deployments



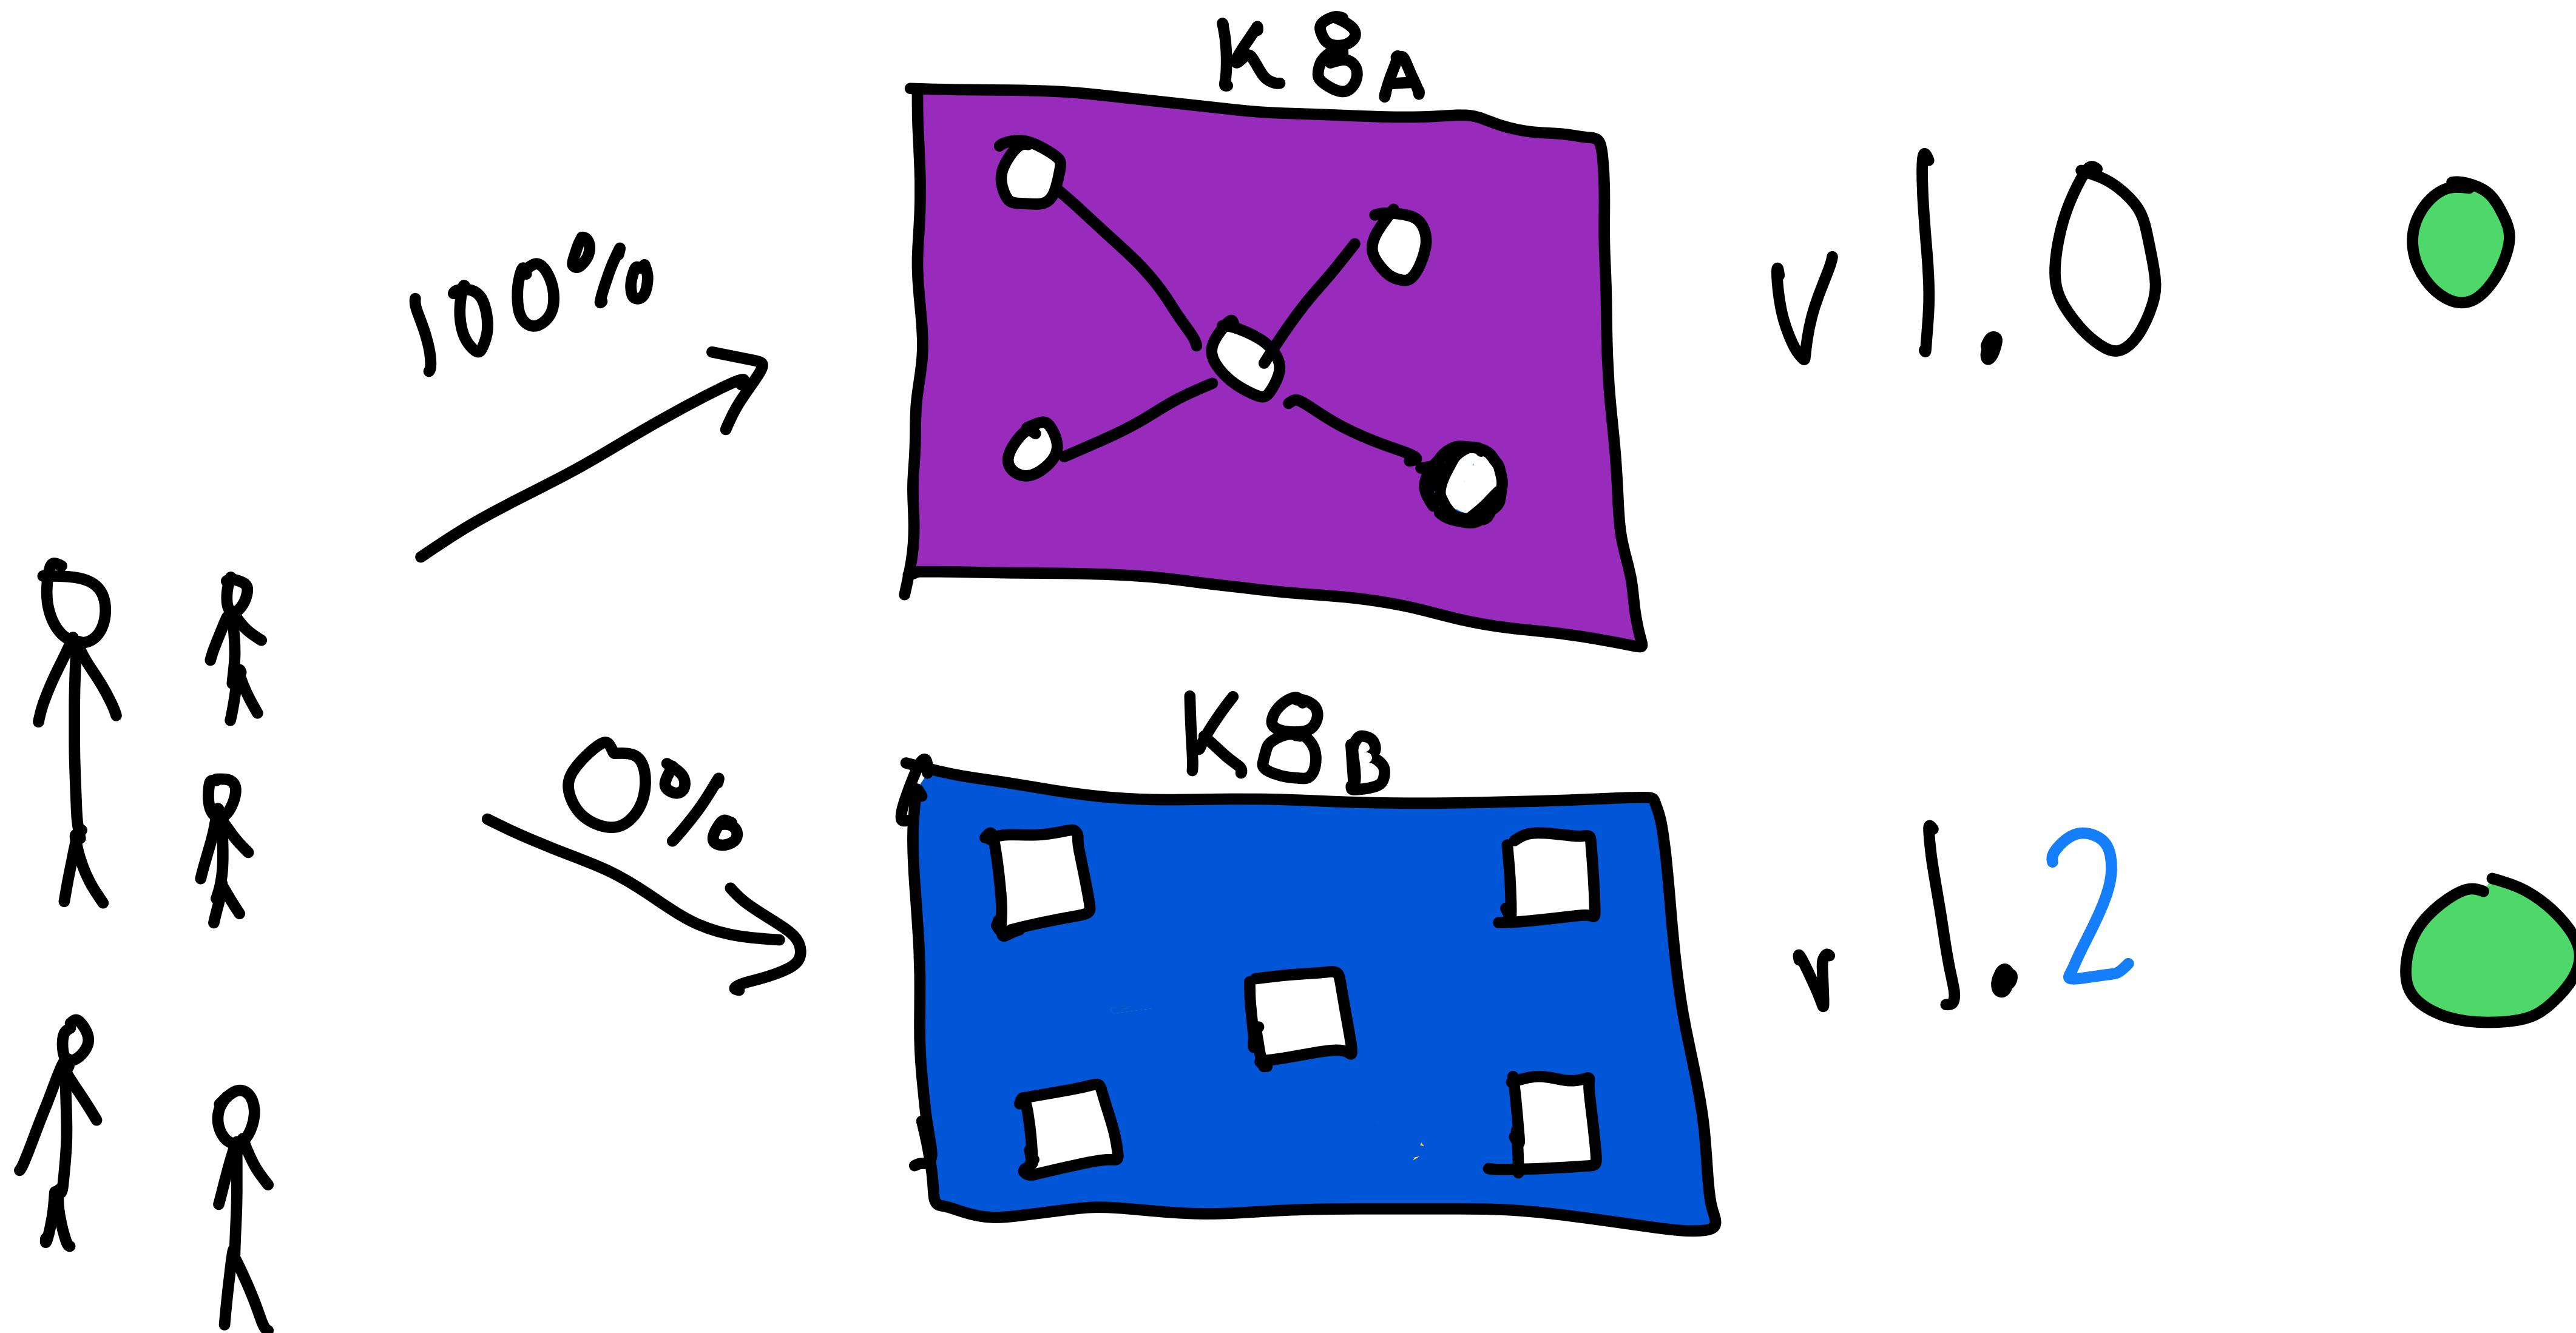


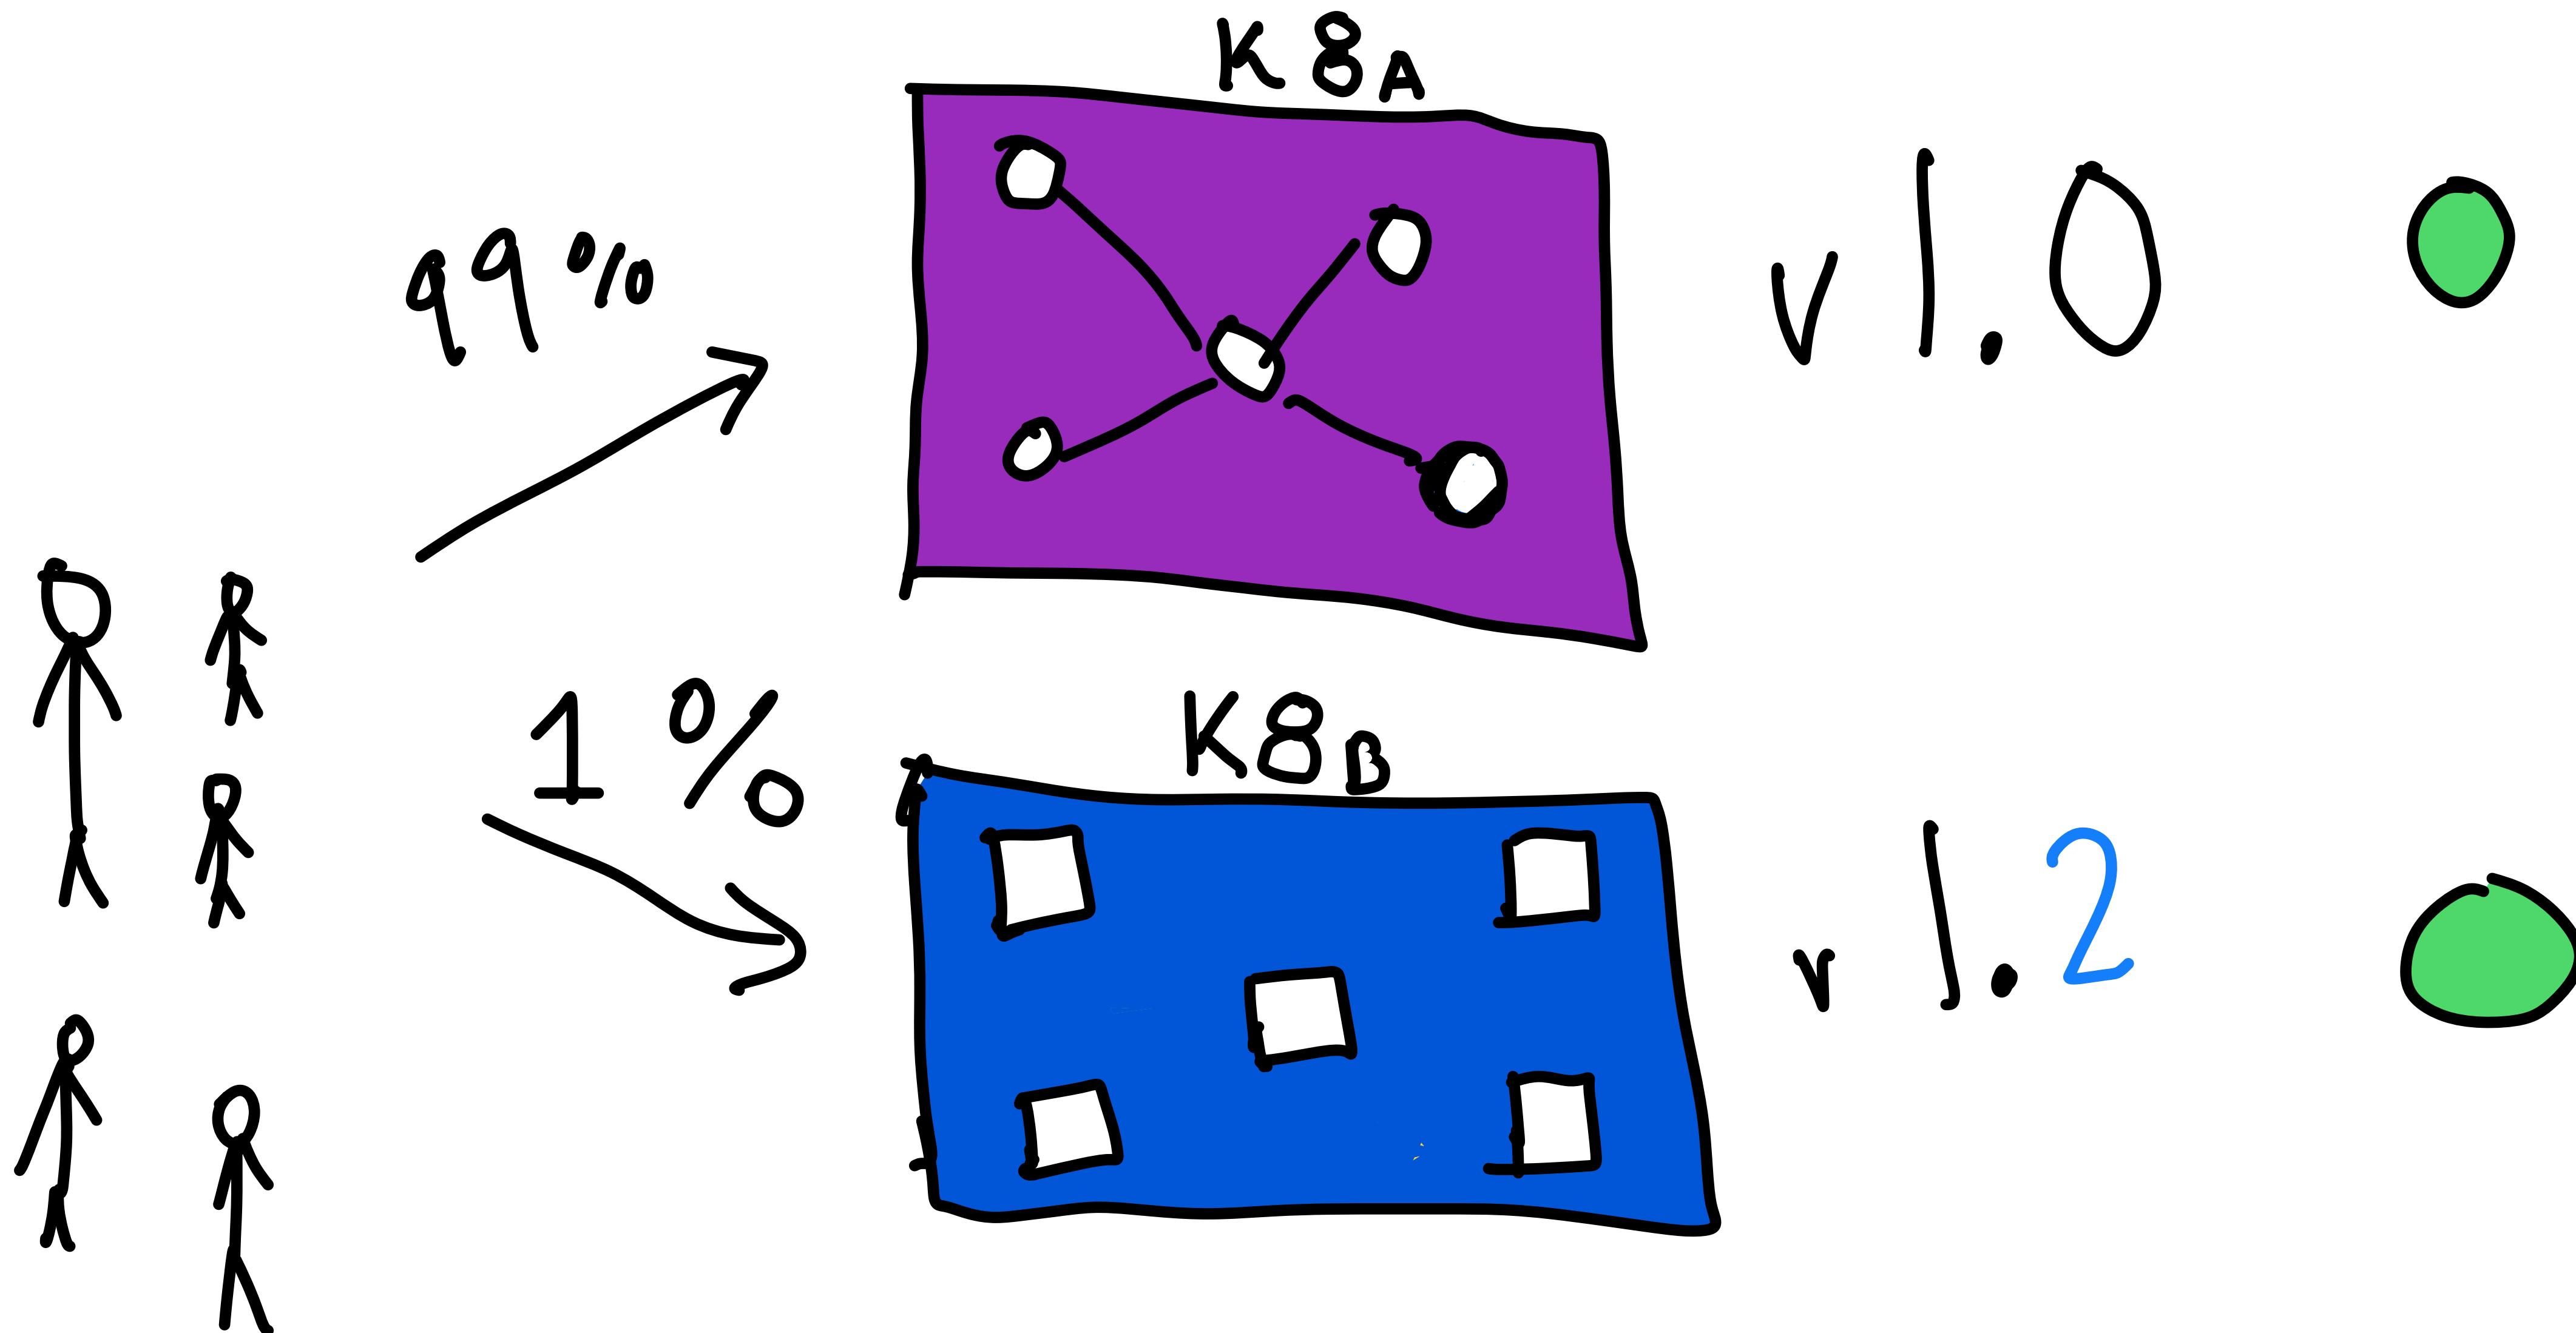


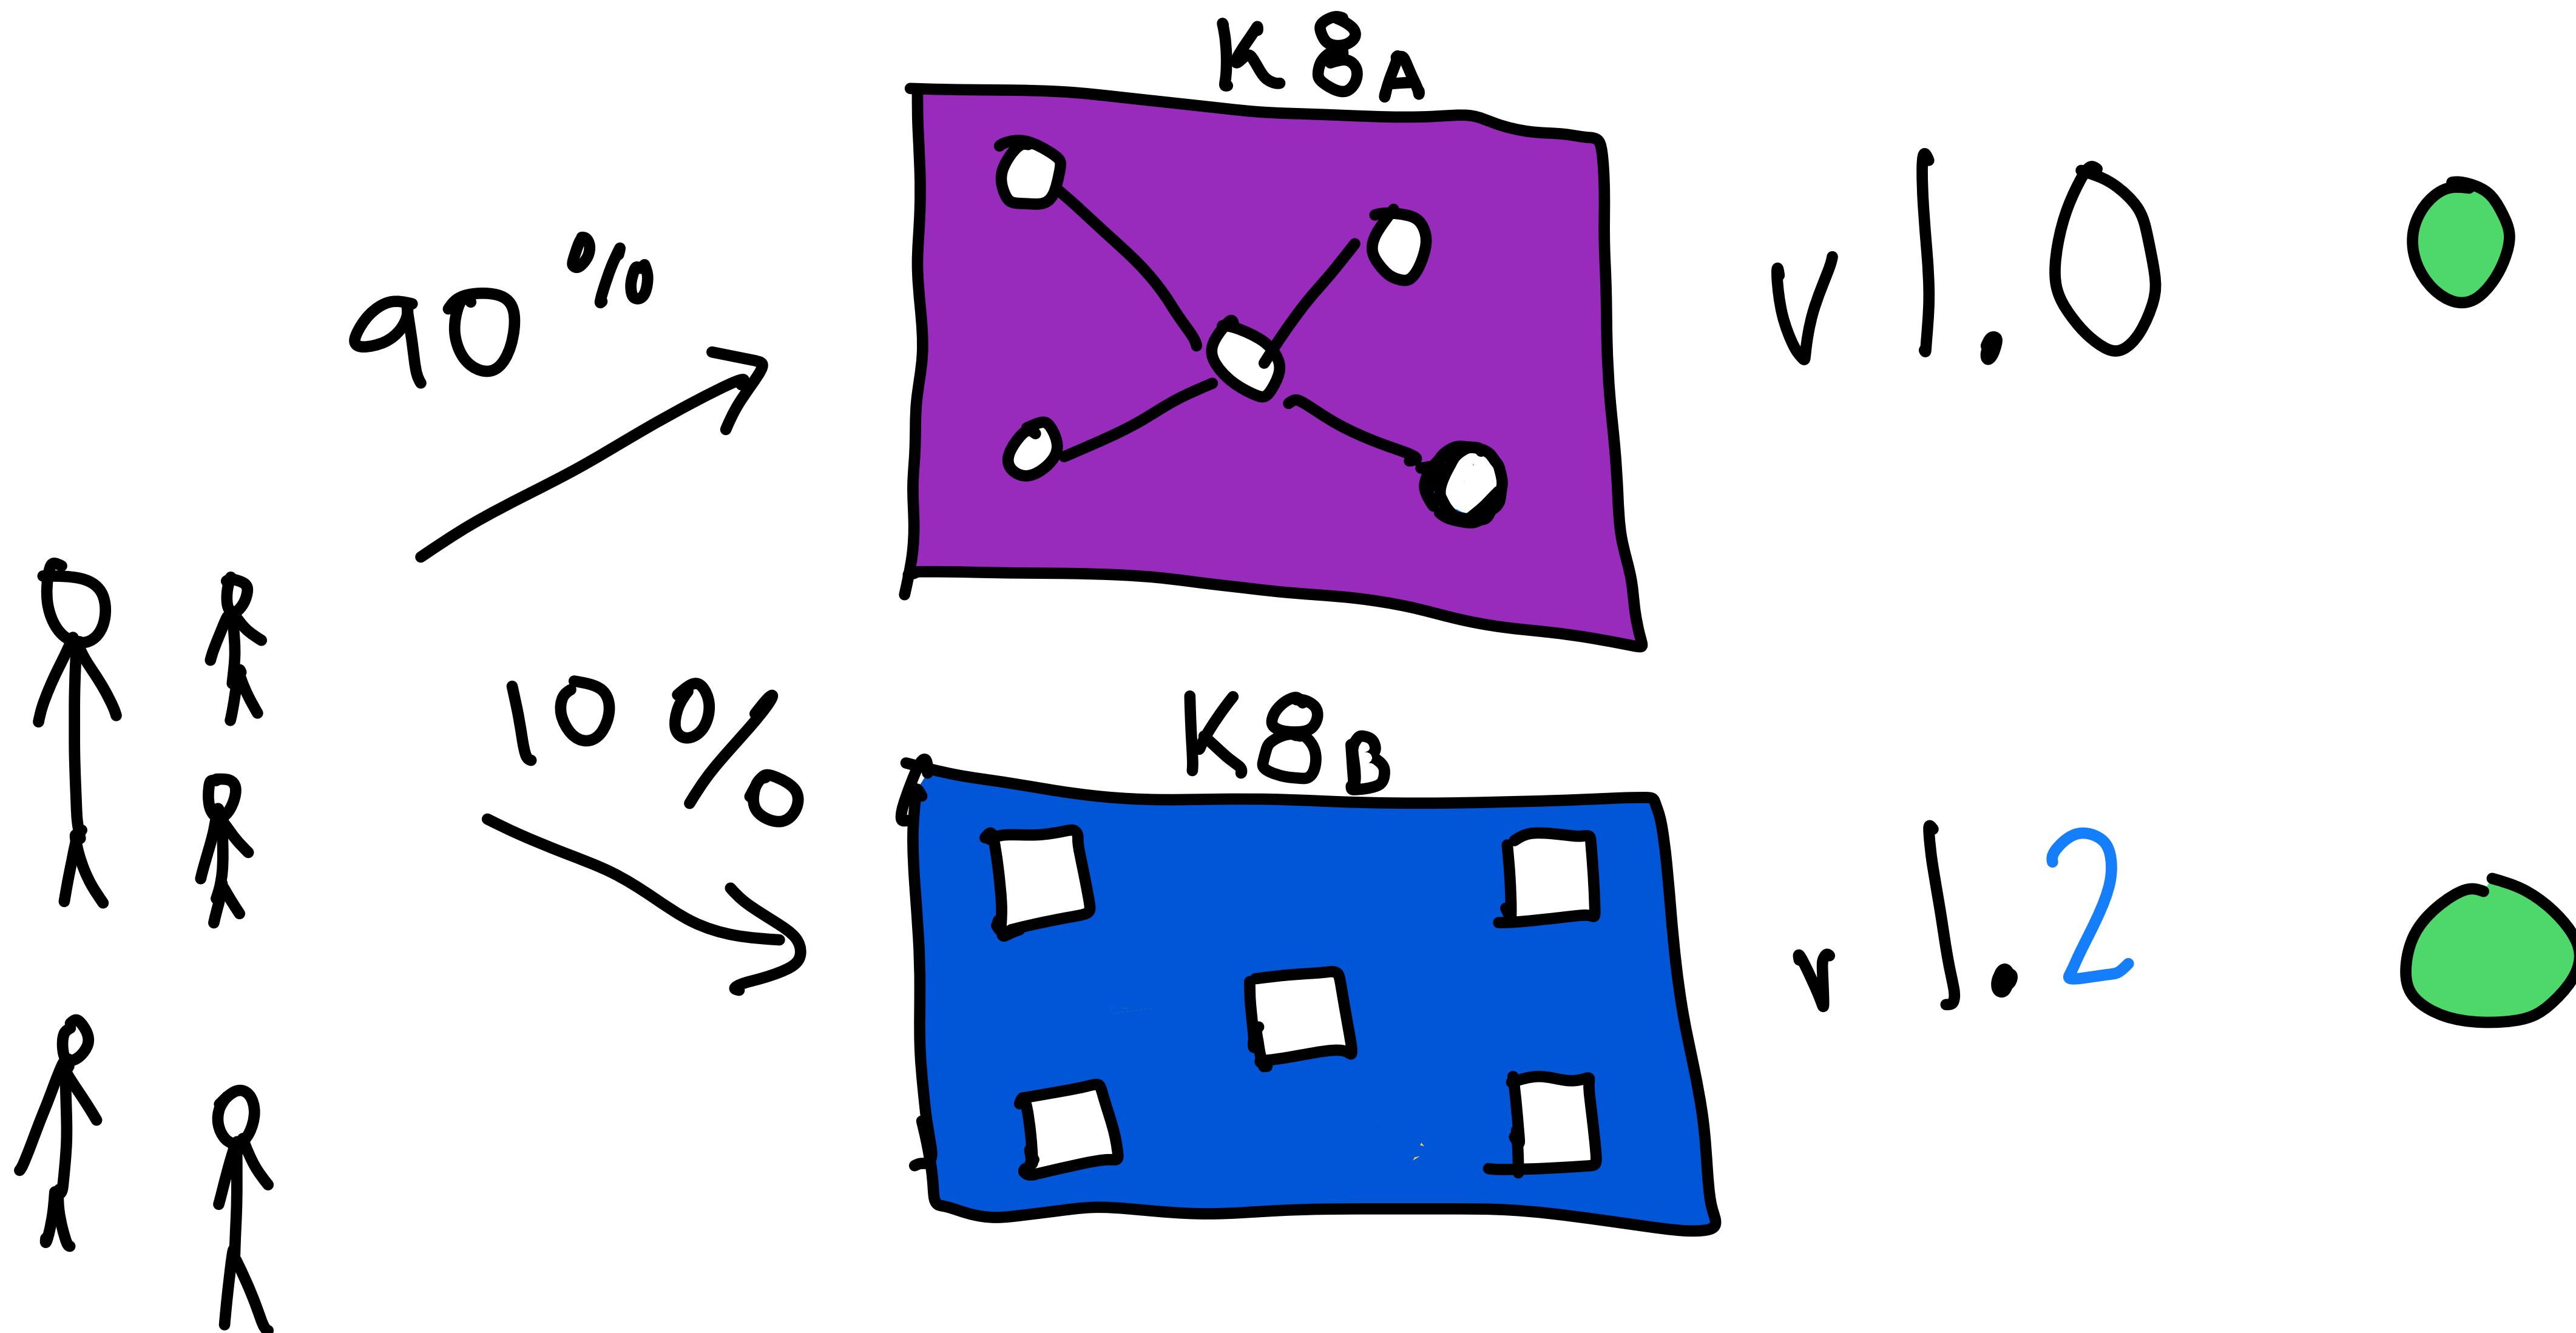


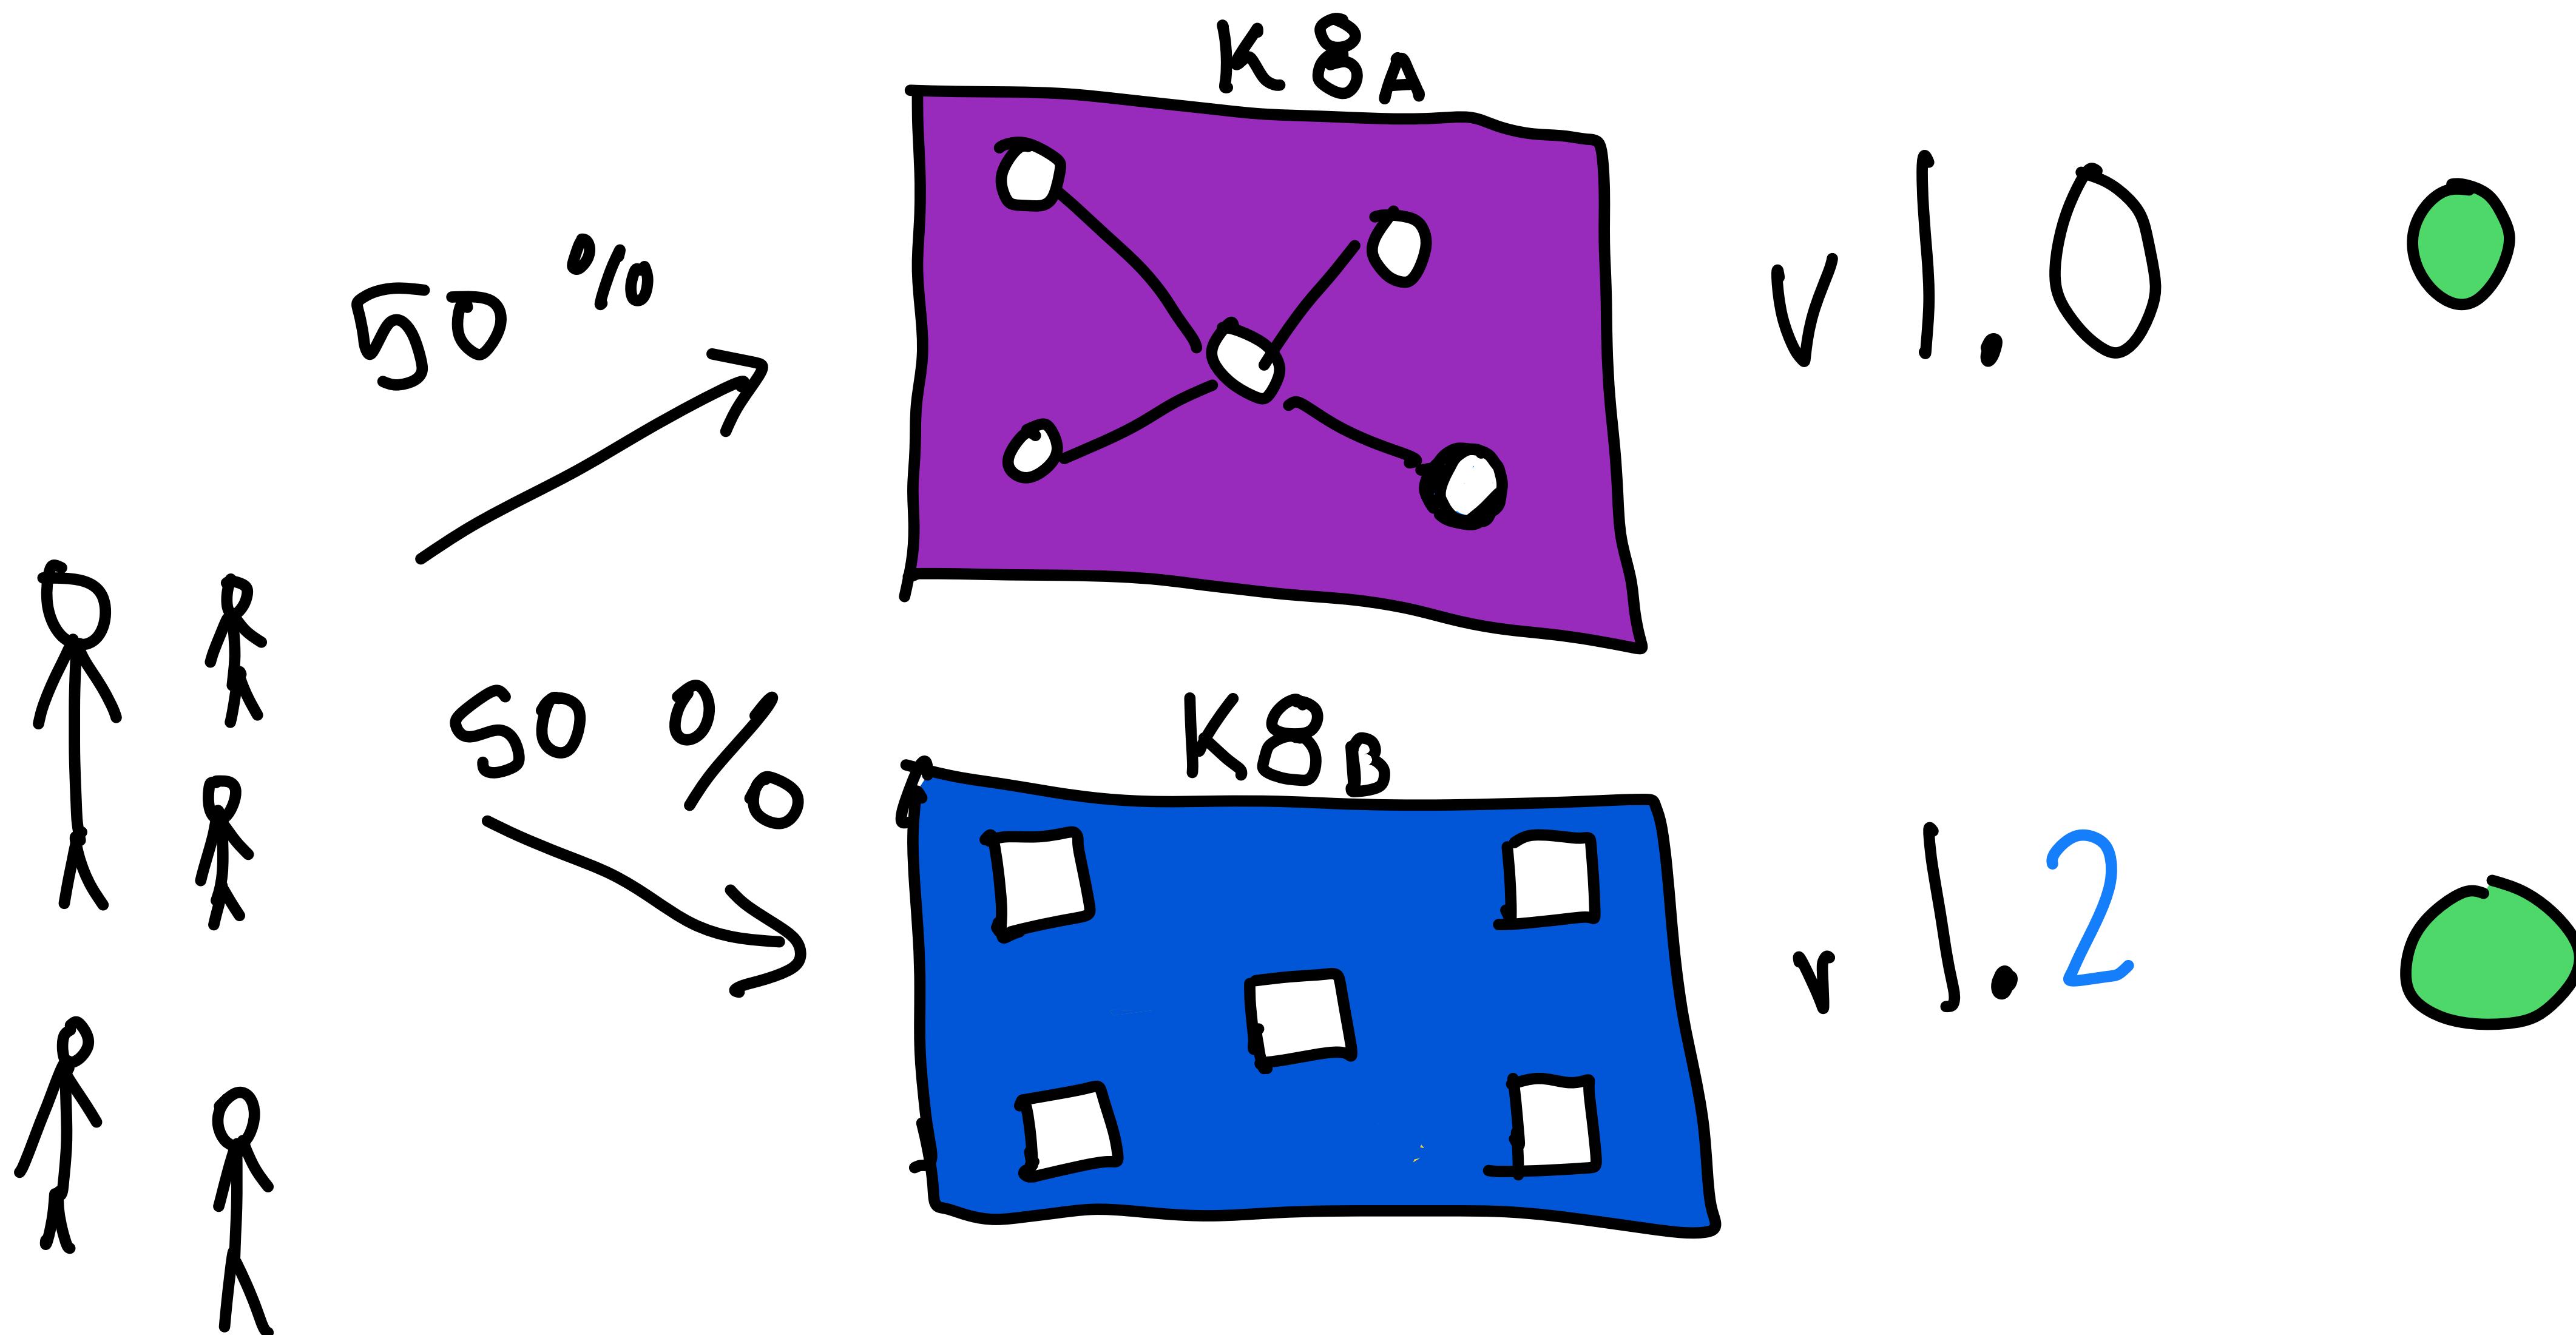


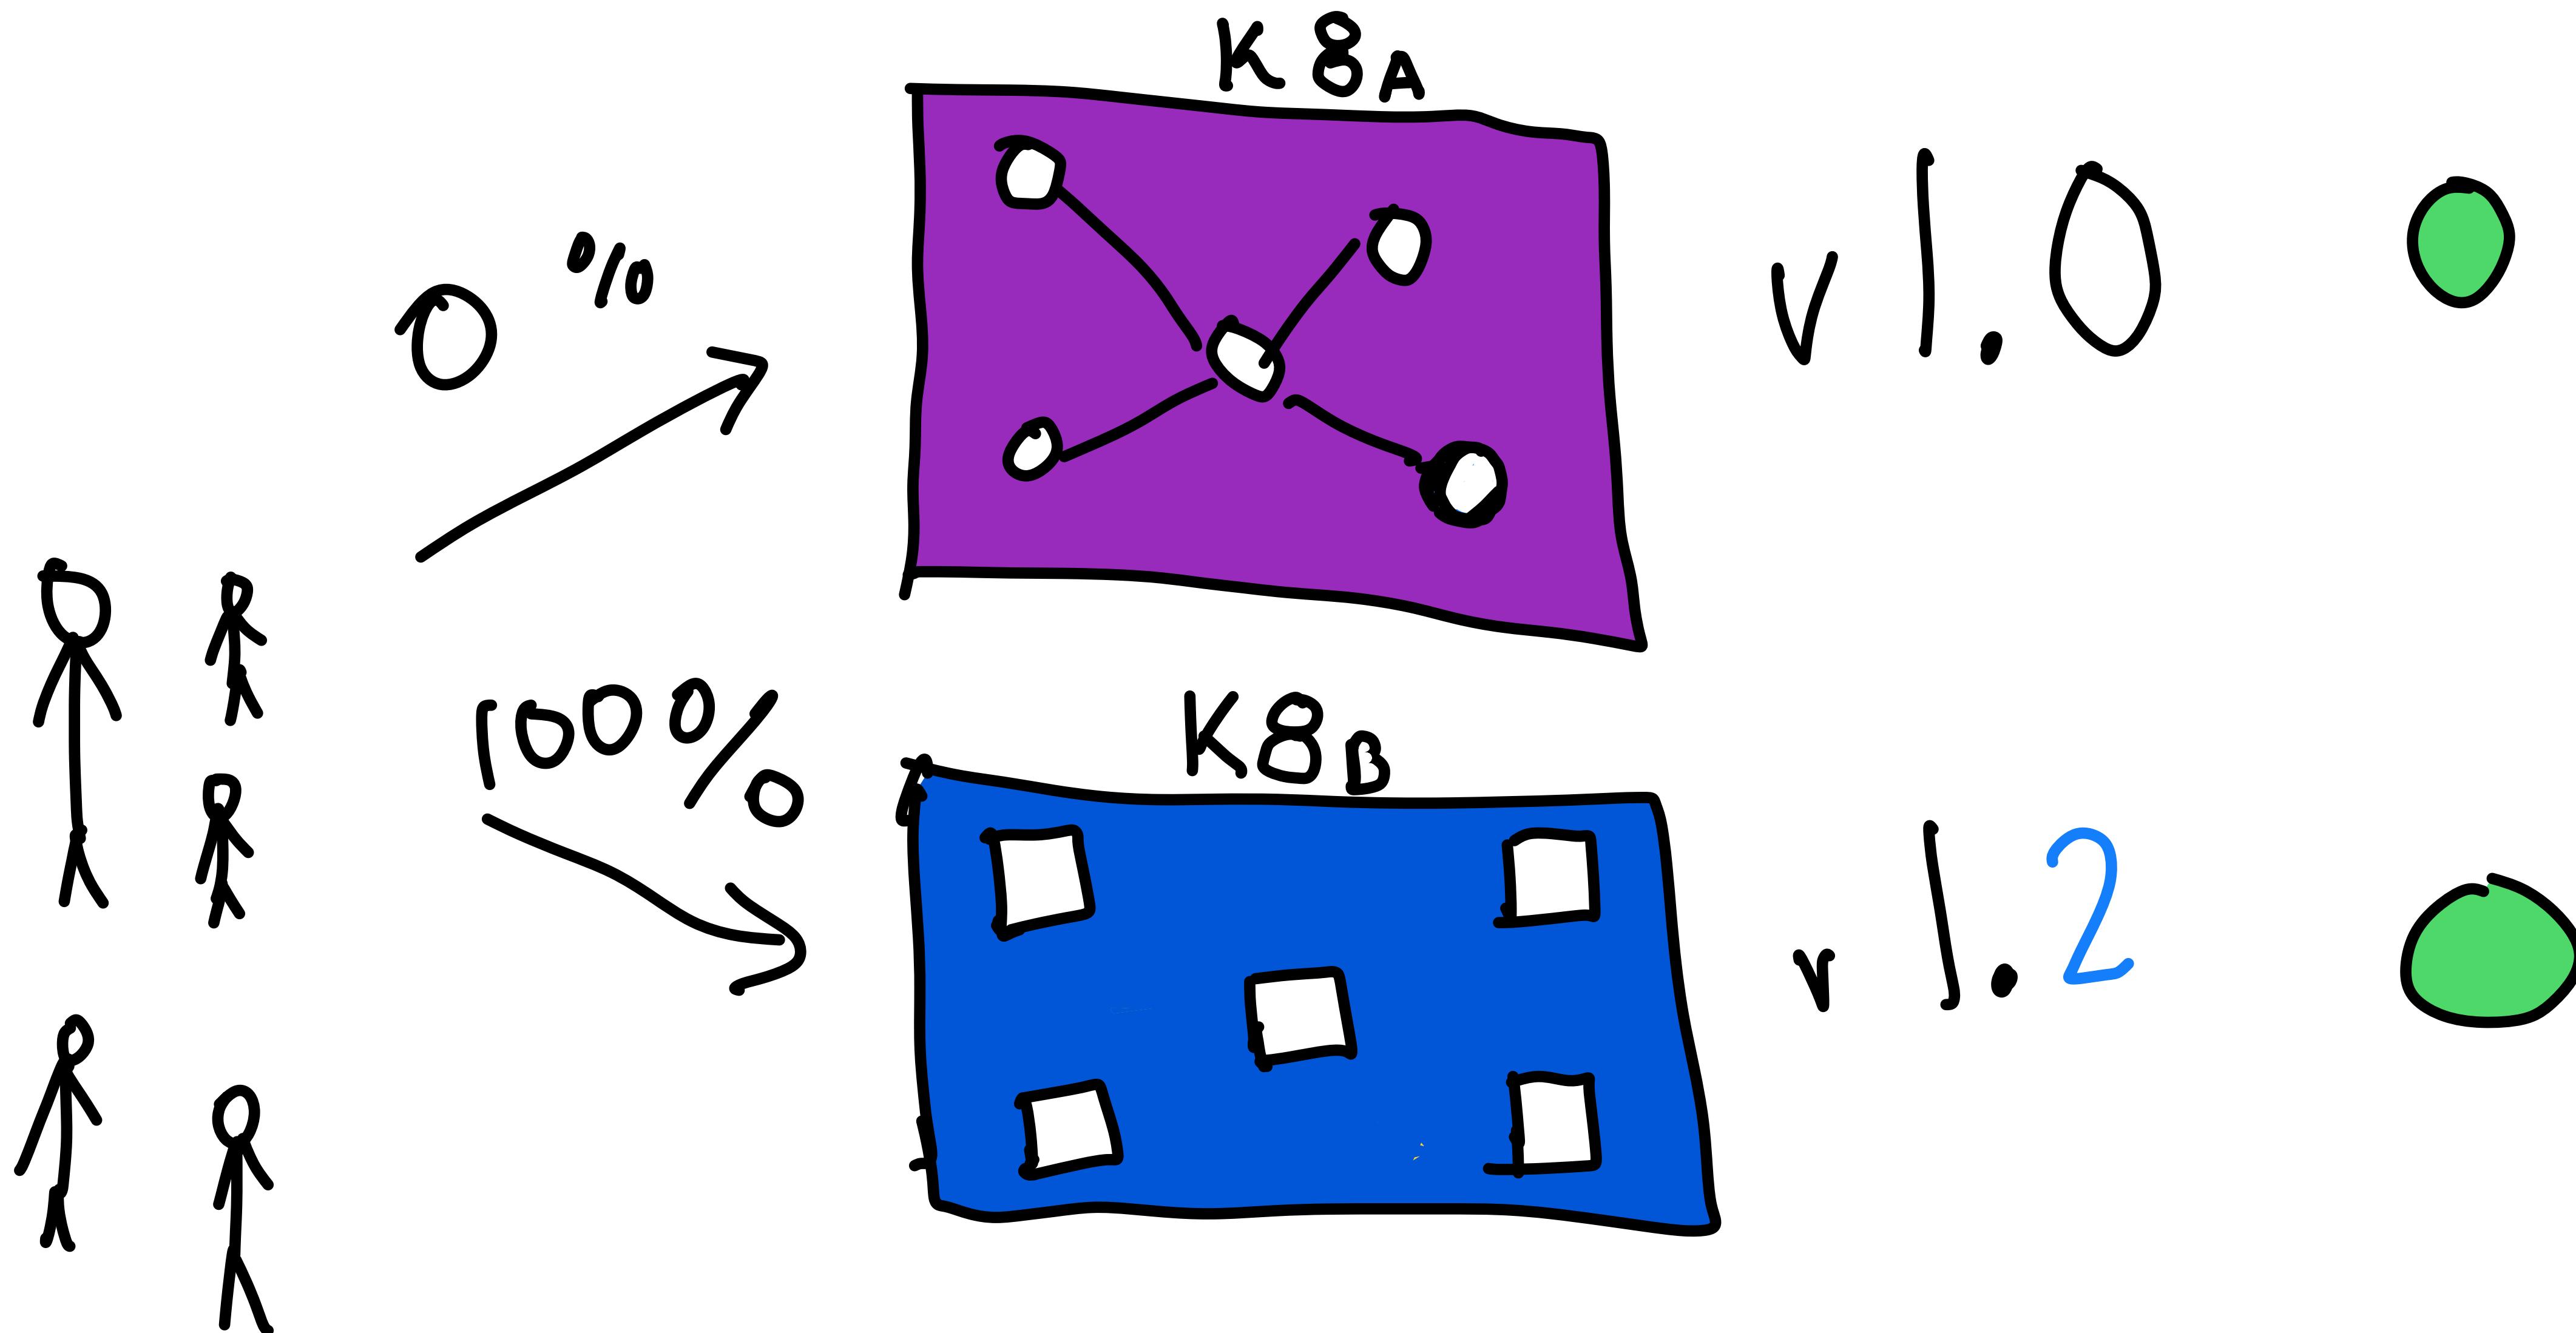
Rollback!



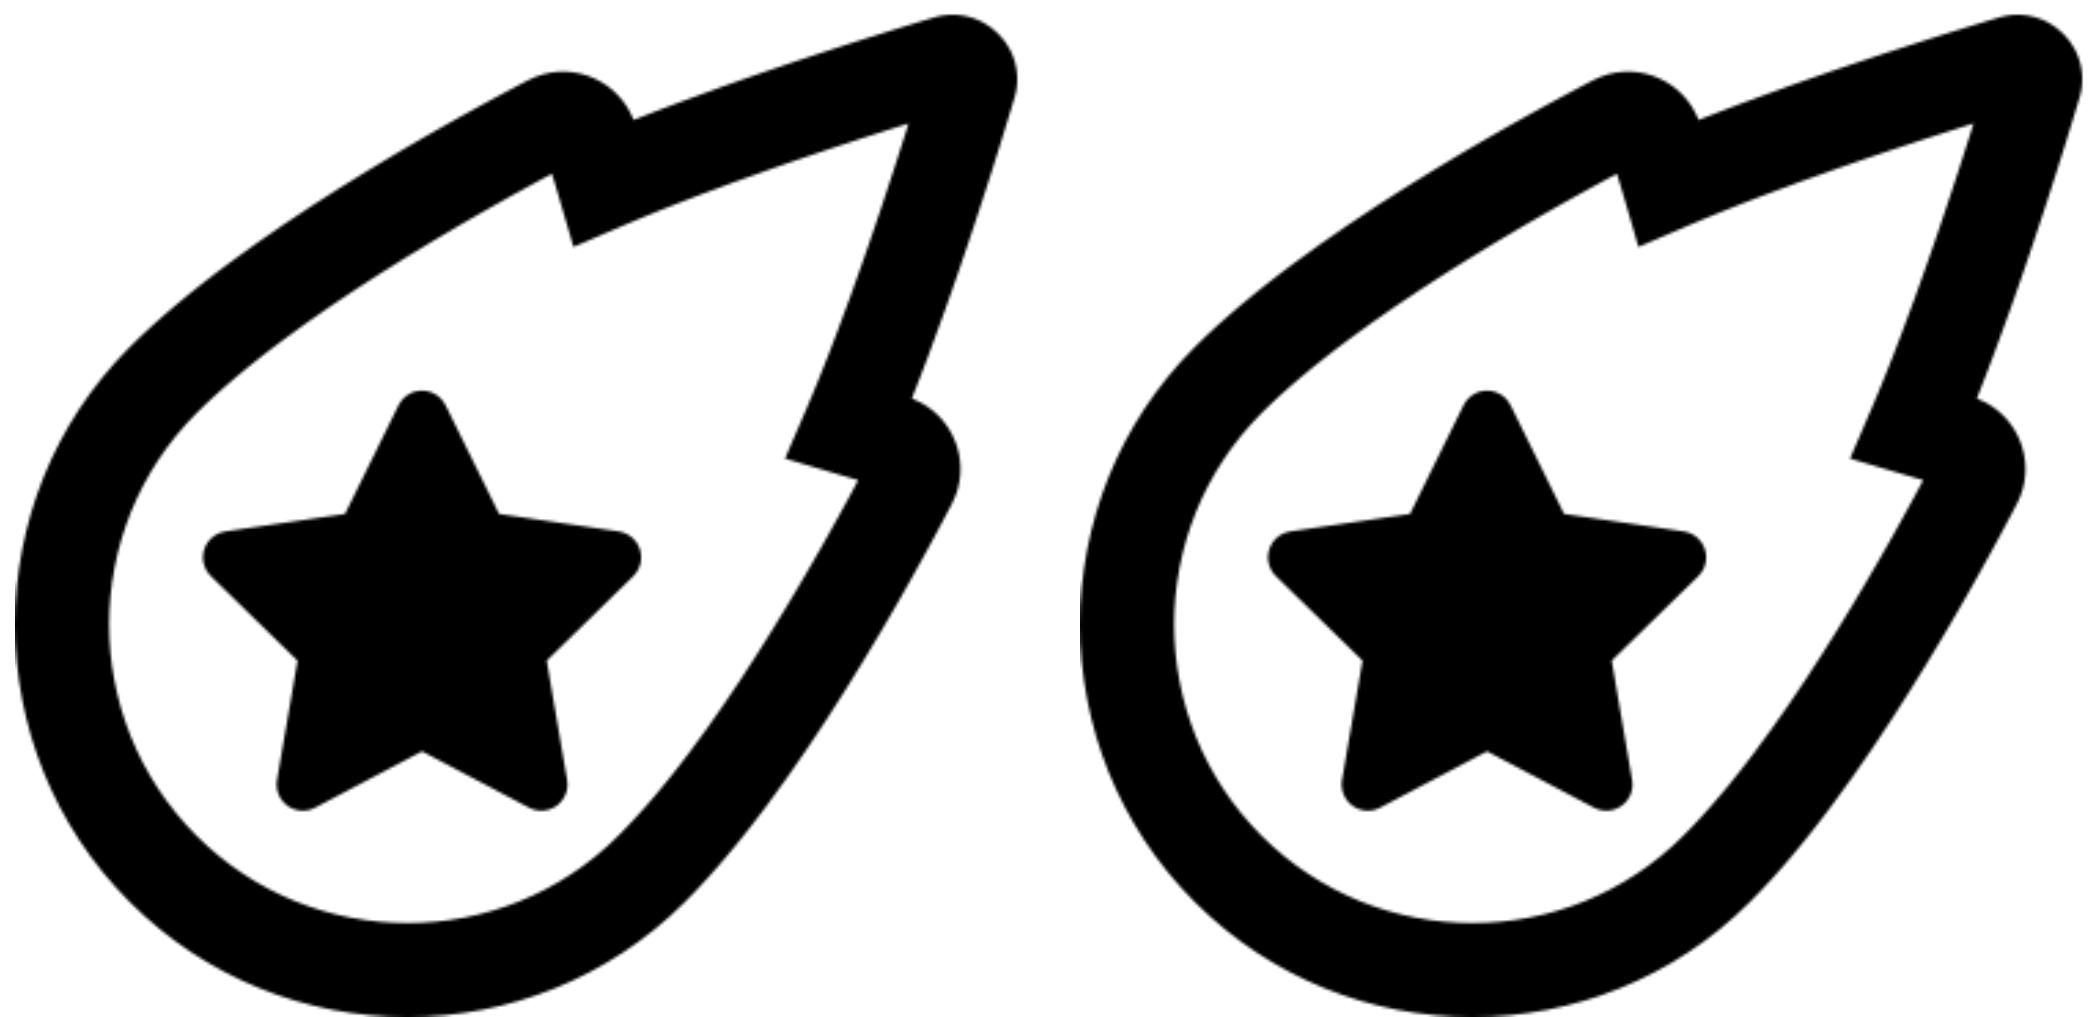








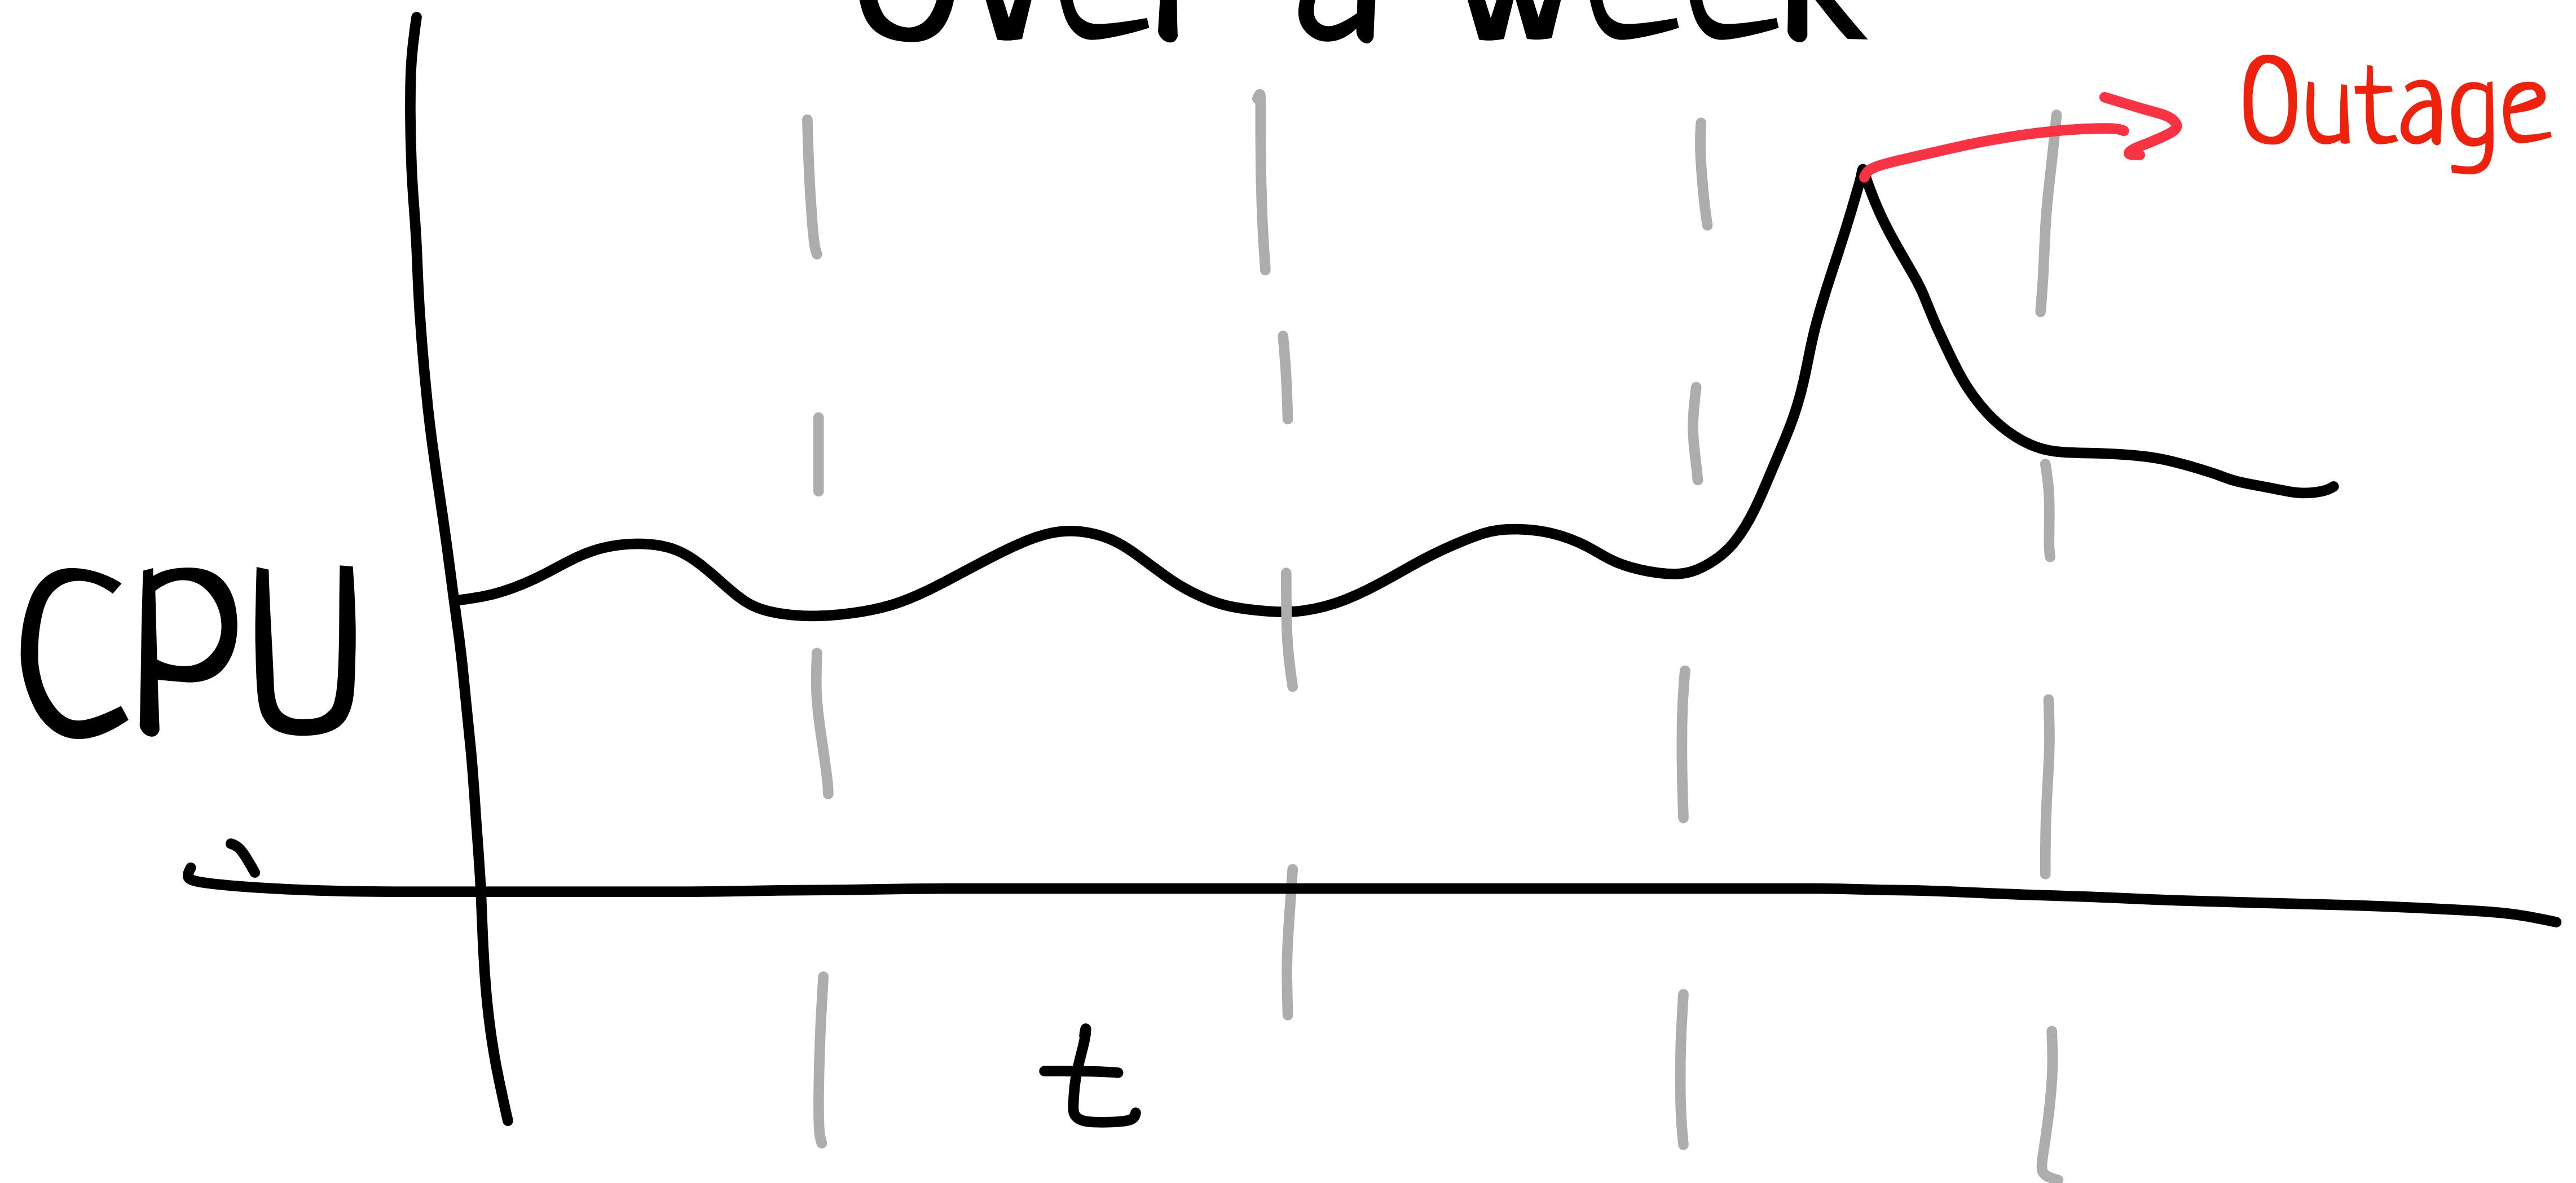
EXPERT

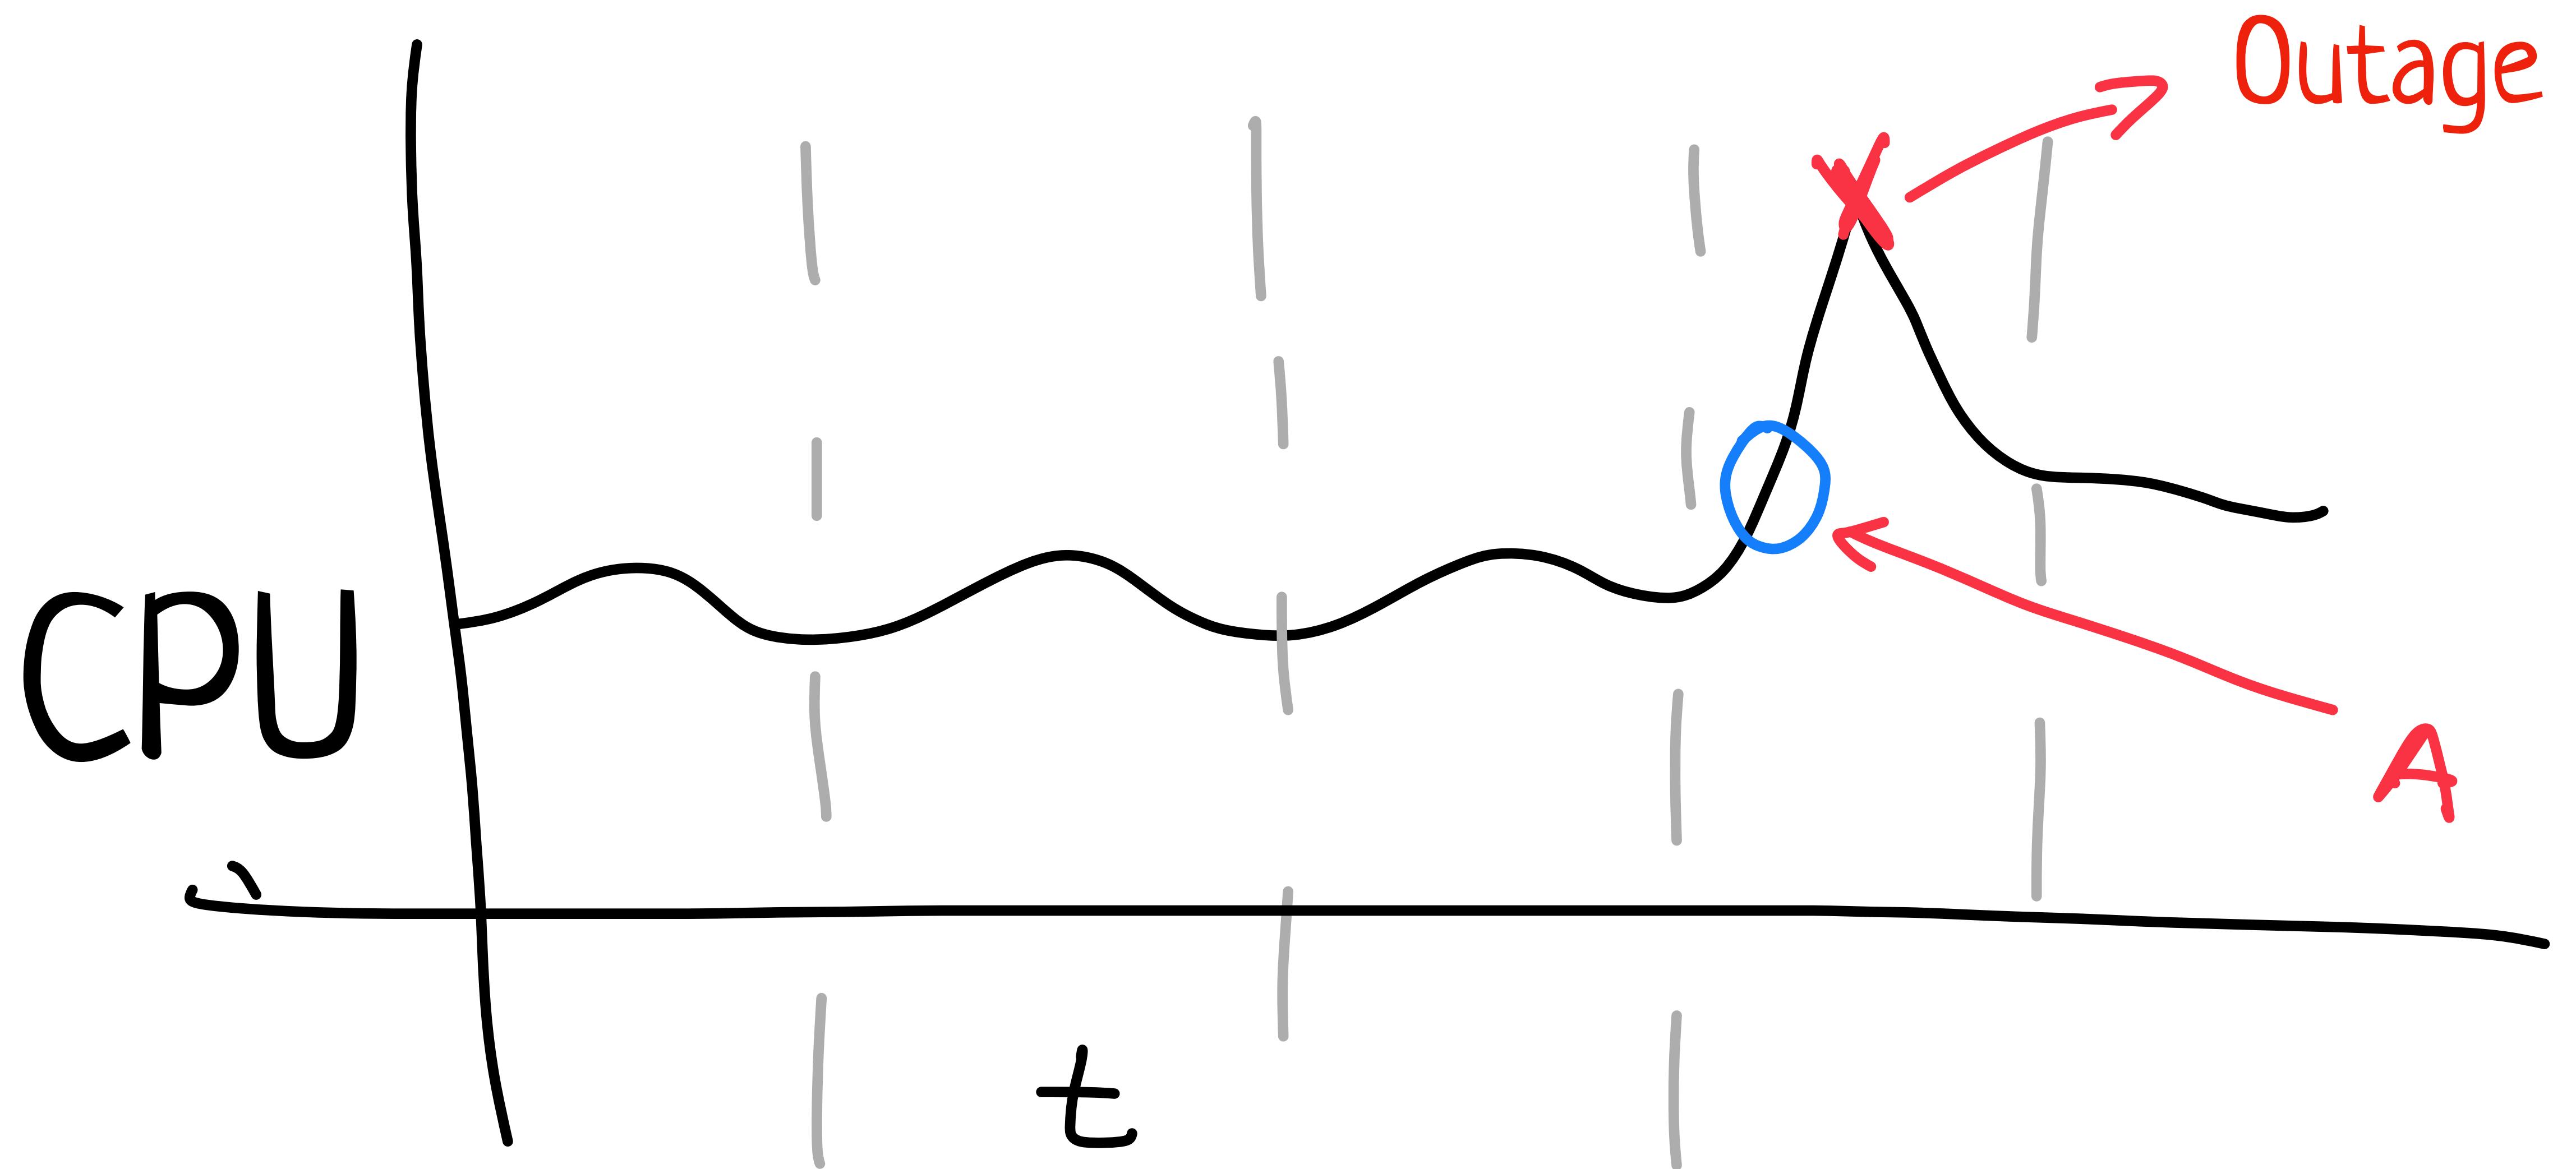


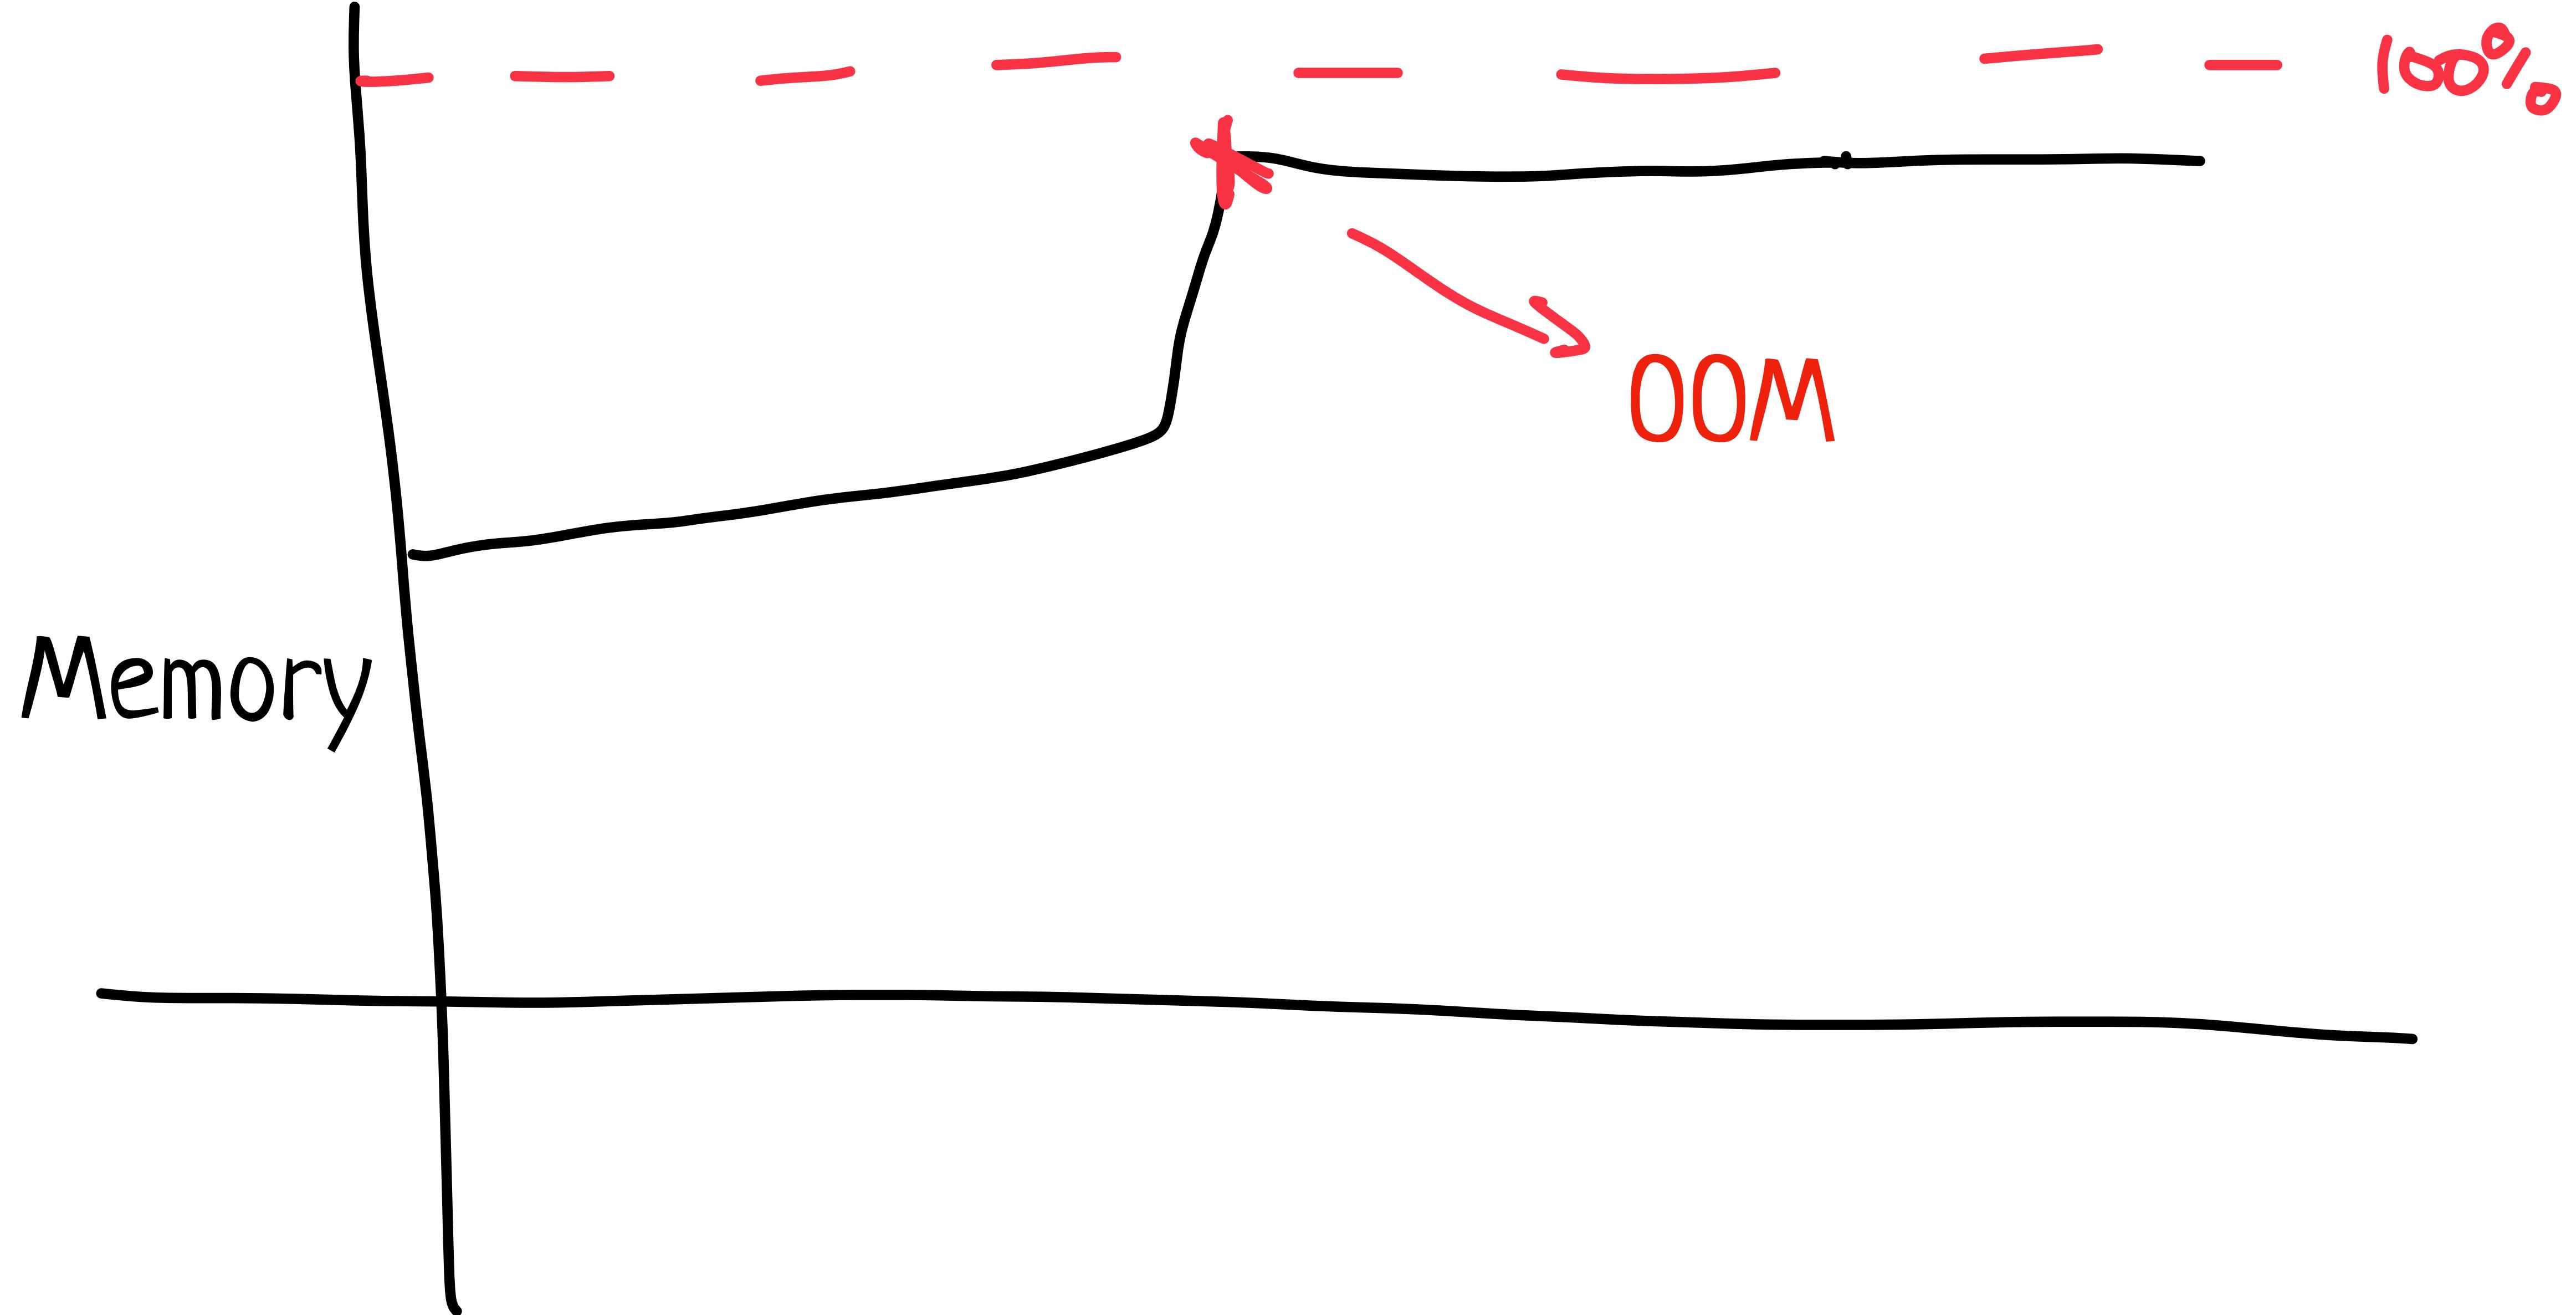
These are hard problems in the field.

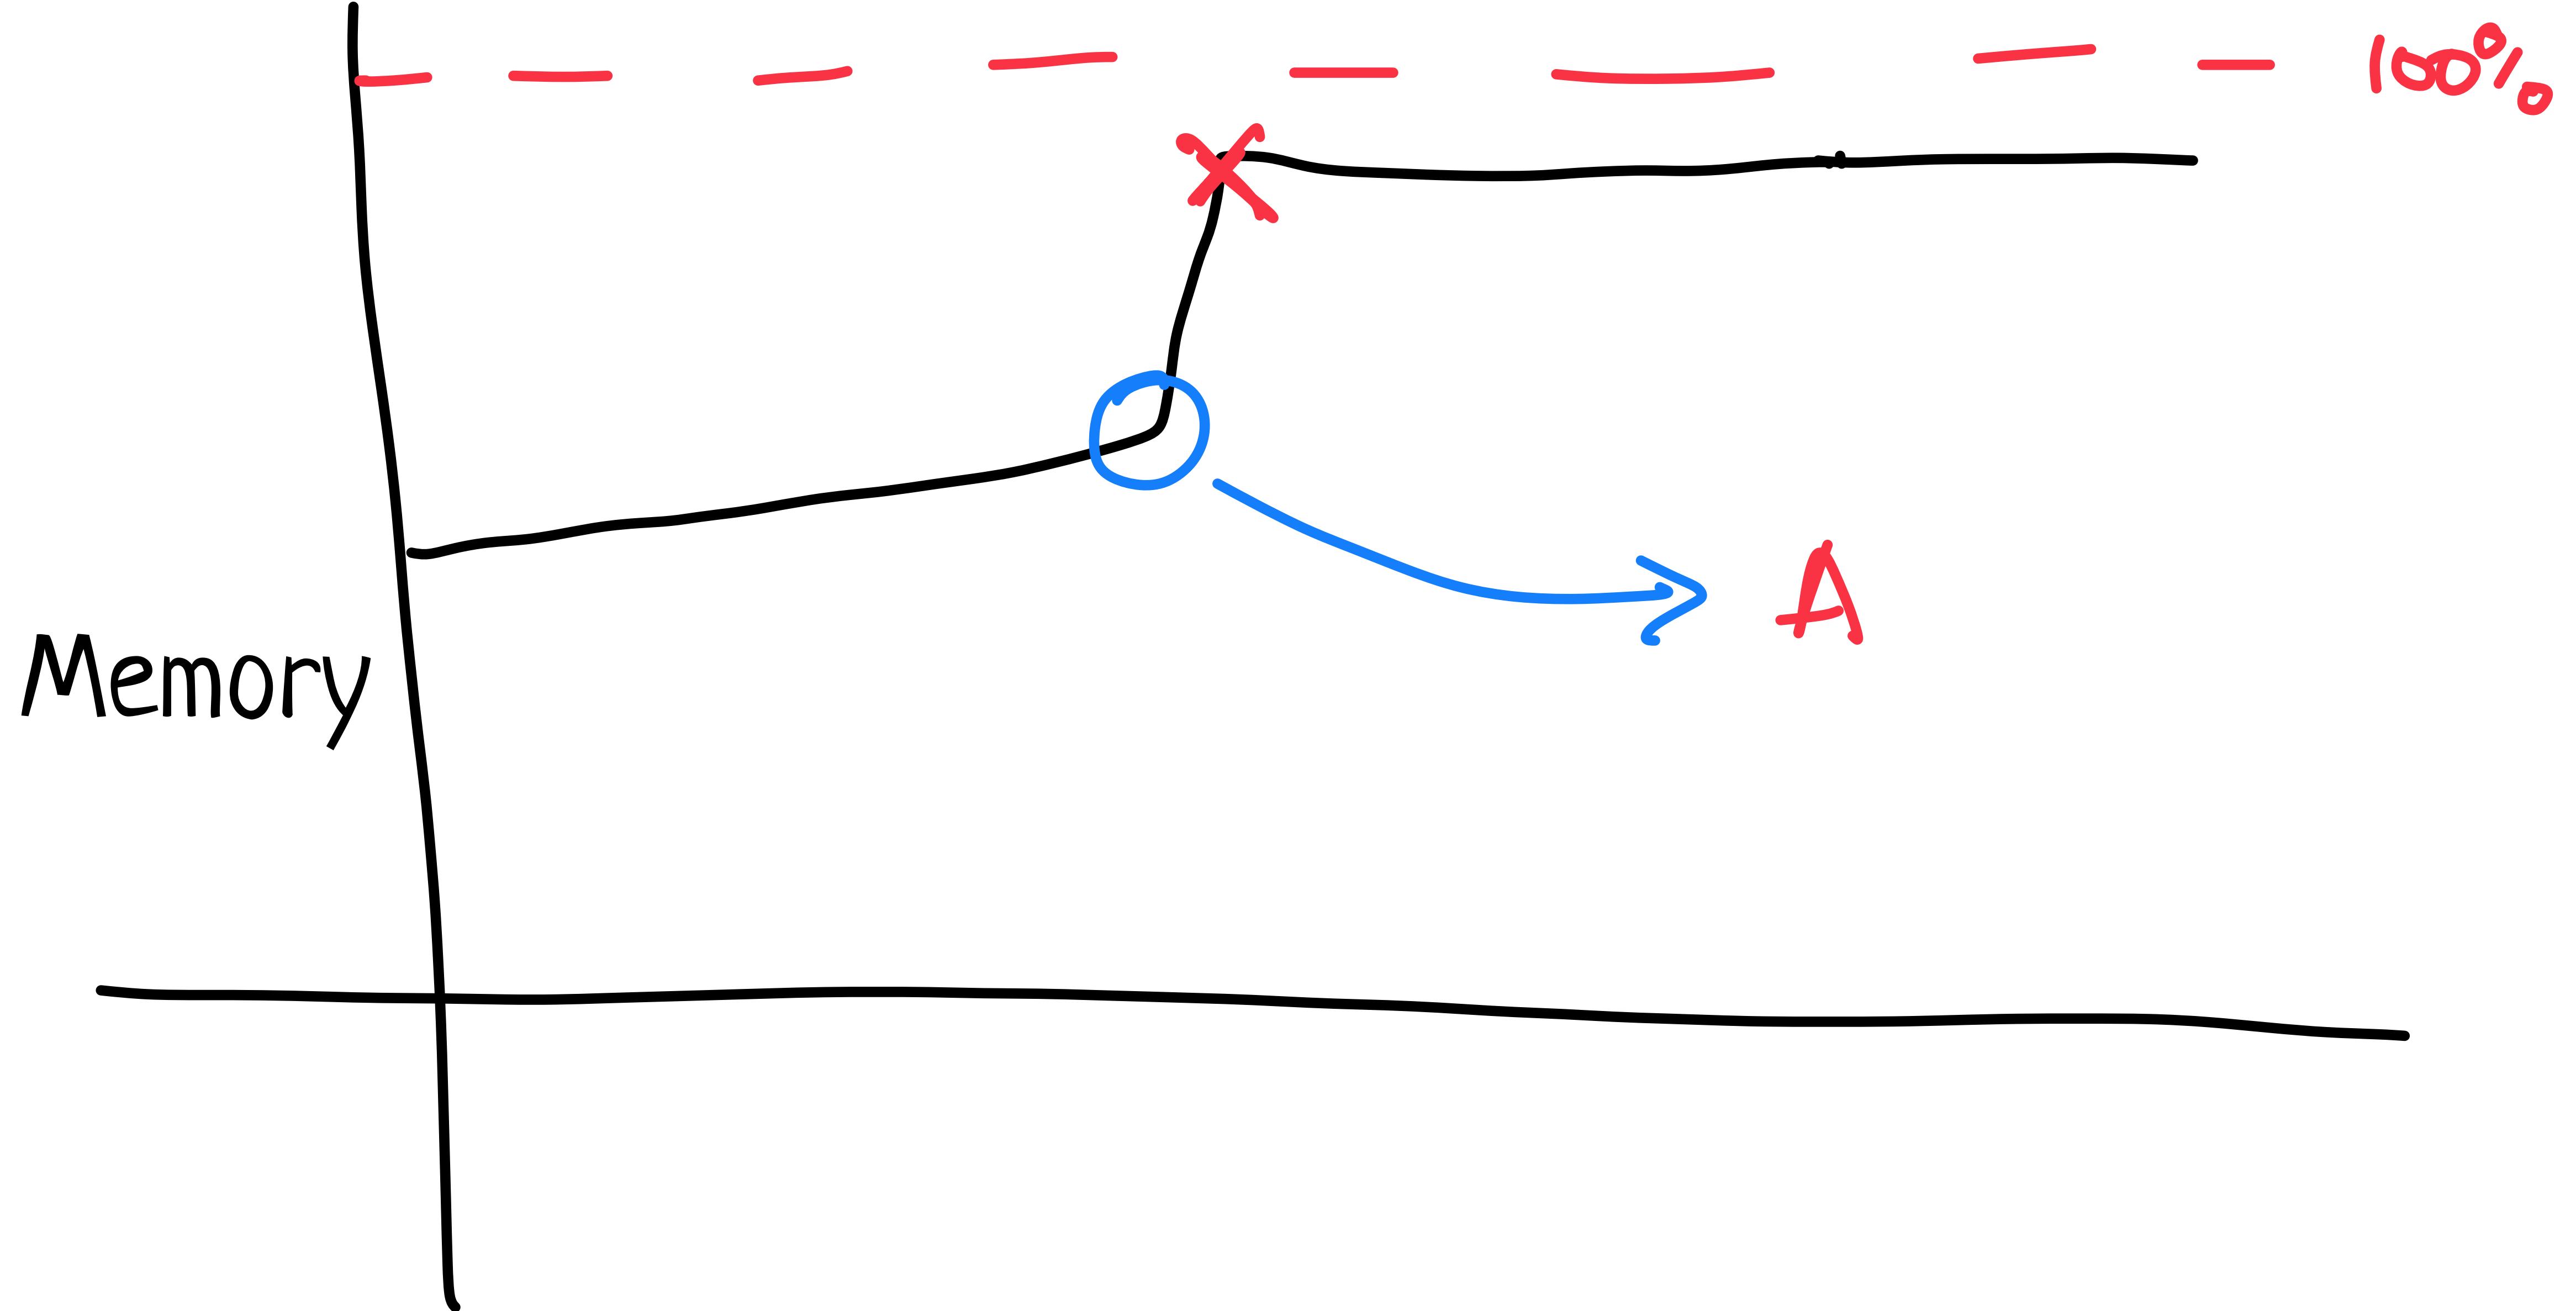
Anomaly detection with automation

Over a week

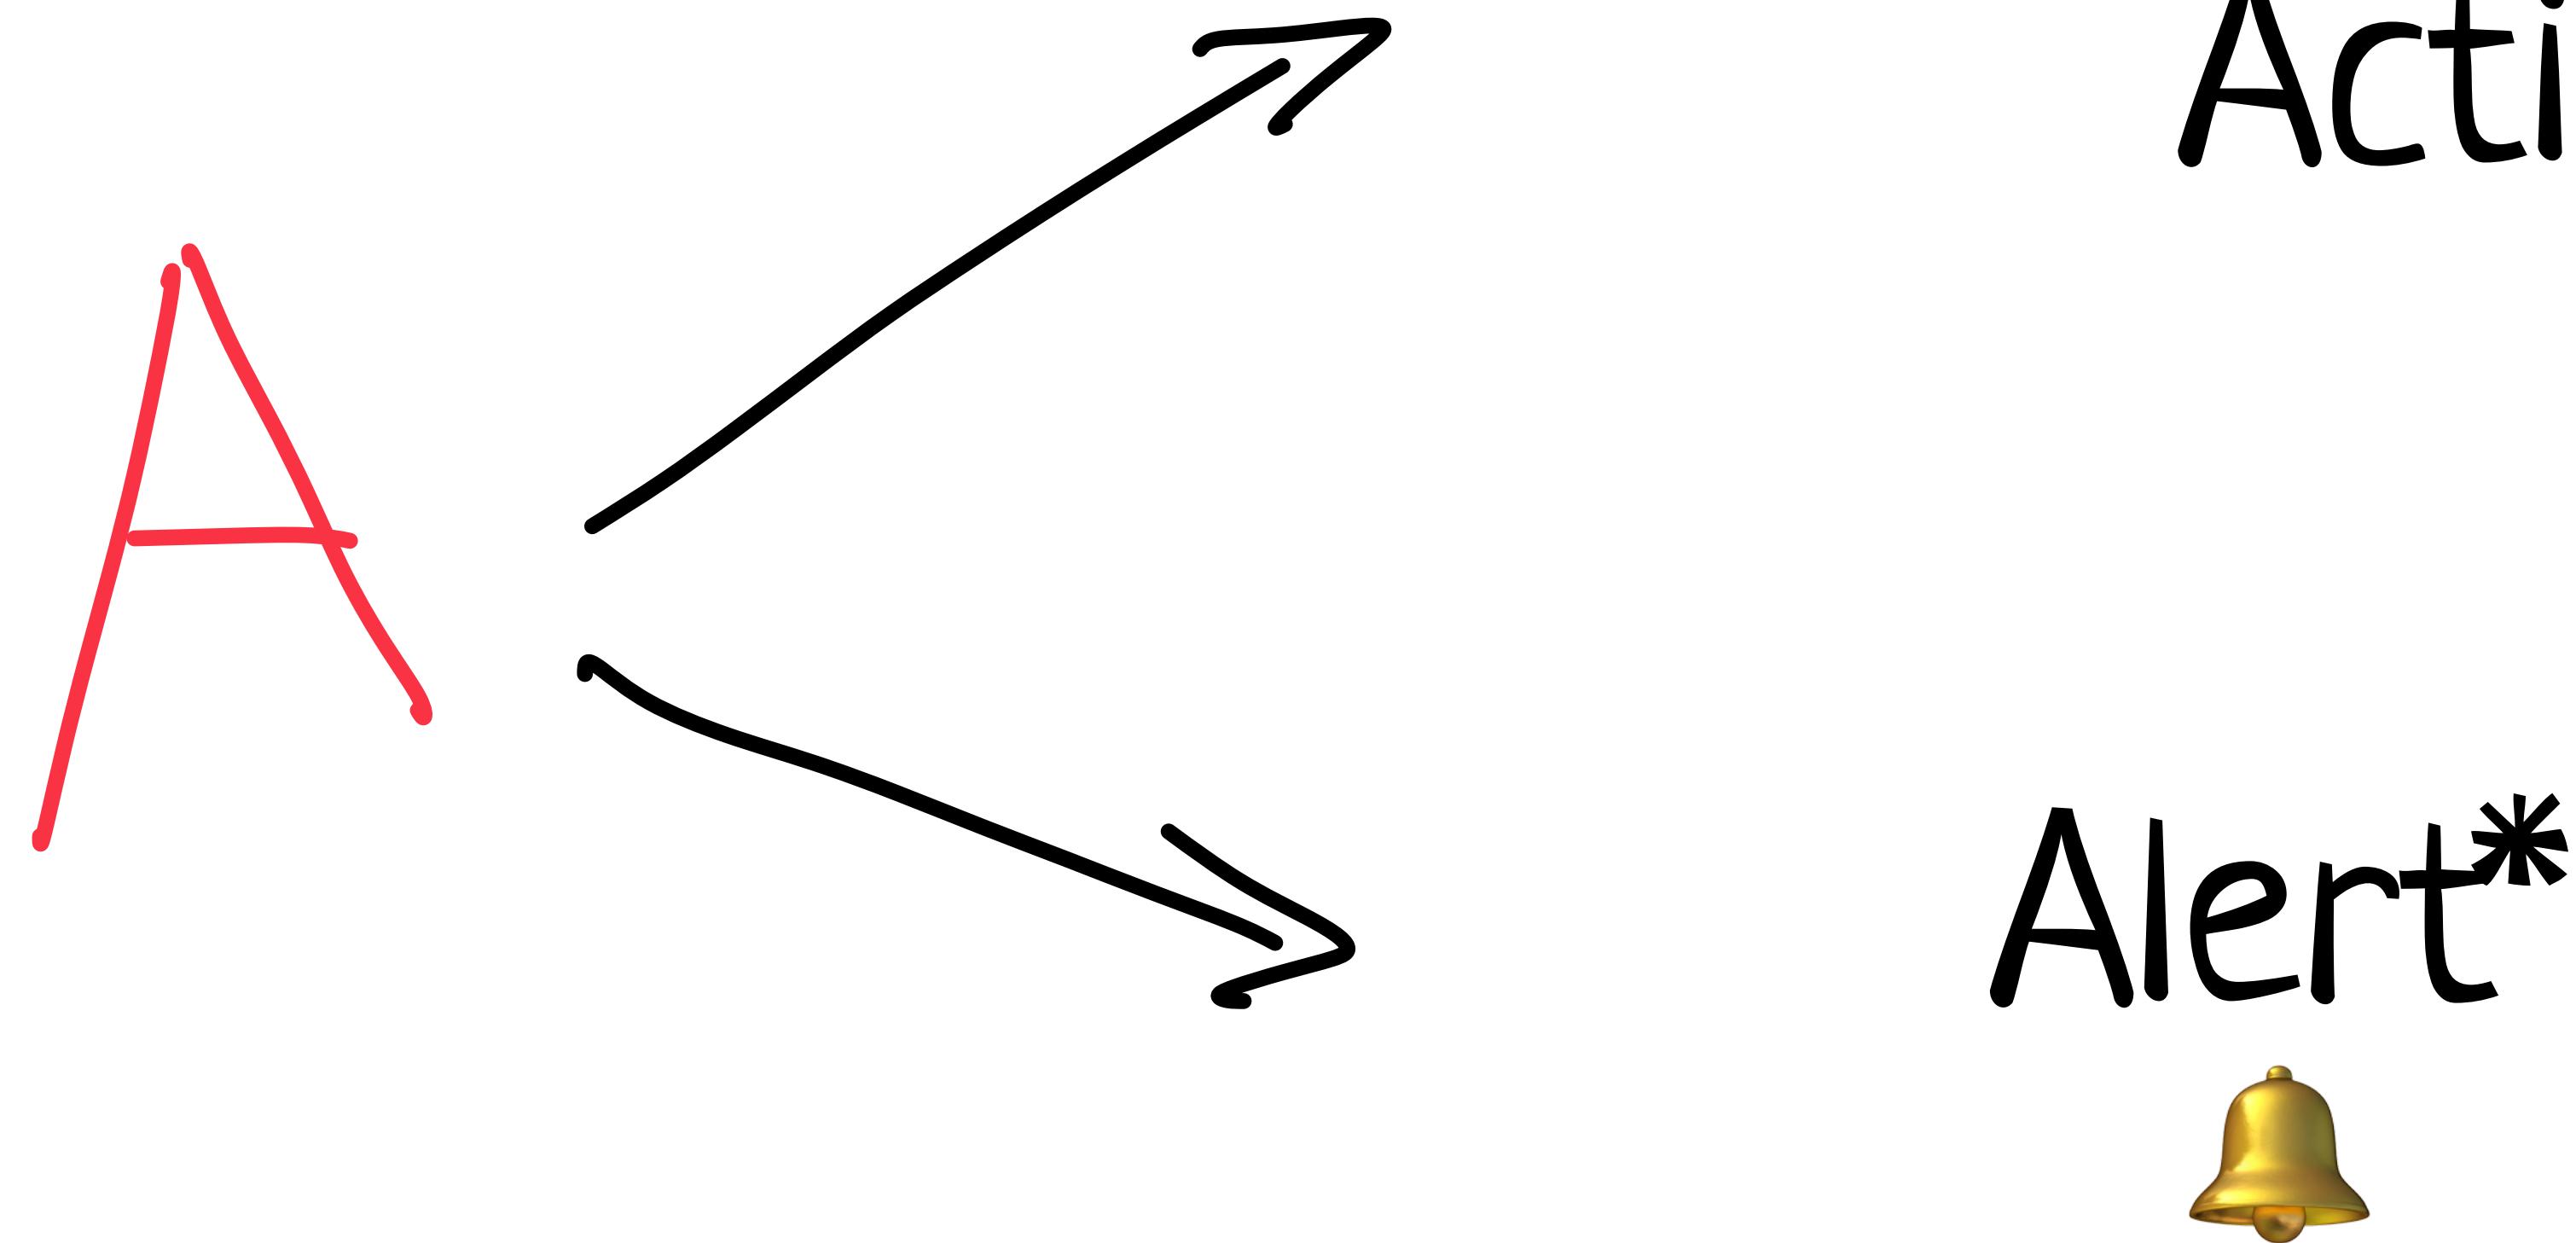






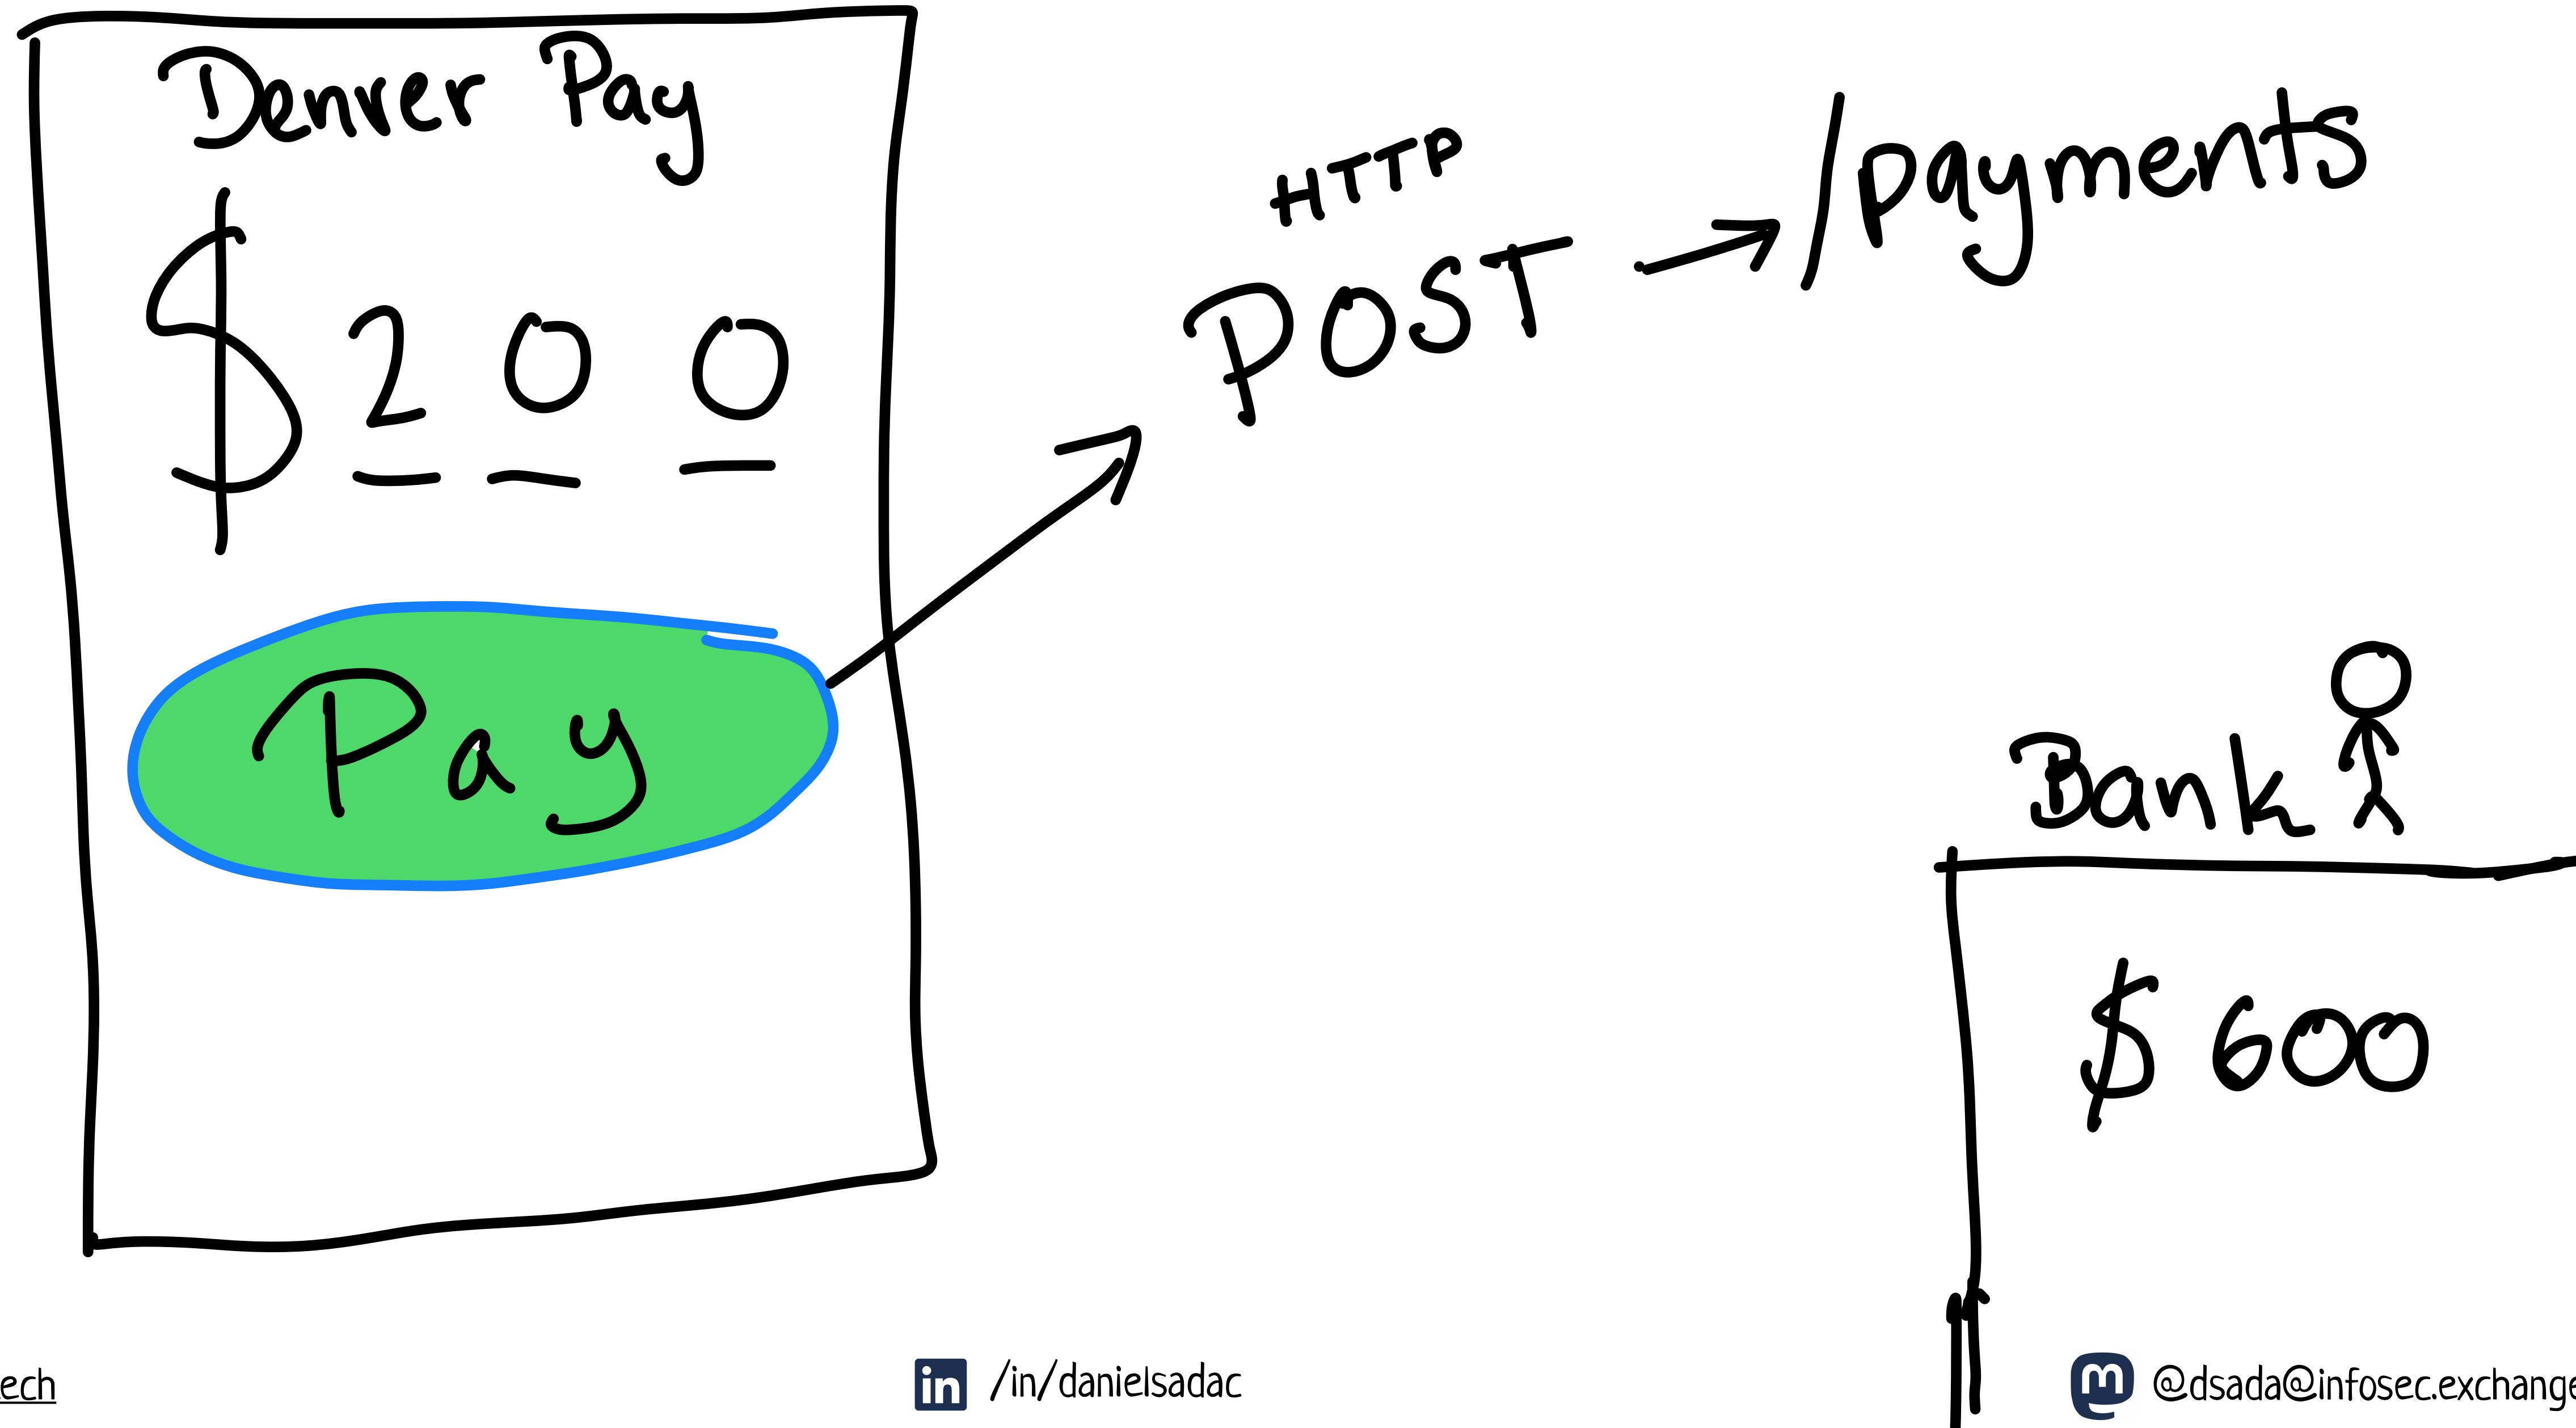


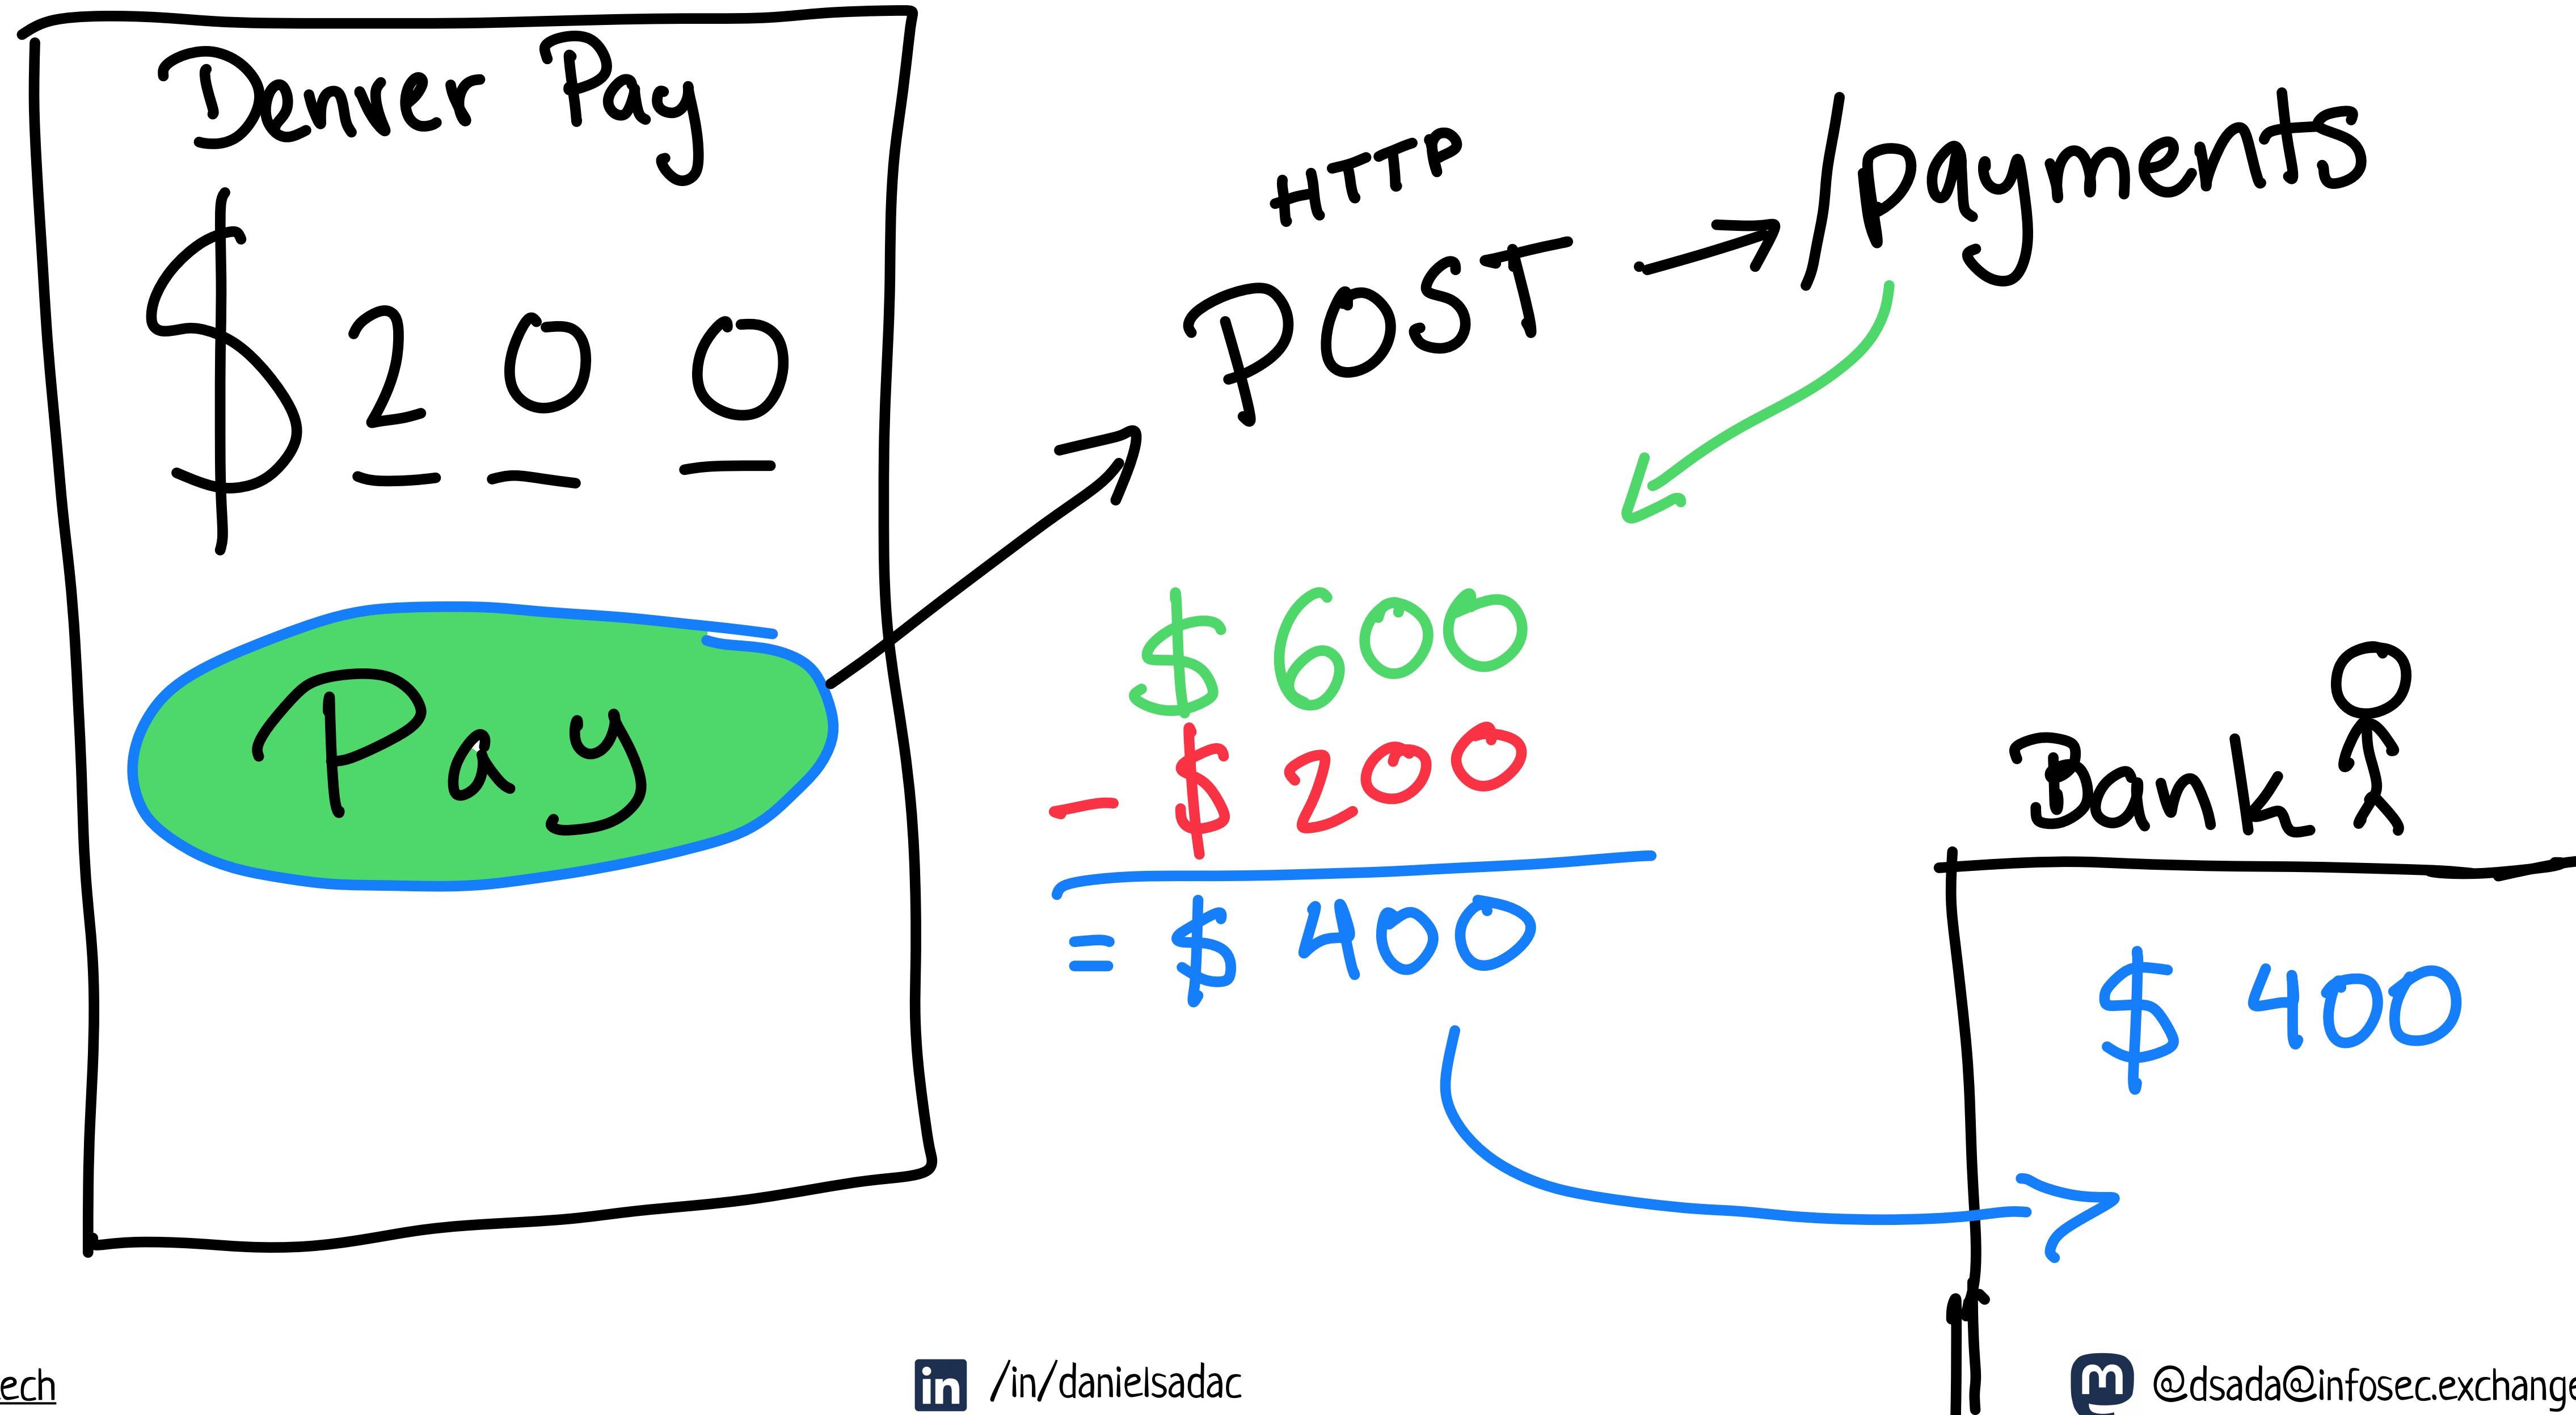
Corrective Action

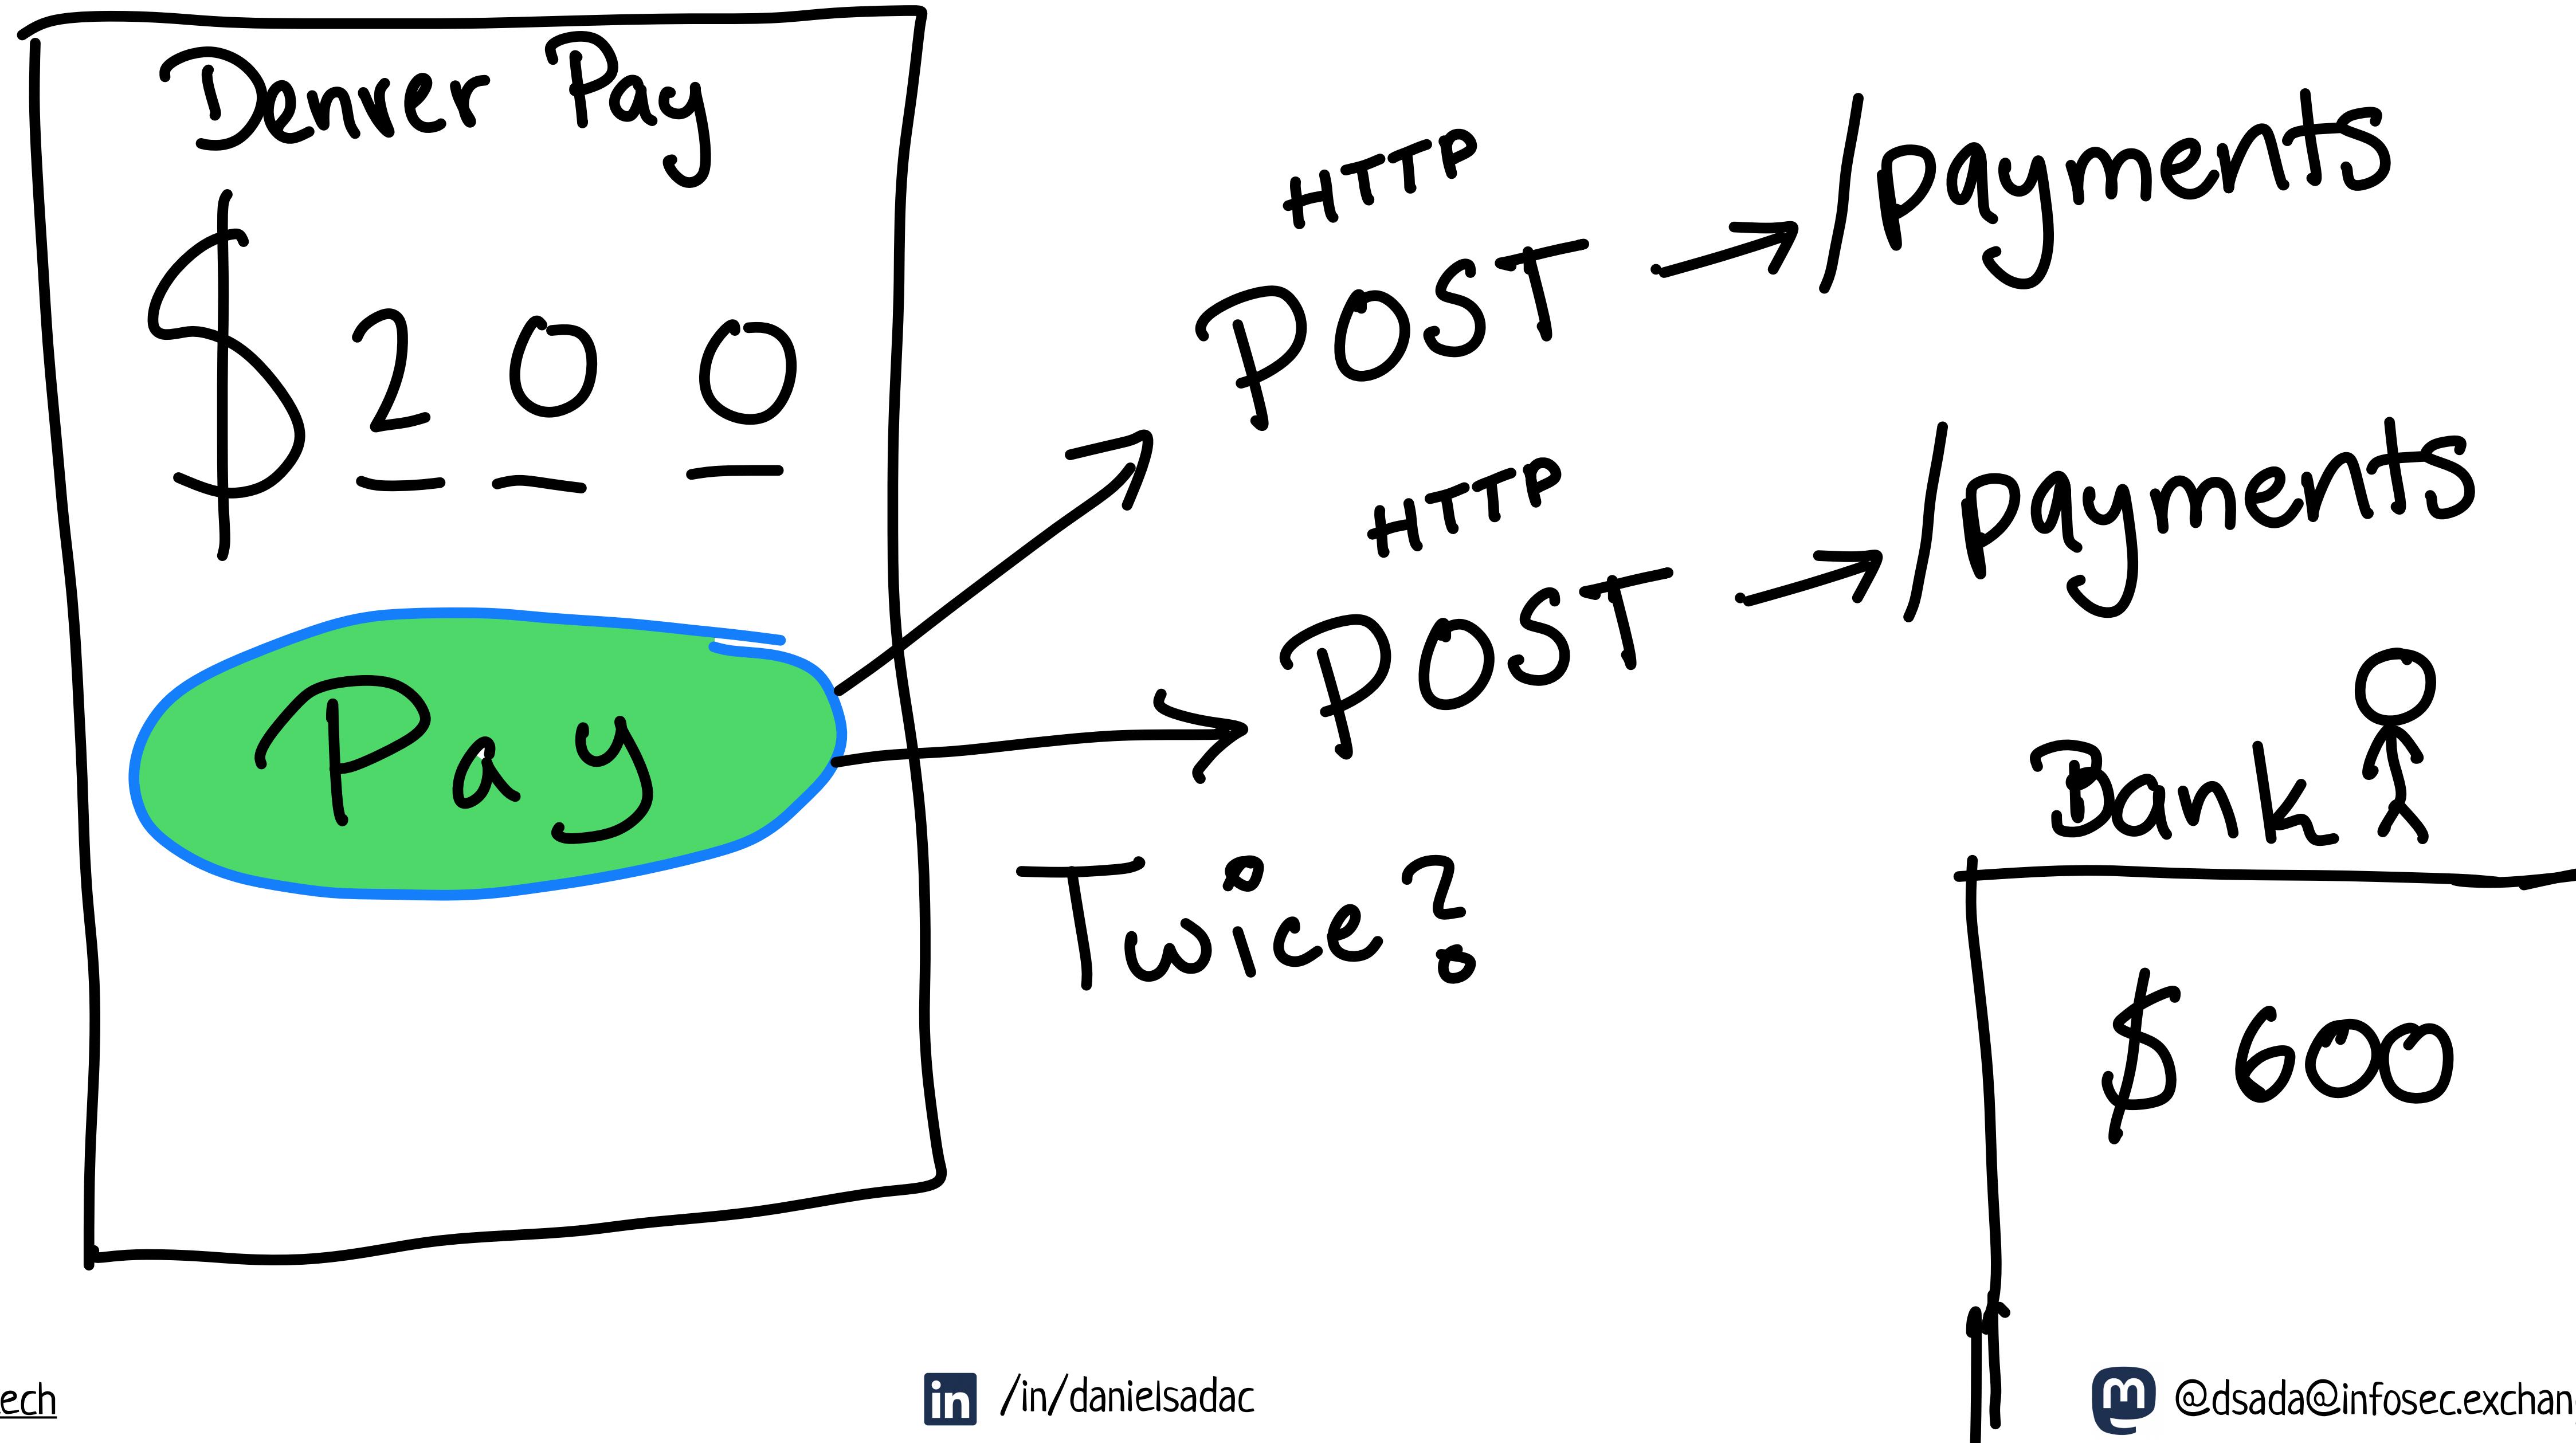


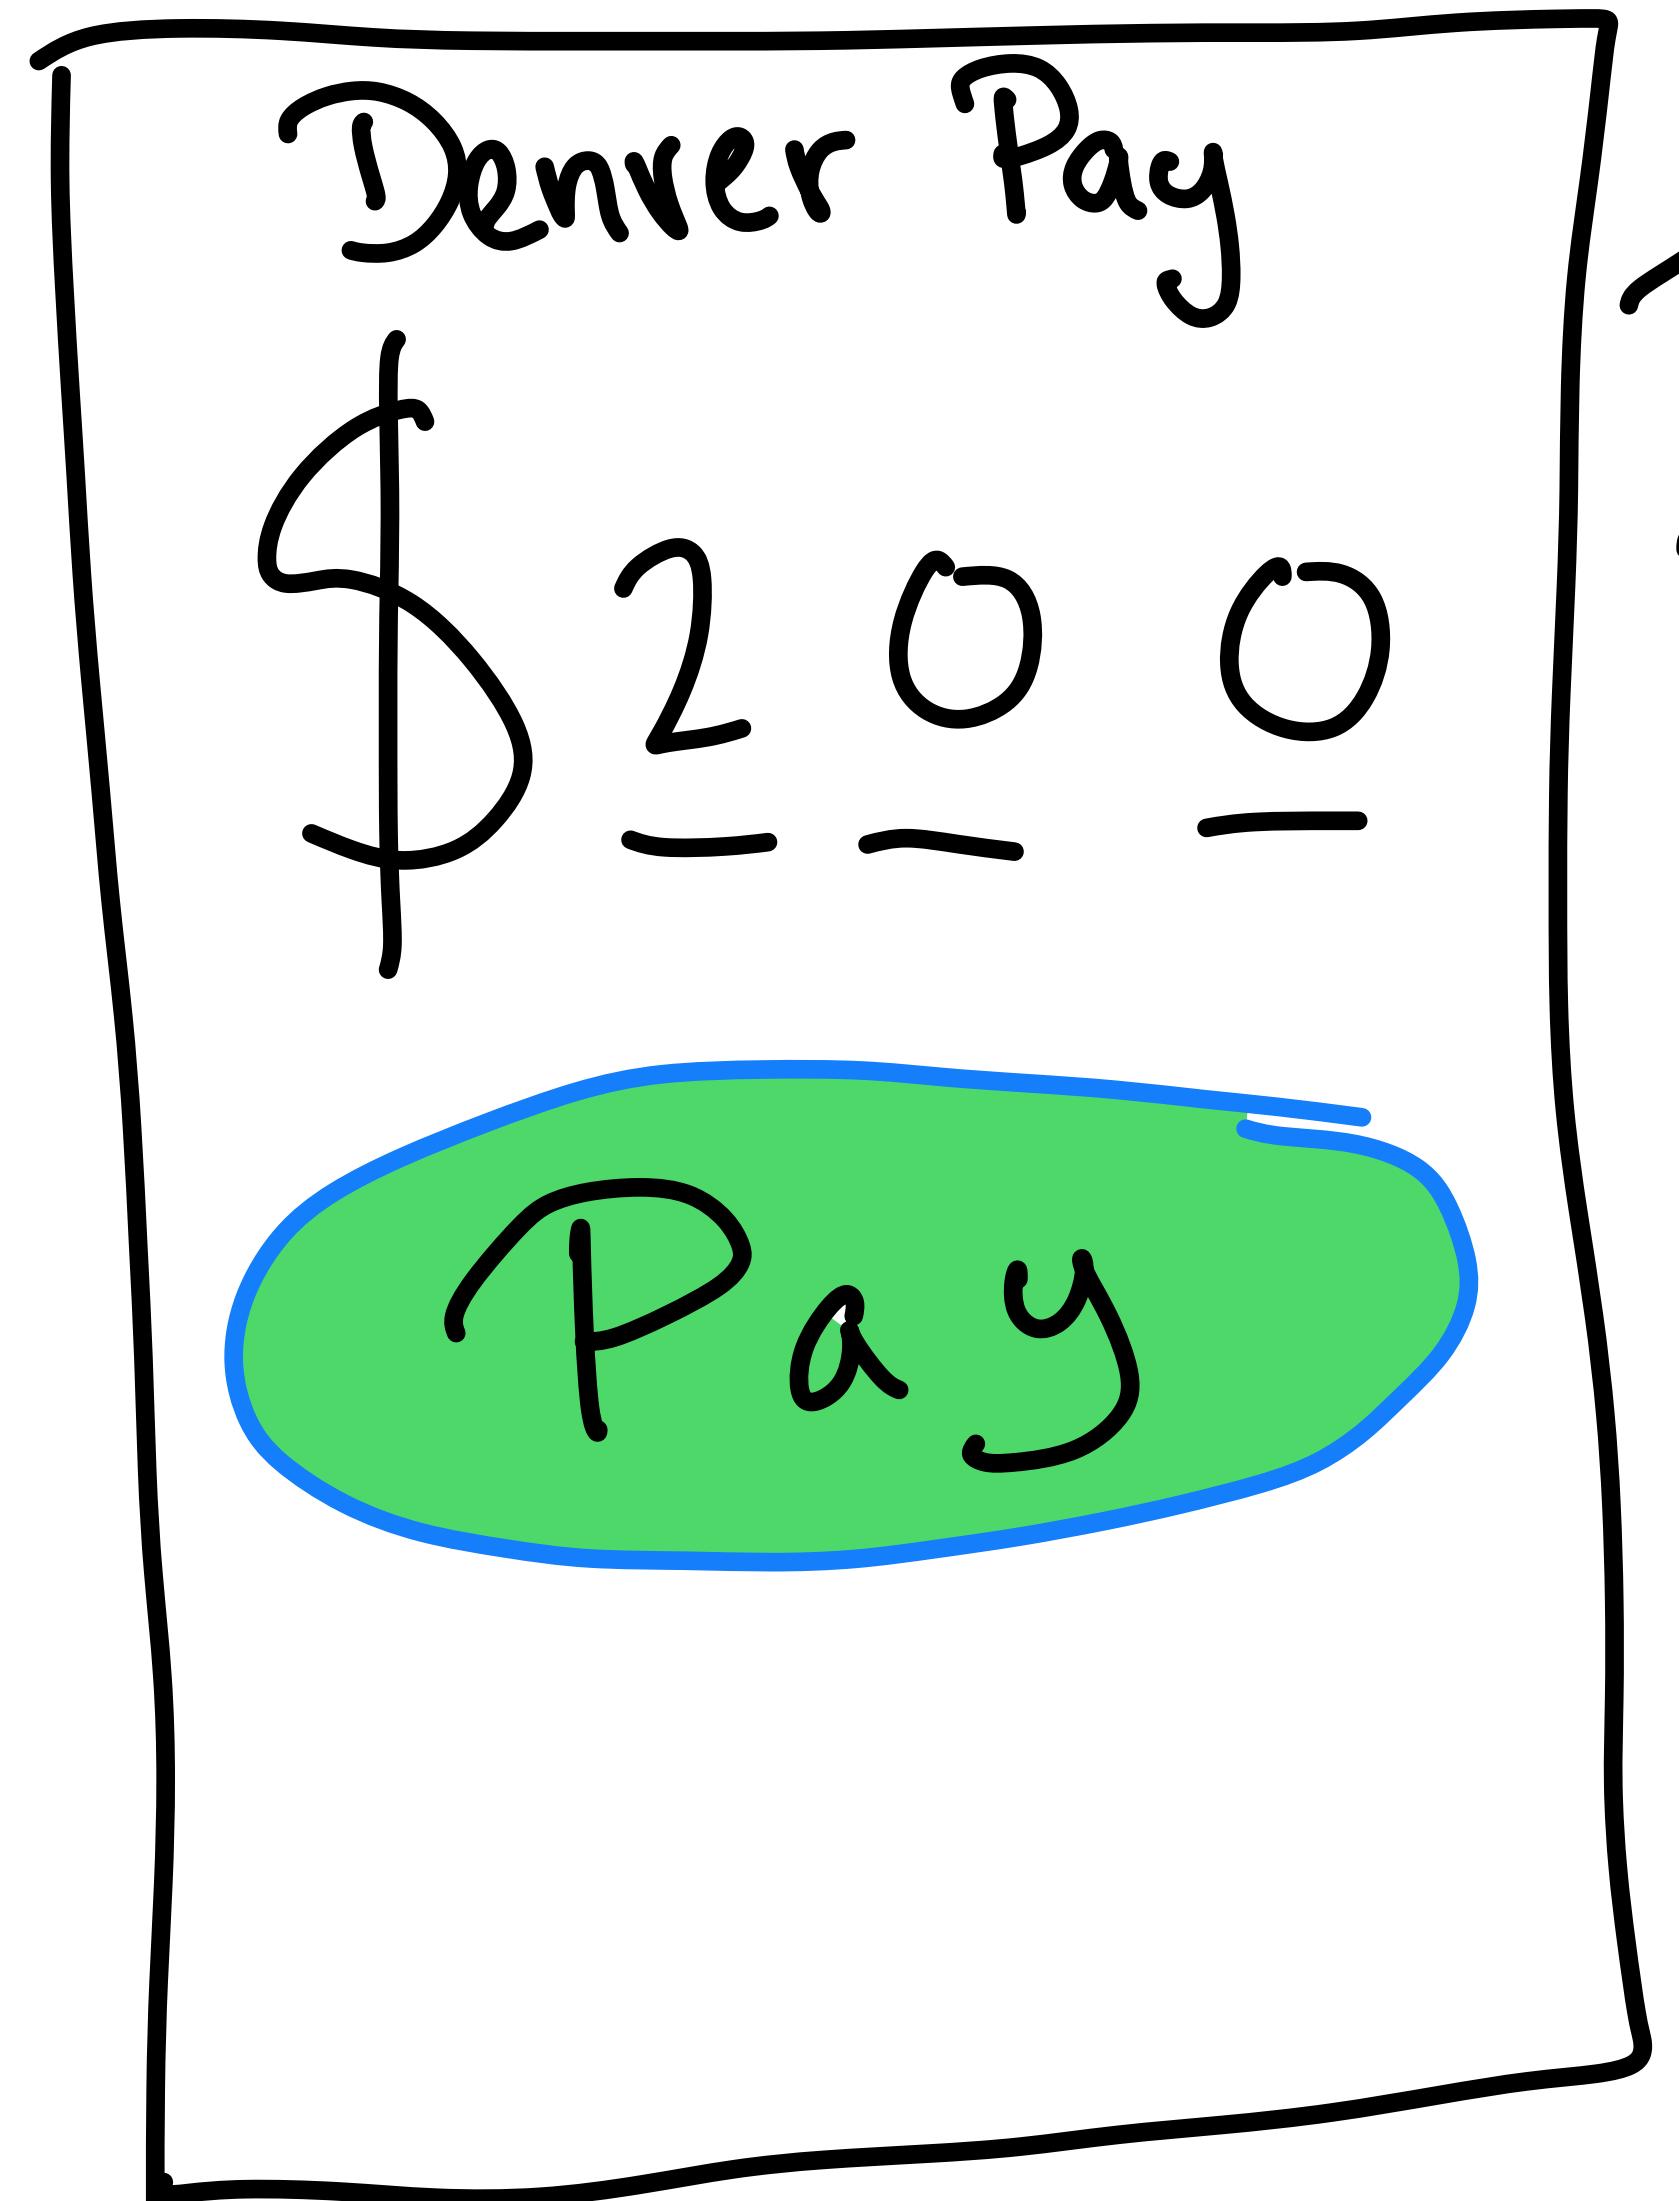
*But alerts don't let you sleep at night :(

Idempotence



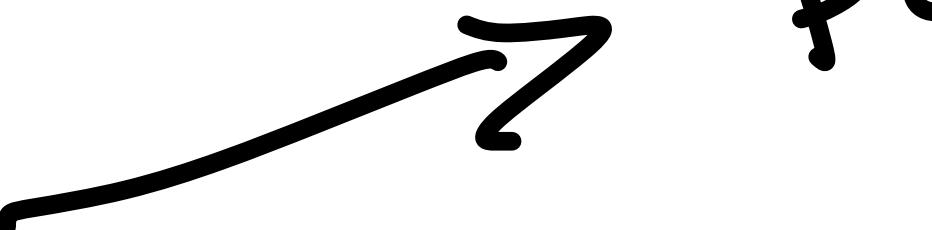






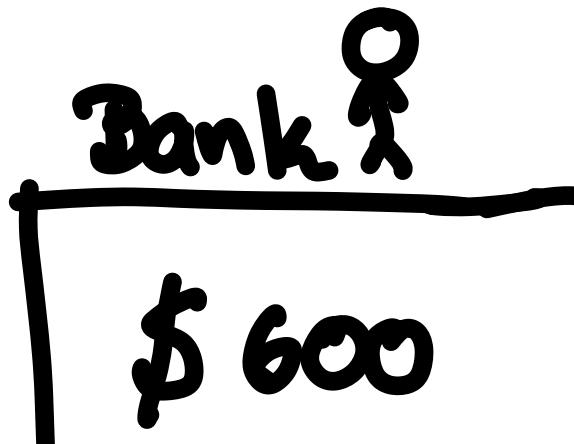
HTTP POST → /payments

HTTP POST → /payments



or

$$\begin{array}{r} \$600 \\ - \$200 \\ \hline = \$400 \end{array}$$



Idempotence

“the property of certain operations in mathematics and computer science that they can be applied multiple times without changing the result beyond the initial application.”

<https://en.wikipedia.org/wiki/Idempotence>

Times ran

Effect

1 UPDATE students SET first_name = 'John' WHERE id = 123;



2 UPDATE students SET first_name = 'John' WHERE id = 123;



3 UPDATE students SET first_name = 'John' WHERE id = 123;



4 UPDATE students SET first_name = 'John' WHERE id = 123;



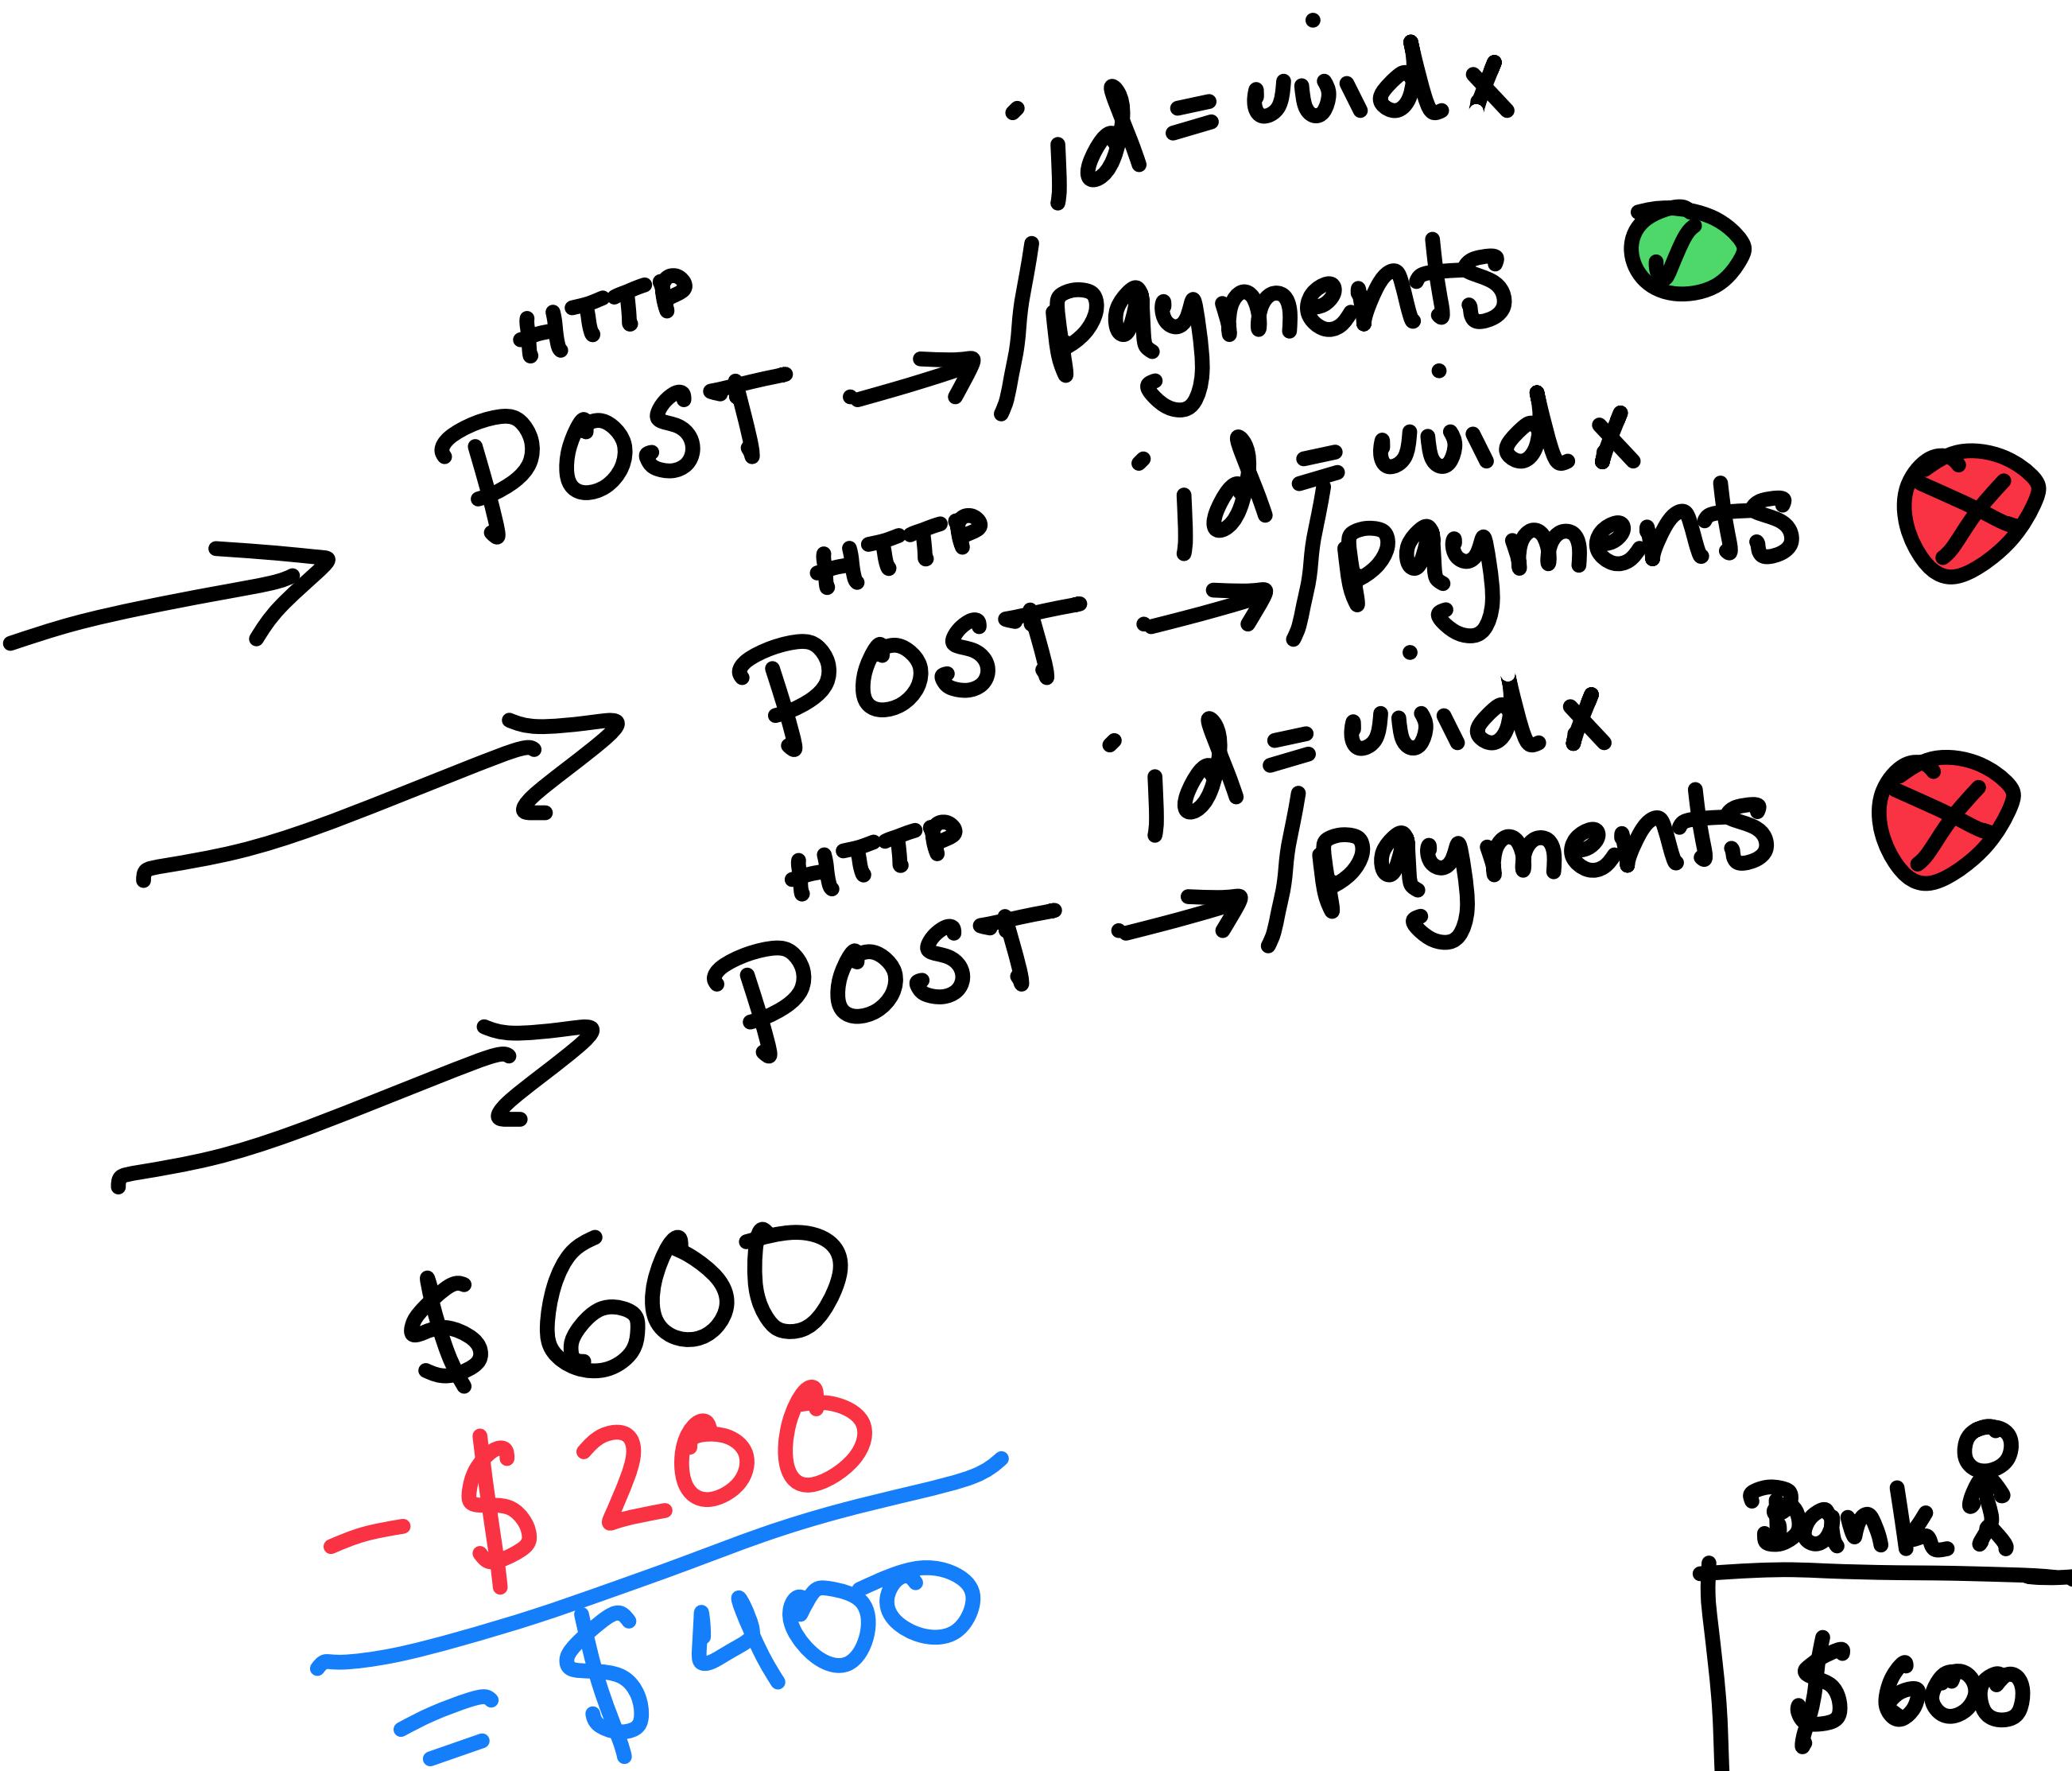
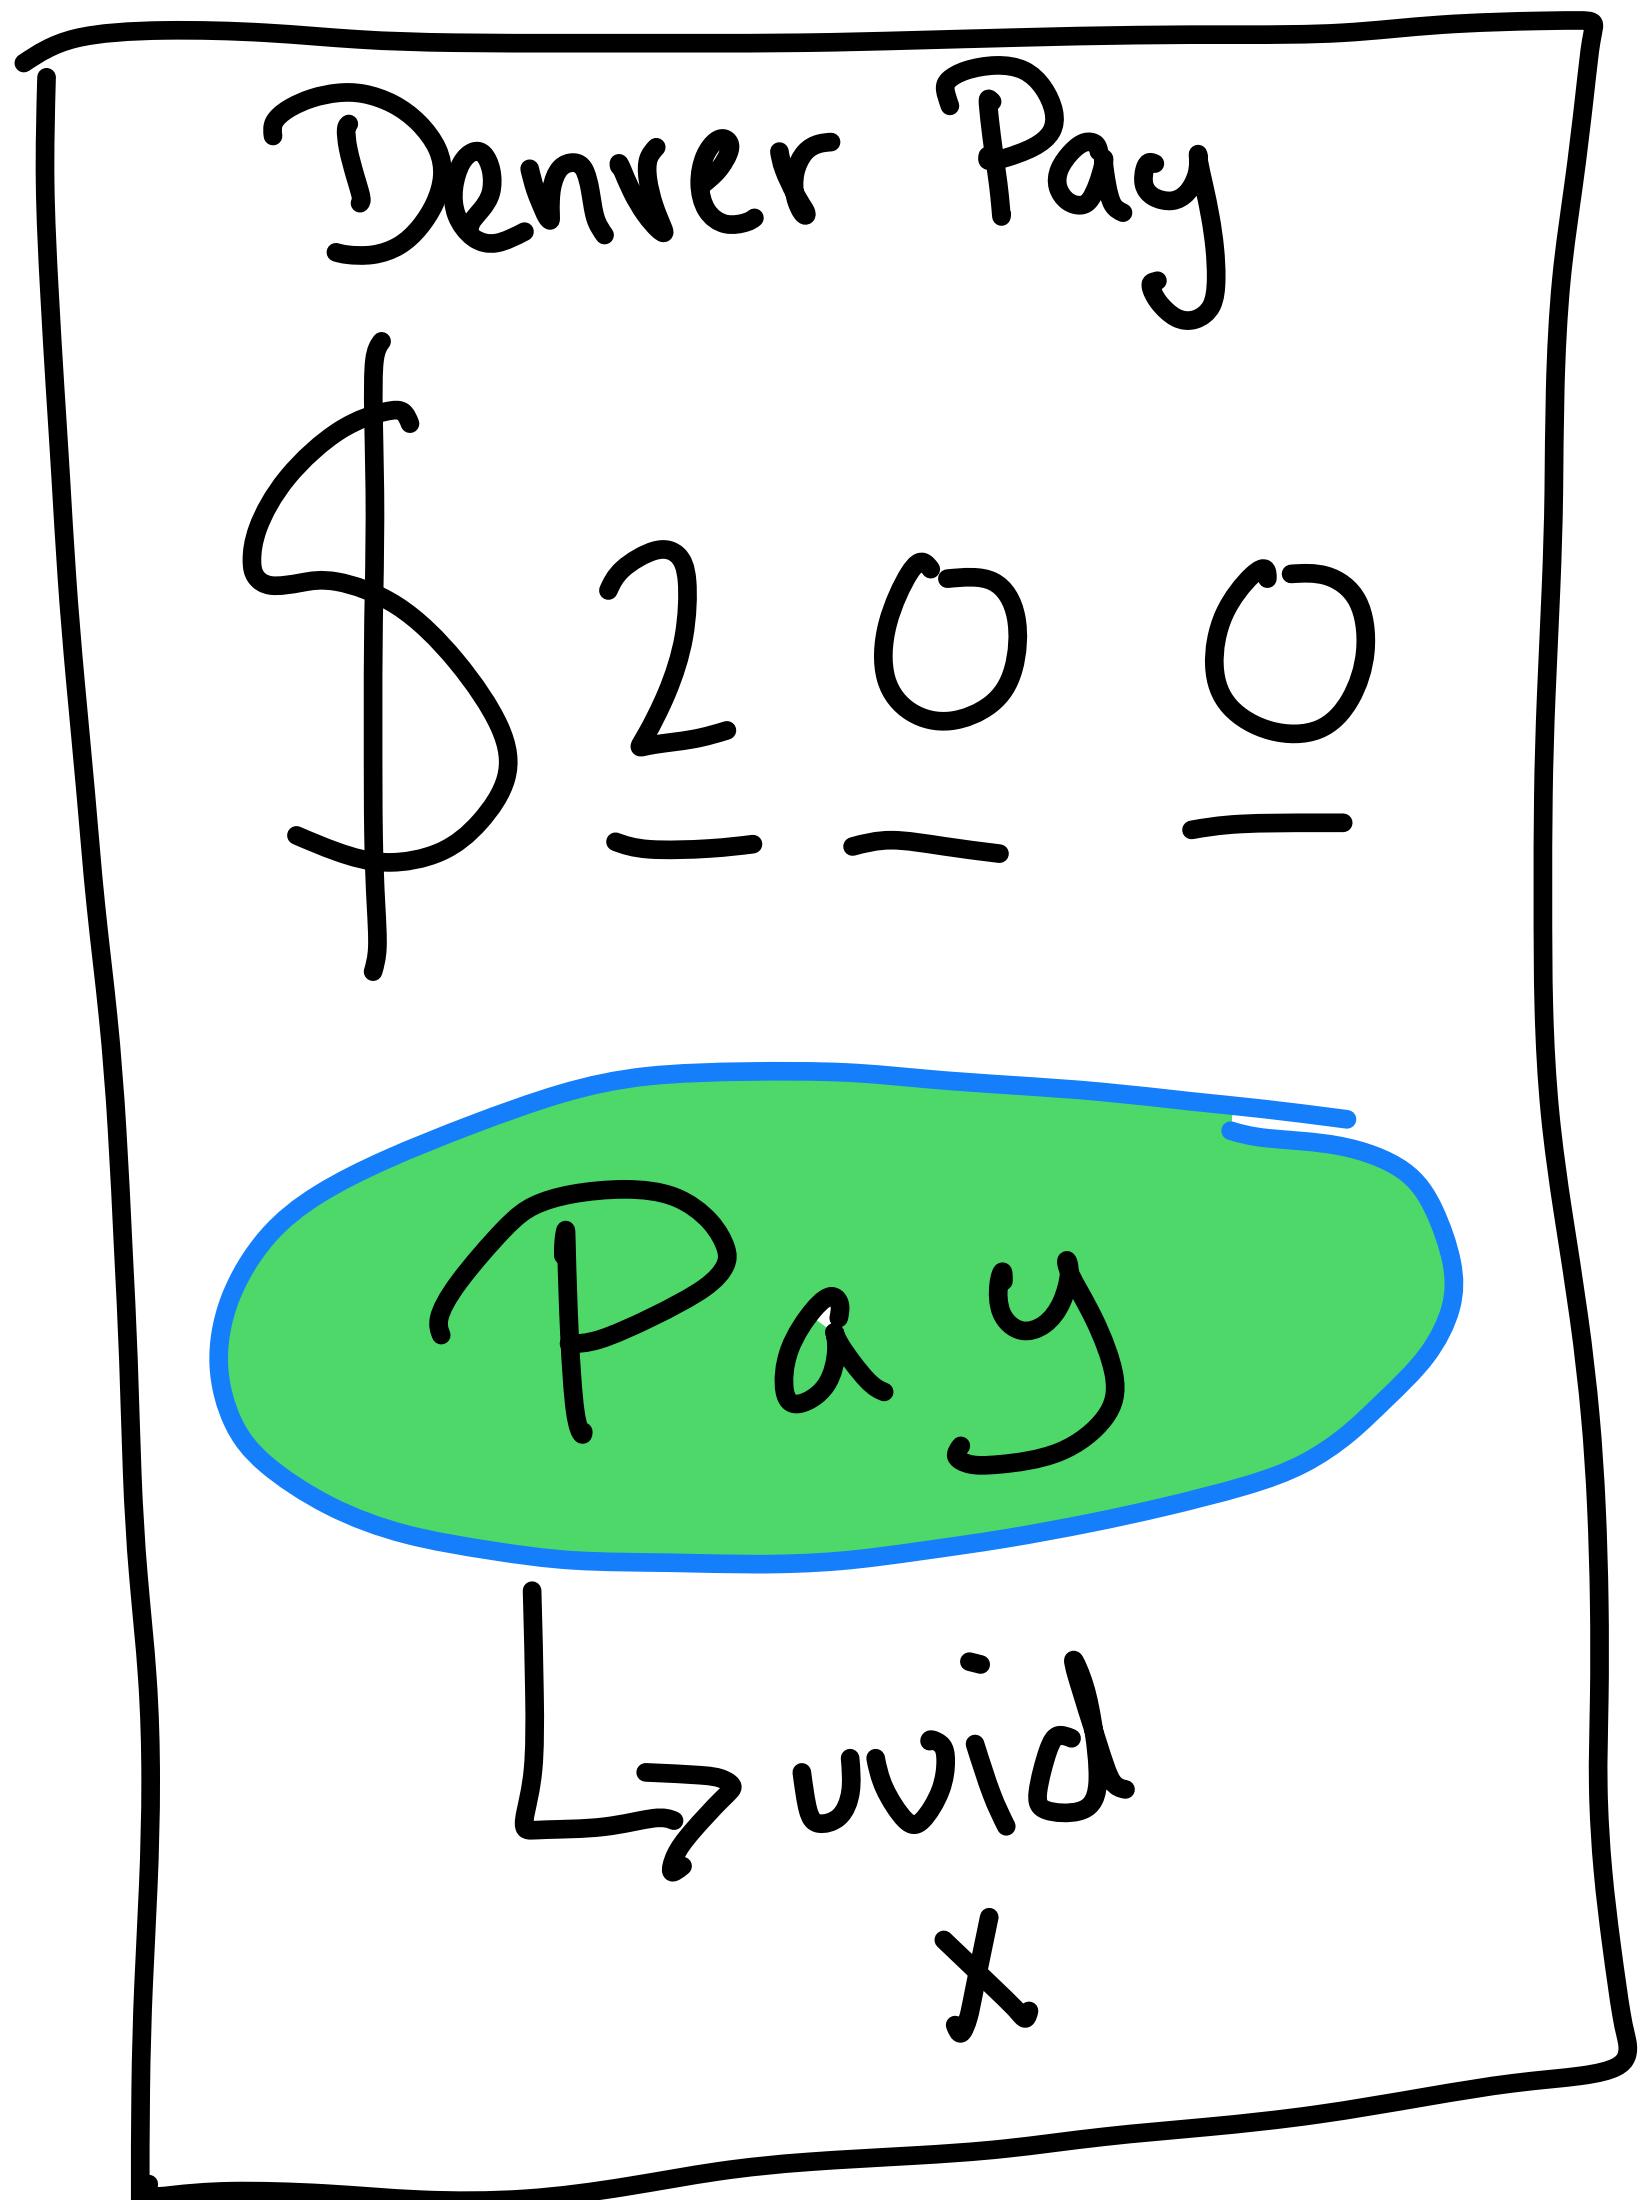
5 UPDATE students SET first_name = 'John' WHERE id = 123;



...

<https://ieftimov.com/posts/understand-how-why-add-idempotent-requests-api/>

One of many solutions: Idempotency Keys



We could still talk about:

Fault tolerant applications

Chaos engineering

How orchestration platforms solve some of
this problems

Observability.

But that is left as an exercise to the viewer.

And with all this architecture tips...

I wish you a good night's sleep...
(Maybe?)

DANIEL
SAIDA

Developer Productivity & Culture

Thank you!

